

Computer Architecture

Lecture 13a: Memory Controllers: Service Quality

Ataberk Olgun

Prof. Onur Mutlu

ETH Zürich

Fall 2023

9 November 2023

Recall: Memory Controllers and Shared Resources

DRAM vs. Other Types of Memories

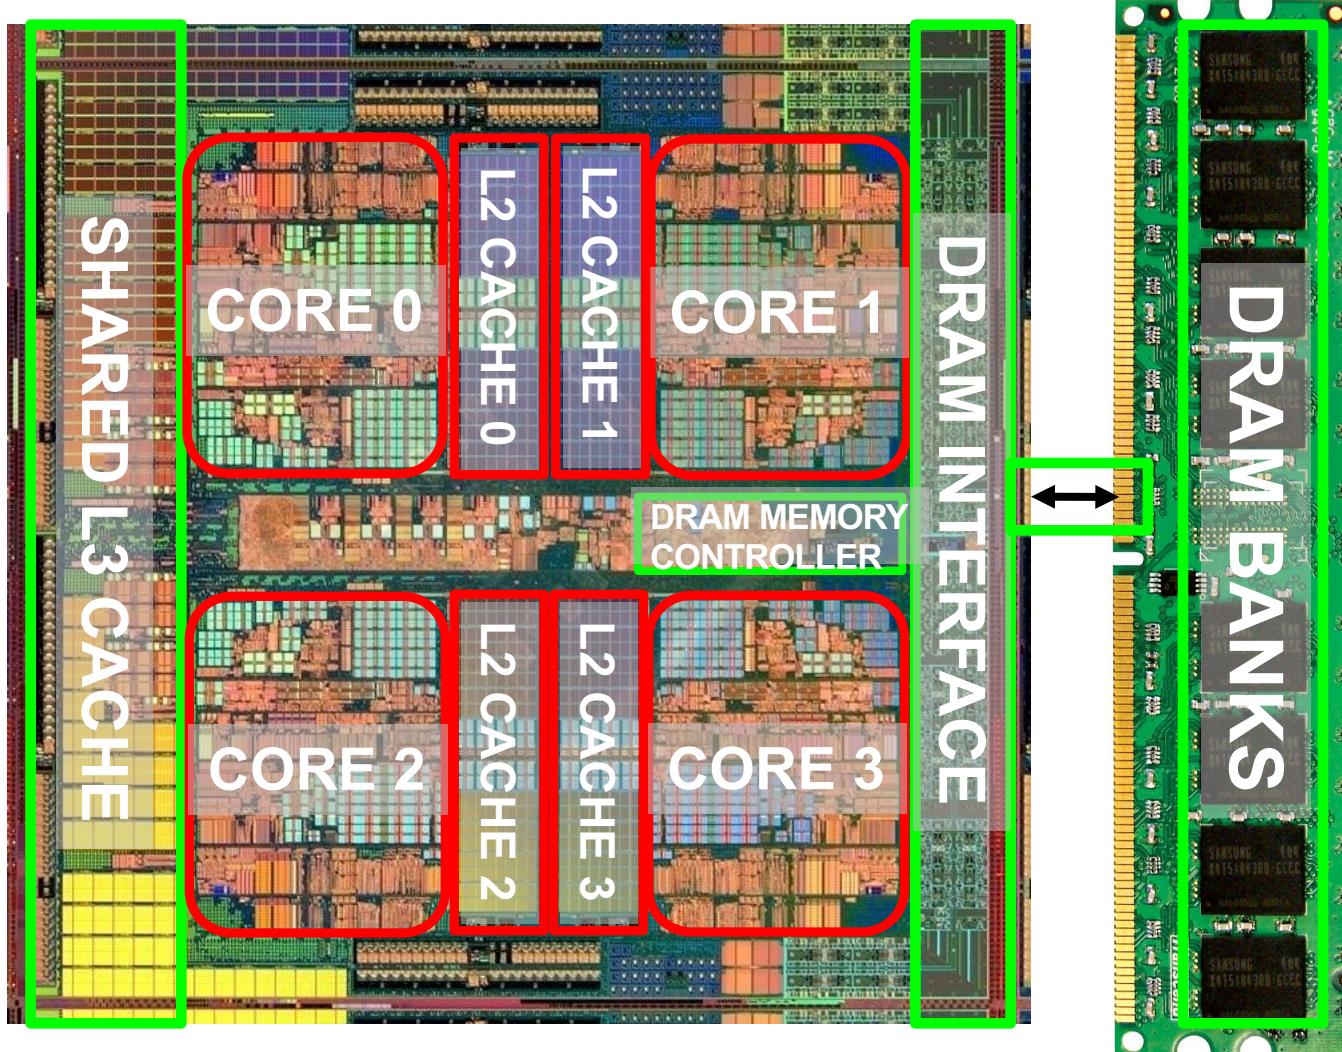
- Long latency memories have similar characteristics that need to be controlled.
- This lecture will use DRAM as an example, but many scheduling and control issues are similar in the design of controllers for other types of memories
 - Flash memory
 - Other emerging memory technologies
 - Phase Change Memory
 - Spin-Transfer Torque Magnetic Memory
 - These other technologies can also place other demands on the controller

DRAM Types

- DRAM has different types with different interfaces optimized for different purposes
 - Commodity: DDR, DDR2, DDR3, DDR4, DDR5, ...
 - Low power (for mobile): LPDDR1, ..., LPDDR5, ...
 - High bandwidth (for graphics): GDDR2, ..., GDDR5, ...
 - Low latency: eDRAM, RLDRAM, ...
 - 3D stacked: WIO, HBM, HMC, HBM2.0, ...
 - ...
- Underlying microarchitecture is fundamentally the same
- A flexible memory controller can support various DRAM types
- This complicates the memory controller
 - Difficult to support all types (and upgrades)
 - Analog interface is different for different DRAM types

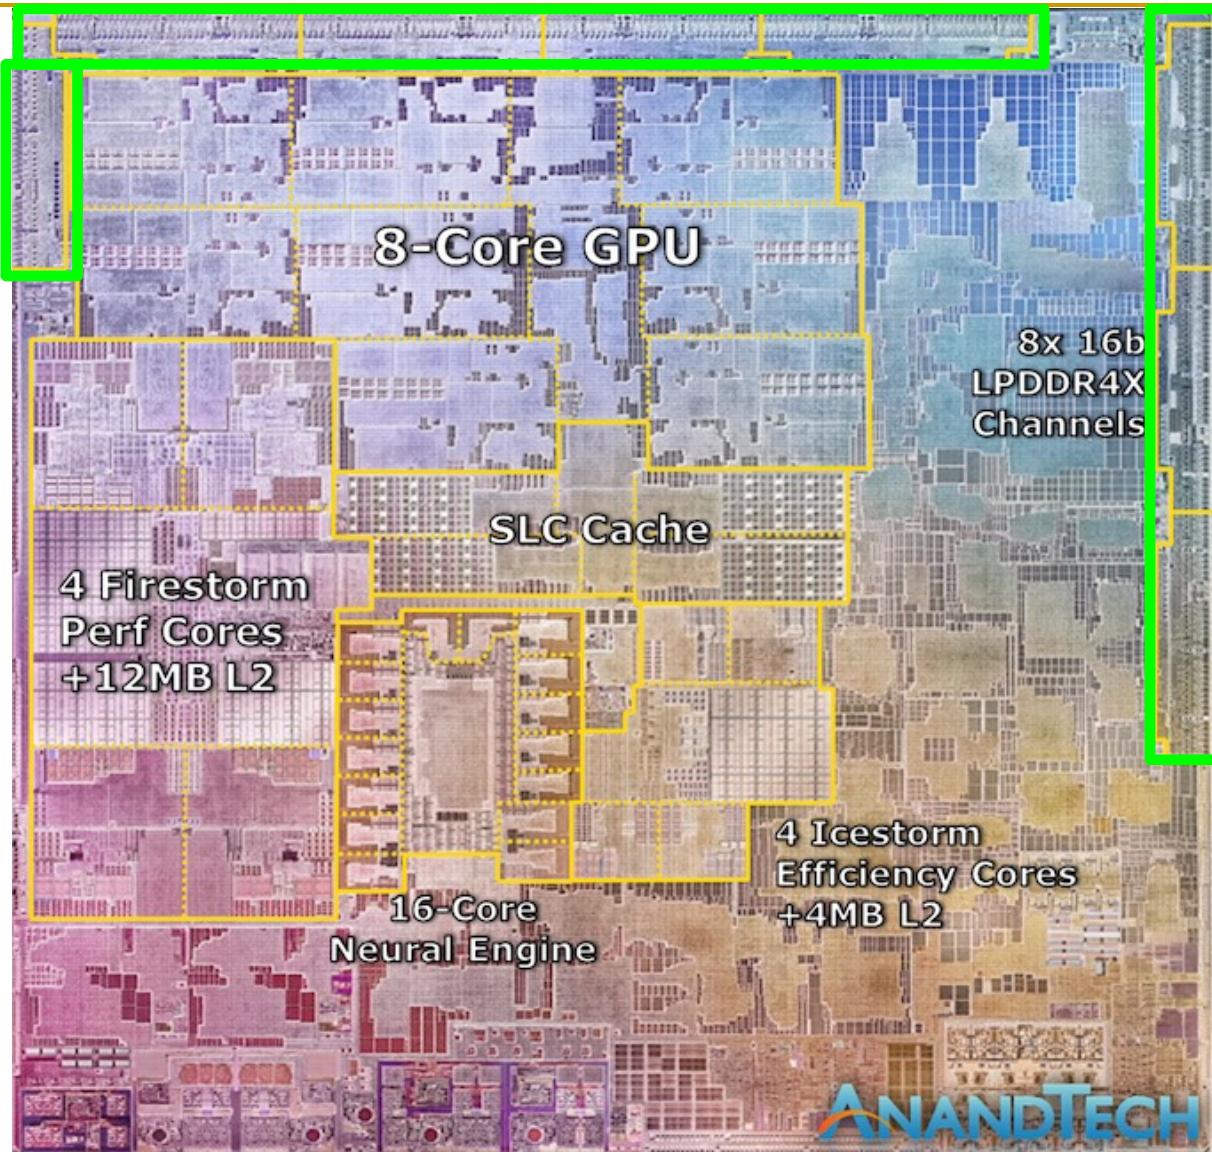
DRAM Control Logic Is Large

Multi-Core
Chip



*Die photo credit: AMD Barcelona

DRAM Control Logic Is Large



Apple M1,
2021

DRAM Controller: Functions

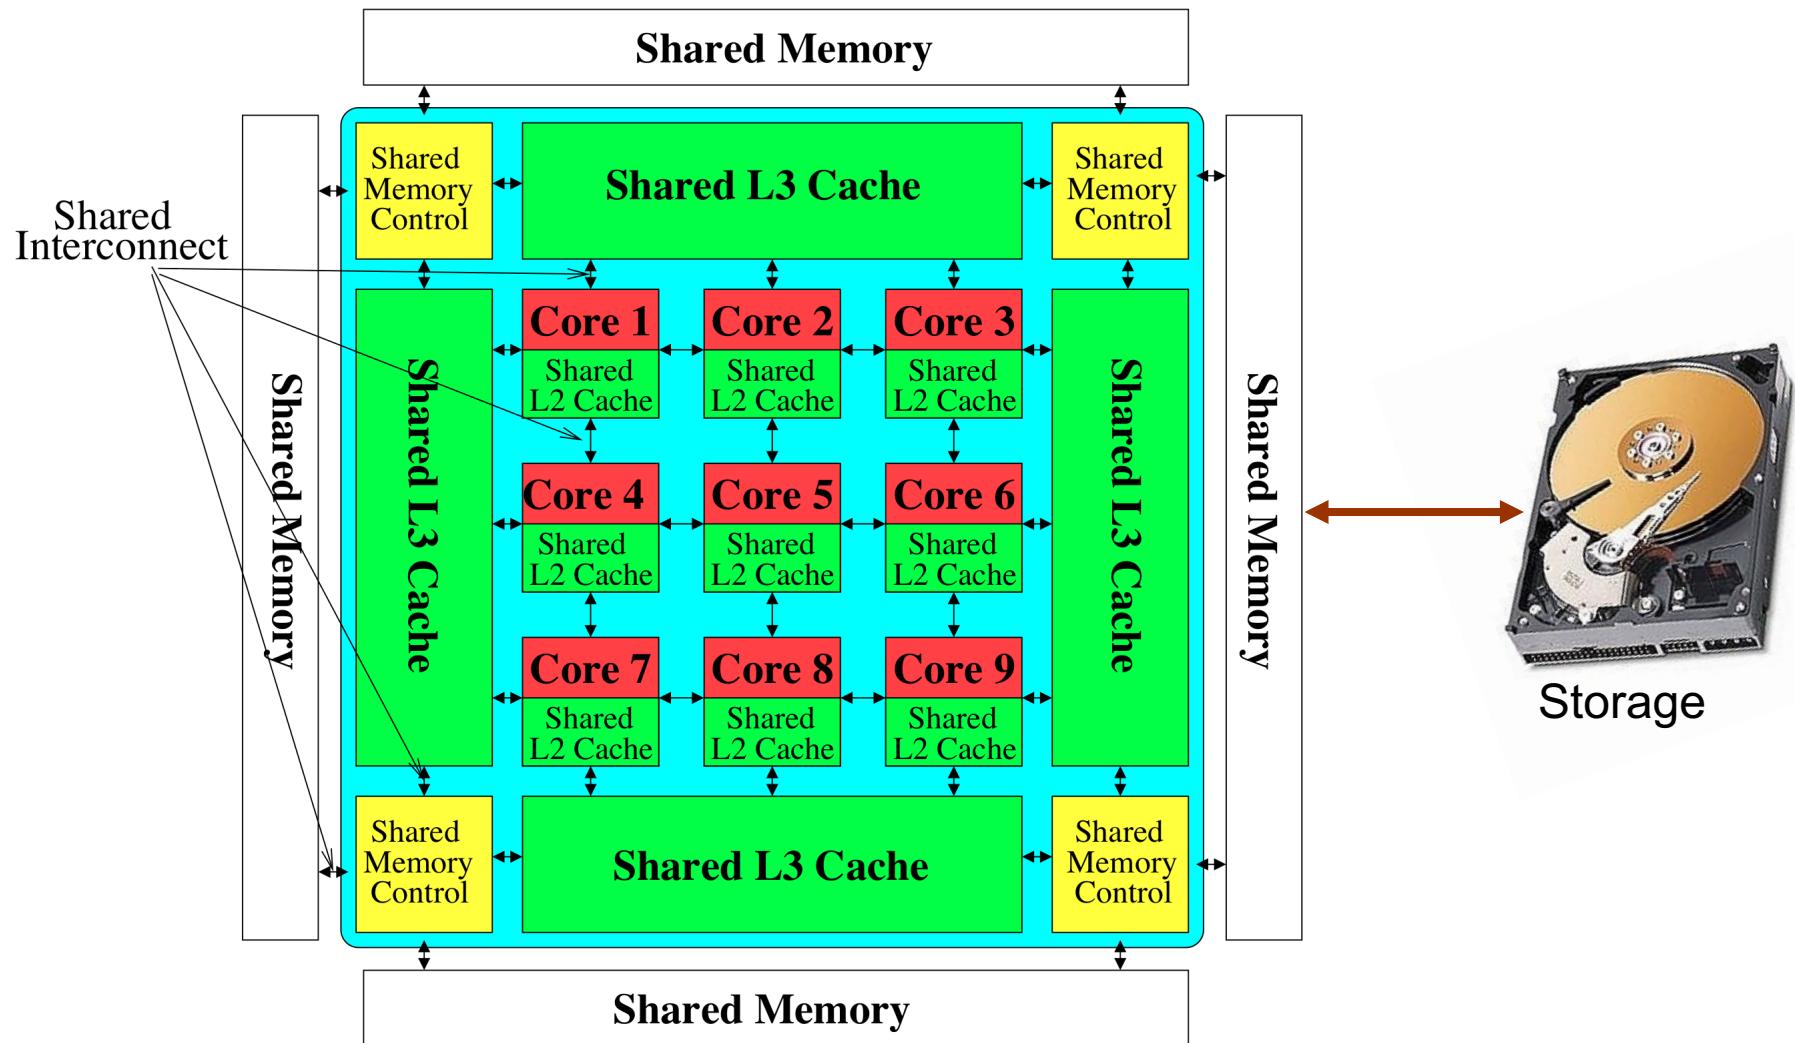
- Ensure correct operation of DRAM (refresh and timing)
- Service DRAM requests while obeying timing constraints of DRAM chips
 - Constraints: resource conflicts (bank, bus, channel), minimum write-to-read delays
 - Translate requests to DRAM command sequences
- Buffer and schedule requests for high performance + QoS
 - Reordering, row-buffer, bank, rank, bus management
- Manage power consumption and thermals in DRAM
 - Turn on/off DRAM chips, manage power modes

DRAM Scheduling Policies

- A scheduling policy is a request prioritization order
- Prioritization can be based on
 - Request age
 - Row buffer hit/miss status
 - Request type (prefetch, read, write)
 - Requestor type (load miss or store miss)
 - Request criticality
 - Oldest miss in the core?
 - How many instructions in core are dependent on it?
 - Will it stall the processor?
 - Interference caused to other cores
 - ...

Shared Resource Design for Multi-Core Systems

Memory System: A *Shared Resource* View



Most of the system is dedicated to storing and moving data

Resource Sharing Concept

- Idea: Instead of dedicating a hardware resource to a hardware context, allow multiple contexts to use it
 - Example resources: functional units, pipeline, caches, buses, memory, interconnects, storage
- Why?
 - + Resource sharing improves utilization/efficiency → throughput
 - When a resource is left idle by one thread, another thread can use it; no need to replicate shared data
 - + Reduces communication latency
 - For example, shared data kept in the same cache in SMT processors
 - + Compatible with the shared memory model

Resource Sharing Disadvantages

- Resource sharing results in **contention for resources**
 - When the resource is not idle, another thread cannot use it
 - If space is occupied by one thread, another thread needs to re-occupy it
- Sometimes reduces each or some thread's performance
 - Thread performance can be worse than when it is run alone
- Eliminates performance isolation → inconsistent performance across runs
 - Thread performance depends on co-executing threads
- Uncontrolled (free-for-all) sharing degrades QoS
 - Causes unfairness, starvation

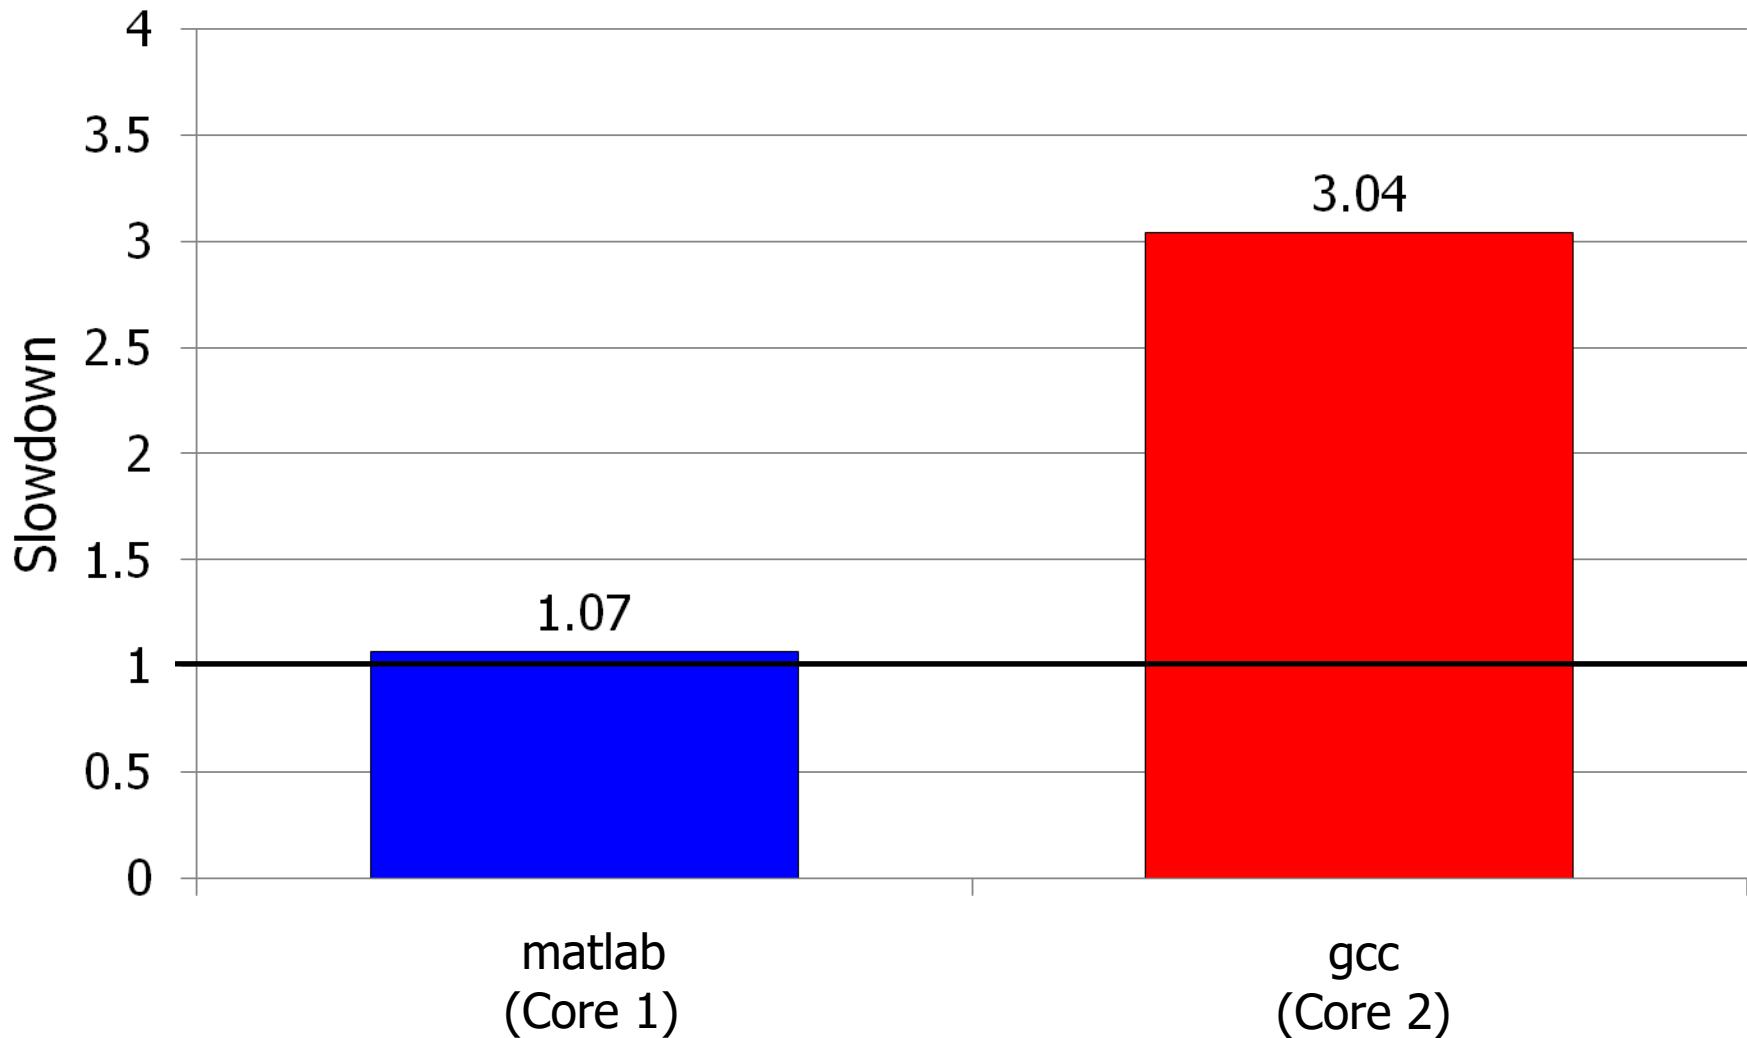
Need to efficiently and fairly utilize shared resources

Inter-Thread/Application Interference

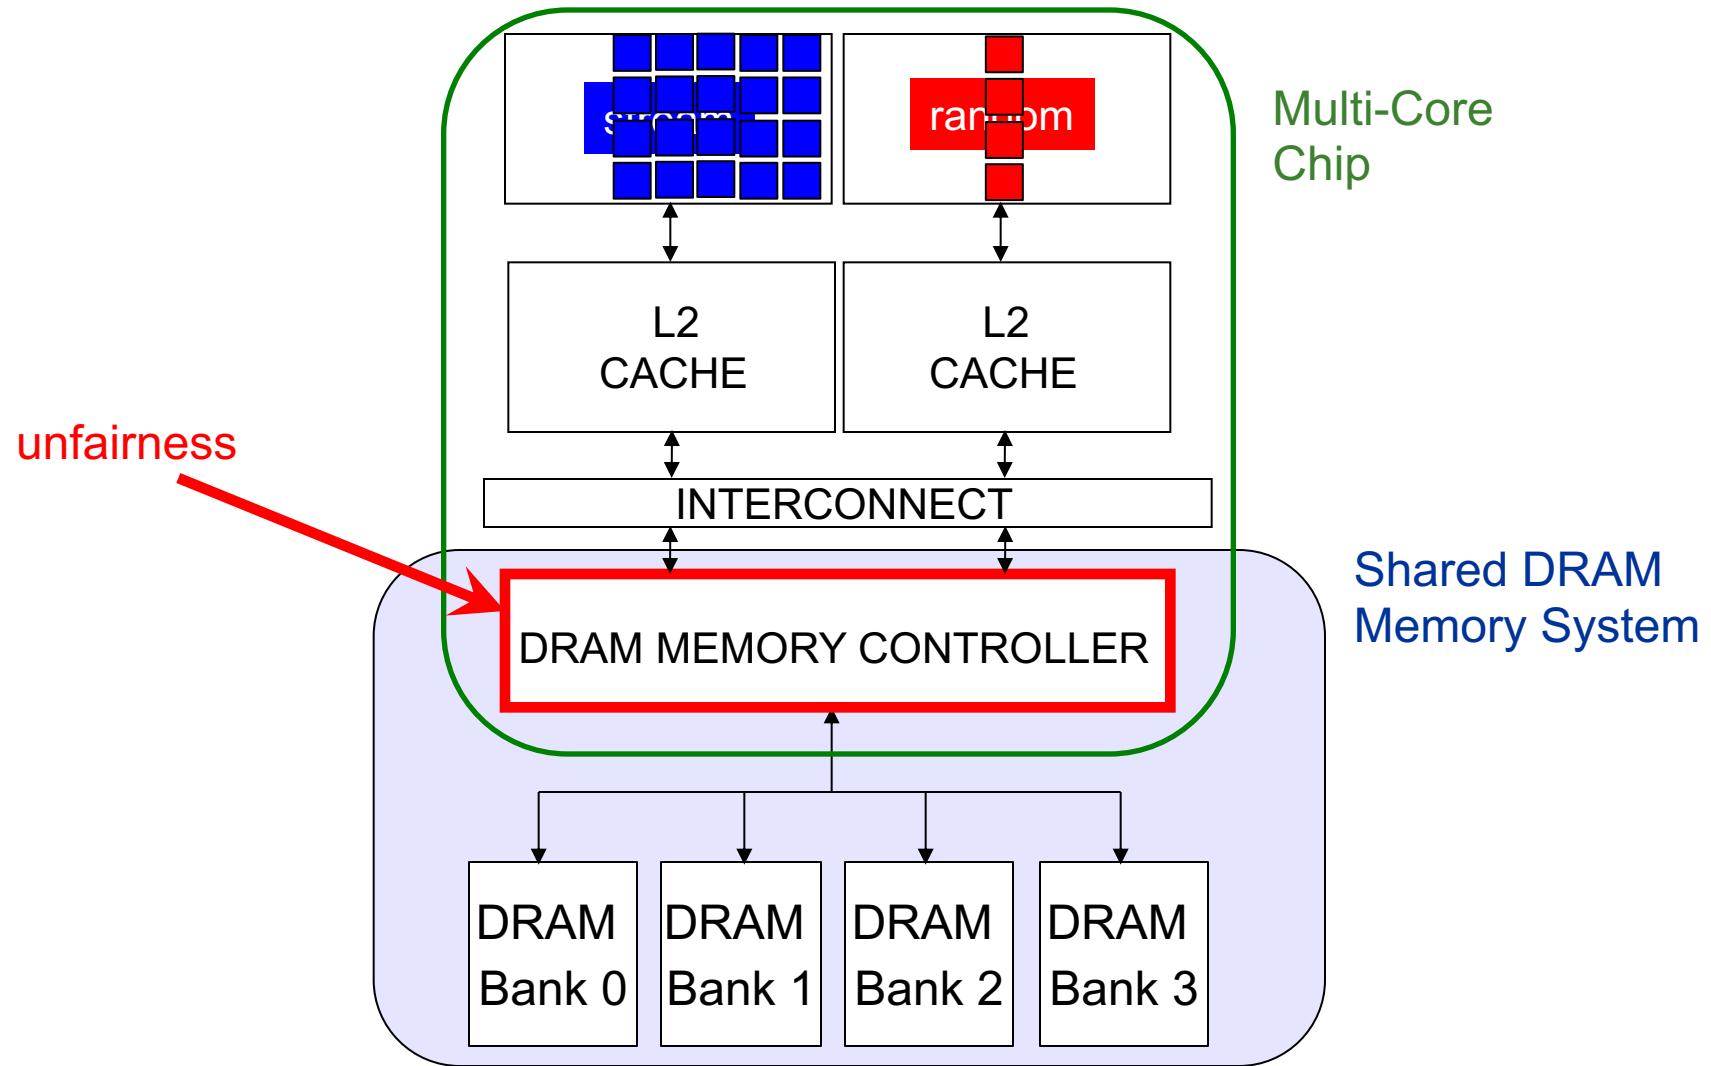
- Problem: Threads share the memory system, but memory system does not distinguish between threads' requests

- Existing memory systems
 - Free-for-all, shared based on demand
 - Control algorithms thread-unaware and thread-unfair
 - Aggressive threads can deny service to others
 - Do not try to reduce or control inter-thread interference

Unfair Slowdowns due to Interference



Uncontrolled Interference: An Example



A Memory Performance Hog

```
// initialize large arrays A, B  
  
for (j=0; j<N; j++) {  
    index = j*linesize; streaming  
    A[index] = B[index];  
    ...  
}
```

STREAM

- Sequential memory access
- Very high row buffer locality (96% hit rate)
- Memory intensive

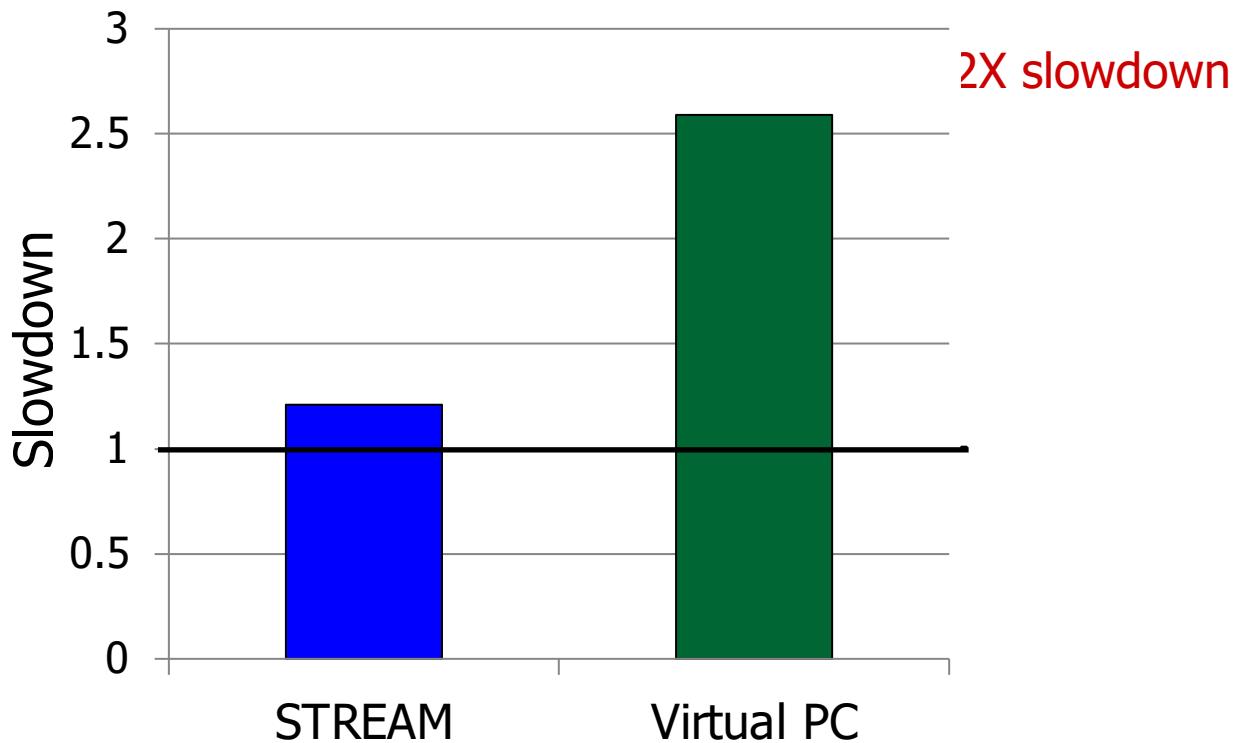
```
// initialize large arrays A, B  
  
for (j=0; j<N; j++) {  
    index = rand(); random  
    A[index] = B[index];  
    ...  
}
```

RANDOM

- Random memory access
- Very low row buffer locality (3% hit rate)
- Similarly memory intensive

Moscibroda and Mutlu, “Memory Performance Attacks,” USENIX Security 2007.

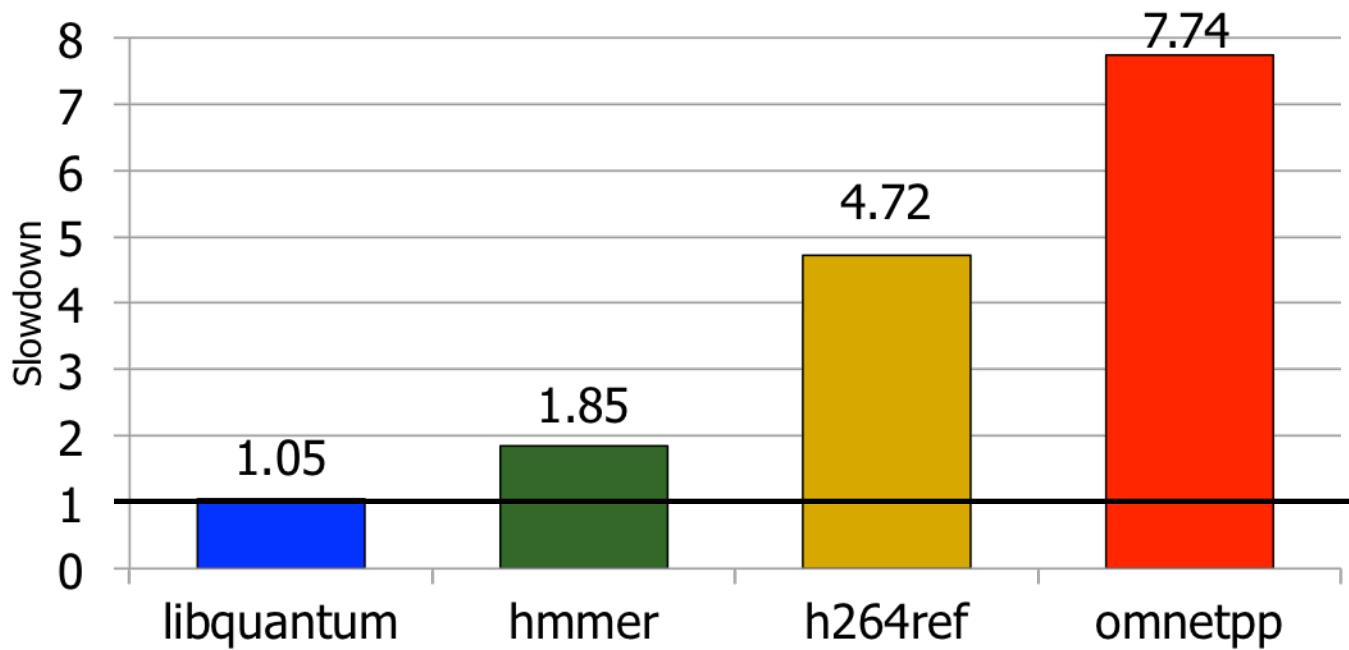
Effect of the Memory Performance Hog



Results on Intel Pentium D running Windows XP
(Similar results for Intel Core Duo and AMD Turion, and on Fedora Linux)

Moscibroda and Mutlu, “Memory Performance Attacks,” USENIX Security 2007.

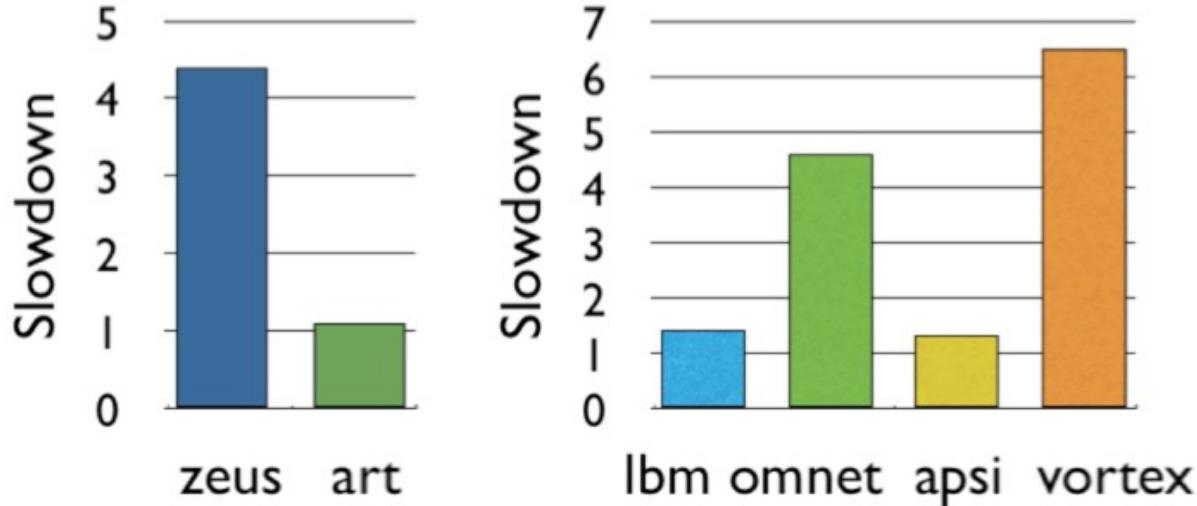
Greater Problem with More Cores



- Vulnerable to denial of service (DoS)
- Unable to enforce priorities or SLAs
- Low system performance

Uncontrollable, unpredictable system

Greater Problem with More Cores



- Vulnerable to denial of service (DoS)
- Unable to enforce priorities or SLAs
- Low system performance

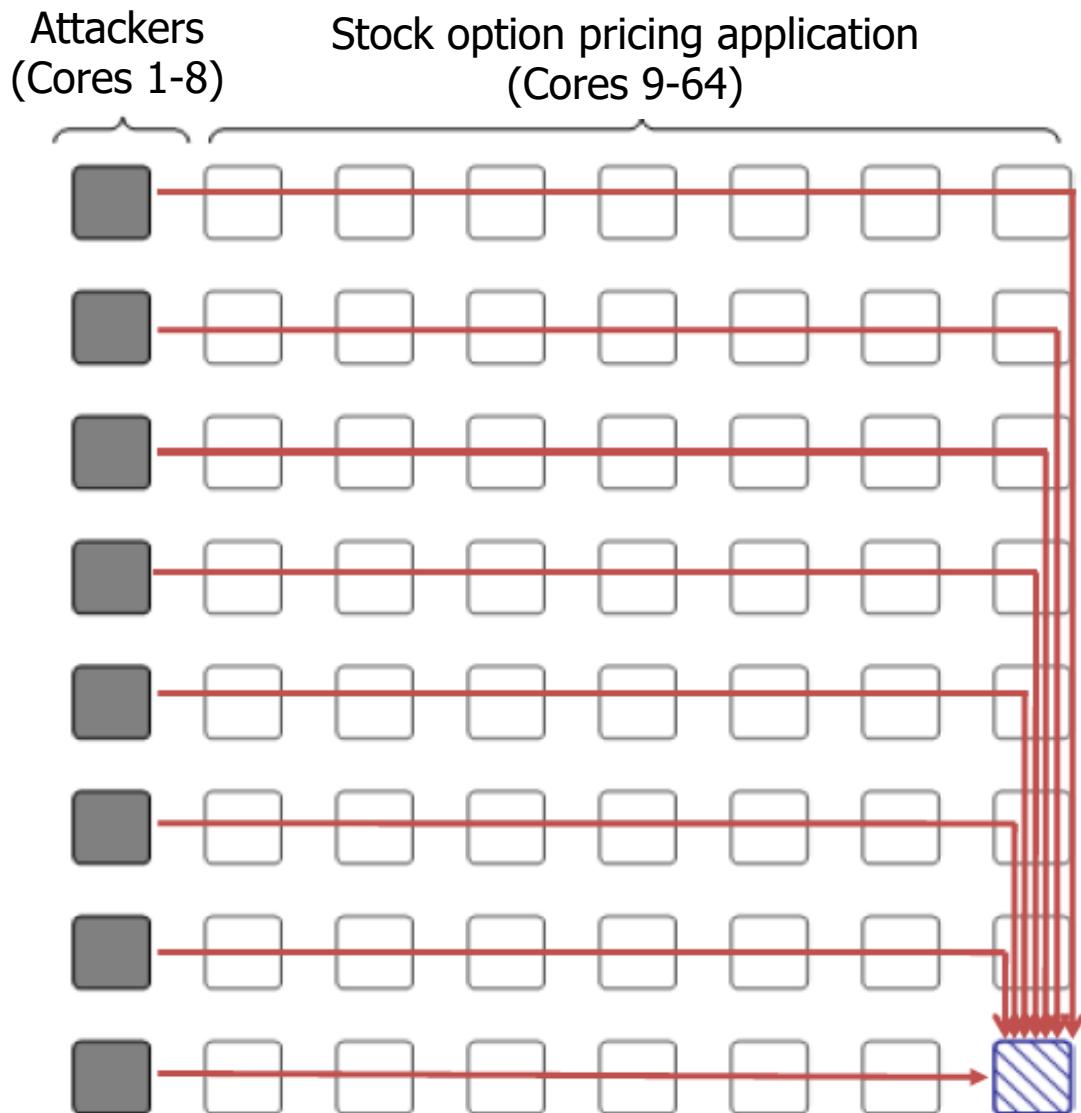
Uncontrollable, unpredictable system

Distributed DoS in Networked Multi-Core Systems

Cores connected via
packet-switched
routers on chip

~5000X latency increase

Grot, Hestness, Keckler, Mutlu,
“Preemptive virtual clock: A Flexible,
Efficient, and Cost-effective QOS
Scheme for Networks-on-Chip,”
MICRO 2009.



More on Memory Performance Attacks

- Thomas Moscibroda and Onur Mutlu,
**"Memory Performance Attacks: Denial of Memory Service
in Multi-Core Systems"**
Proceedings of the 16th USENIX Security Symposium (USENIX SECURITY), pages 257-274, Boston, MA, August 2007. Slides
(ppt)

Memory Performance Attacks: Denial of Memory Service in Multi-Core Systems

*Thomas Moscibroda Onur Mutlu
Microsoft Research
{moscitho,onur}@microsoft.com*

More on Interconnect Based Starvation

- Boris Grot, Stephen W. Keckler, and Onur Mutlu,
"Preemptive Virtual Clock: A Flexible, Efficient, and Cost-effective QOS Scheme for Networks-on-Chip"
Proceedings of the 42nd International Symposium on Microarchitecture (MICRO), pages 268-279, New York, NY, December 2009. [Slides \(pdf\)](#)

Preemptive Virtual Clock: A Flexible, Efficient, and Cost-effective QOS Scheme for Networks-on-Chip

Boris Grot

Stephen W. Keckler

Onur Mutlu[†]

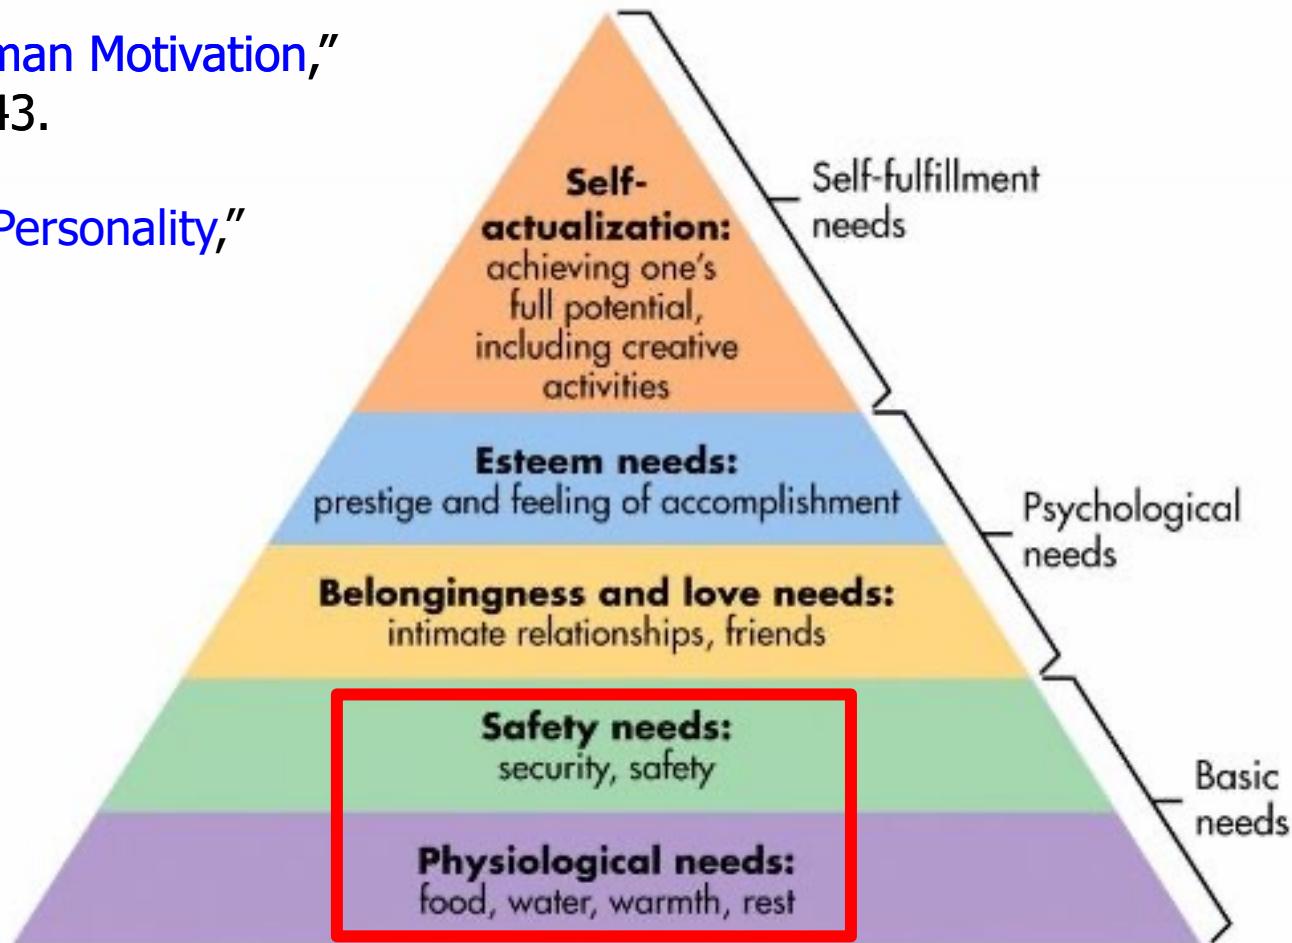
Department of Computer Sciences
The University of Texas at Austin
{bgrot, skeckler}@cs.utexas.edu}

[†]Computer Architecture Laboratory (CALCM)
Carnegie Mellon University
onur@cmu.edu

Maslow's (Human) Hierarchy of Needs

Maslow, "A Theory of Human Motivation,"
Psychological Review, 1943.

Maslow, "Motivation and Personality,"
Book, 1954-1970.



- Lack of QoS can be a **safety and security** problem

How Do We Solve The Problem?

- Inter-thread interference is uncontrolled in all memory resources
 - Memory controller
 - Interconnect
 - Caches
- We need to control it
 - i.e., design an interference-aware (QoS-aware) memory system

Fundamental Interference Control Techniques

- Goal: to reduce/control inter-thread memory interference

1. Prioritization or request scheduling

2. Data mapping to banks/channels/ranks

3. Core/source throttling

4. Application/thread scheduling

QoS-Aware Memory Scheduling: Revolution & Evolution

QoS-Aware Memory Scheduling: Evolution

- **Stall-time fair memory scheduling** [Mutlu+ MICRO'07]
 - Idea: Estimate and balance thread slowdowns
 - Takeaway: **Proportional thread progress improves performance, especially when threads are “heavy”** (memory intensive)
- **Parallelism-aware batch scheduling** [Mutlu+ ISCA'08, Top Picks'09]
 - Idea: Rank threads and service in rank order (to preserve bank parallelism); batch requests to prevent starvation
 - Takeaway: **Preserving within-thread bank-parallelism improves performance**; request batching improves fairness
- **ATLAS memory scheduler** [Kim+ HPCA'10]
 - Idea: Prioritize threads that have attained the least service from the memory scheduler
 - Takeaway: **Prioritizing “light” threads improves performance**

QoS-Aware Memory Scheduling: Evolution

- Thread cluster memory scheduling [Kim+ MICRO'10, Top Picks'11]
 - Idea: Cluster threads into two groups (latency vs. bandwidth sensitive); prioritize the latency-sensitive ones; employ a fairness policy in the bandwidth sensitive group
 - Takeaway: **Heterogeneous scheduling policy that is different based on thread behavior maximizes both performance and fairness**
- Integrated Memory Channel Partitioning and Scheduling [Muralidhara+ MICRO'11]
 - Idea: Only prioritize very latency-sensitive threads in the scheduler; mitigate all other applications' interference via channel partitioning
 - Takeaway: **Intelligently combining application-aware channel partitioning and memory scheduling provides better performance than either**

QoS-Aware Memory Scheduling: Evolution

- Parallel application memory scheduling [Ebrahimi+ MICRO'11]
 - Idea: Identify and prioritize limiter threads of a multithreaded application in the memory scheduler; provide fast and fair progress to non-limiter threads
 - Takeaway: Carefully prioritizing between limiter and non-limiter threads of a parallel application improves performance
- Staged memory scheduling [Ausavarungnirun+ ISCA'12]
 - Idea: Divide the functional tasks of an application-aware memory scheduler into multiple distinct stages, where each stage is significantly simpler than a monolithic scheduler
 - Takeaway: Staging enables the design of a scalable and relatively simpler application-aware memory scheduler that works on very large request buffers

QoS-Aware Memory Scheduling: Evolution

- MISE: Memory Slowdown Model [Subramanian+ HPCA'13]
 - Idea: Estimate the performance of a thread by estimating its change in memory request service rate when run alone vs. shared → use this simple model to estimate slowdown to design a scheduling policy that provides predictable performance or fairness
 - Takeaway: Request service rate of a thread is a good proxy for its performance; alone request service rate can be estimated by giving high priority to the thread in memory scheduling for a while
- ASM: Application Slowdown Model [Subramanian+ MICRO'15]
 - Idea: Extend MISE to take into account cache+memory interference
 - Takeaway: Cache access rate of an application can be estimated accurately and is a good proxy for application performance

QoS-Aware Memory Scheduling: Evolution

- **BLISS: Blacklisting Memory Scheduler** [Subramanian+ ICCD'14, TPDS'16]
 - Idea: Deprioritize (i.e., blacklist) a thread that has consecutively serviced a large number of requests
 - Takeaway: **Blacklisting greatly reduces interference enables the scheduler to be simple without requiring full thread ranking**
- **DASH: Deadline-Aware Memory Scheduler** [Usui+ TACO'16]
 - Idea: Balance prioritization between CPUs, GPUs and Hardware Accelerators (HWA) by keeping HWA progress in check vs. deadlines such that HWAs do not hog performance and appropriately distinguishing between latency-sensitive vs. bandwidth-sensitive CPU workloads
 - Takeaway: **Proper control of HWA progress and application-aware CPU prioritization leads to better system performance while meeting HWA deadlines**

QoS-Aware Memory Scheduling: Evolution

- Prefetch-aware shared resource management [Ebrahimi+ ISCA'11] [Ebrahimi+ MICRO'09] [Ebrahimi+ HPCA'09] [Lee+ MICRO'08'09]
 - Idea: Prioritize prefetches depending on how they affect system performance; even accurate prefetches can degrade performance of the system
 - Takeaway: Carefully controlling and prioritizing prefetch requests improves performance and fairness
- DRAM-Aware last-level cache policies and write scheduling [Lee+ HPS Tech Report'10] [Seshadri+ ISCA'14]
 - Idea: Design cache eviction and replacement policies such that they proactively exploit the state of the memory controller and DRAM (e.g., proactively evict data from the cache that hit in open rows)
 - Takeaway: Coordination of last-level cache and DRAM policies improves performance and fairness; writes should not be ignored

QoS-Aware Memory Scheduling: Evolution

- **FIRM: Memory Scheduling for NVM** [Zhao+ MICRO'14]
 - Idea: Carefully handle write-read prioritization with coarse-grained batching and application-aware scheduling
 - Takeaway: Carefully controlling and prioritizing write requests improves performance and fairness; write requests are especially critical in NVMs
- **Criticality-Aware Memory Scheduling for GPUs** [Jog+ SIGMETRICS'16]
 - Idea: Prioritize latency-critical cores' requests in a GPU system
 - Takeaway: Need to carefully balance locality and criticality to make sure performance improves by taking advantage of both
- **Worst-case Execution Time Based Memory Scheduling for Real-Time Systems** [Kim+ RTAS'14, JRTS'16]

Stall-Time Fair Memory Scheduling

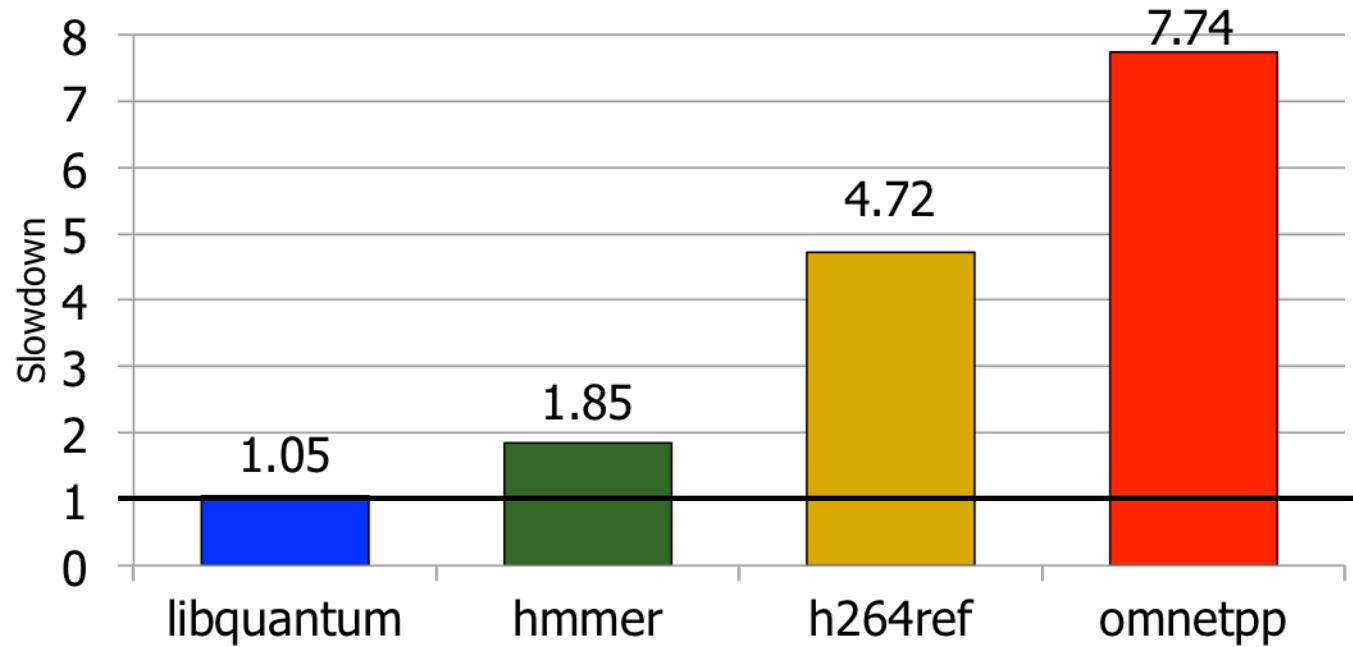
Onur Mutlu and Thomas Moscibroda,

"Stall-Time Fair Memory Access Scheduling for Chip Multiprocessors"

40th International Symposium on Microarchitecture (MICRO),

pages 146-158, Chicago, IL, December 2007. Slides (ppt)

The Problem: Unfairness



- Vulnerable to denial of service (DoS)
- Unable to enforce priorities or SLAs
- Low system performance

Uncontrollable, unpredictable system

STFM Scheduling Algorithm [MICRO'07]

- For each thread, the DRAM controller
 - Tracks ST_{shared}
 - Estimates ST_{alone}
- Each cycle, the DRAM controller
 - Computes Slowdown = $ST_{\text{shared}}/ST_{\text{alone}}$ for threads with legal requests
 - Computes **unfairness** = MAX Slowdown / MIN Slowdown
- If unfairness < α
 - Use DRAM throughput oriented scheduling policy
- If **unfairness** $\geq \alpha$
 - Use fairness-oriented scheduling policy
 - (1) requests from thread with MAX Slowdown first
 - (2) row-hit first , (3) oldest-first

More on STFM

- Onur Mutlu and Thomas Moscibroda,
"Stall-Time Fair Memory Access Scheduling for Chip Multiprocessors"

Proceedings of the 40th International Symposium on Microarchitecture (MICRO), pages 146-158, Chicago, IL, December 2007. [[Summary](#)] [[Slides \(ppt\)](#)]

Stall-Time Fair Memory Access Scheduling for Chip Multiprocessors

Onur Mutlu Thomas Moscibroda

Microsoft Research
{onur,moscitho}@microsoft.com

Parallelism-Aware Batch Scheduling

Onur Mutlu and Thomas Moscibroda,

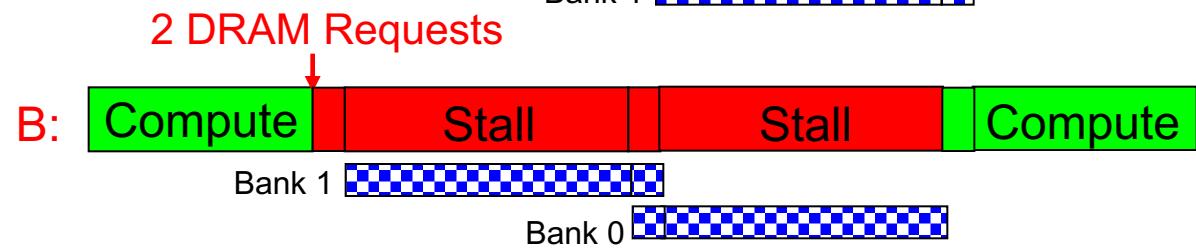
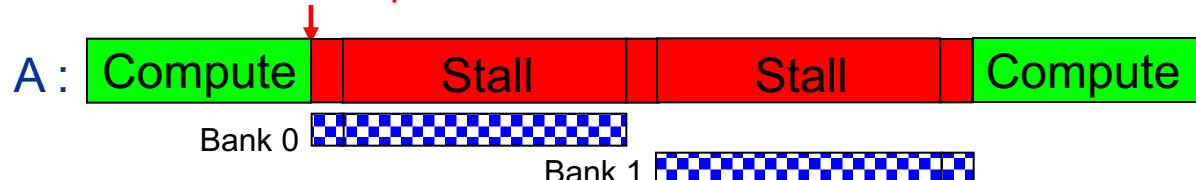
"Parallelism-Aware Batch Scheduling: Enhancing both Performance and Fairness of Shared DRAM Systems"

35th International Symposium on Computer Architecture (ISCA),
pages 63-74, Beijing, China, June 2008. [Slides \(ppt\)](#)

Parallelism-Aware Scheduler

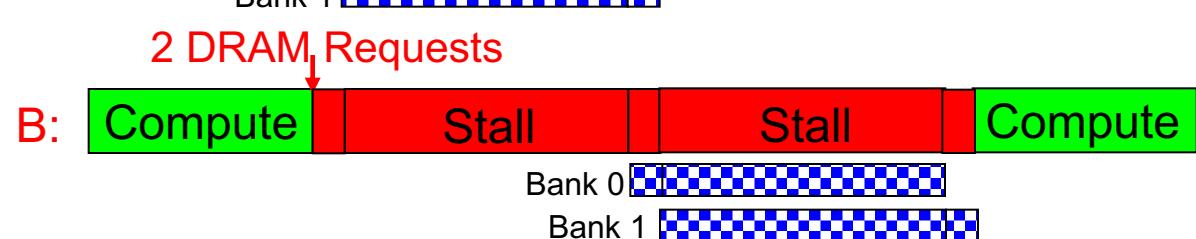
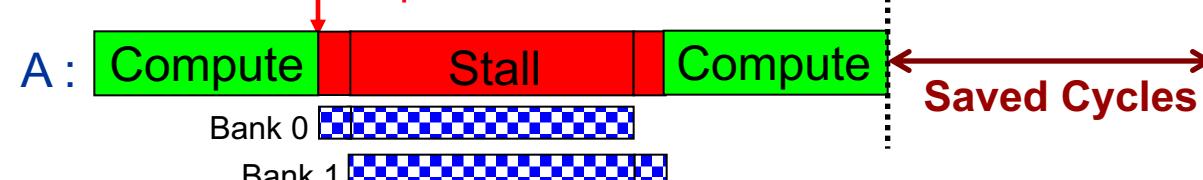
Baseline Scheduler:

2 DRAM Requests

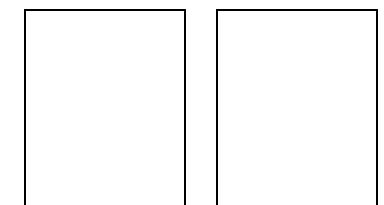


Parallelism-aware Scheduler:

2 DRAM Requests



Bank 0 Bank 1



Thread A: Bank 0, Row 1

Thread B: Bank 1, Row 99

Thread B: Bank 0, Row 99

Thread A: Bank 1, Row 1

Average stall-time:
~1.5 bank access
latencies

More on PAR-BS

- Onur Mutlu and Thomas Moscibroda,
"Parallelism-Aware Batch Scheduling: Enhancing both Performance and Fairness of Shared DRAM Systems"
Proceedings of the 35th International Symposium on Computer Architecture (ISCA), pages 63-74, Beijing, China, June 2008.
[Summary] [Slides (ppt)]
One of the 12 computer architecture papers of 2008 selected as Top Picks by IEEE Micro.

Parallelism-Aware Batch Scheduling:

Enhancing both Performance and Fairness of Shared DRAM Systems

Onur Mutlu Thomas Moscibroda
Microsoft Research
{onur,moscitho}@microsoft.com

More on PAR-BS

- Onur Mutlu and Thomas Moscibroda,

"Parallelism-Aware Batch Scheduling: Enabling High-Performance and Fair Memory Controllers"

*IEEE Micro, Special Issue: Micro's Top Picks from 2008 Computer Architecture Conferences (**MICRO TOP PICKS**)*, Vol. 29, No. 1, pages 22-32, January/February 2009.

PARALLELISM-AWARE BATCH SCHEDULING: ENABLING HIGH-PERFORMANCE AND FAIR SHARED MEMORY CONTROLLERS

UNCONTROLLED INTERTHREAD INTERFERENCE IN MAIN MEMORY CAN DESTROY INDIVIDUAL THREADS' MEMORY-LEVEL PARALLELISM, EFFECTIVELY SERIALIZING THE MEMORY REQUESTS OF A THREAD WHOSE LATENCIES WOULD OTHERWISE HAVE LARGELY OVERLAPPED, THEREBY REDUCING SINGLE-THREAD PERFORMANCE. THE PARALLELISM-AWARE BATCH SCHEDULER PRESERVES EACH THREAD'S MEMORY-LEVEL PARALLELISM, ENSURES FAIRNESS AND STARVATION FREEDOM, AND SUPPORTS SYSTEM-LEVEL THREAD PRIORITIES.

ATLAS Memory Scheduler

Yoongu Kim, Dongsu Han, Onur Mutlu, and Mor Harchol-Balter,

**"ATLAS: A Scalable and High-Performance
Scheduling Algorithm for Multiple Memory Controllers"**

*16th International Symposium on High-Performance Computer Architecture (HPCA),
Bangalore, India, January 2010.* [Slides \(pptx\)](#)

ATLAS: Summary

- Goal: To maximize system performance
- Main idea: Prioritize the thread that has attained the least service from the memory controllers (Adaptive per-Thread Least Attained Service Scheduling)
 - Rank threads based on attained service in the past time interval(s)
 - Enforce thread ranking in the memory scheduler during the current interval
- Why it works: Prioritizes “light” (memory non-intensive) threads that are more likely to keep their cores busy

More on ATLAS Memory Scheduler

- Yoongu Kim, Dongsu Han, Onur Mutlu, and Mor Harchol-Balter,
"ATLAS: A Scalable and High-Performance Scheduling Algorithm for Multiple Memory Controllers"

Proceedings of the 16th International Symposium on High-Performance Computer Architecture (HPCA), Bangalore, India, January 2010. Slides (pptx)

Best paper session. One of the four papers nominated for the Best Paper Award by the Program Committee.

ATLAS: A Scalable and High-Performance Scheduling Algorithm for Multiple Memory Controllers

Yoongu Kim Dongsu Han Onur Mutlu Mor Harchol-Balter

Carnegie Mellon University

End of Recall: Memory Controllers and Shared Resources

Fundamental Interference Control Techniques

- Goal: to reduce/control inter-thread memory interference

1. Prioritization or request scheduling
 2. Data mapping to banks/channels/ranks
 3. Core/source throttling
 4. Application/thread scheduling
-
- The diagram consists of four numbered items, each enclosed in a red rectangular box. To the right of the first item is the text "More of these" with an arrow pointing to the box. To the right of the second item is the text "Some of these" with an arrow pointing to the box. The third item has no accompanying text or arrow.

TCM: Thread Cluster Memory Scheduling

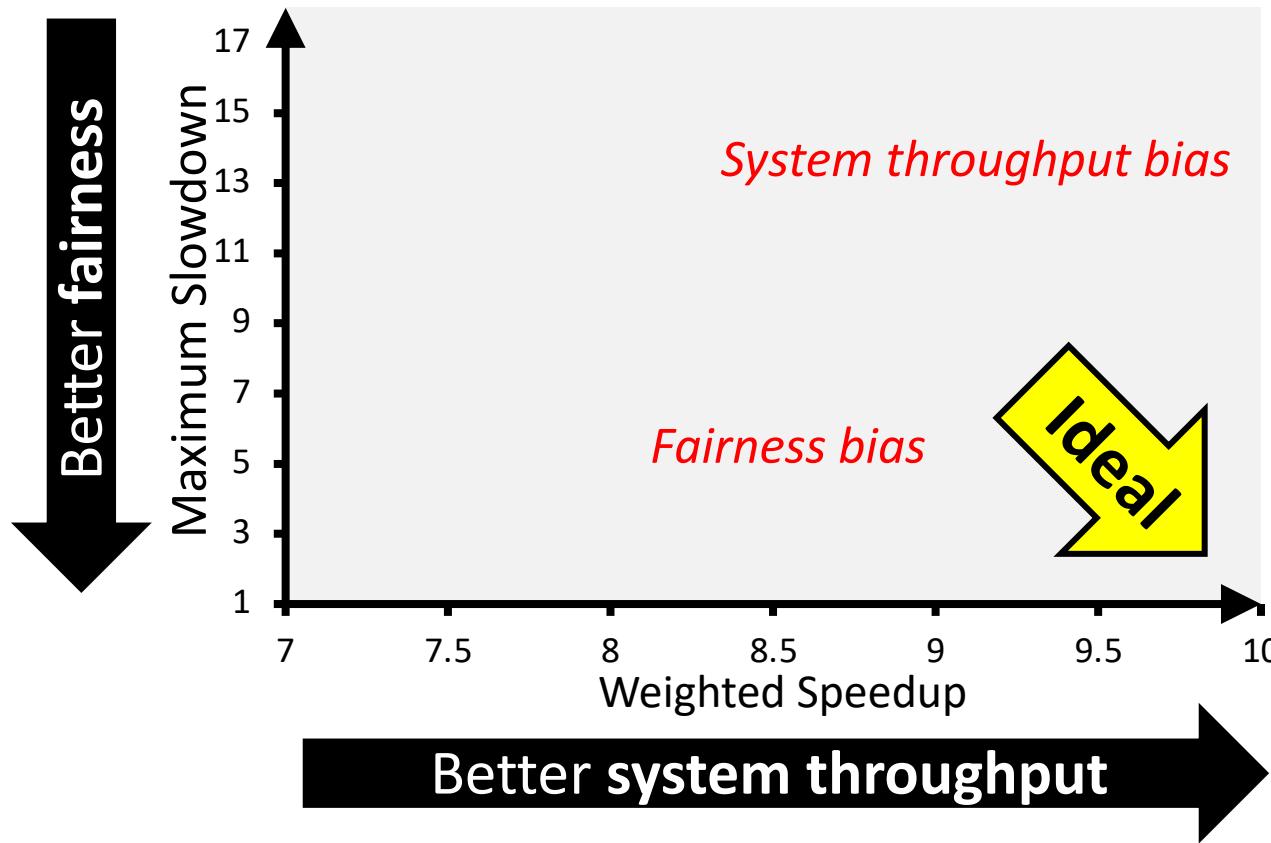
Yoongu Kim, Michael Papamichael, Onur Mutlu, and Mor Harchol-Balter,

**"Thread Cluster Memory Scheduling:
Exploiting Differences in Memory Access Behavior"**

43rd International Symposium on Microarchitecture (MICRO),
pages 65-76, Atlanta, GA, December 2010. [Slides \(pptx\)](#) [\(pdf\)](#)

Previous Scheduling Algorithms are Biased

24 cores, 4 memory controllers, 96 workloads

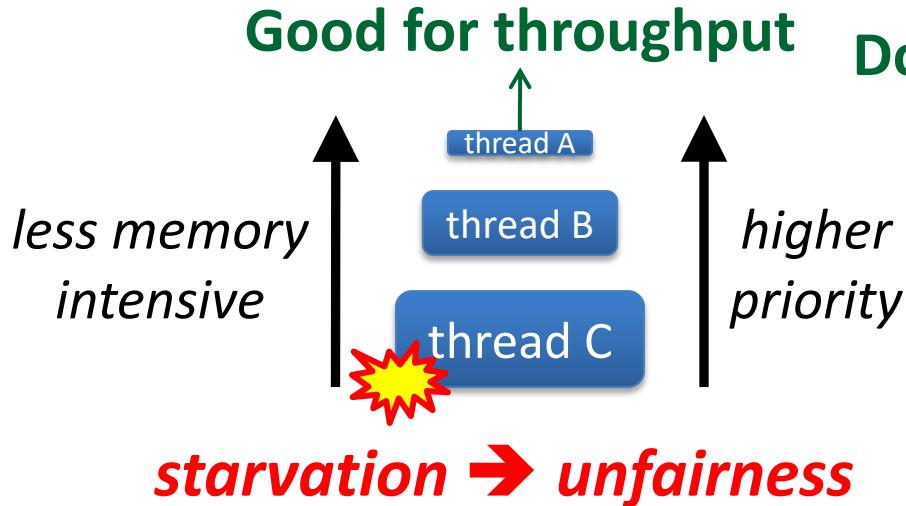


No previous memory scheduling algorithm provides both the best fairness and system throughput

Throughput vs. Fairness

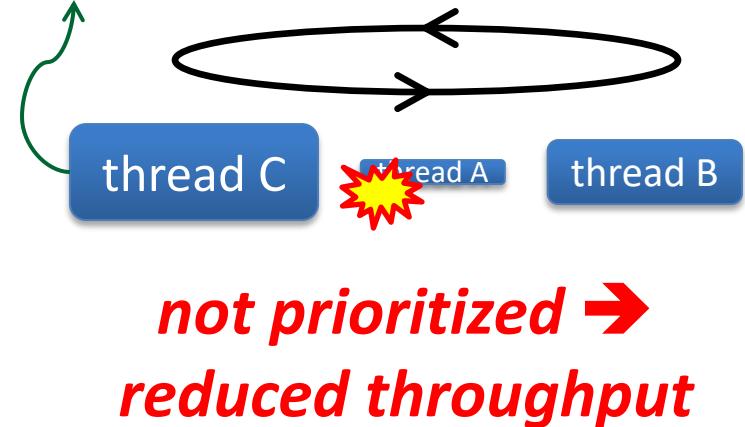
Throughput biased approach

Prioritize less memory-intensive threads



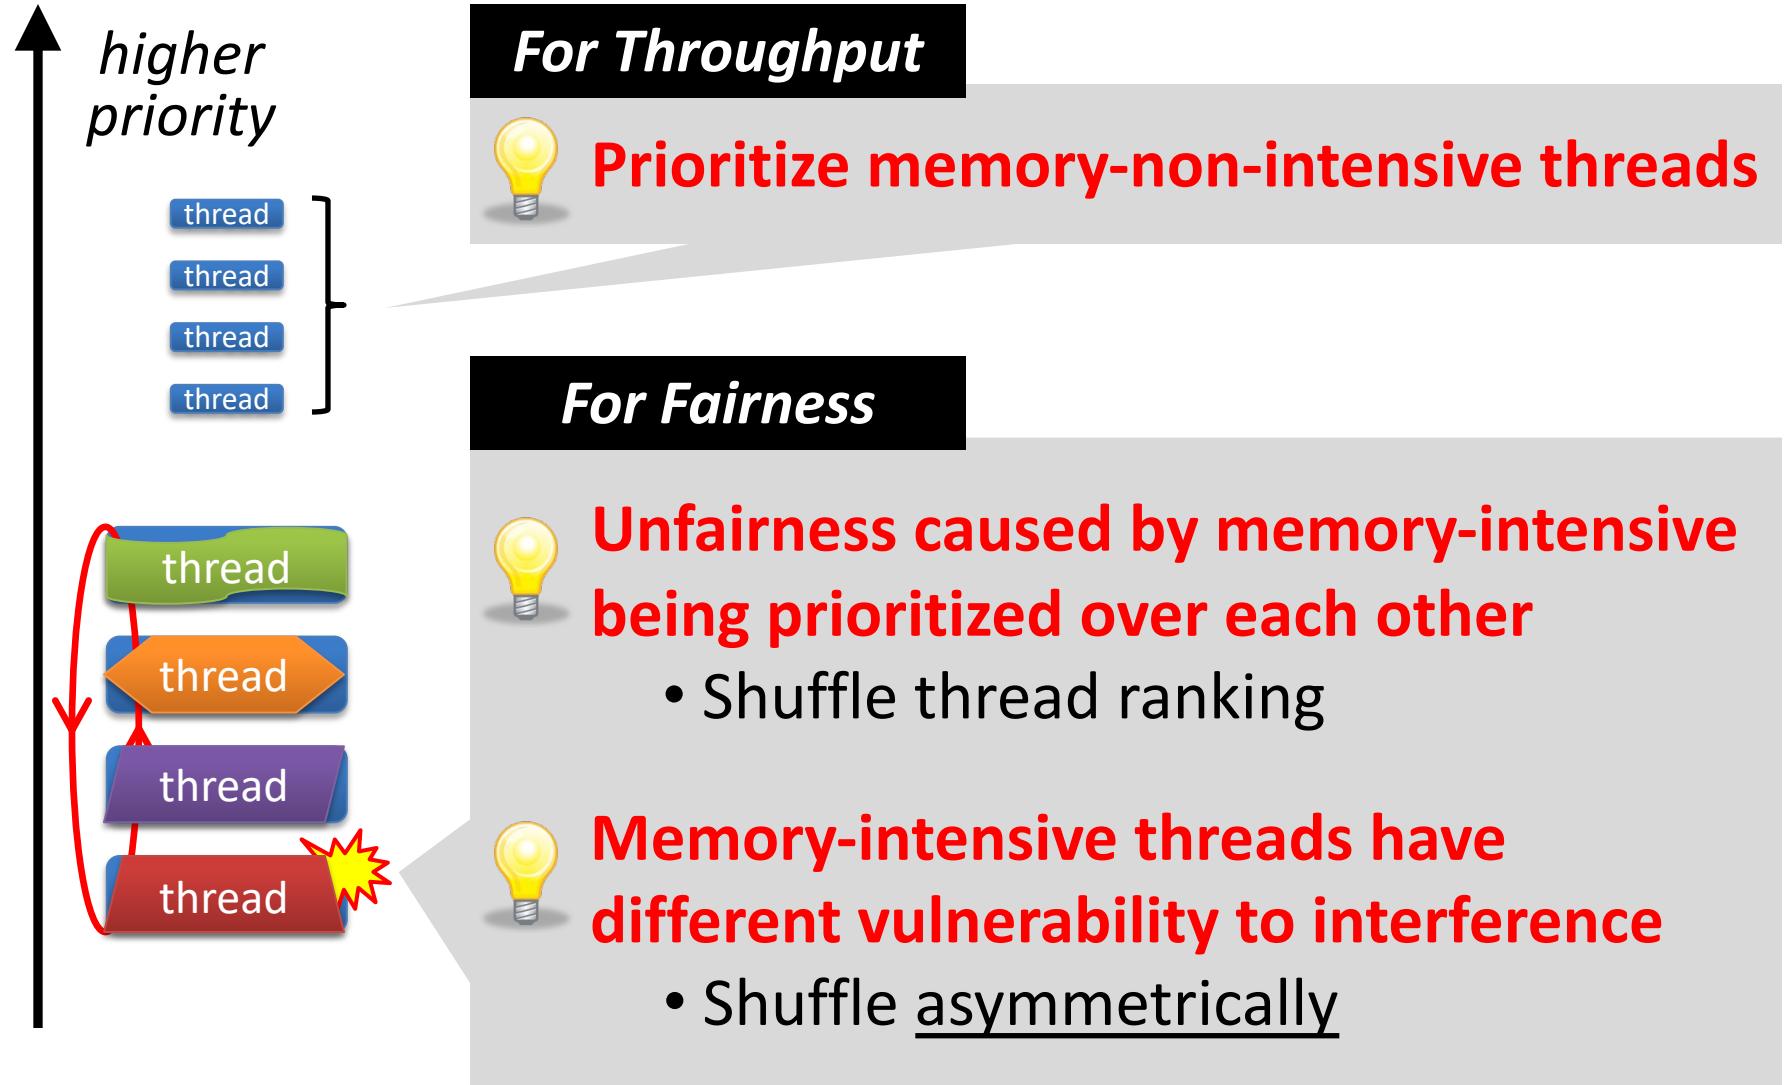
Fairness biased approach

Take turns accessing memory



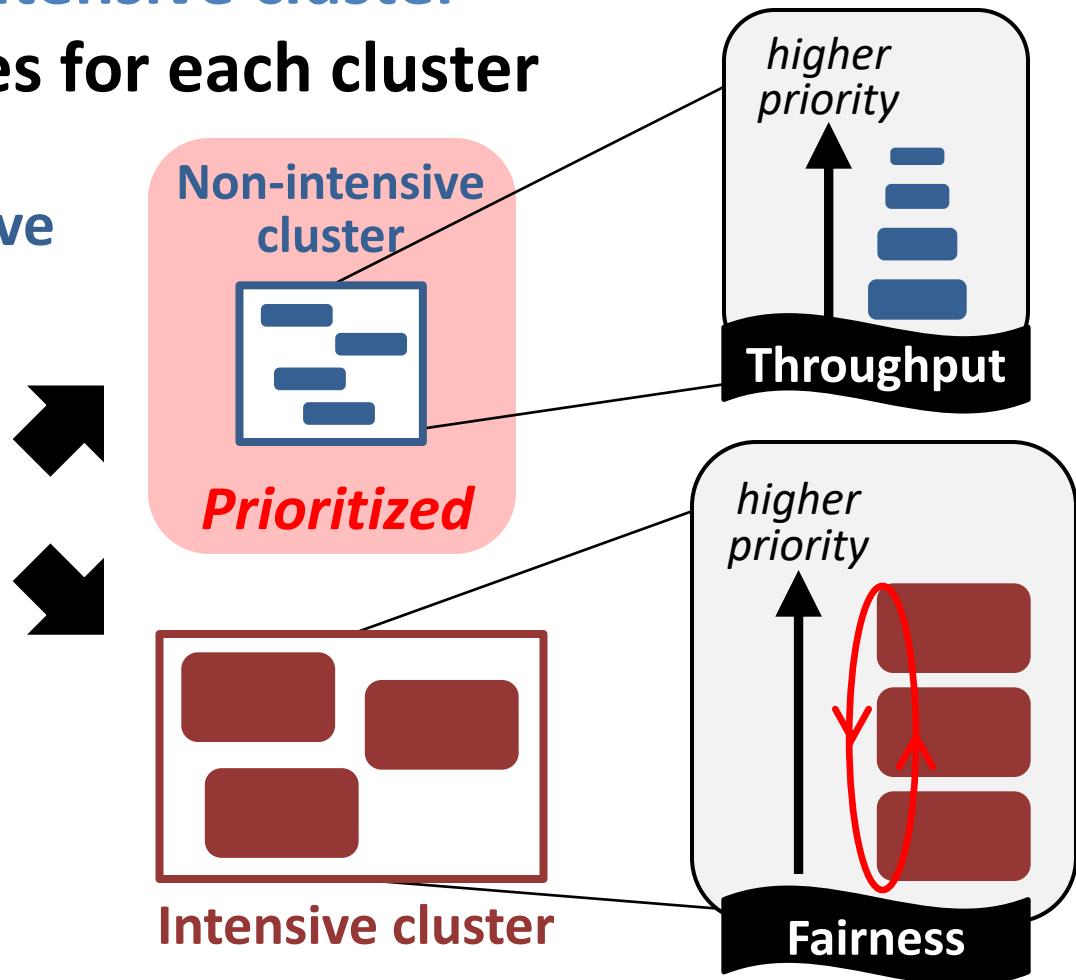
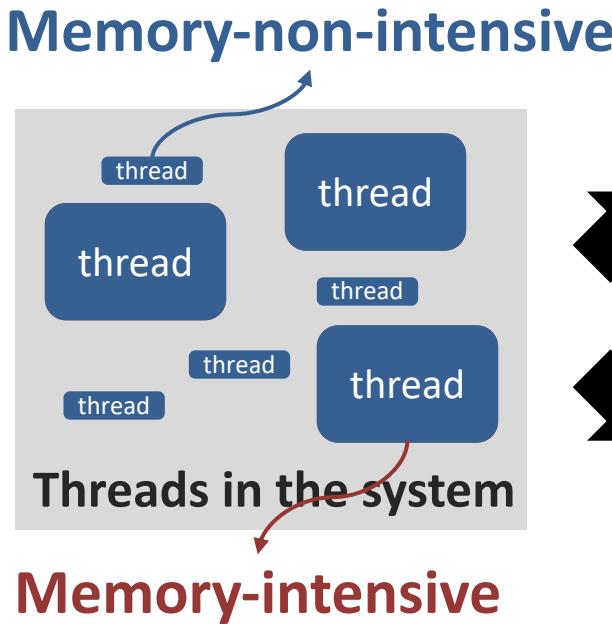
Single policy for all threads is insufficient

Achieving the Best of Both Worlds

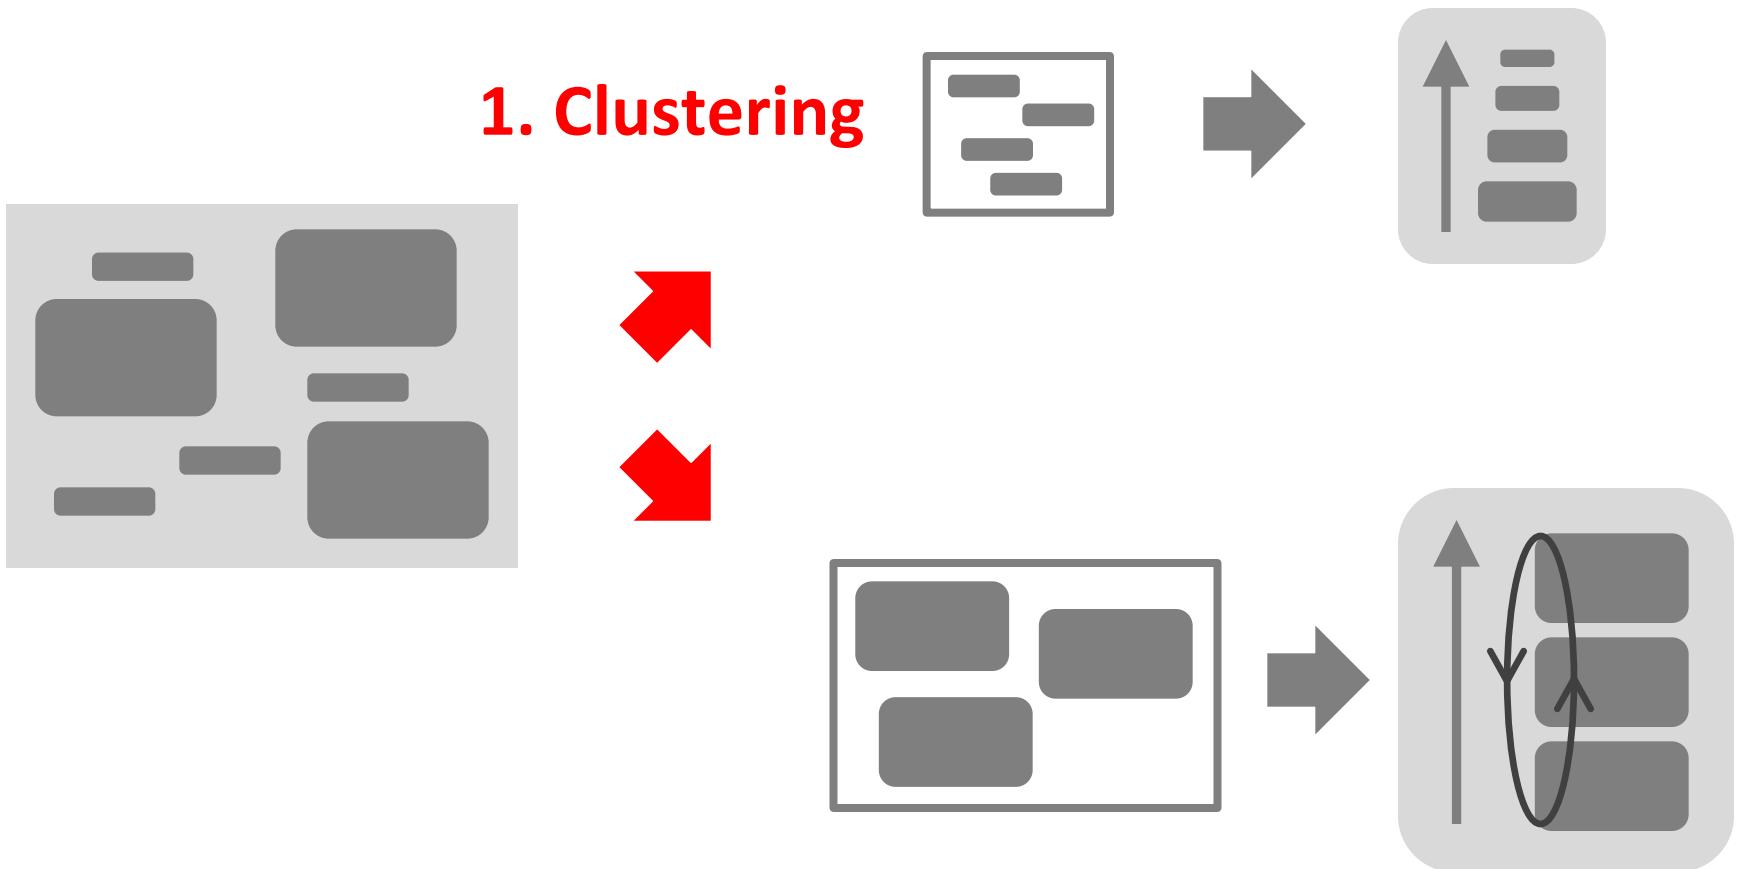


Thread Cluster Memory Scheduling [Kim+ MICRO'10]

1. Group threads into two *clusters*
2. Prioritize **non-intensive cluster**
3. Different policies for each cluster

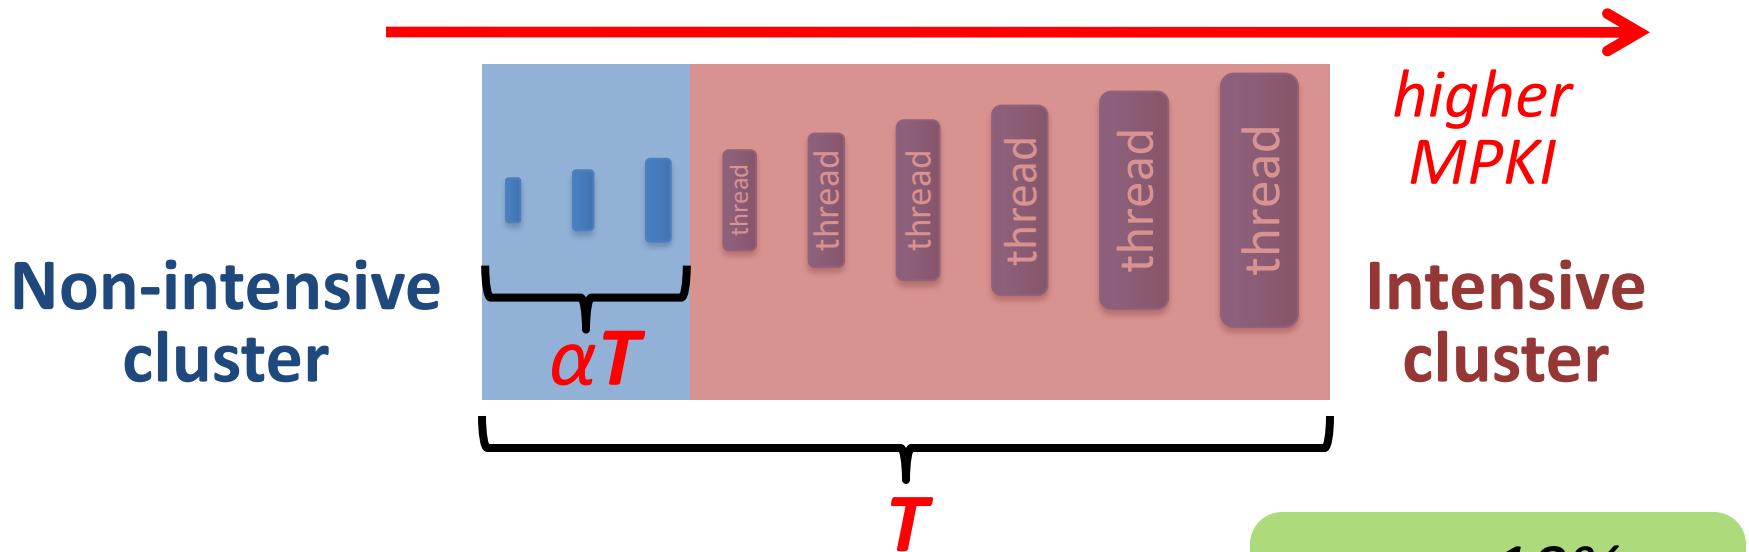


TCM Outline



Clustering Threads

Step1 Sort threads by **MPKI** (misses per kiloinstruction)

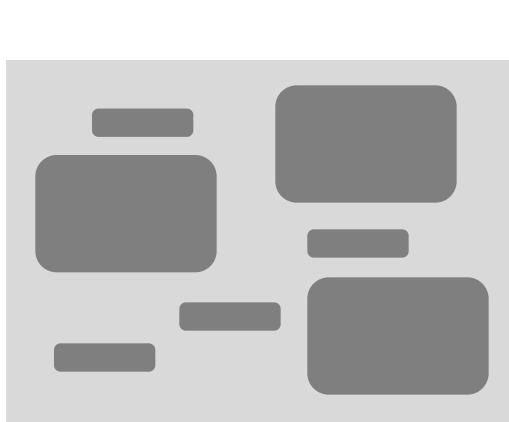


T = Total *memory bandwidth usage*

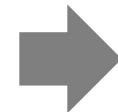
$\alpha < 10\%$
ClusterThreshold

Step2 Memory bandwidth usage αT divides clusters

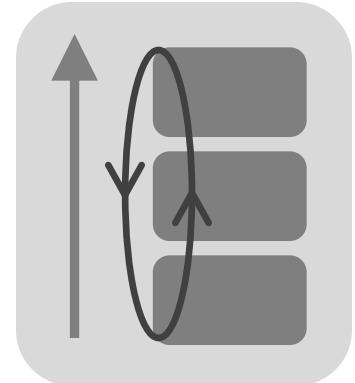
TCM Outline



1. Clustering

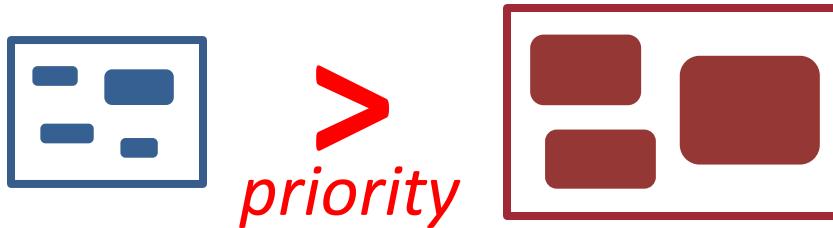


2. Between
Clusters



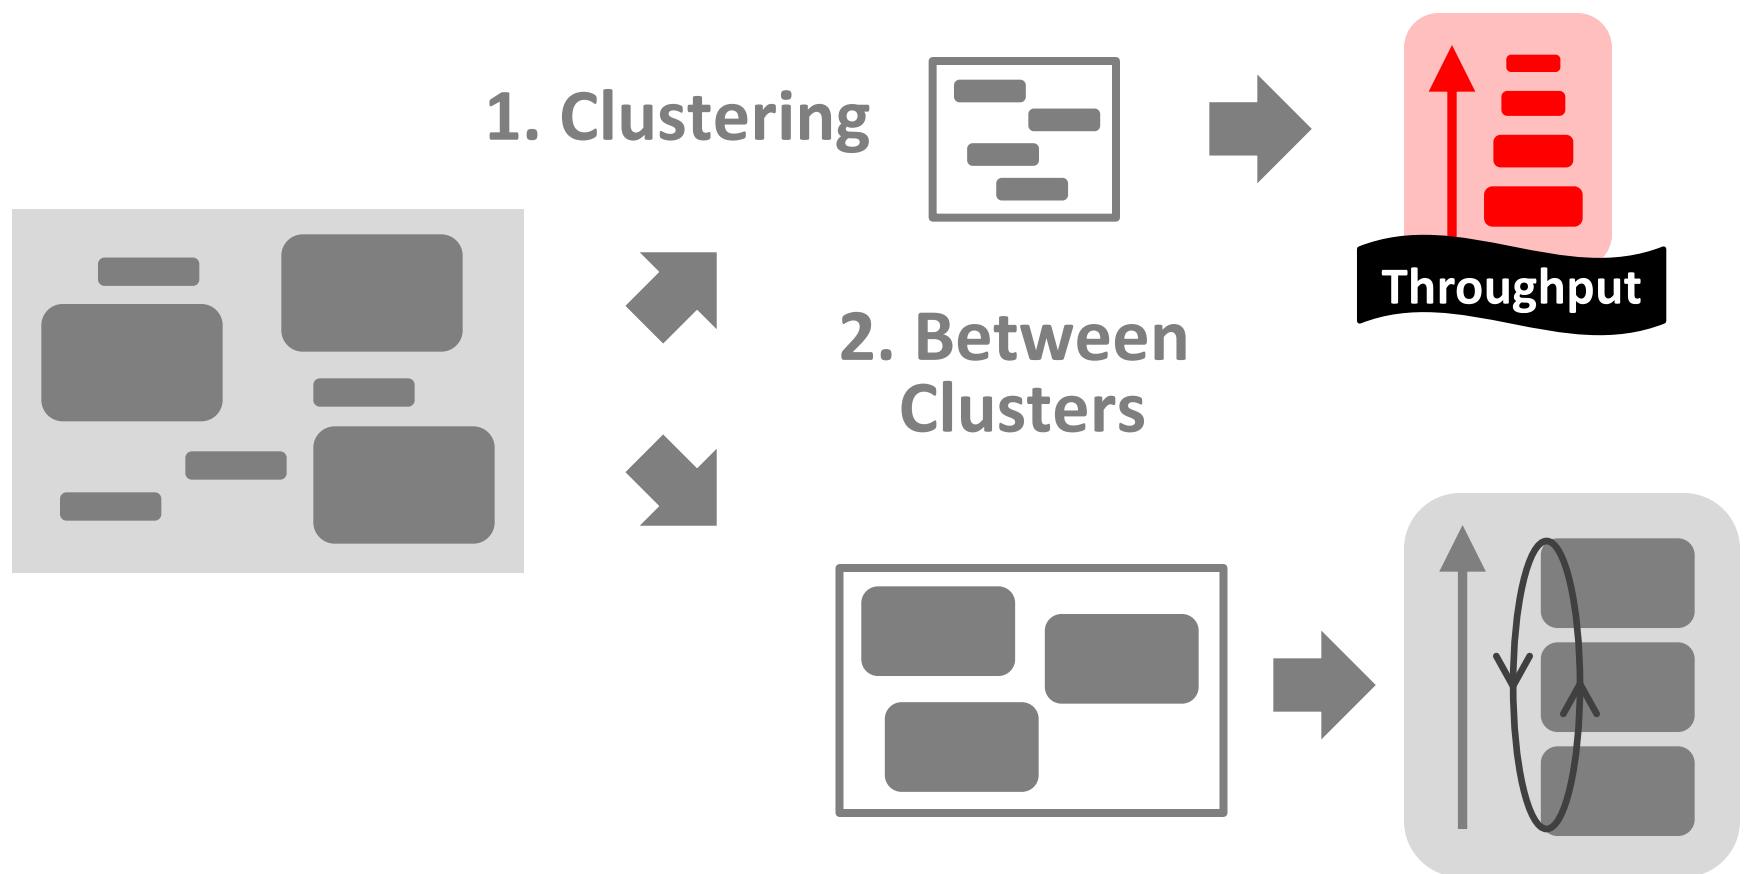
Prioritization Between Clusters

Prioritize non-intensive cluster



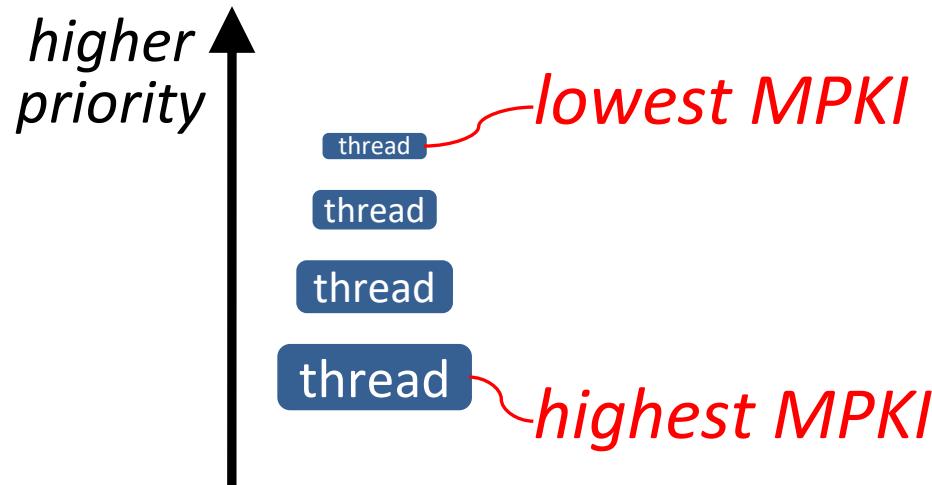
- **Increases system throughput**
 - Non-intensive threads have greater potential for making progress
- **Does not degrade fairness**
 - Non-intensive threads are “light”
 - Rarely interfere with intensive threads

TCM Outline



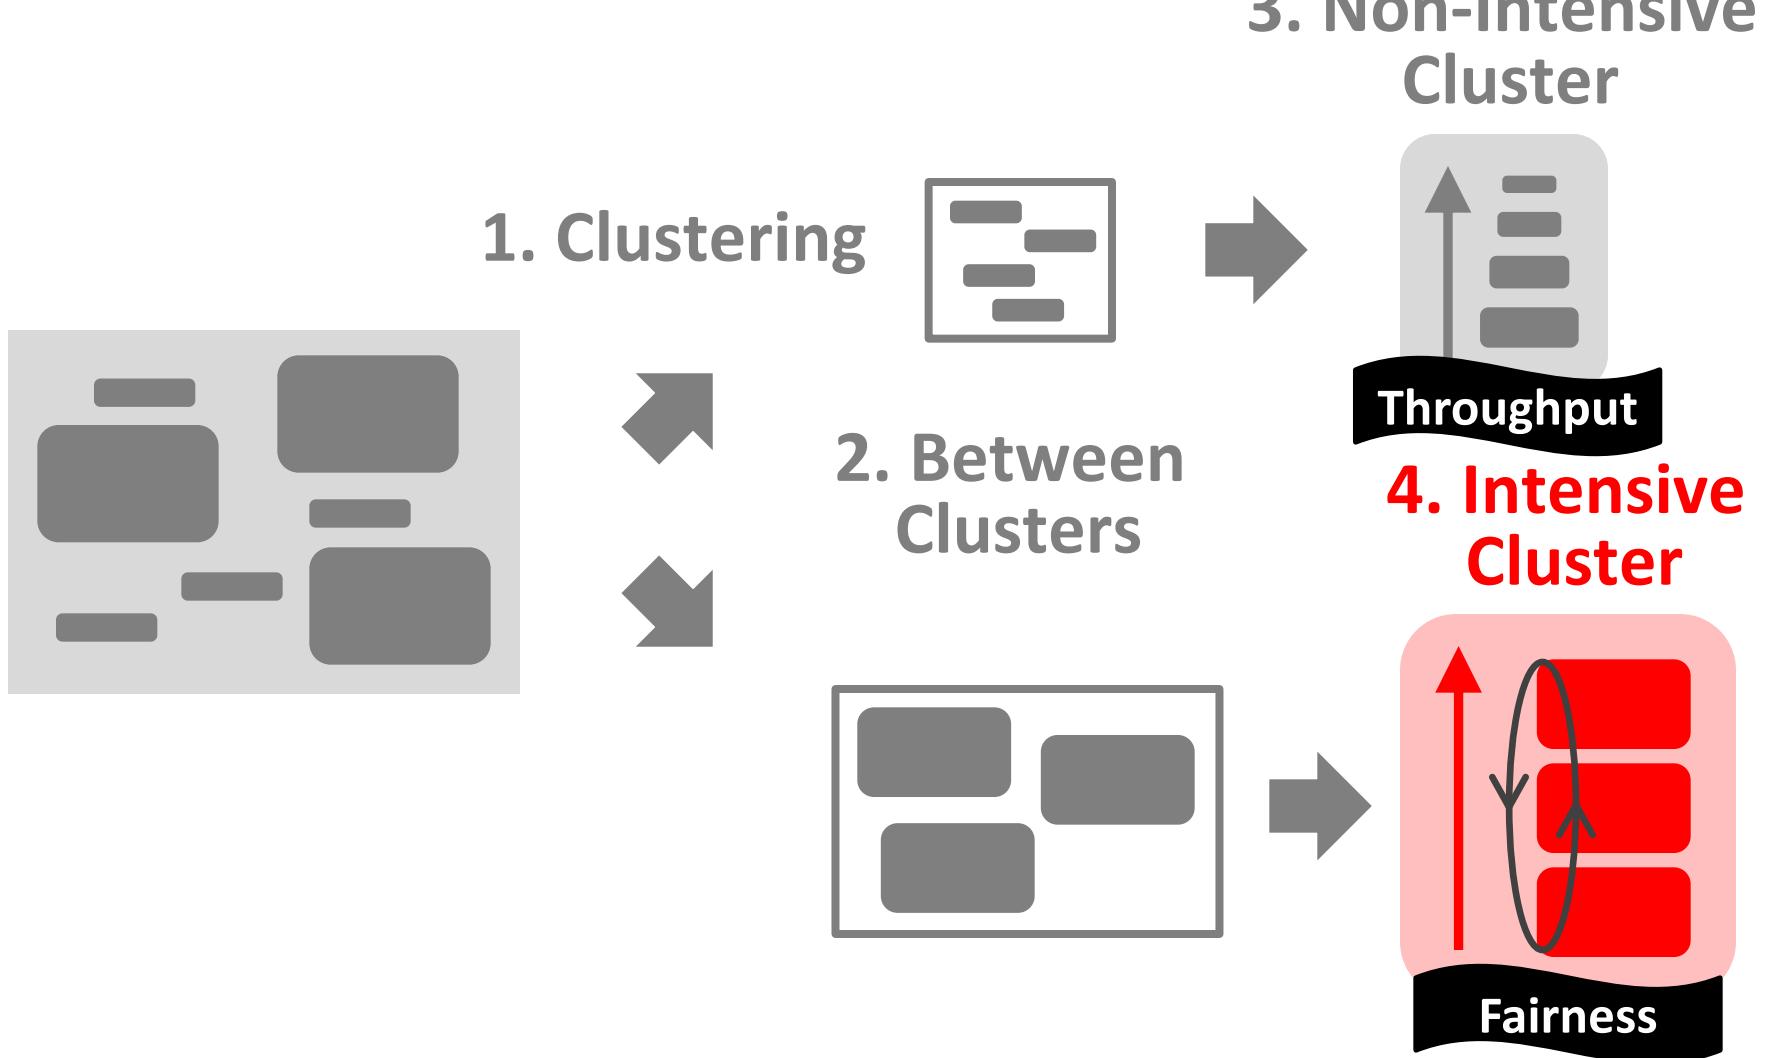
Non-Intensive Cluster

Prioritize threads according to MPKI



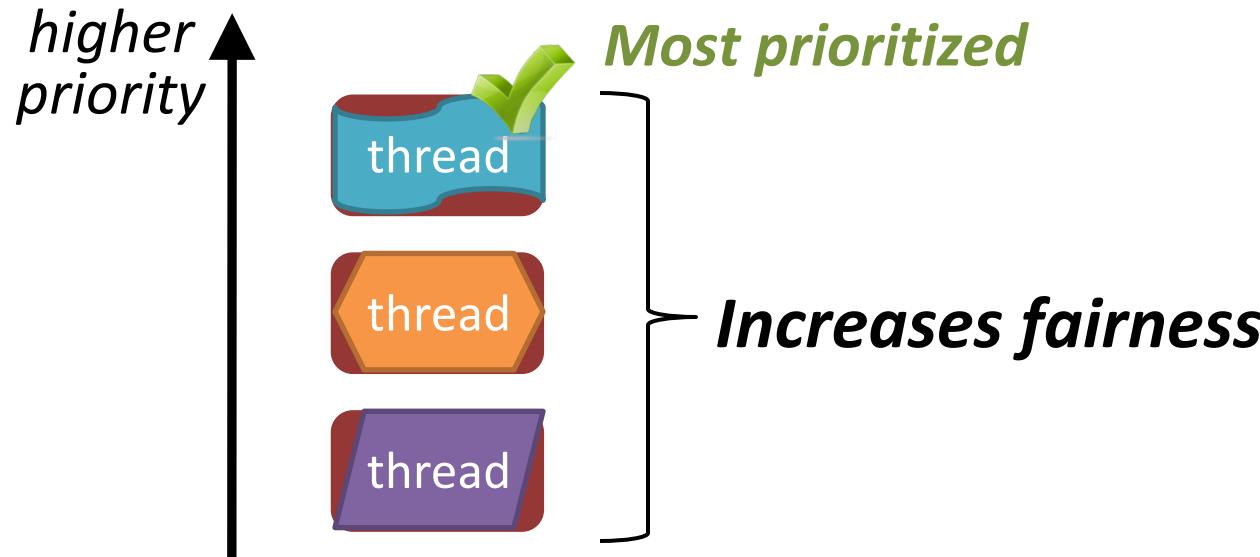
- **Increases system throughput**
 - Least intensive thread has the greatest potential for making progress in the processor

TCM Outline



Intensive Cluster

Periodically shuffle the priority of threads



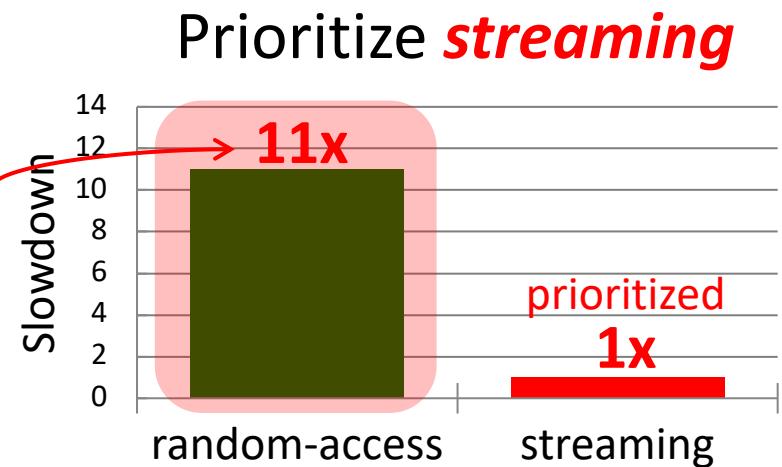
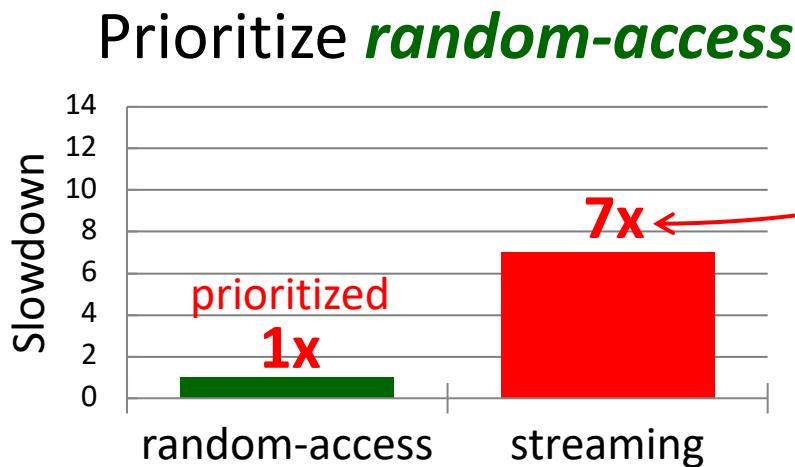
- Is treating all threads equally good enough?
- ***BUT: Equal turns ≠ Same slowdown***

Case Study: A Tale of Two Threads

Case Study: Two intensive threads contending

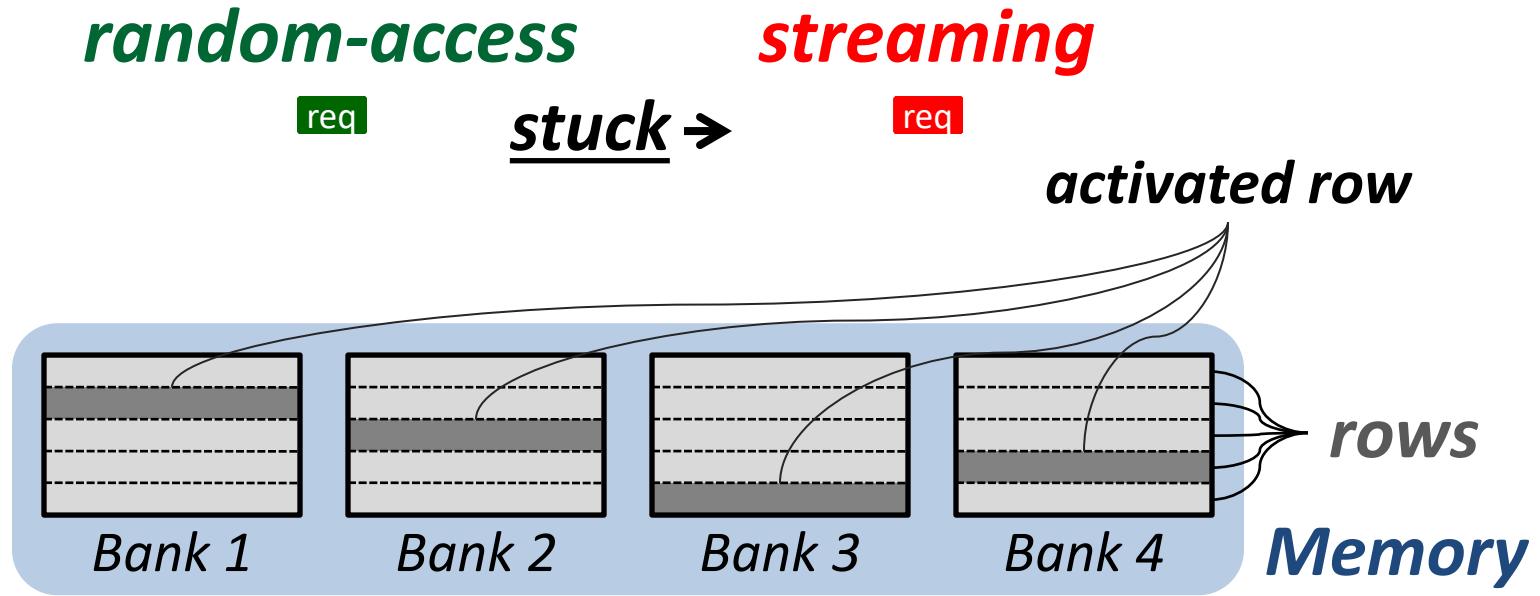
- 1. *random-access*
- 2. *streaming*

} *Which is slowed down more easily?*



random-access thread is more easily slowed down

Why are Threads Different?



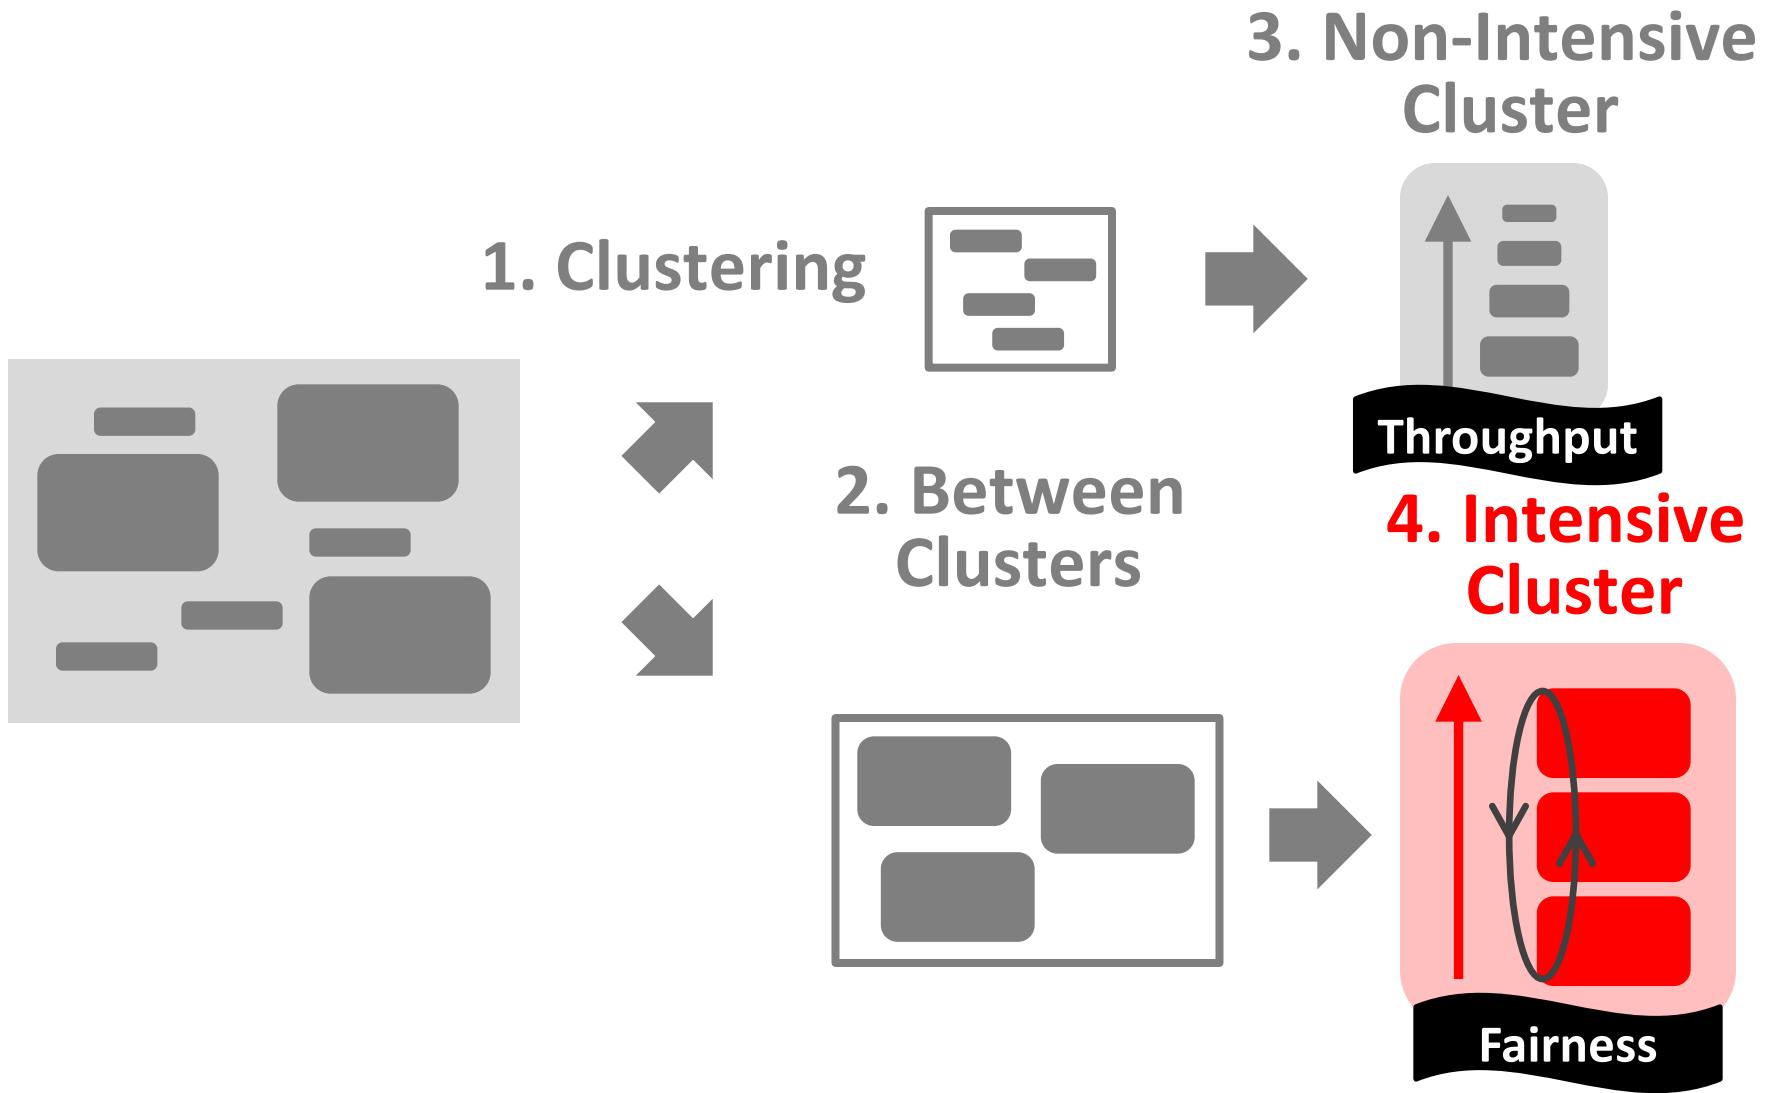
- All requests parallel
- High **bank-level parallelism**

- All requests → Same row
- High **row-buffer locality**



Vulnerable to interference

TCM Outline



Niceness

How to quantify difference between threads?

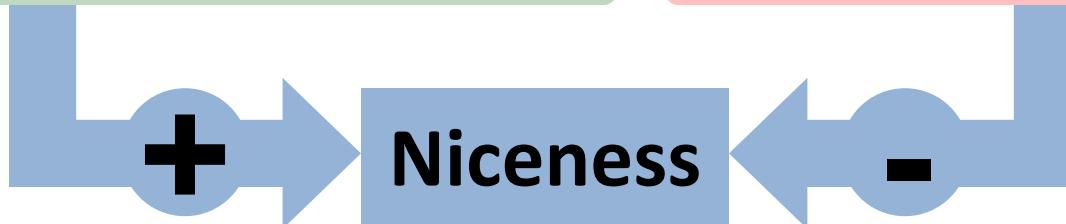


Bank-level parallelism

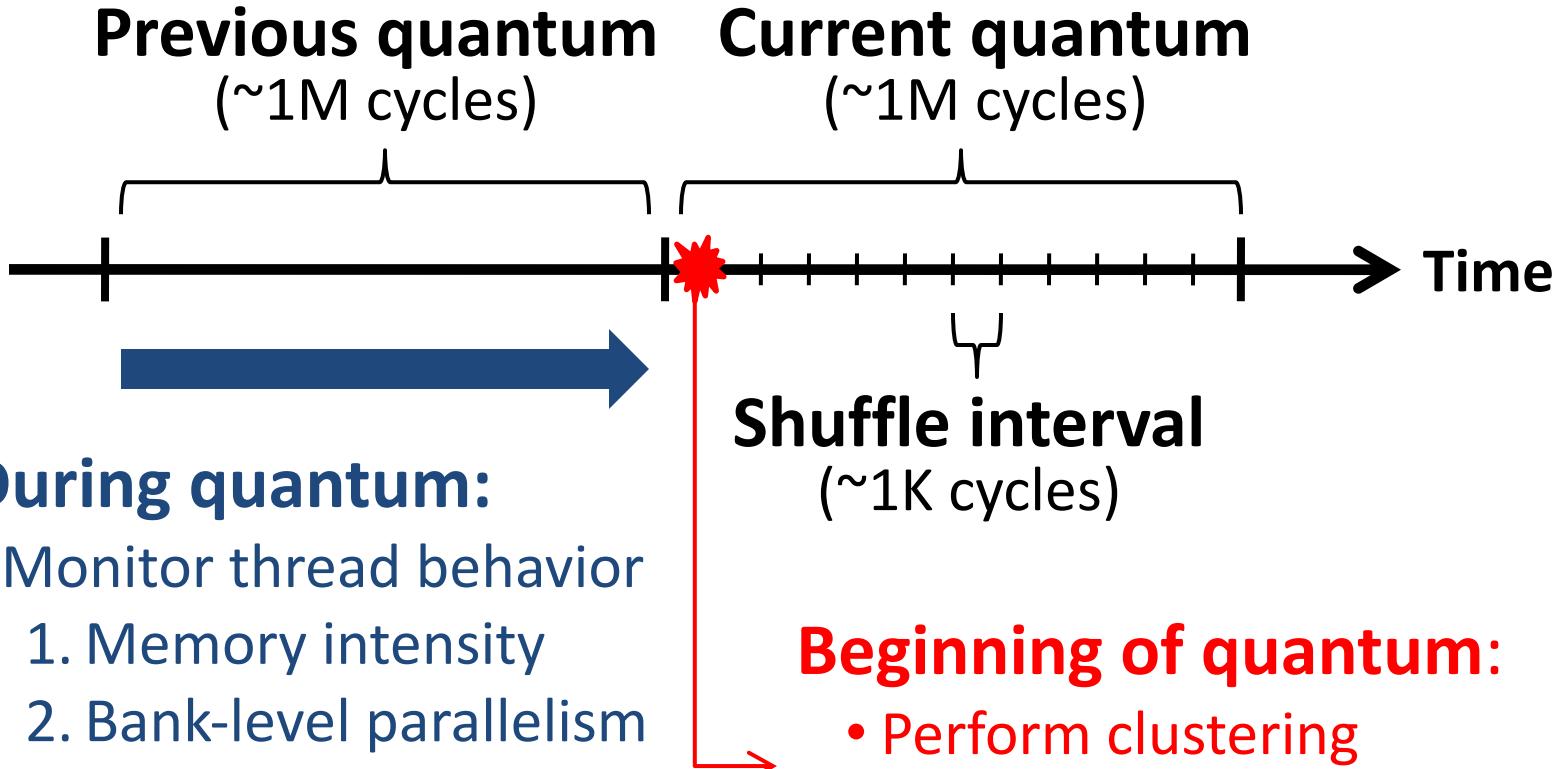
Vulnerability to interference

Row-buffer locality

Causes interference



TCM: Quantum-Based Operation



TCM: Scheduling Algorithm

1. *Highest-rank*: Requests from higher ranked threads prioritized

- Non-Intensive cluster > Intensive cluster
- Non-Intensive cluster: lower intensity → higher rank
- Intensive cluster: rank shuffling

2. *Row-hit*: Row-buffer hit requests are prioritized

3. *Oldest*: Older requests are prioritized

TCM: Implementation Cost

Required storage at memory controller (24 cores)

| Thread memory behavior | Storage |
|------------------------|----------|
| MPKI | ~0.2kb |
| Bank-level parallelism | ~0.6kb |
| Row-buffer locality | ~2.9kb |
| Total | < 4kbits |

- No computation is on the critical path

Previous Work

FRFCFS [Rixner et al., ISCA00]: Prioritizes row-buffer hits

- Thread-oblivious → Low throughput & Low fairness

STFM [Mutlu et al., MICRO07]: Equalizes thread slowdowns

- Non-intensive threads not prioritized → Low throughput

PAR-BS [Mutlu et al., ISCA08]: Prioritizes oldest batch of requests while preserving bank-level parallelism

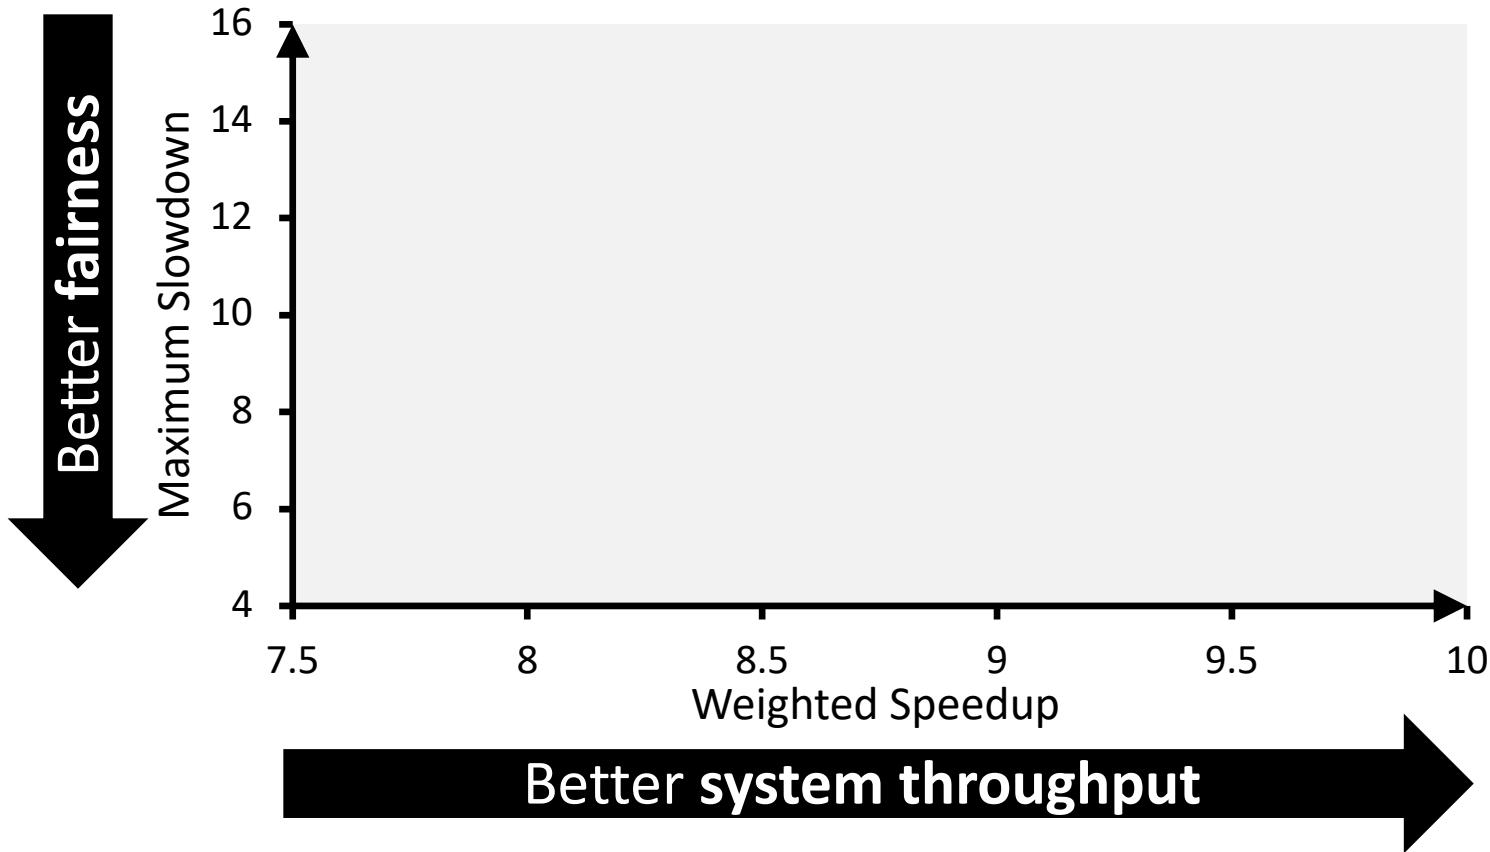
- Non-intensive threads not always prioritized → Low throughput

ATLAS [Kim et al., HPCA10]: Prioritizes threads with less memory service

- Most intensive thread starves → Low fairness

TCM: Throughput and Fairness

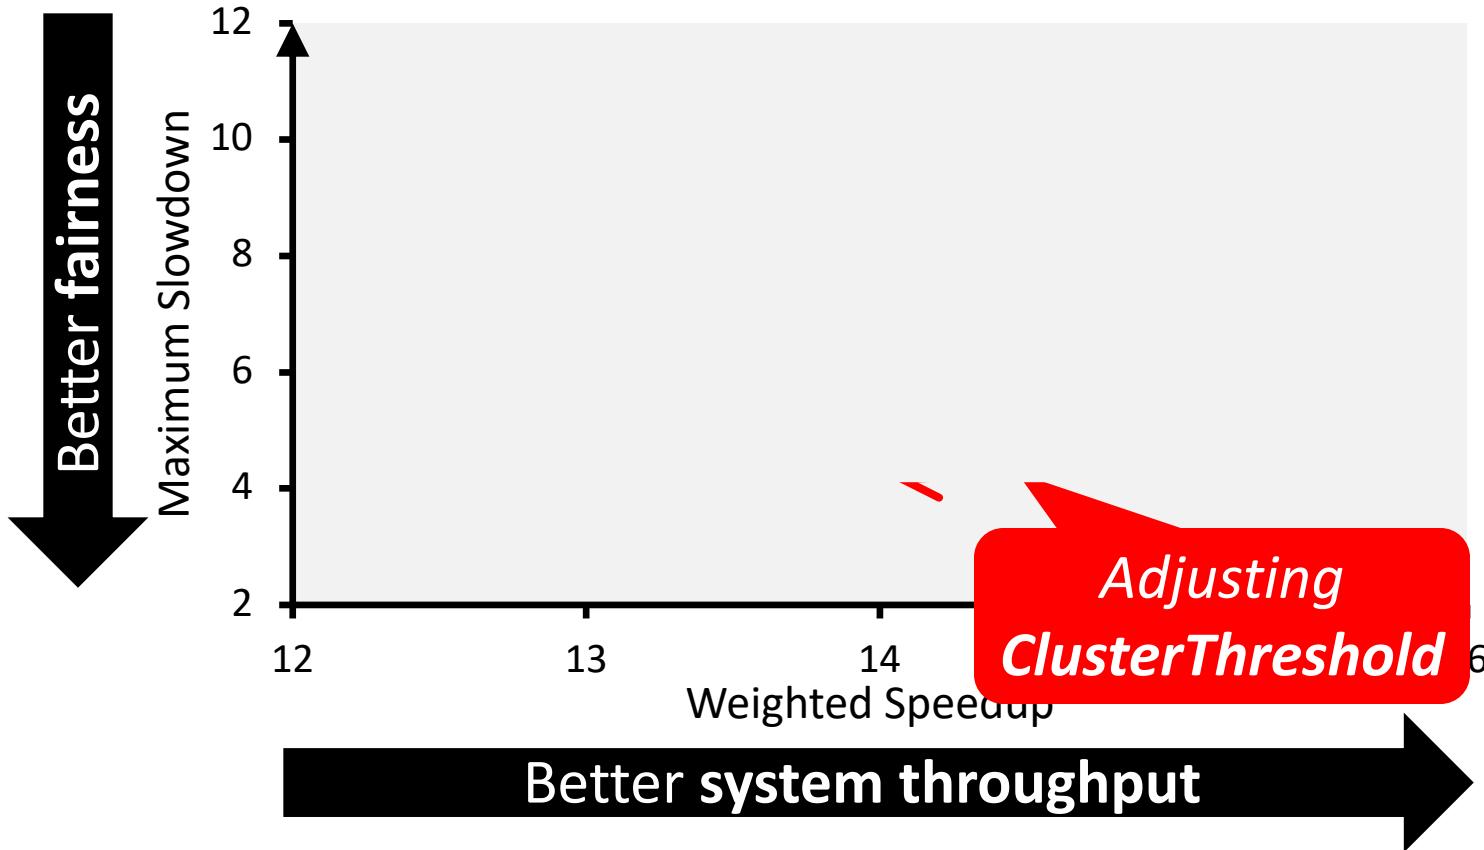
24 cores, 4 memory controllers, 96 workloads



*TCM, a heterogeneous scheduling policy,
provides best fairness and system throughput*

TCM: Fairness-Throughput Tradeoff

When configuration parameter is varied...



TCM allows robust fairness-throughput tradeoff

Operating System Support

- ***ClusterThreshold*** is a tunable knob
 - OS can trade off between fairness and throughput
- Enforcing thread weights
 - OS assigns weights to threads
 - TCM enforces thread weights within each cluster

Conclusion

- No previous memory scheduling algorithm provides both high *system throughput* and *fairness*
 - **Problem:** They use a single policy for all threads
- TCM groups threads into two *clusters*
 1. Prioritize *non-intensive* cluster → throughput
 2. Shuffle priorities in *intensive* cluster → fairness
 3. Shuffling should favor *nice* threads → fairness
- *TCM provides the best system throughput and fairness*

TCM Pros and Cons

- Upsides:
 - Provides both high fairness and high performance
 - Caters to the needs for different types of threads (latency vs. bandwidth sensitive)
 - (Relatively) simple

- Downsides:
 - Scalability to large buffer sizes?
 - Robustness of clustering and shuffling algorithms?
 - Ranking is still too complex?

More on TCM

- Yoongu Kim, Michael Papamichael, Onur Mutlu, and Mor Harchol-Balter,
**"Thread Cluster Memory Scheduling: Exploiting Differences in
Memory Access Behavior"**

*Proceedings of the 43rd International Symposium on
Microarchitecture (MICRO), pages 65-76, Atlanta, GA, December
2010.* Slides (pptx) (pdf)

***One of the 11 computer architecture papers of 2010 selected
as Top Picks by IEEE Micro.***

Thread Cluster Memory Scheduling: Exploiting Differences in Memory Access Behavior

Yoongu Kim

yoonguk@ece.cmu.edu

Michael Papamichael

papamix@cs.cmu.edu

Onur Mutlu

onur@cmu.edu

Mor Harchol-Balter

harchol@cs.cmu.edu

Carnegie Mellon University

More on TCM

- Yoongu Kim, Michael Papamichael, Onur Mutlu, and Mor Harchol-Balter,
"Thread Cluster Memory Scheduling"

IEEE Micro, Special Issue: *Micro's Top Picks from 2010 Computer Architecture Conferences (MICRO TOP PICKS)*, Vol. 31, No. 1, pages 78-89, January/February 2011.

THREAD CLUSTER MEMORY SCHEDULING

MEMORY SCHEDULERS IN MULTICORE SYSTEMS SHOULD CAREFULLY SCHEDULE MEMORY REQUESTS FROM DIFFERENT THREADS TO ENSURE HIGH SYSTEM PERFORMANCE AND FAIR, FAST PROGRESS OF EACH THREAD. NO EXISTING MEMORY SCHEDULER PROVIDES BOTH THE HIGHEST SYSTEM PERFORMANCE AND HIGHEST FAIRNESS. THREAD CLUSTER MEMORY SCHEDULING IS A NEW ALGORITHM THAT ACHIEVES THE BEST OF BOTH WORLDS BY DIFFERENTIATING LATENCY-SENSITIVE THREADS FROM BANDWIDTH-SENSITIVE ONES AND EMPLOYING DIFFERENT SCHEDULING POLICIES FOR EACH.

The Blacklisting Memory Scheduler

Lavanya Subramanian, Donghyuk Lee, Vivek Seshadri, Harsha Rastogi, and Onur Mutlu,

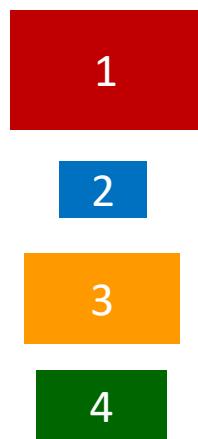
"The Blacklisting Memory Scheduler: Achieving High Performance and Fairness at Low Cost"

Proceedings of the 32nd IEEE International Conference on Computer Design (ICCD),

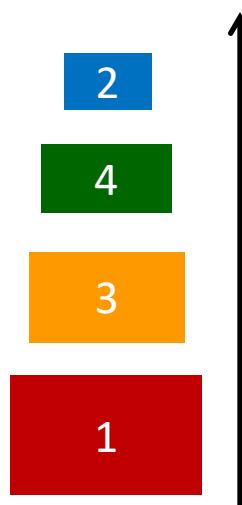
Seoul, South Korea, October 2014. [[Slides \(pptx\)](#) [\(pdf\)](#)]

Tackling Inter-Application Interference: Application-aware Memory Scheduling

Monitor



Rank



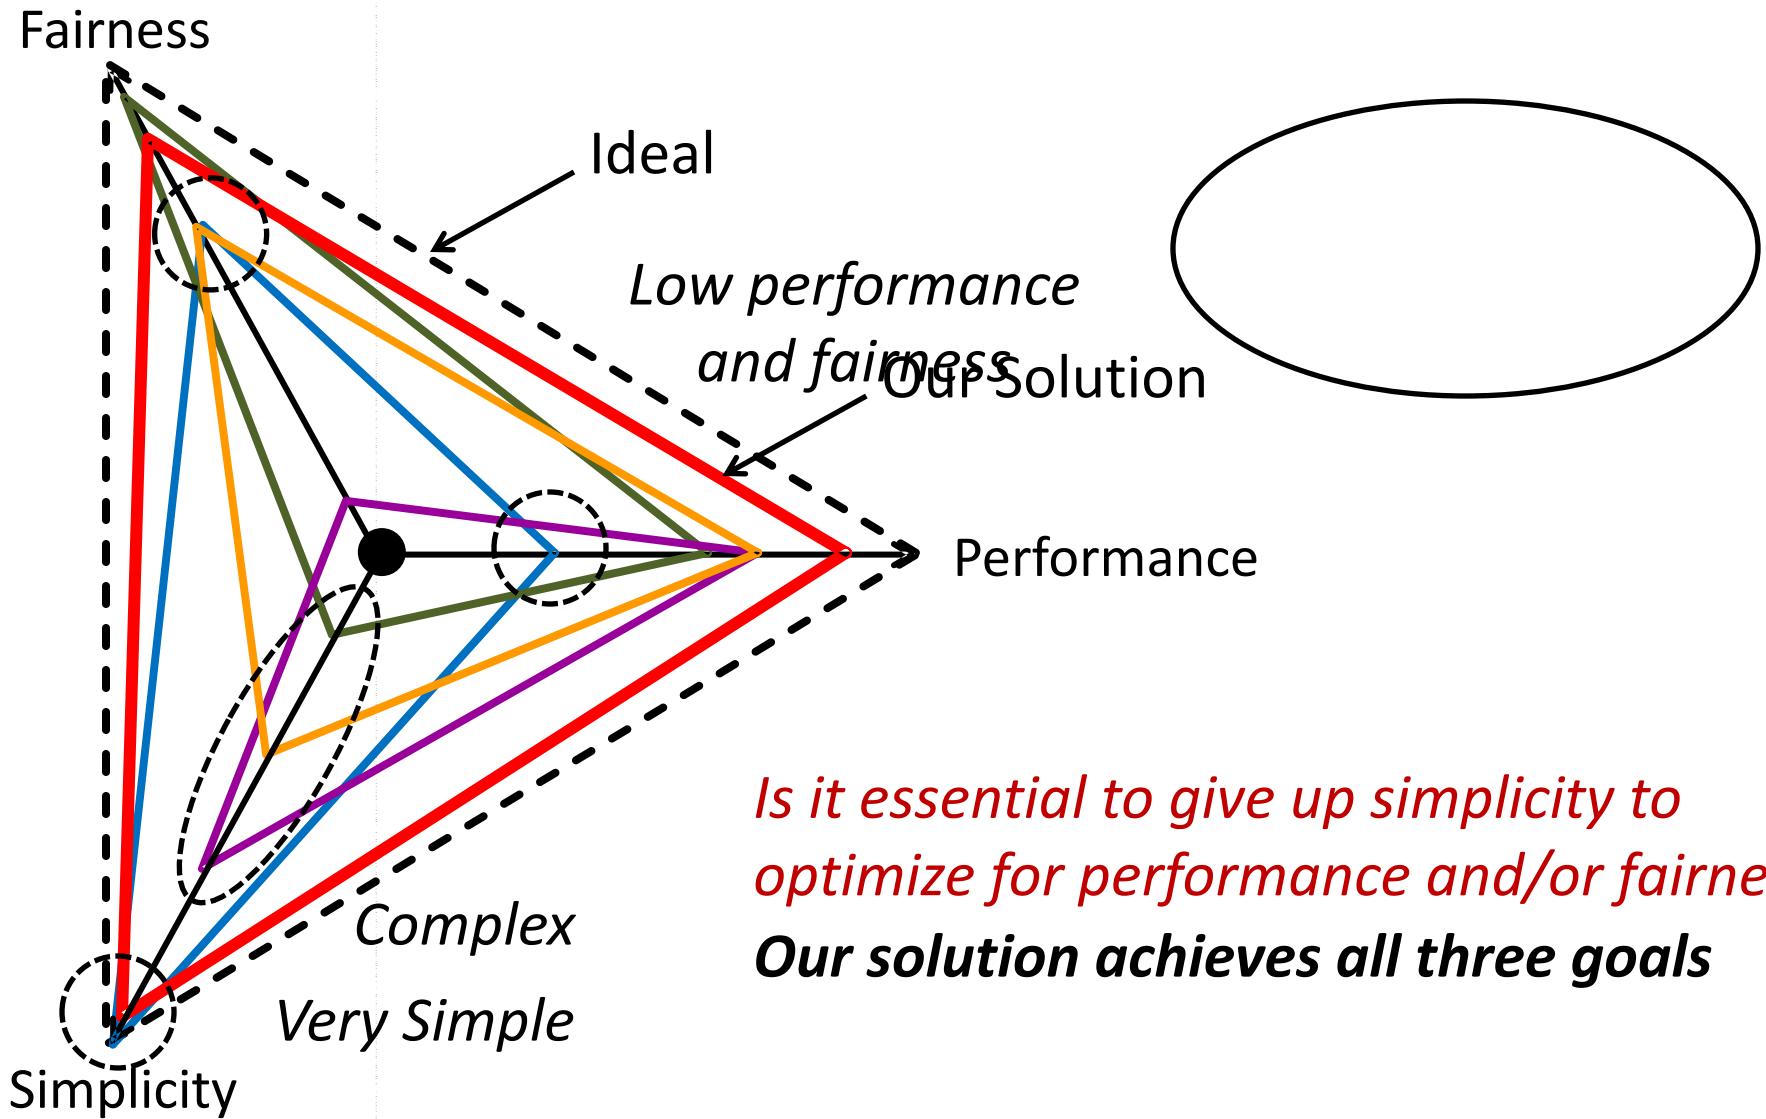
*Enforce
Ranks*

Request Buffer
App. ID
(AID)

| Request | AID | Highest Ranked AID |
|---------|-----|--------------------|
| Req 1 | 1 | = |
| Req 2 | 4 | = |
| Req 3 | 1 | = |
| Req 4 | 1 | = |
| Req 5 | 3 | = |
| Req 5 | 2 | = |
| Req 7 | 1 | = |
| Req 8 | 3 | = |

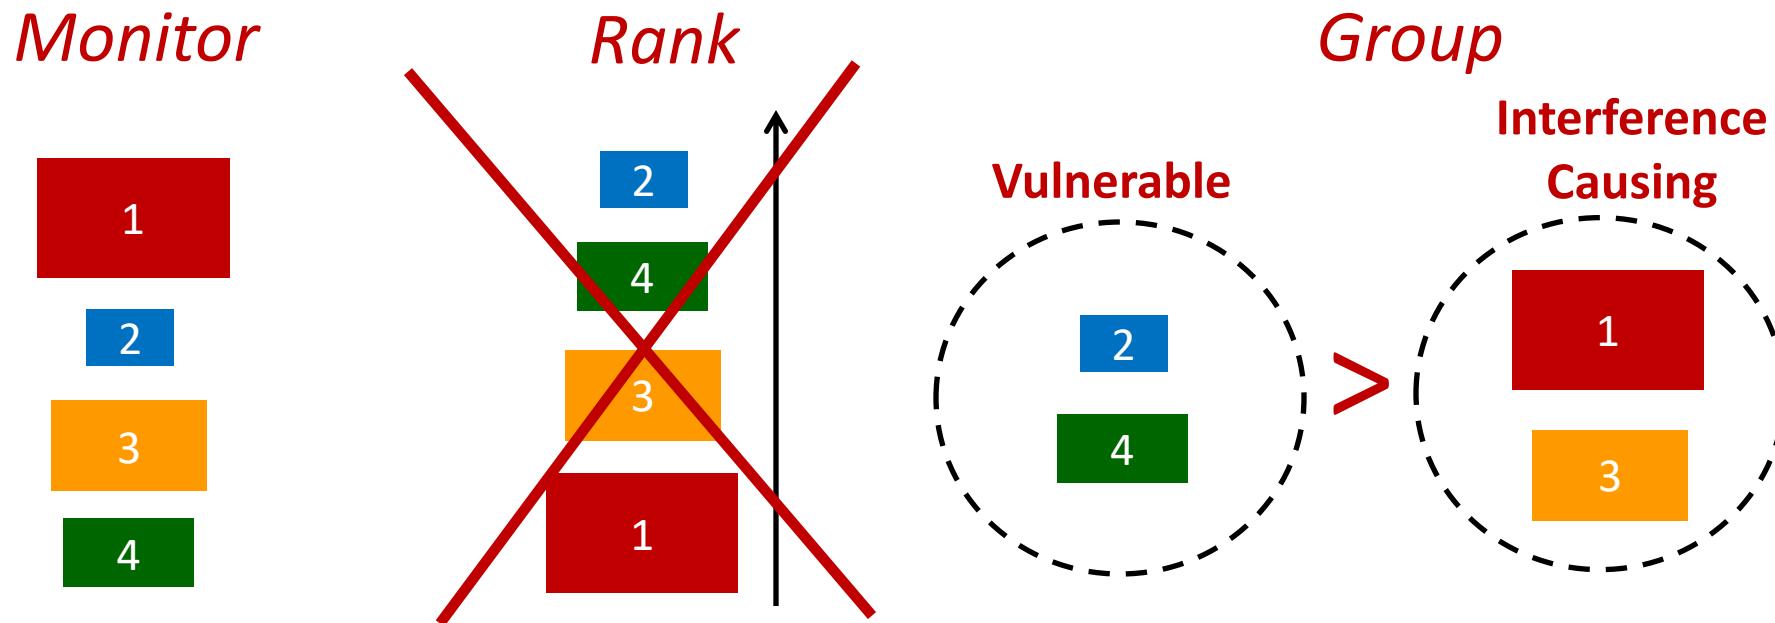
*Full ranking increases
critical path latency and area
significantly to improve
performance and fairness*

Performance vs. Fairness vs. Simplicity



Key Observation 1: Group Rather Than Rank

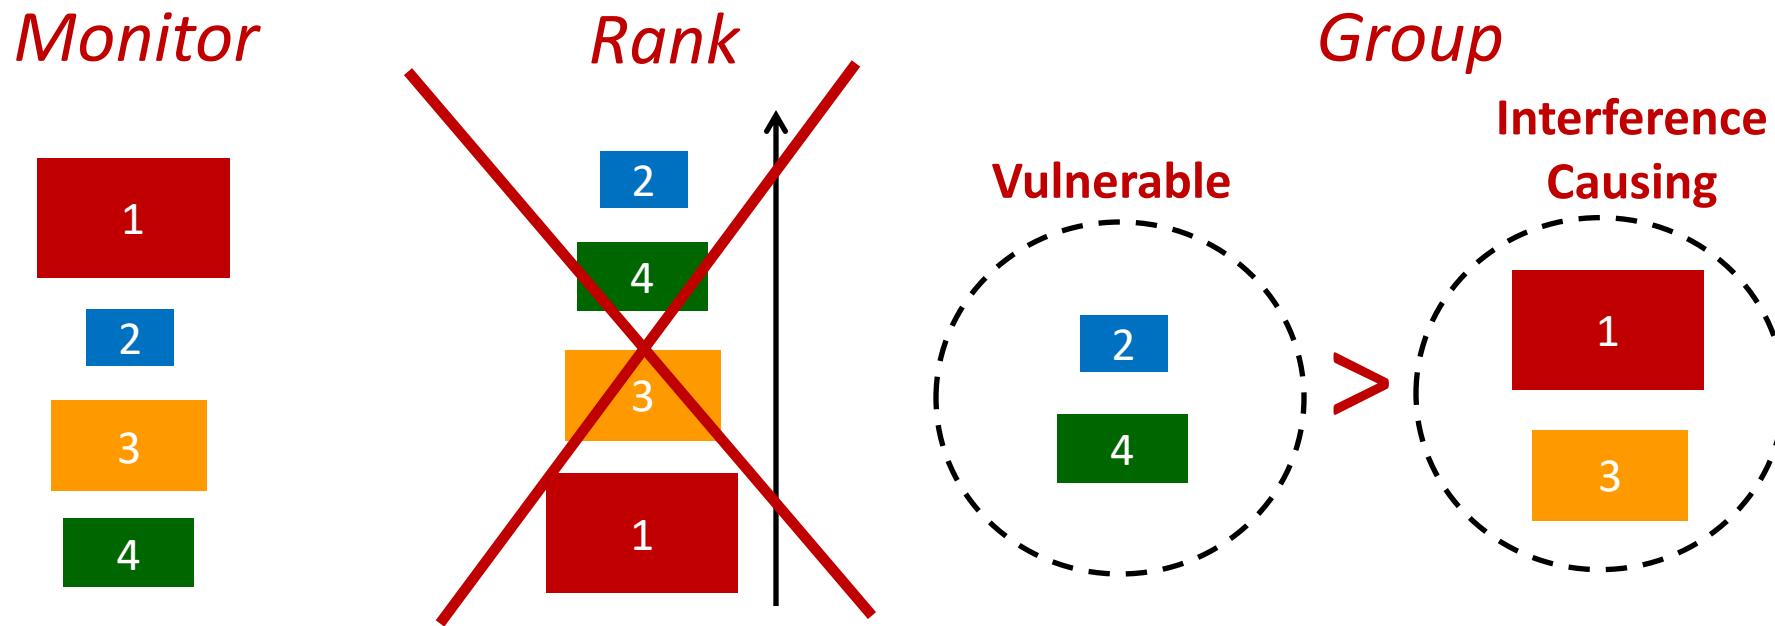
Observation 1: Sufficient to separate applications into two groups, rather than do full ranking



Benefit 2: Lower slowdowns than ranking

Key Observation 1: Group Rather Than Rank

Observation 1: Sufficient to separate applications into two groups, rather than do full ranking



How to classify applications into groups?

Key Observation 2

Observation 2: Serving a large number of consecutive requests from an application causes interference

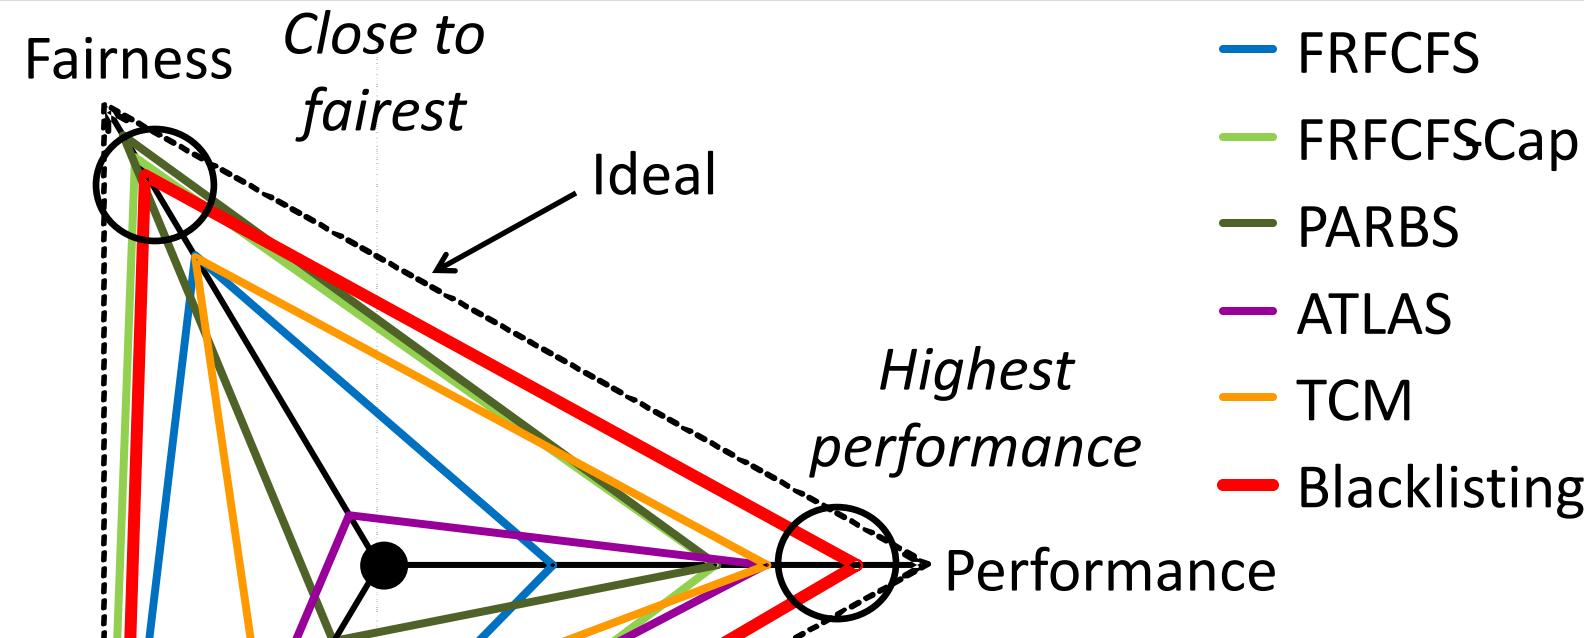
Basic Idea:

- *Group* applications with a large number of consecutive requests as *interference-causing* → *Blacklisting*
- *Deprioritize* blacklisted applications
- *Clear* blacklist periodically (1000s of cycles)

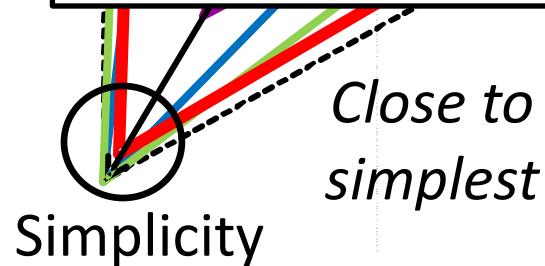
Benefits:

- *Lower complexity*
- *Finer grained grouping decisions* → *Lower unfairness*

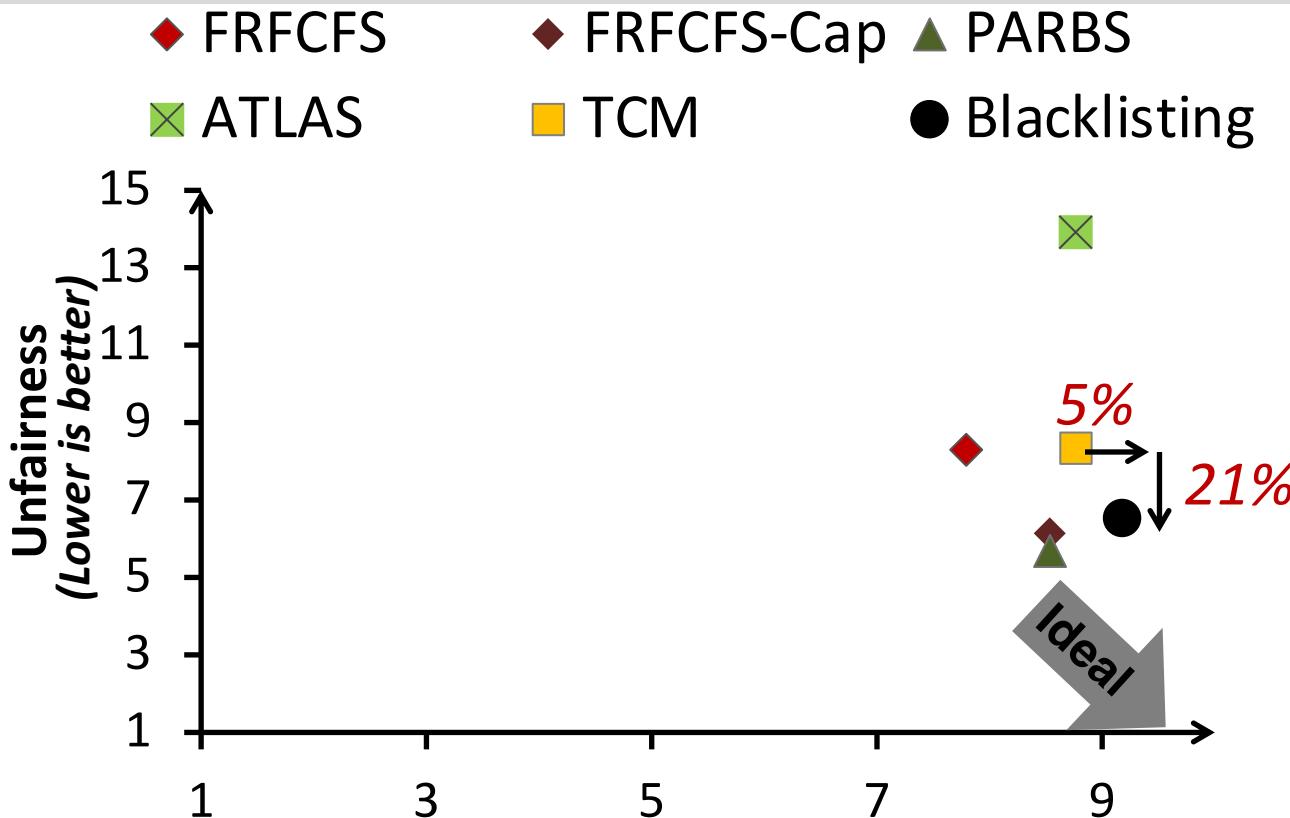
Performance vs. Fairness vs. Simplicity



Blacklisting is the closest scheduler to ideal

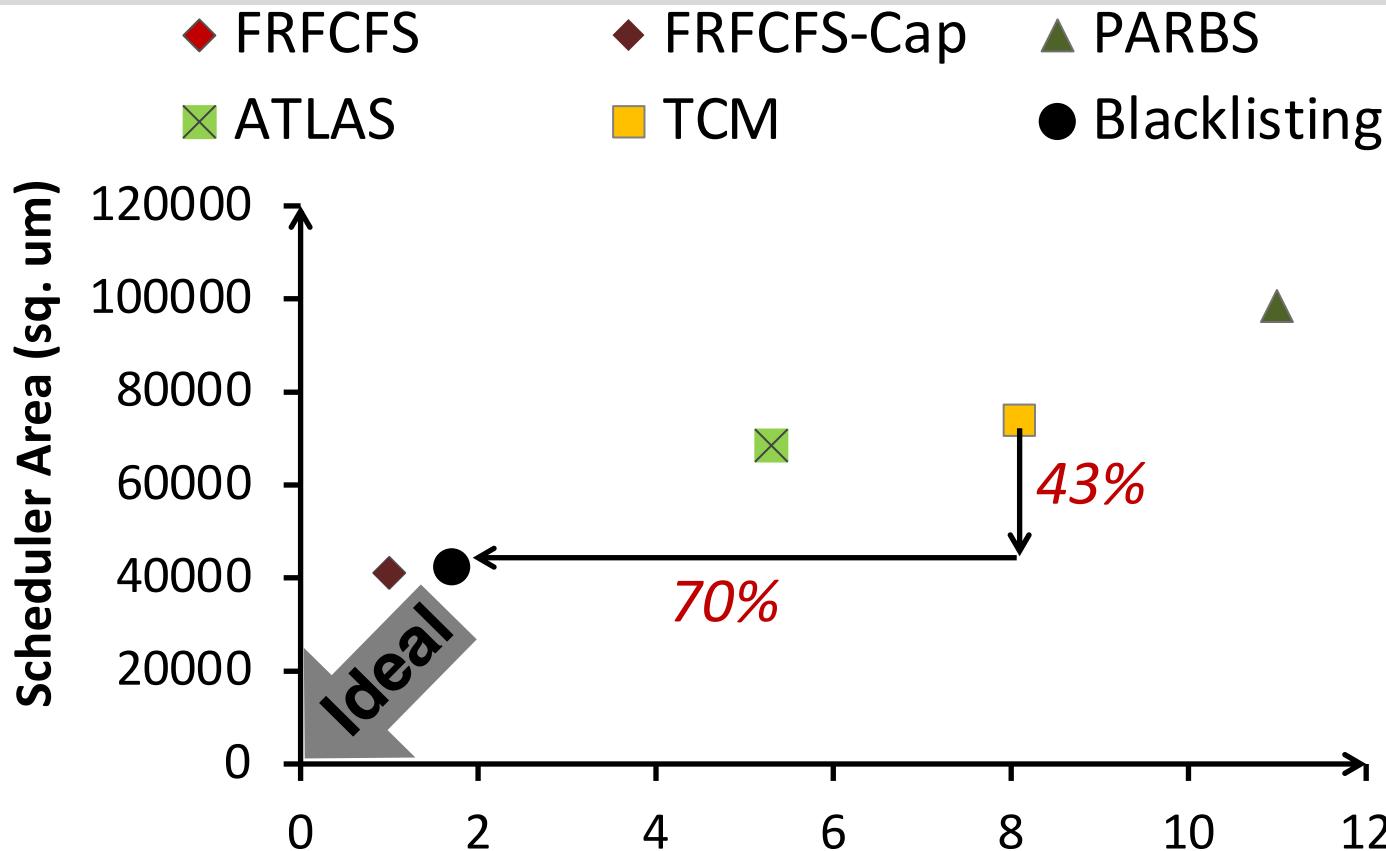


Performance and Fairness



1. *Blacklisting achieves the highest performance*
2. *Blacklisting balances performance and fairness*

Complexity



Blacklisting reduces complexity significantly

More on BLISS (I)

- Lavanya Subramanian, Donghyuk Lee, Vivek Seshadri, Harsha Rastogi, and Onur Mutlu,

"The Blacklisting Memory Scheduler: Achieving High Performance and Fairness at Low Cost"

Proceedings of the 32nd IEEE International Conference on Computer Design (ICCD), Seoul, South Korea, October 2014.
[Slides (pptx) (pdf)]

The Blacklisting Memory Scheduler: Achieving High Performance and Fairness at Low Cost

Lavanya Subramanian, Donghyuk Lee, Vivek Seshadri, Harsha Rastogi, Onur Mutlu
Carnegie Mellon University
`{lsubrama,donghyu1,visesh,harshar,onur}@cmu.edu`

More on BLISS: Longer Version

- Lavanya Subramanian, Donghyuk Lee, Vivek Seshadri, Harsha Rastogi, and Onur Mutlu,

"BLISS: Balancing Performance, Fairness and Complexity in Memory Access Scheduling"

IEEE Transactions on Parallel and Distributed Systems (TPDS), to appear in 2016. [arXiv.org version](#), April 2015.

An earlier version as *SAFARI Technical Report*, TR-SAFARI-2015-004, Carnegie Mellon University, March 2015.

[\[Source Code\]](#)

BLISS: Balancing Performance, Fairness and Complexity in Memory Access Scheduling

Lavanya Subramanian, Donghyuk Lee, Vivek Seshadri, Harsha Rastogi, and Onur Mutlu

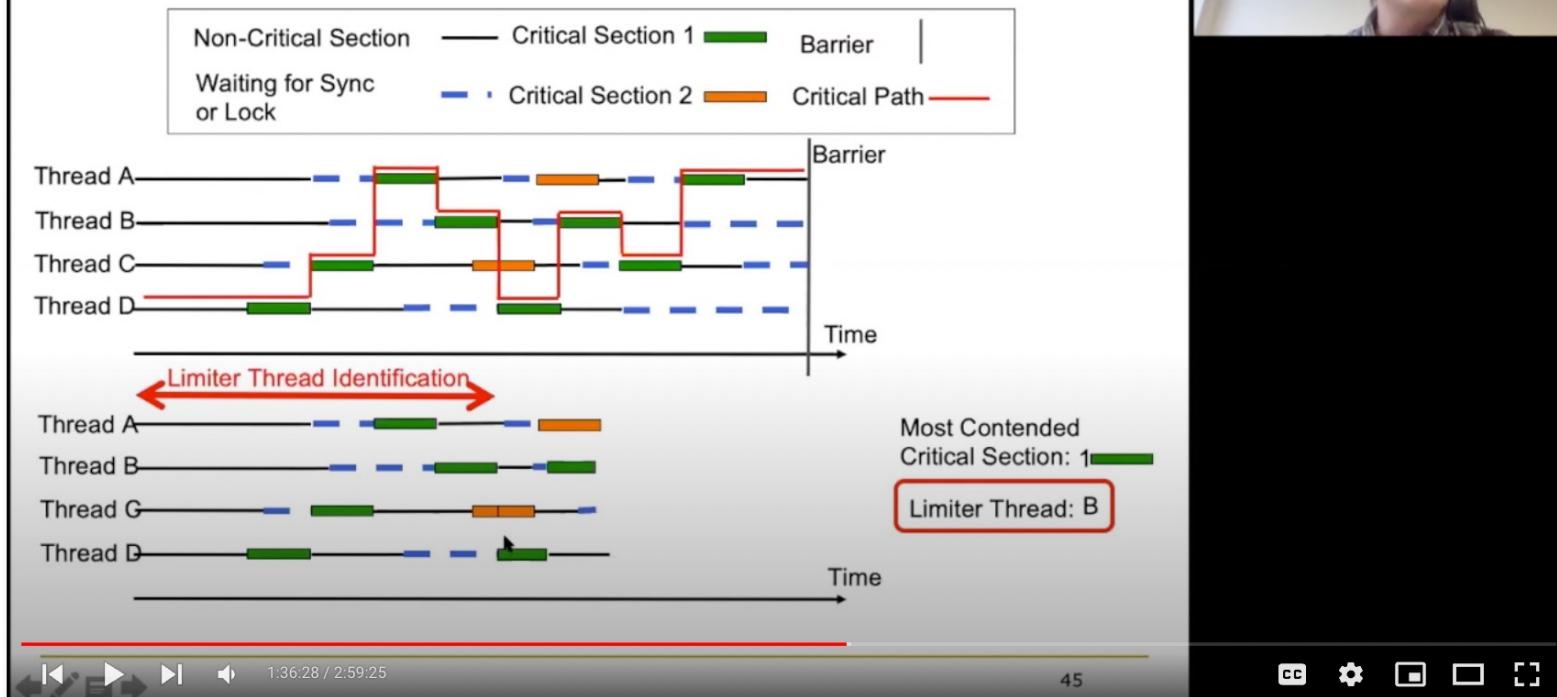
Handling Memory Interference In Multithreaded Applications

Eiman Ebrahimi, Rustam Miftakhutdinov, Chris Fallin,
Chang Joo Lee, Onur Mutlu, and Yale N. Patt,
"Parallel Application Memory Scheduling"

*Proceedings of the 44th International Symposium on Microarchitecture (**MICRO**),
Porto Alegre, Brazil, December 2011.* [Slides \(pptx\)](#)

Lecture on Parallel Application Scheduling

Prioritizing Requests from Limiter Threads



ETH ZENTRUM

Computer Arch - Lecture 13: Memory Interference and Quality of Service II (ETH Zürich, Spring 2020)

1,006 views • Nov 7, 2020

23 likes 0 dislikes SHARE SAVE ...



Onur Mutlu Lectures
19.8K subscribers

ANALYTICS

EDIT VIDEO

Lecture on Bottleneck Acceleration

Observation: Limiting Bottlenecks Change Over Time

A=full linked list; B=empty linked list
repeat

Lock A

Traverse list A

Remove X from A

Unlock A

Compute on X

Lock B

Traverse list B

Insert X into B

Unlock B

32 threads

Contention (# of threads waiting)

time [Mcycles]

until A is empty

55:08 / 2:30:02

◀ ▶ ⏪ 🔍 ⚙️

Computer Architecture - Lecture 17: Bottleneck Acceleration (ETH Zürich, Fall 2020)

880 views • Nov 20, 2020

13 0 SHARE SAVE ...



Onur Mutlu Lectures
19.8K subscribers

ANALYTICS

EDIT VIDEO

Multithreaded (Parallel) Applications

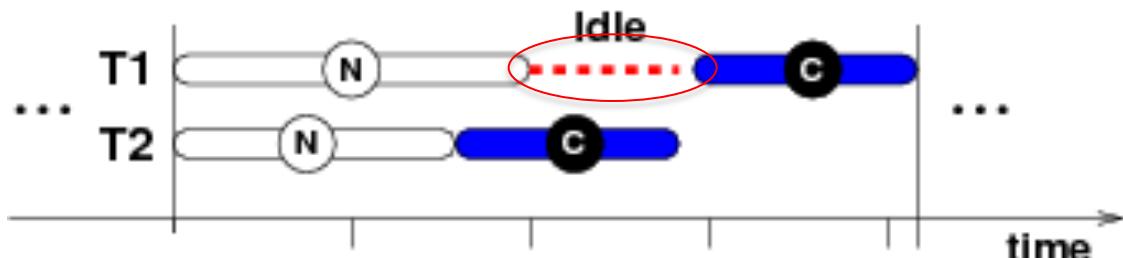
- Threads in a multi-threaded application can be inter-dependent
 - As opposed to threads from different applications
- Such threads can synchronize with each other
 - Locks, barriers, pipeline stages, condition variables, semaphores, ...
- Some threads can be on the critical path of execution due to synchronization; some threads are not
- Even within a thread, some “code segments” may be on the critical path of execution; some are not

Critical Sections

- Enforce mutually exclusive access to shared data
- Only one thread can be executing it at a time
- Contended critical sections make threads wait → threads causing serialization can be on the critical path

Each thread:

```
loop {  
    Compute  
    lock(A)  
    Update shared data  
    unlock(A)  
}
```

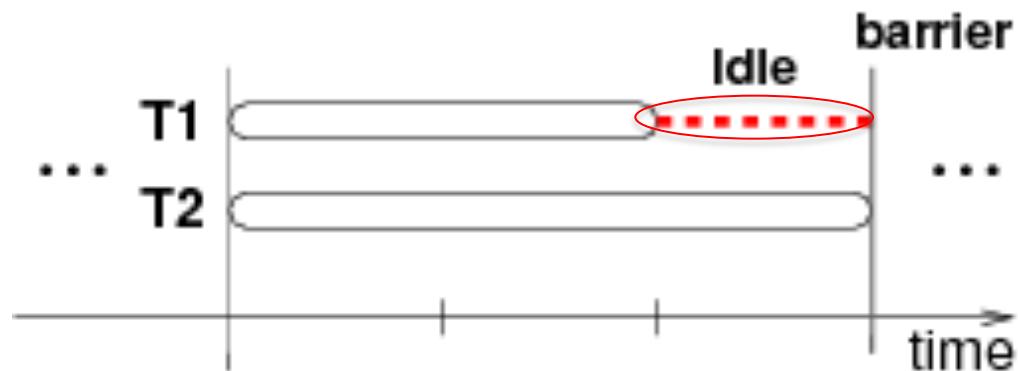


Barriers

- Synchronization point
- Threads have to wait until all threads reach the barrier
- Last thread arriving at the barrier is on the critical path

Each thread:

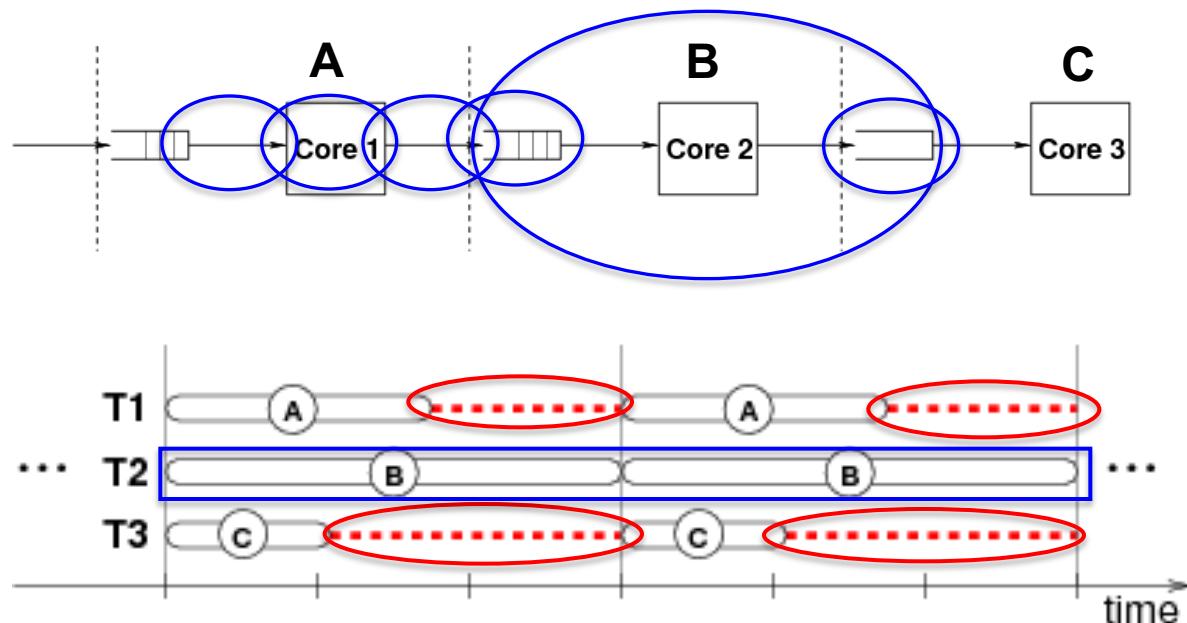
```
loop1 {  
    Compute  
}  
barrier  
loop2 {  
    Compute  
}
```



Stages of Pipelined Programs

- Loop iterations are statically divided into code segments called *stages*
- Threads execute stages on different cores
- Thread executing the slowest stage is on the critical path

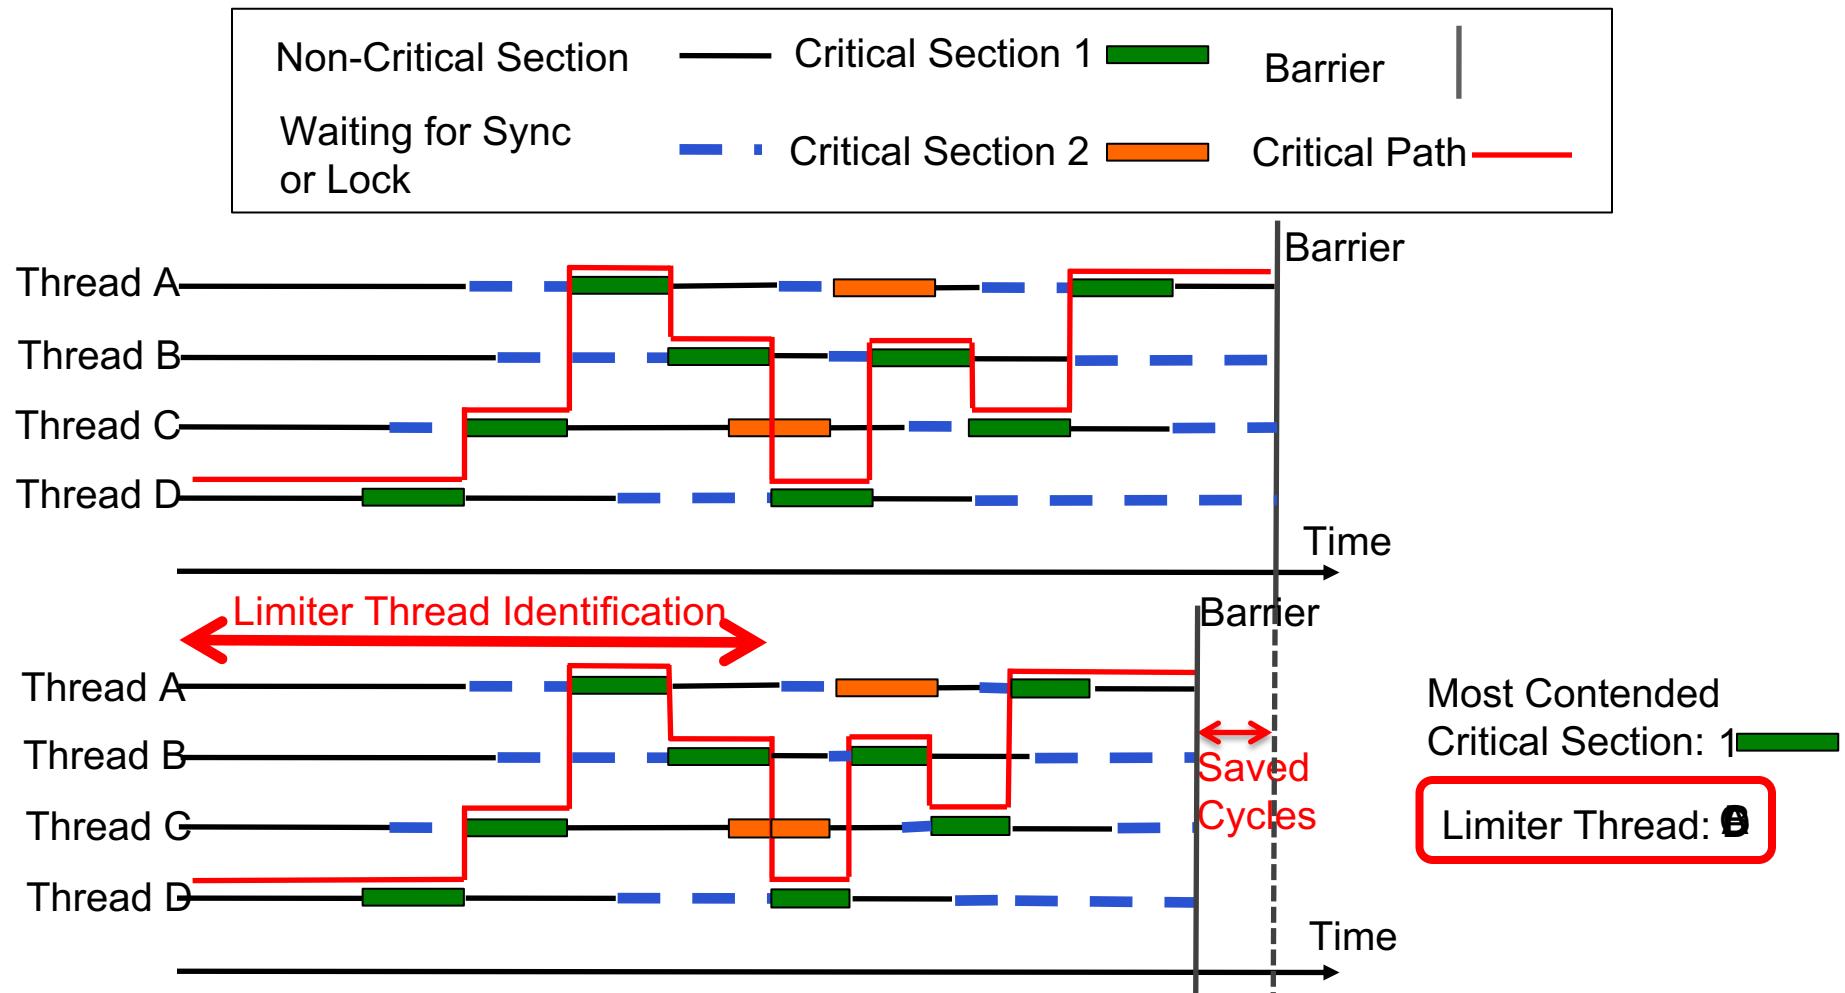
```
loop {  
    Compute1 A  
    Compute2 B  
    Compute3 C  
}
```



Handling Interference in Parallel Applications

- Threads in a multithreaded application are inter-dependent
 - Some threads can be on the critical path of execution due to synchronization; some threads are not
 - How do we schedule requests of inter-dependent threads to maximize multithreaded application performance?
-
- Idea: **Estimate limiter threads** likely to be on the critical path and prioritize their requests; **shuffle priorities of non-limiter threads** to reduce memory interference among them [Ebrahimi+, MICRO'11]
 - Hardware/software cooperative limiter thread estimation:
 - Thread executing the most contended critical section
 - Thread executing the slowest pipeline stage
 - Thread that is falling behind the most in reaching a barrier

Prioritizing Requests from Limiter Threads



Parallel App Mem Scheduling: Pros and Cons

- Upsides:
 - Improves the performance of multi-threaded applications
 - Provides a mechanism for estimating “limiter threads”
 - Opens a path for slowdown estimation for multi-threaded applications
- Downsides:
 - What if there are multiple multi-threaded applications running together?
 - Limiter thread estimation can be complex
 - “Joao+, Bottleneck Identification and Scheduling, ASPLOS 2012”

More on PAMS

- Eiman Ebrahimi, Rustam Miftakhutdinov, Chris Fallin, Chang Joo Lee, Onur Mutlu, and Yale N. Patt,

"Parallel Application Memory Scheduling"

Proceedings of the 44th International Symposium on Microarchitecture (MICRO), Porto Alegre, Brazil, December 2011. Slides (pptx)

Parallel Application Memory Scheduling

Eiman Ebrahimi[†] Rustam Miftakhutdinov[†] Chris Fallin[§]
Chang Joo Lee[‡] José A. Joao[†] Onur Mutlu[§] Yale N. Patt[‡]

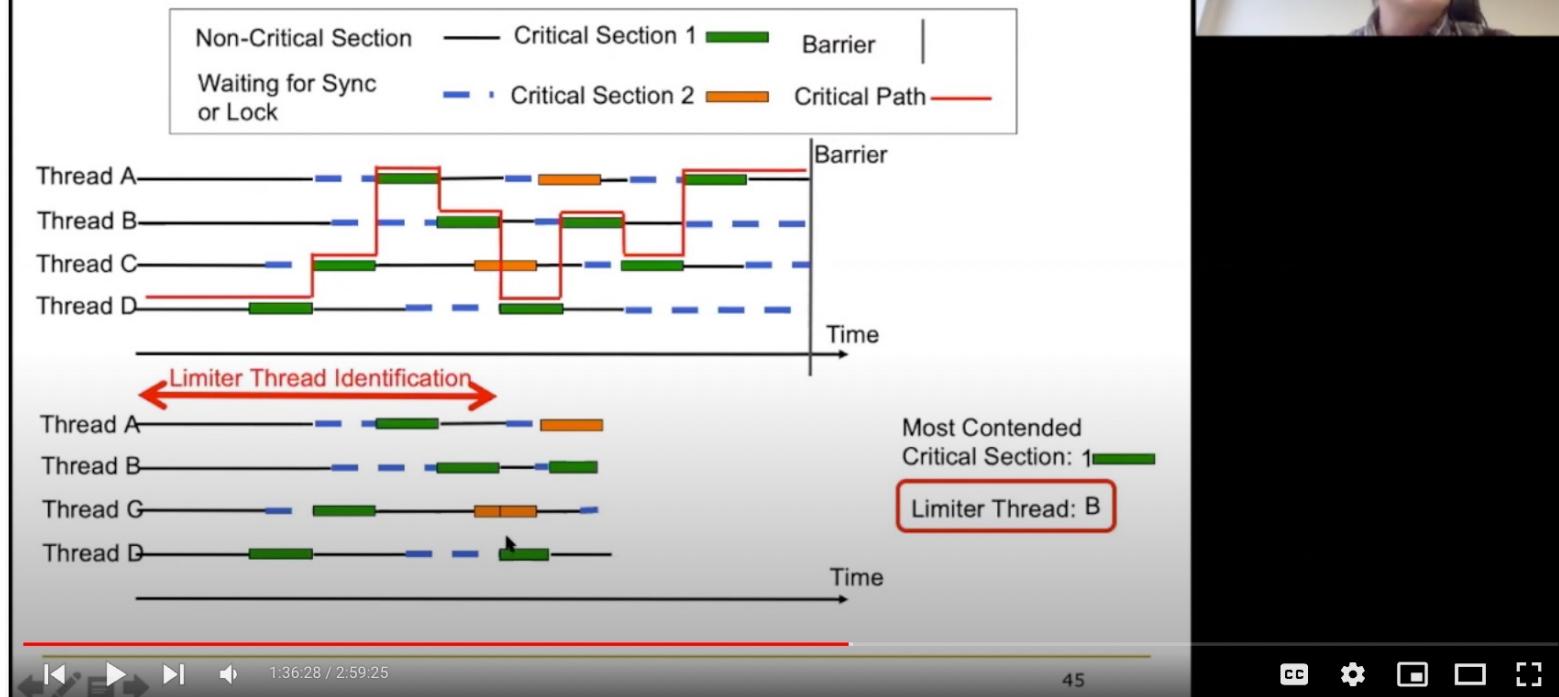
[†]Department of Electrical and Computer Engineering
The University of Texas at Austin
{ebrahimi,rustam,joao,patt}@ece.utexas.edu

[§]Carnegie Mellon University
{cfallin,onur}@cmu.edu

[‡]Intel Corporation
chang.joo.lee@intel.com

Lecture on Parallel Application Scheduling

Prioritizing Requests from Limiter Threads



ETH ZENTRUM

Computer Arch - Lecture 13: Memory Interference and Quality of Service II (ETH Zürich, Spring 2020)

1,006 views • Nov 7, 2020

23 0 SHARE SAVE ...



Onur Mutlu Lectures
19.8K subscribers

ANALYTICS

EDIT VIDEO

Bottleneck Identification & Scheduling

- Jose A. Joao, M. Aater Suleman, Onur Mutlu, and Yale N. Patt,
"[Bottleneck Identification and Scheduling in Multithreaded Applications](#)"

Proceedings of the [17th International Conference on Architectural Support for Programming Languages and Operating Systems \(ASPLOS\)](#), London, UK, March 2012. [Slides \(ppt\)](#) [\(pdf\)](#)

Bottleneck Identification and Scheduling in Multithreaded Applications

José A. Joao

ECE Department
The University of Texas at Austin
joao@ece.utexas.edu

M. Aater Suleman

Calxeda Inc.
aater.suleman@calxeda.com

Onur Mutlu

Computer Architecture Lab.
Carnegie Mellon University
onur@cmu.edu

Yale N. Patt

ECE Department
The University of Texas at Austin
patt@ece.utexas.edu

Utility-Based Bottleneck Acceleration

- Jose A. Joao, M. Aater Suleman, Onur Mutlu, and Yale N. Patt,
**"Utility-Based Acceleration of Multithreaded Applications
on Asymmetric CMPs"**

Proceedings of the 40th International Symposium on Computer Architecture (ISCA), Tel-Aviv, Israel, June 2013. [Slides \(ppt\)](#) [Slides \(pdf\)](#)

Utility-Based Acceleration of Multithreaded Applications on Asymmetric CMPs

José A. Joao [†] M. Aater Suleman ^{‡†} Onur Mutlu [§] Yale N. Patt [†]

[†] ECE Department
The University of Texas at Austin
Austin, TX, USA
{joao,patt}@ece.utexas.edu

[‡] Flux7 Consulting
Austin, TX, USA
suleman@hps.utexas.edu

[§] Computer Architecture Laboratory
Carnegie Mellon University
Pittsburgh, PA, USA
onur@cmu.edu

Lecture on Bottleneck Acceleration

Observation: Limiting Bottlenecks Change Over Time

A=full linked list; B=empty linked list
repeat

Lock A

Traverse list A

Remove X from A

Unlock A

Compute on X

Lock B

Traverse list B

Insert X into B

Unlock B

32 threads

Contention (# of threads waiting)

time [Mcycles]

until A is empty

55:08 / 2:30:02

◀ ▶ ⏪ 🔍 ⚙️

Computer Architecture - Lecture 17: Bottleneck Acceleration (ETH Zürich, Fall 2020)

880 views • Nov 20, 2020

13 0 SHARE SAVE ...



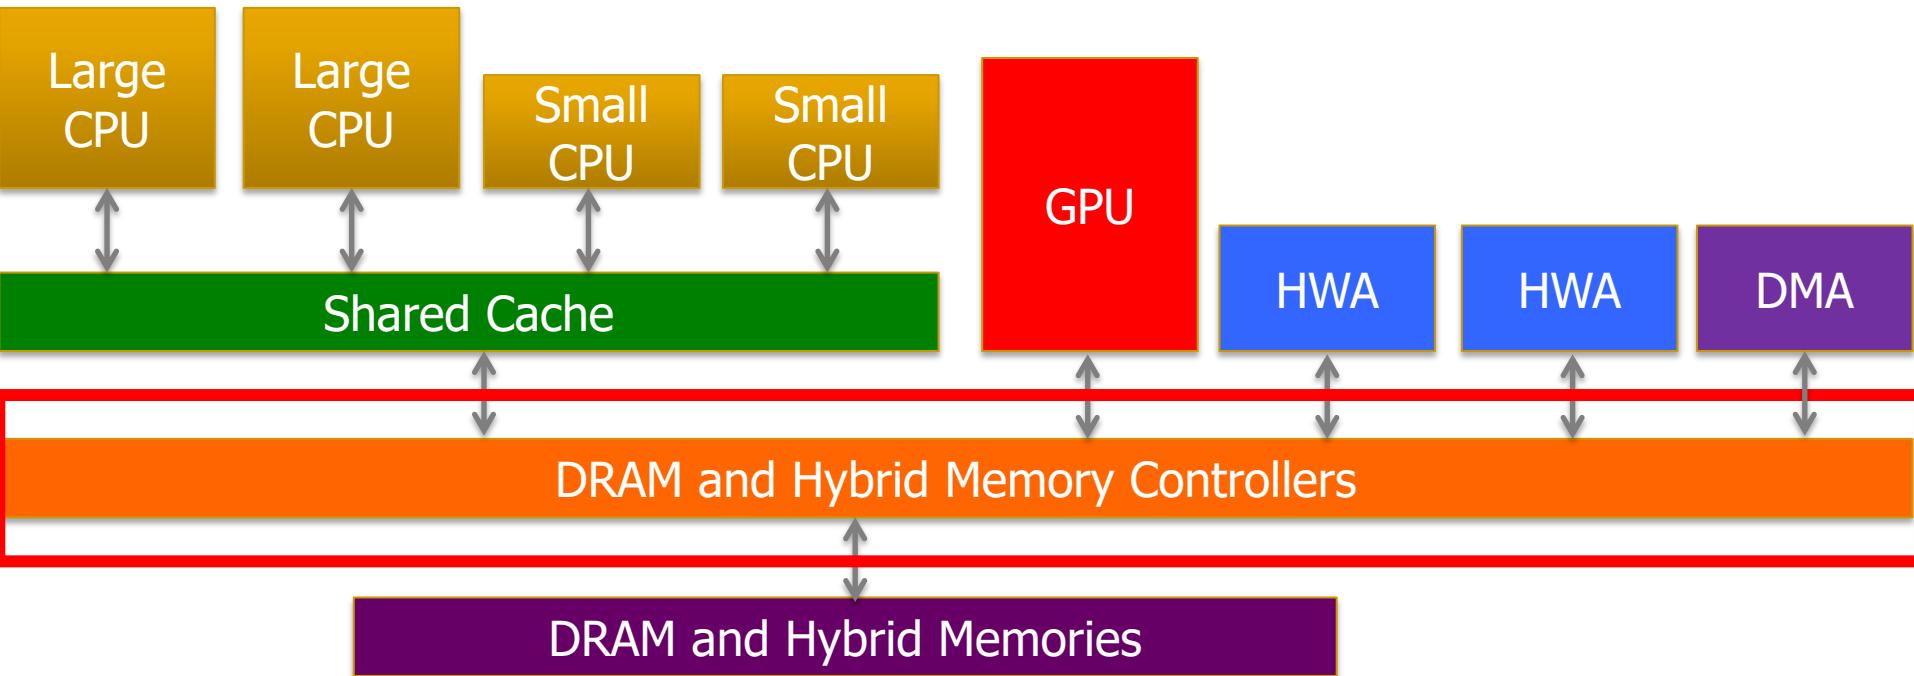
Onur Mutlu Lectures
19.8K subscribers

ANALYTICS

EDIT VIDEO

Memory Scheduling for Heterogeneous Systems

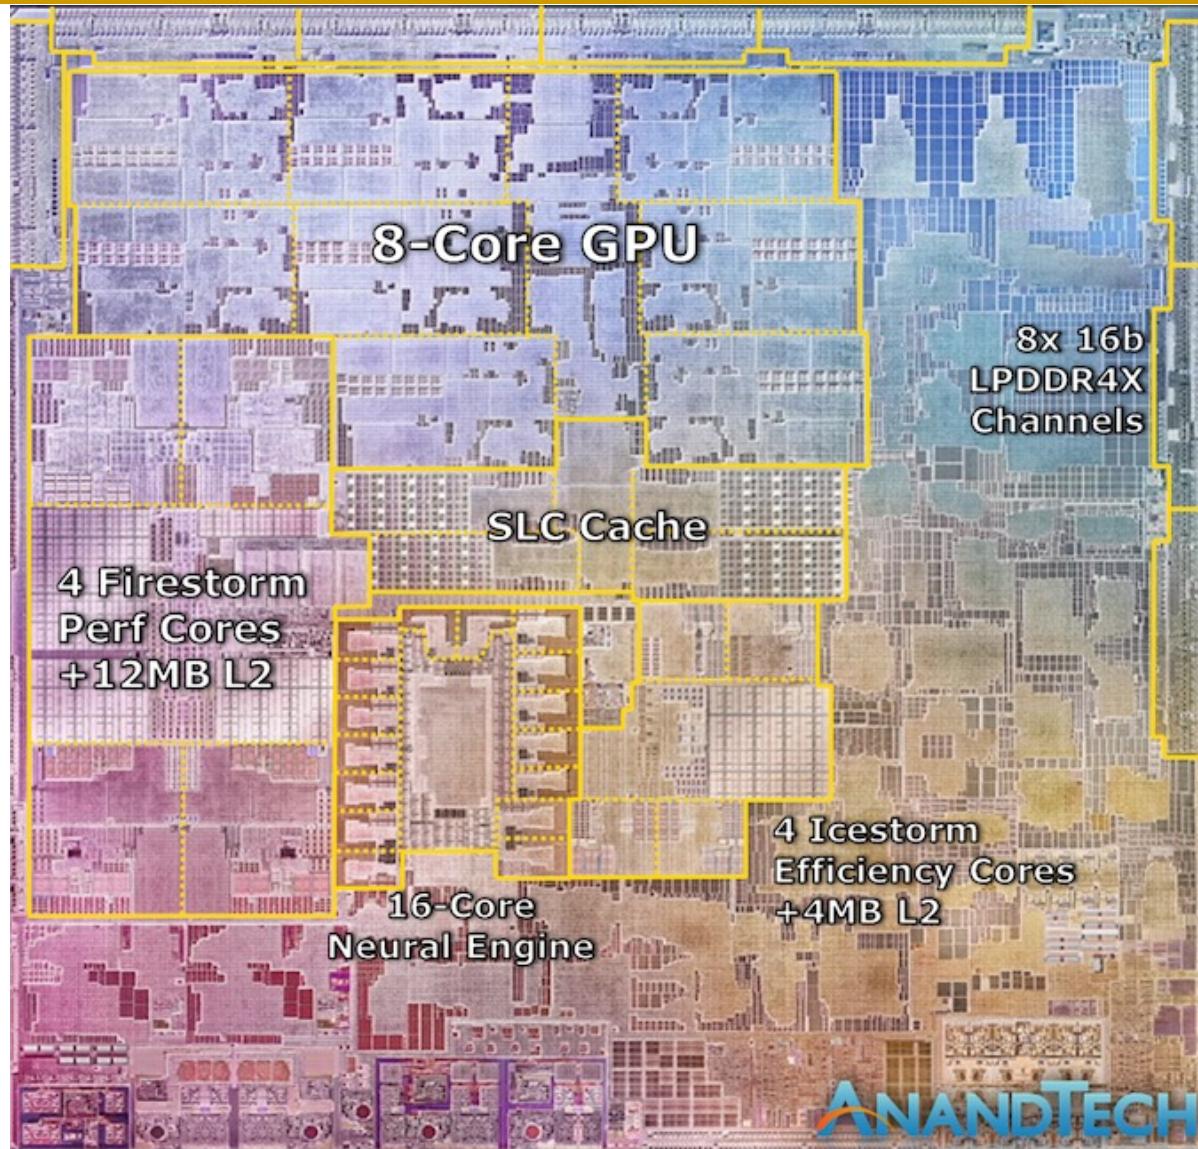
Current SoC Architectures: Heterogeneity



- Heterogeneous agents: CPUs, GPUs, HWAs, DMA engine, ...
- Memory resources shared by CPUs/GPUs/HWAs → Interference

How to allocate resources to heterogeneous agents
to mitigate interference and provide predictable performance?

Current SoC Architectures: Heterogeneity



Apple M1,
2021

Lecture on Heterogeneous System Scheduling

SMS: Staged Memory Scheduling

The diagram illustrates the SMS (Staged Memory Scheduling) process across three stages:

- Stage 1 (Batch Formation):** Five cores (Core 1, Core 2, Core 3, Core 4, GPU) each produce a single data block. These blocks are collected by a central scheduler.
- Stage 2 (Batch Scheduler):** The collected blocks are grouped into four batches, which are then assigned to four memory banks (Bank 1, Bank 2, Bank 3, Bank 4).
- Stage 3 (DRAM Command Scheduler):** The four banks send commands to DRAM simultaneously.

Below the diagram is a video player interface showing the video is at 16:23 / 2:59:25, has 15 likes, and includes standard video controls like play/pause, volume, and full screen.

ETH ZENTRUM
Computer Arch - Lecture 13: Memory Interference and Quality of Service II (ETH Zürich, Spring 2020)
1,006 views • Nov 7, 2020

Onur Mutlu Lectures
19.8K subscribers

ANALYTICS EDIT VIDEO

Lecture on Heterogeneous Computing Systems

Asymmetry Enables Customization

| | | | |
|---|---|---|---|
| c | c | c | c |
| c | c | c | c |
| c | c | c | c |
| c | c | c | c |

Symmetric

| | | | |
|----|----|----|----|
| C1 | | C2 | |
| | | C3 | |
| C4 | C4 | C4 | C4 |
| C5 | C5 | C5 | C5 |

Asymmetric



- Symmetric: One size fits all
 - Energy and performance suboptimal for different “workload” behaviors
- Asymmetric: Enables customization and adaptation
 - Processing requirements vary across workloads (applications and phases)
 - Execute code on best-fit resources (minimal energy, adequate perf.)

13:13 / 1:09:33

16

II CC HD □ □ ☰

Computer Architecture - Lecture 16b: Parallelism and Heterogeneity (ETH Zürich, Fall 2020)

840 views • Nov 20, 2020

17 DISLIKE SHARE SAVE ...



Onur Mutlu Lectures
19.9K subscribers

ANALYTICS

EDIT VIDEO

Staged Memory Scheduling

- Rachata Ausavarungnirun, Kevin Chang, Lavanya Subramanian, Gabriel Loh, and Onur Mutlu,

"Staged Memory Scheduling: Achieving High Performance and Scalability in Heterogeneous Systems"

Proceedings of the 39th International Symposium on Computer Architecture (ISCA), Portland, OR, June 2012. Slides (pptx)

Staged Memory Scheduling: Achieving High Performance and Scalability in Heterogeneous Systems

Rachata Ausavarungnirun[†] Kevin Kai-Wei Chang[†] Lavanya Subramanian[†] Gabriel H. Loh[‡] Onur Mutlu[†]

[†]Carnegie Mellon University

{rachata,kevincha,lsubrama,onur}@cmu.edu

[‡]Advanced Micro Devices, Inc.
gabe.loh@amd.com

DASH: Deadline-Aware Memory Scheduler

- Hiroyuki Usui, Lavanya Subramanian, Kevin Kai-Wei Chang, and Onur Mutlu,

"DASH: Deadline-Aware High-Performance Memory Scheduler for Heterogeneous Systems with Hardware Accelerators"

ACM Transactions on Architecture and Code Optimization (TACO),

Vol. 12, January 2016.

Presented at the 11th HiPEAC Conference, Prague, Czech Republic, January 2016.

[Slides (pptx) (pdf)]

[Source Code]

DASH: Deadline-Aware High-Performance Memory Scheduler for Heterogeneous Systems with Hardware Accelerators

HIROYUKI USUI, LAVANYA SUBRAMANIAN, KEVIN KAI-WEI CHANG,
and ONUR MUTLU, Carnegie Mellon University

Handling CPU-IO Interference

Donghyuk Lee, Lavanya Subramanian, Rachata Ausavarungnirun, Jongmoo Choi,
and Onur Mutlu,

"Decoupled Direct Memory Access: Isolating CPU and IO Traffic by Leveraging a Dual-Data-Port DRAM"

Proceedings of the 24th International Conference on Parallel Architectures and Compilation Techniques (PACT), San Francisco, CA, USA, October 2015.

[[Slides \(pptx\)](#) ([pdf](#))]

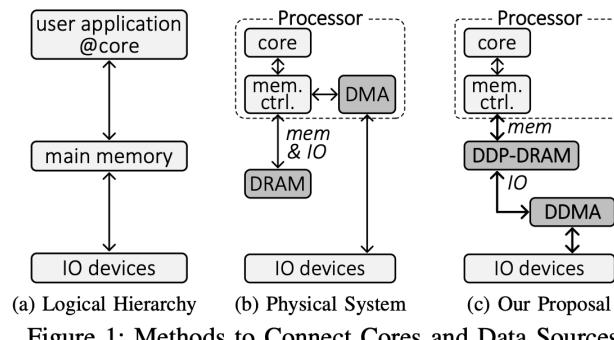


Figure 1: Methods to Connect Cores and Data Sources

Decoupled Direct Memory Access: Isolating CPU and IO Traffic by Leveraging a Dual-Data-Port DRAM

Donghyuk Lee* Lavanya Subramanian* Rachata Ausavarungnirun* Jongmoo Choi† Onur Mutlu*

*Carnegie Mellon University

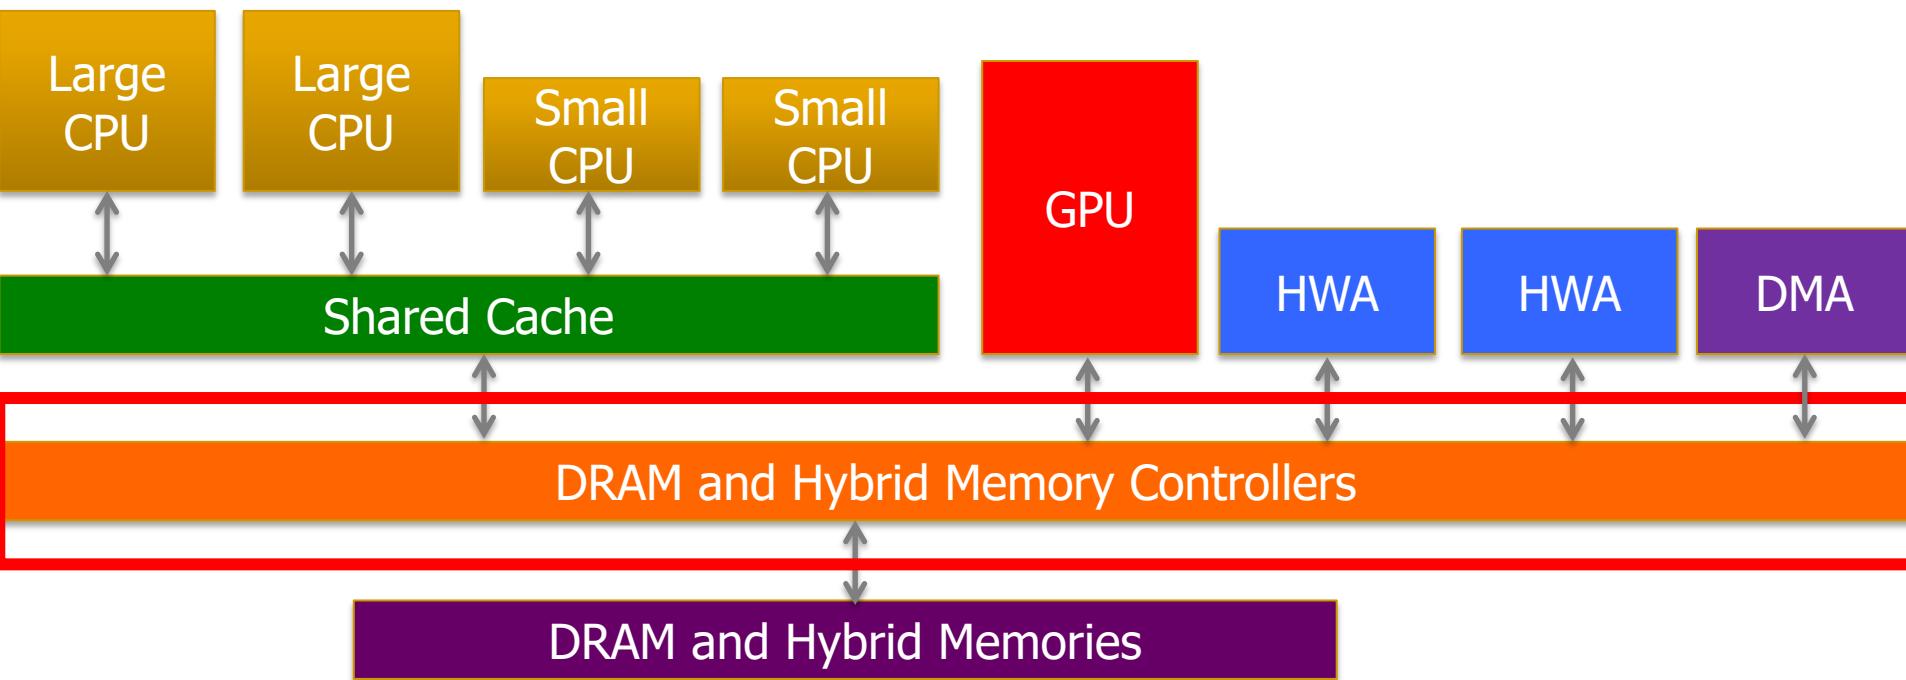
{donghyu1, lsubrama, rachata, onur}@cmu.edu

†Dankook University

choijm@dankook.ac.kr

Predictable Performance: Strong Memory Service Guarantees

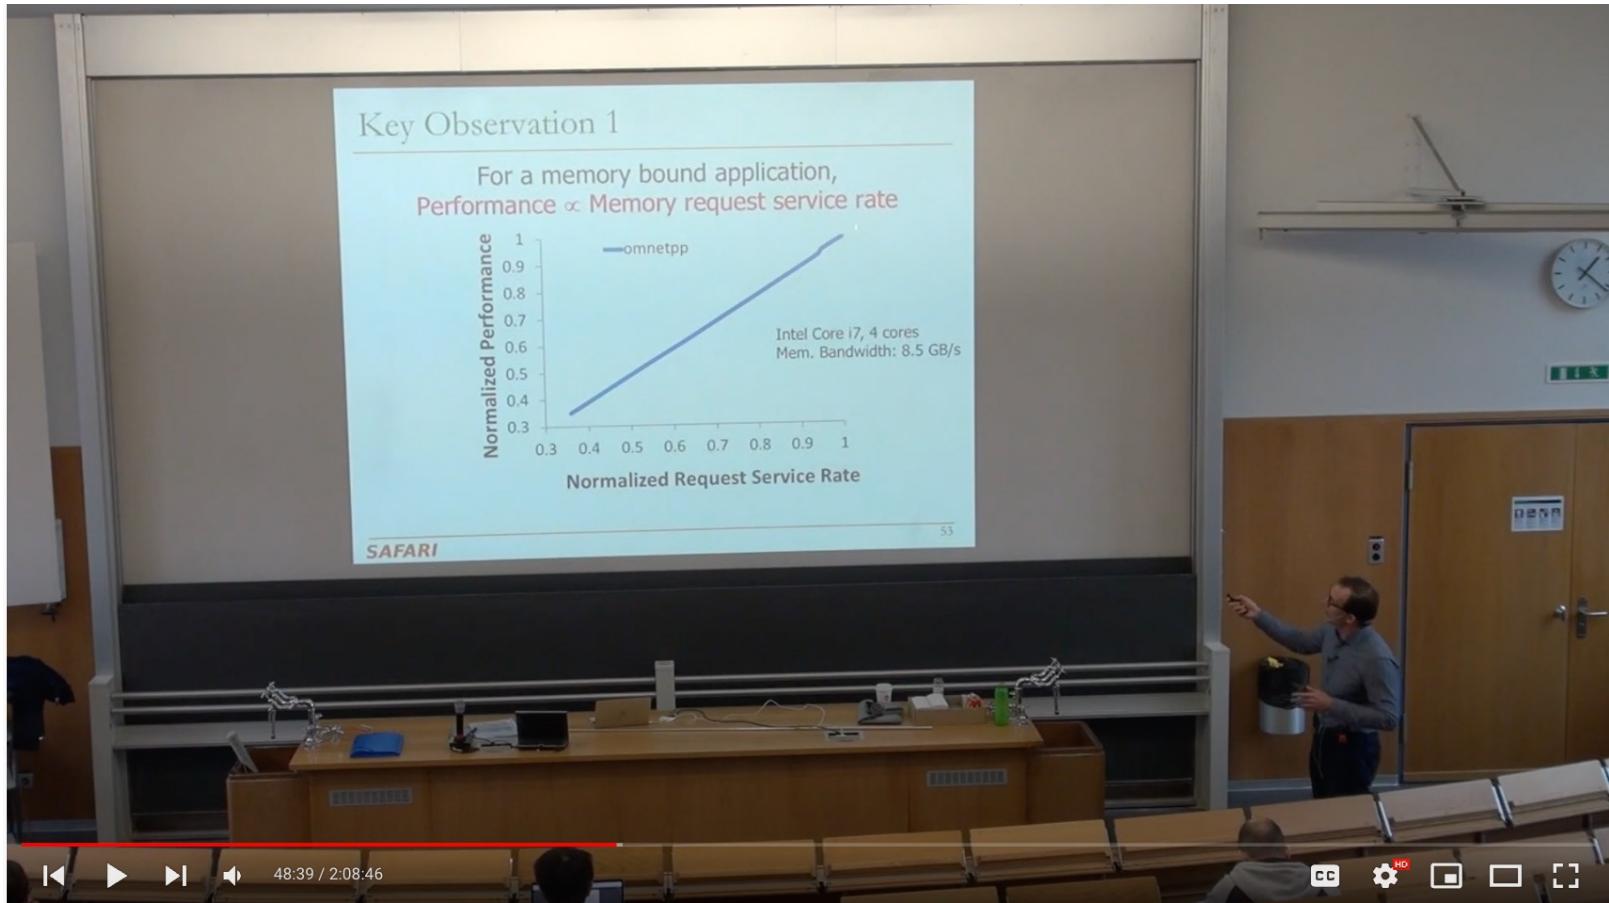
Goal: Predictable Performance in Complex Systems



- Heterogeneous agents: CPUs, GPUs, HWAs, DMA engine, ...
- Memory resources shared by CPUs/GPUs/HWAs → Interference

How to allocate resources to heterogeneous agents
to mitigate interference and provide predictable performance?

Lecture on Predictable Performance



ETH ZÜRICH HAUPTGEBÄUDE

Computer Architecture - Lecture 17: Memory Interference and QoS II (ETH Zürich, Fall 2018)

274 views • Nov 23, 2018

1 like 4 dislikes SHARE SAVE ...



Onur Mutlu Lectures
19.8K subscribers

ANALYTICS

EDIT VIDEO

Lecture on Predictable Performance

The slide title is "Effectiveness of MISE in Enforcing QoS" and it states "Across 3000 data points". It contains a 2x2 table:

| | Predicted Met | Predicted Not Met |
|-------------------|---------------|-------------------|
| QoS Bound Met | 78.8% | 2.1% |
| QoS Bound Not Met | 2.2% | 16.9% |

A callout highlights the 78.8% value. A box at the bottom states: "MISE-QoS correctly predicts whether or not the bound is met for 95.7% of workloads". The slide footer includes "SAFARI" and the number 161. The YouTube player controls show a progress bar at 25:29 / 1:11:14, and icons for CC, settings, and sharing.

Memory Systems - Lecture 6.4: Memory Interface and QoS (Technion, Summer 2018)

163 views • Oct 12, 2018

like 5 dislike 0 SHARE SAVE ...



Onur Mutlu Lectures
19.8K subscribers

ANALYTICS

EDIT VIDEO

Predictable Performance Readings (I)

Eiman Ebrahimi, Chang Joo Lee, Onur Mutlu, and Yale N. Patt,

"Fairness via Source Throttling: A Configurable and High-Performance Fairness Substrate for Multi-Core Memory Systems"

Proceedings of the 15th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), pages 335-346, Pittsburgh, PA, March 2010. [Slides \(pdf\)](#)

Best paper award.

Fairness via Source Throttling: A Configurable and High-Performance Fairness Substrate for Multi-Core Memory Systems

Eiman Ebrahimi† Chang Joo Lee† Onur Mutlu§ Yale N. Patt†

†Department of Electrical and Computer Engineering
The University of Texas at Austin
{ebrahimi, cjlee, patt}@ece.utexas.edu

§Computer Architecture Laboratory (CALCM)
Carnegie Mellon University
onur@cmu.edu

Predictable Performance Readings (II)

- Lavanya Subramanian, Vivek Seshadri, Yoongu Kim, Ben Jaiyen, and Onur Mutlu,

"MISE: Providing Performance Predictability and Improving Fairness in Shared Main Memory Systems"

Proceedings of the 19th International Symposium on High-Performance Computer Architecture (HPCA), Shenzhen, China, February 2013. [Slides \(pptx\)](#)

MISE: Providing Performance Predictability and Improving Fairness in Shared Main Memory Systems

Lavanya Subramanian

Vivek Seshadri

Yoongu Kim

Ben Jaiyen

Onur Mutlu

Carnegie Mellon University

Predictable Performance Readings (III)

- Lavanya Subramanian, Vivek Seshadri, Arnab Ghosh, Samira Khan, and Onur Mutlu,

"The Application Slowdown Model: Quantifying and Controlling the Impact of Inter-Application Interference at Shared Caches and Main Memory"

Proceedings of the 48th International Symposium on Microarchitecture (MICRO), Waikiki, Hawaii, USA, December 2015.

[[Slides \(pptx\)](#) ([pdf](#))] [[Lightning Session Slides \(pptx\)](#) ([pdf](#))] [[Poster \(pptx\)](#) ([pdf](#))]

[[Source Code](#)]

The Application Slowdown Model: Quantifying and Controlling the Impact of Inter-Application Interference at Shared Caches and Main Memory

Lavanya Subramanian*§ Vivek Seshadri* Arnab Ghosh*†
Samira Khan*‡ Onur Mutlu*

*Carnegie Mellon University §Intel Labs †IIT Kanpur ‡University of Virginia

Predictable Performance Readings (IV)

- Hiroyuki Usui, Lavanya Subramanian, Kevin Kai-Wei Chang, and Onur Mutlu,
"DASH: Deadline-Aware High-Performance Memory Scheduler for Heterogeneous Systems with Hardware Accelerators"
ACM Transactions on Architecture and Code Optimization (TACO),
Vol. 12, January 2016.
Presented at the 11th HiPEAC Conference, Prague, Czech Republic,
January 2016.
[\[Slides \(pptx\) \(pdf\)\]](#)
[\[Source Code\]](#)

DASH: Deadline-Aware High-Performance Memory Scheduler for Heterogeneous Systems with Hardware Accelerators

HIROYUKI USUI, LAVANYA SUBRAMANIAN, KEVIN KAI-WEI CHANG,
and ONUR MUTLU, Carnegie Mellon University

Handling CPU-IO Interference

Donghyuk Lee, Lavanya Subramanian, Rachata Ausavarungnirun, Jongmoo Choi,
and Onur Mutlu,

"Decoupled Direct Memory Access: Isolating CPU and IO Traffic by Leveraging a Dual-Data-Port DRAM"

Proceedings of the 24th International Conference on Parallel Architectures and Compilation Techniques (PACT), San Francisco, CA, USA, October 2015.

[[Slides \(pptx\)](#) ([pdf](#))]

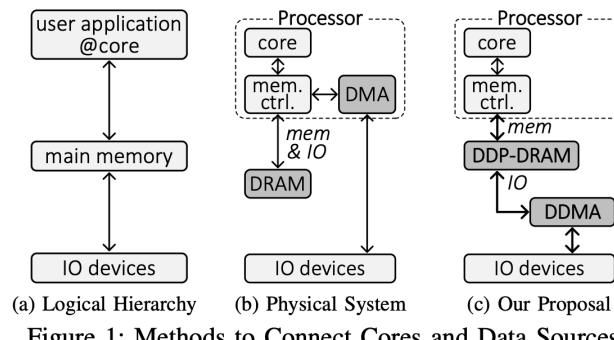


Figure 1: Methods to Connect Cores and Data Sources

Decoupled Direct Memory Access: Isolating CPU and IO Traffic by Leveraging a Dual-Data-Port DRAM

Donghyuk Lee* Lavanya Subramanian* Rachata Ausavarungnirun* Jongmoo Choi† Onur Mutlu*

*Carnegie Mellon University

{donghyu1, lsubrama, rachata, onur}@cmu.edu

†Dankook University

choijm@dankook.ac.kr

Other QoS Approaches

Recall: Fundamental Interference Control Techniques

- Goal: to reduce/control inter-thread memory interference

1. Prioritization or request scheduling

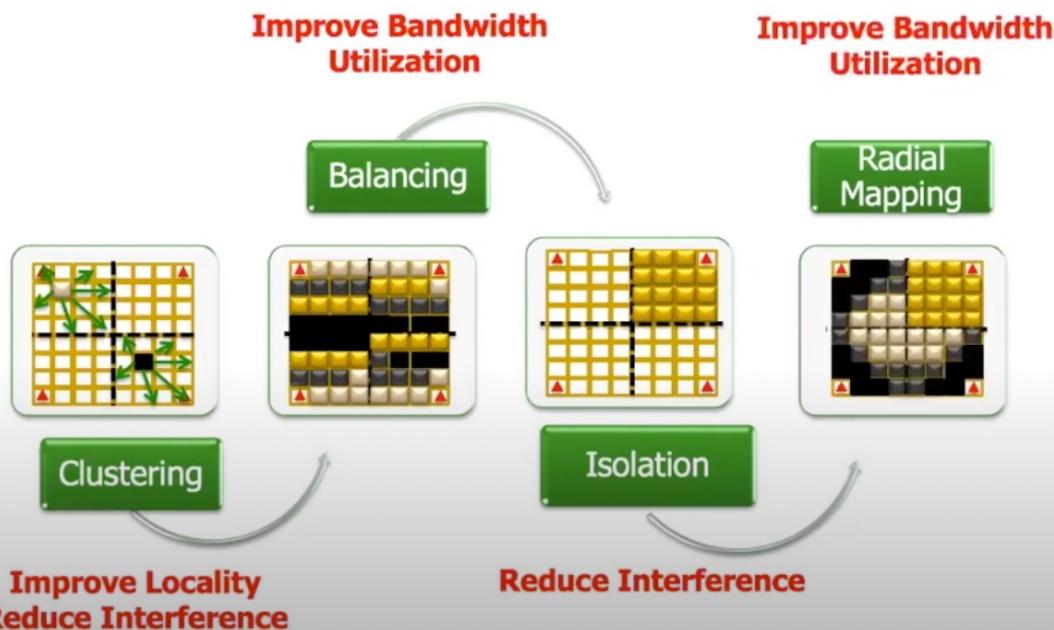
2. Data mapping to banks/channels/ranks

3. Core/source throttling

4. Application/thread scheduling

Lecture on Other QoS Techniques

Application-to-Core Mapping



ETH ZENTRUM

Computer Arch - Lecture 13: Memory Interference and Quality of Service II (ETH Zürich, Spring 2020)

1,006 views • Nov 7, 2020

89

CC G S D E

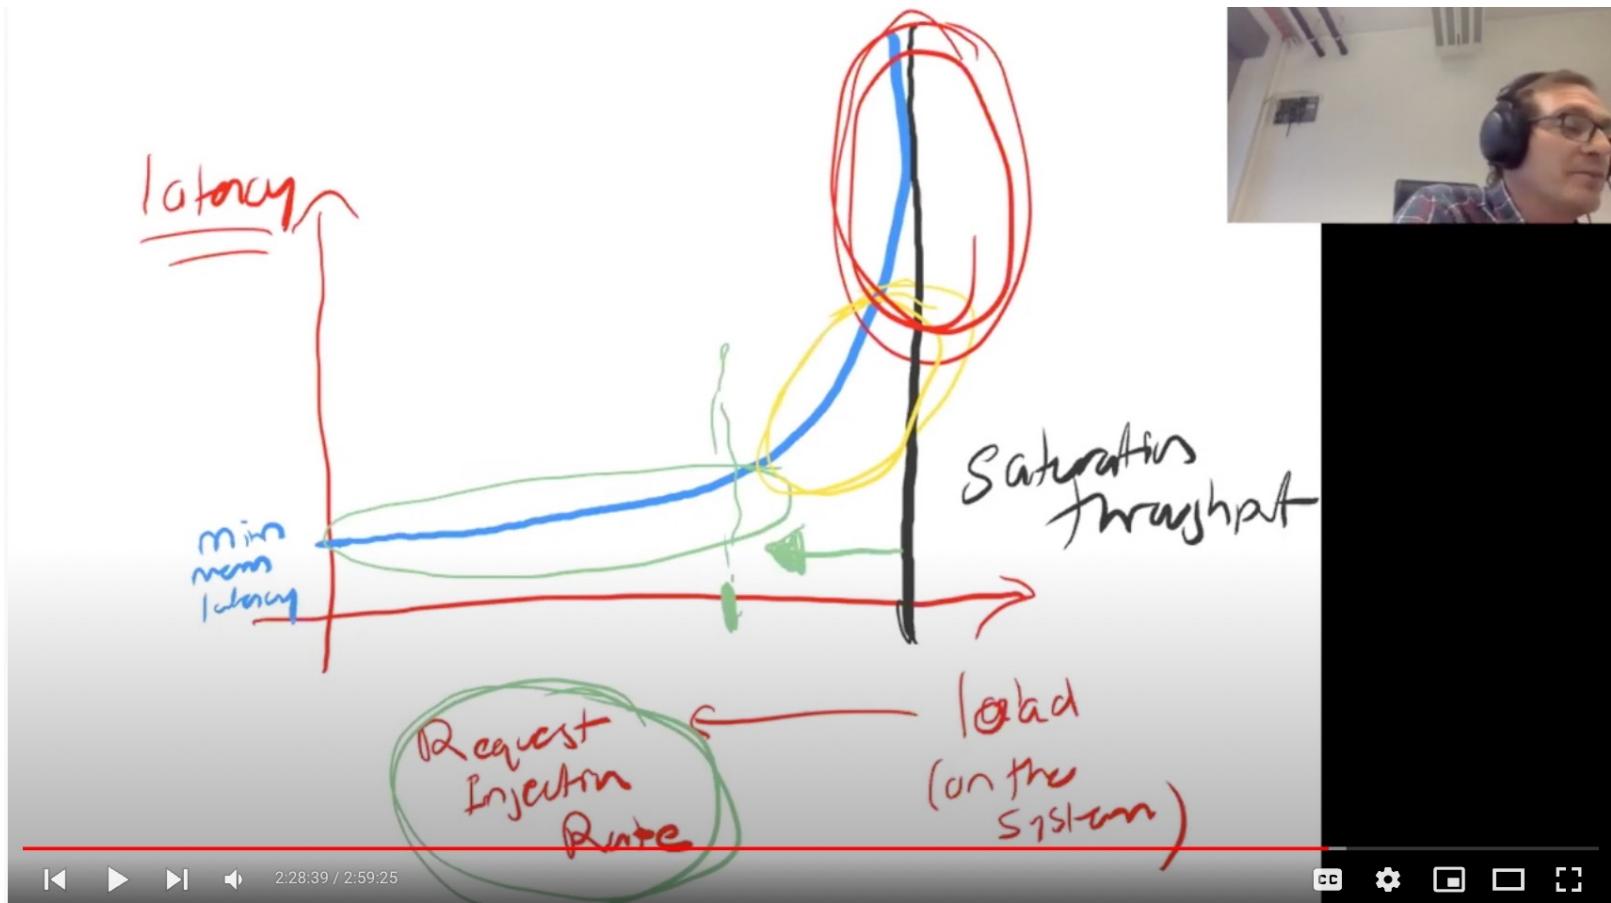


Onur Mutlu Lectures
19.8K subscribers

ANALYTICS

EDIT VIDEO

Lecture on Other QoS Techniques



ETH ZENTRUM

Computer Arch - Lecture 13: Memory Interference and Quality of Service II (ETH Zürich, Spring 2020)

1,006 views • Nov 7, 2020

23 likes 0 dislikes SHARE SAVE ...



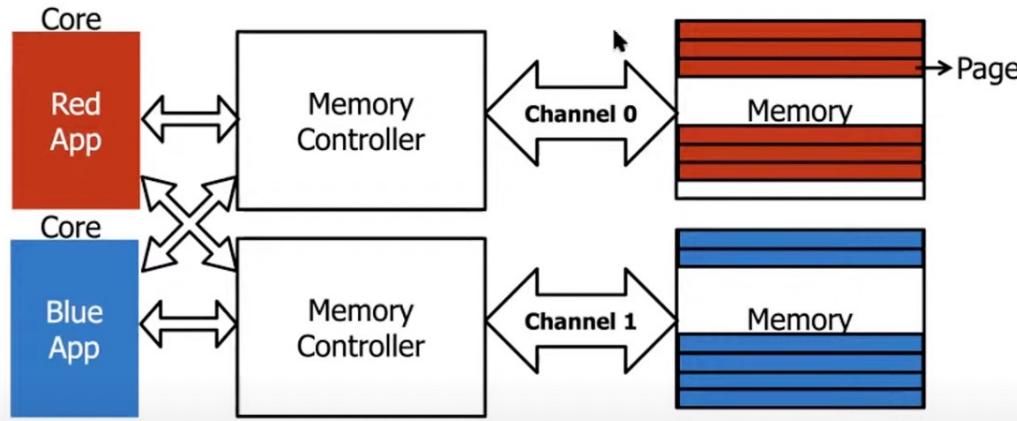
Onur Mutlu Lectures
19.8K subscribers

ANALYTICS

EDIT VIDEO

Memory Channel Partitioning

Partitioning Channels Between Applications



Eliminates interference between applications' requests



ETH ZURICH D-ITET

Seminar in Computer Architecture - Lecture 4: Memory Channel Partitioning (Fall 2021)

379 views • Streamed live on Oct 14, 2021

19 likes 0 dislikes SHARE SAVE ...



Onur Mutlu Lectures
19.8K subscribers

ANALYTICS

EDIT VIDEO

Memory Channel Partitioning

- Sai Prashanth Muralidhara, Lavanya Subramanian, Onur Mutlu, Mahmut Kandemir, and Thomas Moscibroda,
"Reducing Memory Interference in Multicore Systems via Application-Aware Memory Channel Partitioning"
Proceedings of the 44th International Symposium on Microarchitecture (MICRO), Porto Alegre, Brazil, December 2011. Slides (pptx)

Reducing Memory Interference in Multicore Systems via Application-Aware Memory Channel Partitioning

Sai Prashanth Muralidhara
Pennsylvania State University
smuralid@cse.psu.edu

Lavanya Subramanian
Carnegie Mellon University
lsubrama@ece.cmu.edu

Onur Mutlu
Carnegie Mellon University
onur@cmu.edu

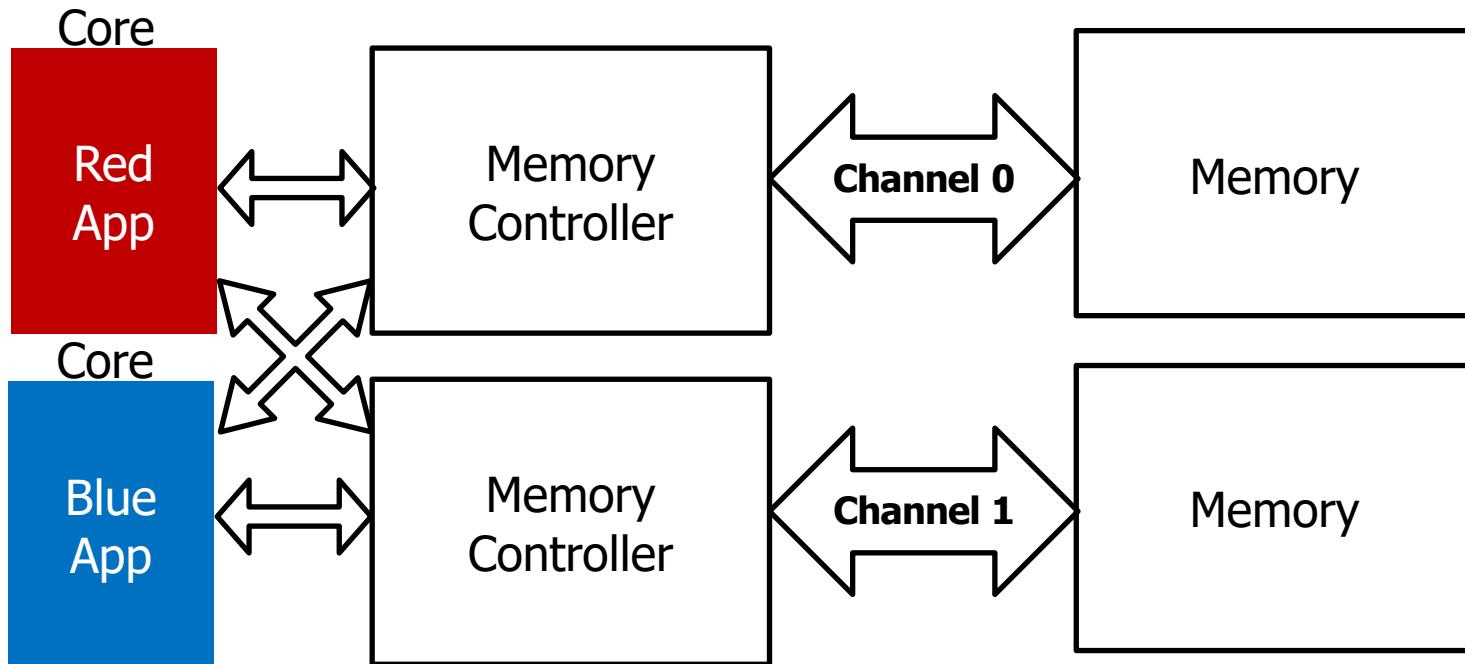
Mahmut Kandemir
Pennsylvania State University
kandemir@cse.psu.edu

Thomas Moscibroda
Microsoft Research Asia
moscitho@microsoft.com

Memory Channel Partitioning

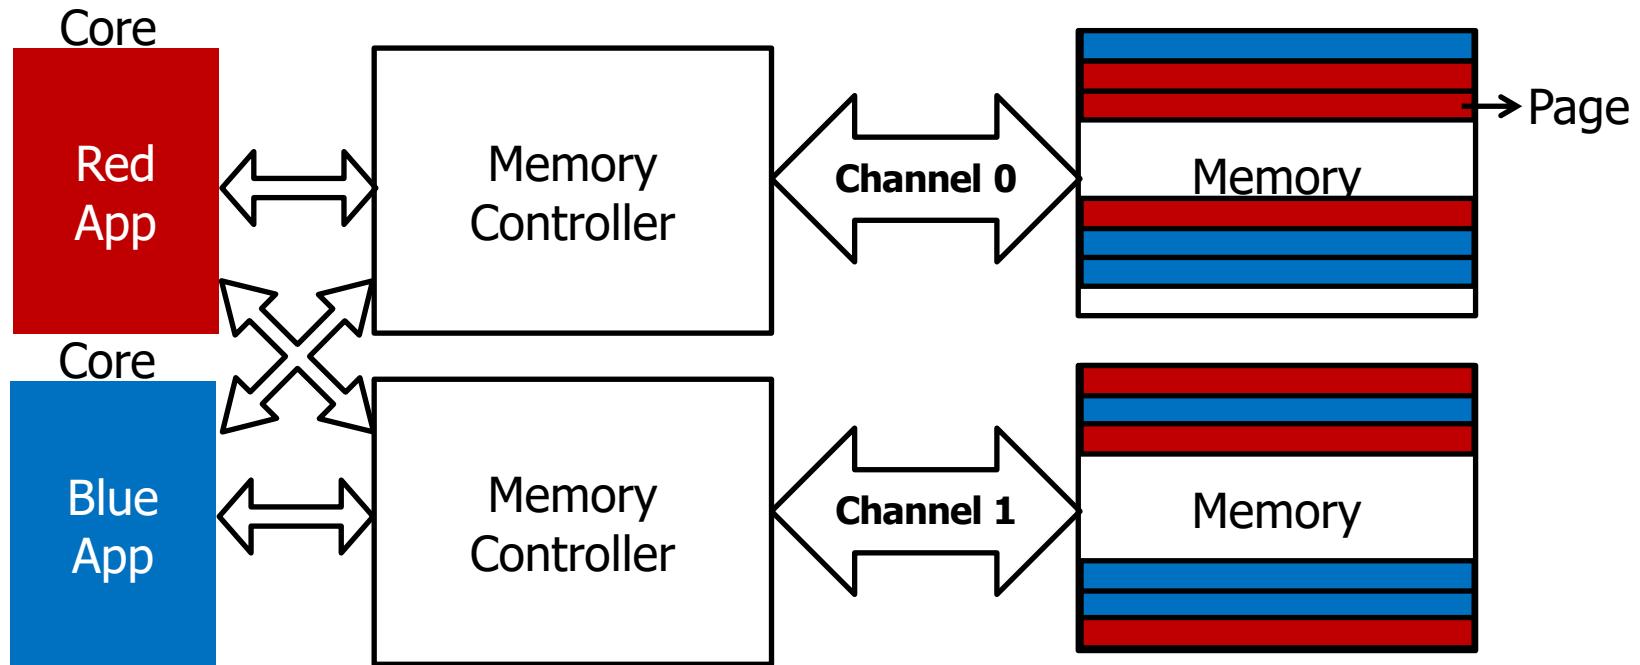
Sai Prashanth Muralidhara, Lavanya Subramanian, Onur Mutlu, Mahmut Kandemir, and Thomas Moscibroda,
**"Reducing Memory Interference in Multicore Systems via
Application-Aware Memory Channel Partitioning"**
44th International Symposium on Microarchitecture (MICRO),
Porto Alegre, Brazil, December 2011. [Slides \(pptx\)](#)

Observation: Modern Systems Have Multiple Channels



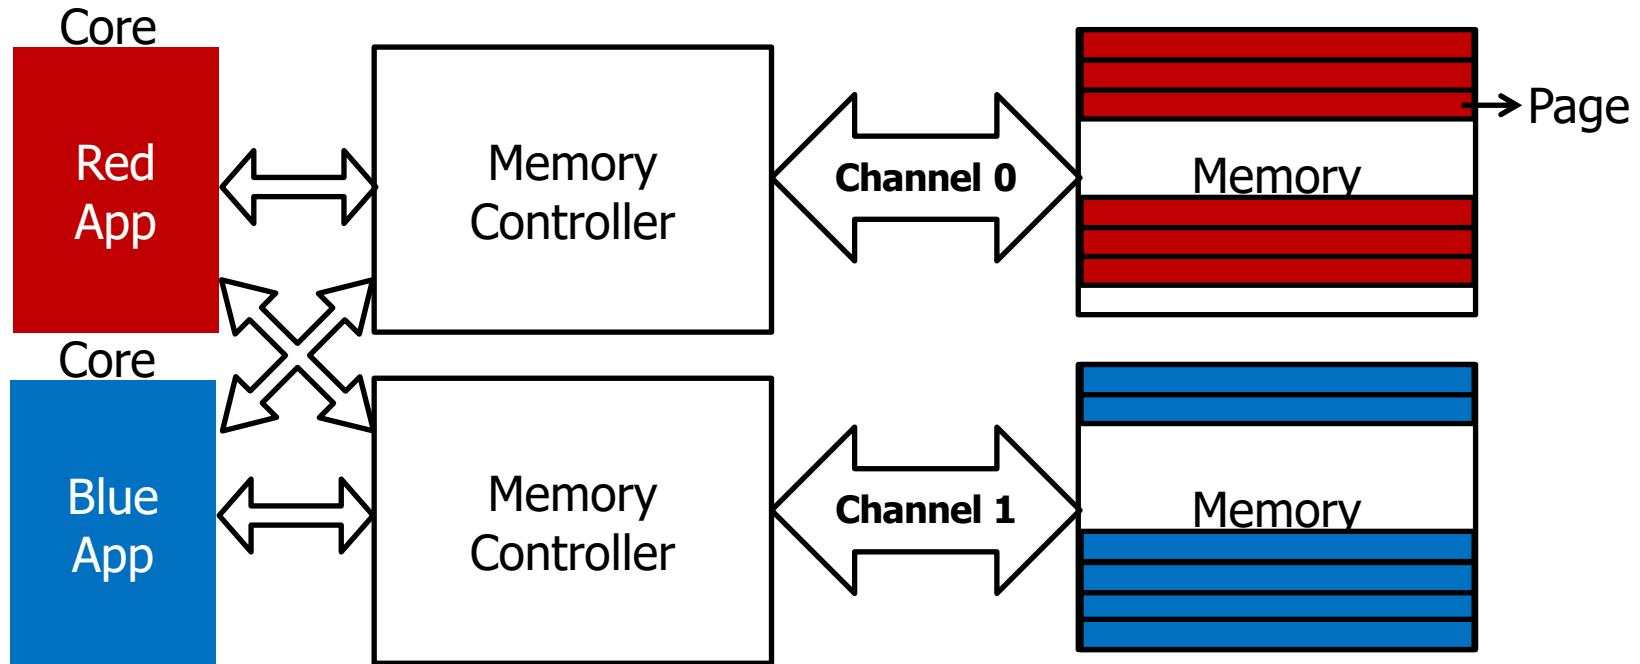
A new degree of freedom
Mapping data across multiple channels

Data Mapping in Current Systems



Causes interference between applications' requests

Partitioning Channels Between Applications



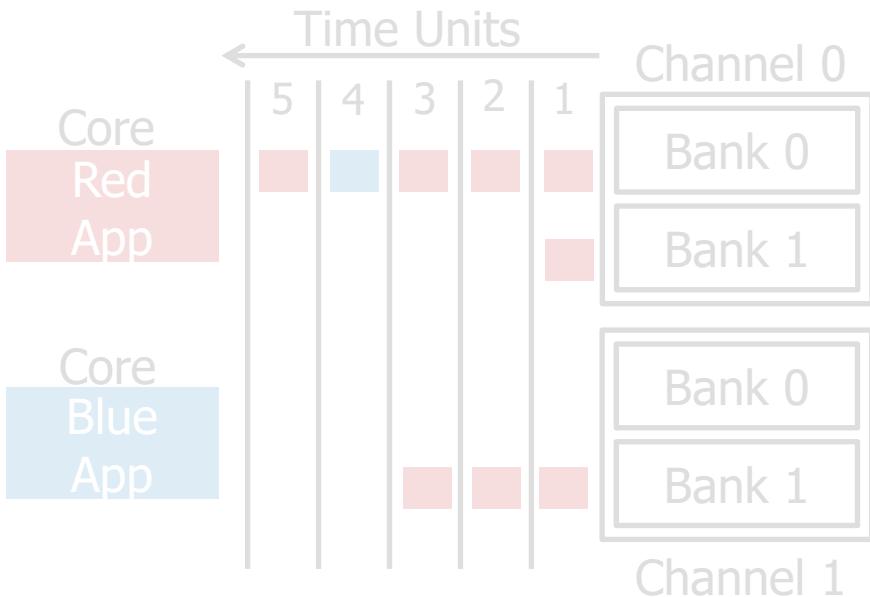
Eliminates interference between applications' requests

Overview: Memory Channel Partitioning (MCP)

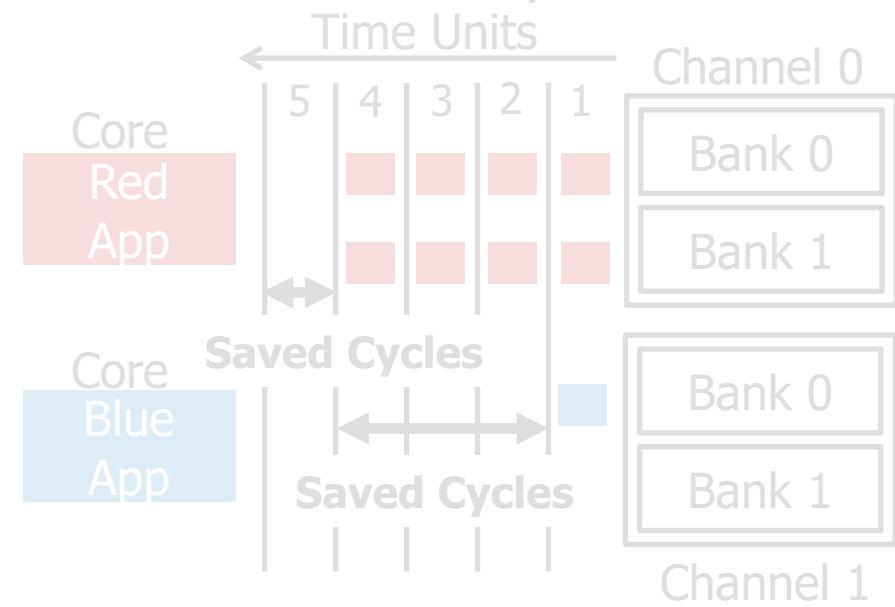
- Goal
 - Eliminate harmful interference between applications
- Basic Idea
 - Map the data of **badly-interfering applications** to different channels
- Key Principles
 - Separate **low and high memory-intensity applications**
 - Separate **low and high row-buffer locality applications**

Key Insight 1: Separate by Memory Intensity

High memory-intensity applications interfere with low memory-intensity applications in shared memory channels



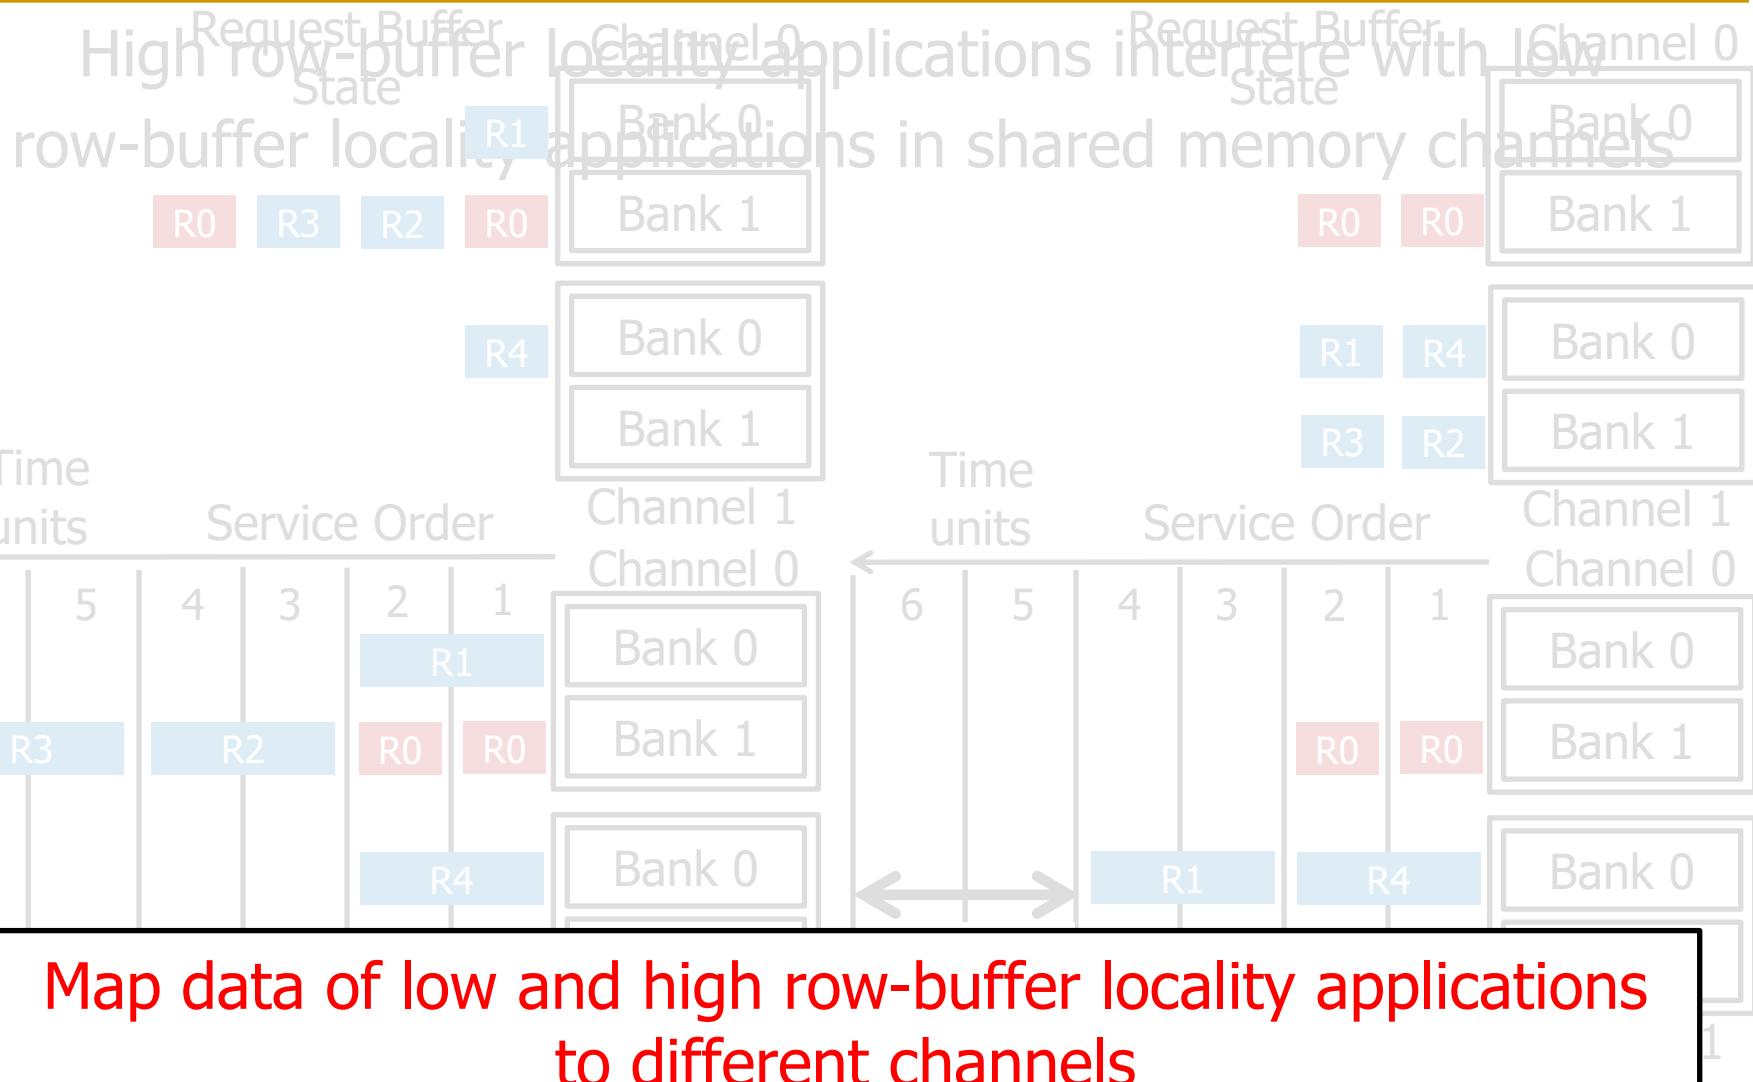
Conventional Page Mapping



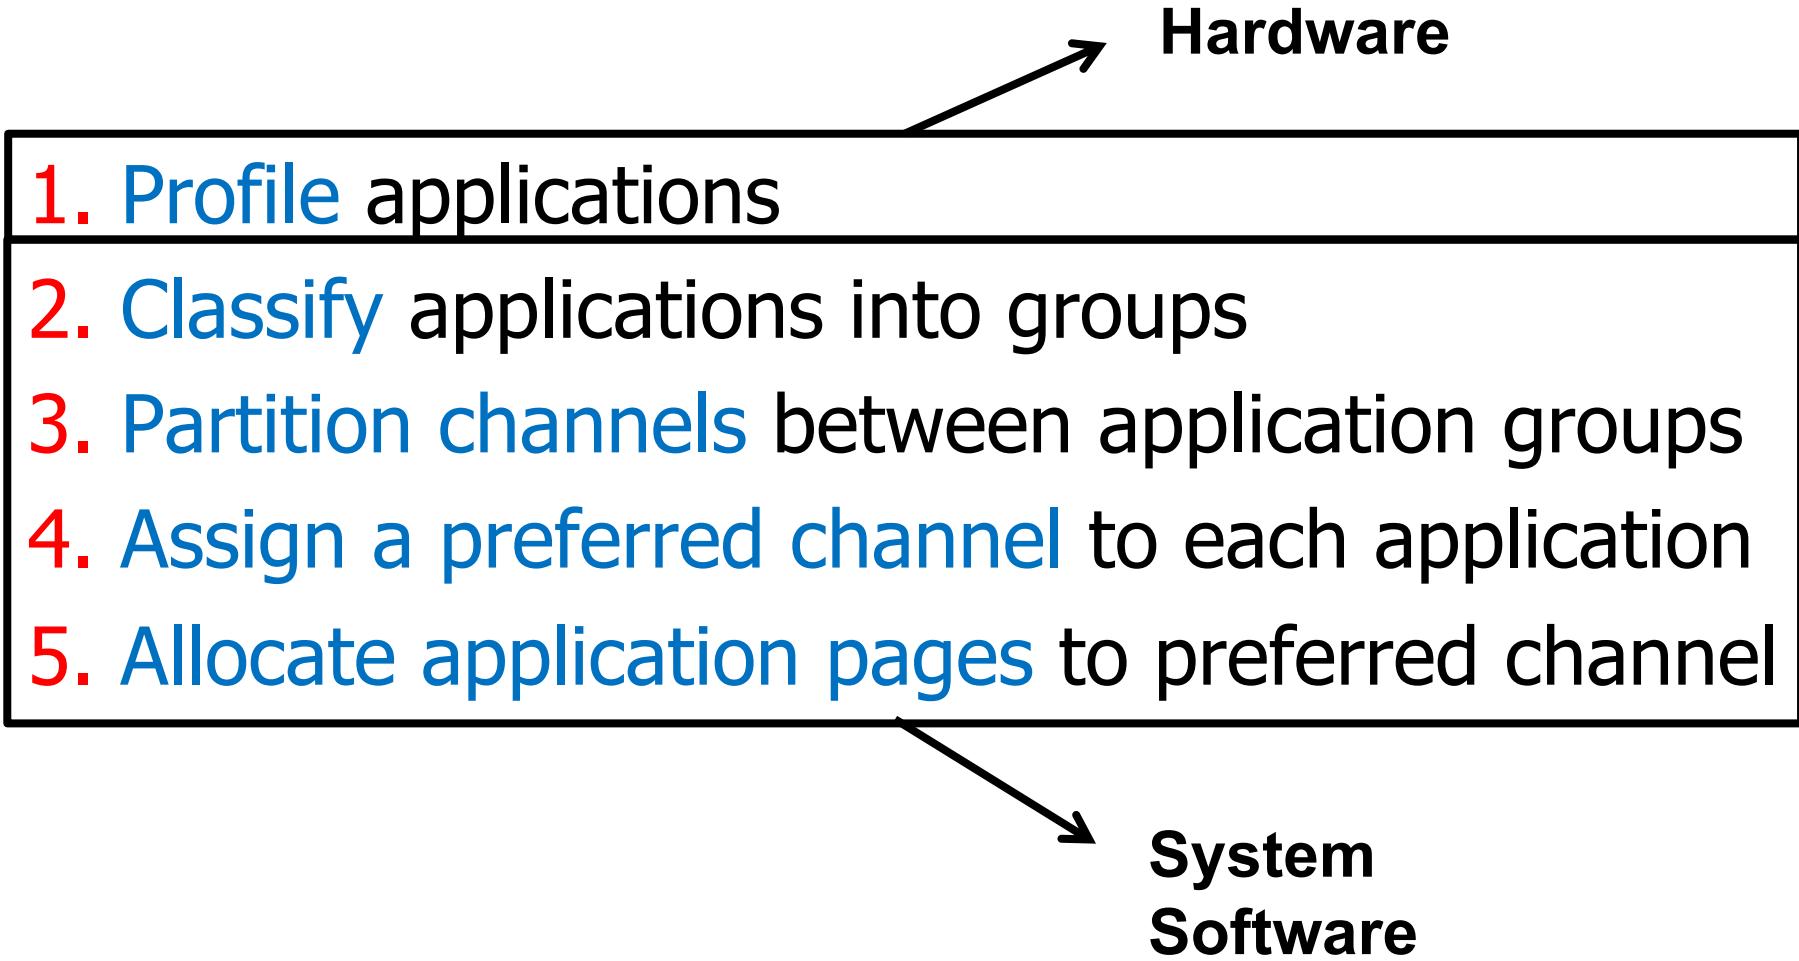
Channel Partitioning

Map data of low and high memory-intensity applications to different channels

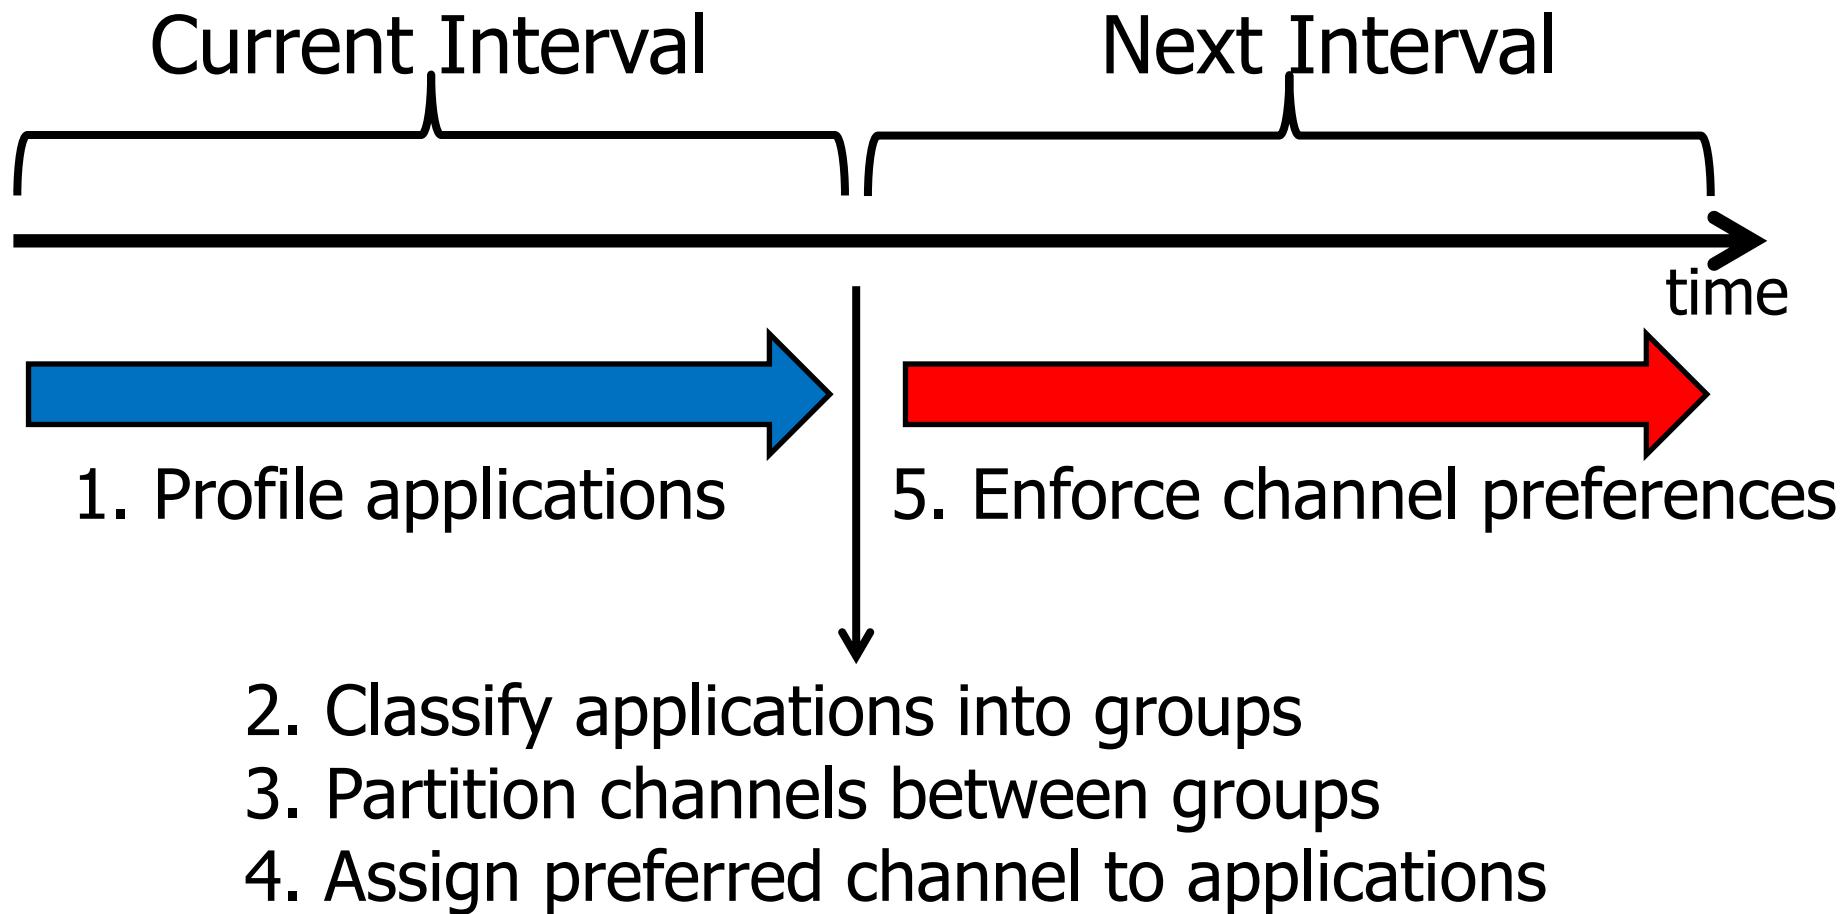
Key Insight 2: Separate by Row-Buffer Locality



Memory Channel Partitioning (MCP) Mechanism



Interval Based Operation



Observations

- Applications with very low memory-intensity rarely access memory
 - Dedicating channels to them results in precious memory bandwidth waste
- They have the most potential to keep their cores busy
 - We would really like to prioritize them
- They interfere minimally with other applications
 - Prioritizing them does not hurt others

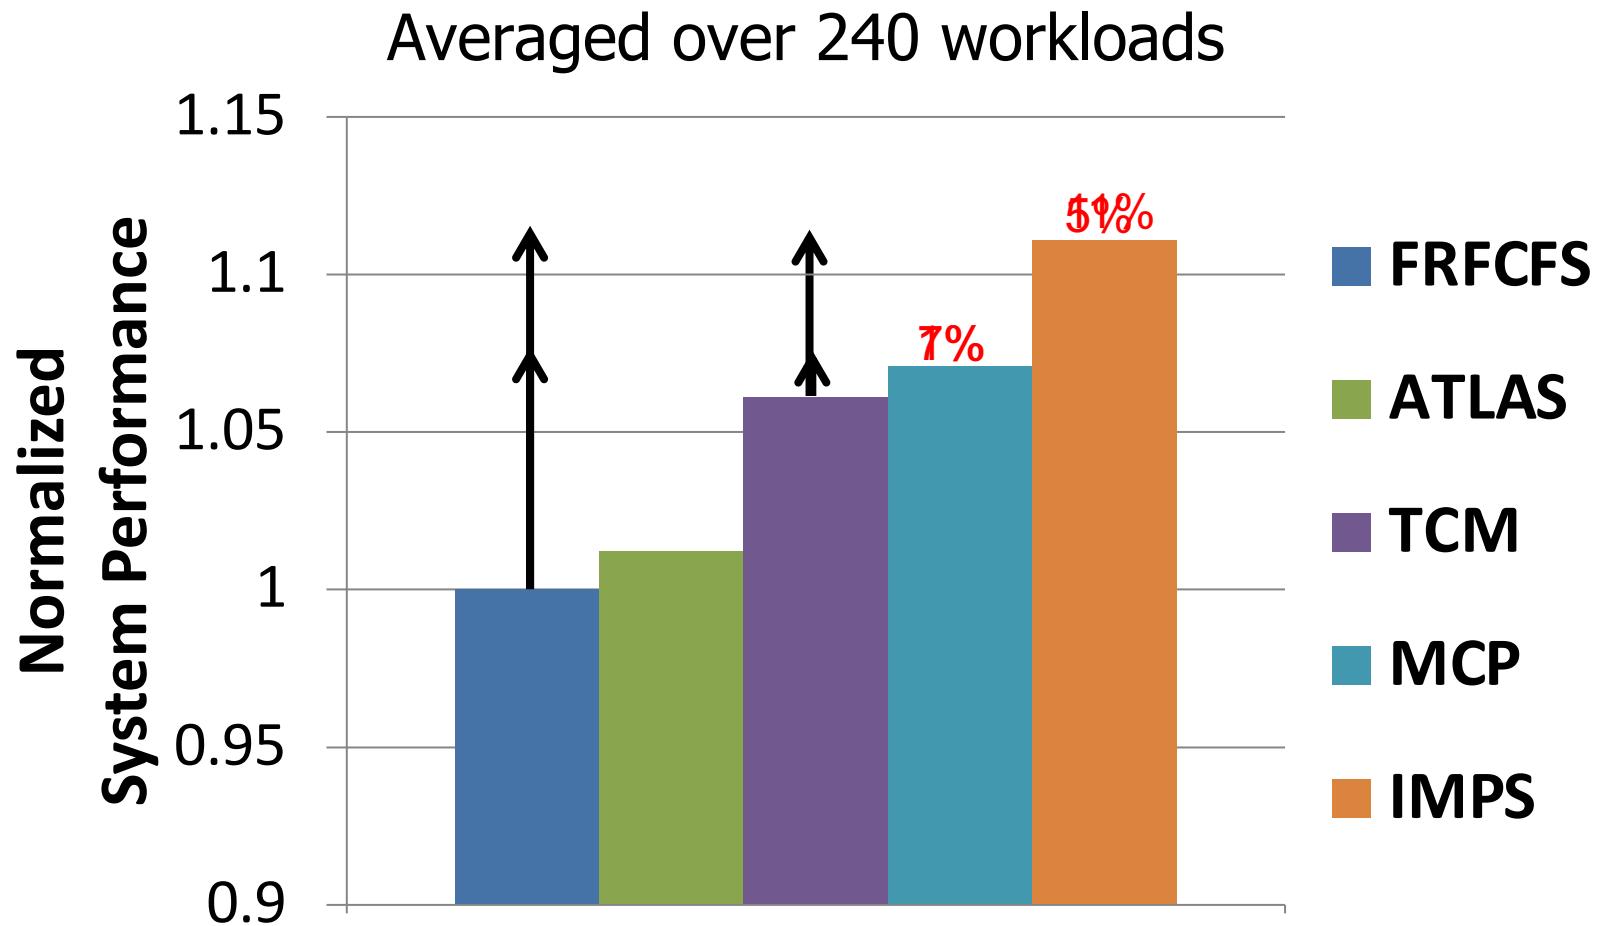
Integrated Memory Partitioning and Scheduling (IMPS)

- Always prioritize very low memory-intensity applications in the memory scheduler
- Use memory channel partitioning to mitigate interference between other applications

Hardware Cost

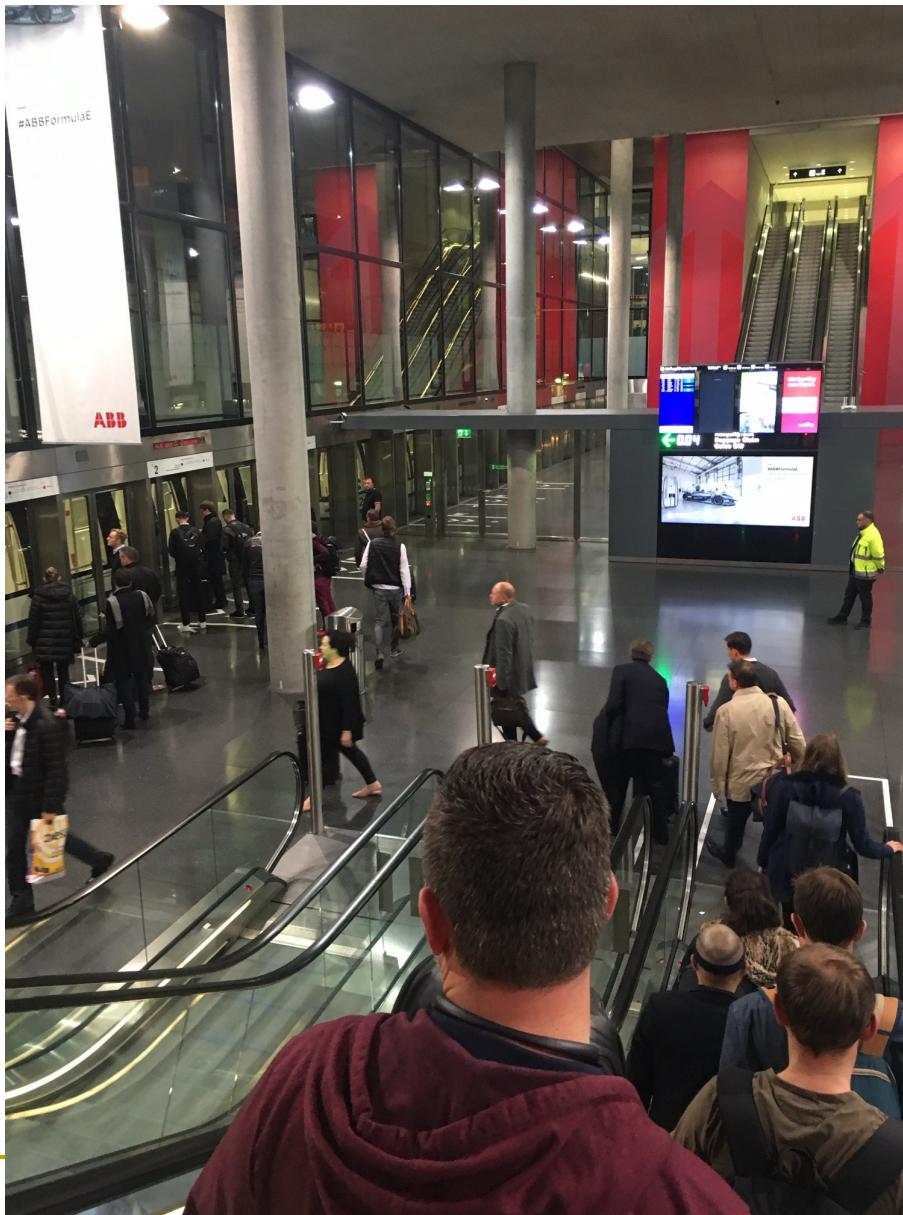
- Memory Channel Partitioning (MCP)
 - Only profiling counters in hardware
 - No modifications to memory scheduling logic
 - 1.5 KB storage cost for a 24-core, 4-channel system
- Integrated Memory Partitioning and Scheduling (IMPS)
 - A single bit per request
 - Scheduler prioritizes based on this single bit

Performance of Channel Partitioning



Better system performance than the best previous scheduler
at lower hardware cost

An Example of Bad Channel Partitioning



Combining Multiple Interference Control Techniques

- Combined interference control techniques can mitigate interference much more than a single technique alone can do
- The key challenge is:
 - Deciding what technique to apply when
 - Partitioning work appropriately between software and hardware

MCP and IMPS: Pros and Cons

- Upsides:
 - Keeps the memory scheduling hardware simple
 - Combines multiple interference reduction techniques
 - Can provide performance isolation across applications mapped to different channels
 - General idea of partitioning can be extended to smaller granularities in the memory hierarchy: banks, subarrays, etc.

- Downsides:
 - Reacting is difficult if workload changes behavior after profiling
 - Overhead of moving pages between channels restricts benefits

More on Memory Channel Partitioning

- Sai Prashanth Muralidhara, Lavanya Subramanian, Onur Mutlu, Mahmut Kandemir, and Thomas Moscibroda,
"Reducing Memory Interference in Multicore Systems via Application-Aware Memory Channel Partitioning"
Proceedings of the 44th International Symposium on Microarchitecture (MICRO), Porto Alegre, Brazil, December 2011. Slides (pptx)

Reducing Memory Interference in Multicore Systems via Application-Aware Memory Channel Partitioning

Sai Prashanth Muralidhara
Pennsylvania State University
smuralid@cse.psu.edu

Lavanya Subramanian
Carnegie Mellon University
lsubrama@ece.cmu.edu

Onur Mutlu
Carnegie Mellon University
onur@cmu.edu

Mahmut Kandemir
Pennsylvania State University
kandemir@cse.psu.edu

Thomas Moscibroda
Microsoft Research Asia
moscitho@microsoft.com

Source Throttling

Eiman Ebrahimi, Chang Joo Lee, Onur Mutlu, and Yale N. Patt,

"Fairness via Source Throttling: A Configurable and High-Performance Fairness Substrate for Multi-Core Memory Systems"

Proceedings of the 15th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), pages 335-346, Pittsburgh, PA, March 2010. [Slides \(pdf\)](#)

Best paper award.

Fairness via Source Throttling: A Configurable and High-Performance Fairness Substrate for Multi-Core Memory Systems

Eiman Ebrahimi† Chang Joo Lee† Onur Mutlu§ Yale N. Patt†

†Department of Electrical and Computer Engineering
The University of Texas at Austin
{ebrahimi, cjlee, patt}@ece.utexas.edu

§Computer Architecture Laboratory (CALCM)
Carnegie Mellon University
onur@cmu.edu

Fairness via Source Throttling

Eiman Ebrahimi, Chang Joo Lee, Onur Mutlu, and Yale N. Patt,

"Fairness via Source Throttling: A Configurable and High-Performance Fairness Substrate for Multi-Core Memory Systems"

15th Intl. Conf. on Architectural Support for Programming Languages and Operating Systems (ASPLOS),
pages 335-346, Pittsburgh, PA, March 2010. [Slides \(pdf\)](#)

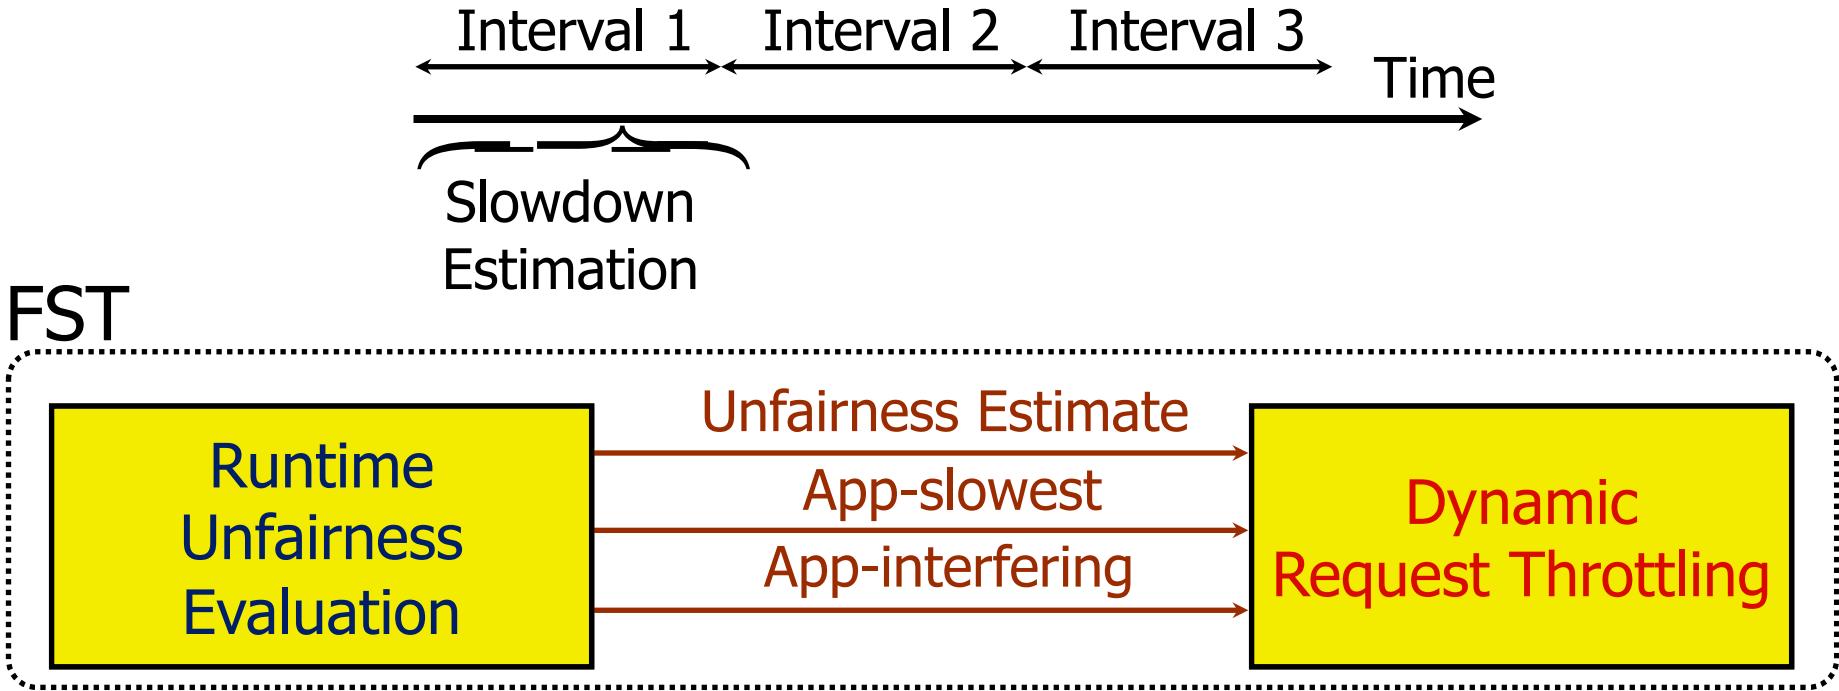
Source Throttling: A Fairness Substrate

- Key idea: Manage inter-thread interference at the **cores (sources)**, not at the **shared resources**
 - Dynamically estimate unfairness in the memory system
 - Feed back this information into a controller
 - Throttle cores' memory access rates accordingly
 - Whom to throttle and by how much depends on performance target (throughput, fairness, per-thread QoS, etc)
 - E.g., if unfairness > system-software-specified target then **throttle down** core causing unfairness & **throttle up** core that was unfairly treated
 - Ebrahimi et al., "Fairness via Source Throttling," ASPLOS'10, TOCS'12.
-

Fairness via Source Throttling (FST)

- Two components (interval-based)
- Run-time unfairness evaluation (in hardware)
 - Dynamically estimates the unfairness (application slowdowns) in the memory system
 - Estimates which application is slowing down which other
- Dynamic request throttling (hardware or software)
 - Adjusts how aggressively each core makes requests to the shared resources
 - Throttles down request rates of cores causing unfairness
 - Limit miss buffers, limit injection rate

Fairness via Source Throttling (FST) [ASPLOS'10]



1- Estimating system unfairness
2- Find app. with the highest slowdown (App-slowest)
3- Find app. causing most interference for App-slowest (App-interfering)

```
if (Unfairness Estimate > Target)
{
    1-Throttle down App-interfering
        (limit injection rate and parallelism)
    2-Throttle up App-slowest
}
```

Dynamic Request Throttling

- Goal: Adjust **how aggressively** each core makes requests to the shared memory system
- Mechanisms:
 - Miss Status Holding Register (MSHR) quota
 - Controls the **number of concurrent requests** accessing shared resources from each application
 - Request injection frequency
 - Controls **how often memory requests are issued** to the last level cache from the MSHRs

Dynamic Request Throttling

- Throttling level assigned to each core determines both MSHR quota and request injection rate

| Throttling level | MSHR quota | Request Injection Rate |
|------------------|------------|------------------------|
| 100% | 128 | Every cycle |
| 50% | 64 | Every other cycle |
| 25% | 32 | Once every 4 cycles |
| 10% | 12 | Once every 10 cycles |
| 5% | 6 | Once every 20 cycles |
| 4% | 5 | Once every 25 cycles |
| 3% | 3 | Once every 30 cycles |
| 2% | 2 | Once every 50 cycles |

Total # of MSHRs: 128

System Software Support

- Different fairness objectives can be configured by system software
 - Keep maximum slowdown in check
 - Estimated Max Slowdown < Target Max Slowdown
 - Keep slowdown of particular applications in check to achieve a particular performance target
 - Estimated Slowdown(i) < Target Slowdown(i)
- Support for thread priorities
 - Weighted Slowdown(i) =
Estimated Slowdown(i) x Weight(i)

Source Throttling Results: Takeaways

- Source throttling alone provides better performance than a combination of “smart” memory scheduling and fair caching
 - Decisions made at the memory scheduler and the cache sometimes contradict each other
- Neither source throttling alone nor “smart resources” alone provides the best performance
- Combined approaches are even more powerful
 - Source throttling and resource-based interference control

Source Throttling: Ups and Downs

- Advantages
 - + Core/request throttling is easy to implement: no need to change the memory scheduling algorithm
 - + Can be a general way of handling shared resource contention
 - + Can reduce overall load/contention in the memory system

- Disadvantages
 - Requires slowdown estimations → difficult to estimate
 - Thresholds can become difficult to optimize
 - throughput loss due to too much throttling
 - can be difficult to find an overall-good configuration

More on Source Throttling (I)

Eiman Ebrahimi, Chang Joo Lee, Onur Mutlu, and Yale N. Patt,

"Fairness via Source Throttling: A Configurable and High-Performance Fairness Substrate for Multi-Core Memory Systems"

Proceedings of the 15th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), pages 335-346, Pittsburgh, PA, March 2010. [Slides \(pdf\)](#)

Best paper award.

Fairness via Source Throttling: A Configurable and High-Performance Fairness Substrate for Multi-Core Memory Systems

Eiman Ebrahimi† Chang Joo Lee† Onur Mutlu§ Yale N. Patt†

†Department of Electrical and Computer Engineering
The University of Texas at Austin
{ebrahimi, cjlee, patt}@ece.utexas.edu

§Computer Architecture Laboratory (CALCM)
Carnegie Mellon University
onur@cmu.edu

More on Source Throttling (II)

- Kevin Chang, Rachata Ausavarungnirun, Chris Fallin, and Onur Mutlu,
"HAT: Heterogeneous Adaptive Throttling for On-Chip Networks"

Proceedings of the 24th International Symposium on Computer Architecture and High Performance Computing (SBAC-PAD), New York, NY, October 2012. Slides (pptx) (pdf)

HAT: Heterogeneous Adaptive Throttling for On-Chip Networks

Kevin Kai-Wei Chang, Rachata Ausavarungnirun, Chris Fallin, Onur Mutlu
Carnegie Mellon University
`{kevincha, rachata, cfallin, onur}@cmu.edu`

More on Source Throttling (III)

- George Nychis, Chris Fallin, Thomas Moscibroda, Onur Mutlu, and Srinivasan Seshan,
**"On-Chip Networks from a Networking Perspective:
Congestion and Scalability in Many-core Interconnects"**
*Proceedings of the 2012 ACM SIGCOMM Conference
(SIGCOMM)*, Helsinki, Finland, August 2012. [Slides \(pptx\)](#)

On-Chip Networks from a Networking Perspective: Congestion and Scalability in Many-Core Interconnects

George Nychis[†], Chris Fallin[†], Thomas Moscibroda[§], Onur Mutlu[†], Srinivasan Seshan[†]

[†] Carnegie Mellon University
{gnychis,cfallin,onur,srini}@cmu.edu

[§] Microsoft Research Asia
moscitho@microsoft.com

Application-to-Core Mapping to Reduce Interference

- Reetuparna Das, Rachata Ausavarungnirun, Onur Mutlu, Akhilesh Kumar, and Mani Azimi,

"Application-to-Core Mapping Policies to Reduce Memory System Interference in Multi-Core Systems"

Proceedings of the 19th International Symposium on High-Performance Computer Architecture (HPCA), Shenzhen, China, February 2013.
Slides (pptx)

Application-to-Core Mapping Policies to Reduce Memory System Interference in Multi-Core Systems

Reetuparna Das* Rachata Ausavarungnirun† Onur Mutlu† Akhilesh Kumar‡ Mani Azimi‡
University of Michigan* Carnegie Mellon University† Intel Labs‡

Architecture-Aware DRM

- Hui Wang, Canturk Isci, Lavanya Subramanian, Jongmoo Choi, Depei Qian, and Onur Mutlu,

"A-DRM: Architecture-aware Distributed Resource Management of Virtualized Clusters"

Proceedings of the 11th ACM SIGPLAN/SIGOPS International Conference on Virtual Execution Environments (VEE), Istanbul, Turkey, March 2015.

[Slides (pptx) (pdf)]

A-DRM: Architecture-aware Distributed Resource Management of Virtualized Clusters

Hui Wang^{†*}, Canturk Isci[‡], Lavanya Subramanian*, Jongmoo Choi^{§*}, Depei Qian[†], Onur Mutlu*

[†]Beihang University, [‡]IBM Thomas J. Watson Research Center, ^{*}Carnegie Mellon University, [§]Dankook University

{hui.wang, depeiq}@buaa.edu.cn, canturk@us.ibm.com, {lsubrama, onur}@cmu.edu, choijm@dankook.ac.kr

Summary

Summary: Fundamental Interference Control Techniques

- Goal: to reduce/control interference
 - 1. Prioritization or request scheduling
 - 2. Data mapping to banks/channels/ranks
 - 3. Core/source throttling
 - 4. Application/thread scheduling

Best is to combine all. How would you do that?

Summary: Memory QoS Approaches and Techniques

- Approaches: Smart vs. dumb resources
 - Smart resources: QoS-aware memory scheduling
 - Dumb resources: Source throttling; channel partitioning
 - Both approaches are effective at reducing interference
 - No single best approach for all workloads
- Techniques: Request/thread scheduling, source throttling, memory partitioning
 - All approaches are effective at reducing interference
 - Can be applied at different levels: hardware vs. software
 - No single best technique for all workloads
- Combined approaches and techniques are the most powerful
 - Integrated Memory Channel Partitioning and Scheduling [MICRO'11]

Summary: Memory Interference and QoS

- QoS-unaware memory → uncontrollable and unpredictable system
 - Providing QoS awareness improves performance, predictability, fairness, and utilization of the memory system
 - Discussed many new techniques to:
 - Minimize memory interference
 - Provide predictable performance
 - Many new research ideas needed for integrated techniques and closing the interaction with software
-

What Did We Not Cover?

- Prefetch-aware shared resource management
- DRAM-controller co-design
- Cache interference management
- Interconnect interference management
- Write-read scheduling
- DRAM designs to reduce interference
- Interference & QoS in processing-in-memory
- ...

What the Future May Bring

- Memory QoS techniques for heterogeneous SoC systems
 - Many accelerators, processing in/near memory, better predictability, higher performance
- Combinations of memory QoS/performance techniques
 - E.g., data mapping and scheduling
- **Use of machine learning techniques to manage resources**
- Real prototypes

Computer Architecture

Lecture 13a: Memory Controllers: Service Quality

Ataberk Olgun

Prof. Onur Mutlu

ETH Zürich

Fall 2023

9 November 2023

Backup Slides

Memory Scheduling for Heterogeneous Systems

Lecture on Heterogeneous System Scheduling

SMS: Staged Memory Scheduling

The diagram illustrates the SMS (Staged Memory Scheduling) process across three stages:

- Stage 1 (Batch Formation):** Four cores (Core 1, Core 2, Core 3, Core 4) and a GPU feed into individual batch formation units. Each unit contains a yellow base layer and a colored (blue, red, green, purple) upper layer. A dashed line separates Stage 1 from Stage 2.
- Stage 2 (Batch Scheduler):** A central green "Batch Scheduler" unit receives inputs from Stage 1. It then outputs to four separate memory banks (Bank 1, Bank 2, Bank 3, Bank 4). A dashed line separates Stage 2 from Stage 3.
- Stage 3 (DRAM Command Scheduler):** The four memory banks (Bank 1, Bank 2, Bank 3, Bank 4) are shown, each containing a blue base layer and a colored (blue, green, red, yellow) upper layer. An arrow labeled "To DRAM" points from the bottom right of Stage 3 towards the right edge of the slide.

Below the diagram is a video player interface with controls, a timestamp (16:23 / 2:59:25), and a progress bar at 15%.

ETH ZENTRUM
Computer Arch - Lecture 13: Memory Interference and Quality of Service II (ETH Zürich, Spring 2020)
1,006 views • Nov 7, 2020

Onur Mutlu Lectures
19.8K subscribers

ANALYTICS EDIT VIDEO

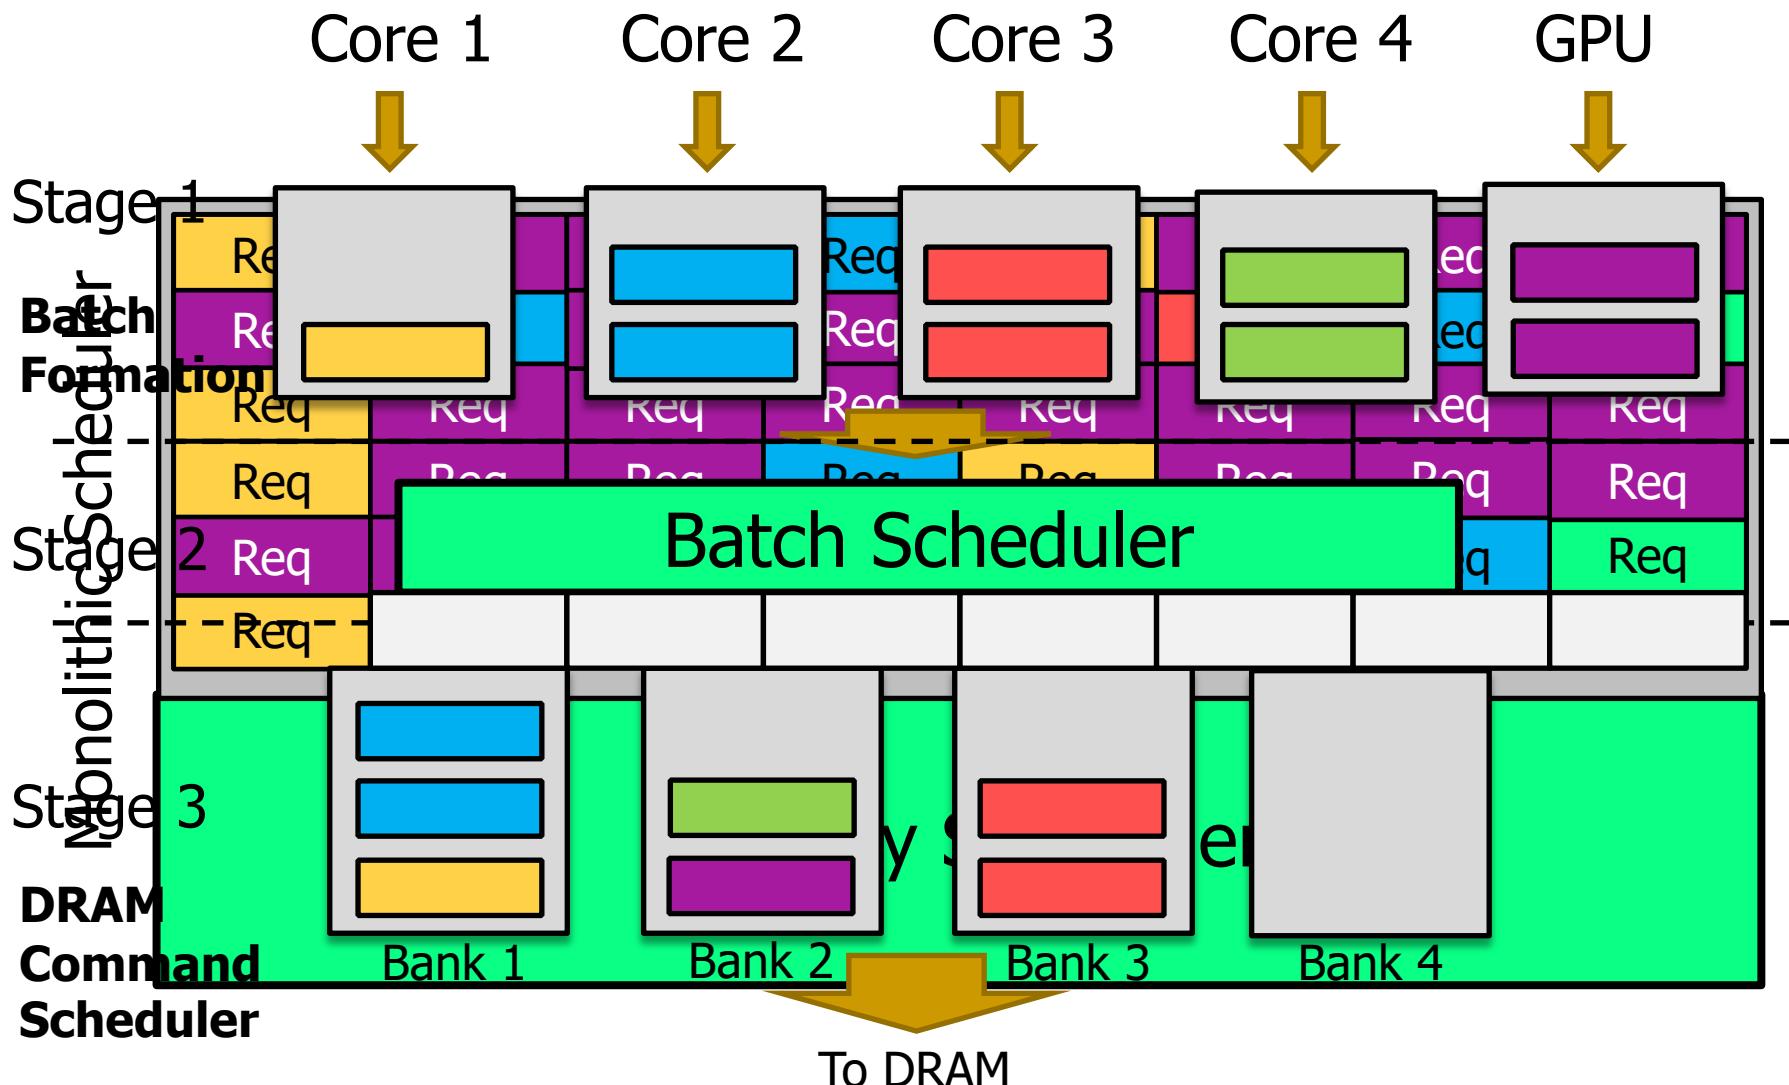
Staged Memory Scheduling

Rachata Ausavarungnirun, Kevin Chang, Lavanya Subramanian, Gabriel Loh, and Onur Mutlu,
**"Staged Memory Scheduling: Achieving High Performance
and Scalability in Heterogeneous Systems"**
39th International Symposium on Computer Architecture (ISCA),
Portland, OR, June 2012.

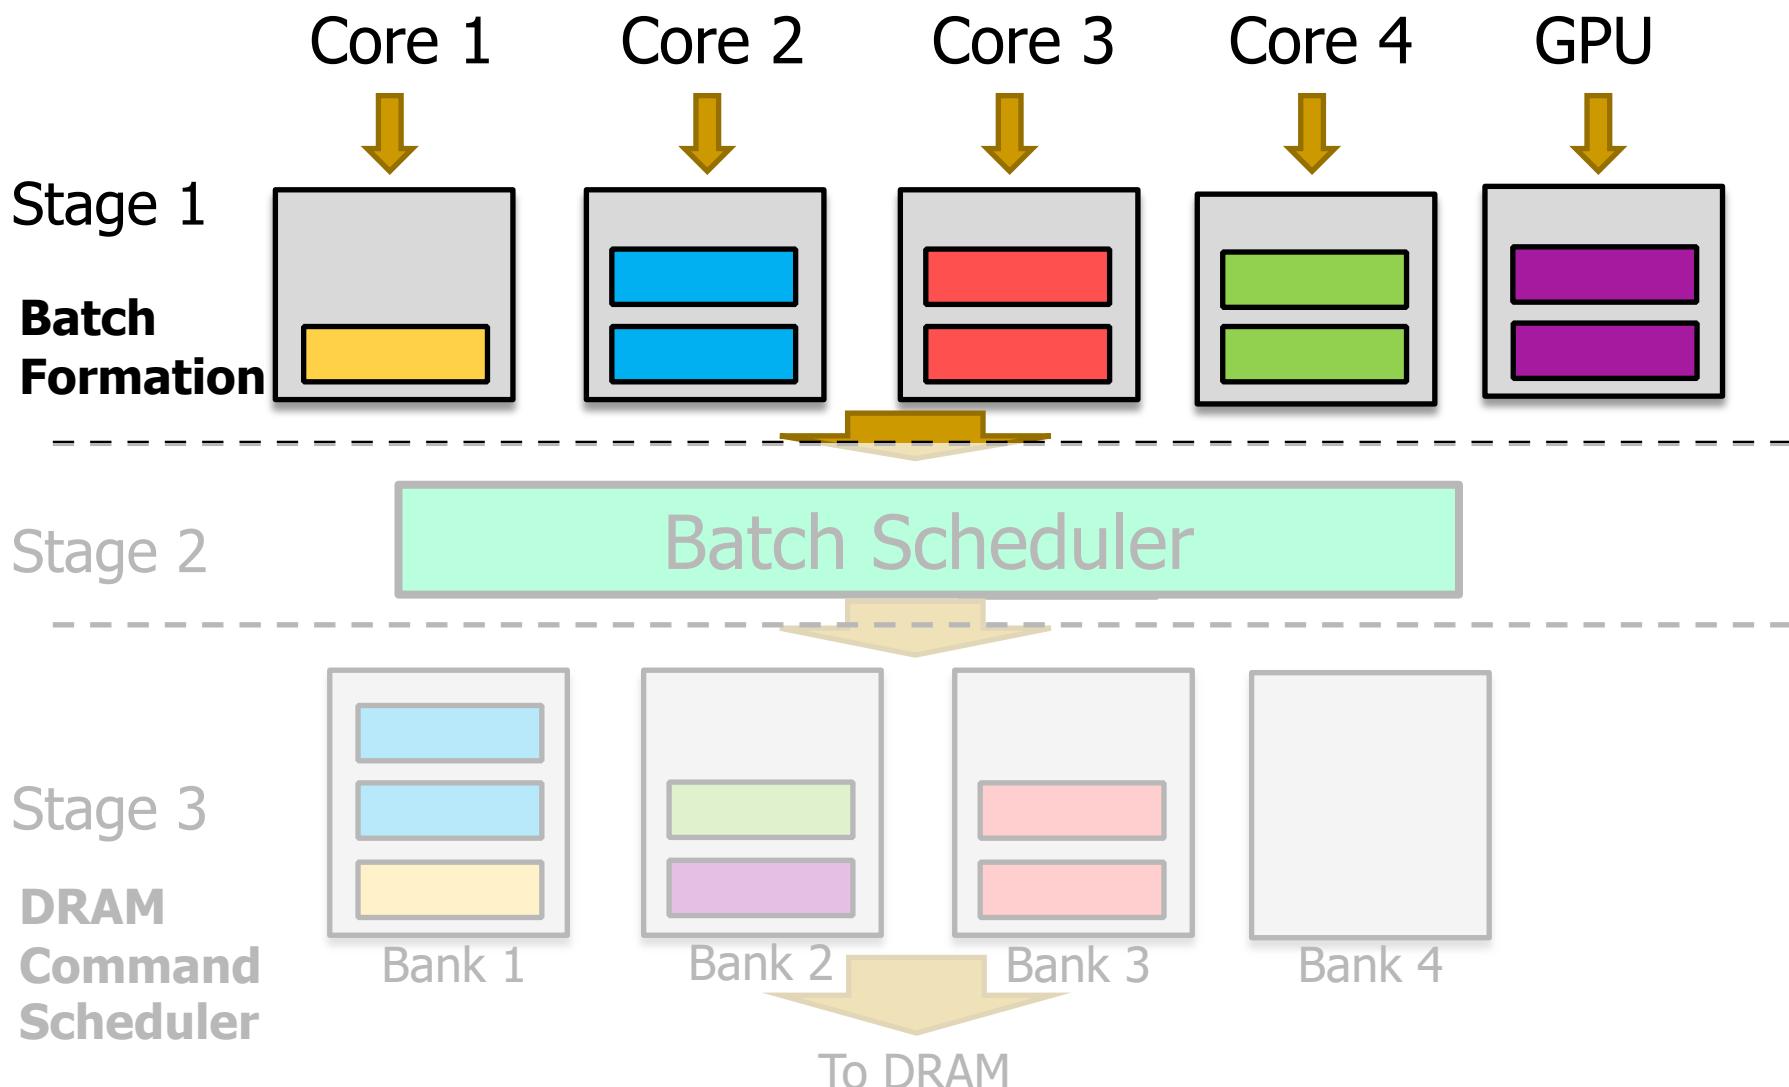
SMS: Executive Summary

- **Observation:** Heterogeneous CPU-GPU systems require memory schedulers with **large request buffers**
- **Problem:** Existing monolithic application-aware memory scheduler designs are **hard to scale** to large request buffer sizes
- **Solution:** Staged Memory Scheduling (SMS)
decomposes the memory controller into three simple stages:
 - 1) Batch formation: maintains row buffer locality
 - 2) Batch scheduler: reduces interference between applications
 - 3) DRAM command scheduler: issues requests to DRAM
- Compared to state-of-the-art memory schedulers:
 - SMS is significantly simpler and more scalable
 - SMS provides higher performance and fairness

SMS: Staged Memory Scheduling

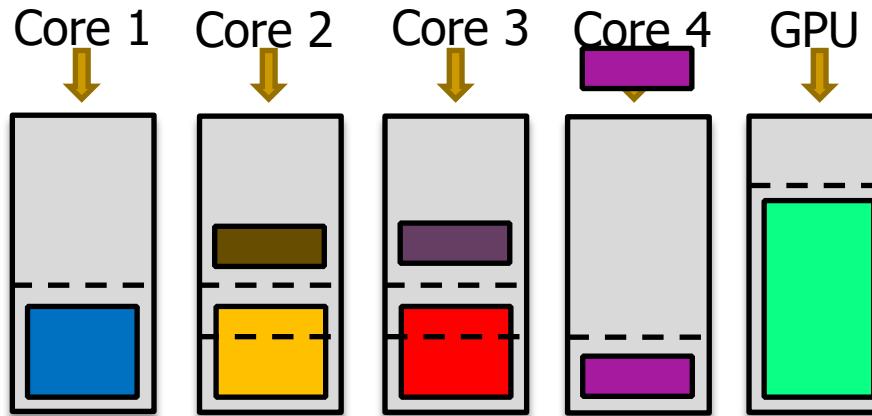


SMS: Staged Memory Scheduling



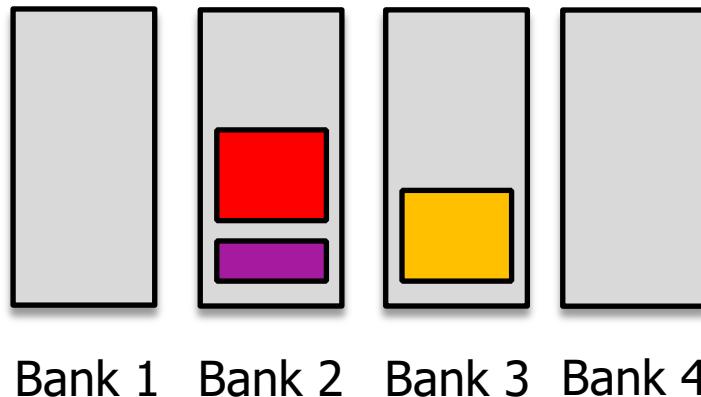
Putting Everything Together

Stage 1:
Batch
Formation



Stage 2: Batch Scheduler

Stage 3:
DRAM
Command
Scheduler



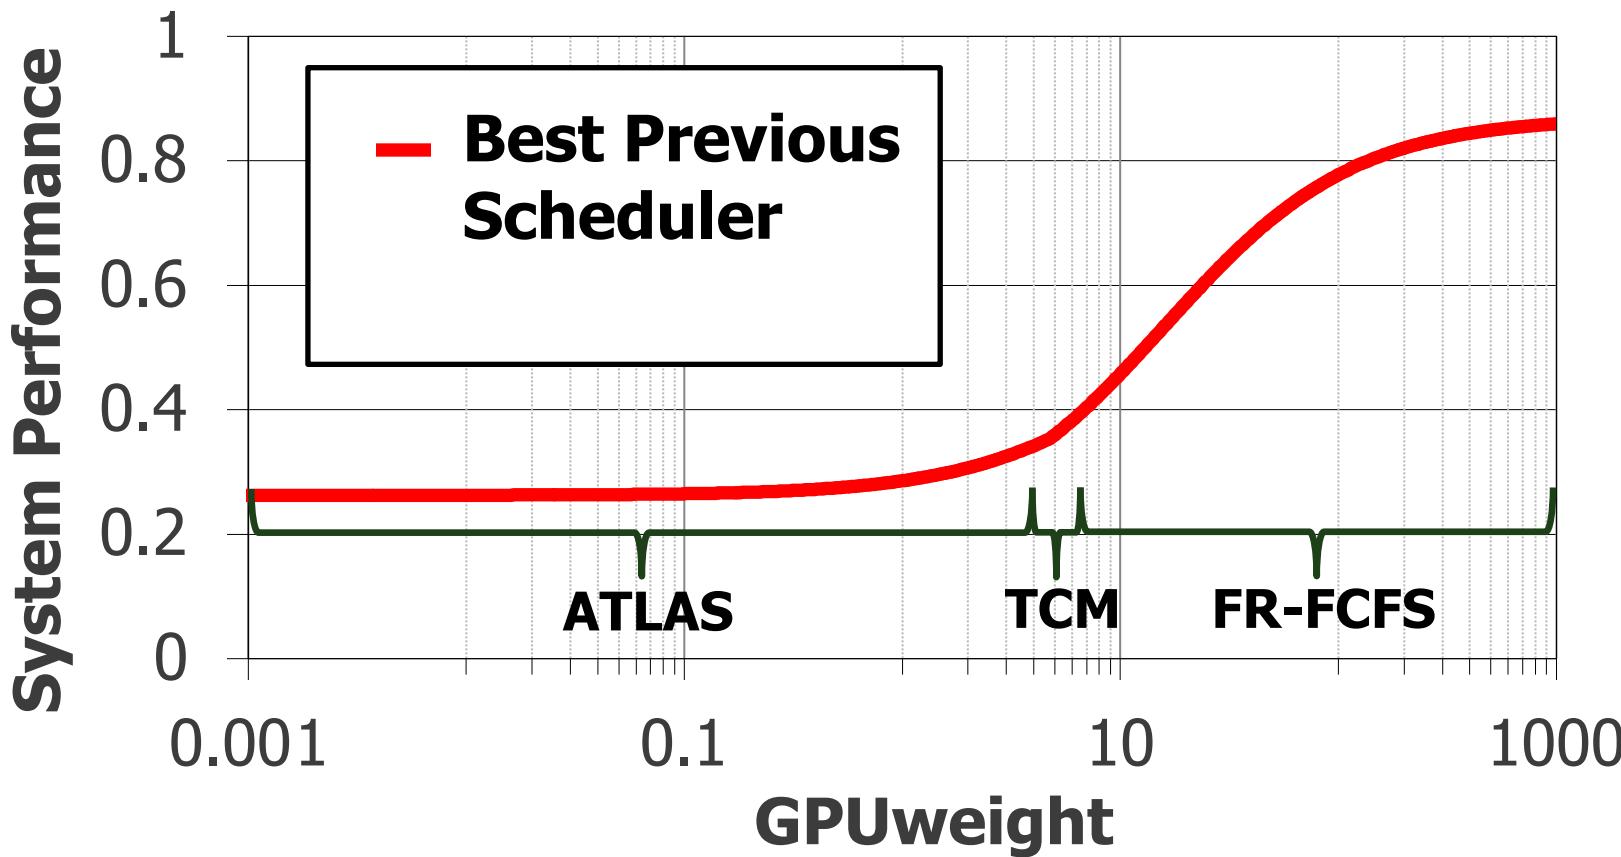
Current Batch
Scheduling
Policy
RR

Complexity

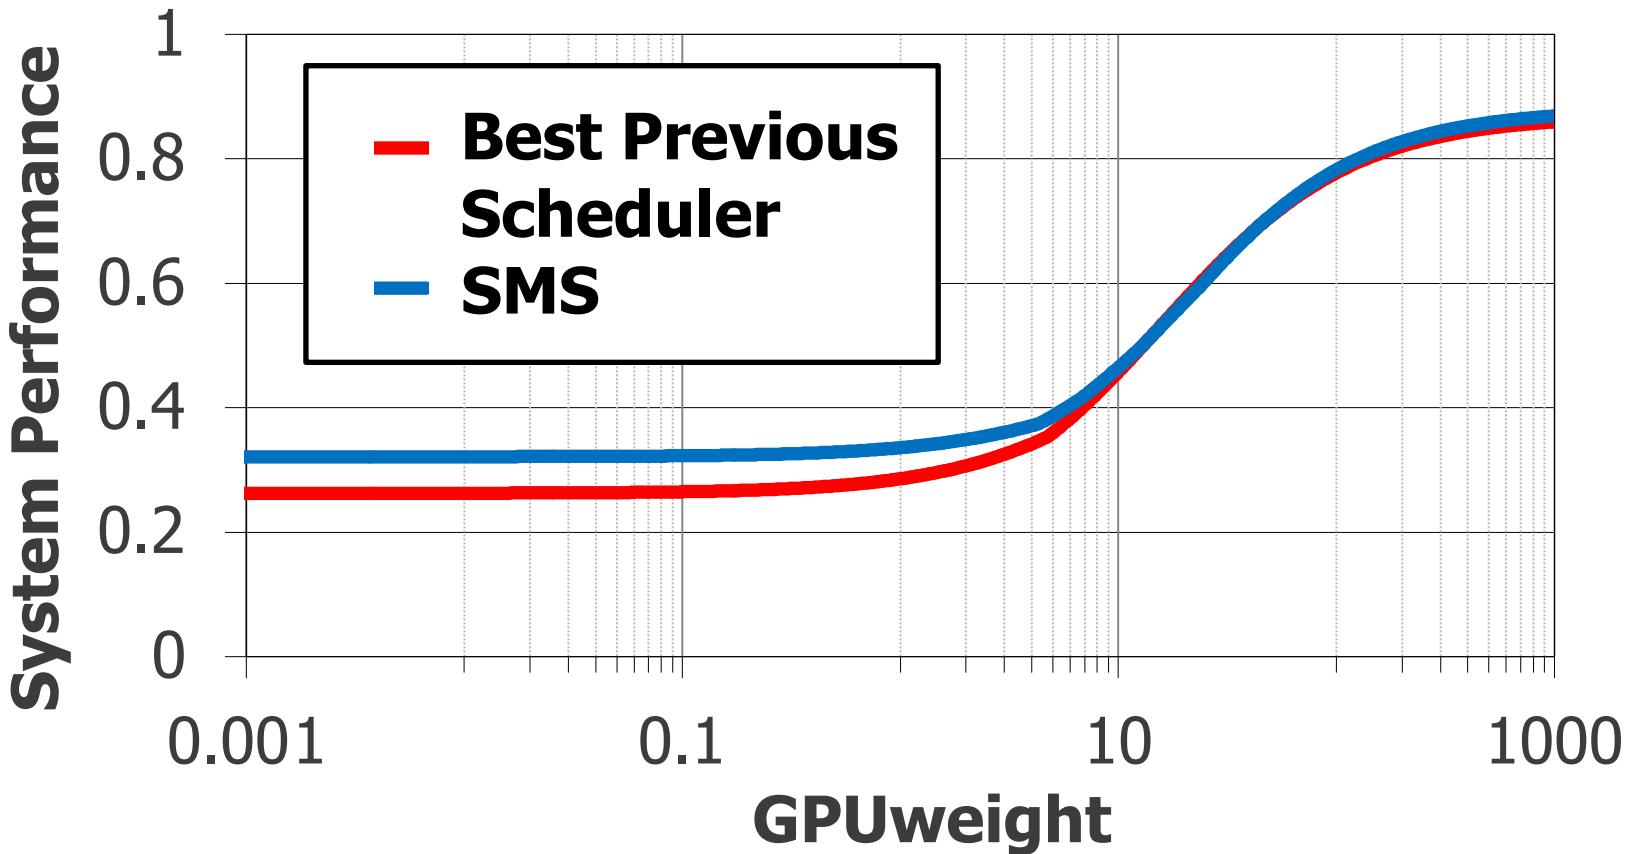
- Compared to a row hit first scheduler, SMS consumes*
 - 66% less area
 - 46% less static power

- Reduction comes from:
 - Monolithic scheduler → stages of simpler schedulers
 - Each stage has a simpler scheduler (considers fewer properties at a time to make the scheduling decision)
 - Each stage has simpler buffers (FIFO instead of out-of-order)
 - Each stage has a portion of the total buffer size (buffering is distributed across stages)

Performance at Different GPU Weights



Performance at Different GPU Weights



- At every GPU weight, SMS outperforms the best previous scheduling algorithm for that weight

More on SMS

- Rachata Ausavarungnirun, Kevin Chang, Lavanya Subramanian, Gabriel Loh, and Onur Mutlu,

"Staged Memory Scheduling: Achieving High Performance and Scalability in Heterogeneous Systems"

Proceedings of the 39th International Symposium on Computer Architecture (ISCA), Portland, OR, June 2012. Slides (pptx)

Staged Memory Scheduling: Achieving High Performance and Scalability in Heterogeneous Systems

Rachata Ausavarungnirun[†] Kevin Kai-Wei Chang[†] Lavanya Subramanian[†] Gabriel H. Loh[‡] Onur Mutlu[†]

[†]Carnegie Mellon University

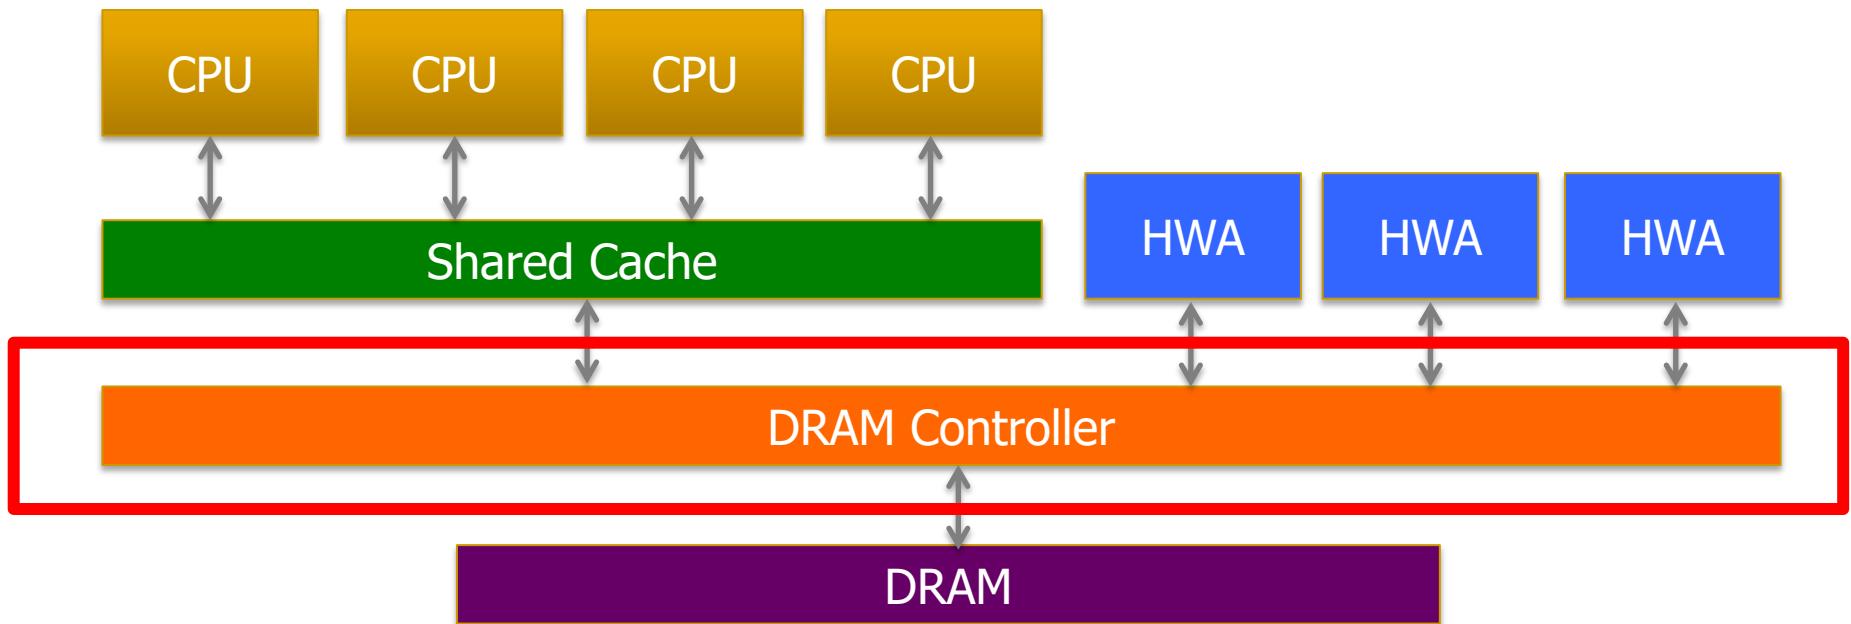
{rachata,kevincha,lsubrama,onur}@cmu.edu

[‡]Advanced Micro Devices, Inc.
gabe.loh@amd.com

DASH Memory Scheduler

[TACO 2016]

Current SoC Architectures



- Heterogeneous agents: CPUs and HWAs
 - HWA : Hardware Accelerator
- Main memory is shared by CPUs and HWAs → Interference
How to schedule memory requests from CPUs and HWAs
to mitigate interference?

DASH Scheduler: Executive Summary

- Problem: Hardware accelerators (HWAs) and CPUs share the same memory subsystem and interfere with each other in main memory
- Goal: Design a memory scheduler that improves CPU performance while meeting HWAs' deadlines
- Challenge: Different HWAs have different memory access characteristics and different deadlines, which current schedulers do not smoothly handle
 - Memory-intensive and long-deadline HWAs significantly degrade CPU performance *when they become high priority* (due to slow progress)
 - Short-deadline HWAs sometimes miss their deadlines *despite high priority*
- Solution: DASH Memory Scheduler
 - Prioritize HWAs over CPU anytime when the HWA is not making good progress
 - Application-aware scheduling for CPUs and HWAs
- Key Results:
 - 1) Improves CPU performance for a wide variety of workloads by 9.5%
 - 2) Meets 100% deadline met ratio for HWAs
- DASH source code freely available on our GitHub

Goal of Our Scheduler (DASH)

- **Goal:** Design a memory scheduler that
 - Meets GPU/accelerators' frame rates/deadlines *and*
 - Achieves high CPU performance
- **Basic Idea:**
 - *Different CPU applications and hardware accelerators have different memory requirements*
 - Track progress of different agents and prioritize accordingly

Key Observation: Distribute Priority for Accelerators

- GPU/accelerators need priority to meet deadlines
- Worst case prioritization not always the best
- Prioritize when they are **not** on track to meet a deadline

*Distributing priority over time mitigates impact
of accelerators on CPU cores' requests*

Key Observation: Not All Accelerators are Equal

- **Long-deadline** accelerators are more likely to **meet** their deadlines
- **Short-deadline** accelerators are more likely to **miss** their deadlines

*Schedule short-deadline accelerators
based on worst-case memory access time*

Key Observation: Not All CPU cores are Equal

- **Memory-intensive** cores are much **less vulnerable** to interference
- **Memory non-intensive** cores are much **more vulnerable** to interference

*Prioritize accelerators over memory-intensive cores
to ensure accelerators do not become urgent*

DASH Summary: Key Ideas and Results

- *Distribute priority for HWAs*
- *Prioritize HWAs over memory-intensive CPU cores even when not urgent*
- *Prioritize short-deadline-period HWAs based on worst case estimates*

Improves CPU performance by 7-21%

Meets (almost) 100% of deadlines for HWAs

DASH: Scheduling Policy

- DASH scheduling policy
 1. Short-deadline-period HWAs with high priority
 2. Long-deadline-period HWAs with high priority
 3. Memory non-intensive CPU applications
 4. Long-deadline-period HWAs with low priority
 5. Memory-intensive CPU applications
 6. Short-deadline-period HWAs with low priority
- } Switch
probabilistically

More on DASH

- Hiroyuki Usui, Lavanya Subramanian, Kevin Kai-Wei Chang, and Onur Mutlu,

"DASH: Deadline-Aware High-Performance Memory Scheduler for Heterogeneous Systems with Hardware Accelerators"

ACM Transactions on Architecture and Code Optimization (TACO),
Vol. 12, January 2016.

Presented at the 11th HiPEAC Conference, Prague, Czech Republic,
January 2016.

[Slides (pptx) (pdf)]

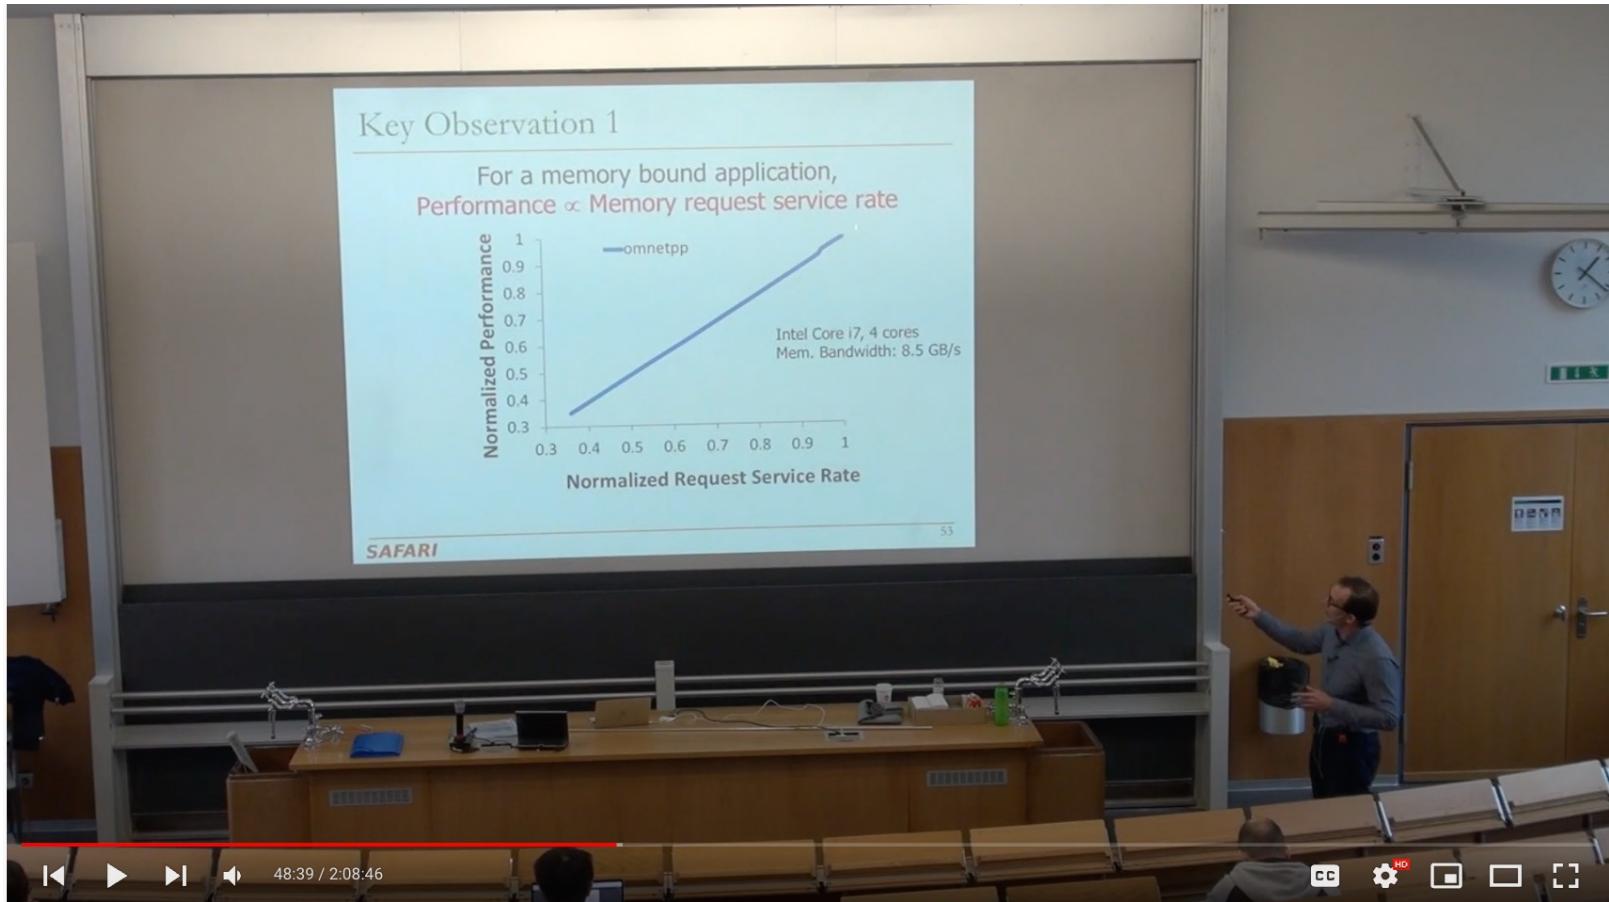
[Source Code]

DASH: Deadline-Aware High-Performance Memory Scheduler for Heterogeneous Systems with Hardware Accelerators

HIROYUKI USUI, LAVANYA SUBRAMANIAN, KEVIN KAI-WEI CHANG,
and ONUR MUTLU, Carnegie Mellon University

Predictable Performance: Strong Memory Service Guarantees

Lecture on Predictable Performance



ETH ZÜRICH HAUPTGEBÄUDE

Computer Architecture - Lecture 17: Memory Interference and QoS II (ETH Zürich, Fall 2018)

274 views • Nov 23, 2018

1 like 4 dislikes SHARE SAVE ...



Onur Mutlu Lectures
19.8K subscribers

ANALYTICS

EDIT VIDEO

Lecture on Predictable Performance

The slide title is "Effectiveness of MISE in Enforcing QoS" and it states "Across 3000 data points". It contains a 2x2 table:

| | Predicted Met | Predicted Not Met |
|-------------------|---------------|-------------------|
| QoS Bound Met | 78.8% | 2.1% |
| QoS Bound Not Met | 2.2% | 16.9% |

A callout highlights the "Predicted Met" cell for "QoS Bound Met" at 78.8%. Another callout highlights the "Predicted Not Met" cell for "QoS Bound Not Met" at 16.9%.

In a callout box, the text reads: "MISE-QoS correctly predicts whether or not the bound is met for 95.7% of workloads".

The slide footer includes the word "SAFARI" and the number 161. The video player interface shows a progress bar at 25:29 / 1:11:14, and various control icons like back, forward, volume, and settings.

Memory Systems - Lecture 6.4: Memory Interface and QoS (Technion, Summer 2018)

163 views • Oct 12, 2018

like 5 dislike 0 share save ...

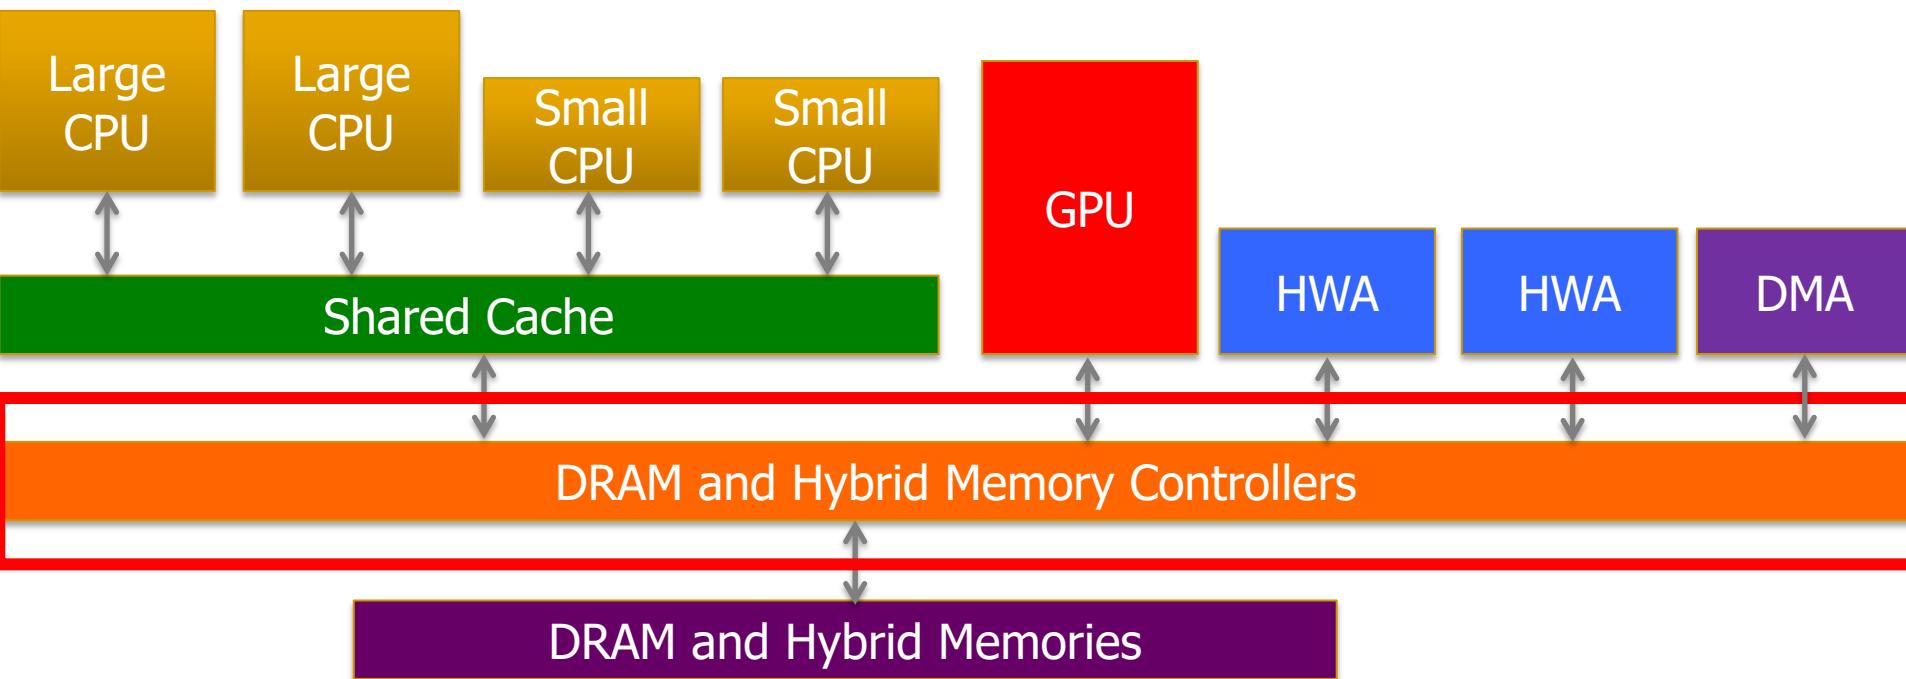


Onur Mutlu Lectures
19.8K subscribers

ANALYTICS

EDIT VIDEO

Goal: Predictable Performance in Complex Systems



- Heterogeneous agents: CPUs, GPUs, HWAs, DMA engine, ...
- Memory resources shared by CPUs/GPUs/HWAs → Interference

How to allocate resources to heterogeneous agents
to mitigate interference and provide predictable performance?

Strong Memory Service Guarantees

- Goal: Satisfy performance/SLA requirements in the presence of shared main memory, heterogeneous agents, and hybrid memory/storage
- Approach:
 - Develop techniques/models to accurately estimate the performance loss of an application/agent in the presence of resource sharing
 - Develop mechanisms (hardware and software) to enable the resource partitioning/prioritization needed to achieve the required performance levels for all applications
 - All the while providing high system performance
- Subramanian et al., "MISE: Providing Performance Predictability and Improving Fairness in Shared Main Memory Systems," HPCA 2013.
- Subramanian et al., "The Application Slowdown Model," MICRO 2015.

Predictable Performance Readings (I)

Eiman Ebrahimi, Chang Joo Lee, Onur Mutlu, and Yale N. Patt,

"Fairness via Source Throttling: A Configurable and High-Performance Fairness Substrate for Multi-Core Memory Systems"

Proceedings of the 15th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), pages 335-346, Pittsburgh, PA, March 2010. [Slides \(pdf\)](#)

Best paper award.

Fairness via Source Throttling: A Configurable and High-Performance Fairness Substrate for Multi-Core Memory Systems

Eiman Ebrahimi† Chang Joo Lee† Onur Mutlu§ Yale N. Patt†

†Department of Electrical and Computer Engineering
The University of Texas at Austin
{ebrahimi, cjlee, patt}@ece.utexas.edu

§Computer Architecture Laboratory (CALCM)
Carnegie Mellon University
onur@cmu.edu

Predictable Performance Readings (II)

- Lavanya Subramanian, Vivek Seshadri, Yoongu Kim, Ben Jaiyen, and Onur Mutlu,

"MISE: Providing Performance Predictability and Improving Fairness in Shared Main Memory Systems"

Proceedings of the 19th International Symposium on High-Performance Computer Architecture (HPCA), Shenzhen, China, February 2013. [Slides \(pptx\)](#)

MISE: Providing Performance Predictability and Improving Fairness in Shared Main Memory Systems

Lavanya Subramanian

Vivek Seshadri

Yoongu Kim

Ben Jaiyen

Onur Mutlu

Carnegie Mellon University

Predictable Performance Readings (III)

- Lavanya Subramanian, Vivek Seshadri, Arnab Ghosh, Samira Khan, and Onur Mutlu,

"The Application Slowdown Model: Quantifying and Controlling the Impact of Inter-Application Interference at Shared Caches and Main Memory"

Proceedings of the 48th International Symposium on Microarchitecture (MICRO), Waikiki, Hawaii, USA, December 2015.

[[Slides \(pptx\)](#) ([pdf](#))] [[Lightning Session Slides \(pptx\)](#) ([pdf](#))] [[Poster \(pptx\)](#) ([pdf](#))]

[[Source Code](#)]

The Application Slowdown Model: Quantifying and Controlling the Impact of Inter-Application Interference at Shared Caches and Main Memory

Lavanya Subramanian*§ Vivek Seshadri* Arnab Ghosh*†
Samira Khan*‡ Onur Mutlu*

*Carnegie Mellon University §Intel Labs †IIT Kanpur ‡University of Virginia

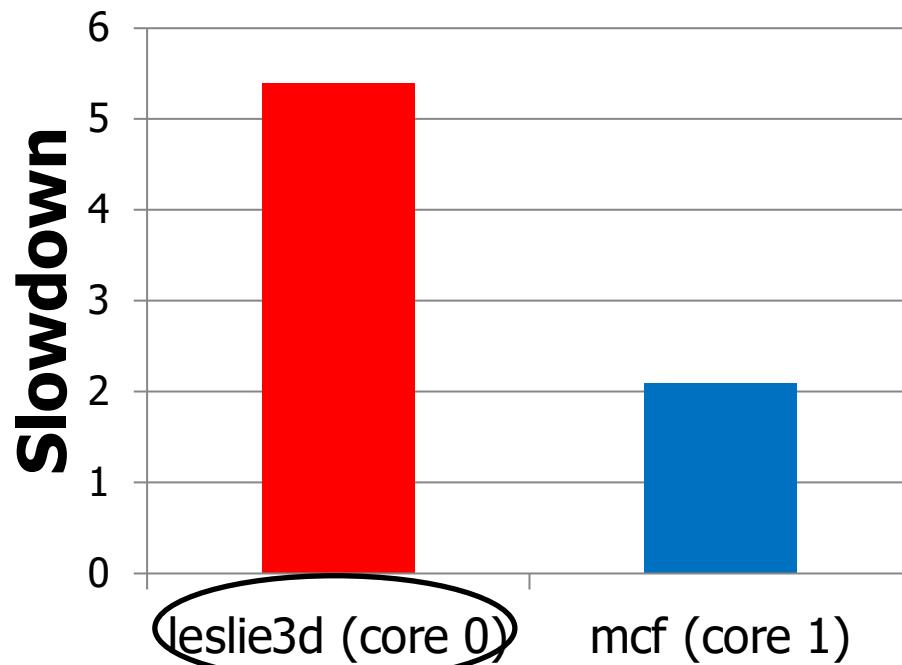
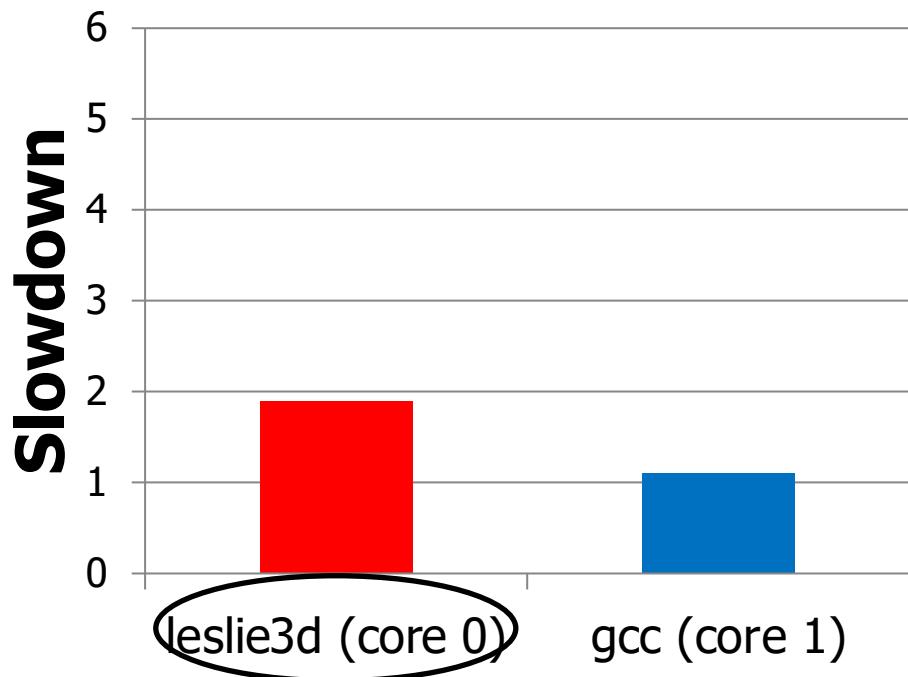
MISE: Providing Performance Predictability in Shared Main Memory Systems

Lavanya Subramanian, Vivek Seshadri,
Yoongu Kim, Ben Jaiyen, Onur Mutlu

SAFARI

Carnegie Mellon

Unpredictable Application Slowdowns



An application's performance depends on
which application it is running with

Need for Predictable Performance

- There is a need for predictable performance
 - When multiple applications share resources
 - Especially if some applications require performance guarantees

Our Goal: Predictable performance
in the presence of memory interference

- Example 2: In server systems
 - Different users' jobs consolidated onto the same server
 - Need to provide bounded slowdowns to critical jobs

Outline

1. Estimate Slowdown

2. Control Slowdown

Outline

1. Estimate Slowdown

- ❑ Key Observations
- ❑ Implementation
- ❑ MISE Model: Putting it All Together
- ❑ Evaluating the Model

2. Control Slowdown

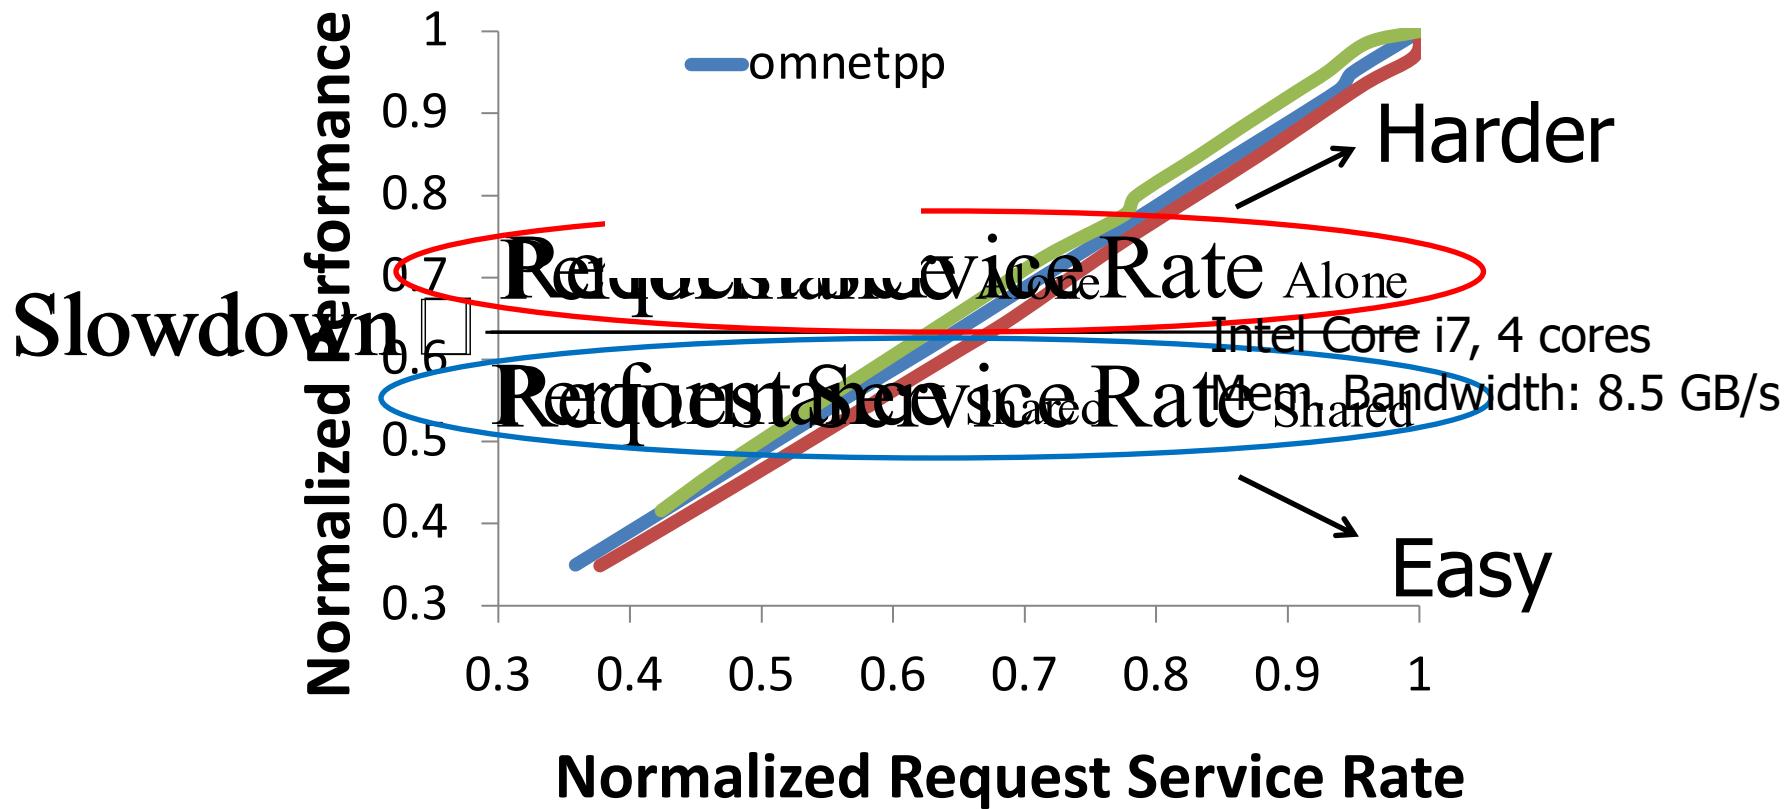
- ❑ Providing Soft Slowdown Guarantees
- ❑ Minimizing Maximum Slowdown

Slowdown: Definition

Slowdown $\square \frac{\text{Performance Alone}}{\text{Performance Shared}}$

Key Observation 1

For a memory bound application,
Performance \propto Memory request service rate



Key Observation 2

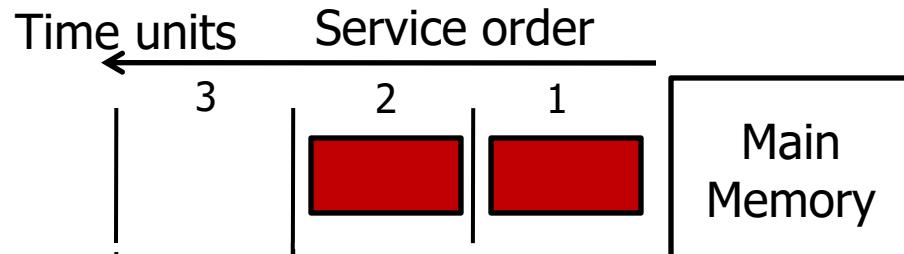
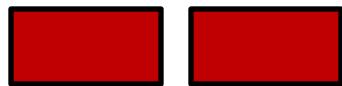
Request Service Rate _{Alone} (RSR_{Alone}) of an application can be estimated by giving the application highest priority in accessing memory

Highest priority → Little interference
(almost as if the application were run alone)

Key Observation 2

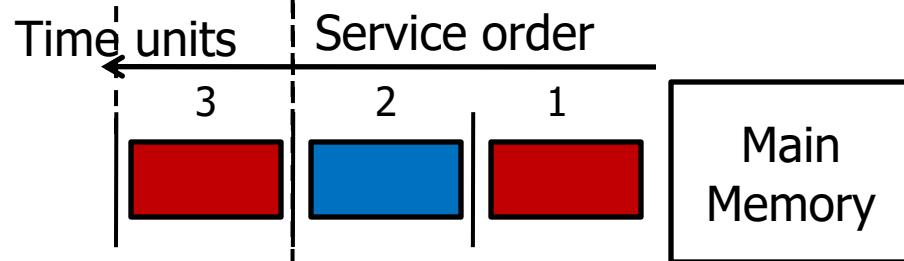
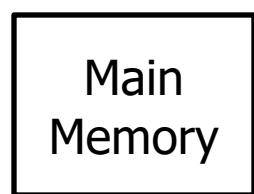
1. Run alone

Request Buffer State



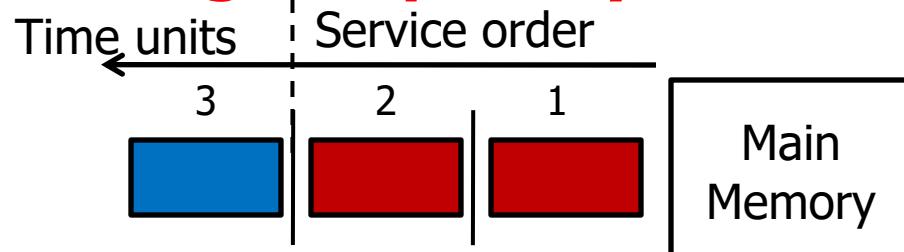
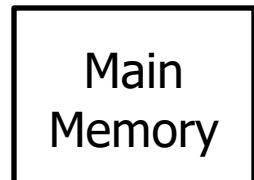
2. Run with another application

Request Buffer State



3. Run with another application: highest priority

Request Buffer State

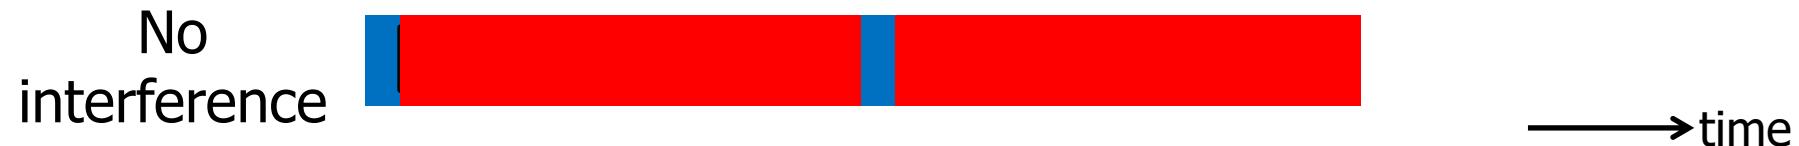
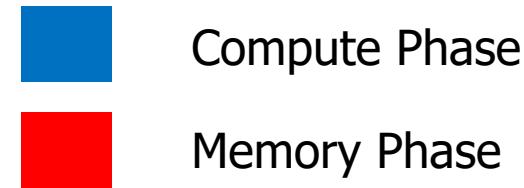


Memory Interference-induced Slowdown Estimation (MISE) model for **memory bound** applications

$$\text{Slowdown} \square \frac{\text{Request Service Rate Alone (RSR}_{\text{Alone}}\text{)}}{\text{Request Service Rate Shared (RSR}_{\text{Shared}}\text{)}}$$

Key Observation 3

■ Memory-bound application



Memory phase slowdown dominates overall slowdown

Key Observation 3

- Non-memory-bound application

Memory Interference-induced Slowdown Estimation (MISE) model for **non-memory bound** applications

$$\text{Slowdown} \square (1 - \alpha) \square \alpha \frac{\text{RSR}_{\text{Alone}}}{\text{RSR}_{\text{Shared}}}$$

Only memory fraction (α) slows down with interference

Outline

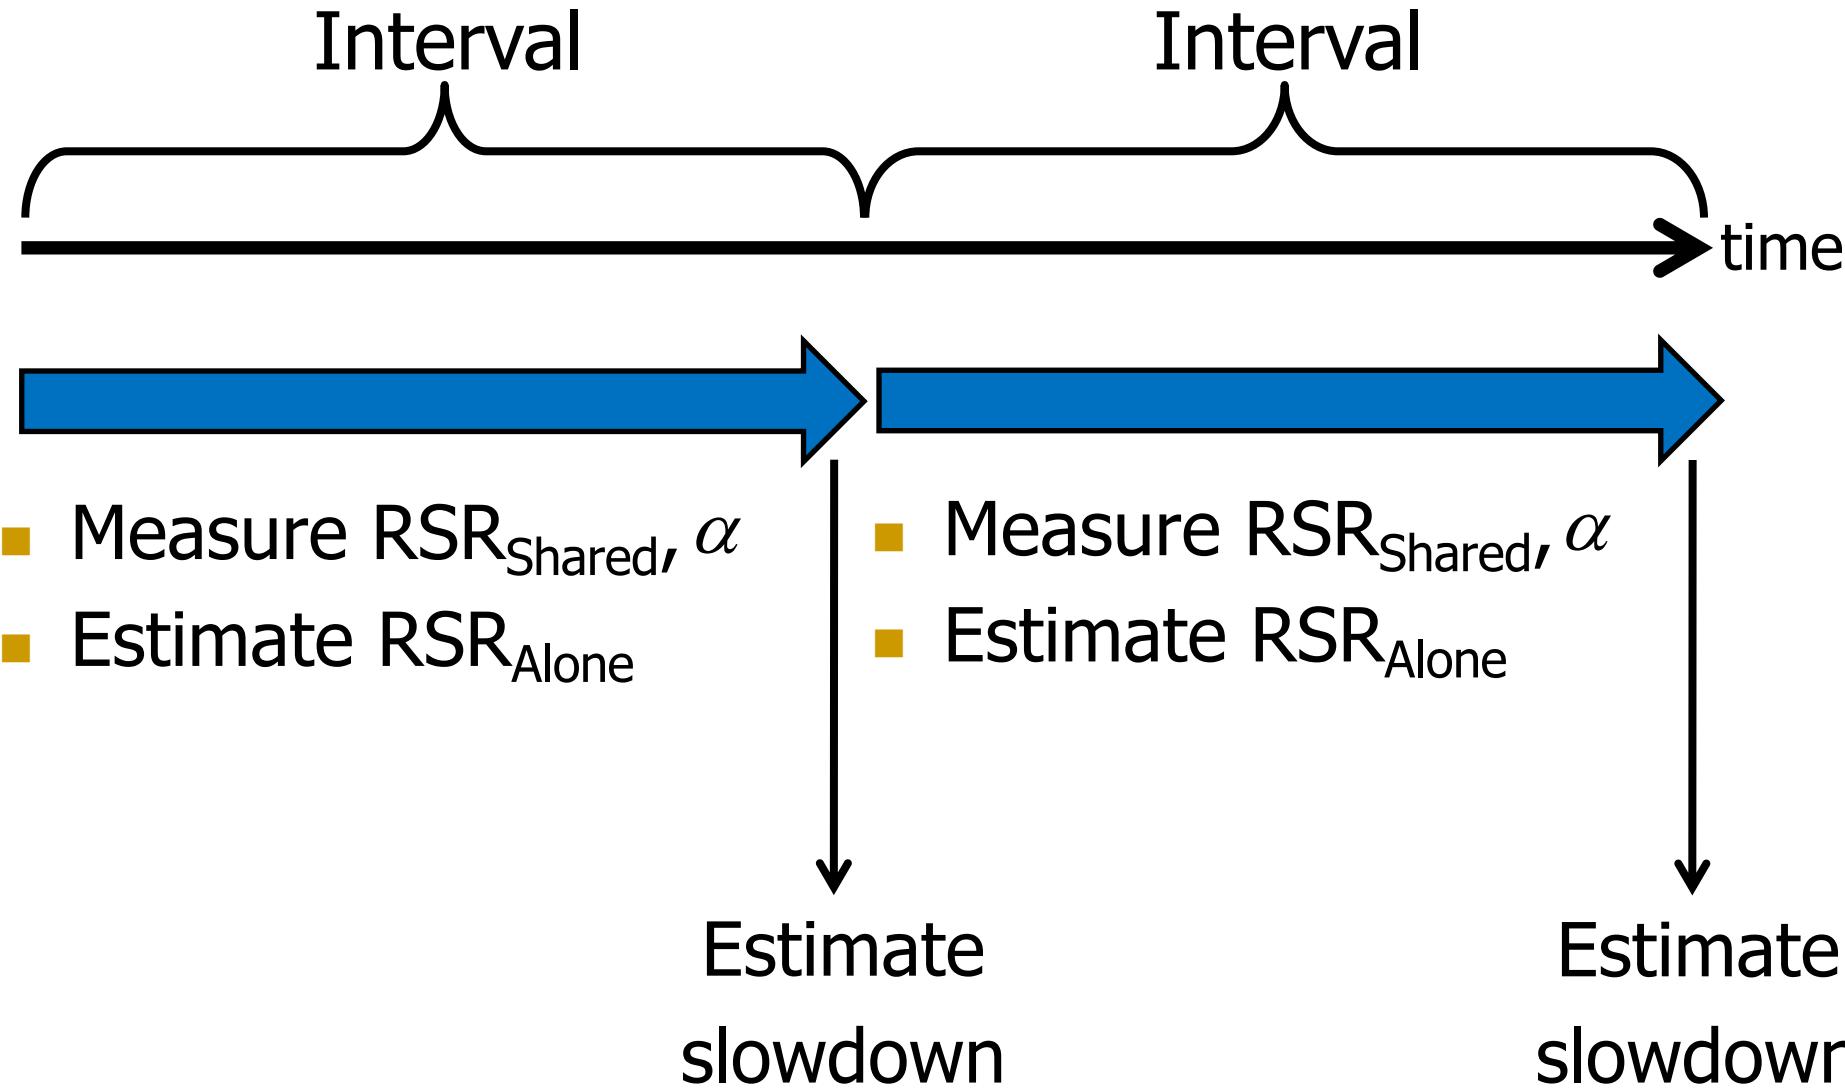
1. Estimate Slowdown

- ❑ Key Observations
- ❑ Implementation
- ❑ MISE Model: Putting it All Together
- ❑ Evaluating the Model

2. Control Slowdown

- ❑ Providing Soft Slowdown Guarantees
- ❑ Minimizing Maximum Slowdown

Interval Based Operation



Measuring RSR_{Shared} and α

- Request Service Rate_{Shared} (RSR_{Shared})
 - Per-core counter to track number of requests serviced
 - At the end of each interval, measure

$$\text{RSR}_{\text{Shared}} \square \frac{\text{Number of Requests Serviced}}{\text{Interval Length}}$$

- Memory Phase Fraction (α)
 - Count number of stall cycles at the core
 - Compute fraction of cycles stalled for memory

Estimating Request Service Rate $_{\text{Alone}}$ ($\text{RSR}_{\text{Alone}}$)

- Divide each interval into shorter epochs
- At the beginning of each epoch
 - Memory controller randomly picks an application as the highest priority application

Goal: Estimate $\text{RSR}_{\text{Alone}}$

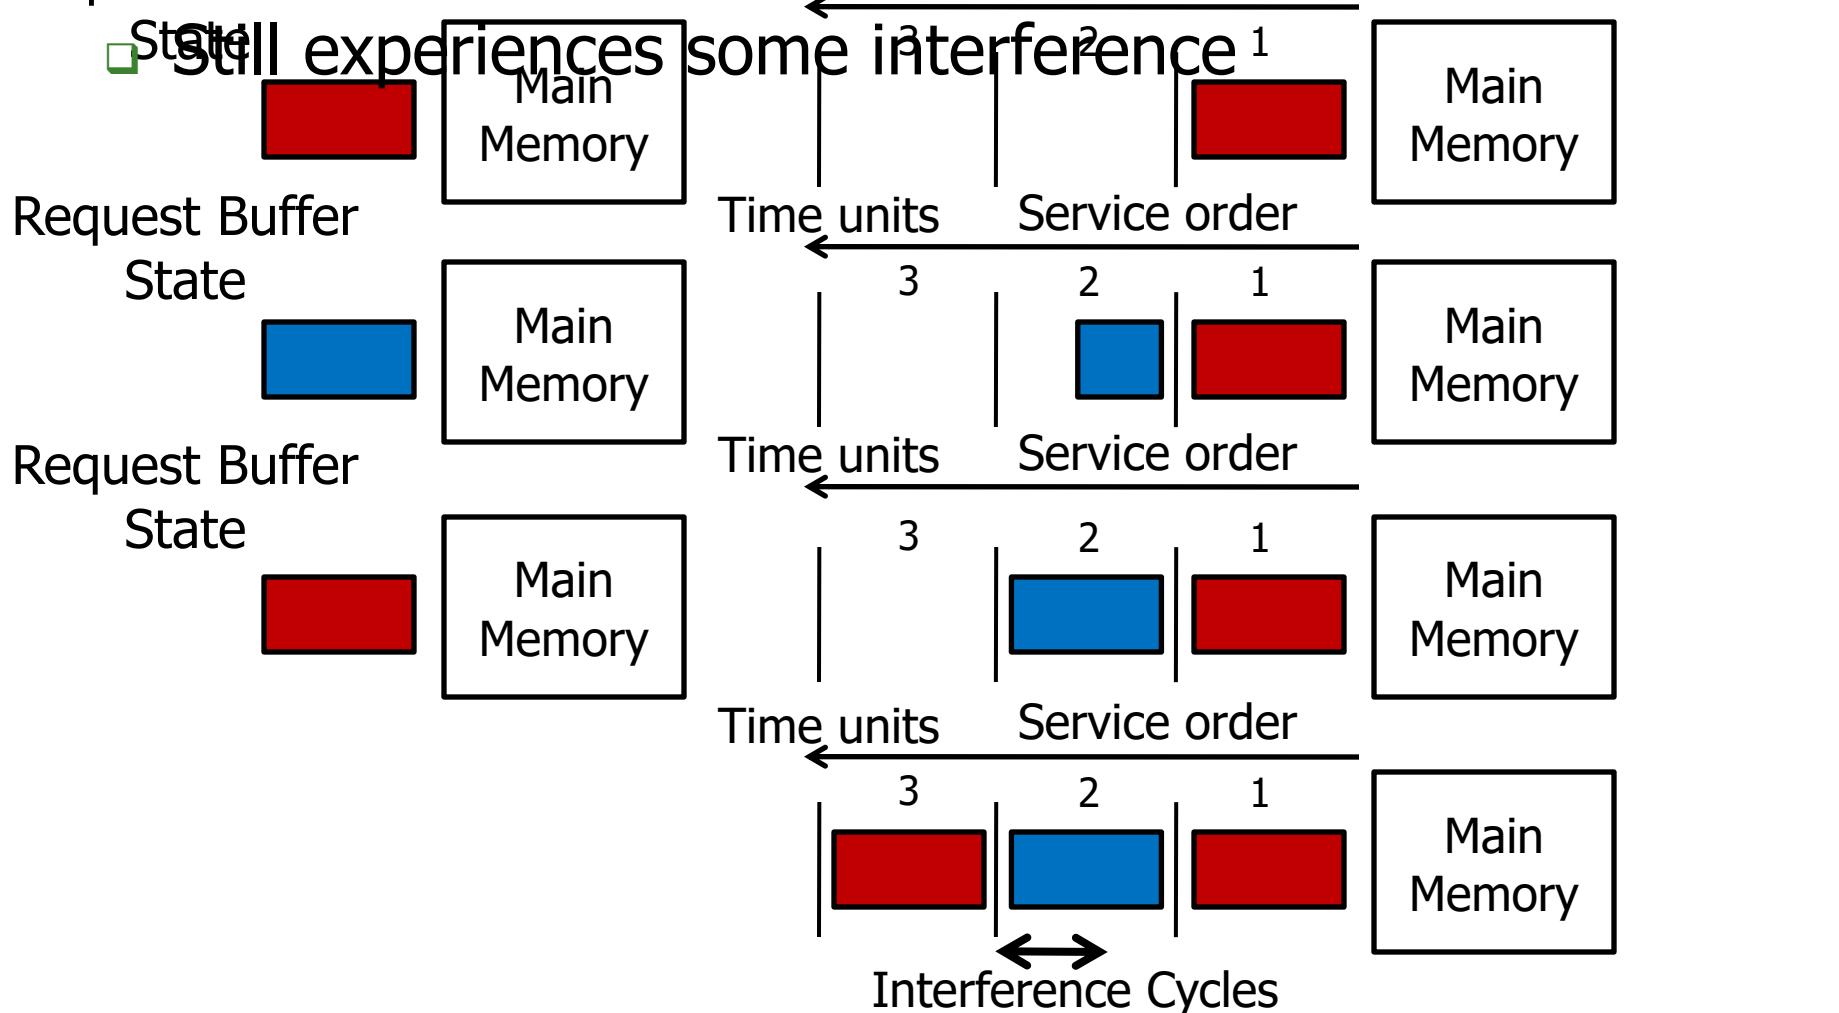
How: Periodically give each application

- At the end of an interval, for each application, estimate

$$\text{RSR}_{\text{Alone}} \square \frac{\text{Number of Requests During High Priority Epochs}}{\text{Number of Cycles Application Given High Priority}}$$

Inaccuracy in Estimating RSR_{Alone}

- When an application has highest priority



Accounting for Interference in RSR_{Alone} Estimation

- Solution: Determine and remove interference cycles from RSR_{Alone} calculation

$$RSR_{Alone} \square \frac{\text{Number of Requests During High Priority Epochs}}{\text{Number of Cycles Application Given High Priority} - \text{Interference Cycles}}$$

- A cycle is an interference cycle if
 - a request from the highest priority application is waiting in the request buffer *and*
 - another application's request was issued previously

Outline

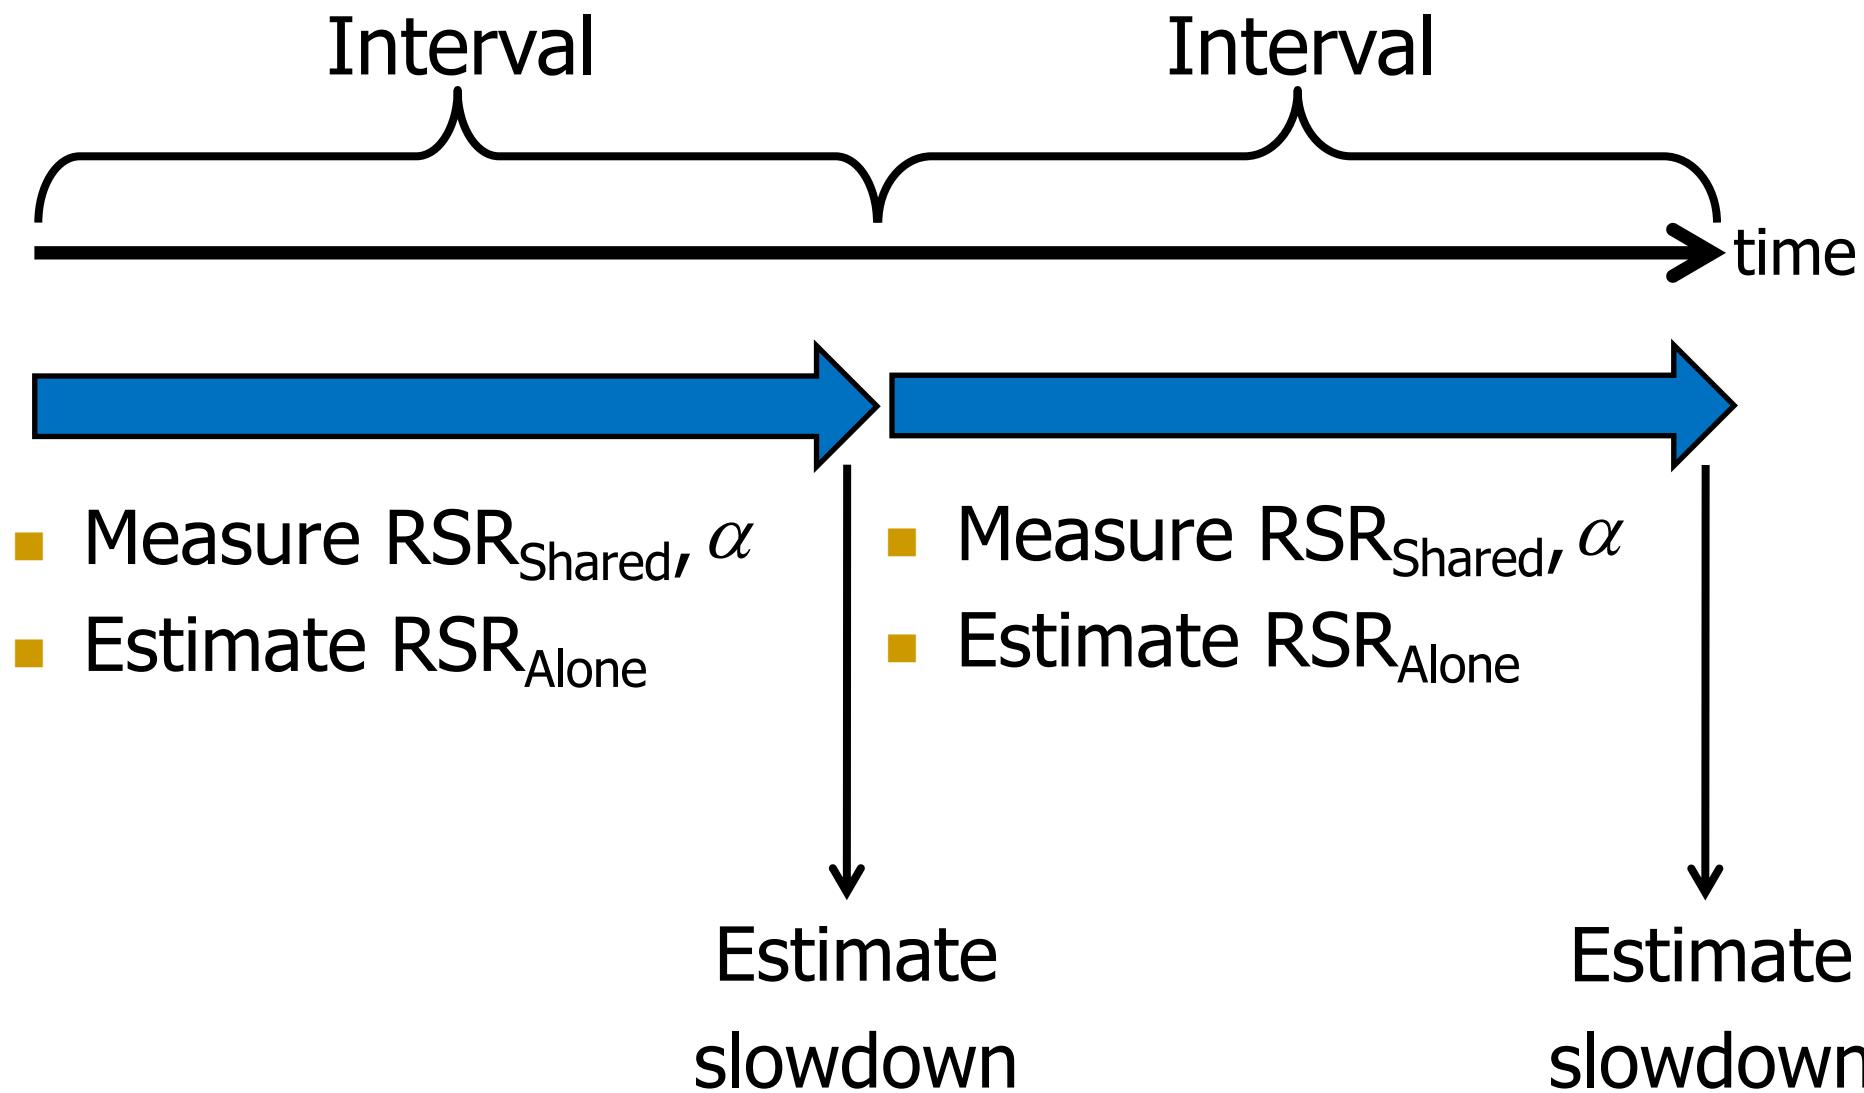
1. Estimate Slowdown

- ❑ Key Observations
- ❑ Implementation
- ❑ MISE Model: Putting it All Together
- ❑ Evaluating the Model

2. Control Slowdown

- ❑ Providing Soft Slowdown Guarantees
- ❑ Minimizing Maximum Slowdown

MISE Model: Putting it All Together



Outline

1. Estimate Slowdown

- ❑ Key Observations
- ❑ Implementation
- ❑ MISE Model: Putting it All Together
- ❑ Evaluating the Model

2. Control Slowdown

- ❑ Providing Soft Slowdown Guarantees
- ❑ Minimizing Maximum Slowdown

Previous Work on Slowdown Estimation

- Previous work on slowdown estimation
 - **STFM** (Stall Time Fair Memory) Scheduling [Mutlu+, MICRO '07]
 - **FST** (Fairness via Source Throttling) [Ebrahimi+, ASPLOS '10]
 - **Per-thread Cycle Accounting** [Du Bois+, HiPEAC '13]
- Basic Idea:



Count number of cycles application receives interference

Two Major Advantages of MISE Over STFM

- Advantage 1:
 - STFM estimates alone performance while an application is receiving interference → Hard
 - MISE estimates alone performance while giving an application the highest priority → Easier

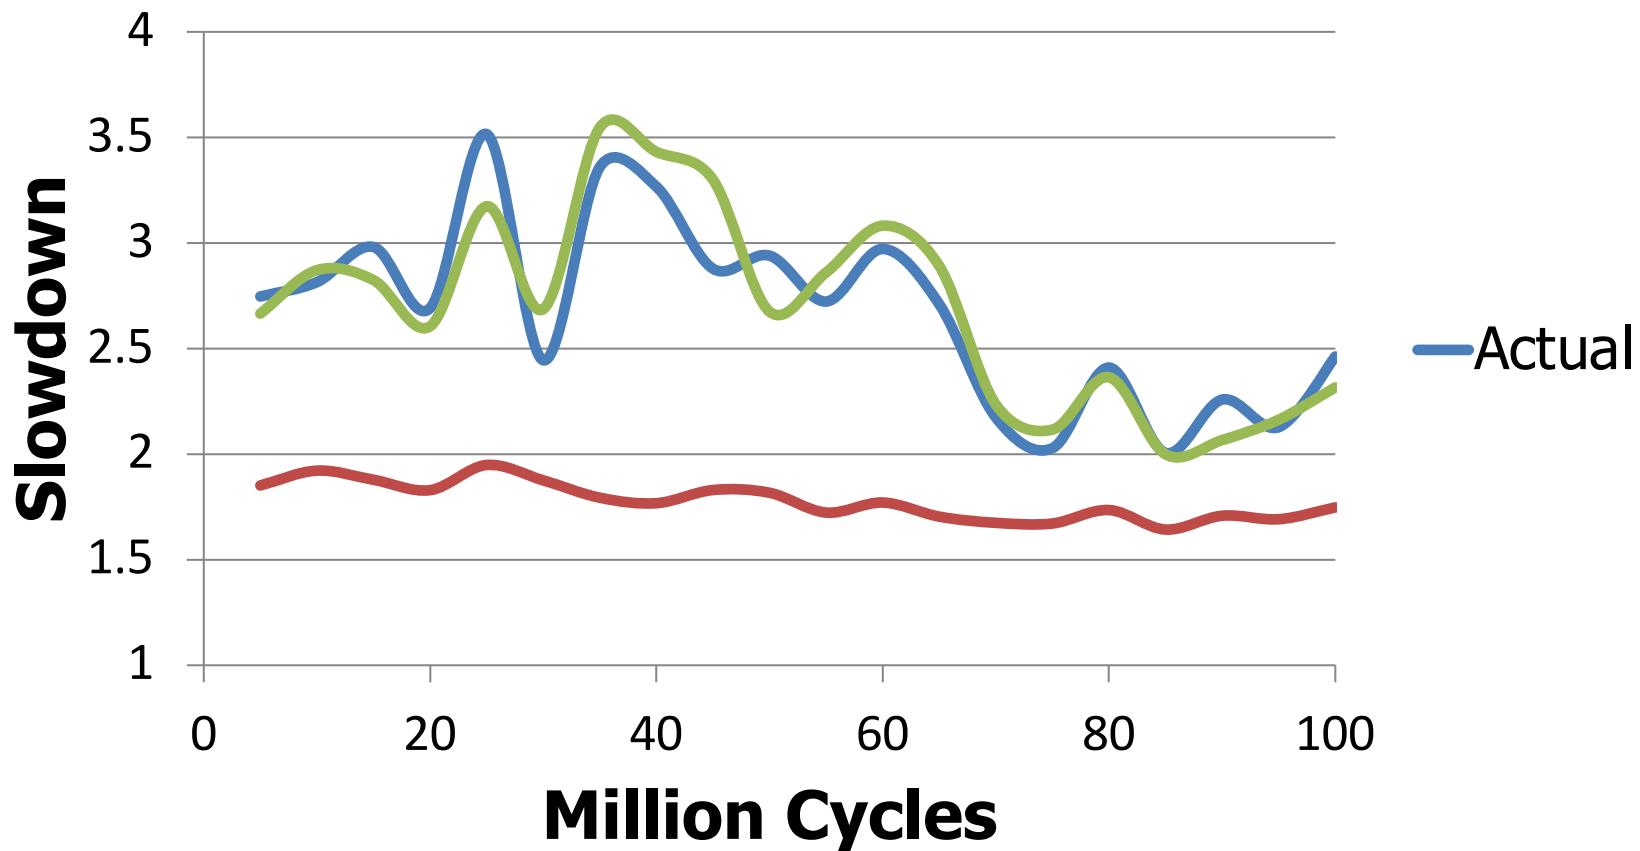
- Advantage 2:
 - STFM does not take into account compute phase for non-memory-bound applications
 - MISE accounts for compute phase → Better accuracy

Methodology

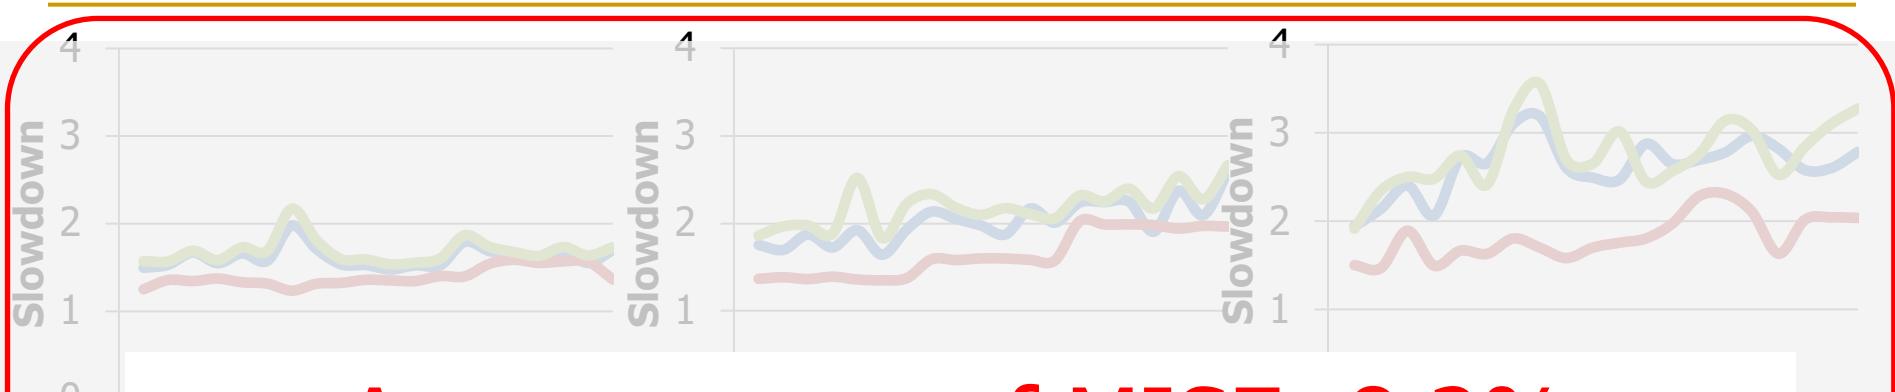
- Configuration of our simulated system
 - 4 cores
 - 1 channel, 8 banks/channel
 - DDR3 1066 DRAM
 - 512 KB private cache/core
- Workloads
 - SPEC CPU2006
 - 300 multi programmed workloads

Quantitative Comparison

SPEC CPU 2006 application
leslie3d



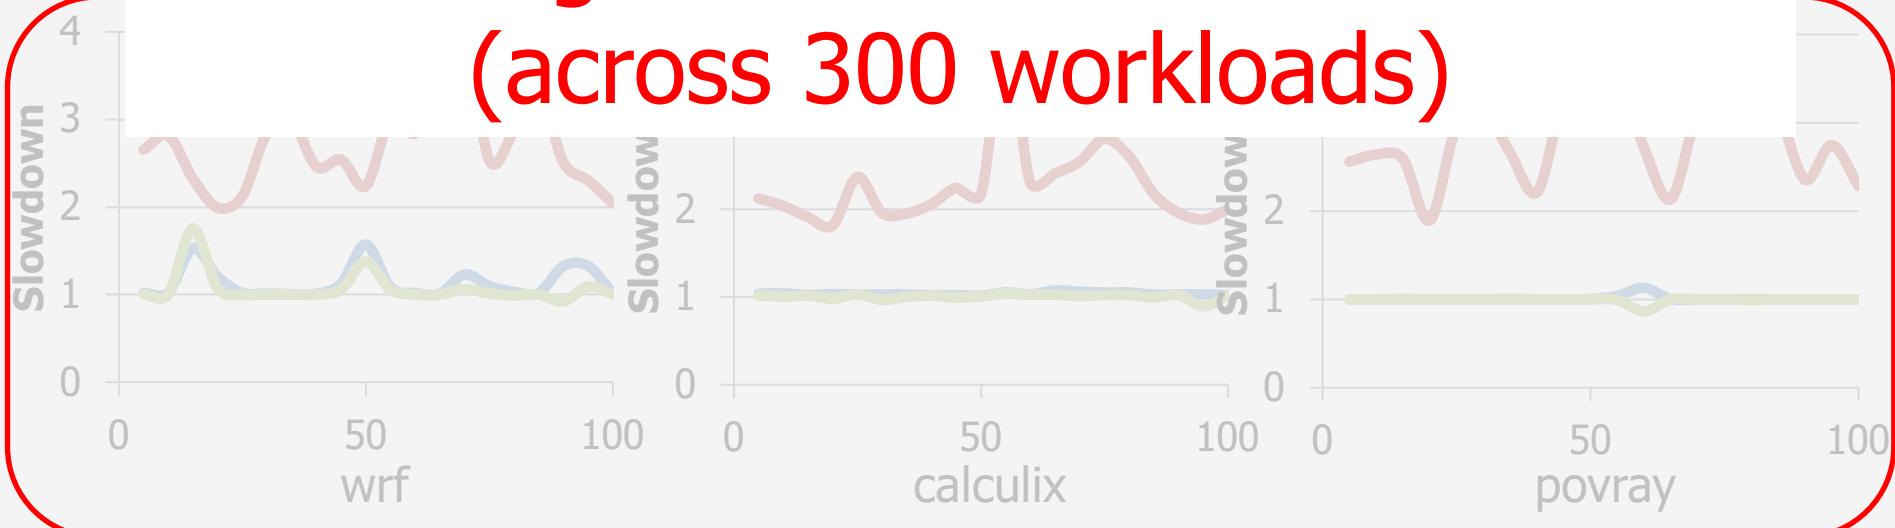
Comparison to STFM



Average error of MISE: 8.2%

Average error of STFM: 29.4%

(across 300 workloads)



Outline

1. Estimate Slowdown

- ❑ Key Observations
- ❑ Implementation
- ❑ MISE Model: Putting it All Together
- ❑ Evaluating the Model

2. Control Slowdown

- ❑ Providing Soft Slowdown Guarantees
- ❑ Minimizing Maximum Slowdown

Providing “Soft” Slowdown Guarantees

- Goal
 - 1. Ensure QoS-critical applications meet a prescribed slowdown bound
 - 2. Maximize system performance for other applications

- Basic Idea
 - Allocate just enough bandwidth to QoS-critical application
 - Assign remaining bandwidth to other applications

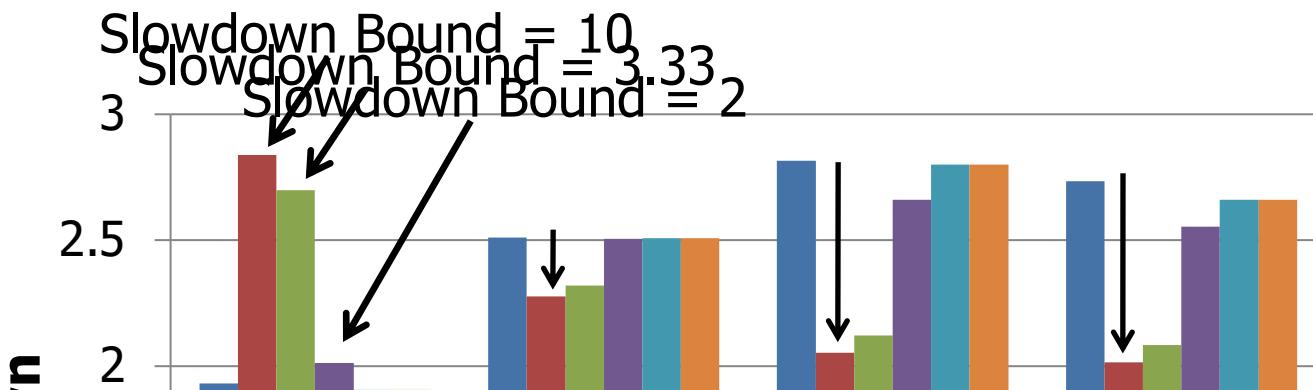
MISE-QoS: Mechanism to Provide Soft QoS

- Assign an initial bandwidth allocation to QoS-critical application
- Estimate slowdown of QoS-critical application using the MISE model
- After every N intervals
 - If slowdown $>$ bound $B +/\!- \varepsilon$, increase bandwidth allocation
 - If slowdown $<$ bound $B +/\!- \varepsilon$, decrease bandwidth allocation
- When slowdown bound not met for N intervals
 - Notify the OS so it can migrate/de-schedule jobs

Methodology

- Each application (25 applications in total) considered the QoS-critical application
- Run with **12 sets of co-runners** of different memory intensities
- Total of **300 multiprogrammed workloads**
- Each workload run with **10 slowdown bound values**
- Baseline memory scheduling mechanism
 - Always prioritize QoS-critical application
[Iyer+, SIGMETRICS 2007]
 - Other applications' requests scheduled in FRFCFS order
[Zuravleff +, US Patent 1997, Rixner+, ISCA 2000]

A Look at One Workload



MISE is effective in

1. meeting the slowdown bound for the QoS-critical application
2. improving performance of non-QoS-critical applications



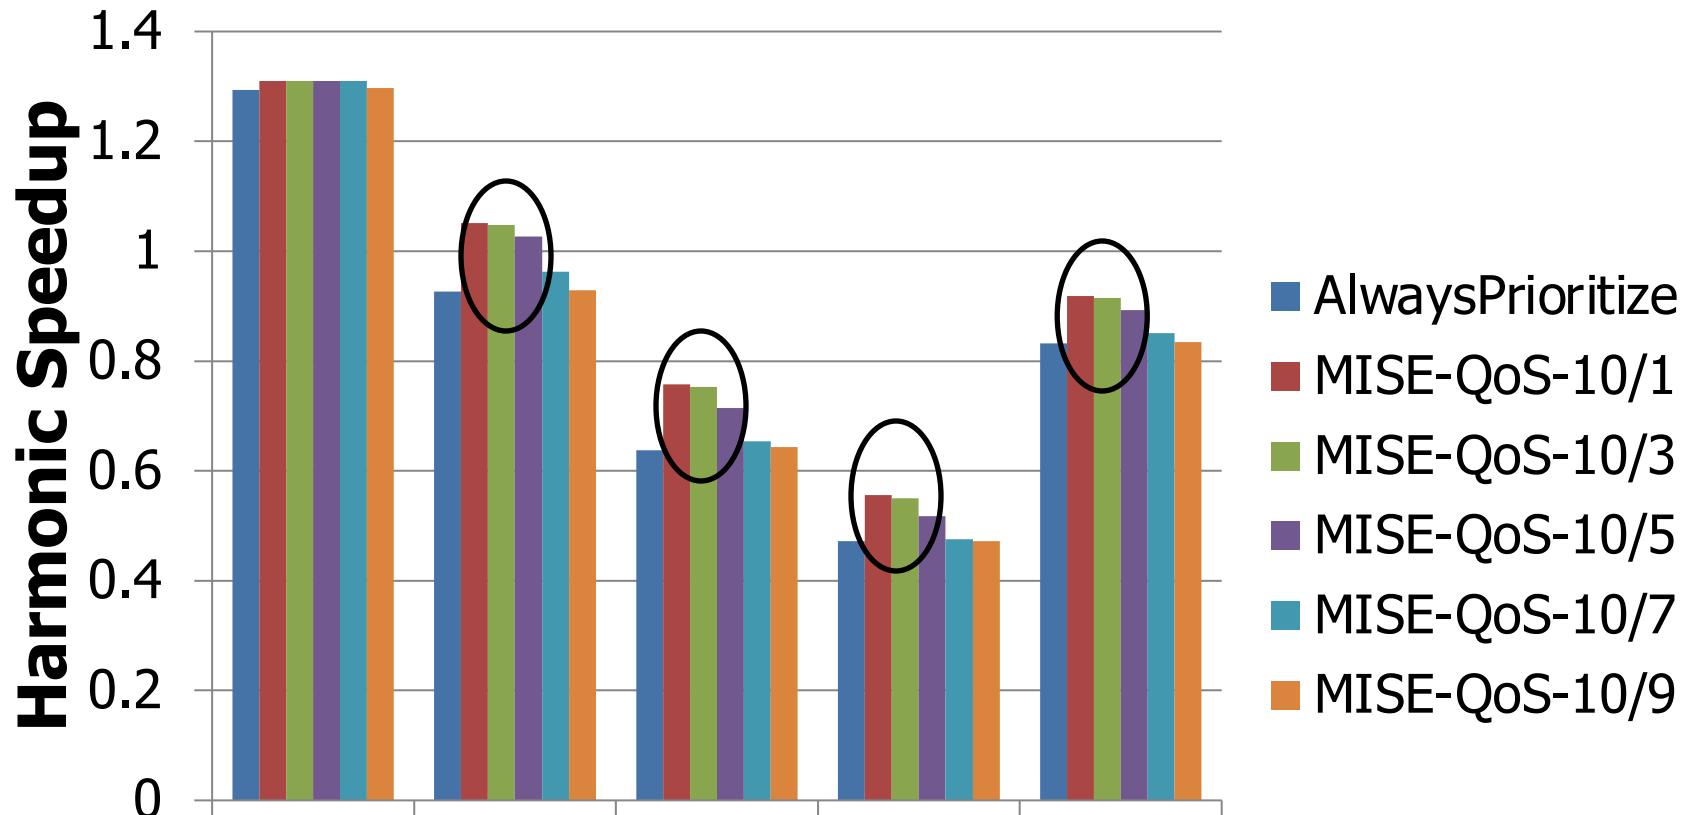
Effectiveness of MISE in Enforcing QoS

Across 3000 data points

| | Predicted Met | Predicted Not Met |
|-------------------|---------------|-------------------|
| QoS Bound Met | 78.8% | 2.1% |
| QoS Bound Not Met | 2.2% | 16.9% |

MISE-QoS correctly predicts whether or not the bound is met for 95.7% of workloads

Performance of Non-QoS-Critical Applications



When slowdown bound is $10/3$
MISE-QoS improves system performance by 10%

Outline

1. Estimate Slowdown

- ❑ Key Observations
- ❑ Implementation
- ❑ MISE Model: Putting it All Together
- ❑ Evaluating the Model

2. Control Slowdown

- ❑ Providing Soft Slowdown Guarantees
- ❑ Minimizing Maximum Slowdown

Other Results in the Paper

- Sensitivity to model parameters
 - Robust across different values of model parameters
- Comparison of STFM and MISE models in enforcing soft slowdown guarantees
 - MISE significantly more effective in enforcing guarantees
- Minimizing maximum slowdown
 - MISE improves fairness across several system configurations

Summary

- Uncontrolled memory interference slows down applications unpredictably
- Goal: Estimate and control slowdowns
- Key contribution
 - MISE: An accurate slowdown estimation model
 - Average error of MISE: 8.2%
- Key Idea
 - Request Service Rate is a proxy for performance
 - Request Service Rate _{Alone} estimated by giving an application highest priority in accessing memory
- Leverage slowdown estimates to control slowdowns
 - Providing soft slowdown guarantees
 - Minimizing maximum slowdown

MISE: Pros and Cons

- Upsides:
 - Simple new insight to estimate slowdown
 - Much more accurate slowdown estimations than prior techniques (STFM, FST)
 - Enables a number of QoS mechanisms that can use slowdown estimates to satisfy performance requirements

- Downsides:
 - Slowdown estimation is not perfect - there are still errors
 - Does not take into account caches and other shared resources in slowdown estimation

More on MISE

- Lavanya Subramanian, Vivek Seshadri, Yoongu Kim, Ben Jaiyen, and Onur Mutlu,

"MISE: Providing Performance Predictability and Improving Fairness in Shared Main Memory Systems"

Proceedings of the 19th International Symposium on High-Performance Computer Architecture (HPCA), Shenzhen, China, February 2013. [Slides \(pptx\)](#)

MISE: Providing Performance Predictability and Improving Fairness in Shared Main Memory Systems

Lavanya Subramanian

Vivek Seshadri

Yoongu Kim

Ben Jaiyen

Onur Mutlu

Carnegie Mellon University

Extending MISE to Shared Caches: ASM

- Lavanya Subramanian, Vivek Seshadri, Arnab Ghosh, Samira Khan, and Onur Mutlu,

"The Application Slowdown Model: Quantifying and Controlling the Impact of Inter-Application Interference at Shared Caches and Main Memory"

Proceedings of the 48th International Symposium on Microarchitecture (MICRO), Waikiki, Hawaii, USA, December 2015.

[[Slides \(pptx\)](#) ([pdf](#))] [[Lightning Session Slides \(pptx\)](#) ([pdf](#))] [[Poster \(pptx\)](#) ([pdf](#))]

[[Source Code](#)]

The Application Slowdown Model: Quantifying and Controlling the Impact of Inter-Application Interference at Shared Caches and Main Memory

Lavanya Subramanian*§ Vivek Seshadri* Arnab Ghosh*†
Samira Khan*‡ Onur Mutlu*

*Carnegie Mellon University §Intel Labs †IIT Kanpur ‡University of Virginia

Other Ways of Handling Memory Interference

Fundamental Interference Control Techniques

- Goal: to reduce/control inter-thread memory interference

1. Prioritization or request scheduling

2. Data mapping to banks/channels/ranks

3. Core/source throttling

4. Application/thread scheduling

Designing QoS-Aware Memory Systems: Approaches

- **Smart resources:** Design each shared resource to have a configurable interference control/reduction mechanism
 - QoS-aware memory controllers
 - QoS-aware interconnects
 - QoS-aware caches
- **Dumb resources:** Keep each resource free-for-all, but reduce/control interference by injection control or data mapping
 - Source throttling to control access to memory system
 - QoS-aware data mapping to memory controllers
 - QoS-aware thread scheduling to cores

Fundamental Interference Control Techniques

- Goal: to reduce/control inter-thread memory interference

1. Prioritization or request scheduling

2. Data mapping to banks/channels/ranks

3. Core/source throttling

4. Application/thread scheduling

More on Source Throttling (I)

Eiman Ebrahimi, Chang Joo Lee, Onur Mutlu, and Yale N. Patt,

"Fairness via Source Throttling: A Configurable and High-Performance Fairness Substrate for Multi-Core Memory Systems"

Proceedings of the 15th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), pages 335-346, Pittsburgh, PA, March 2010. [Slides \(pdf\)](#)

Best paper award.

Fairness via Source Throttling: A Configurable and High-Performance Fairness Substrate for Multi-Core Memory Systems

Eiman Ebrahimi† Chang Joo Lee† Onur Mutlu§ Yale N. Patt†

†Department of Electrical and Computer Engineering
The University of Texas at Austin
{ebrahimi, cjlee, patt}@ece.utexas.edu

§Computer Architecture Laboratory (CALCM)
Carnegie Mellon University
onur@cmu.edu

More on Source Throttling (II)

- Kevin Chang, Rachata Ausavarungnirun, Chris Fallin, and Onur Mutlu,
"HAT: Heterogeneous Adaptive Throttling for On-Chip Networks"

Proceedings of the 24th International Symposium on Computer Architecture and High Performance Computing (SBAC-PAD), New York, NY, October 2012. Slides (pptx) (pdf)

HAT: Heterogeneous Adaptive Throttling for On-Chip Networks

Kevin Kai-Wei Chang, Rachata Ausavarungnirun, Chris Fallin, Onur Mutlu
Carnegie Mellon University
`{kevincha, rachata, cfallin, onur}@cmu.edu`

More on Source Throttling (III)

- George Nychis, Chris Fallin, Thomas Moscibroda, Onur Mutlu, and Srinivasan Seshan,
**"On-Chip Networks from a Networking Perspective:
Congestion and Scalability in Many-core Interconnects"**
*Proceedings of the 2012 ACM SIGCOMM Conference
(SIGCOMM)*, Helsinki, Finland, August 2012. [Slides \(pptx\)](#)

On-Chip Networks from a Networking Perspective: Congestion and Scalability in Many-Core Interconnects

George Nychis[†], Chris Fallin[†], Thomas Moscibroda[§], Onur Mutlu[†], Srinivasan Seshan[†]

[†] Carnegie Mellon University
{gnychis,cfallin,onur,srini}@cmu.edu

[§] Microsoft Research Asia
moscitho@microsoft.com

Fundamental Interference Control Techniques

- Goal: to reduce/control interference
1. Prioritization or request scheduling
 2. Data mapping to banks/channels/ranks
 3. Core/source throttling
 4. Application/thread scheduling

Idea: Pick threads that do not badly interfere with each other to be scheduled together on cores sharing the memory system

Application-to-Core Mapping to Reduce Interference

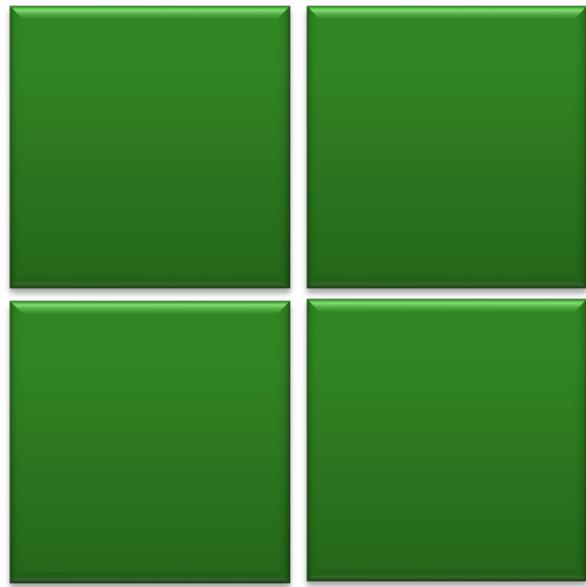
- Reetuparna Das, Rachata Ausavarungnirun, Onur Mutlu, Akhilesh Kumar, and Mani Azimi,

"Application-to-Core Mapping Policies to Reduce Memory System Interference in Multi-Core Systems"

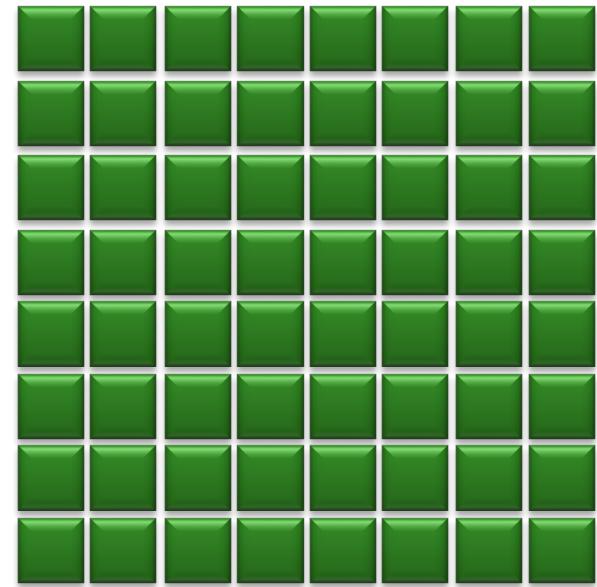
Proceedings of the 19th International Symposium on High-Performance Computer Architecture (HPCA), Shenzhen, China, February 2013.
Slides (pptx)

- Key ideas:
 - Cluster threads to memory controllers (to reduce across chip interference)
 - Isolate interference-sensitive (low-intensity) applications in a separate cluster (to reduce interference from high-intensity applications)
 - Place applications that benefit from memory bandwidth closer to the controller

Multi-Core to Many-Core



Multi-Core



Many-Core

Many-Core On-Chip Communication

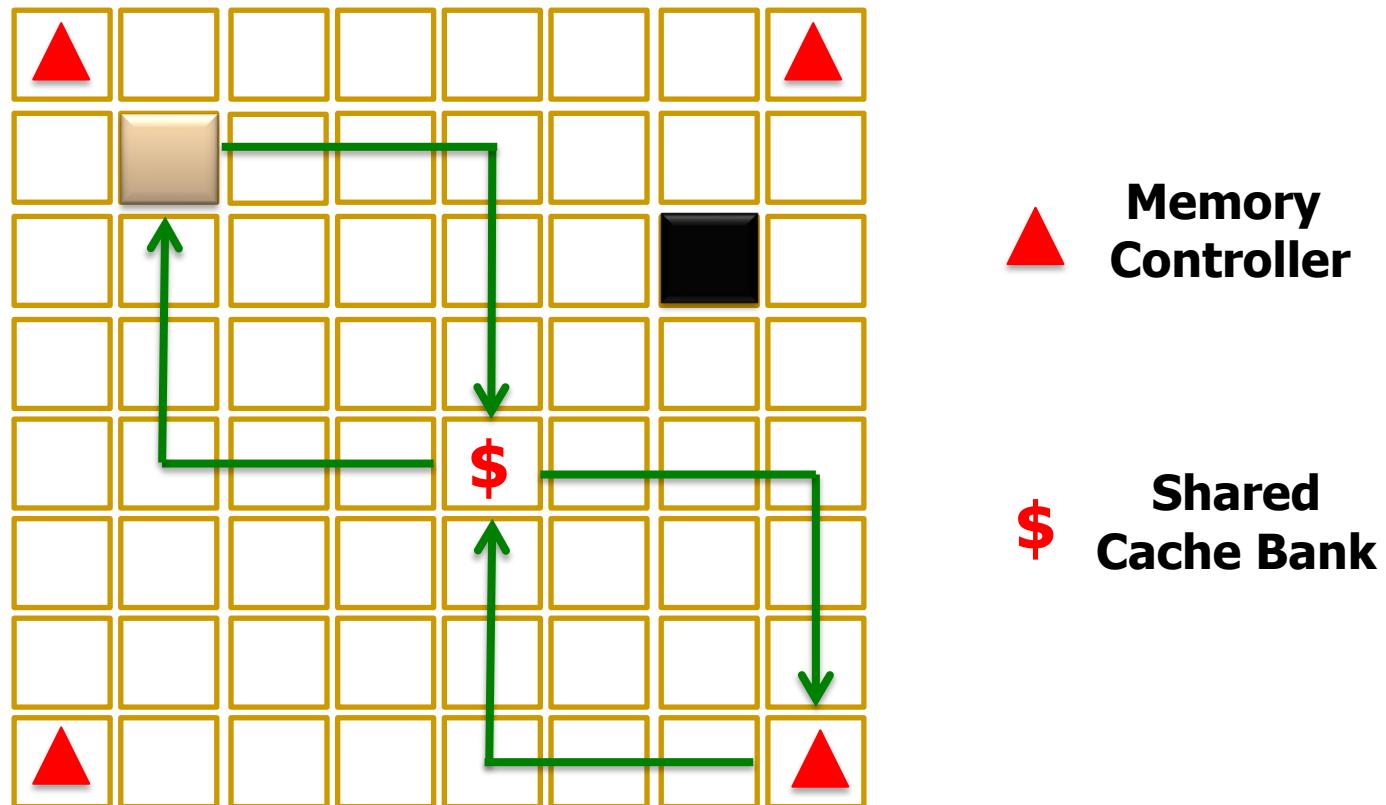
Applications



Light



Heavy

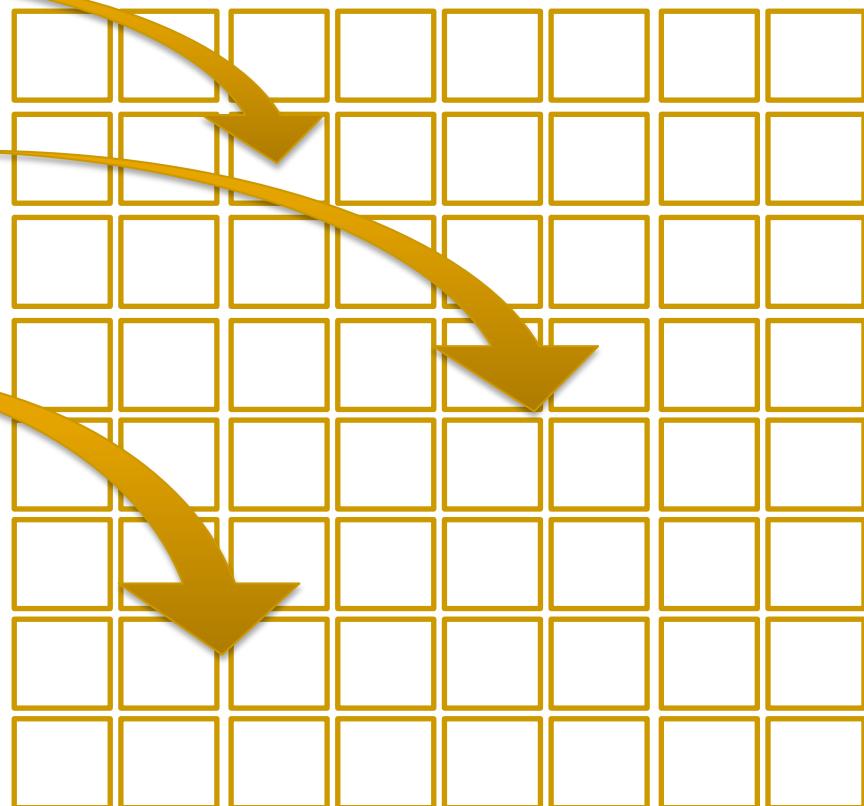


Problem: Spatial Task Scheduling

Applications



Cores



How to map applications to cores?

Challenges in Spatial Task Scheduling

Applications

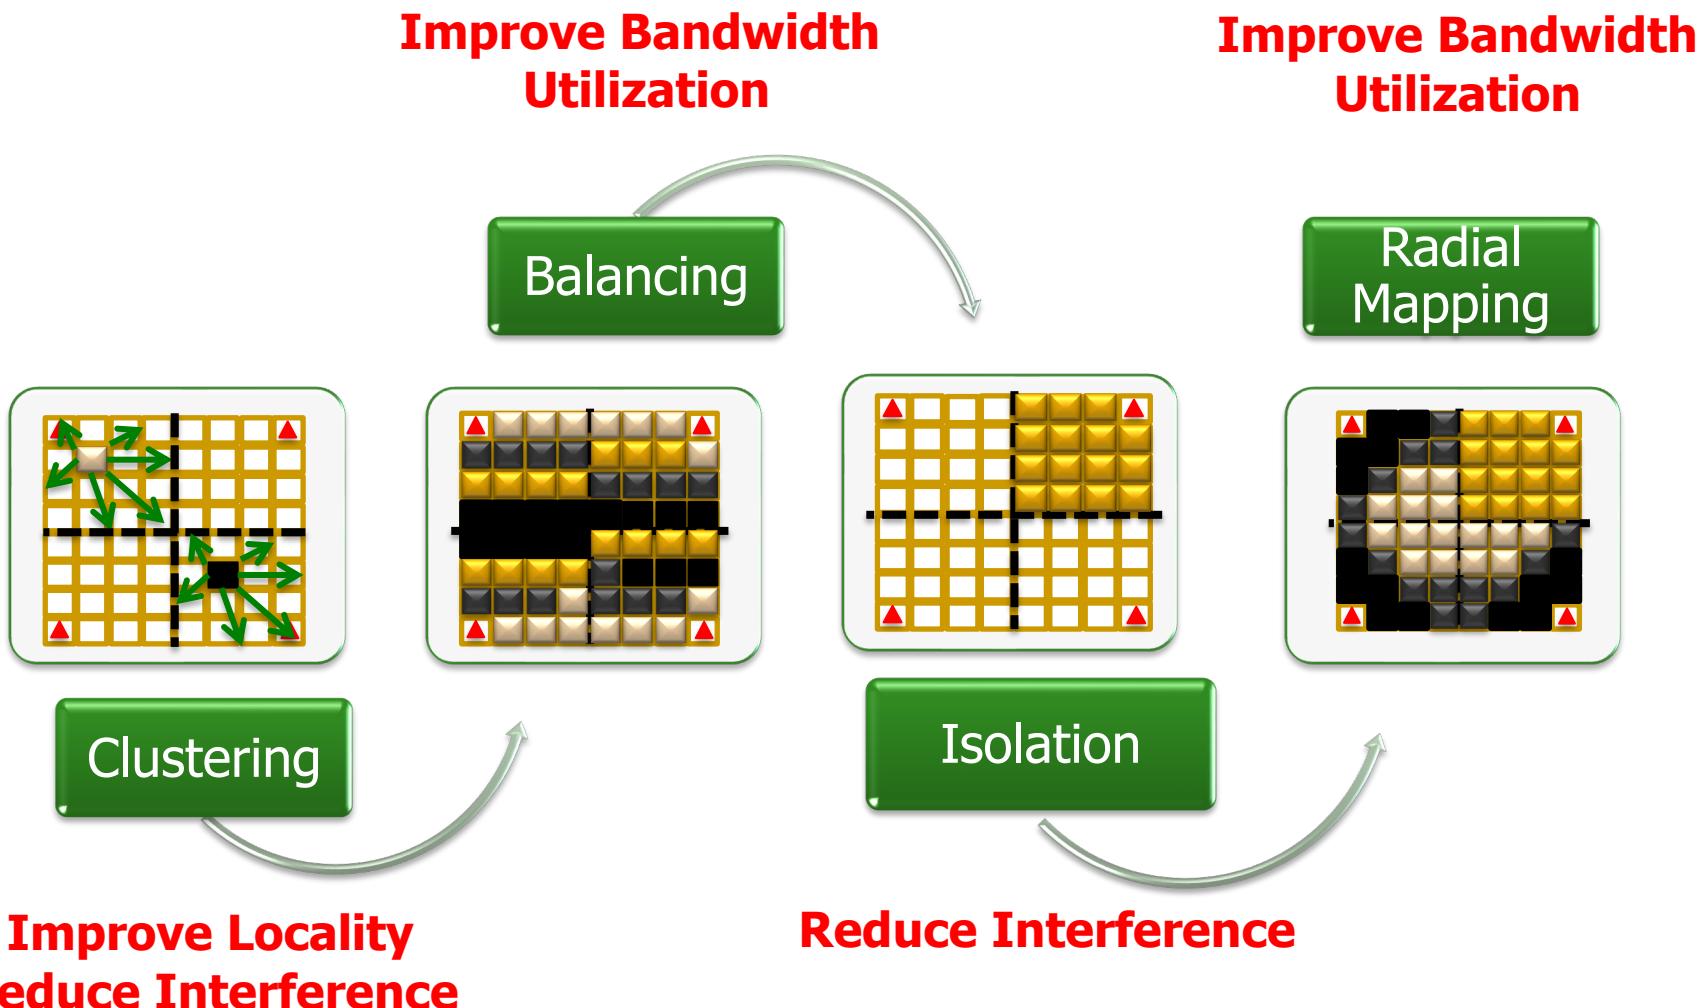
Cores

How to reduce communication distance?

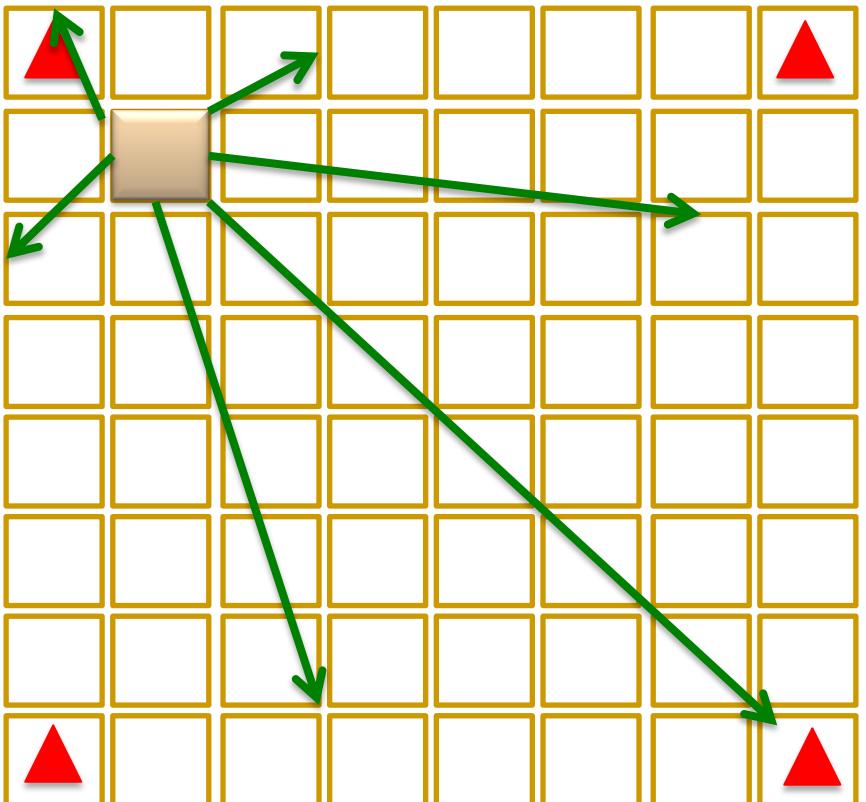
How to reduce destructive interference between applications?

How to prioritize applications to improve throughput?

Application-to-Core Mapping



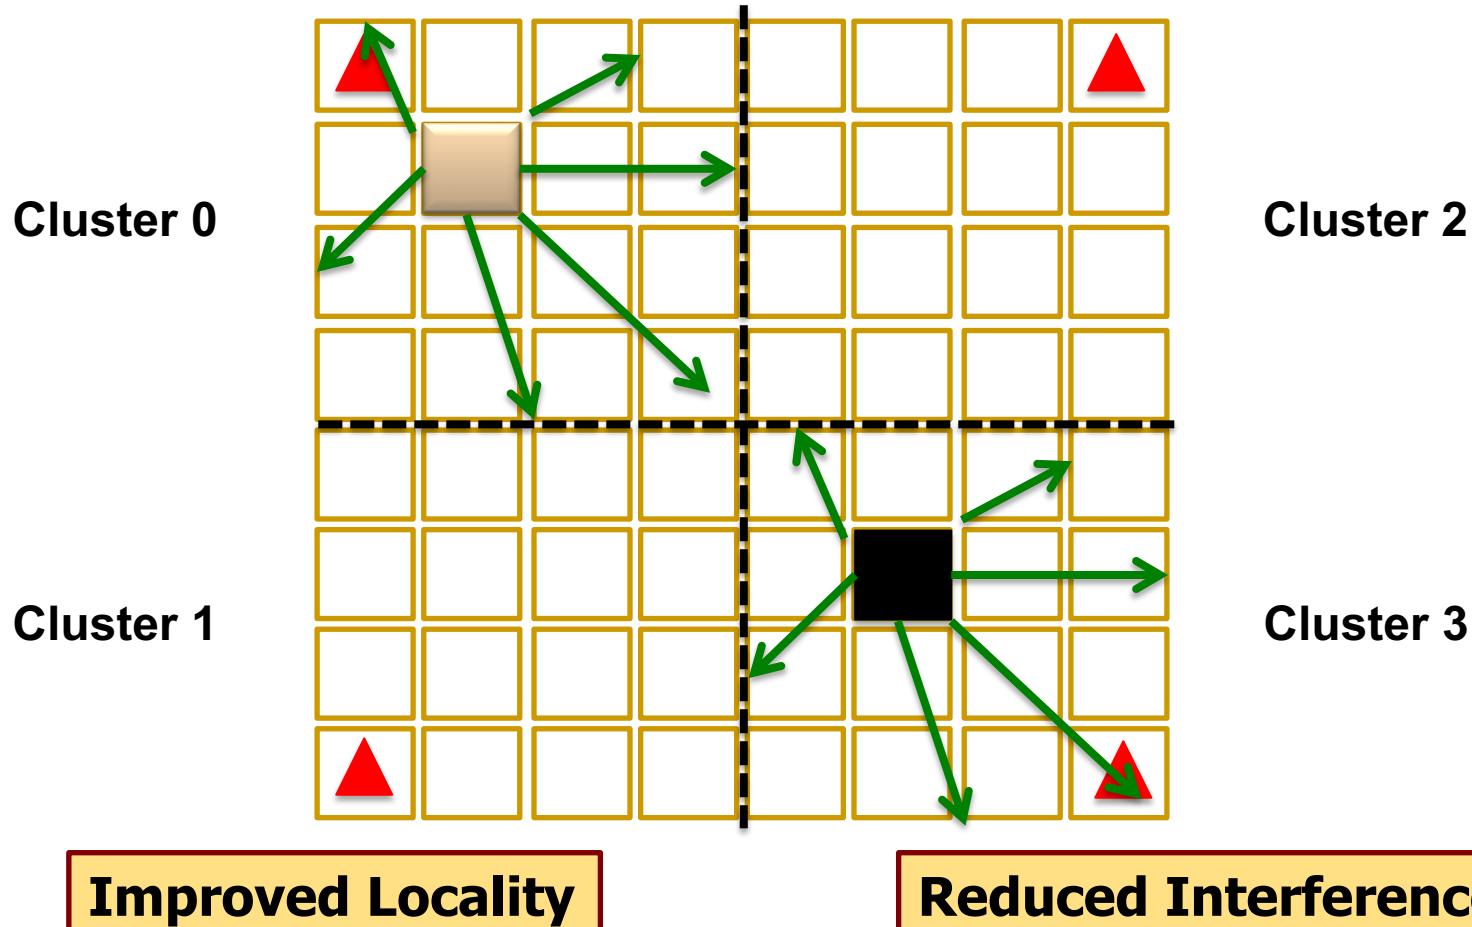
Step 1 — Clustering



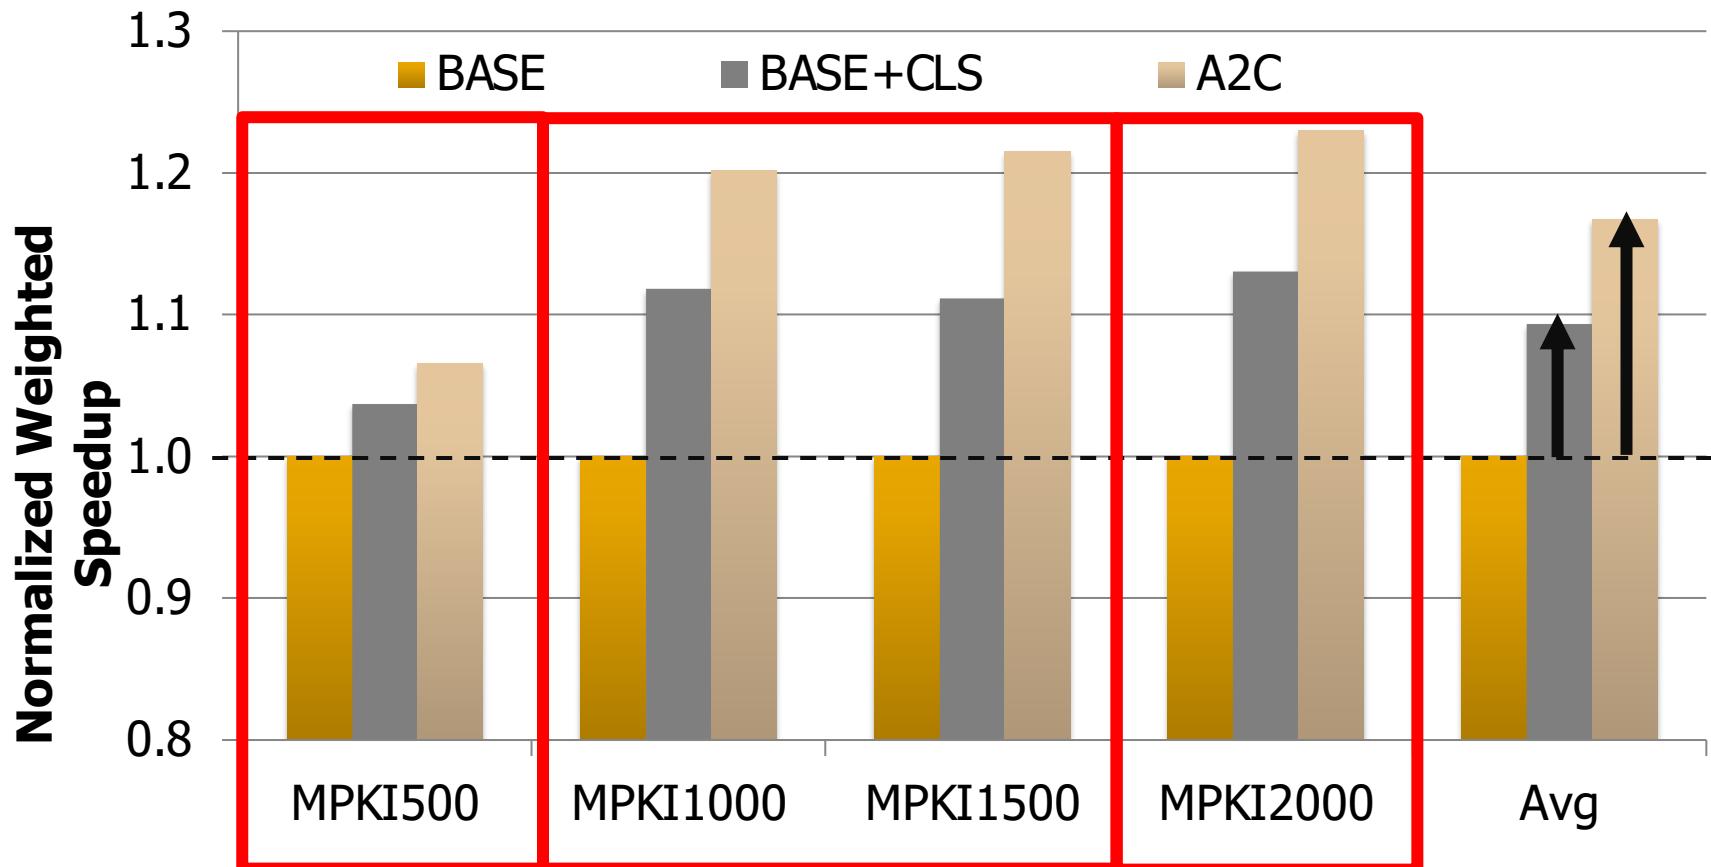
Memory
Controller

Inefficient data mapping to memory and caches

Step 1 — Clustering

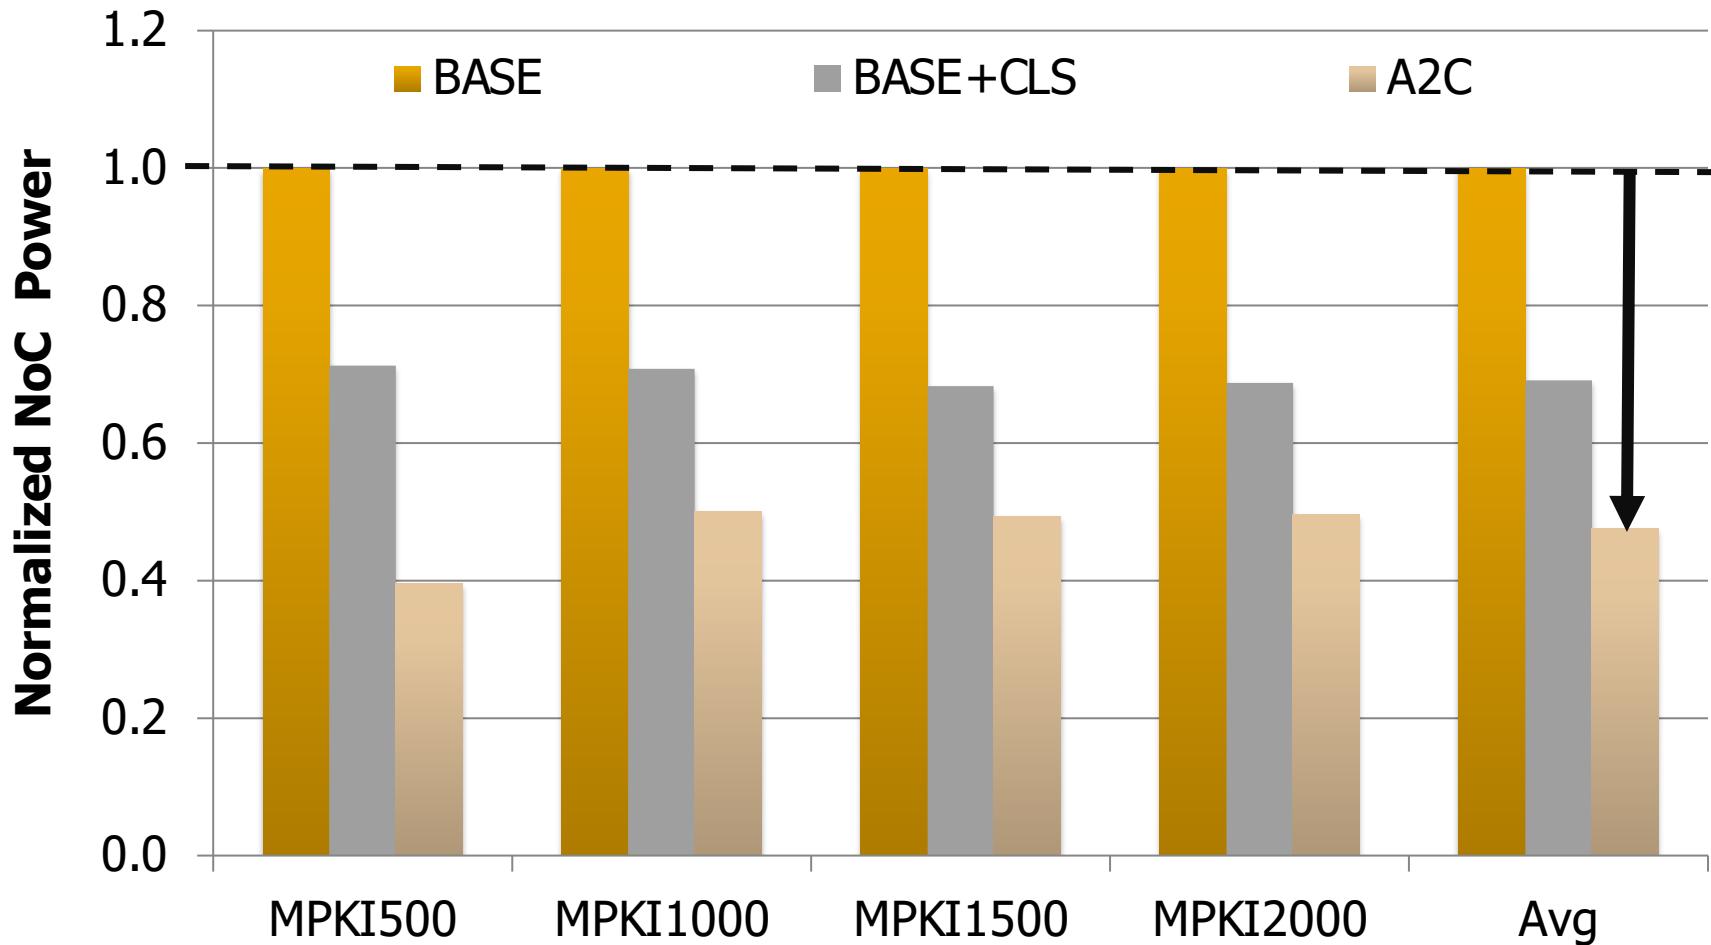


System Performance



System performance improves by 17%

Network Power



Average network power consumption reduces by 52%

More on App-to-Core Mapping

- Reetuparna Das, Rachata Ausavarungnirun, Onur Mutlu, Akhilesh Kumar, and Mani Azimi,

"Application-to-Core Mapping Policies to Reduce Memory System Interference in Multi-Core Systems"

Proceedings of the 19th International Symposium on High-Performance Computer Architecture (HPCA), Shenzhen, China, February 2013.
Slides (pptx)

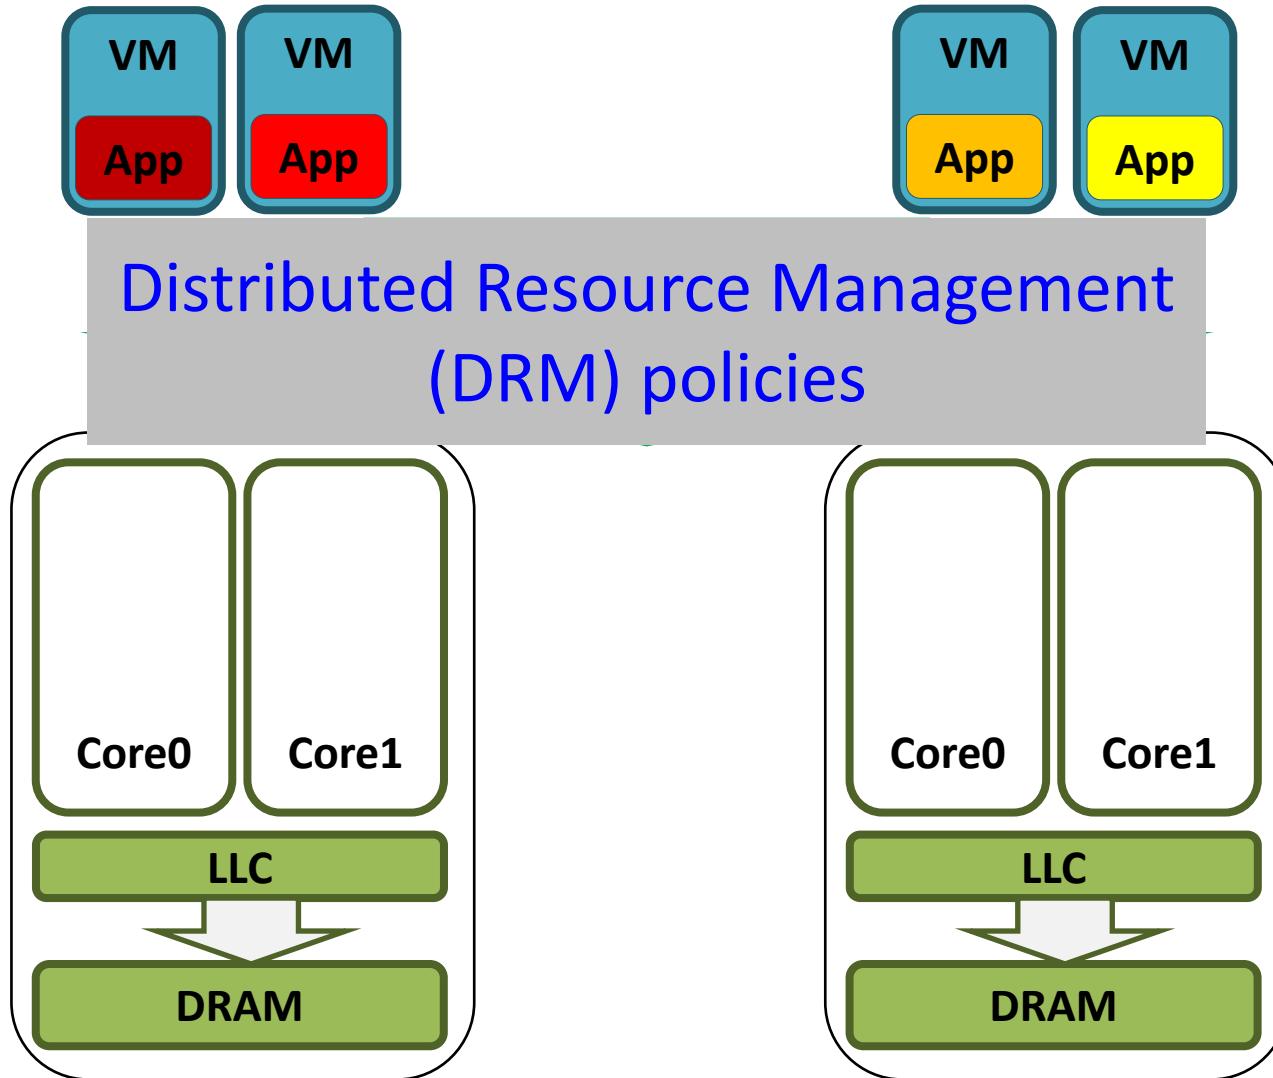
Application-to-Core Mapping Policies to Reduce Memory System Interference in Multi-Core Systems

Reetuparna Das* Rachata Ausavarungnirun† Onur Mutlu† Akhilesh Kumar‡ Mani Azimi‡
University of Michigan* Carnegie Mellon University† Intel Labs‡

Interference-Aware Thread Scheduling

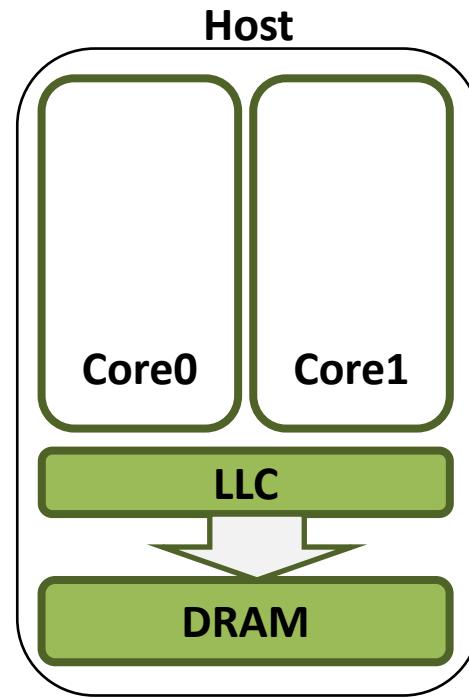
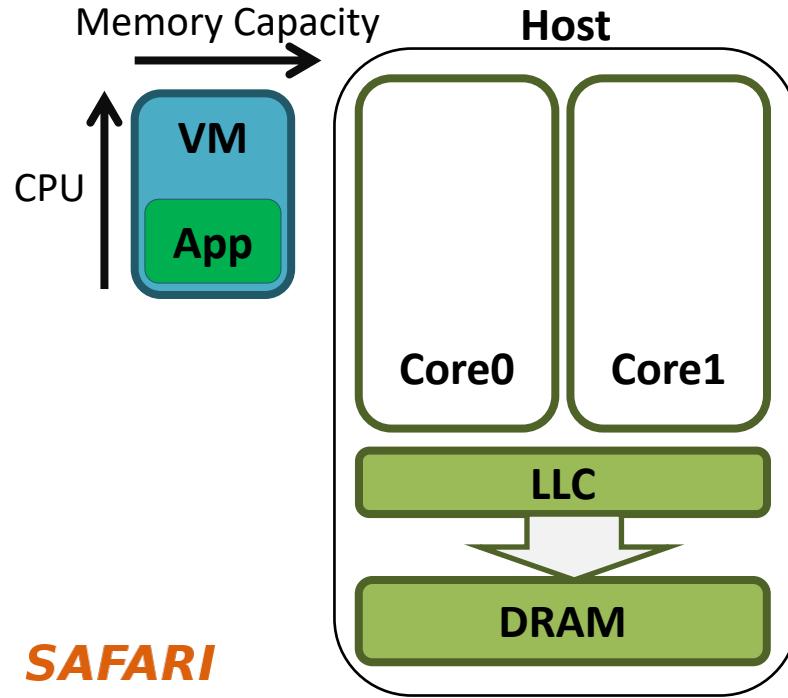
- An example from scheduling in compute clusters (data centers)
- Data centers can be running virtual machines

Virtualized Cluster



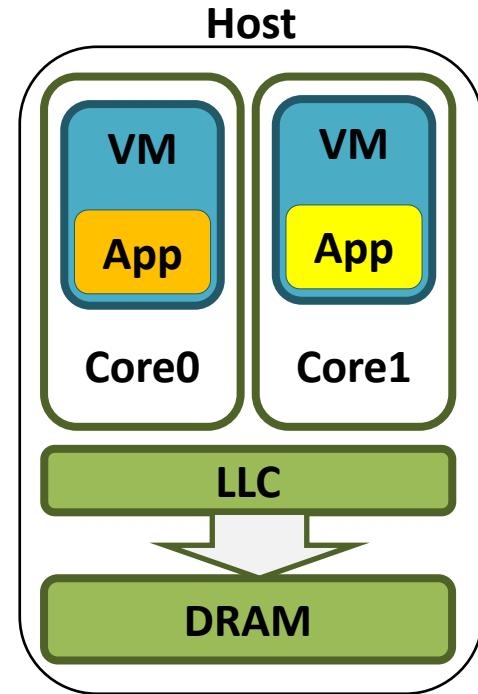
Conventional DRM Policies

Based on operating-system-level metrics
e.g., CPU utilization, memory capacity demand



Microarchitecture-level Interference

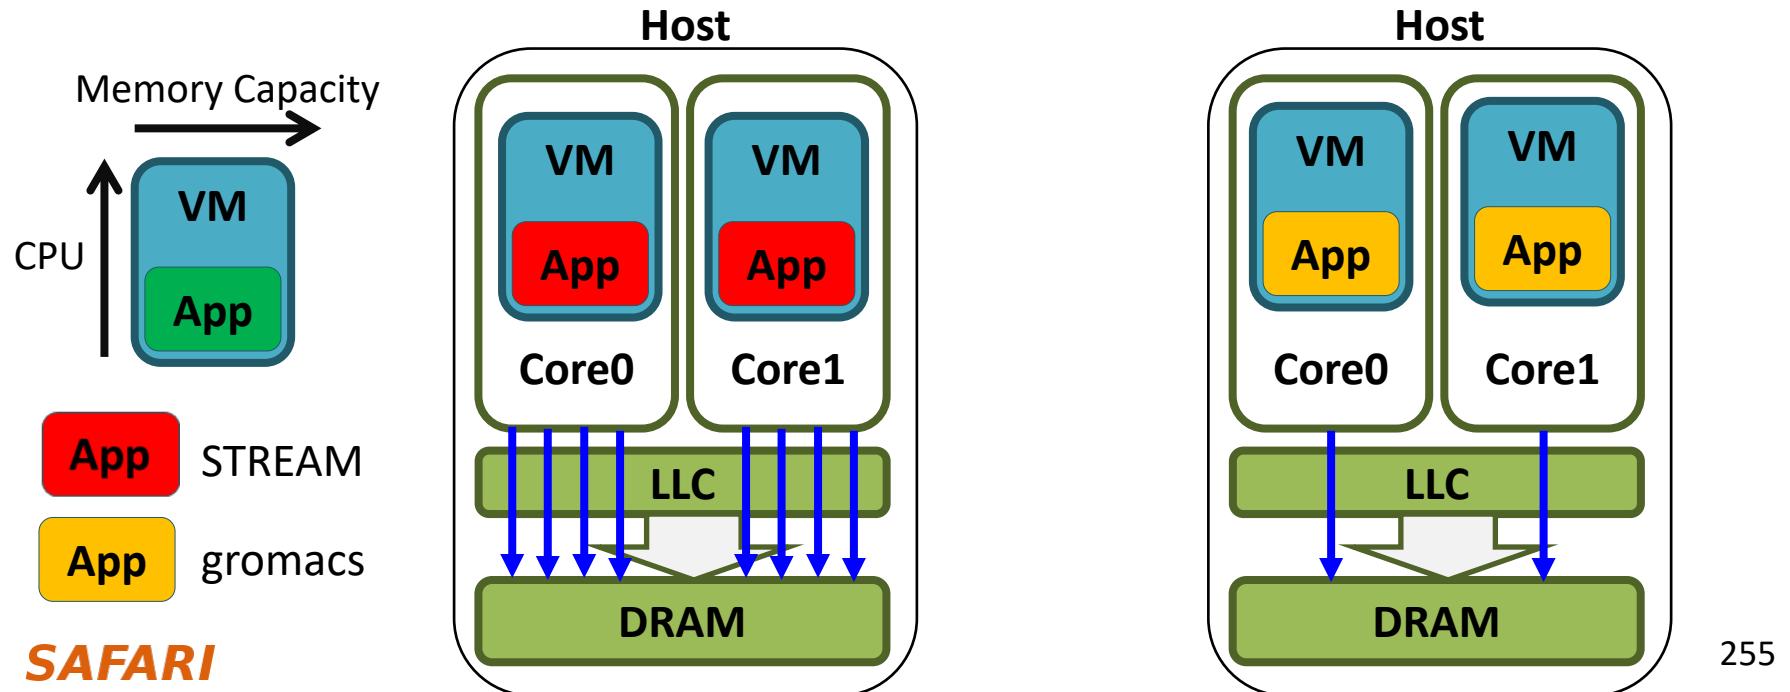
- VMs within a host compete for:
 - Shared cache capacity
 - Shared memory bandwidth



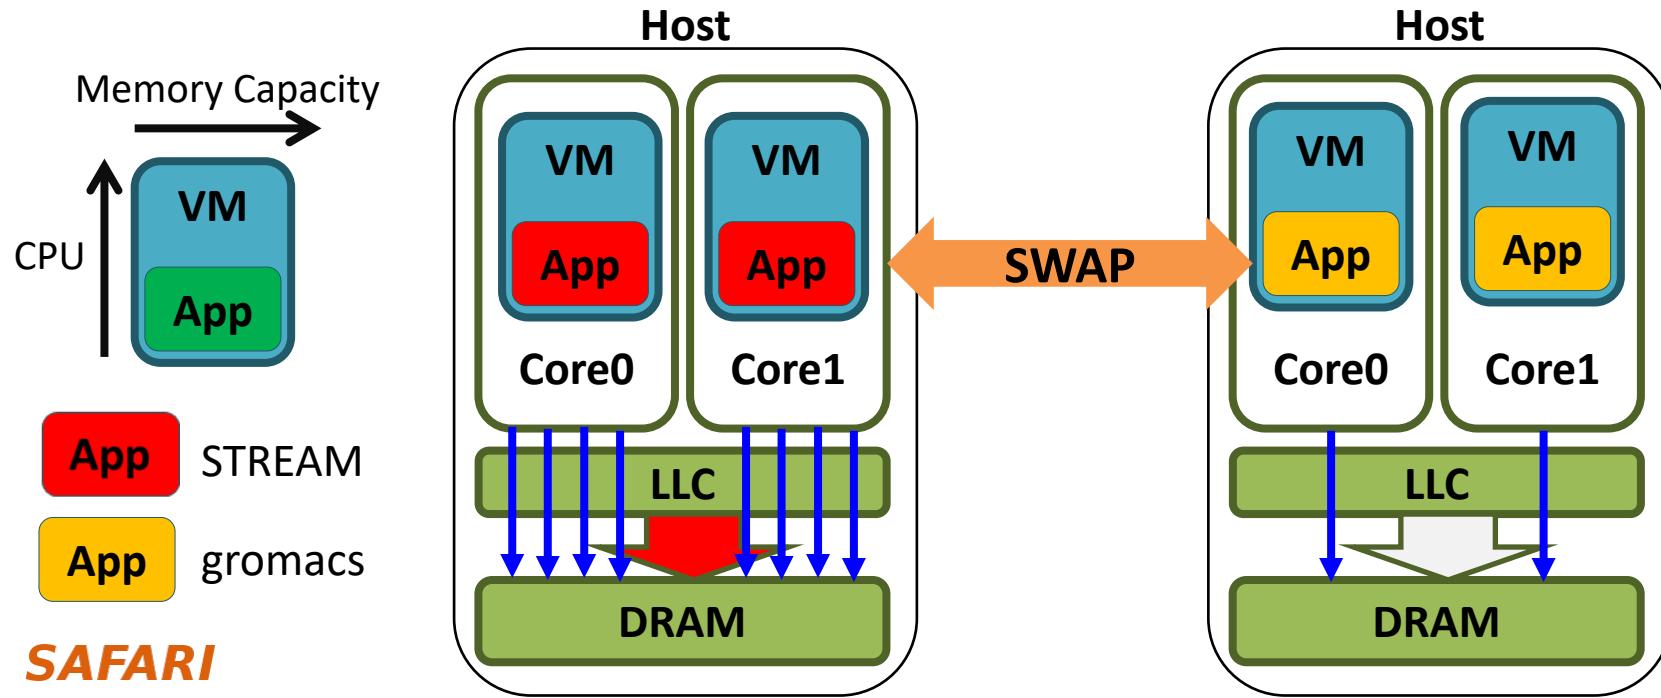
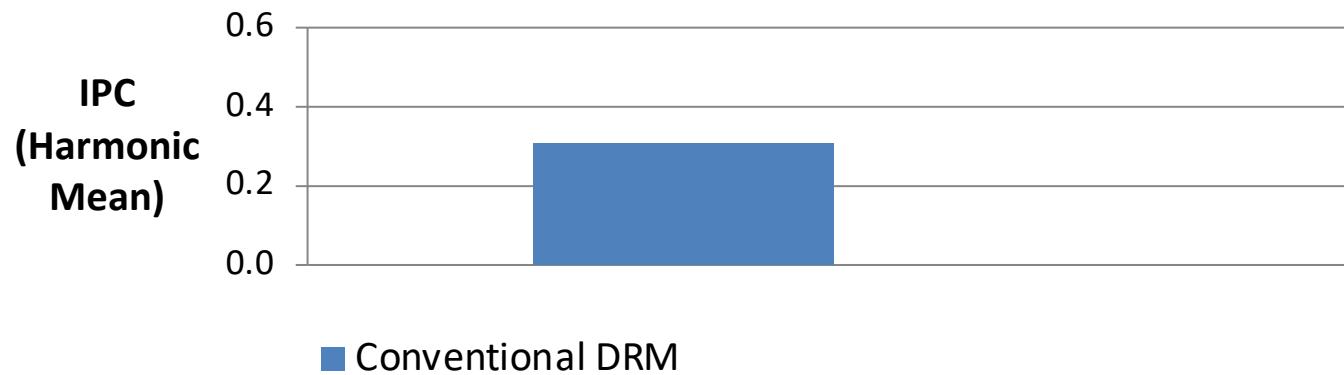
Can operating-system-level metrics capture the microarchitecture-level resource interference?

Microarchitecture Unawareness

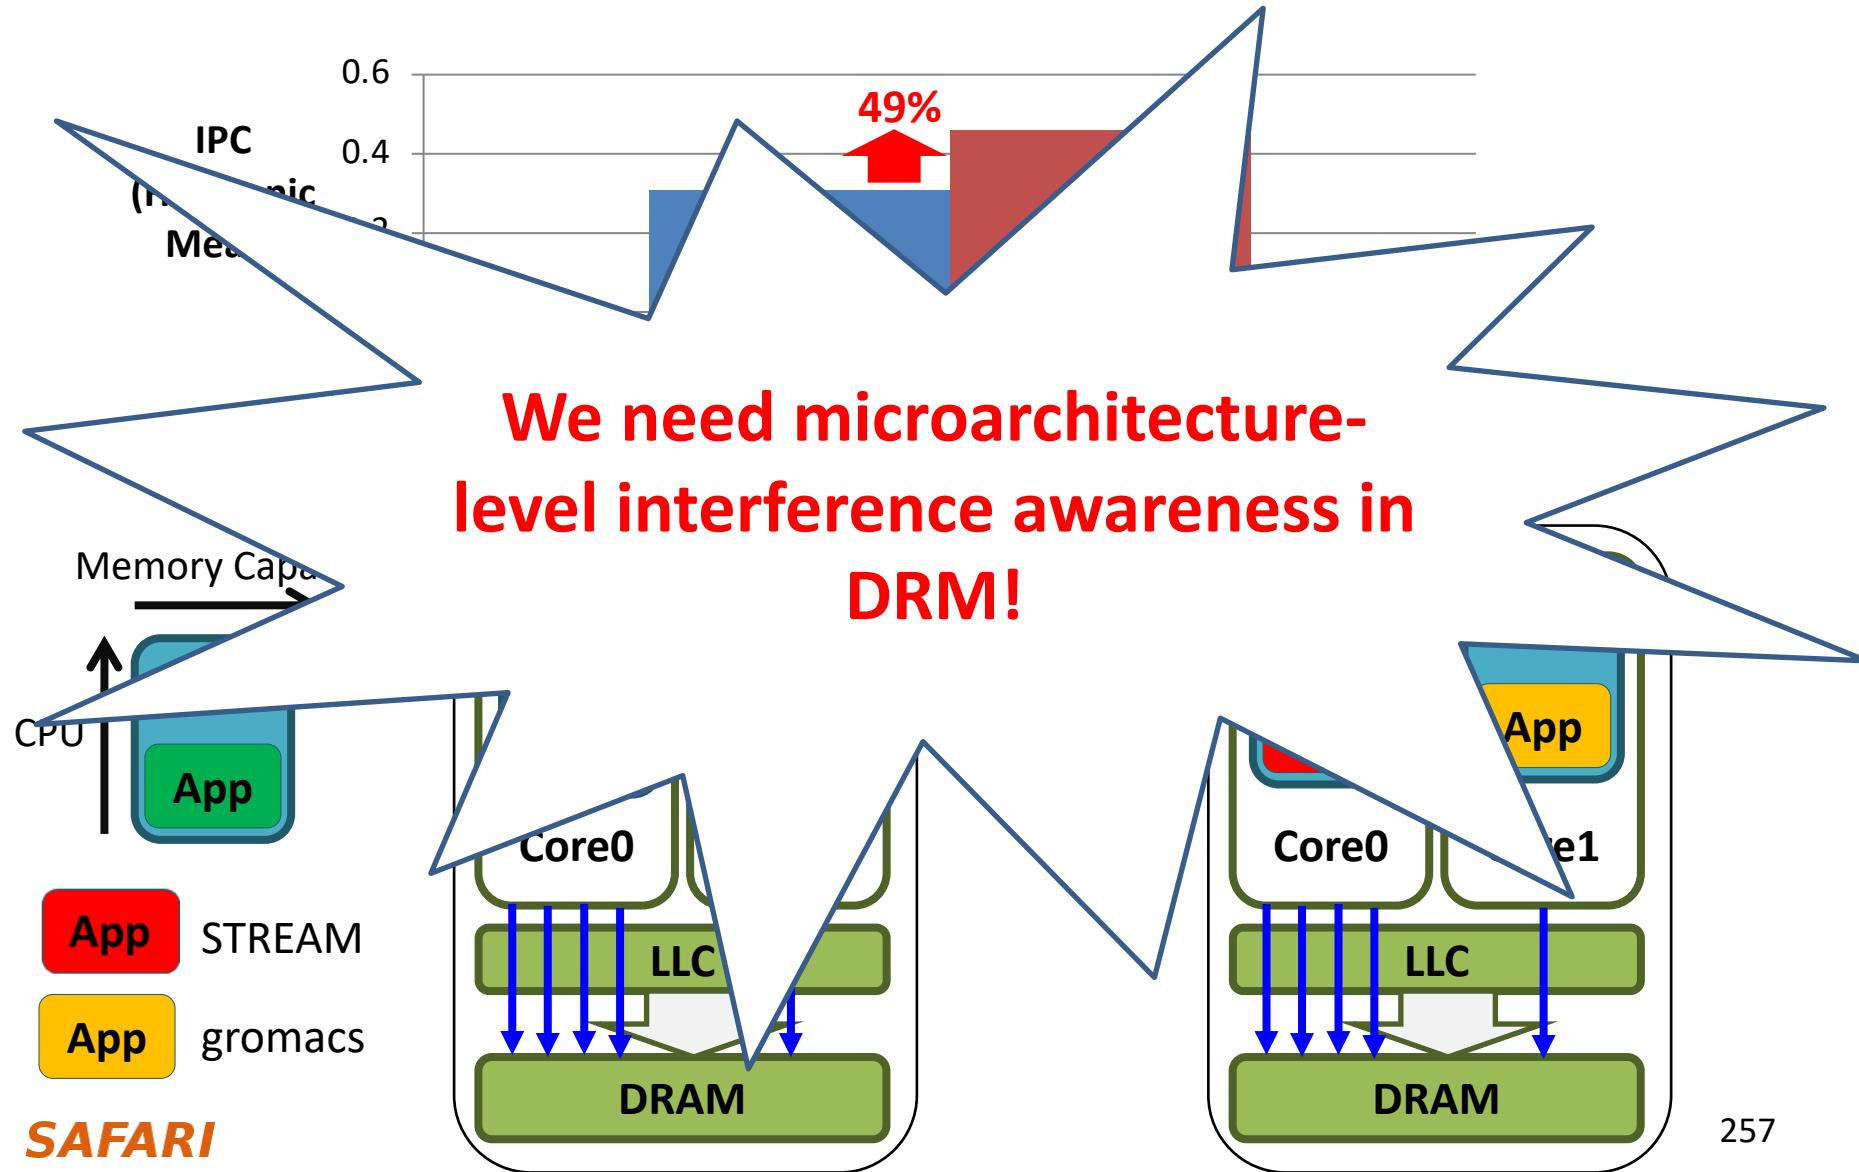
| VM | Operating-system-level metrics | | Microarchitecture-level metrics | |
|-----|--------------------------------|-----------------|---------------------------------|------------------|
| | CPU Utilization | Memory Capacity | LLC Hit Ratio | Memory Bandwidth |
| App | 92% | 369 MB | 2% | 2267 MB/s |
| App | 93% | 348 MB | 98% | 1 MB/s |



Impact on Performance



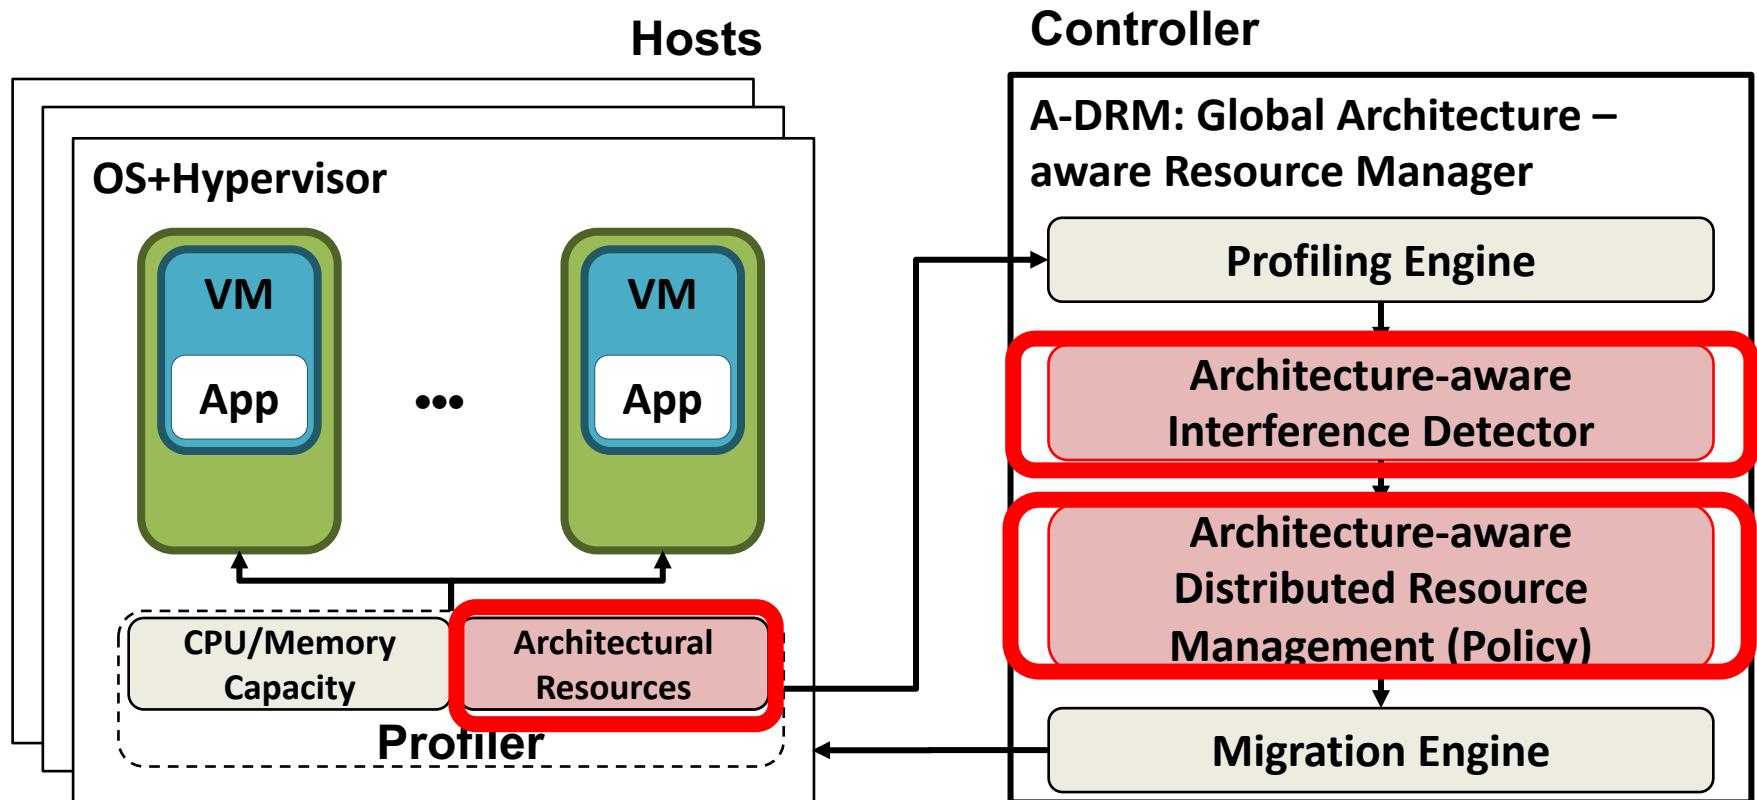
Impact on Performance



A-DRM: Architecture-aware DRM

- **Goal:** Take into account microarchitecture-level shared resource interference
 - Shared cache capacity
 - Shared memory bandwidth
- **Key Idea:**
 - Monitor and detect microarchitecture-level shared resource interference
 - Balance microarchitecture-level resource usage across cluster to minimize memory interference while maximizing system performance

A-DRM: Architecture-aware DRM



More on Architecture-Aware DRM

- Hui Wang, Canturk Isci, Lavanya Subramanian, Jongmoo Choi, Depei Qian, and Onur Mutlu,

"A-DRM: Architecture-aware Distributed Resource Management of Virtualized Clusters"

Proceedings of the 11th ACM SIGPLAN/SIGOPS International Conference on Virtual Execution Environments (VEE), Istanbul, Turkey, March 2015.

[Slides (pptx) (pdf)]

A-DRM: Architecture-aware Distributed Resource Management of Virtualized Clusters

Hui Wang^{†*}, Canturk Isci[‡], Lavanya Subramanian*, Jongmoo Choi^{§*}, Depei Qian[†], Onur Mutlu*

[†]Beihang University, [‡]IBM Thomas J. Watson Research Center, ^{*}Carnegie Mellon University, [§]Dankook University

{hui.wang, depeiq}@buaa.edu.cn, canturk@us.ibm.com, {lsubrama, onur}@cmu.edu, choijm@dankook.ac.kr

Interference-Aware Thread Scheduling

- Advantages
 - + Can eliminate/minimize interference by scheduling “symbiotic applications” together (as opposed to just managing the interference)
 - + Less intrusive to hardware (less need to modify the hardware resources)

- Disadvantages and Limitations
 - High overhead to migrate threads and data between cores and machines
 - Does not work (well) if all threads are similar and they interfere

Summary

Summary: Fundamental Interference Control Techniques

- Goal: to reduce/control interference
 - 1. Prioritization or request scheduling
 - 2. Data mapping to banks/channels/ranks
 - 3. Core/source throttling
 - 4. Application/thread scheduling

Best is to combine all. How would you do that?

Summary: Memory QoS Approaches and Techniques

- Approaches: Smart vs. dumb resources
 - Smart resources: QoS-aware memory scheduling
 - Dumb resources: Source throttling; channel partitioning
 - Both approaches are effective at reducing interference
 - No single best approach for all workloads
- Techniques: Request/thread scheduling, source throttling, memory partitioning
 - All approaches are effective at reducing interference
 - Can be applied at different levels: hardware vs. software
 - No single best technique for all workloads
- Combined approaches and techniques are the most powerful
 - Integrated Memory Channel Partitioning and Scheduling [MICRO'11]

Summary: Memory Interference and QoS

- QoS-unaware memory → uncontrollable and unpredictable system
 - Providing QoS awareness improves performance, predictability, fairness, and utilization of the memory system
 - Discussed many new techniques to:
 - Minimize memory interference
 - Provide predictable performance
 - Many new research ideas needed for integrated techniques and closing the interaction with software
-

What Did We Not Cover?

- Prefetch-aware shared resource management
- DRAM-controller co-design
- Cache interference management
- **Interconnect interference management**
- Write-read scheduling
- **DRAM designs to reduce interference**
- Interference issues in near-memory processing
- ...

What the Future May Bring

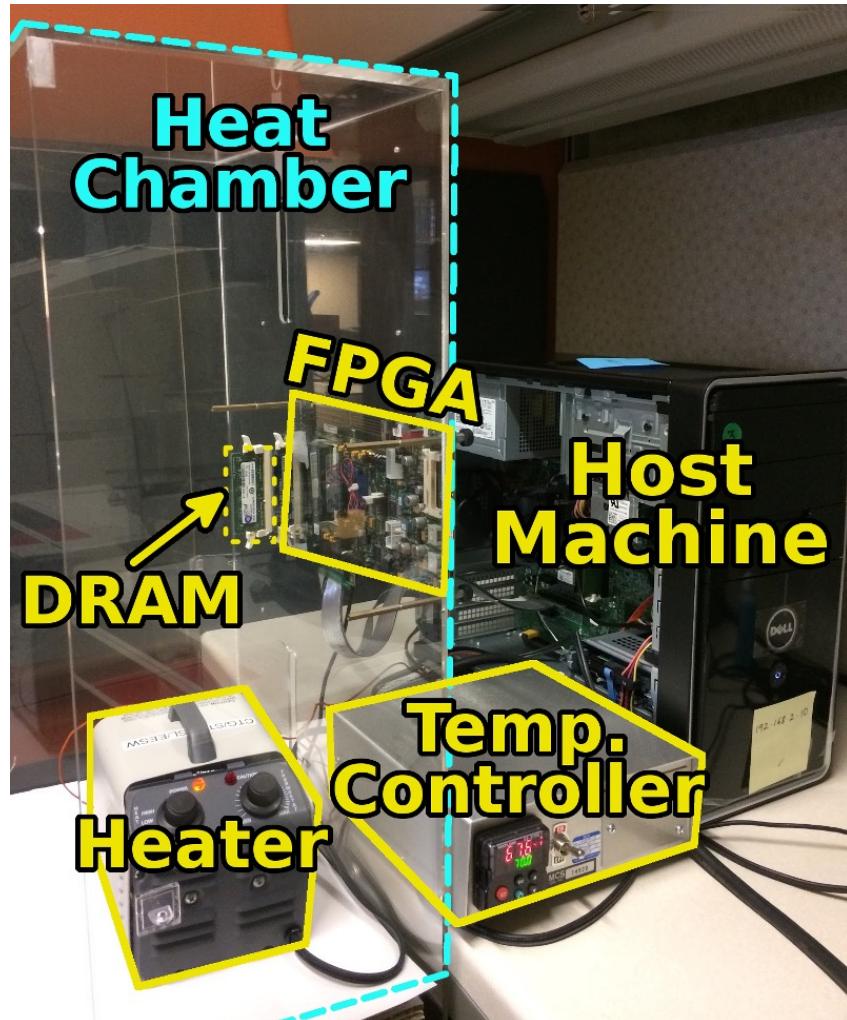
- Memory QoS techniques for heterogeneous SoC systems
 - Many accelerators, processing in/near memory, better predictability, higher performance
- Combinations of memory QoS/performance techniques
 - E.g., data mapping and scheduling
- Fundamentally more intelligent designs that use **machine learning**
- Real prototypes

SoftMC: Open Source DRAM Infrastructure

- Hasan Hassan et al., “SoftMC: A Flexible and Practical Open-Source Infrastructure for Enabling Experimental DRAM Studies,” HPCA 2017.

- **Flexible**
- **Easy to Use (C++ API)**
- **Open-source**

github.com/CMU-SAFARI/SoftMC



- <https://github.com/CMU-SAFARI/SoftMC>

SoftMC: A Flexible and Practical Open-Source Infrastructure for Enabling Experimental DRAM Studies

Hasan Hassan^{1,2,3} Nandita Vijaykumar³ Samira Khan^{4,3} Saugata Ghose³ Kevin Chang³
Gennady Pekhimenko^{5,3} Donghyuk Lee^{6,3} Oguz Ergin² Onur Mutlu^{1,3}

¹*ETH Zürich* ²*TOBB University of Economics & Technology* ³*Carnegie Mellon University*
⁴*University of Virginia* ⁵*Microsoft Research* ⁶*NVIDIA Research*