

Flash-Cosmos

In-Flash Bulk Bitwise Operations Using Inherent Computation Capability of NAND Flash Memory

Jisung Park, Roknoddin Azizi, Geraldo F. Oliveira, Mohammad Sadrosadati,
Rakesh Nadig, David Novo, Juan Gómez Luna, Myungsuk Kim, and Onur Mutlu

MICRO 2022



Executive Summary


- **Background:** Bulk bitwise operations are widely used in many important data-intensive applications, e.g., databases, graph processing, cryptography etc.
- **Problem:**
 - Performance and energy efficiency of bulk bitwise operations are bottlenecked by
 - 1) data movement between storage and the compute unit in traditional systems and in-storage processing (ISP)
 - 2) data sensing (serial reading of operands) in prior in-flash processing (IFP) techniques
 - Prior IFP techniques provide low reliability during computation
- **Goal:** Improve performance, energy efficiency and reliability of bulk bitwise operations in in-flash processing
- **Key Ideas:** Flash-Cosmos (Flash-Computation with One-Shot Multi-Operand Sensing) is an in-flash processing technique that is based on two key ideas:
 - Multi-Wordline Sensing (MWS): Enables multi-operand bulk bitwise operations with a single sensing (read) operation
 - Enhanced SLC-mode Programming (ESP): Increases the voltage margin between the erased and programmed states to provide higher reliability during in-flash computation
- **Key Results:** Flash-Cosmos is evaluated using 160 real 3D NAND flash chips and with a state-of-the-art SSD simulator on three real-world workloads
 - Flash-Cosmos improves the performance and energy efficiency by 3.5x and 3.3x over state-of-the-art IFP technique while providing high reliability during computation

Talk Outline

- Problem, Goals & Key Idea
- Background
- Flash-Cosmos: Computation with One-Shot Multi-Operand Sensing
- Evaluation
- Summary

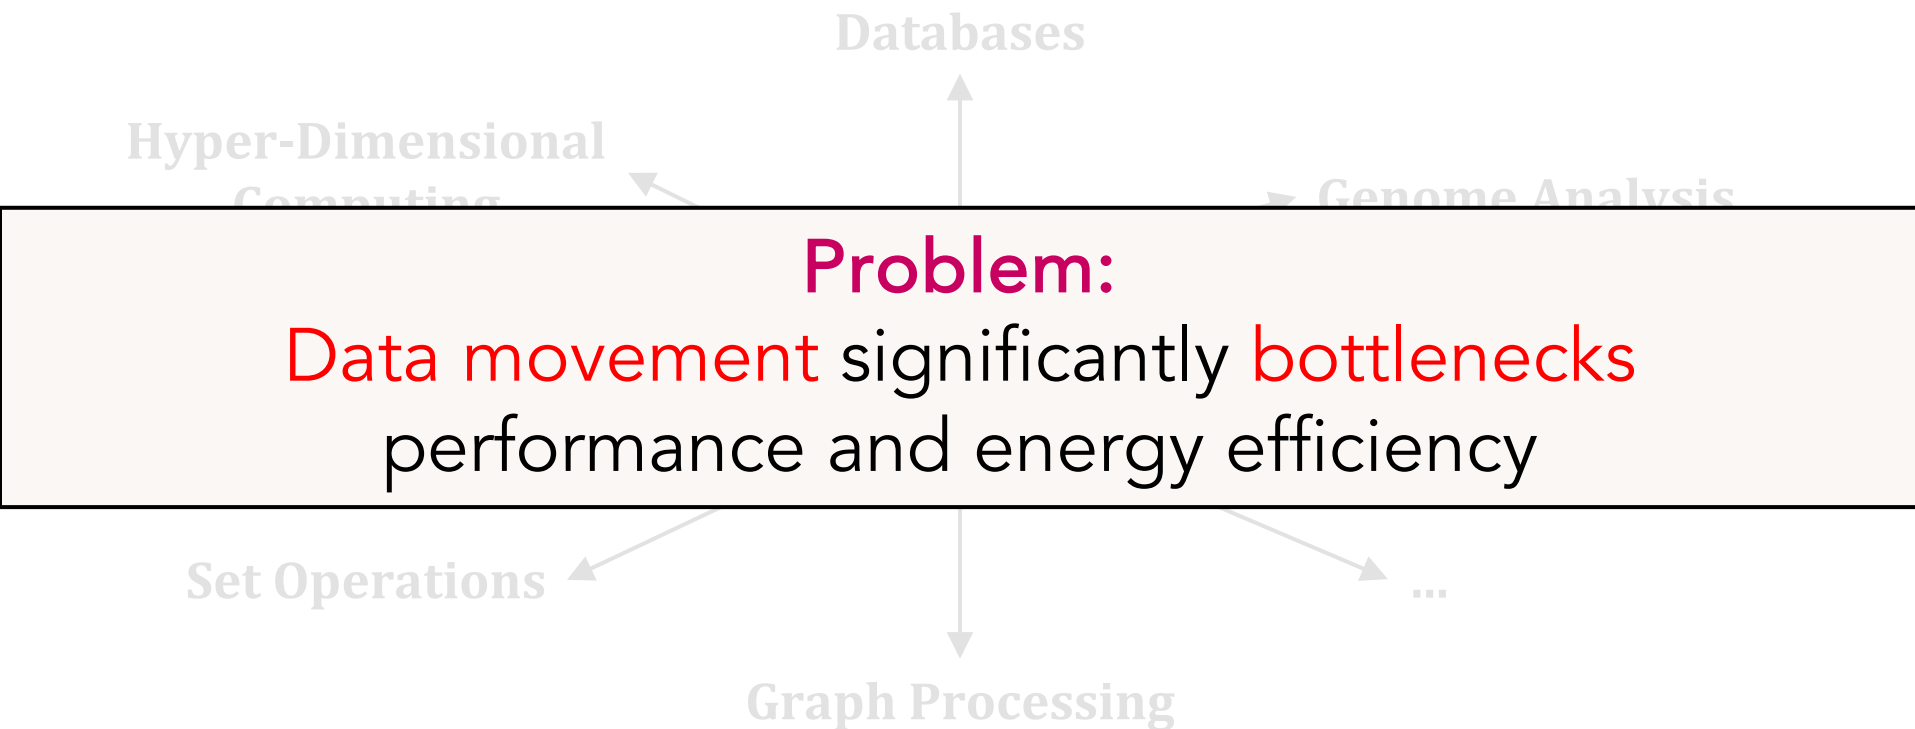
Bulk Bitwise Operations

- Widely-used computation primitive in data-intensive applications



Bulk Bitwise Operations

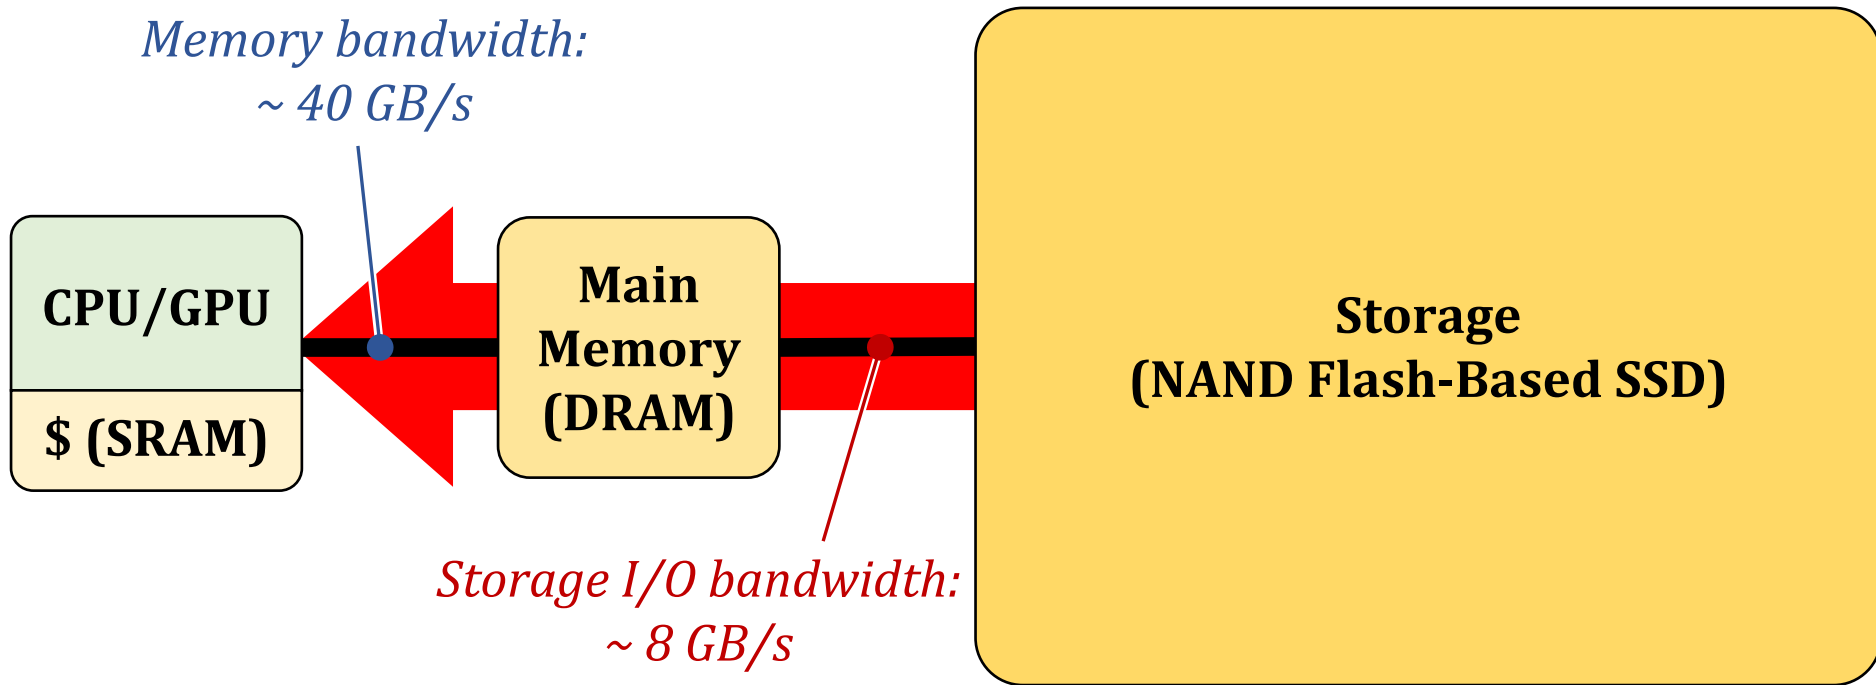
- Widely-used computation primitive in data-intensive applications



Data-Movement Bottleneck

- **Conventional systems:** Outside-storage processing (OSP) that must **move the entire data** to CPUs/GPUs through the memory hierarchy

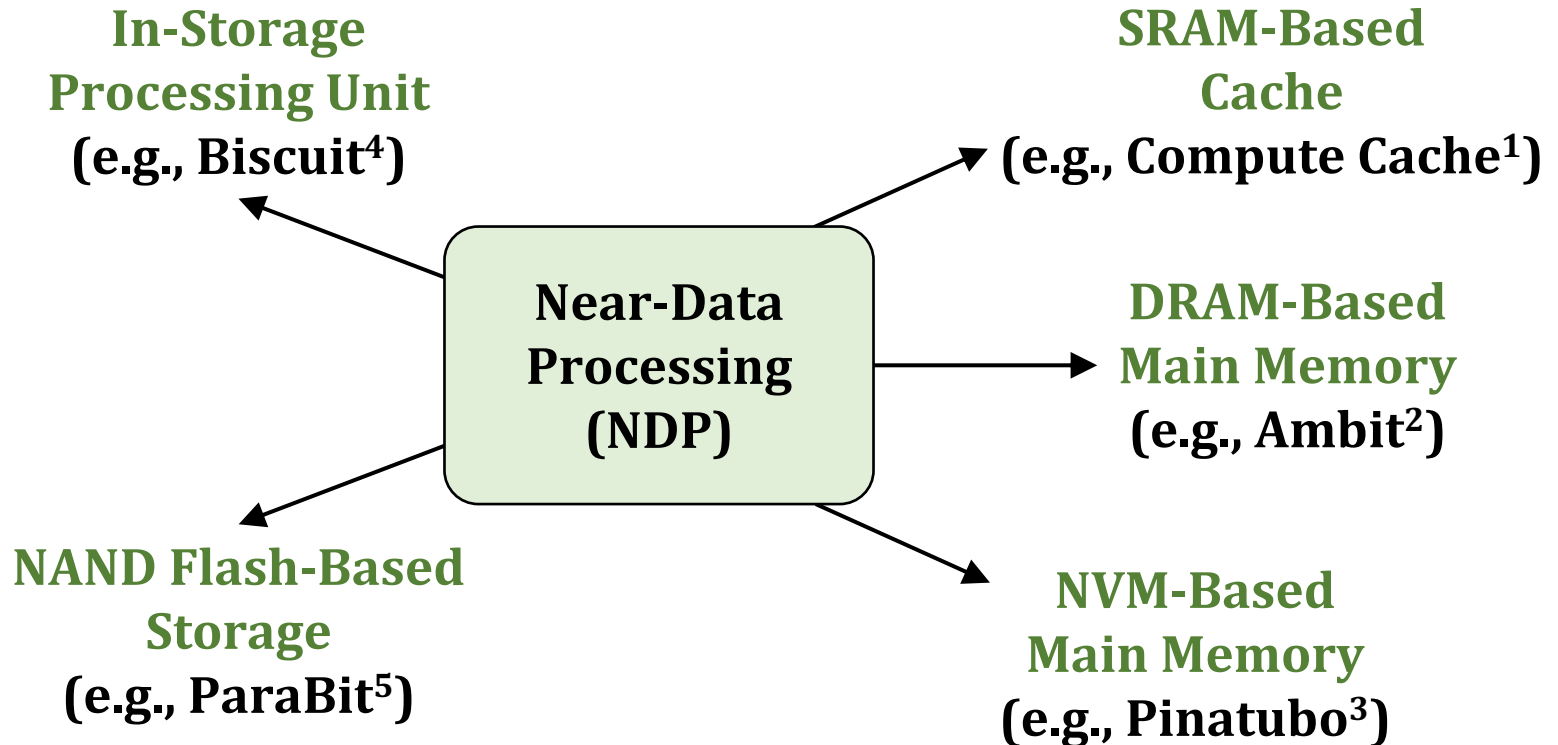
*Memory bandwidth:
~ 40 GB/s*



External I/O bandwidth of storage systems is the **main bottleneck** in conventional systems (OSP)

Near-Data Processing for Bitwise Operations

- Can **effectively** mitigate data movement by performing simple bitwise operations **where the data resides**



¹Aga+, "Compute Caches," HPCA, 2017

²Seshadri+, "Ambit: In-Memory Accelerator for Bulk Bitwise Operations Using Commodity DRAM Technology," MICRO, 2017

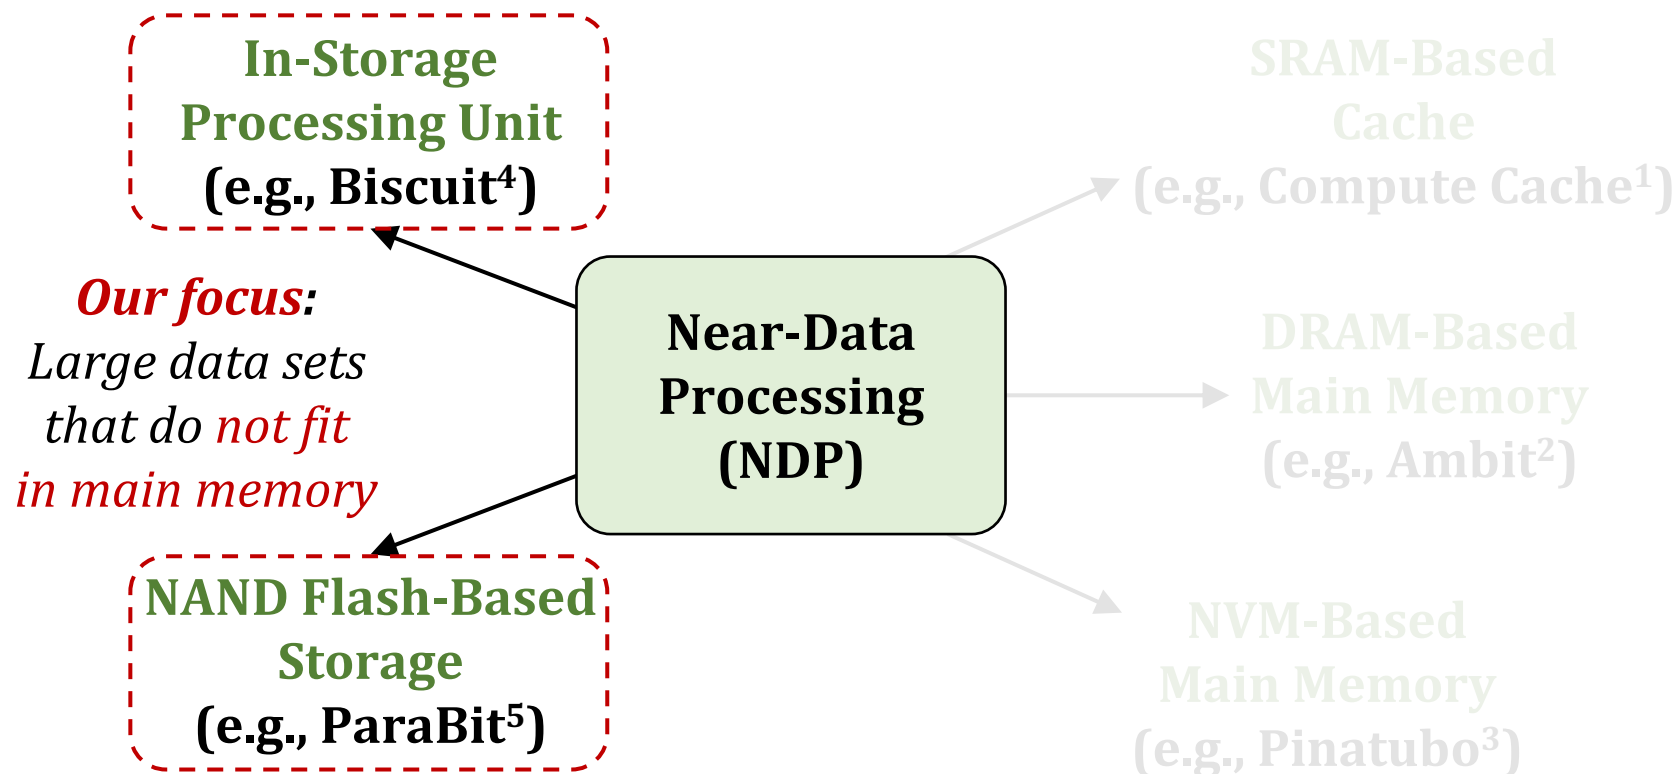
³Li+, "Pinatubo: A Processing-in-Memory Architecture for Bulk Bitwise Operations in Emerging Non-Volatile Memories," DAC, 2016

⁴Gu+, "Biscuit: A Framework for Near-Data Processing of Big Data Workloads," ISCA, 2016

⁵Gao+, "ParaBit: Processing Parallel Bitwise Operations in NAND Flash Memory Based SSDs," MICRO, 2021

Near-Data Processing for Bitwise Operations

- Can **effectively** mitigate data movement by performing simple bitwise operations **where the data resides**



¹Aga+, "Compute Caches," HPCA, 2017

²Seshadri+, "Ambit: In-Memory Accelerator for Bulk Bitwise Operations Using Commodity DRAM Technology," MICRO, 2017

³Li+, "Pinatubo: A Processing-in-Memory Architecture for Bulk Bitwise Operations in Emerging Non-Volatile Memories," DAC, 2016

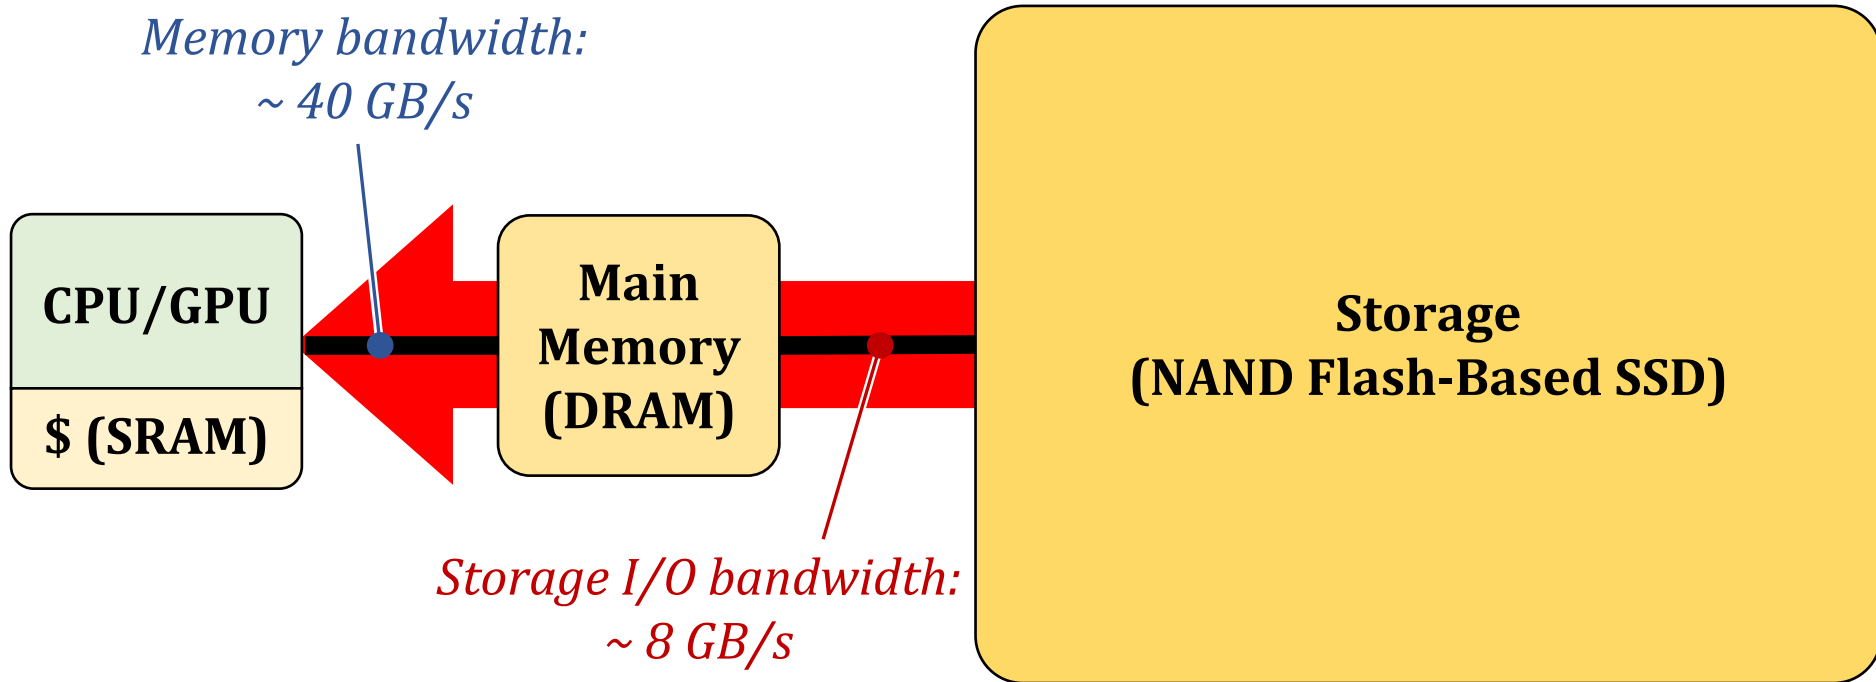
⁴Gu+, "Biscuit: A Framework for Near-Data Processing of Big Data Workloads," ISCA, 2016

⁵Gao+, "ParaBit: Processing Parallel Bitwise Operations in NAND Flash Memory Based SSDs," MICRO, 2021

In-Storage Processing (ISP)

- Uses **in-storage compute units** (embedded cores or FPGA) to send **only the computation results**



*Memory bandwidth:
~ 40 GB/s*



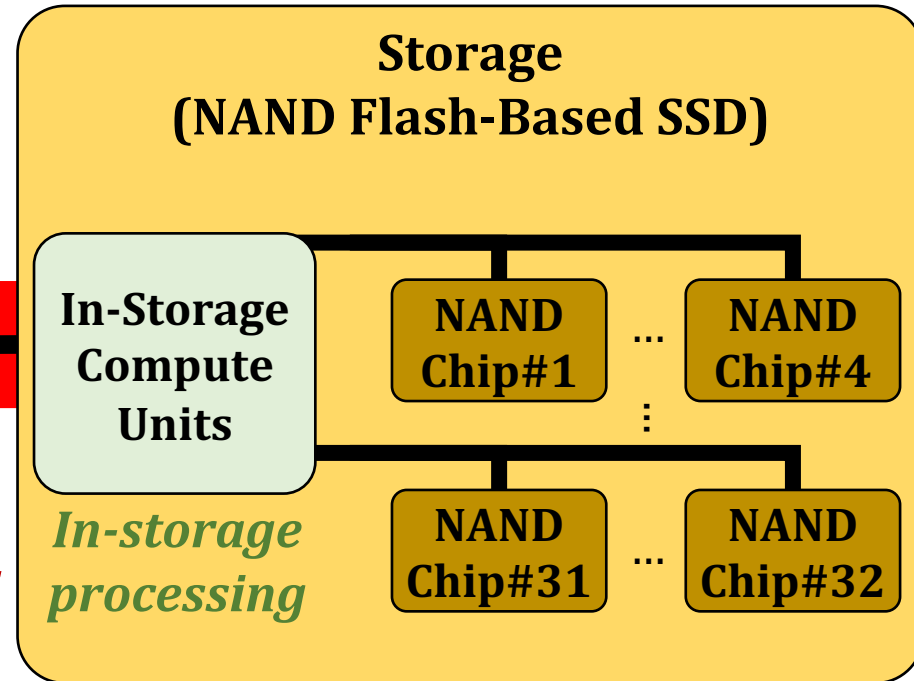
In-Storage Processing (ISP)

- Uses **in-storage compute units** (embedded cores or FPGA) to send **only the computation results**

*Memory bandwidth:
~ 40 GB/s*

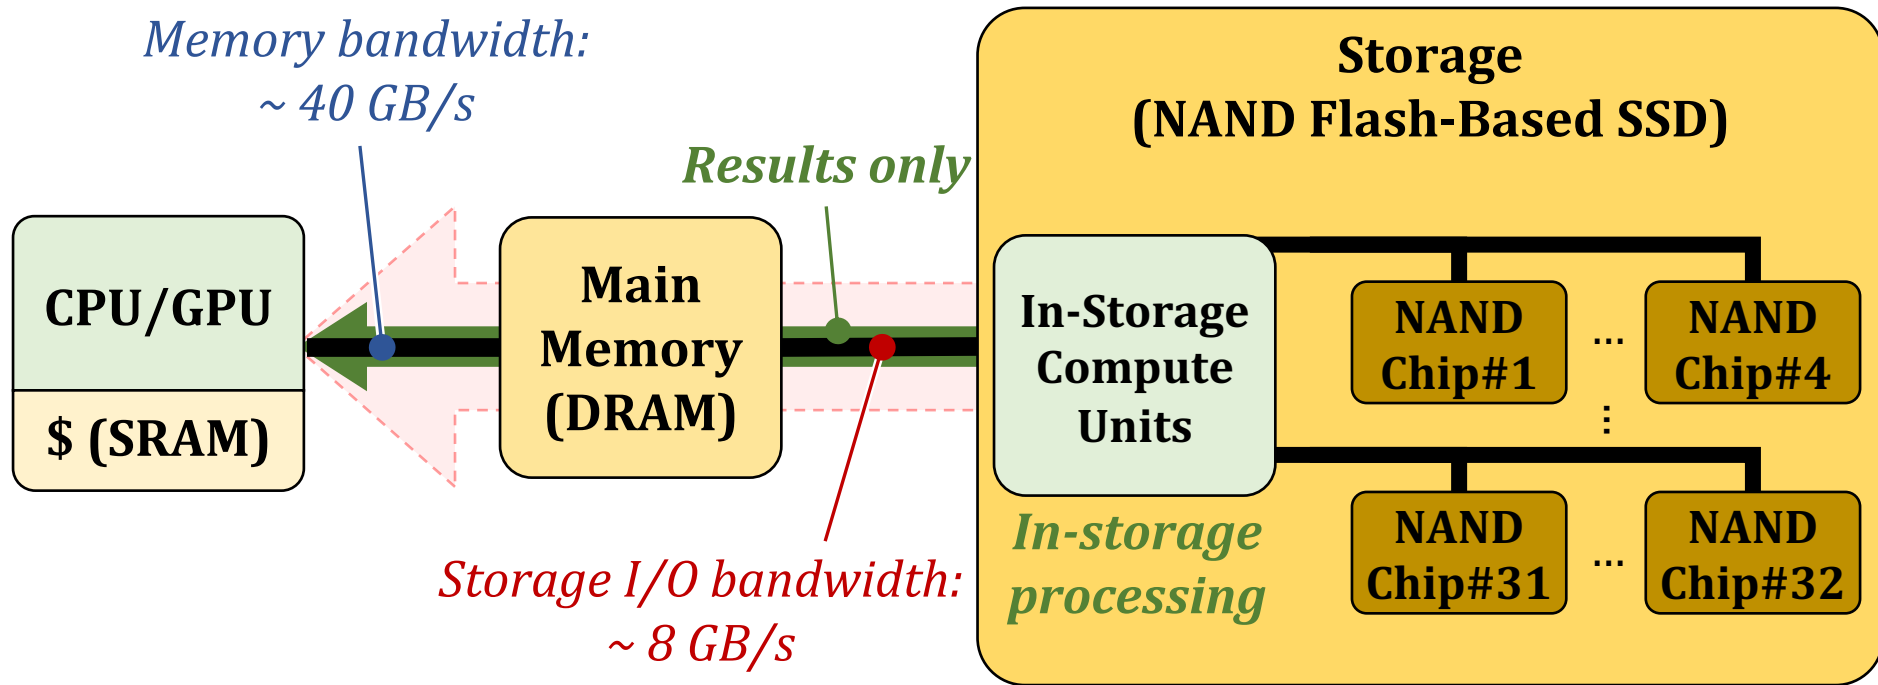


*Storage I/O bandwidth:
~ 8 GB/s*



In-Storage Processing (ISP)

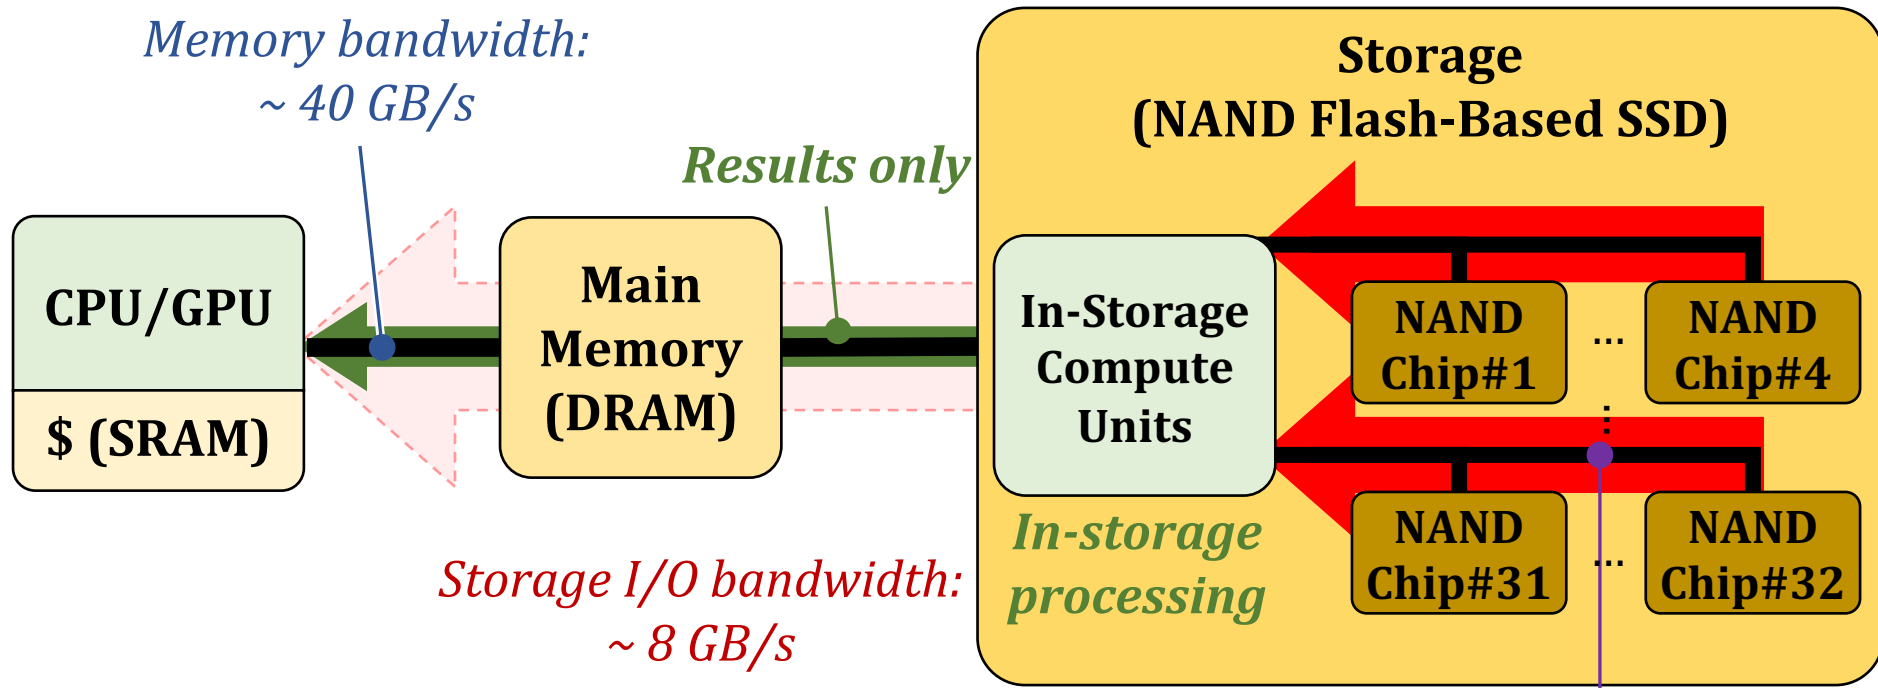
- Uses **in-storage compute units** (embedded cores or FPGA) to send **only the computation results**



ISP can mitigate data movement overhead by **reducing SSD-external data movement**

In-Storage Processing (ISP)

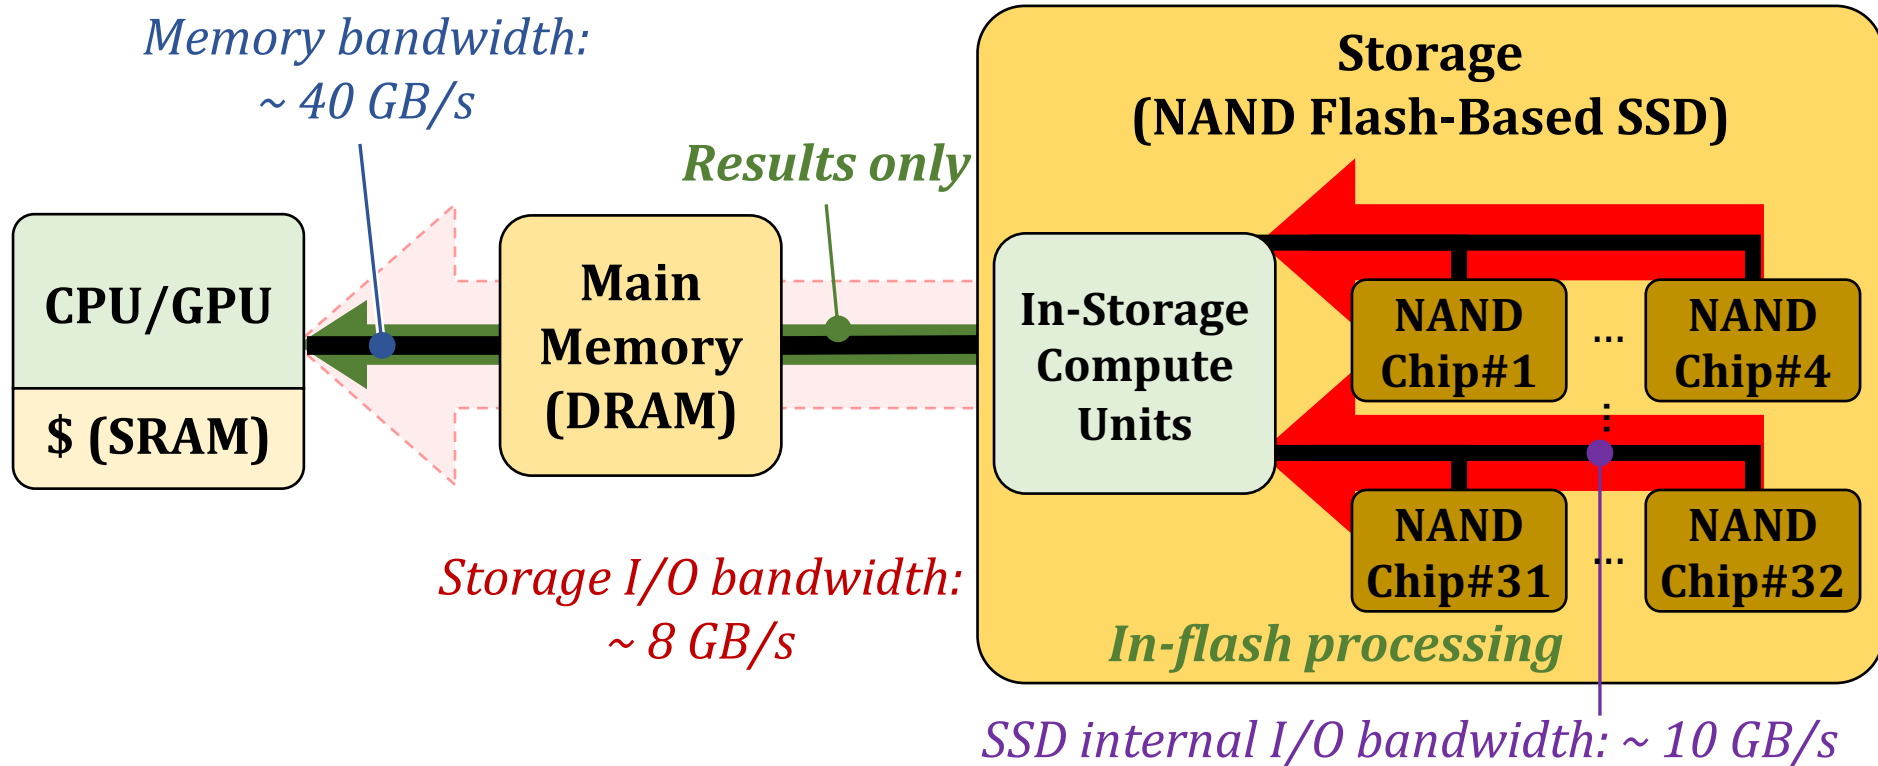
- Uses **in-storage compute units** (embedded cores or FPGA) to send **only the computation results**



SSD-internal bandwidth
becomes the **new bottleneck** in ISP

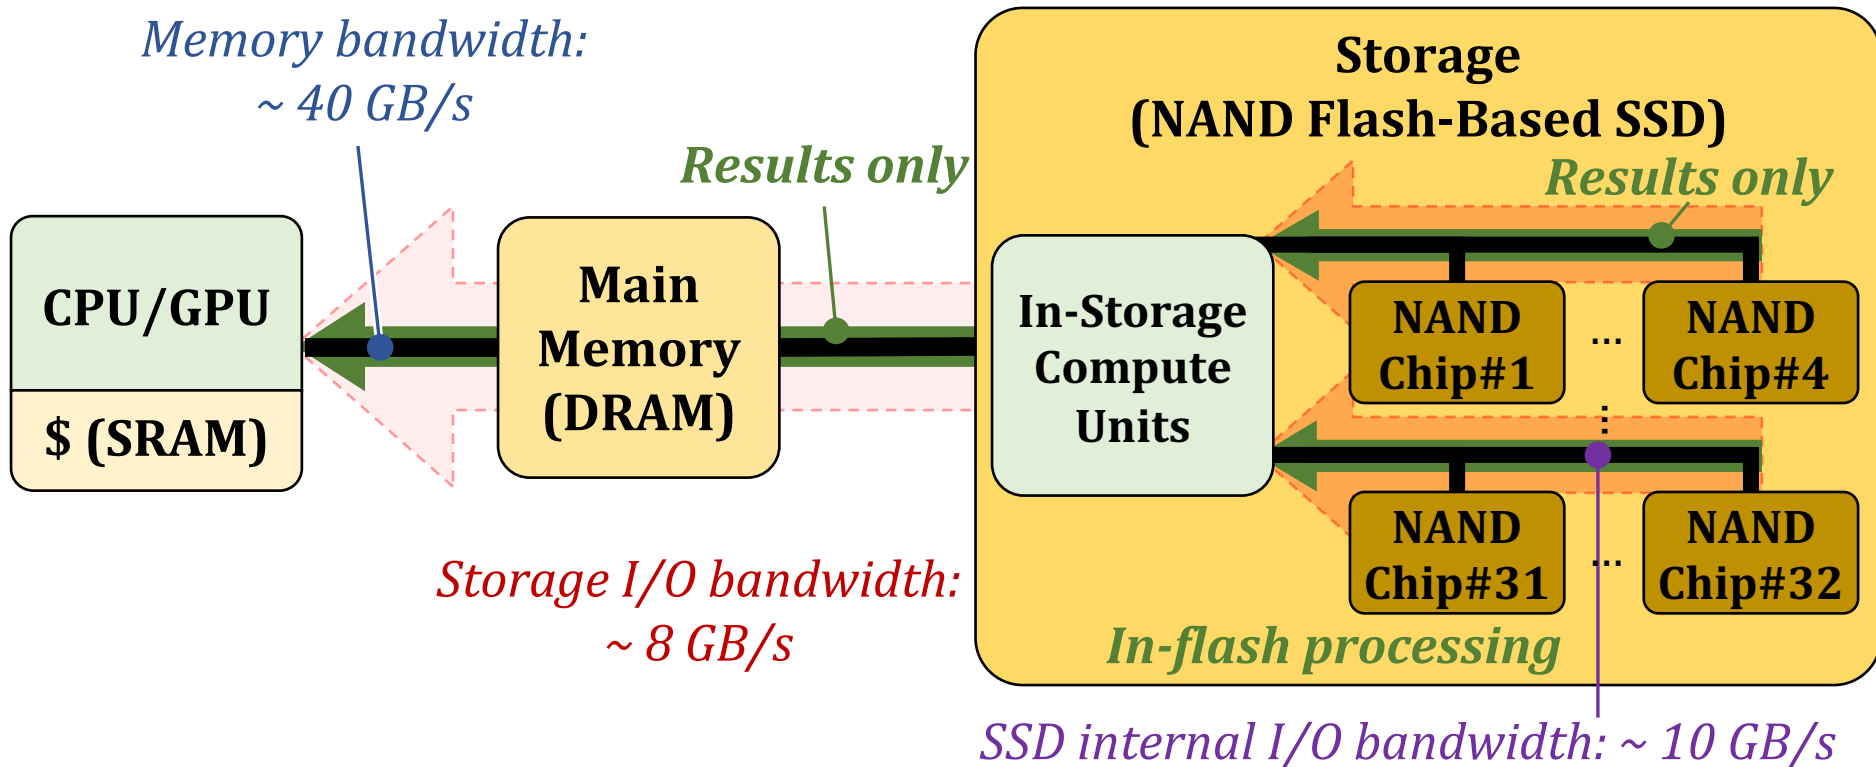
In-Flash Processing (IFP)

- Performs computation **inside** NAND flash chips



In-Flash Processing (IFP)

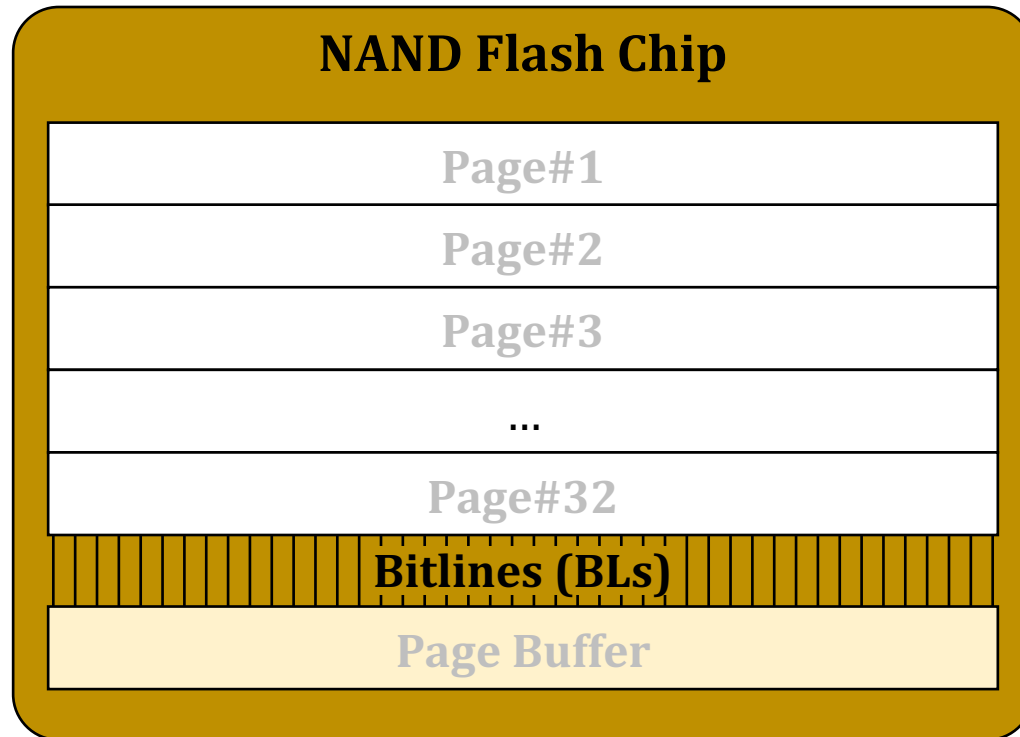
- Performs computation **inside** NAND flash chips



IFP fundamentally mitigates data movement

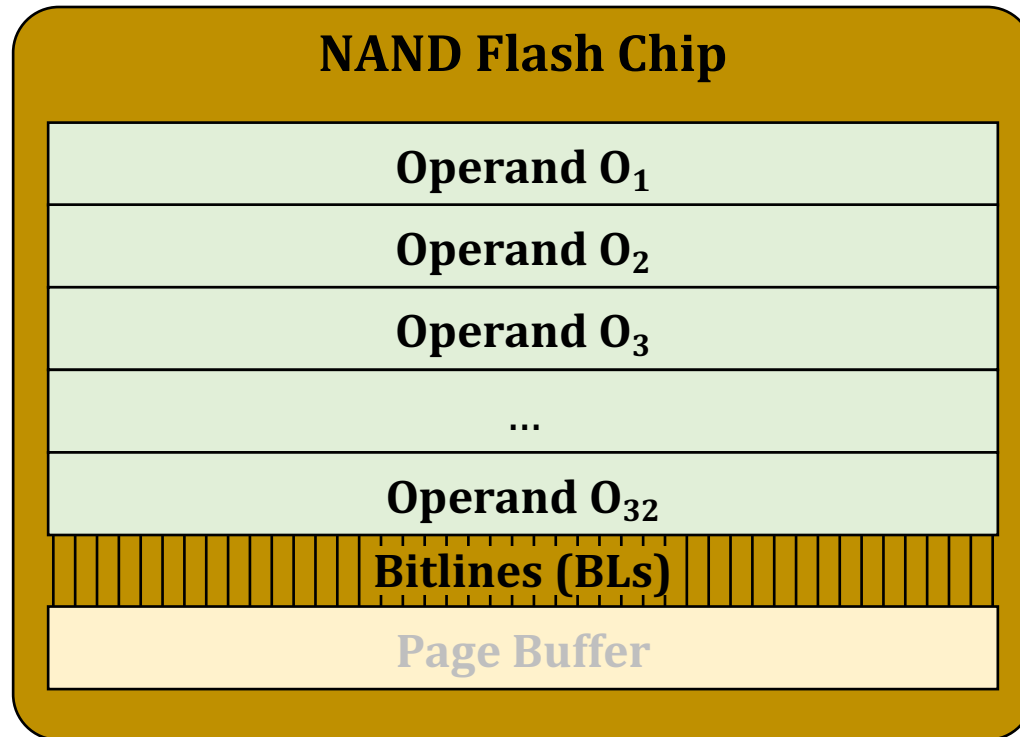
State-of-the-Art IFP Technique (1/2)

- **ParaBit** [Gao+, MICRO 2021]
 - Performs bulk bitwise operations **inside NAND flash chips** by **intelligently controlling the latching circuit of the page buffer**



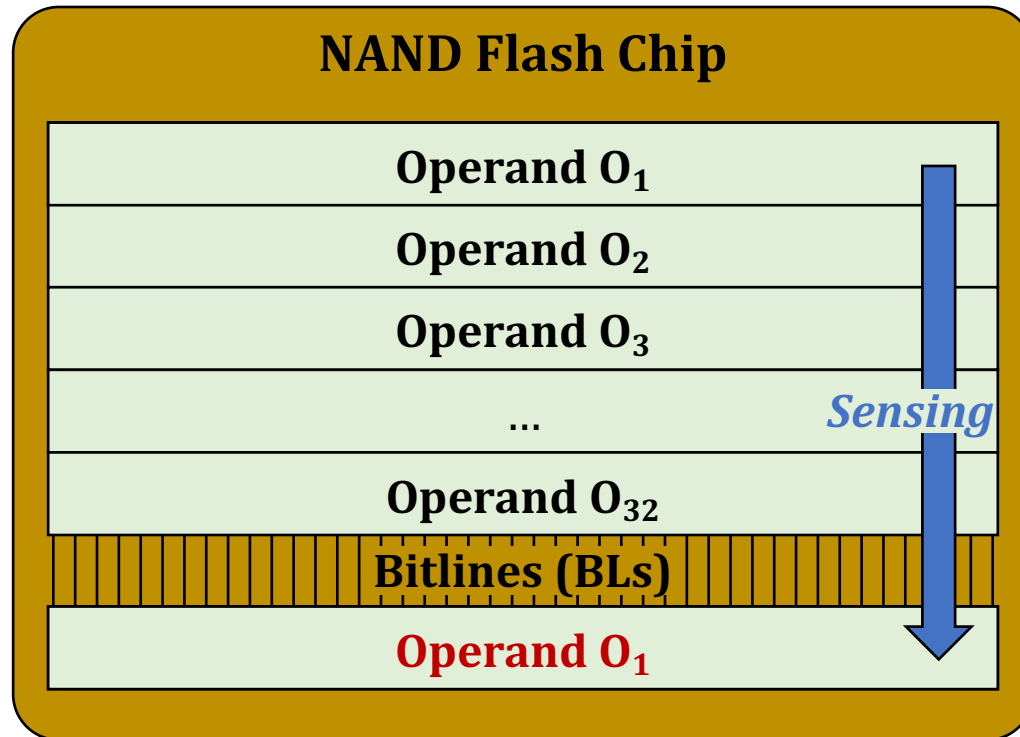
State-of-the-Art IFP Technique (1/2)

- **ParaBit** [Gao+, MICRO 2021]
 - Performs bulk bitwise operations **inside NAND flash chips** by **intelligently controlling the latching circuit of the page buffer**



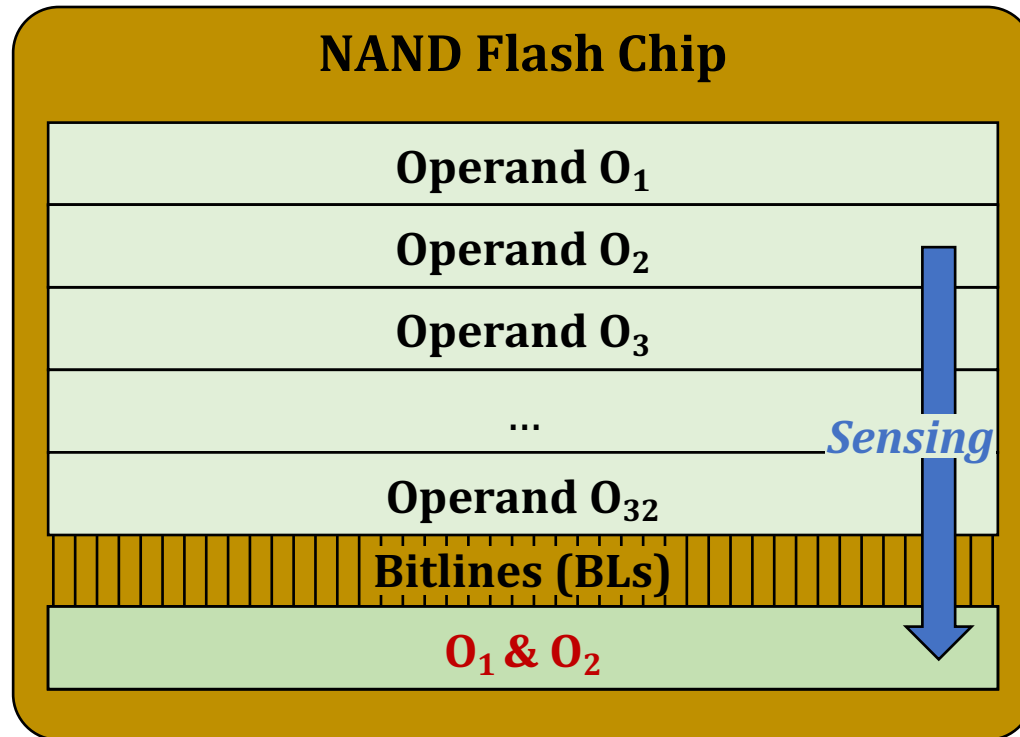
State-of-the-Art IFP Technique (1/2)

- **ParaBit** [Gao+, MICRO 2021]
 - Performs bulk bitwise operations **inside NAND flash chips** by **intelligently controlling the latching circuit of the page buffer**



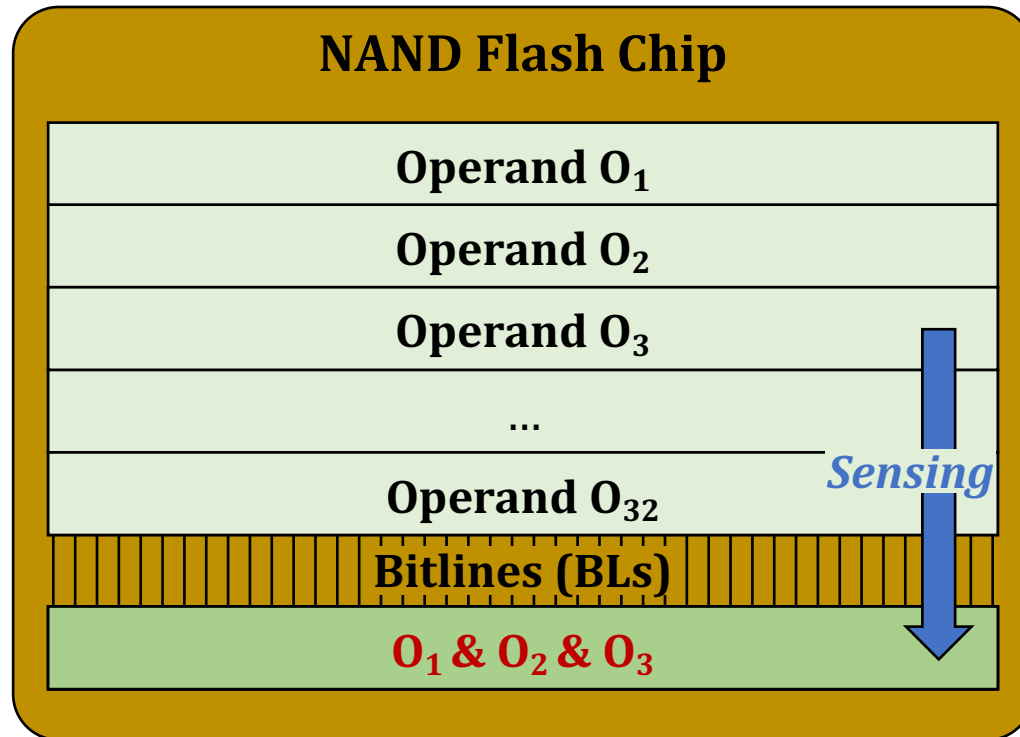
State-of-the-Art IFP Technique (1/2)

- **ParaBit** [Gao+, MICRO 2021]
 - Performs bulk bitwise operations **inside NAND flash chips** by **intelligently controlling the latching circuit of the page buffer**



State-of-the-Art IFP Technique (1/2)

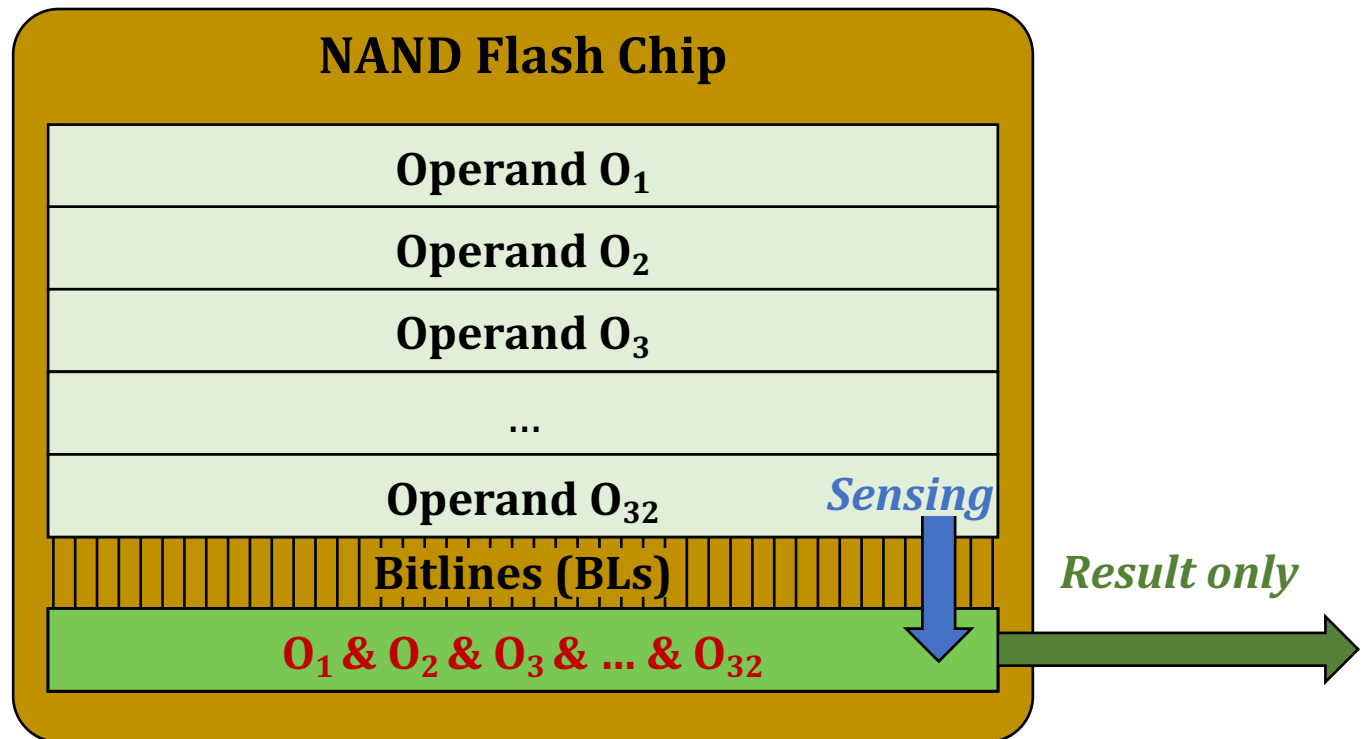
- **ParaBit** [Gao+, MICRO 2021]
 - Performs bulk bitwise operations **inside NAND flash chips** by **intelligently controlling the latching circuit of the page buffer**



State-of-the-Art IFP Technique (1/2)

- **ParaBit** [Gao+, MICRO 2021]

- Performs bulk bitwise operations **inside NAND flash chips** by **intelligently controlling the latching circuit of the page buffer**



State-of-the-Art IFP Technique (1/2)

- ParaBit [Gao+, MICRO 2021]

- Performs bulk bitwise operations inside NAND flash chips by intelligently controlling the latching circuit of the page buffer

NAND Flash Chip

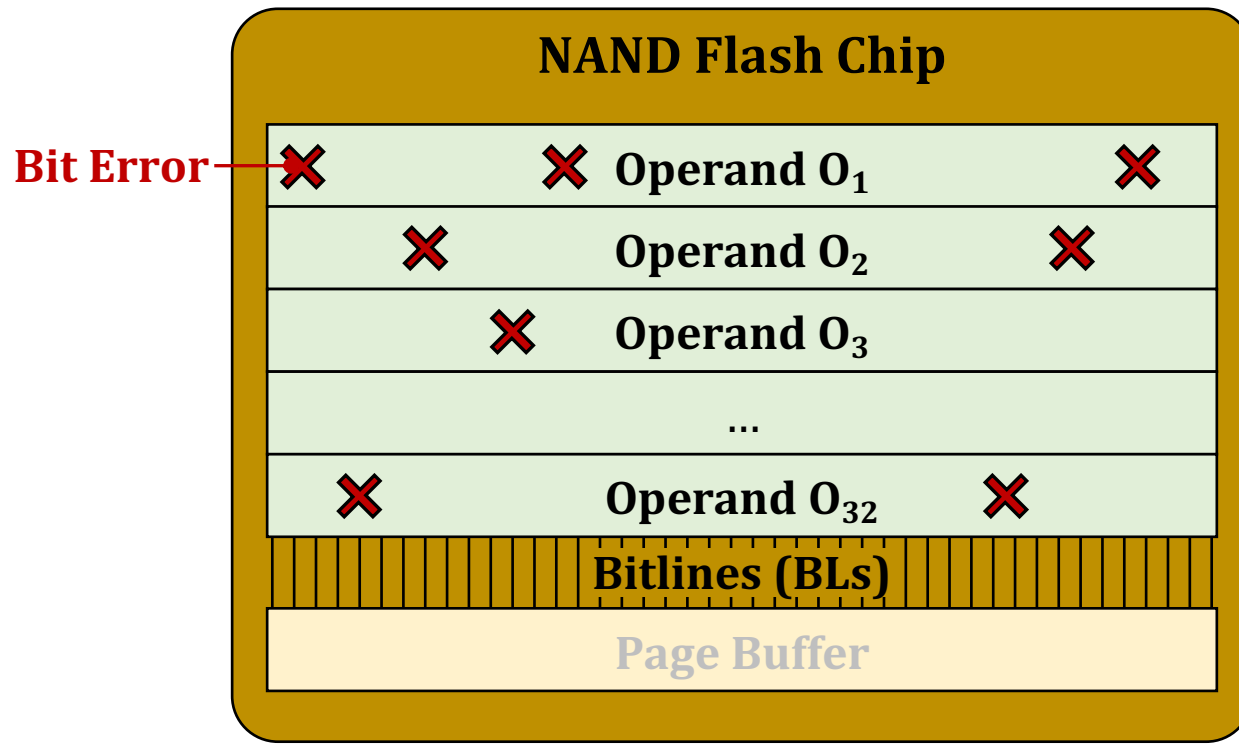
ParaBit significantly reduces data movement out of NAND flash chips

...

Serial sensing operations
(e.g., 32 sensing operations for 32-operand AND)
bottleneck performance and energy in ParaBit

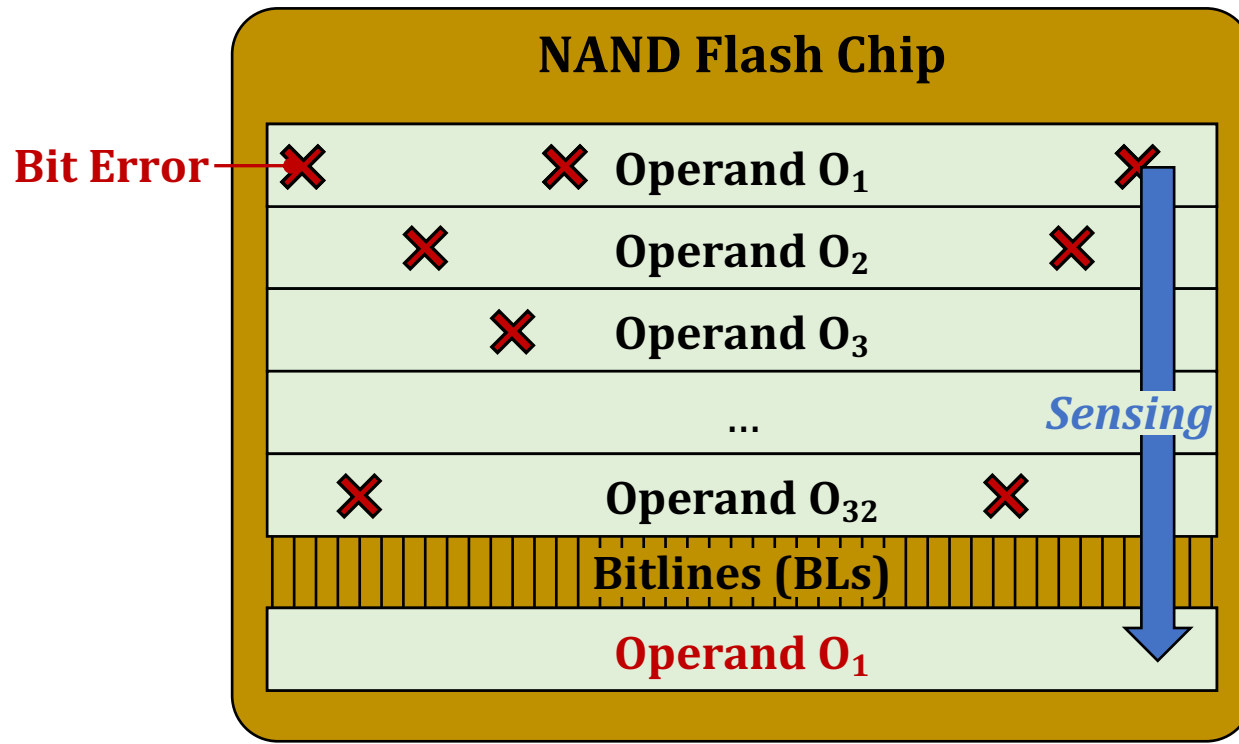
State-of-the-Art IFP Technique (2/2)

- **Limitations** of ParaBit [Gao+, MICRO 2021]
 - **Serial sensing operations** become the new bottleneck
 - **Erroneous computation results** due to the high raw bit-error rate of NAND flash memory



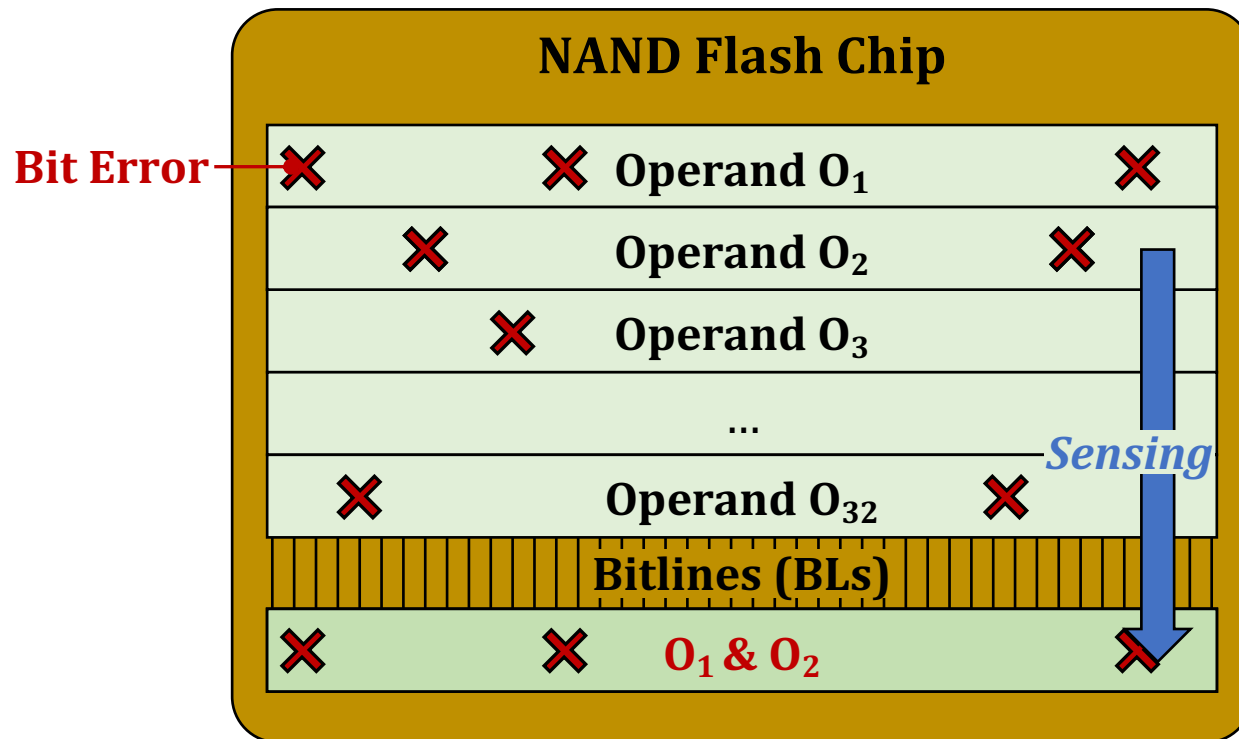
State-of-the-Art IFP Technique (2/2)

- **Limitations** of ParaBit [Gao+, MICRO 2021]
 - **Serial sensing operations** become the new bottleneck
 - **Erroneous computation results** due to the high raw bit-error rate of NAND flash memory



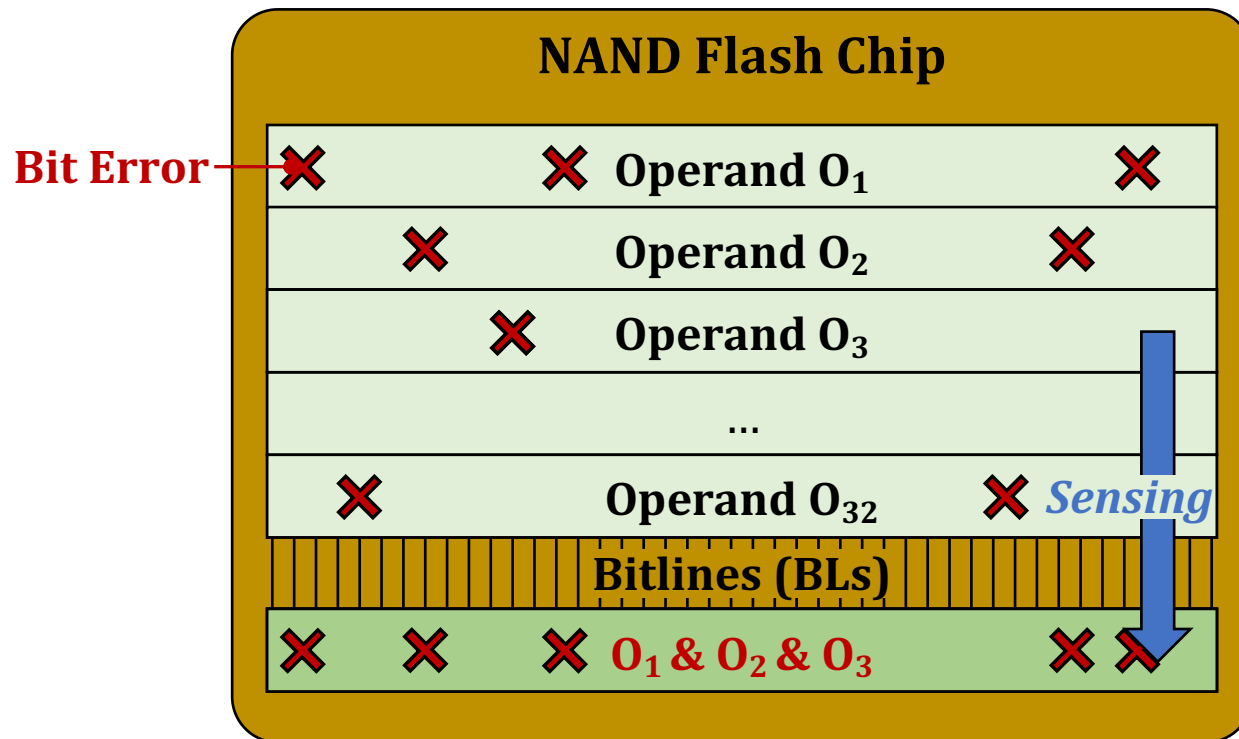
State-of-the-Art IFP Technique (2/2)

- **Limitations** of ParaBit [Gao+, MICRO 2021]
 - **Serial sensing operations** become the new bottleneck
 - **Erroneous computation results** due to the high raw bit-error rate of NAND flash memory



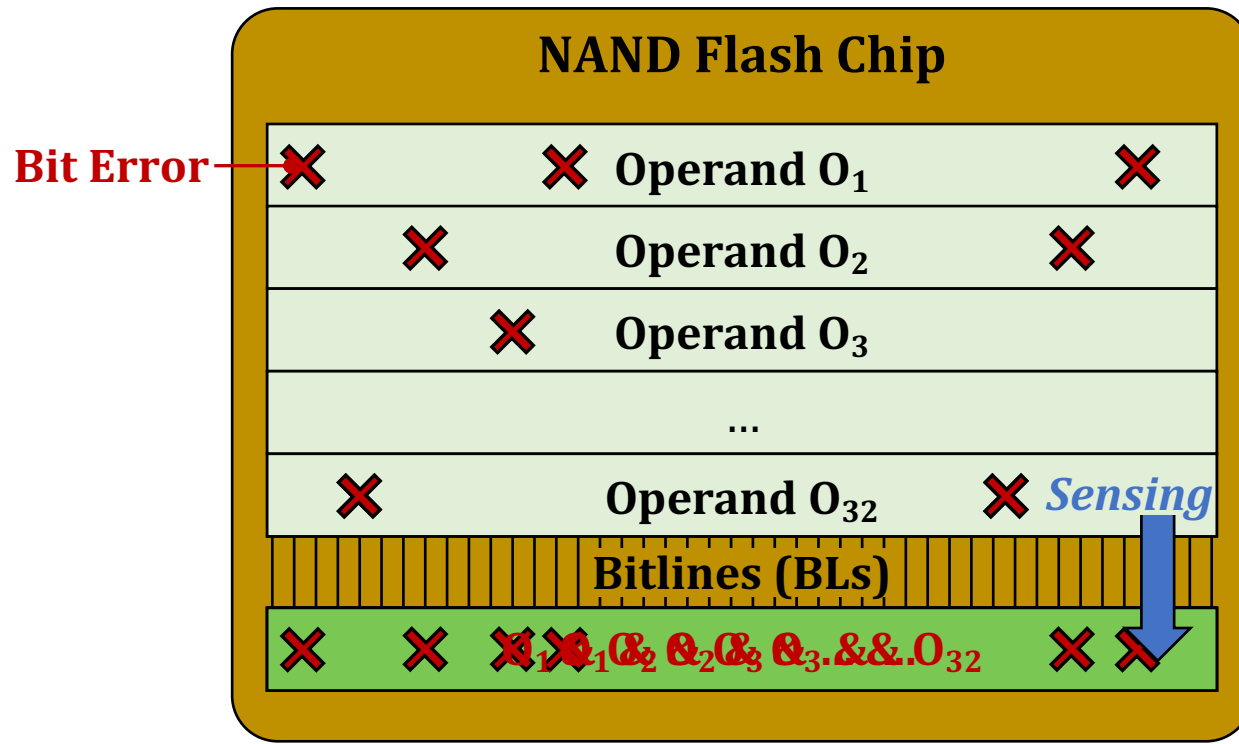
State-of-the-Art IFP Technique (2/2)

- **Limitations** of ParaBit [Gao+, MICRO 2021]
 - **Serial sensing operations** become the new bottleneck
 - **Erroneous computation results** due to the high raw bit-error rate of NAND flash memory



State-of-the-Art IFP Technique (2/2)

- **Limitations** of ParaBit [Gao+, MICRO 2021]
 - **Serial sensing operations** become the new bottleneck
 - **Erroneous computation results** due to the high raw bit-error rate of NAND flash memory



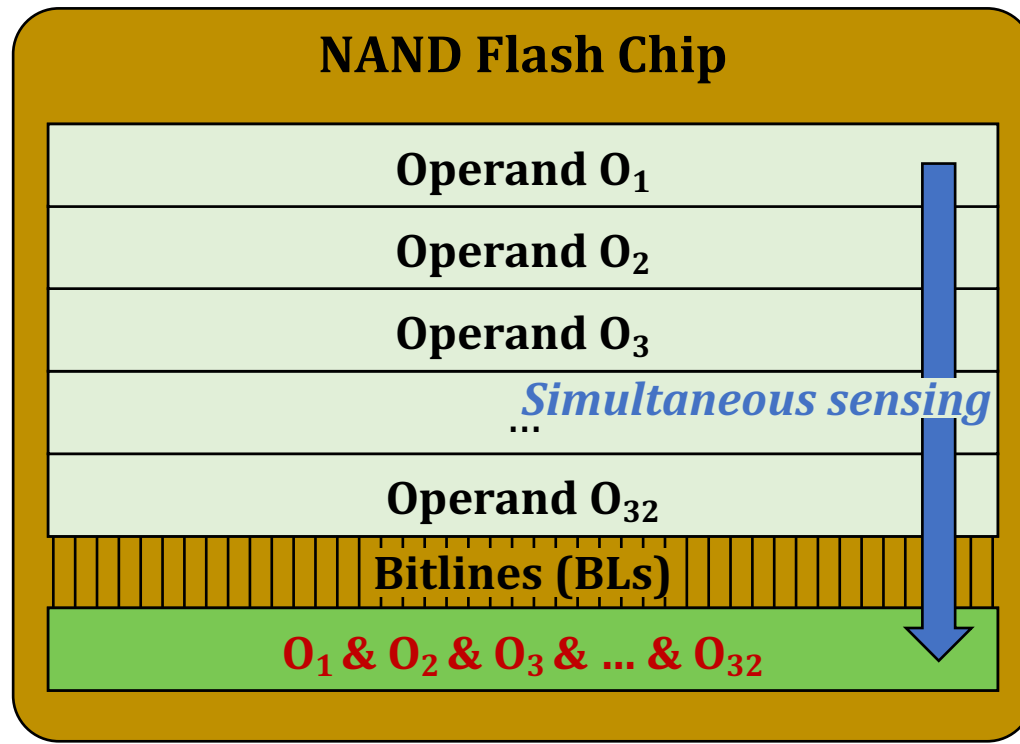
Our Goals

Address the new bottleneck of IFP
(serial sensing of operands)

Make IFP reliable
(provide accurate computation results)

Our Proposal: Flash-Cosmos

- **Flash-Cosmos** enables
 - Computation on **multiple operands** with a **single sensing operation**
 - **Accurate computation results** by eliminating raw bit errors in stored data

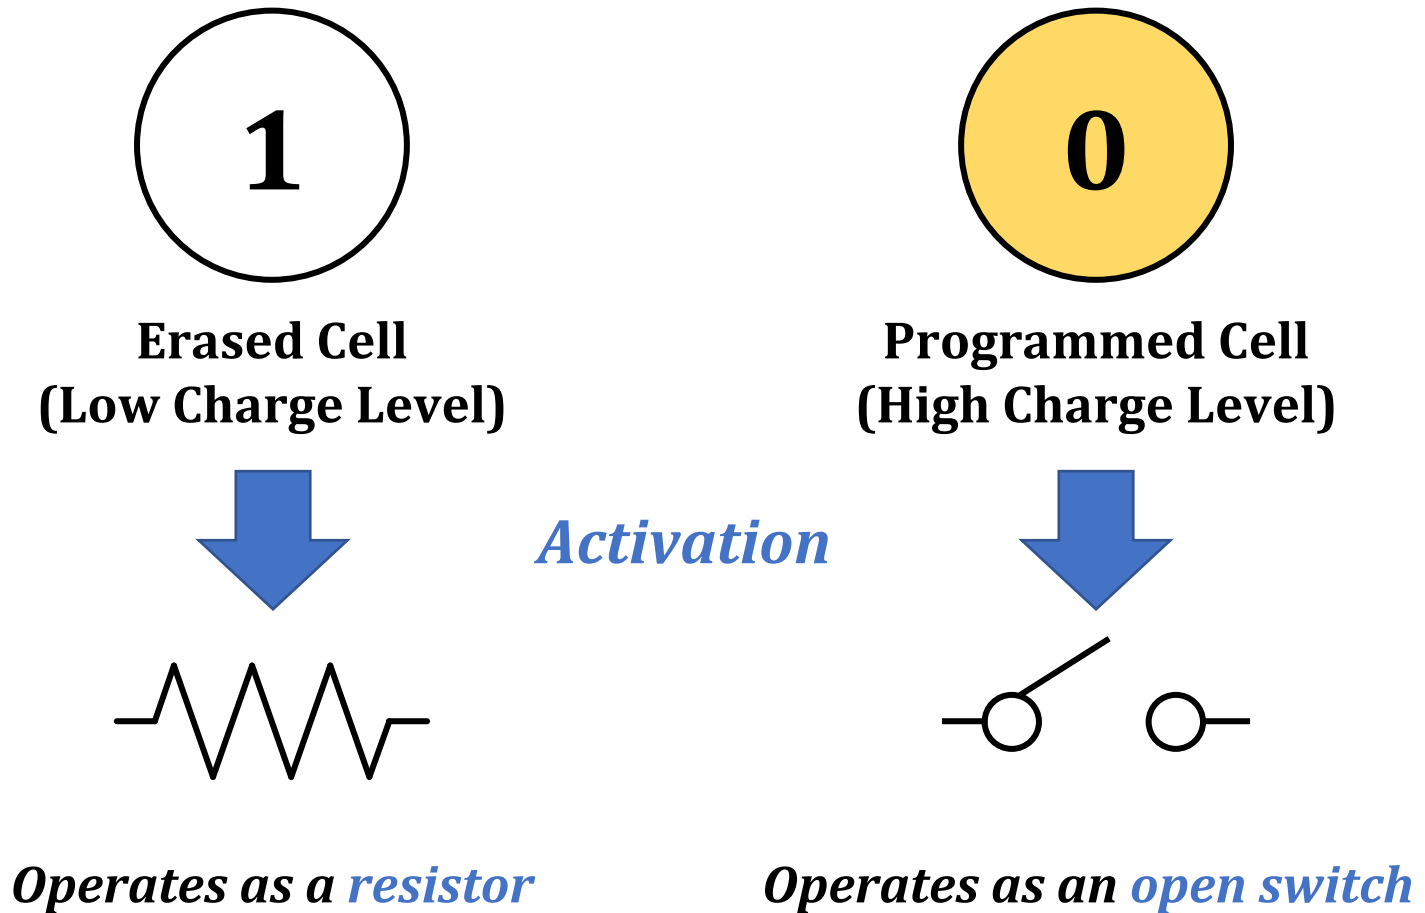


Talk Outline

- Problem, Goals & Key Idea
- Background
- Flash-Cosmos: Computation with One-Shot Multi-Operand Sensing
- Evaluation
- Summary


NAND Flash Basics: A Flash Cell

- A flash cell stores data by adjusting the **amount of charge** in the cell



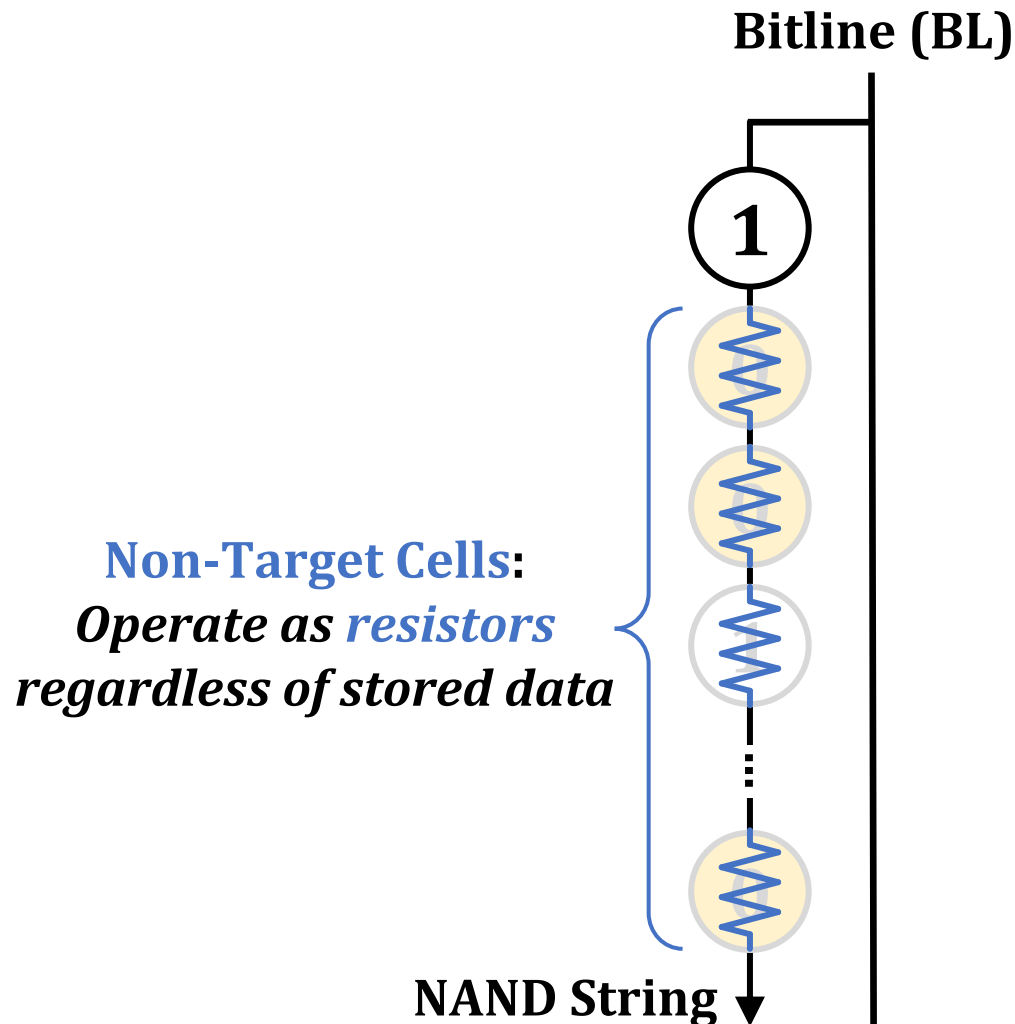
NAND Flash Basics: A NAND String

- A set of flash cells are **serially connected**, forming a **NAND string**



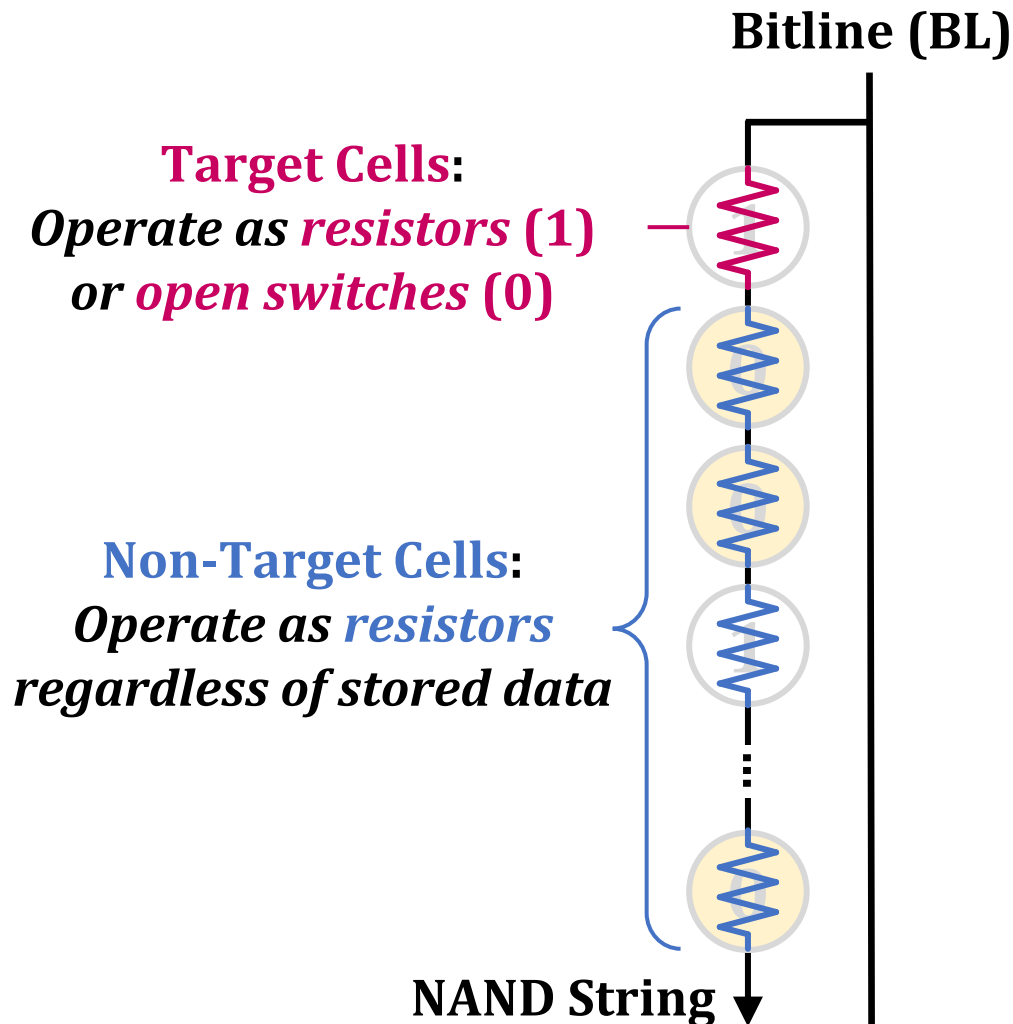
NAND Flash Basics: Read Mechanism (1/3)

- NAND flash memory reads data by [checking the bitline current](#)



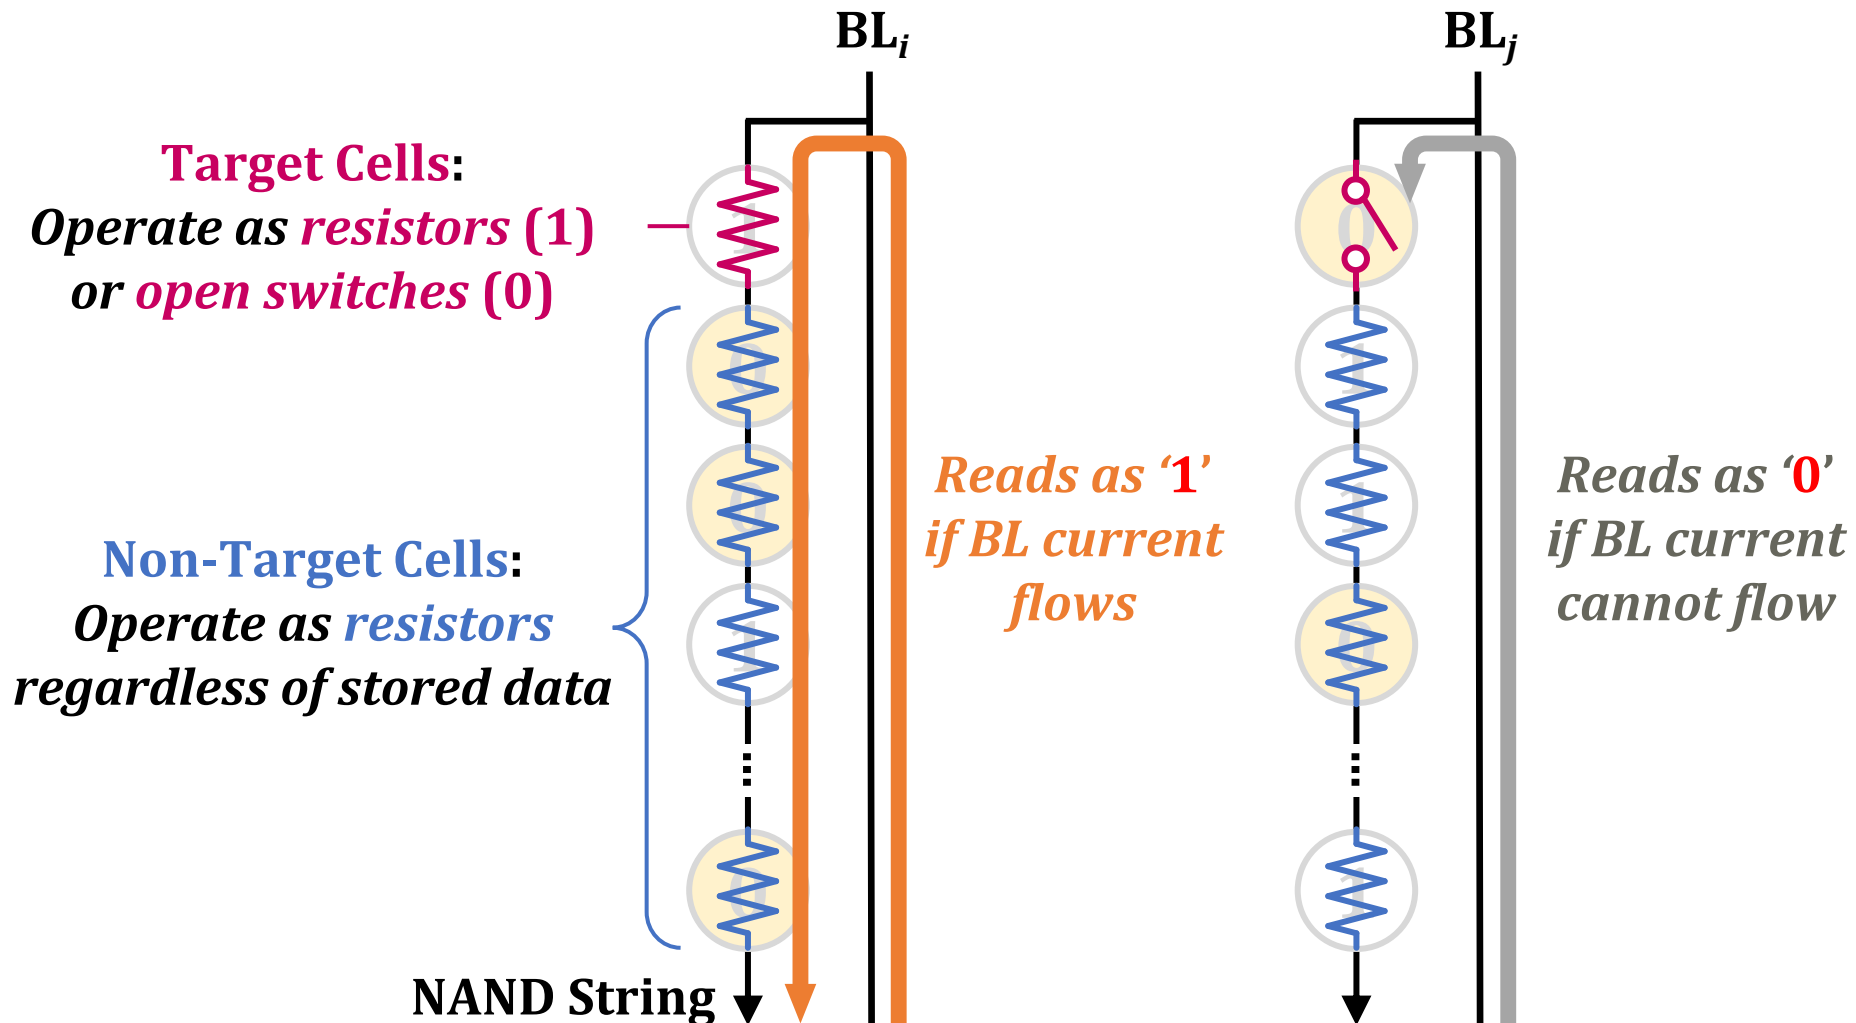
NAND Flash Basics: Read Mechanism (2/3)

- NAND flash memory reads data by [checking the bitline current](#)




NAND Flash Basics: Read Mechanism (3/3)

- NAND flash memory reads data by checking the bitline current




NAND Flash Basics: A NAND Flash Block

- NAND strings connected to different bitlines comprise a **block**



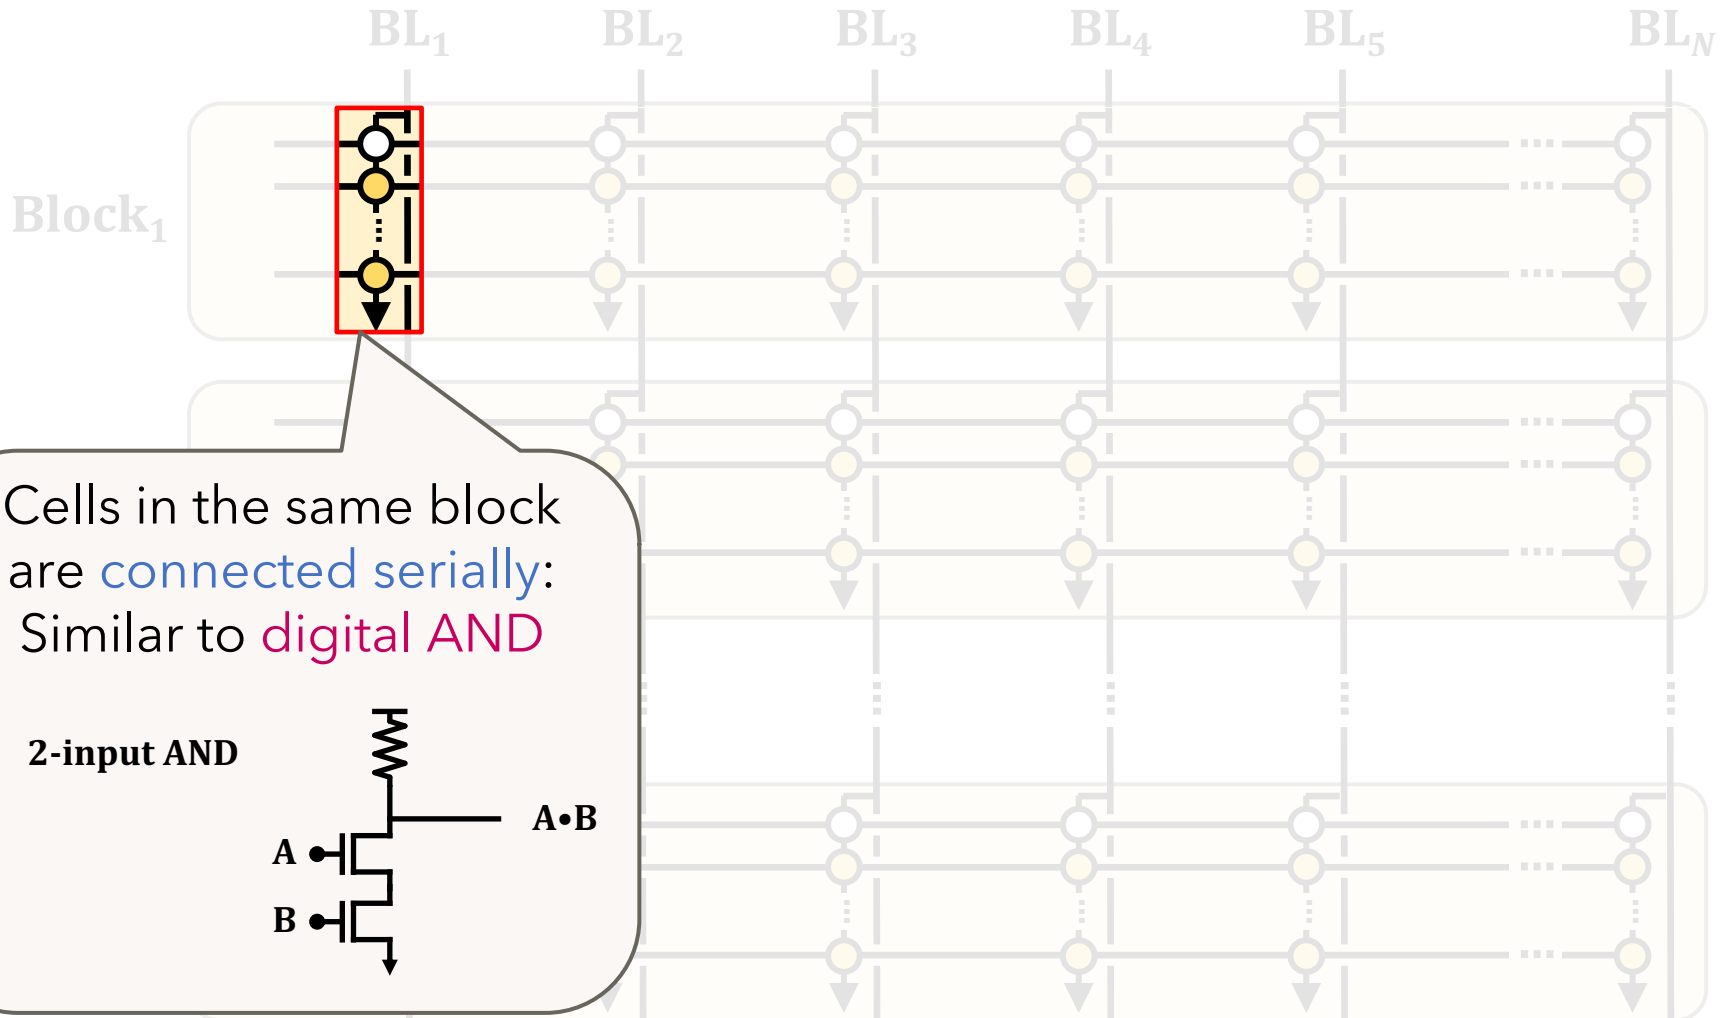
NAND Flash Basics: Block Organization

- A large number of blocks share the same bitlines



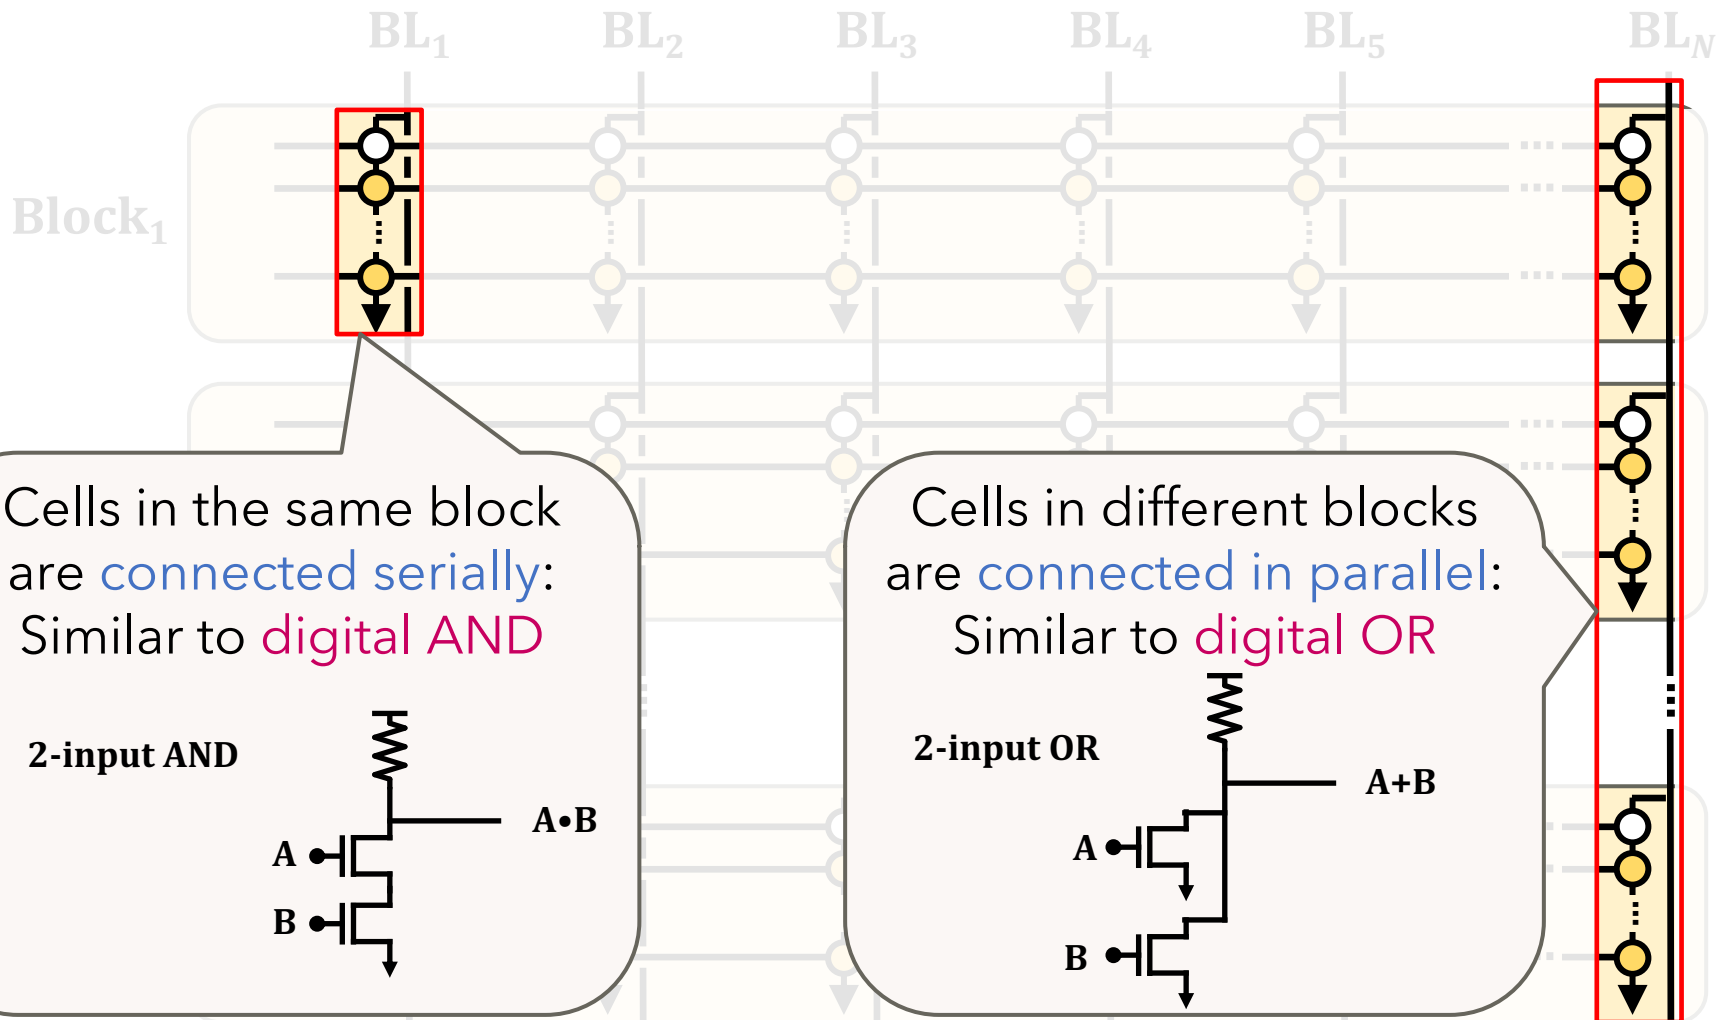
Similarity to Digital Logic Gates

- A large number of blocks share the same bitlines



Similarity to Digital Logic Gates

- A large number of blocks share the same bitlines.



Talk Outline

- Problem, Goals & Key Idea
- Background
- Flash-Cosmos: Computation with One-Shot Multi-Operand Sensing
- Evaluation
- Summary

Key Ideas



Multi-Wordline Sensing (MWS)
to enable in-flash bulk bitwise operations
via a single sensing operation




Enhanced SLC-Mode Programming (ESP)
to eliminate raw bit errors in stored data
(and thus in computation results)

Multi-Wordline Sensing (MWS): Bitwise AND

- **Intra-Block MWS:**

Simultaneously activates multiple WLs in the same block

→ Bitwise AND of the stored data in the WLs

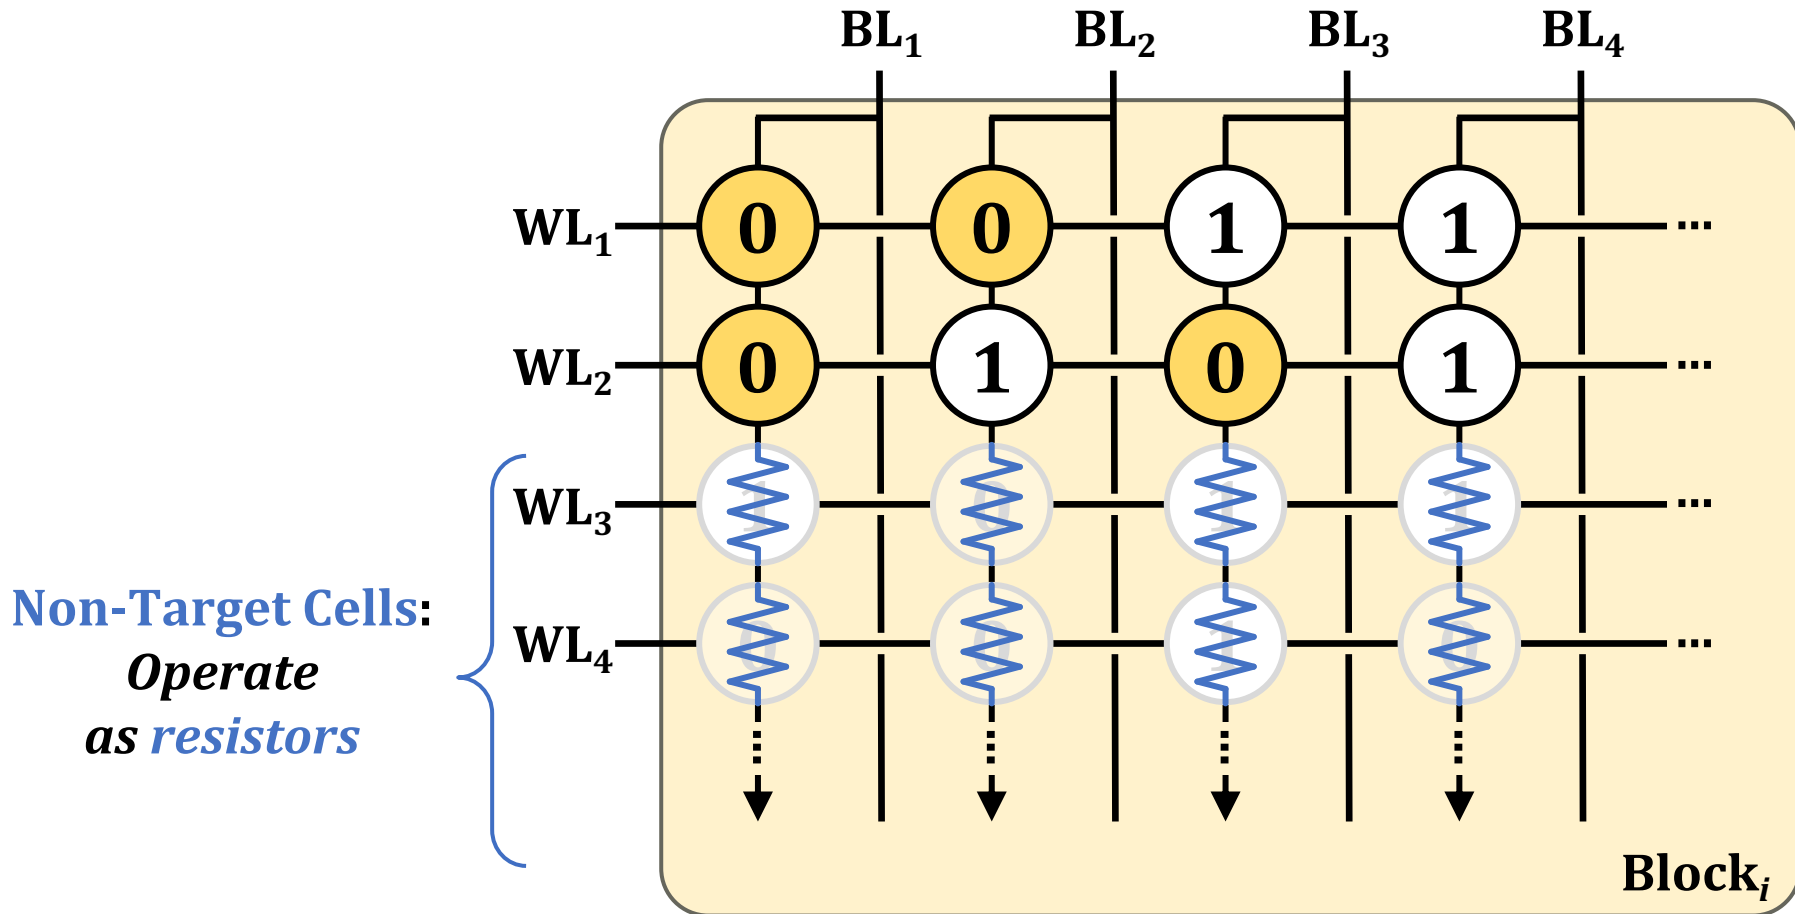


Multi-Wordline Sensing (MWS): Bitwise AND

- **Intra-Block MWS:**

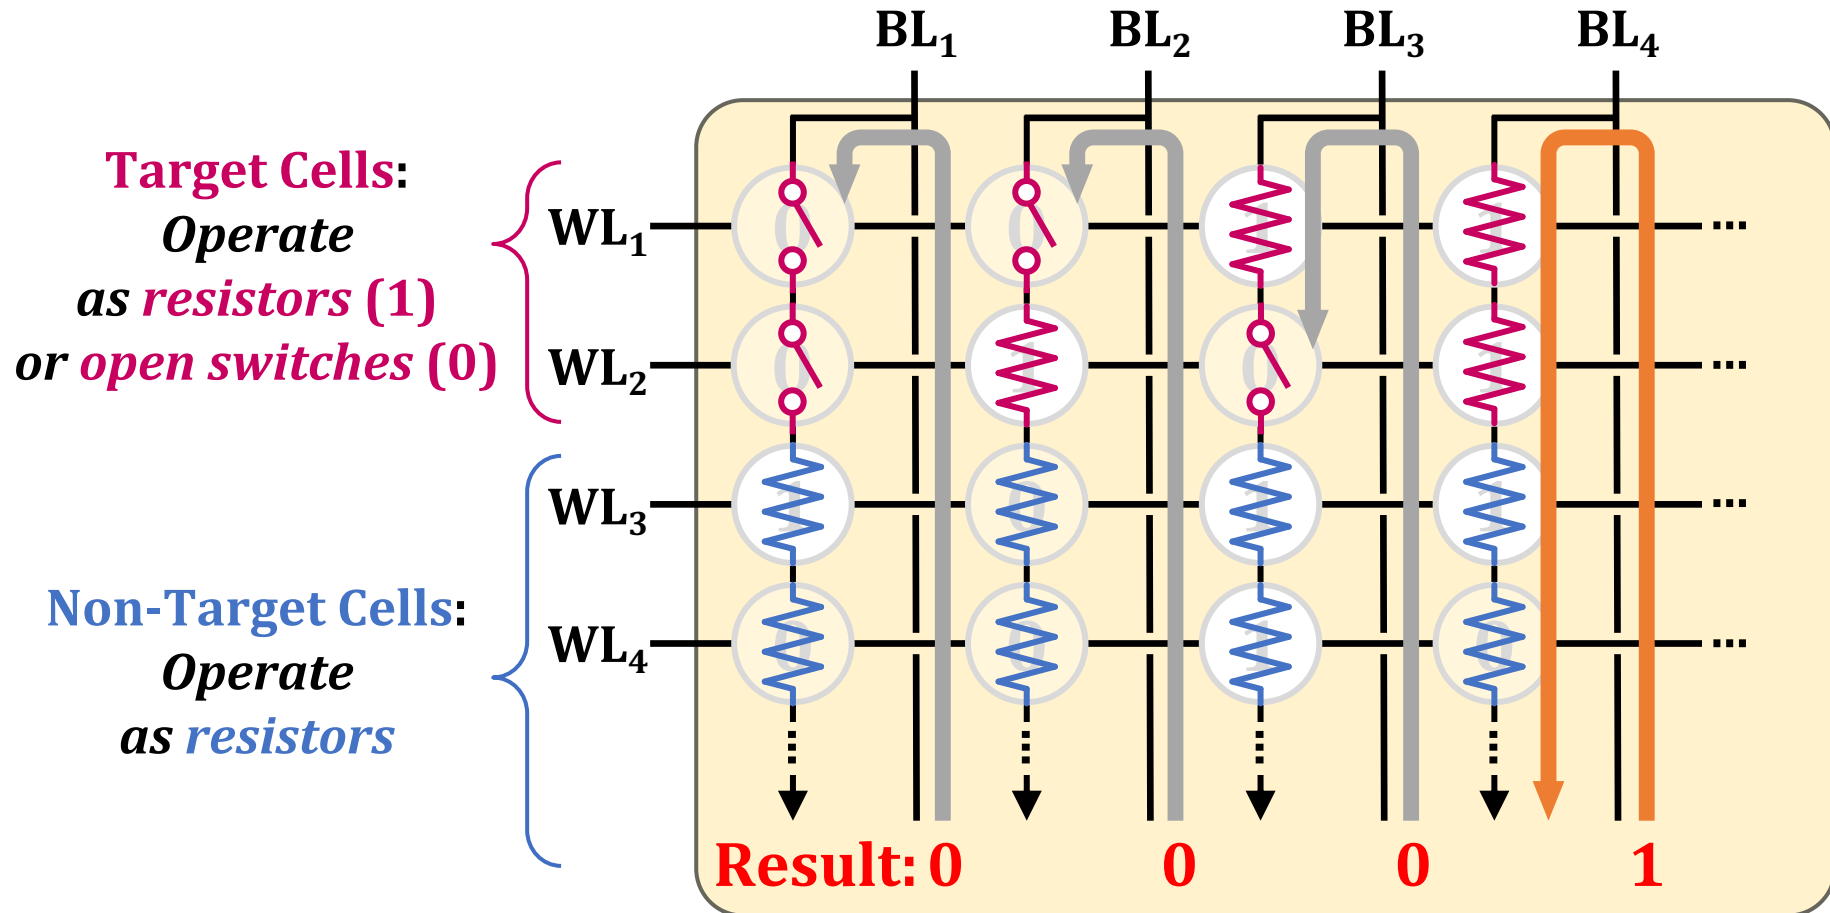
Simultaneously activates multiple WLs in the same block

→ Bitwise AND of the stored data in the WLs



Multi-Wordline Sensing (MWS): Bitwise AND

- **Intra-Block MWS:**
Simultaneously activates multiple WLs in the same block
→ Bitwise AND of the stored data in the WLs



Multi-Wordline Sensing (MWS): Bitwise AND

- **Intra-Block MWS:**


Simultaneously activates multiple WLs in the same block

→ Bitwise AND of the stored data in the WLs

Target Cell:

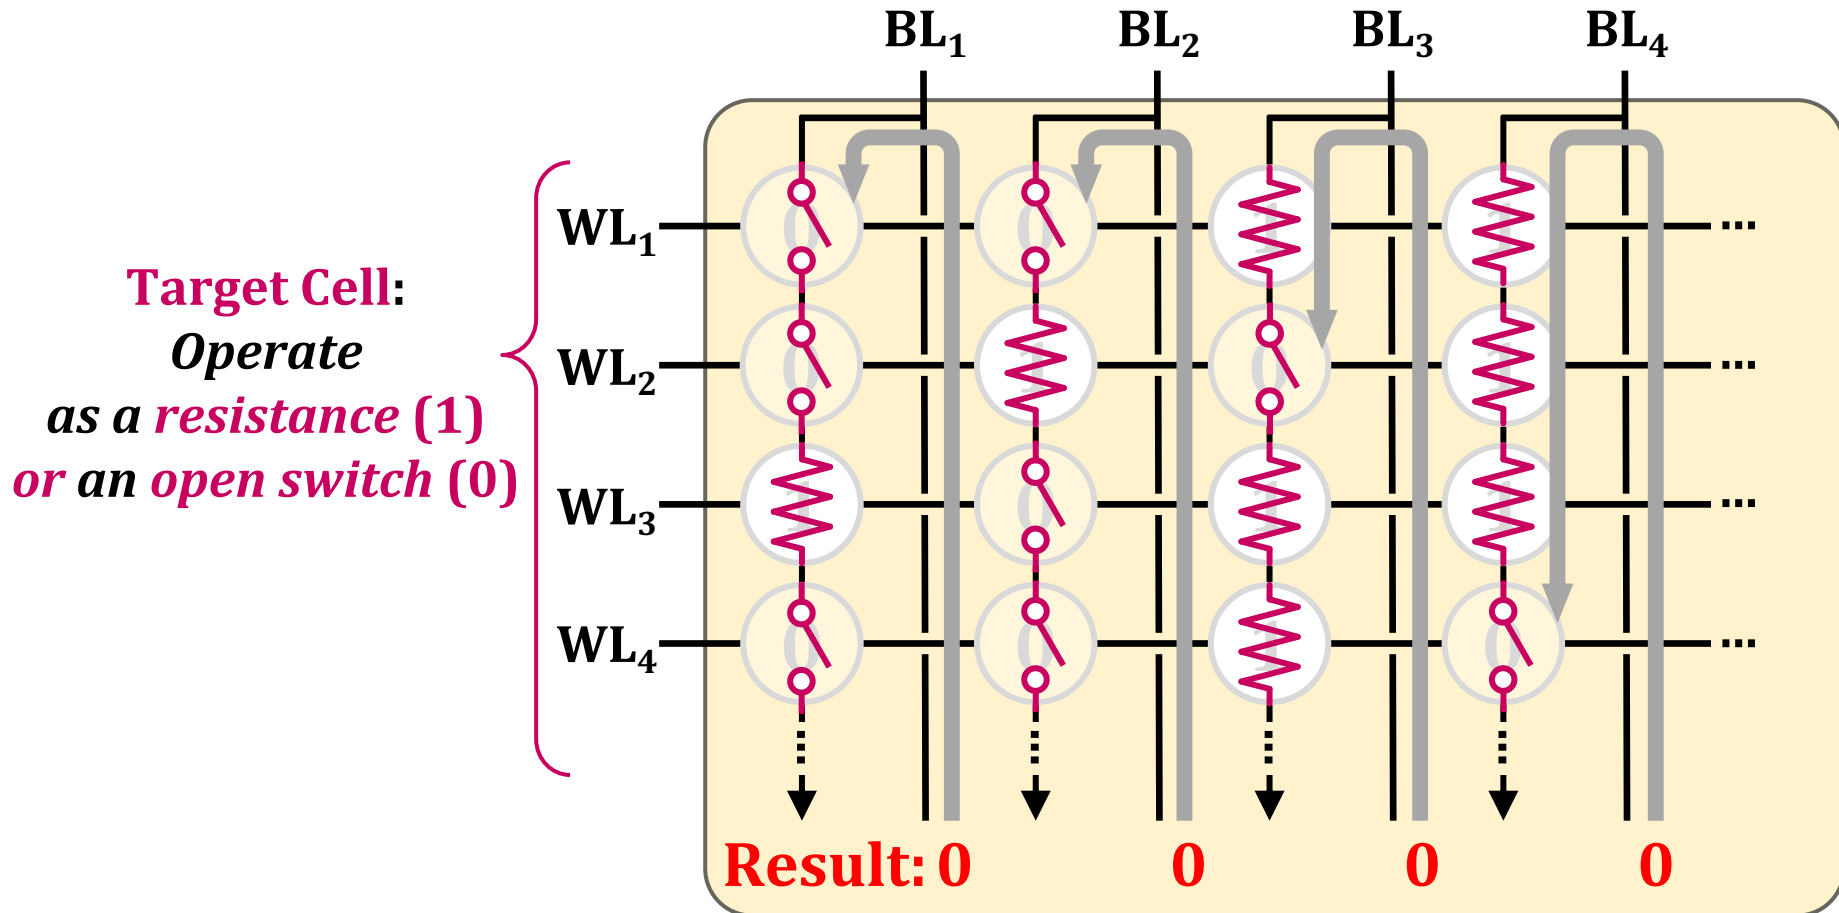
A bitline reads as '1' only when all the target cells store '1'
→ Equivalent to the bitwise AND of all the target cells

Non-Target Cell:
*Operate
as a resistance*



Multi-Wordline Sensing (MWS): Bitwise AND

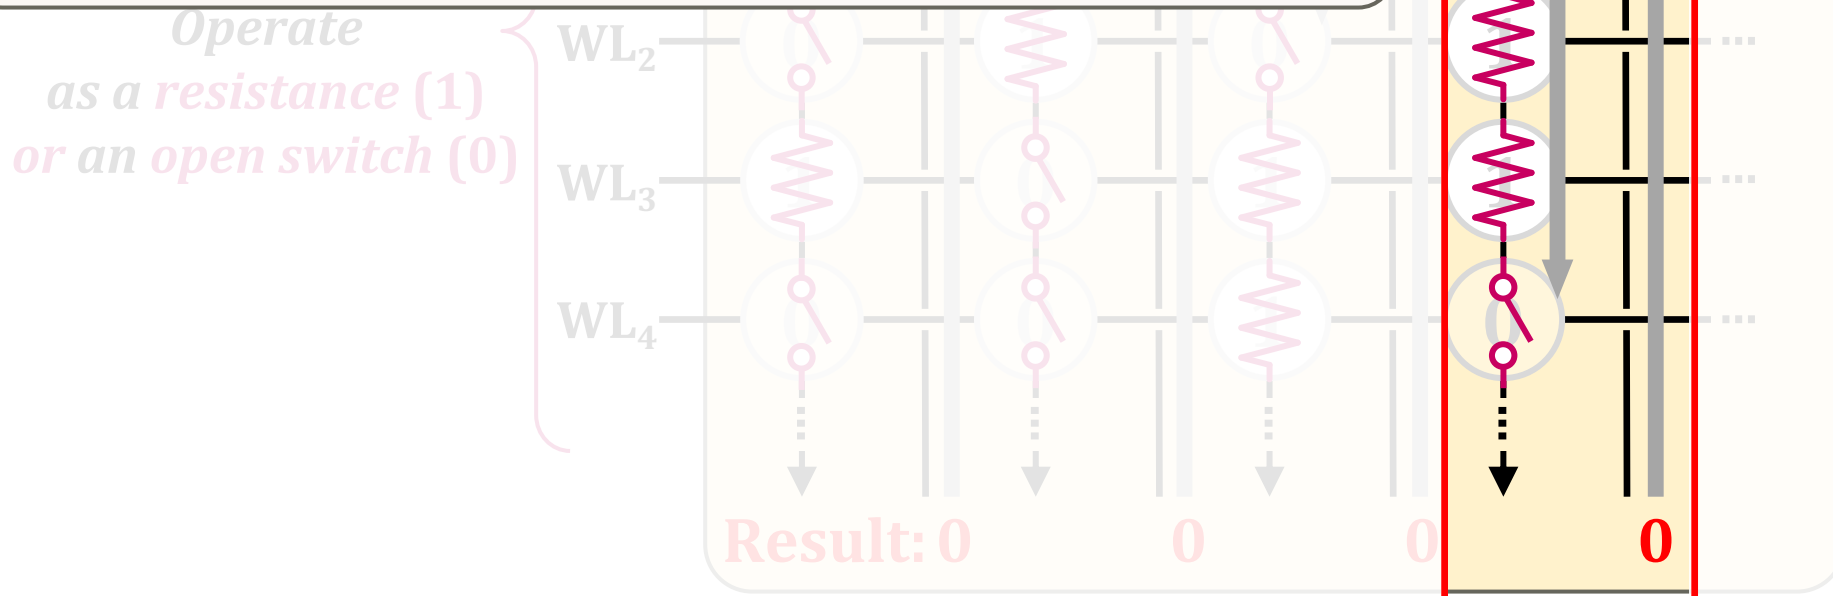
- **Intra-Block MWS:**
Simultaneously activates multiple WLs in the same block
→ Bitwise AND of the stored data in the WLs



Multi-Wordline Sensing (MWS): Bitwise AND

- **Intra-Block MWS:**
Simultaneously activates multiple WLs in the same block
→ Bitwise AND of the stored data in the WLs

A bitline reads as '**0**' when a single target cell stores '**0**'
→ Equivalent to the bitwise AND of all the target cells



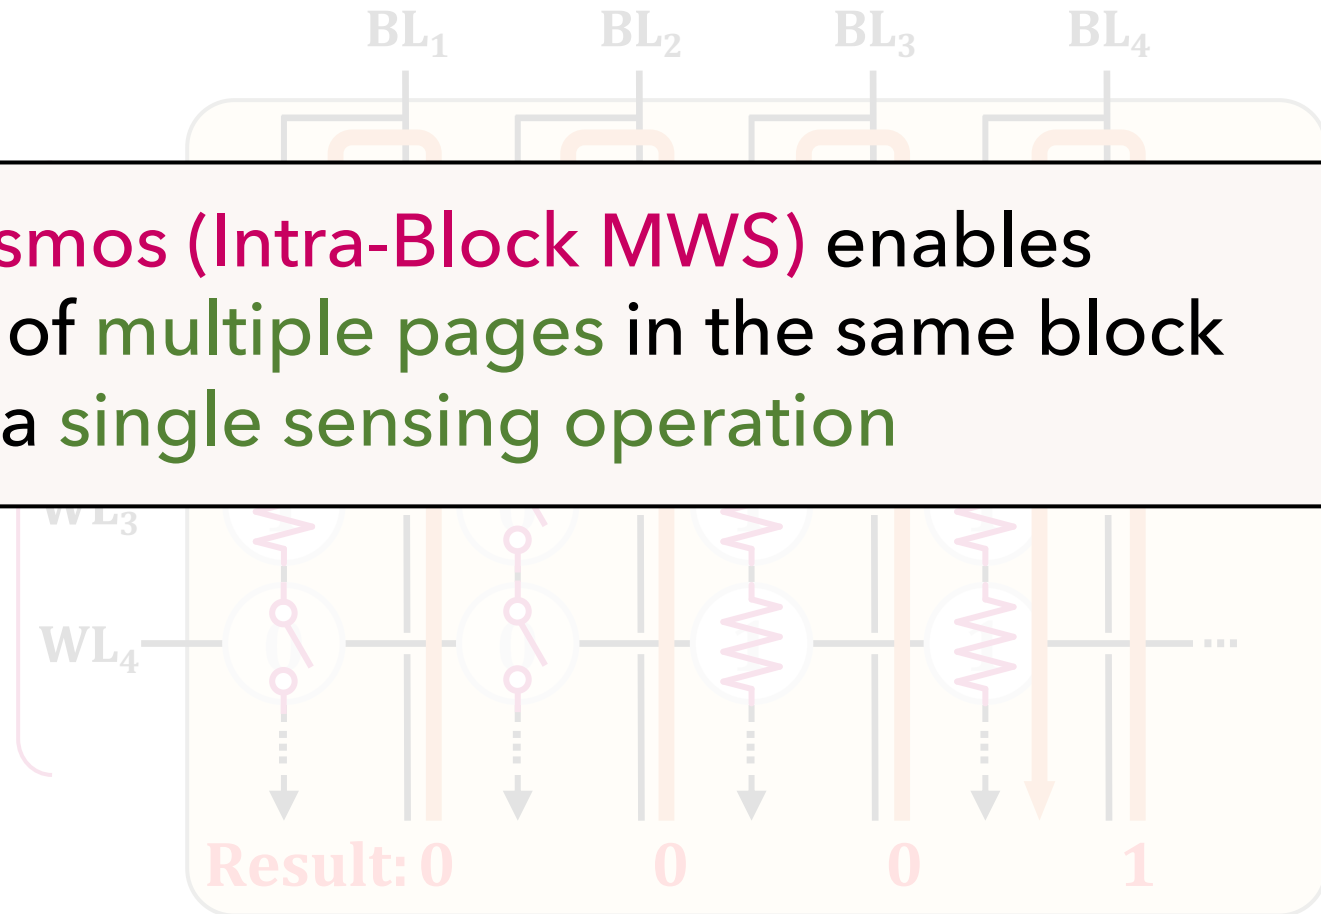
Multi-Wordline Sensing (MWS): Bitwise AND

- **Intra-Block MWS:**

Simultaneously activates multiple WLs in the same block

→ Bitwise AND of the stored data in the WLs

Flash-Cosmos (Intra-Block MWS) enables bitwise AND of multiple pages in the same block via a single sensing operation




Multi-Wordline Sensing (MWS): Bitwise OR

- **Inter-Block MWS:**

Simultaneously activates multiple WLs in different blocks

→ Bitwise OR of the stored data in the WLs




Multi-Wordline Sensing (MWS): Bitwise OR

- **Inter-Block MWS:**

Simultaneously activates multiple WLs in different blocks

→ Bitwise OR of the stored data in the WLs

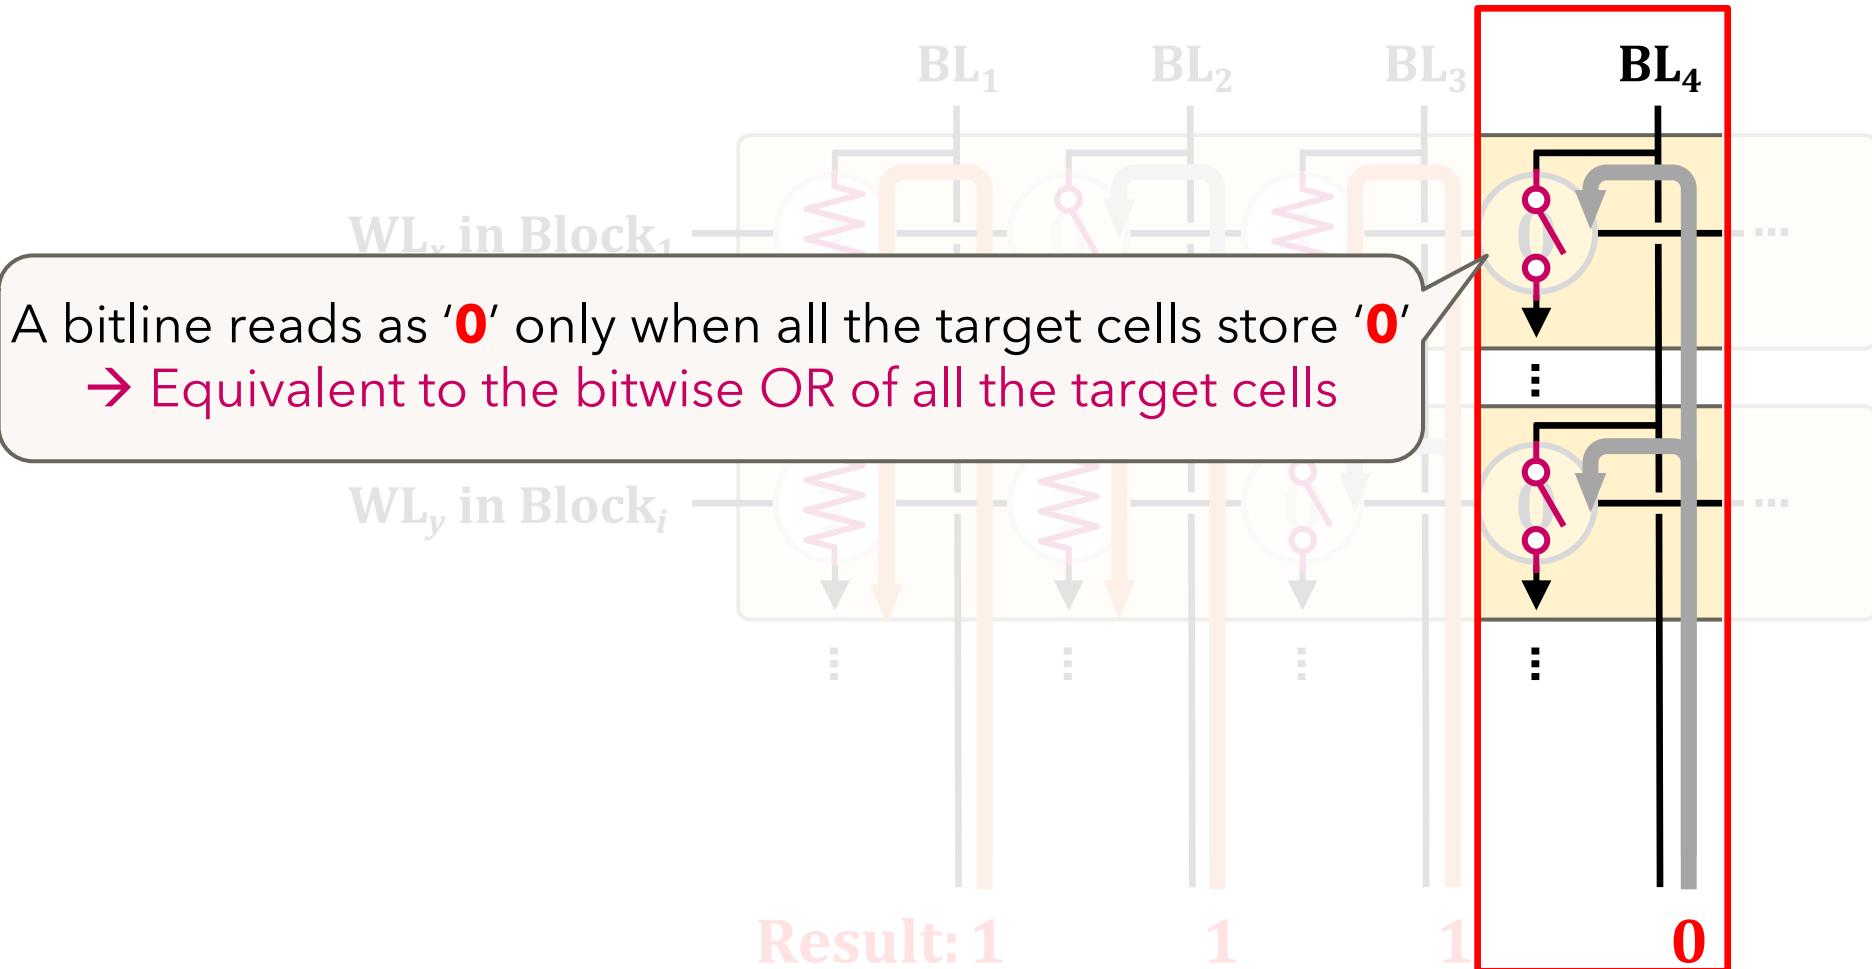


Multi-Wordline Sensing (MWS): Bitwise OR

- **Inter-Block MWS:**

Simultaneously activates multiple WLs in different blocks

→ Bitwise OR of the stored data in the WLs




Multi-Wordline Sensing (MWS): Bitwise OR

- **Inter-Block MWS:**

Simultaneously activates multiple WLs in different blocks

→ Bitwise OR of the stored data in the WLs




Multi-Wordline Sensing (MWS): Bitwise OR

- **Inter-Block MWS:**

Simultaneously activates multiple WLs in different blocks

→ Bitwise OR of the stored data in the WLs

A bitline reads as '**1**' when a single target cell stores '**1**'
→ Equivalent to the bitwise OR of all the target cells




Multi-Wordline Sensing (MWS): Bitwise OR


- Inter-Block MWS:

Simultaneously activates multiple WLs in different blocks

□ Bitwise OR of the stored data in the WLs



Flash-Cosmos (Inter-Block MWS) enables
bitwise OR of multiple pages in different blocks
via a single sensing operation



Other Types of Bitwise Operations

Flash-Cosmos also enables
other types of bitwise operations
(NOT/NAND/NOR/XOR/XNOR)
leveraging **existing features** of NAND flash memory

Flash-Cosmos: In-Flash Bulk Bitwise Operations Using Inherent Computation Capability of NAND Flash Memory

Jisung Park^{§∇} Roknoddin Azizi[§] Geraldo F. Oliveira[§] Mohammad Sadrosadati[§]
Rakesh Nadig[§] David Novo[†] Juan Gómez-Luna[§] Myungsuk Kim[‡] Onur Mutlu[§]

[§]*ETH Zürich* [∇]*POSTECH* [†]*LIRMM, Univ. Montpellier, CNRS* [‡]*Kyungpook National University*



<https://arxiv.org/abs/2209.05566.pdf>

Key Ideas



Multi-Wordline Sensing (MWS)
to enable in-flash bulk bitwise operations
via a single sensing operation



Enhanced SLC-Mode Programming (ESP)
to eliminate raw bit errors in stored data
(and thus in computation results)

Enhanced SLC-Mode Programming (ESP)

- **Goal:** eliminate raw bit errors in stored data (and computation results)
- **Key ideas**
 - Programs only a single bit per cell (SLC-mode programming)
 - Trades storage density for reliable computation
 - Performs more precise programming of the cells
 - Trades programming latency for reliable computation

Maximizes the reliability margin
between the different states of flash cells

Enhanced SLC-Mode Programming (ESP)

- To eliminate raw bit errors in stored data (and computation results)

Flash-Cosmos (ESP) enables
reliable in-flash computation
by trading storage density & programming latency

Storage & latency overheads affect
only data used in in-flash computation

Talk Outline

- Problem, Goals & Key Idea
- Background
- Flash-Cosmos: Computation with One-Shot Multi-Operand Sensing
- Evaluation
- Summary

Evaluation Methodology

▪ Real-device characterization

- To validate the feasibility and reliability of Flash-Cosmos
- Using 160 48-WL-layer 3D Triple-Level Cell NAND flash chips
 - 3,686,400 tested wordlines
- Under worst-case operating conditions
 - Under a 1-year retention time at 10K P/E cycles
 - Worst-case data patterns

▪ System-level evaluation

- Using the state-of-the-art SSD simulator (MQSim [Tavakkol+, FAST'18])
- Three real-world applications
 - Bitmap Indices (BMI): Bitwise AND of up to ~1,000 operands
 - Image Segmentation (IMS): Bitwise AND of 3 operands
 - K-clique Star Listing (KCS): Bitwise OR of up to 32 operands
- Baselines
 - Outside-Storage Processing (OSP): A multi-core CPU (Intel i7-11700K)
 - In-Storage Processing (ISP): An in-storage hardware accelerator
 - ParaBit [Gao+, MICRO'21]: State-of-the-art in-flash processing mechanism


Results: Real-Device Characterization

Both intra- and inter-block MWS operations
require no changes to the cell array
of commodity NAND flash chips

Both MWS operations can activate multiple WLs
(intra: up to 48, inter: up to 4) at the same time
with small increase in sensing latency ($< 10\%$)

ESP significantly improves
the reliability of computation results
(no observed bit error in the tested flash cells)

Results: Performance & Energy



Flash-Cosmos provides **significant performance & energy benefits** over all the baselines

The larger the number of operands,
the higher the performance & energy benefits

More in the Paper

- Support for other types of bitwise operations
 - NOT/NAND/NOR/XOR/XNOR
- More detailed real-device characterization results
- Optimizations to improve bitwise operation performance
- Flash-Cosmos command interface
- System support
- Overhead analysis

Flash-Cosmos: In-Flash Bulk Bitwise Operations Using Inherent Computation Capability of NAND Flash Memory

Jisung Park^{§∇} Roknoddin Azizi[§] Geraldo F. Oliveira[§] Mohammad Sadrosadati[§]
Rakesh Nadig[§] David Novo[†] Juan Gómez-Luna[§] Myungsuk Kim[‡] Onur Mutlu[§]

[§]*ETH Zürich* [∇]*POSTECH* [†]*LIRMM, Univ. Montpellier, CNRS* [‡]*Kyungpook National University*



<https://arxiv.org/abs/2209.05566.pdf>

Summary: Flash-Cosmos



The first work that enables
in-flash multi-operand bulk bitwise operations
with a single sensing operation and high reliability



Improves performance
by 32x/25x/3.5x over OSP/ISP/ParaBit



Improves energy efficiency
by 95x/13.4x/3.3x over OSP/ISP/ParaBit



Low-cost & requires no changes to flash cell arrays

Flash-Cosmos

In-Flash Bulk Bitwise Operations Using Inherent Computation Capability of NAND Flash Memory

Jisung Park, Roknoddin Azizi, Geraldo F. Oliveira, Mohammad Sadrosadati,
Rakesh Nadig, David Novo, Juan Gómez Luna, Myungsuk Kim, and Onur Mutlu

MICRO 2022

