

# Computer Architecture

## Lecture 26: On-Chip Networks

Prof. Onur Mutlu

ETH Zürich

Fall 2023

26 January 2024

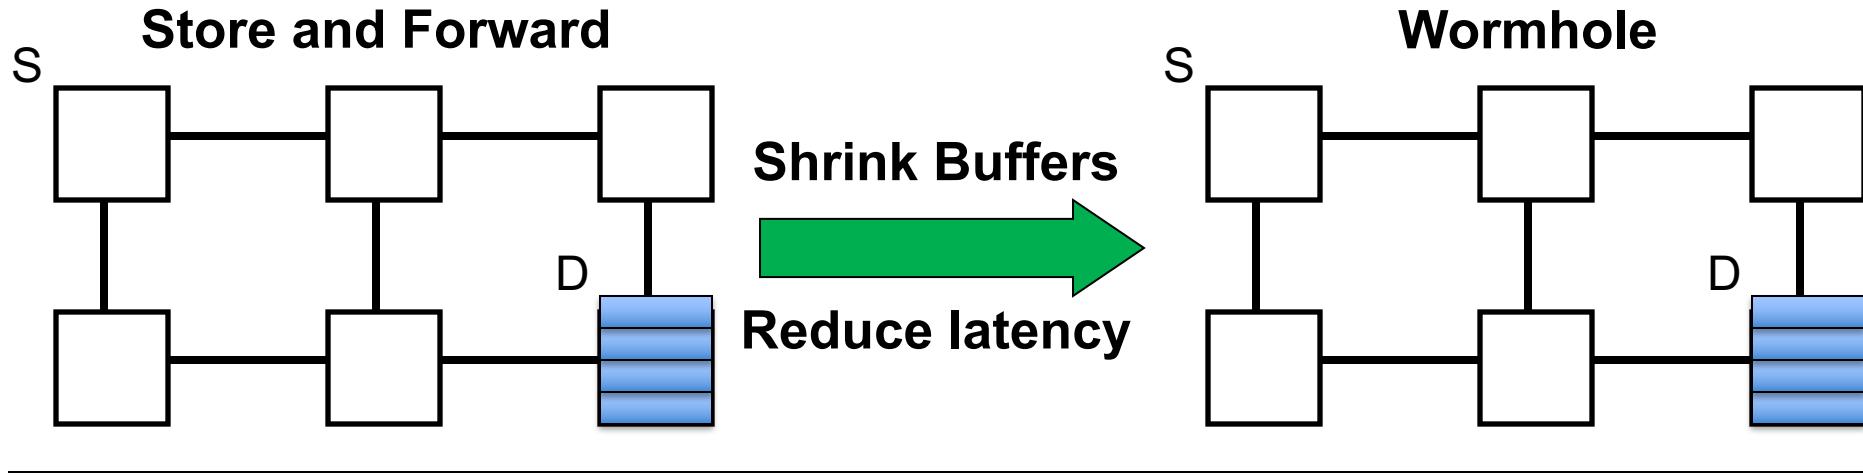
# Recall: Interconnection Network Basics

---

- **Topology**
  - Specifies the way switches are wired
  - Affects routing, reliability, throughput, latency, building ease
- **Routing (algorithm)**
  - How does a message get from source to destination
  - Static or adaptive
- **Buffering and Flow Control**
  - What do we store within the routers & links?
    - Entire packets, parts of packets, etc?
  - How do we throttle during oversubscription?
  - Tightly coupled with routing strategy

# Buffered Flow Control

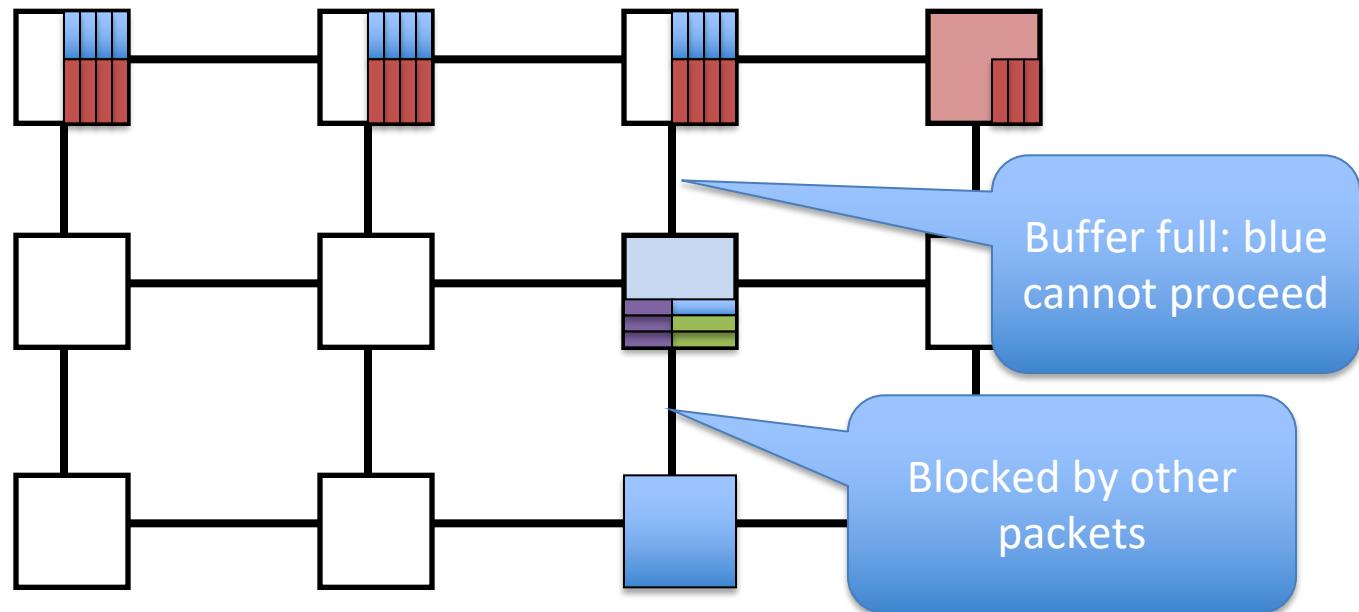
# Review: Buffered Flow Control



**Any other issues?**

**Head-of-Line Blocking**

**Use Virtual Channels**



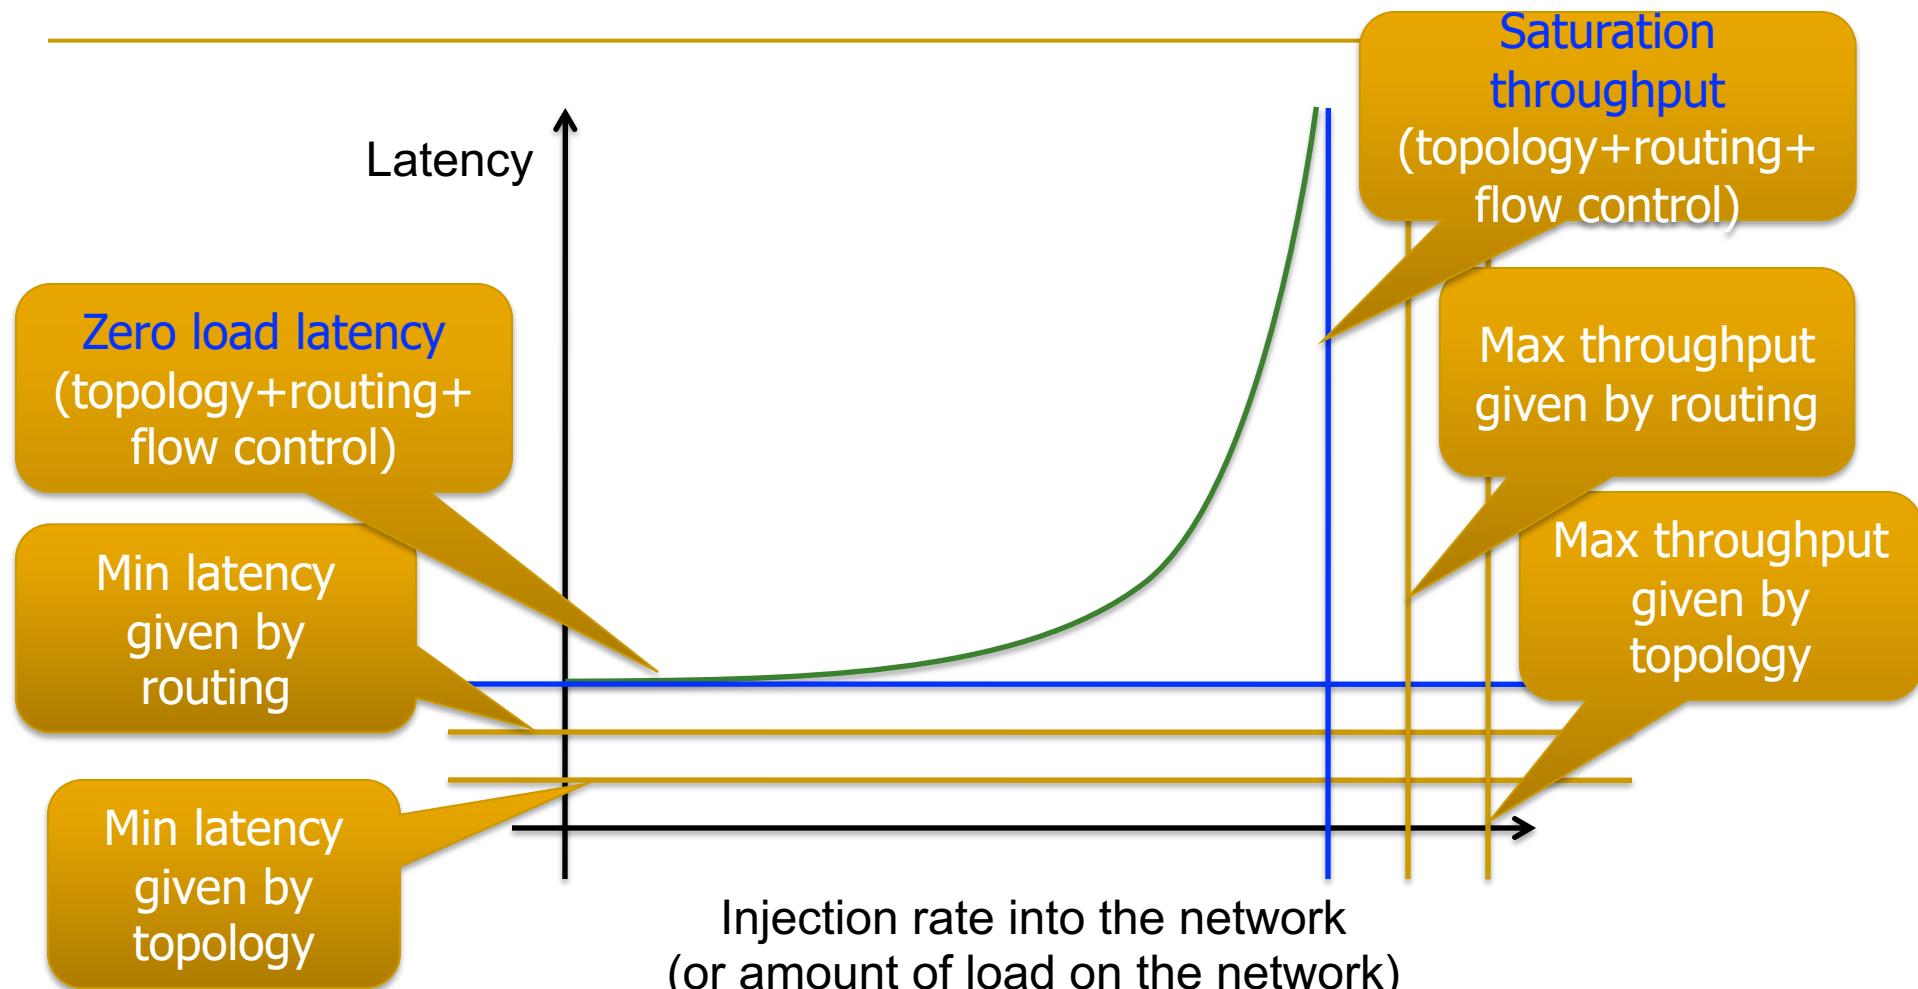
# Recall: Communicating Buffer Availability

---

- Credit-based flow control
  - Upstream knows how many buffers are downstream
  - Downstream passes back credits to upstream
  - Significant upstream signaling (esp. for small flits)
- On/Off (XON/XOFF) flow control
  - Downstream has on/off signal to upstream
- ACK/NACK flow control
  - Upstream optimistically sends downstream
  - Buffer cannot be deallocated until ACK/NACK received
  - Inefficiently utilizes buffer space

# Interconnection Network Performance

# Interconnection Network Performance



**Saturation throughput:** Injection rate at which latency asymptotes  
**“Zero load” latency:** Latency with no contention

# Ideal Latency

---

- Ideal latency
  - Solely due to wire delay between source and destination

$$T_{ideal} = \frac{D}{v} + \frac{L}{b}$$

- D = Manhattan distance
  - The distance between two points measured along axes at right angles.
- v = propagation velocity
- L = packet size
- b = channel bandwidth

# Actual Latency

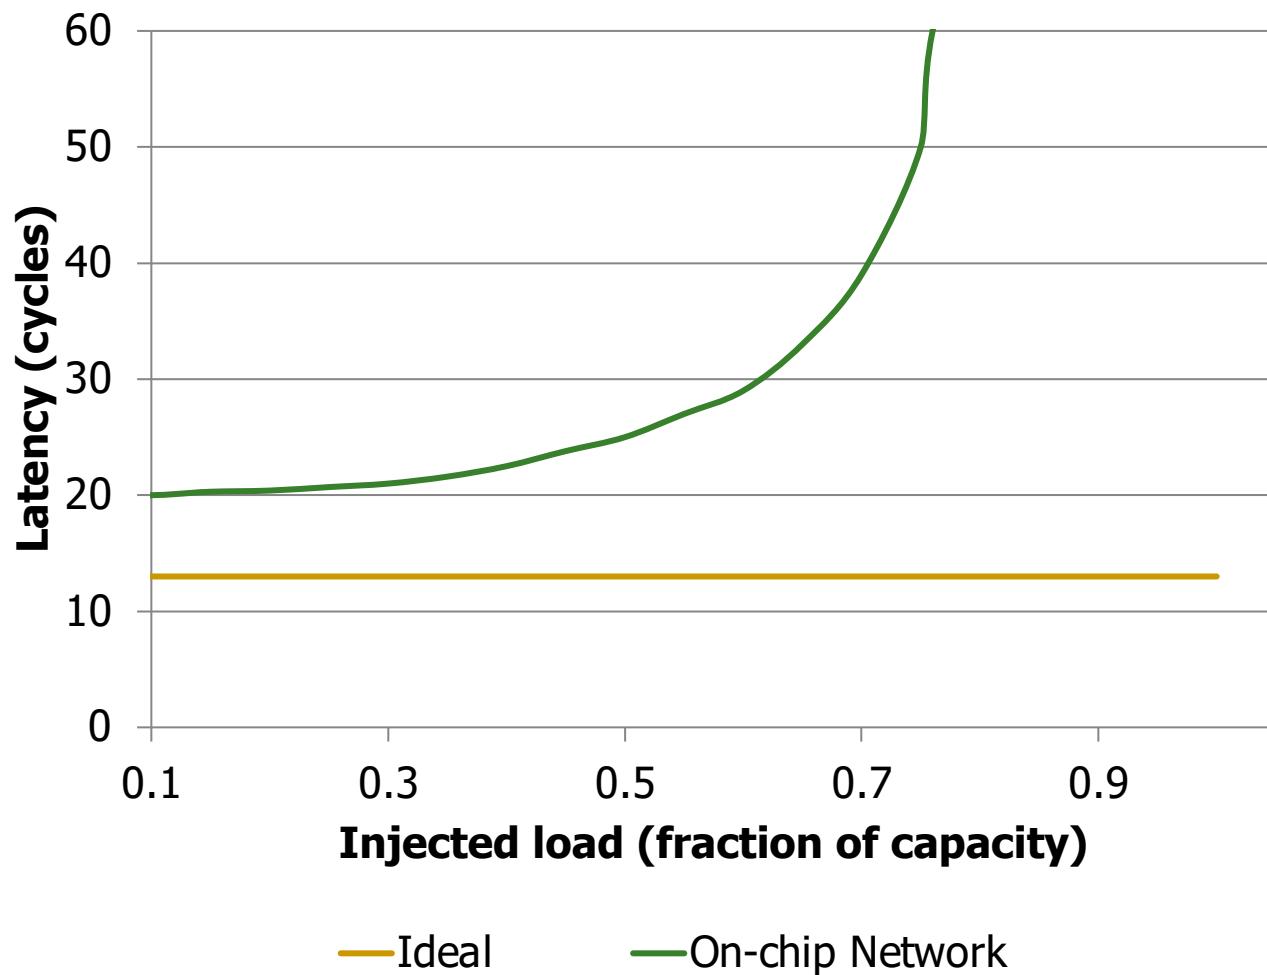
---

- Dedicated wiring impractical
  - Long wires segmented with insertion of routers

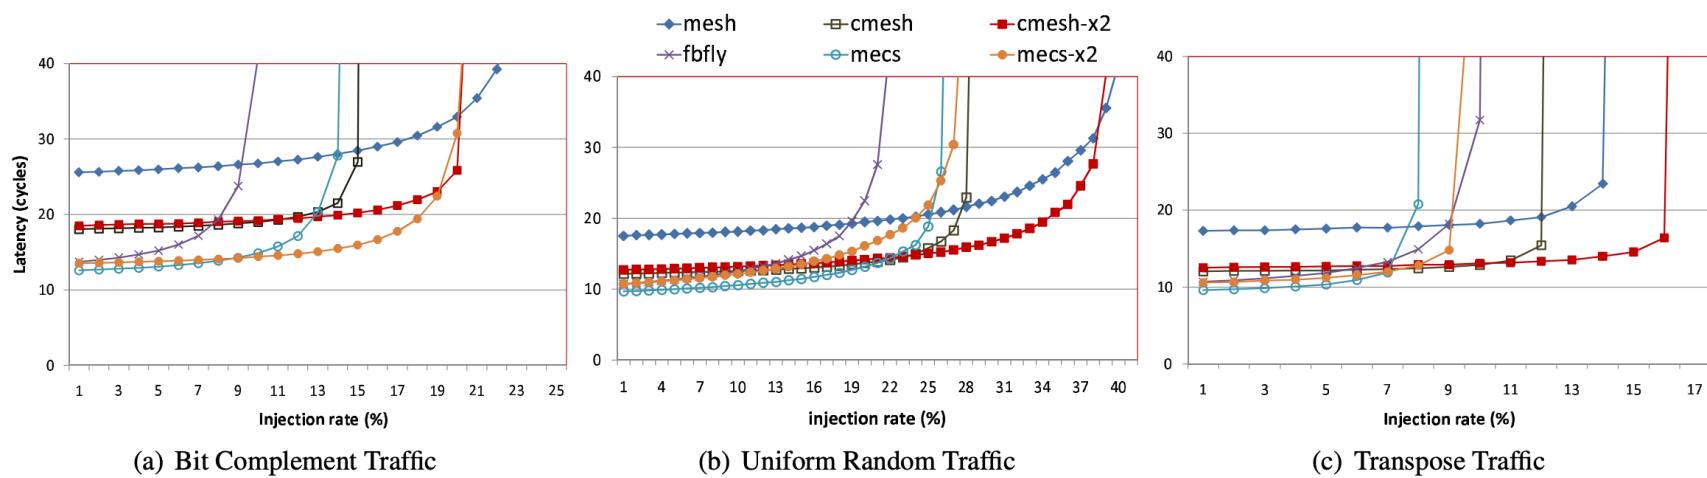
$$T_{actual} = \frac{D}{v} + \frac{L}{b} + H \cdot T_{router} + T_c$$

- D = Manhattan distance
- v = propagation velocity
- L = packet size
- b = channel bandwidth
- H = hops
- $T_{router}$  = router latency
- $T_c$  = latency due to contention

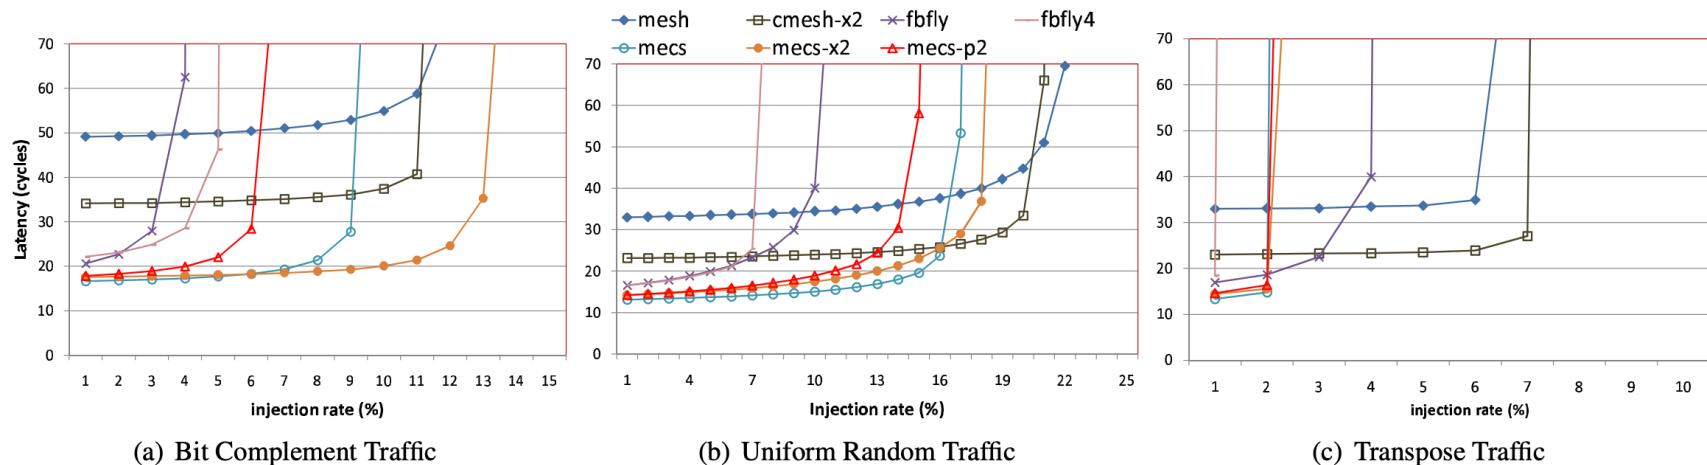
# Load-Latency Curve



# Load-Latency Curve Examples

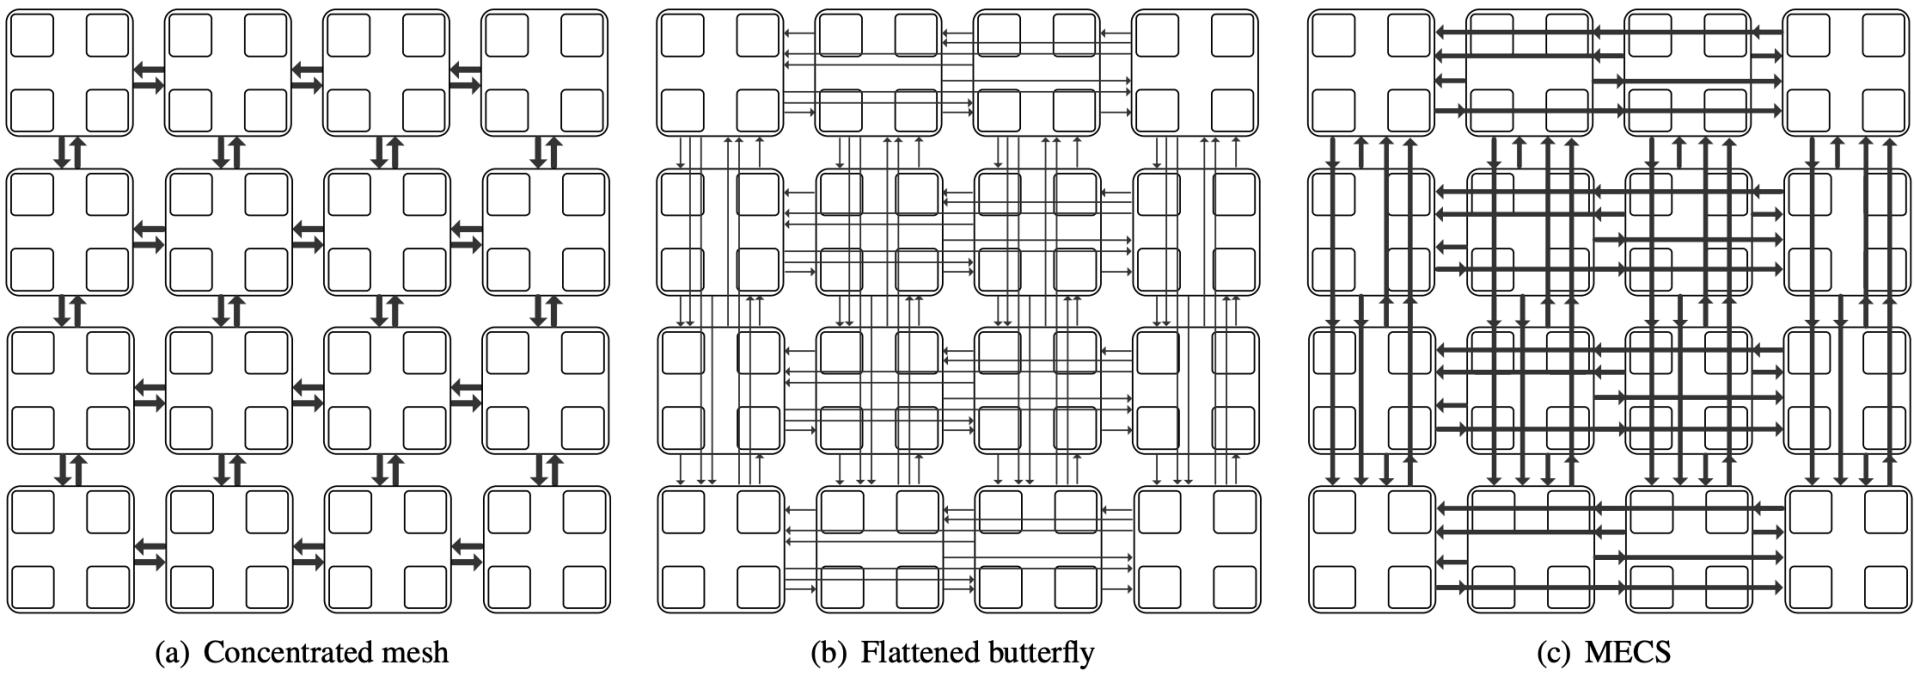


**Figure 4. Load-latency graphs for 64-node mesh, CMesh, flattened butterfly and MECS topologies.**



**Figure 5. Load-latency graphs for 256-node mesh, CMesh, flattened butterfly and MECS topologies.**

# Examined Topologies in Prior Slide



**Figure 1. Concentrated Mesh, Flattened Butterfly and MECS topologies for a 64-terminal network.**

Different topologies work differently for different communication patterns

# Multi-Drop Express Channels (MECS)

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"Express Cube Topologies for On-Chip Interconnects"**  
*Proceedings of the 15th International Symposium on High-Performance Computer Architecture (HPCA)*, pages 163-174, Raleigh, NC, February 2009. [Slides \(ppt\)](#)

## Express Cube Topologies for On-Chip Interconnects

Boris Grot

Joel Hestness

Stephen W. Keckler

Onur Mutlu<sup>†</sup>

Department of Computer Sciences

The University of Texas at Austin

{bgrot, hestness, skeckler}@cs.utexas.edu

<sup>†</sup>Computer Architecture Laboratory (CALCM)

Carnegie Mellon University

onur@cmu.edu

# Kilo-NoC Building on MECS

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"Kilo-NOC: A Heterogeneous Network-on-Chip Architecture for Scalability and Service Guarantees"**

*Proceedings of the 38th International Symposium on Computer Architecture (ISCA), San Jose, CA, June 2011.* Slides (pptx)

***One of the 12 computer architecture papers of 2011 selected as Top Picks by IEEE Micro.***

## Kilo-NOC: A Heterogeneous Network-on-Chip Architecture for Scalability and Service Guarantees

Boris Grot<sup>1</sup>

bgrot@cs.utexas.edu

Joel Hestness<sup>1</sup>

hestness@cs.utexas.edu

Stephen W. Keckler<sup>1,2</sup>

skeckler@nvidia.com

Onur Mutlu<sup>3</sup>

onur@cmu.edu

<sup>1</sup>The University of Texas at Austin  
Austin, TX

<sup>2</sup>NVIDIA  
Santa Clara, CA

<sup>3</sup>Carnegie Mellon University  
Pittsburgh, PA

# Kilo-NoC Building on MECS

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"A QoS-Enabled On-Die Interconnect Fabric for Kilo-Node Chips"**  
*IEEE Micro, Special Issue: Micro's Top Picks from 2011 Computer Architecture Conferences (MICRO TOP PICKS)*, Vol. 32, No. 3, May/June 2012.
- 

## A QoS-ENABLED ON-DIE INTERCONNECT FABRIC FOR KILO-NODE CHIPS

---

TO MEET RAPIDLY GROWING PERFORMANCE DEMANDS AND ENERGY CONSTRAINTS, FUTURE CHIPS WILL LIKELY FEATURE THOUSANDS OF ON-DIE RESOURCES. EXISTING NETWORK-ON-CHIP SOLUTIONS WEREN'T DESIGNED FOR SCALABILITY AND WILL BE UNABLE TO MEET FUTURE INTERCONNECT DEMANDS. A HYBRID NETWORK-ON-CHIP ARCHITECTURE CALLED KILO-NOC CO-OPTIMIZES TOPOLOGY, FLOW CONTROL, AND QUALITY OF SERVICE TO ACHIEVE SIGNIFICANT GAINS IN EFFICIENCY.

---

# Network Performance Metrics

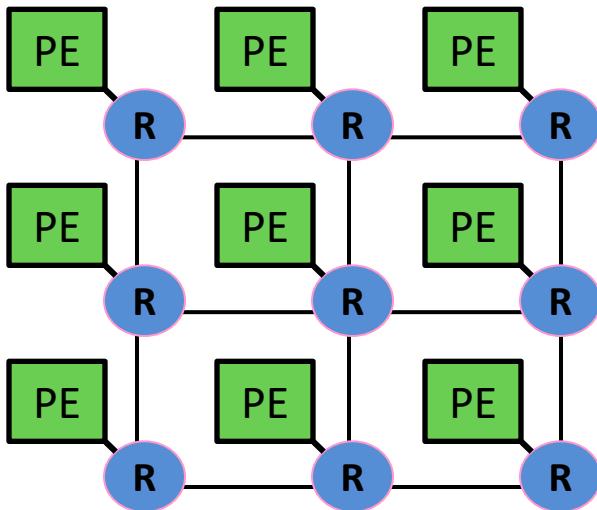
---

- Packet latency (avg/max)
- Round trip latency (avg/max)
- Saturation throughput
- Application-level performance: execution time
- System performance: job throughput
  - Affected by interference among threads/applications

# Buffering and Flow Control in On-Chip Networks

# On-Chip Networks

- Connect **cores, caches, memory controllers, etc**
  - Buses and crossbars are not scalable
- **Usually packet switched**
- **2D mesh:** Commonly used topology
- **XY Routing with FIFO or Round robin port arbitration** common
- **Virtual channel buffering** common

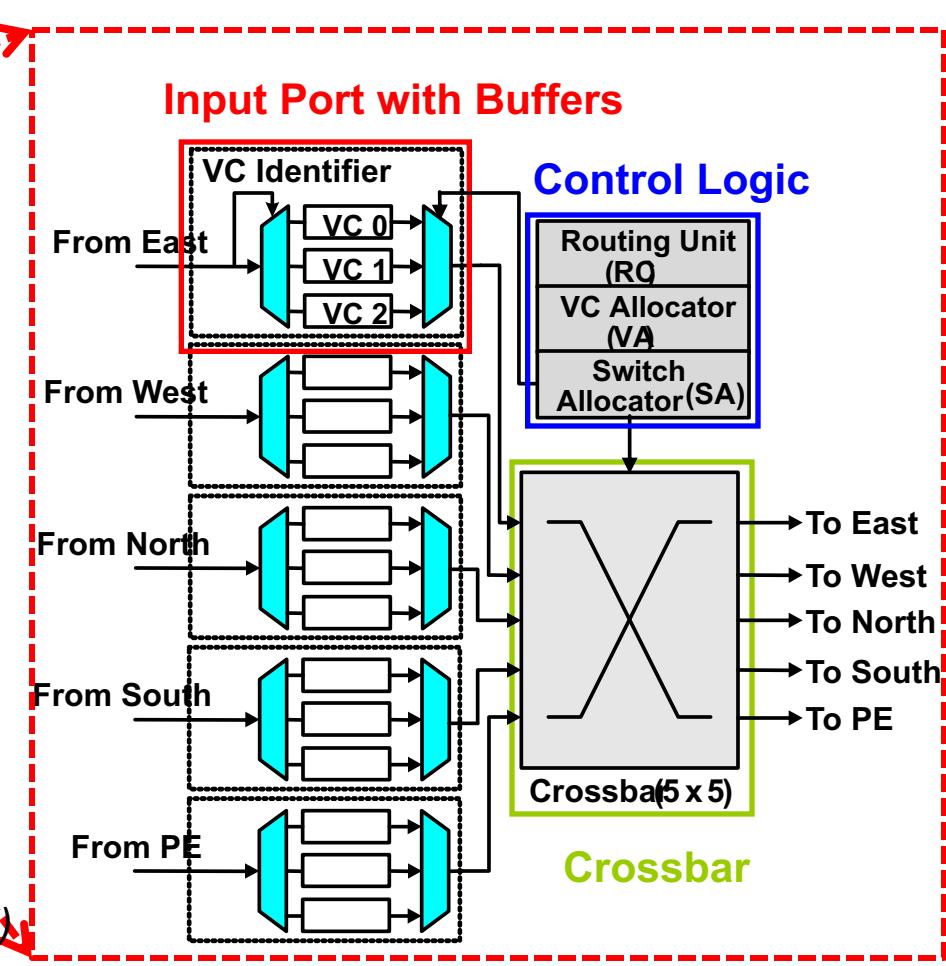
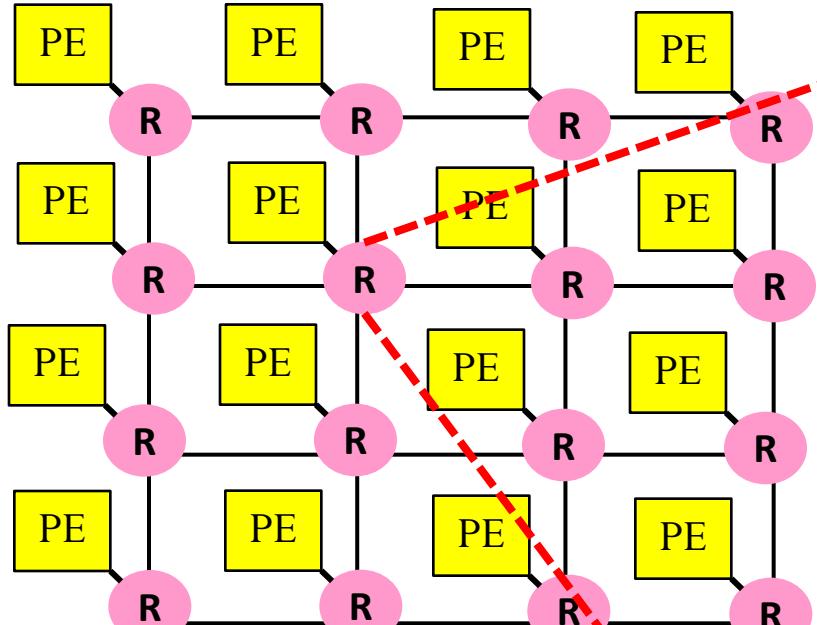


 Router

 Processing Element  
(Cores, L2 Banks,  
Memory Controllers, etc)

- Primarily serve **cache misses and memory requests**

# On-Chip Networks



R Router

PE Processing Element  
(Cores, L2 Banks, Memory Controllers etc)

# On-Chip vs. Off-Chip Interconnects

---

- On-chip advantages
  - Low latency between cores
  - No pin constraints
  - Rich & low-power wiring resources
    - Very high bandwidth
    - Simpler (global) coordination
- On-chip constraints/disadvantages
  - 2D substrate limits easy-to-implement topologies
  - Energy/power consumption a key concern
    - Complex algorithms undesirable
    - Large buffers undesirable
  - Logic area & metal layers constrain use of wiring resources

# On-Chip vs. Off-Chip Interconnects (II)

---

- **Cost**
  - Off-chip: Channels, pins, connectors, cables
  - On-chip: Cost is storage and switches (wires are plentiful)
    - Leads to networks with many wide channels, less buffering
- **Channel characteristics**
  - On chip short distance → low latency
  - On chip RC lines → need repeaters every 1-2mm
    - Can put logic in repeaters
- **Workloads**
  - Off-chip: Large-scale parallel application multi-chip traffic
  - On-chip: Multi-core cache/memory traffic

# On-Chip vs. Off-Chip Tradeoffs

---

- Dally & Towles, “Route Packets, Not Wires: On-Chip Interconnection Networks,” DAC 2001.

## 3 Challenges in architecture and design

While the same principles apply to interconnection networks at all scales, on-chip networks have a number of characteristics that make their design quite different than the inter-chip (and inter-board) networks that have been designed for years. In particular, wires and pins are more abundant than in inter-chip networks and buffers space is less abundant. These differences enable a number of new network topologies, flow control methods, and other techniques. In particular, we identify three areas that are ripe for future research:

### 3.1 What topologies are best matched to the abundant wiring resources available on chip?

On chip networks have enormous wiring resources at their disposal. In the example network described above, there can be up to 6,000 wires on each metal layer crossing each edge of a tile. It is quite easy to achieve over 24,000 ‘pins’ crossing the four edges of a tile. In contrast, inter-chip networks have historically been pin limited, required to limit the connections of one router chip to far less than 1,000 total pins. This large, 24:1, difference between router pin limitations allows the designer to trade wiring resources for network performance, making a qualitative difference in network architecture.

### 3.2 What flow control methods reduce buffer count and hence router overhead?

Buffer space in an on-chip router directly impacts the area overhead of the network and thus must be kept to a minimum. In contrast, most inter-chip network routers are pin limited and thus have ample room for very large buffers.

### 3.3 What circuits best exploit the structured wiring of on-chip networks?

Much of the advantage of on-chip networks derives from the regular, structured nature of their wiring. As described below, the well controlled electrical parameters of this wiring enable the use of high-performance circuits such as pulsed low-swing drivers and receivers to reduce power dissipation, reduce latency, and increase repeater spacing. While these transceivers yield big performance

# On-Chip vs. Off-Chip Tradeoffs

---

- George Nychis, Chris Fallin, Thomas Moscibroda, Onur Mutlu, and Srinivasan Seshan,

**"On-Chip Networks from a Networking Perspective:  
Congestion and Scalability in Many-core Interconnects"**

*Proceedings of the 2012 ACM SIGCOMM*

*Conference (**SIGCOMM**)*, Helsinki, Finland, August 2012. Slides  
(pptx)

## **On-Chip Networks from a Networking Perspective: Congestion and Scalability in Many-Core Interconnects**

George Nychis<sup>†</sup>, Chris Fallin<sup>†</sup>, Thomas Moscibroda<sup>§</sup>, Onur Mutlu<sup>†</sup>, Srinivasan Seshan<sup>†</sup>

<sup>†</sup> Carnegie Mellon University  
[{gnychis,cfallin,onur,srini}@cmu.edu](mailto:{gnychis,cfallin,onur,srini}@cmu.edu)

<sup>§</sup> Microsoft Research Asia  
[moscitho@microsoft.com](mailto:moscitho@microsoft.com)

# On-Chip vs. Off-Chip Tradeoffs (II)

---

- George Nychis, Chris Fallin, Thomas Moscibroda, and Onur Mutlu,  
**"Next Generation On-Chip Networks: What Kind of Congestion Control Do We Need?"**

*Proceedings of the 9th ACM Workshop on Hot Topics in Networks (HOTNETS), Monterey, CA, October 2010.* Slides (ppt) (key)

## Next Generation On-Chip Networks: What Kind of Congestion Control Do We Need?

George Nychis†, Chris Fallin†, Thomas Moscibroda§, Onur Mutlu†

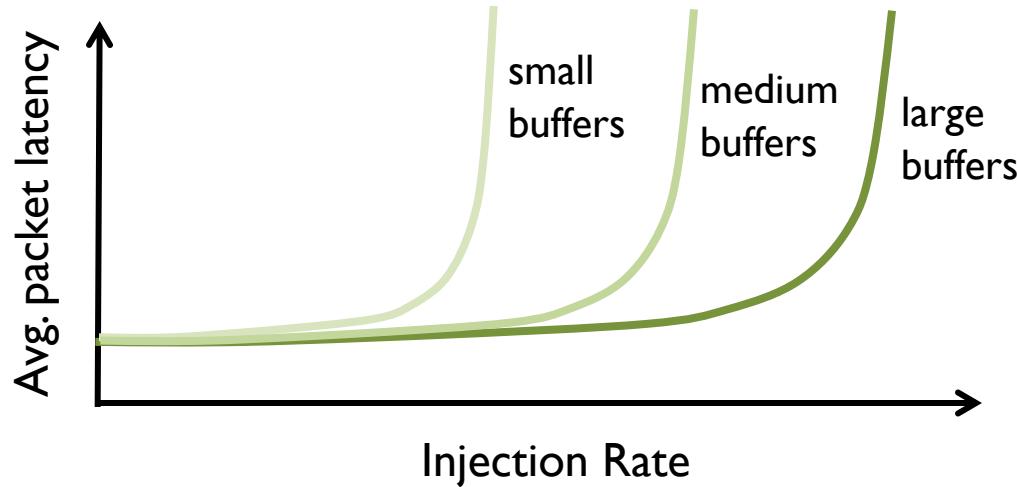
†Carnegie Mellon University

{gnychis,cfallin,onur}@cmu.edu

§Microsoft Research  
moscitho@microsoft.com

# Buffers in NoC Routers

- Buffers are necessary for high network throughput  
→ buffers increase total available bandwidth in network



# Buffers in NoC Routers

- Buffers are necessary for high network performance  
→ buffers increase total available bandwidth



- Buffers consume significant chip area
  - Dynamic energy consumption
  - Static energy consumption
- Buffers reduce system efficiency

Can we get rid of buffers...?

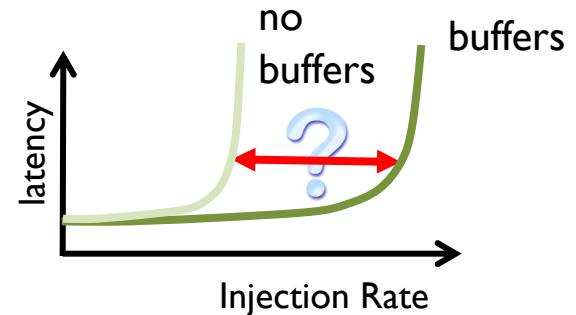


- Buffers require significant chip area
  - e.g., in TRIPS prototype chip, input buffers occupy 75% of total on-chip network area [Gratz et al, ICCD'06]



# Going Bufferless...?

- How much throughput do we lose?  
→ How is latency affected?



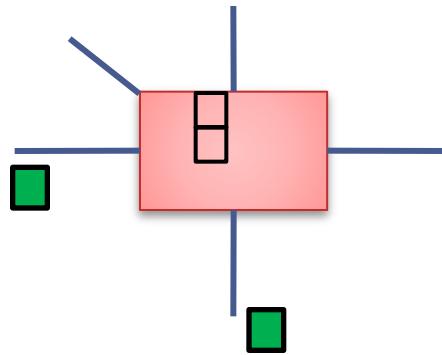
- Up to what **injection rates** can we use bufferless routing?  
→ Are there **realistic scenarios** in which an NoC operates at injection rates below the threshold?
- Can we achieve **energy reduction**?  
→ If so, how much...?
- Can we reduce **area, complexity**, etc...?



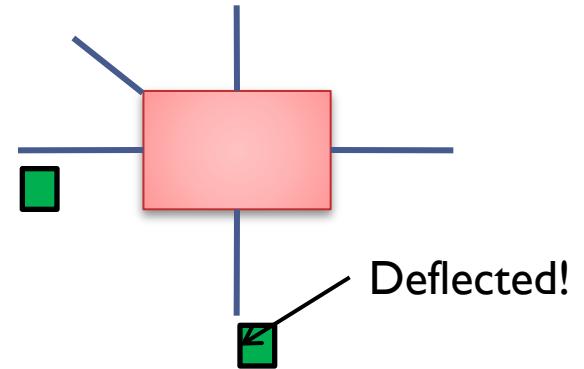
Answers in  
our paper  
(ISCA'09)!

# BLESS: Bufferless Routing

- Always forward *all* incoming flits to some output port
- If no productive direction is available, send to another direction
- → packet is deflected
- → Hot-potato routing [Baran' 62]

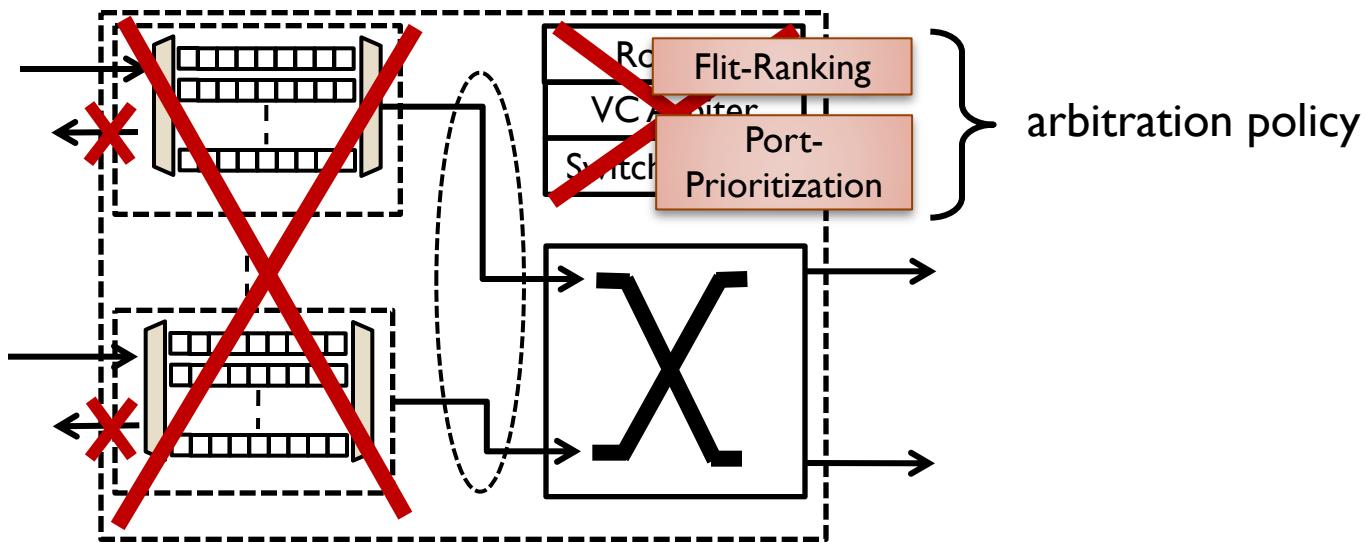


Buffered



BLESS

# BLESS: Bufferless Routing



Flit-Ranking

Port-  
Prioritization

1. Create a ranking over all incoming flits
2. For a given flit in this ranking, find the best free output-port  
Apply to each flit in order of ranking

# FLIT-BLESS: Flit-Level Routing

- Each flit is routed independently.
- **Oldest-first arbitration** (other policies evaluated in paper)

Flit-Ranking

I. Oldest-first ranking

Port-Prioritization

II. Assign flit to productive port, if possible.  
Otherwise, assign to non-productive port.

- **Network Topology:**  
→ Can be applied to most topologies (Mesh, Torus, Hypercube, Trees, ...)
  - 1) #output ports , #input ports at every router
  - 2) every router is reachable from every other router
- **Flow Control & Injection Policy:**  
→ Completely **local**, inject whenever input port is free
- **Absence of Deadlocks:** every flit is always moving
- **Absence of Livelocks:** with oldest-first ranking

# BLESS: Advantages & Disadvantages

## Advantages

- No buffers
- Purely local flow control
- Simplicity
  - no credit-flows
  - no virtual channels
  - simplified router design
- No deadlocks, livelocks
- Adaptivity
  - packets are deflected around congested areas!
- Router latency reduction
- Area savings



## Disadvantages

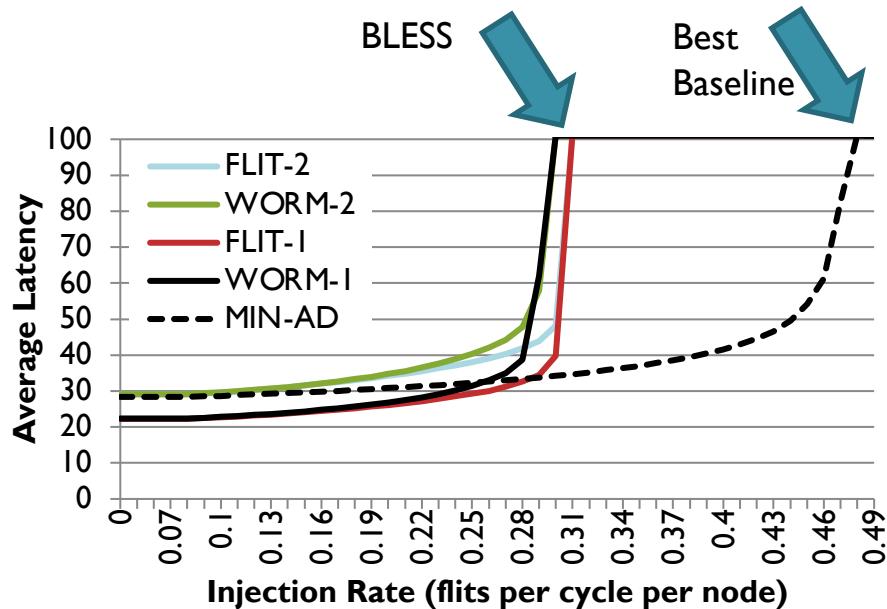
- Increased latency
- Reduced bandwidth
- Increased buffering at receiver
- Header information at each flit
- Oldest-first arbitration complex
- QoS becomes difficult

Impact on energy...?



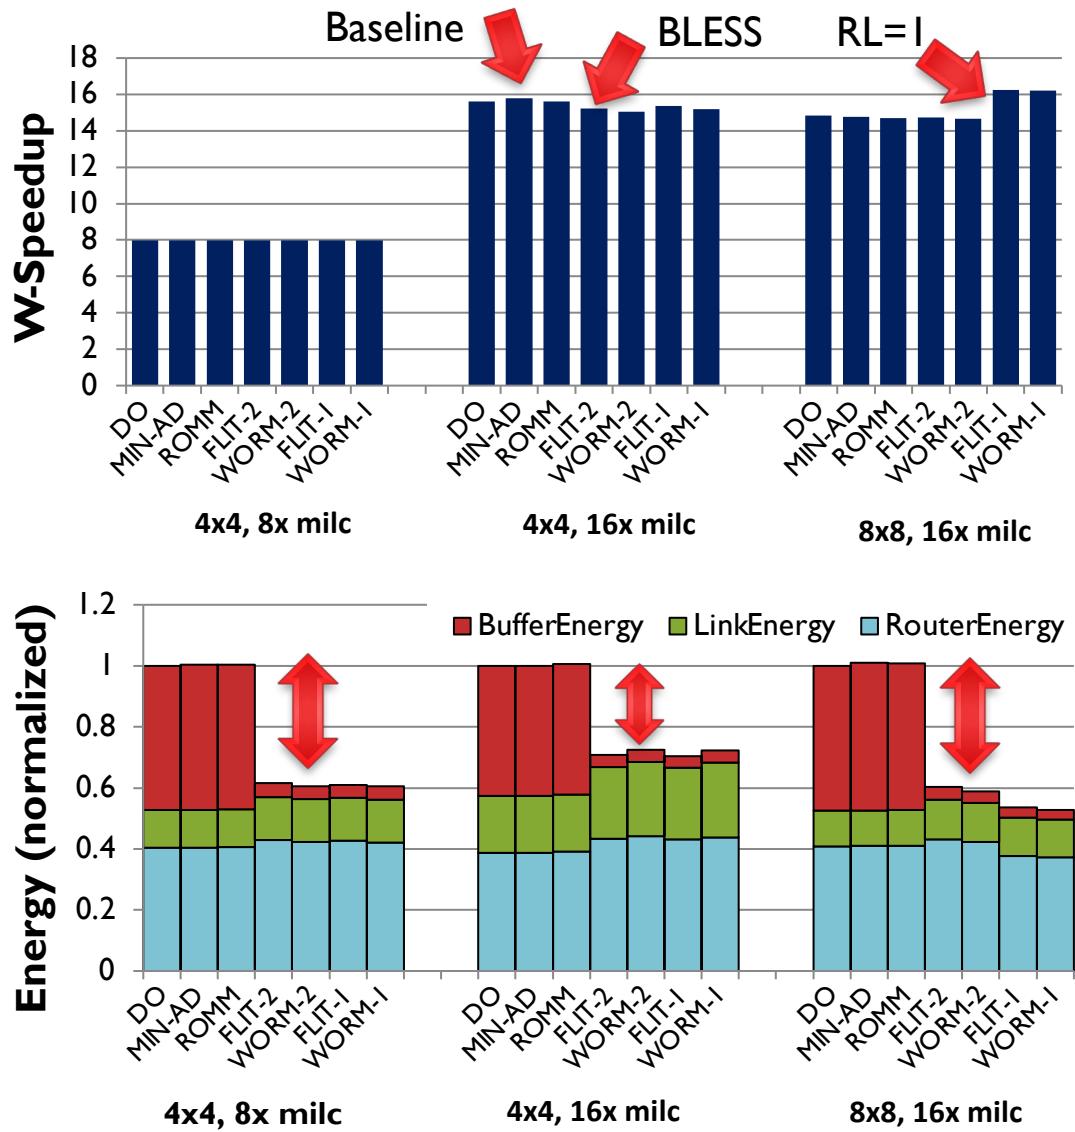
# Evaluation – Synthetic Traces

- First, the bad news ☺
- Uniform random injection
- BLESS has **significantly lower saturation throughput** compared to buffered baseline.



# Evaluation – Homogenous Case Study

- milc benchmarks  
(moderately intensive)
  - Perfect caches!
  - Very little performance degradation with BLESS  
(less than 4% in dense network)
  - With router latency  $l$ ,  
BLESS can even  
outperform baseline  
(by  $\sim 10\%$ )
  - Significant energy improvements  
(almost 40%)



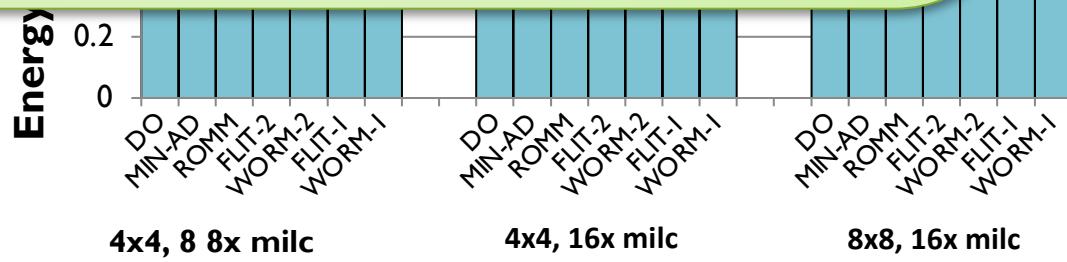
# Evaluation – Homogenous Case Study

- milc benchmarks

## Observations:

- 1) Injection rates not extremely high on average  
→ self-throttling!
- 2) For bursts and temporary hotspots, use network links as buffers!

- Significant energy improvements (almost 40%)



# BLESS Conclusions

---

- For a very wide range of applications and network settings, buffers are not needed in NoC
  - Significant energy savings (32% even in dense networks and perfect caches)
  - Area-savings of 60%
  - Simplified router and network design (flow control, etc...)
  - Performance slowdown is minimal (can even increase!)

➤ A strong case for a **rethinking of NoC design!**

- Future research:
  - Support for quality of service, different traffic classes, energy-management, etc...

# Bufferless Deflection Routing in NoCs

---

- Thomas Moscibroda and Onur Mutlu,  
**"A Case for Bufferless Routing in On-Chip Networks"**  
*Proceedings of the 36th International Symposium on Computer Architecture (ISCA)*, pages 196-207, Austin, TX, June 2009. Slides (pptx)

## A Case for Bufferless Routing in On-Chip Networks

Thomas Moscibroda  
Microsoft Research  
[moscitho@microsoft.com](mailto:moscitho@microsoft.com)

Onur Mutlu  
Carnegie Mellon University  
[onur@cmu.edu](mailto:onur@cmu.edu)

# Issues In Bufferless Deflection Routing

---

- Livelock
- Resulting Router Complexity
- Performance & Congestion at High Loads
- Quality of Service and Fairness
- Chris Fallin, Greg Nazario, Xiangyao Yu, Kevin Chang, Rachata Ausavarungnirun, and Onur Mutlu,  
**"Bufferless and Minimally-Buffered Deflection Routing"**  
*Invited Book Chapter in Routing Algorithms in Networks-on-Chip,  
pp. 241-275, Springer, 2014.*

# Low-Complexity Bufferless Routing

---

- Chris Fallin, Chris Craik, and Onur Mutlu,  
**"CHIPPER: A Low-Complexity Bufferless Deflection Router"**

*Proceedings of the 17th International Symposium on High-Performance Computer Architecture (HPCA), pages 144-155, San Antonio, TX, February 2011.* Slides (pptx)

## CHIPPER: A Low-complexity Bufferless Deflection Router

Chris Fallin cfallin@cmu.edu	Chris Craik craik@cmu.edu	Onur Mutlu onur@cmu.edu
---------------------------------	------------------------------	----------------------------

Computer Architecture Lab (CALCM)  
Carnegie Mellon University

# CHIPPER: A Low-complexity Bufferless Deflection Router

Chris Fallin, Chris Craik, and Onur Mutlu,

**"CHIPPER: A Low-Complexity Bufferless Deflection Router"**

*Proceedings of the 17th International Symposium on High-Performance Computer Architecture (HPCA)*, pages 144-155, San Antonio, TX, February 2011. Slides (pptx)

**SAFARI** Carnegie Mellon

# Motivation

---

- Recent work has proposed **bufferless deflection routing** (BLESS [Moscibroda, ISCA 2009])
  - Energy savings: ~40% in total NoC energy
  - Area reduction: ~40% in total NoC area
  - Minimal performance loss: ~4% on average
  - Unfortunately: unaddressed complexities in router
    - long critical path, large reassembly buffers
- **Goal:** obtain these benefits while simplifying the router in order to **make bufferless NoCs practical.**

# Problems that Bufferless Routers Must Solve

---

## 1. Must provide **livelock freedom**

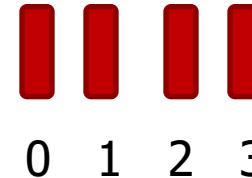
→ A packet should not be deflected forever

## 2. Must **reassemble packets** upon arrival

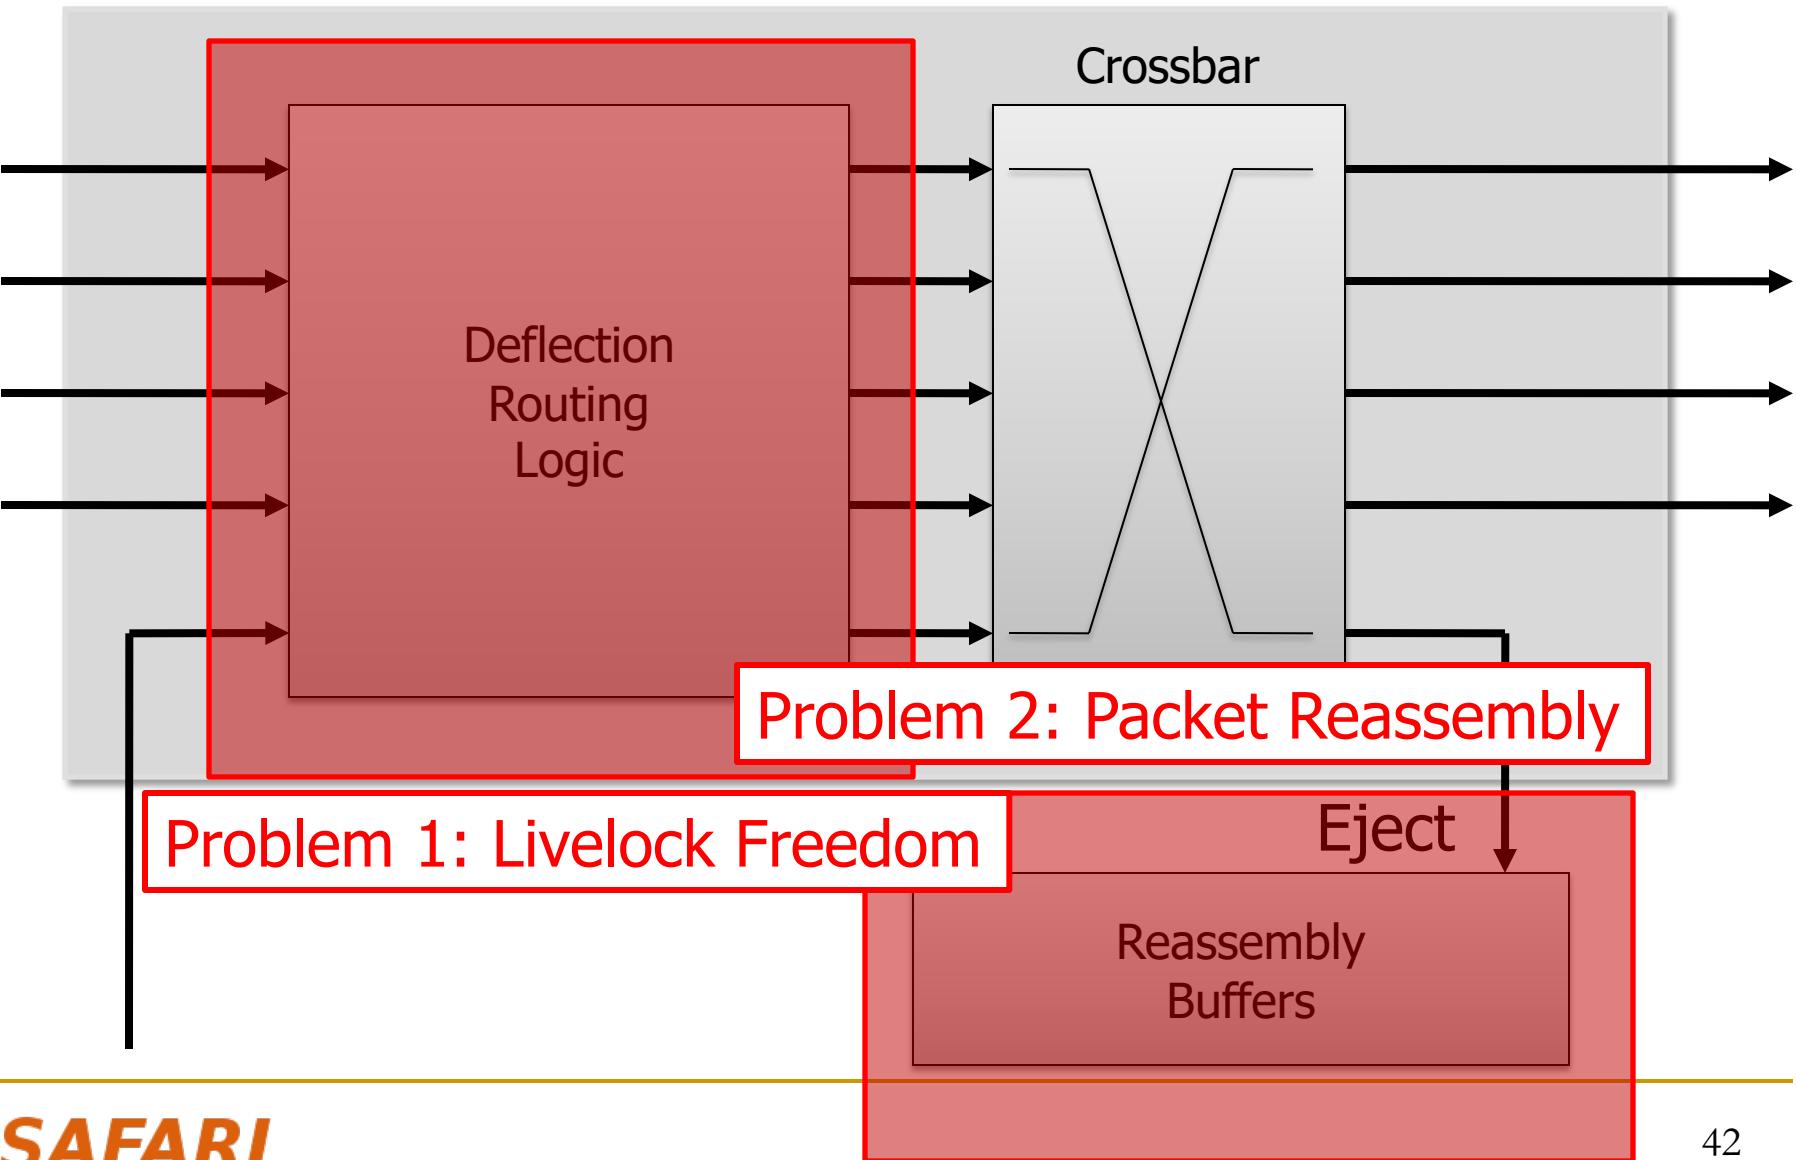
**Flit:** atomic routing unit



**Packet:** one or multiple flits



# A Bufferless Router: A High-Level View



# Complexity in Bufferless Deflection Routers

---

## 1. Must provide livelock freedom

Flits are sorted by age, then assigned in age order to output ports

→ **43% longer critical path than buffered router**

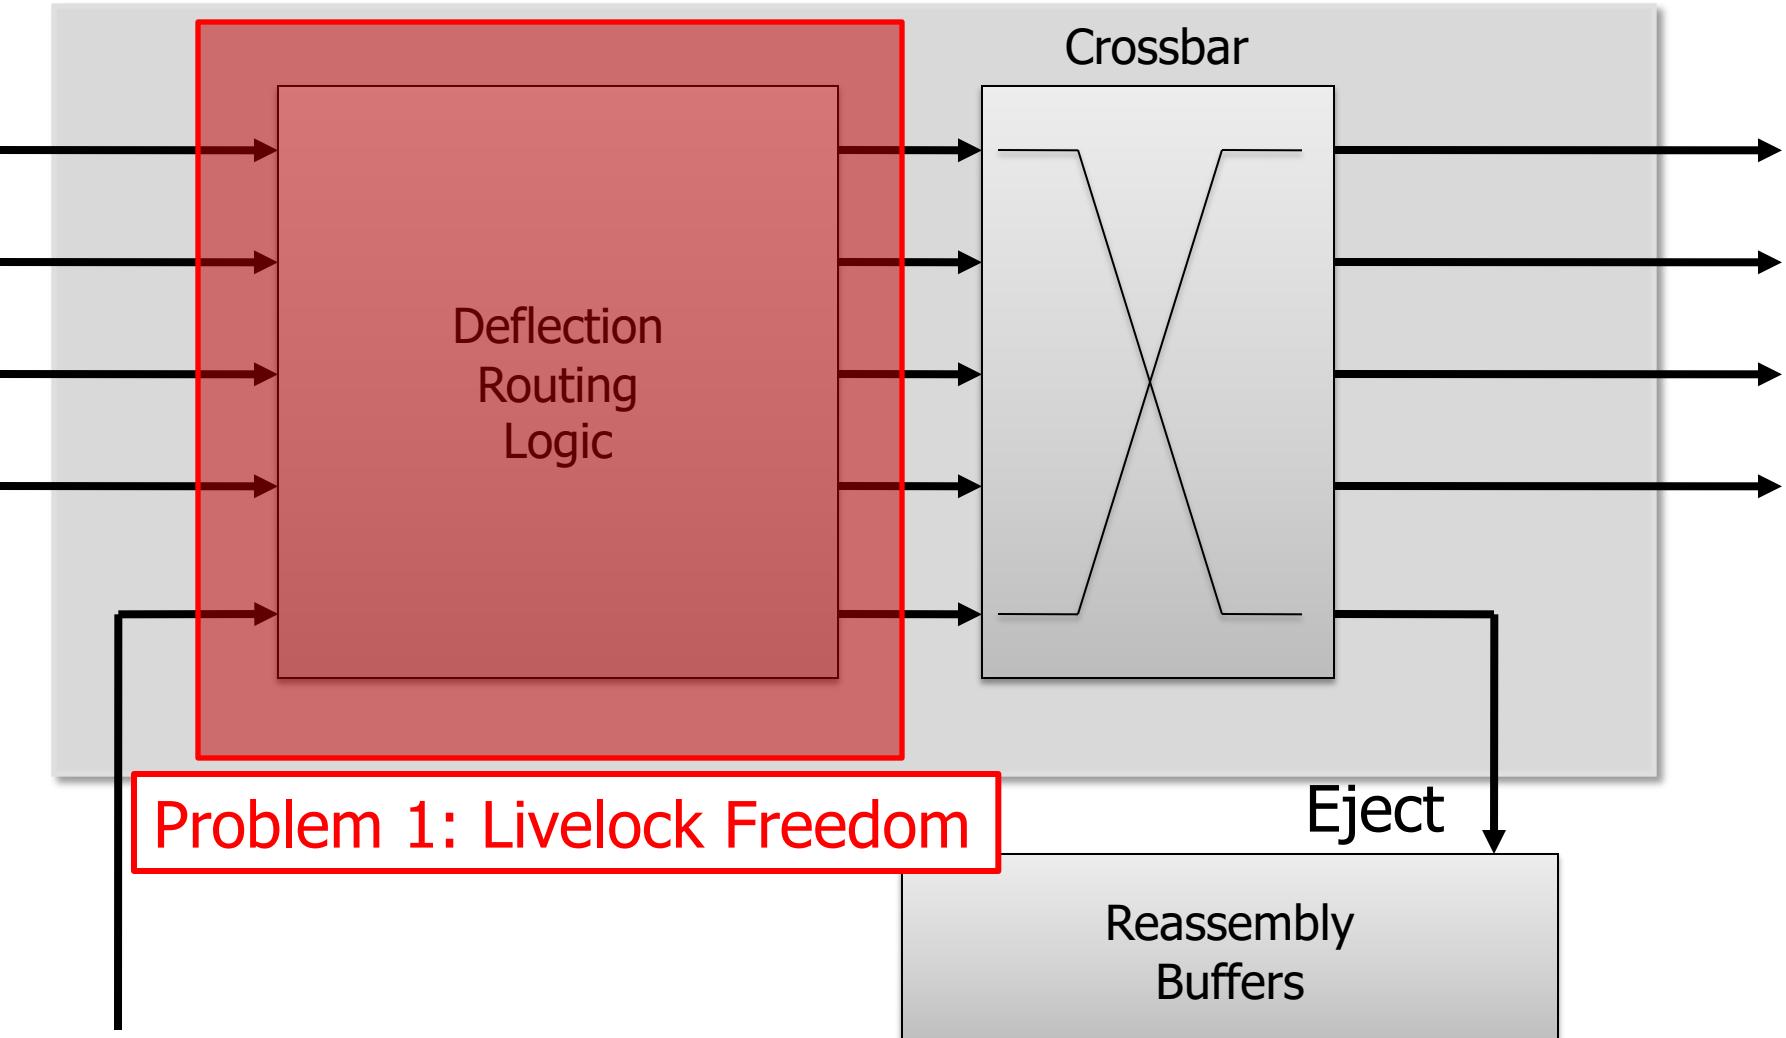
## 2. Must reassemble packets upon arrival

Reassembly buffers must be sized for worst case

→ **4KB per node**

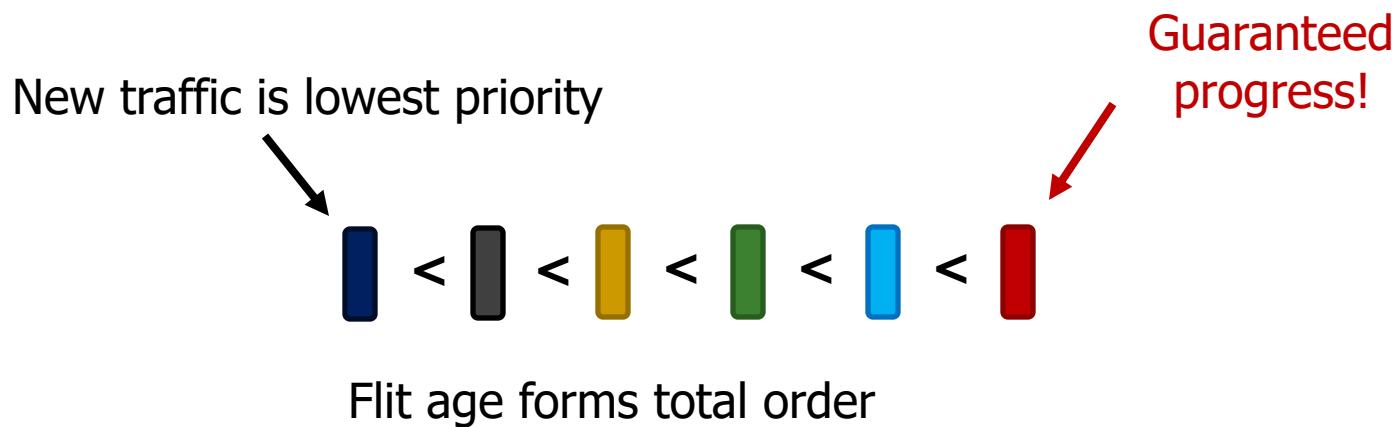
(8x8, 64-byte cache block)

# Problem 1: Livelock Freedom



# Livelock Freedom in Previous Work

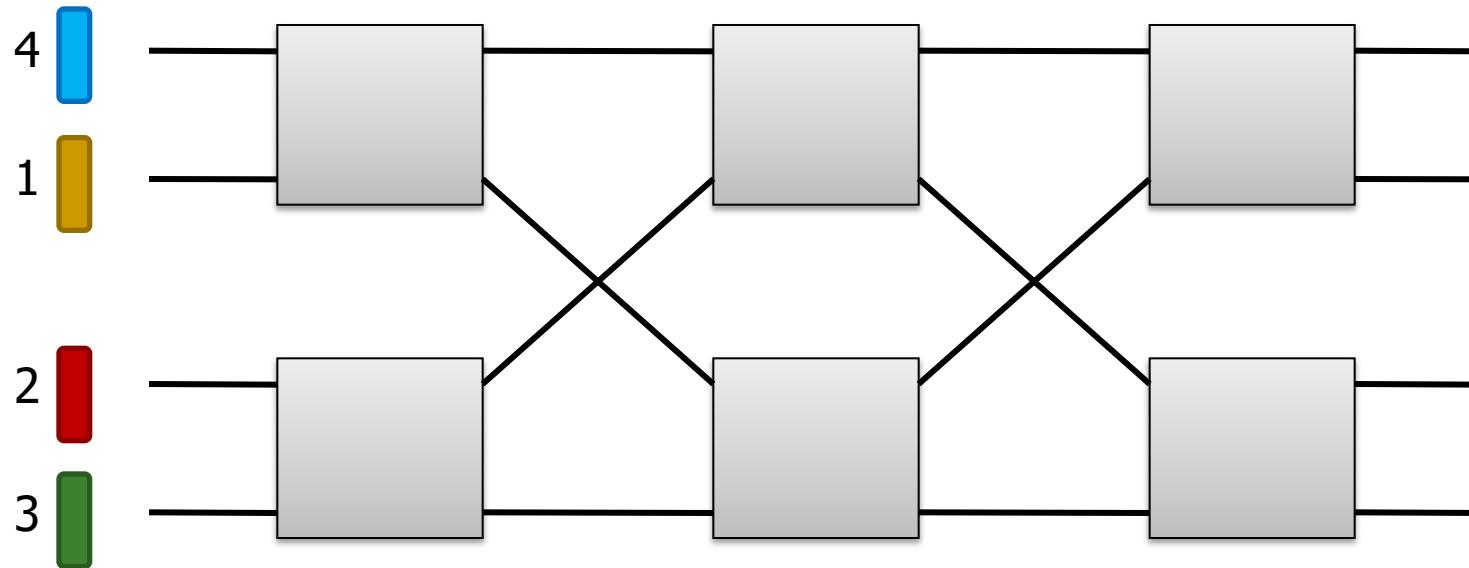
- What stops a flit from deflecting forever?
- All flits are **timestamped**
- **Oldest flits** are assigned their desired ports
- **Total order among flits**



- But what is the **cost** of this?

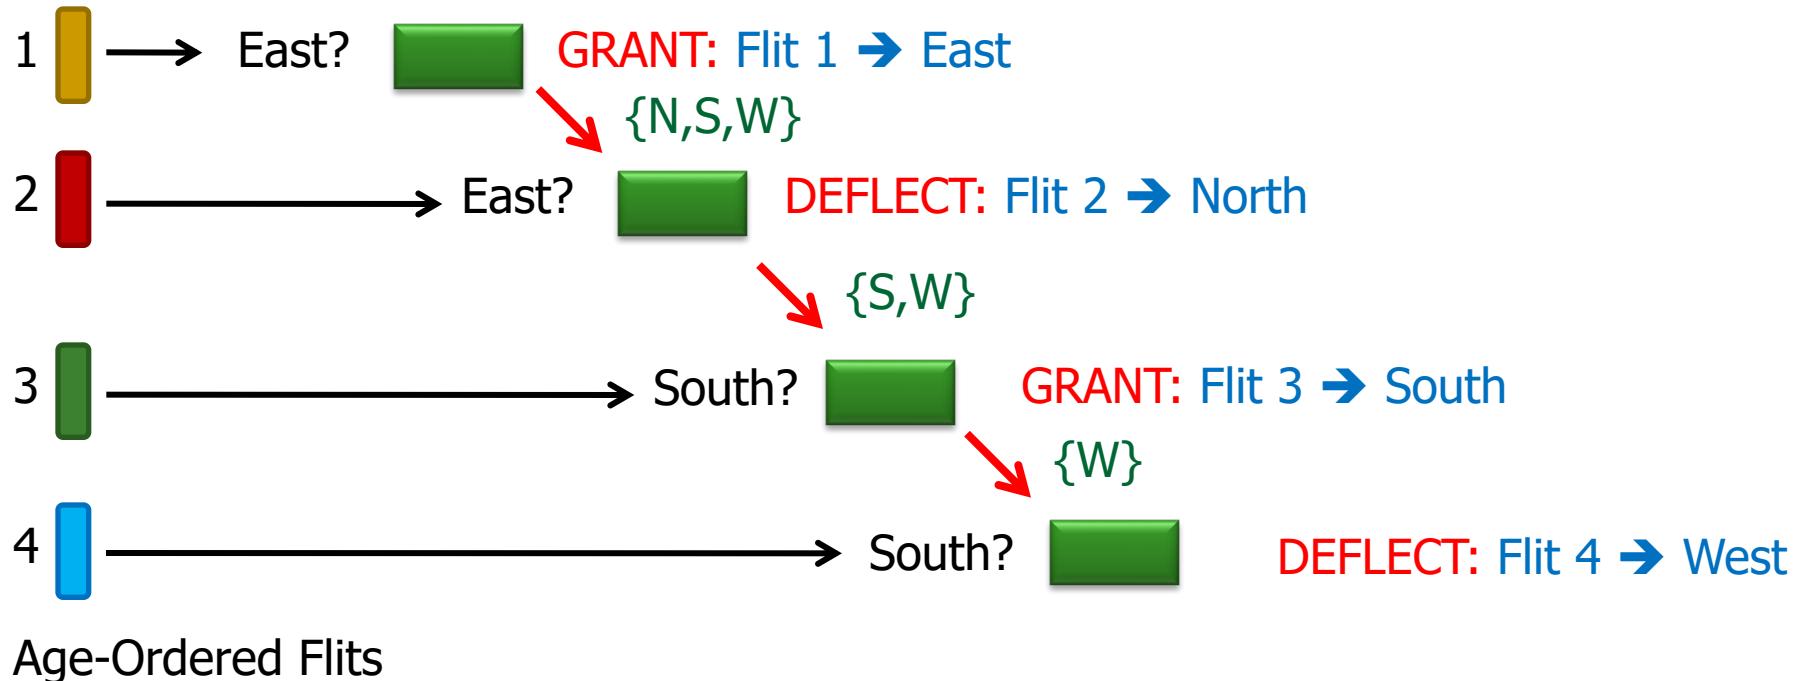
# Age-Based Priorities are Expensive: Sorting

- Router must sort flits by age: long-latency sort network
  - Three comparator stages for 4 flits



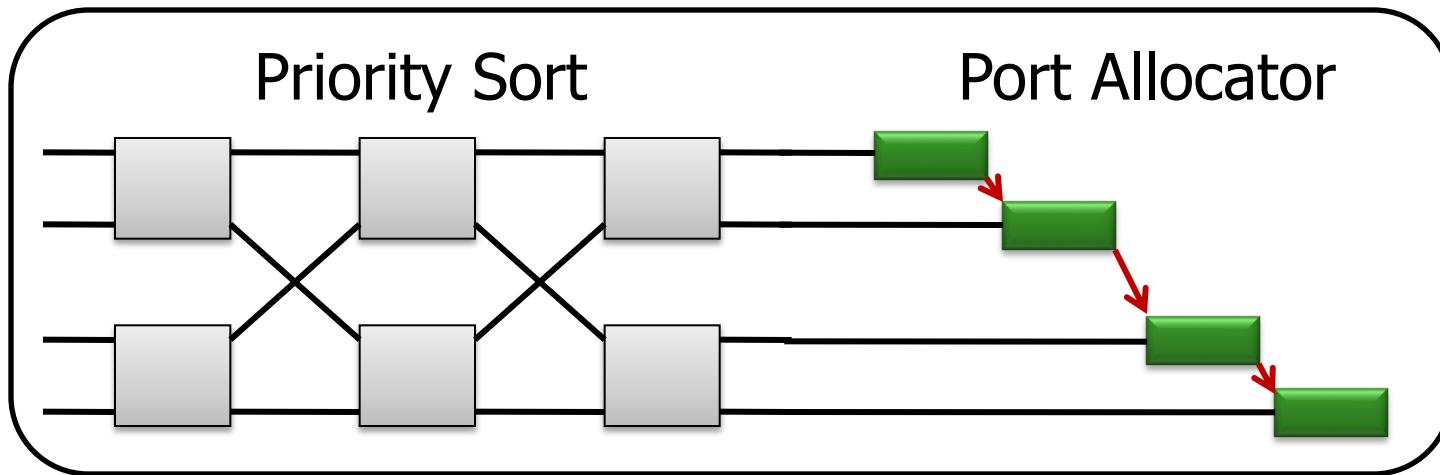
# Age-Based Priorities Are Expensive: Allocation

- After sorting, flits assigned to output ports in priority order
- Port assignment of younger flits depends on that of older flits
  - **sequential dependence** in the port allocator



# Age-Based Priorities Are Expensive

- Overall, **deflection routing logic** based on **Oldest-First** has a **43% longer critical path** than a buffered router



- Question: is there a cheaper way to route while guaranteeing livelock-freedom?

# Solution: Golden Packet for Livelock Freedom

---

- What is *really necessary* for livelock freedom?

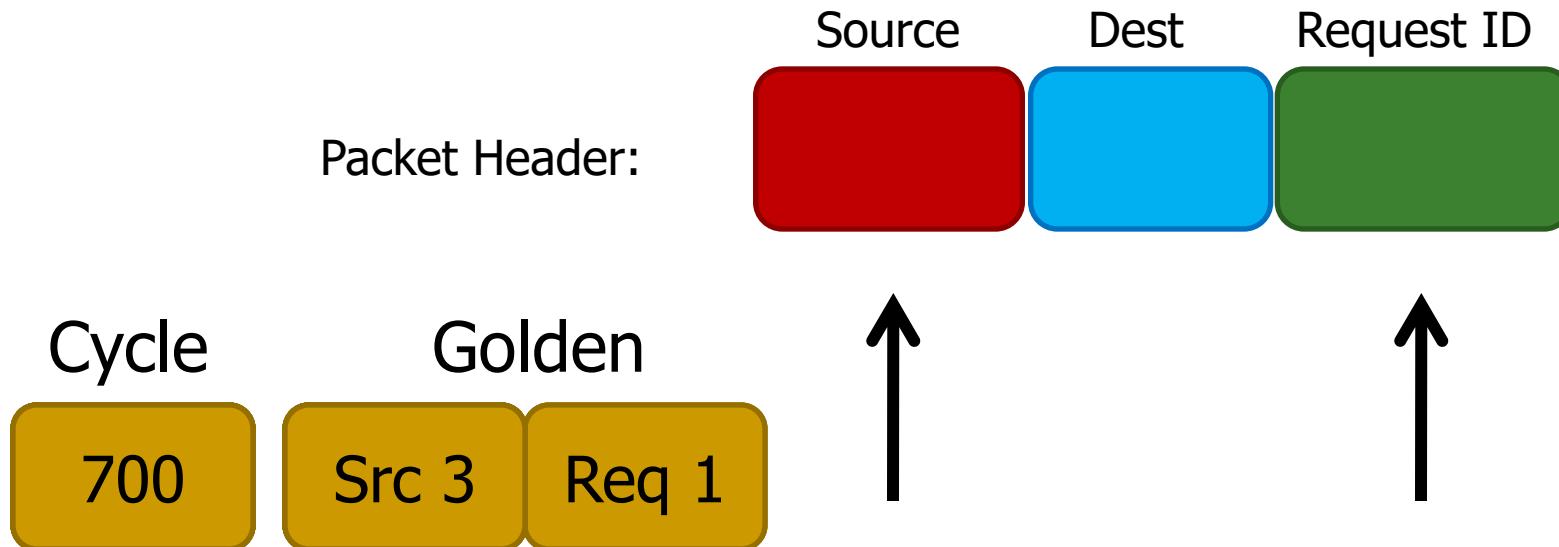
**Key Insight:** No total order. It is enough to:

1. Pick one flit to prioritize until arrival
2. Ensure any flit is eventually picked



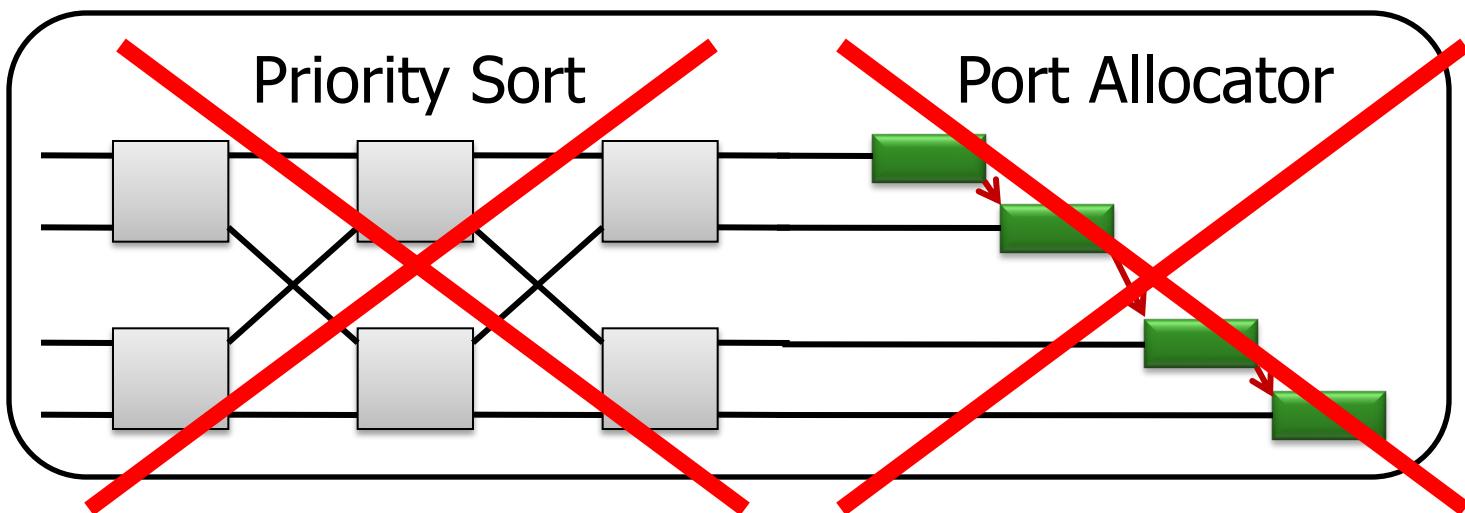
# Which Packet is Golden?

- We select the **Golden Packet** so that:
  1. a given packet stays golden long enough to **ensure arrival**  
→ **maximum no-contention latency**
  2. the selection rotates through **all possible packet IDs**  
→ **static rotation schedule for simplicity**



# What Does Golden Flit Routing Require?

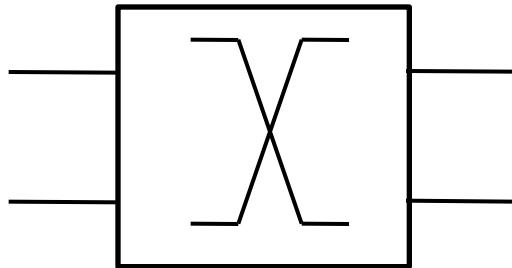
- Only **need** to properly route the Golden Flit
- **First Insight:** no need for full sort
- **Second Insight:** no need for sequential allocation



# Golden Flit Routing With Two Inputs

---

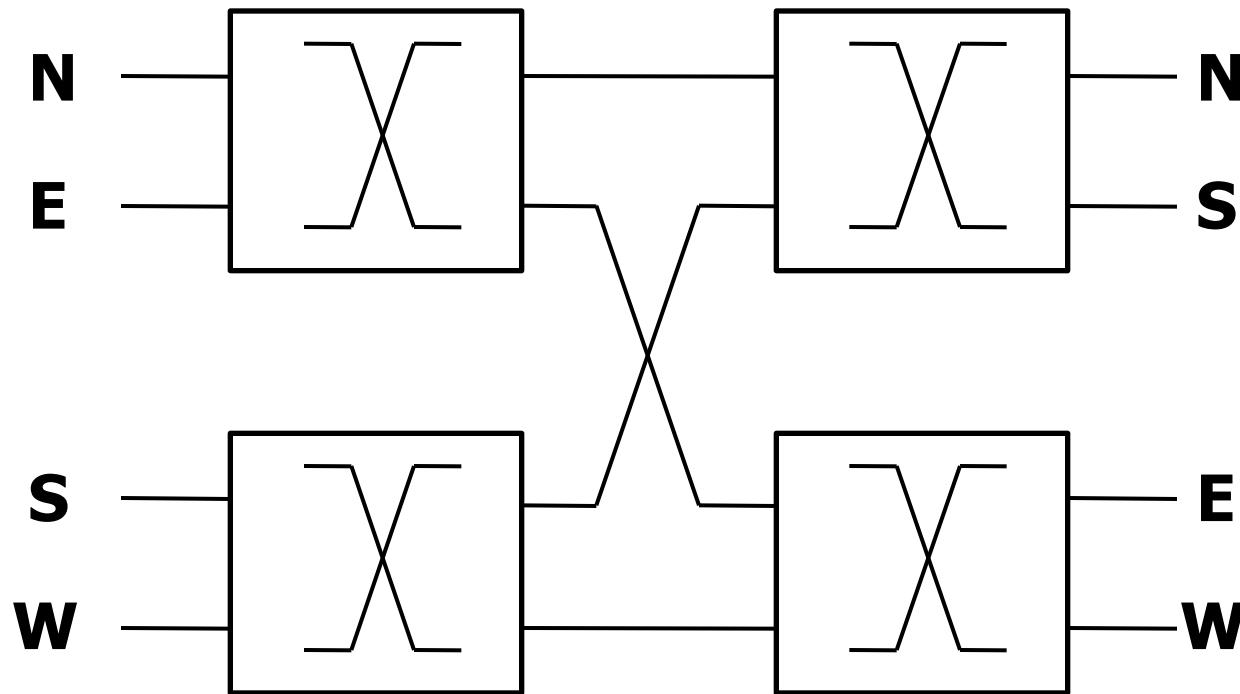
- Let's route the Golden Flit in a two-input router first



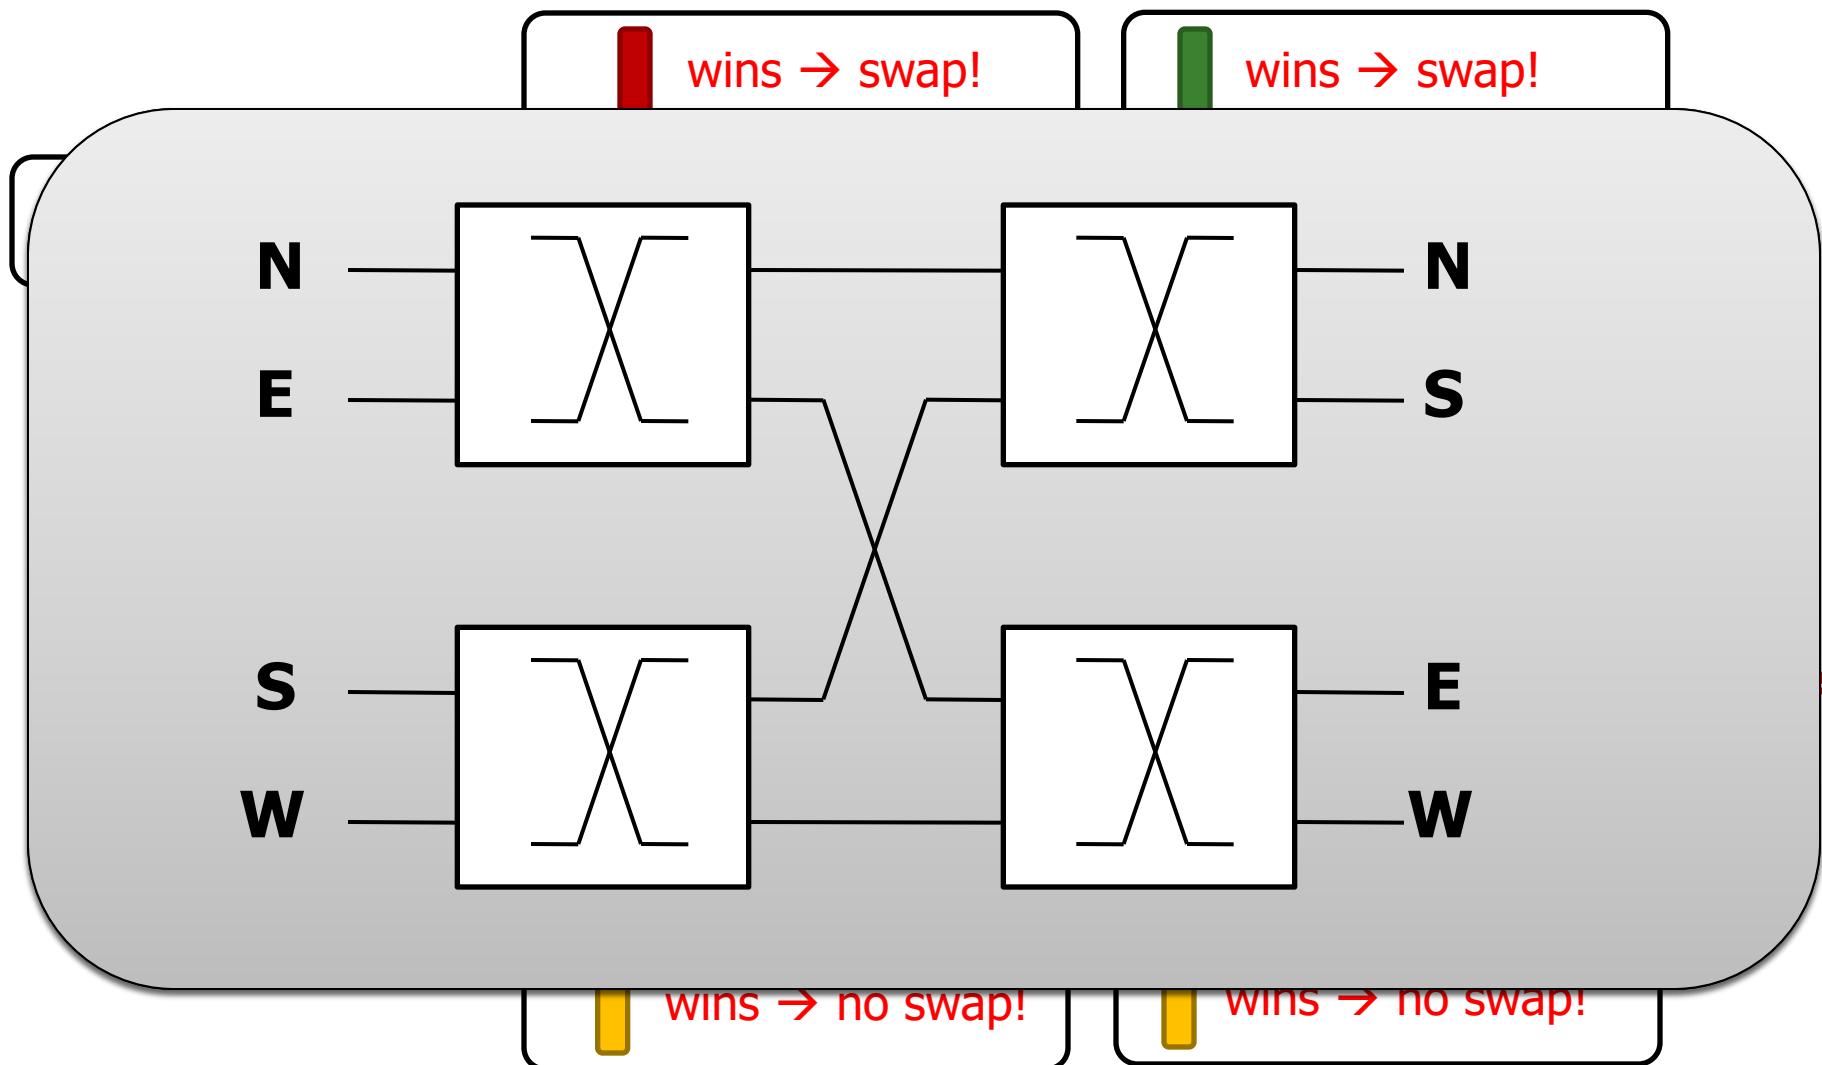
- Step 1:** pick a “winning” flit: Golden Flit, else random
  - Step 2:** steer the winning flit to its desired output and deflect other flit
- **Golden Flit is always routed toward its destination**

# Golden Flit Routing with Four Inputs

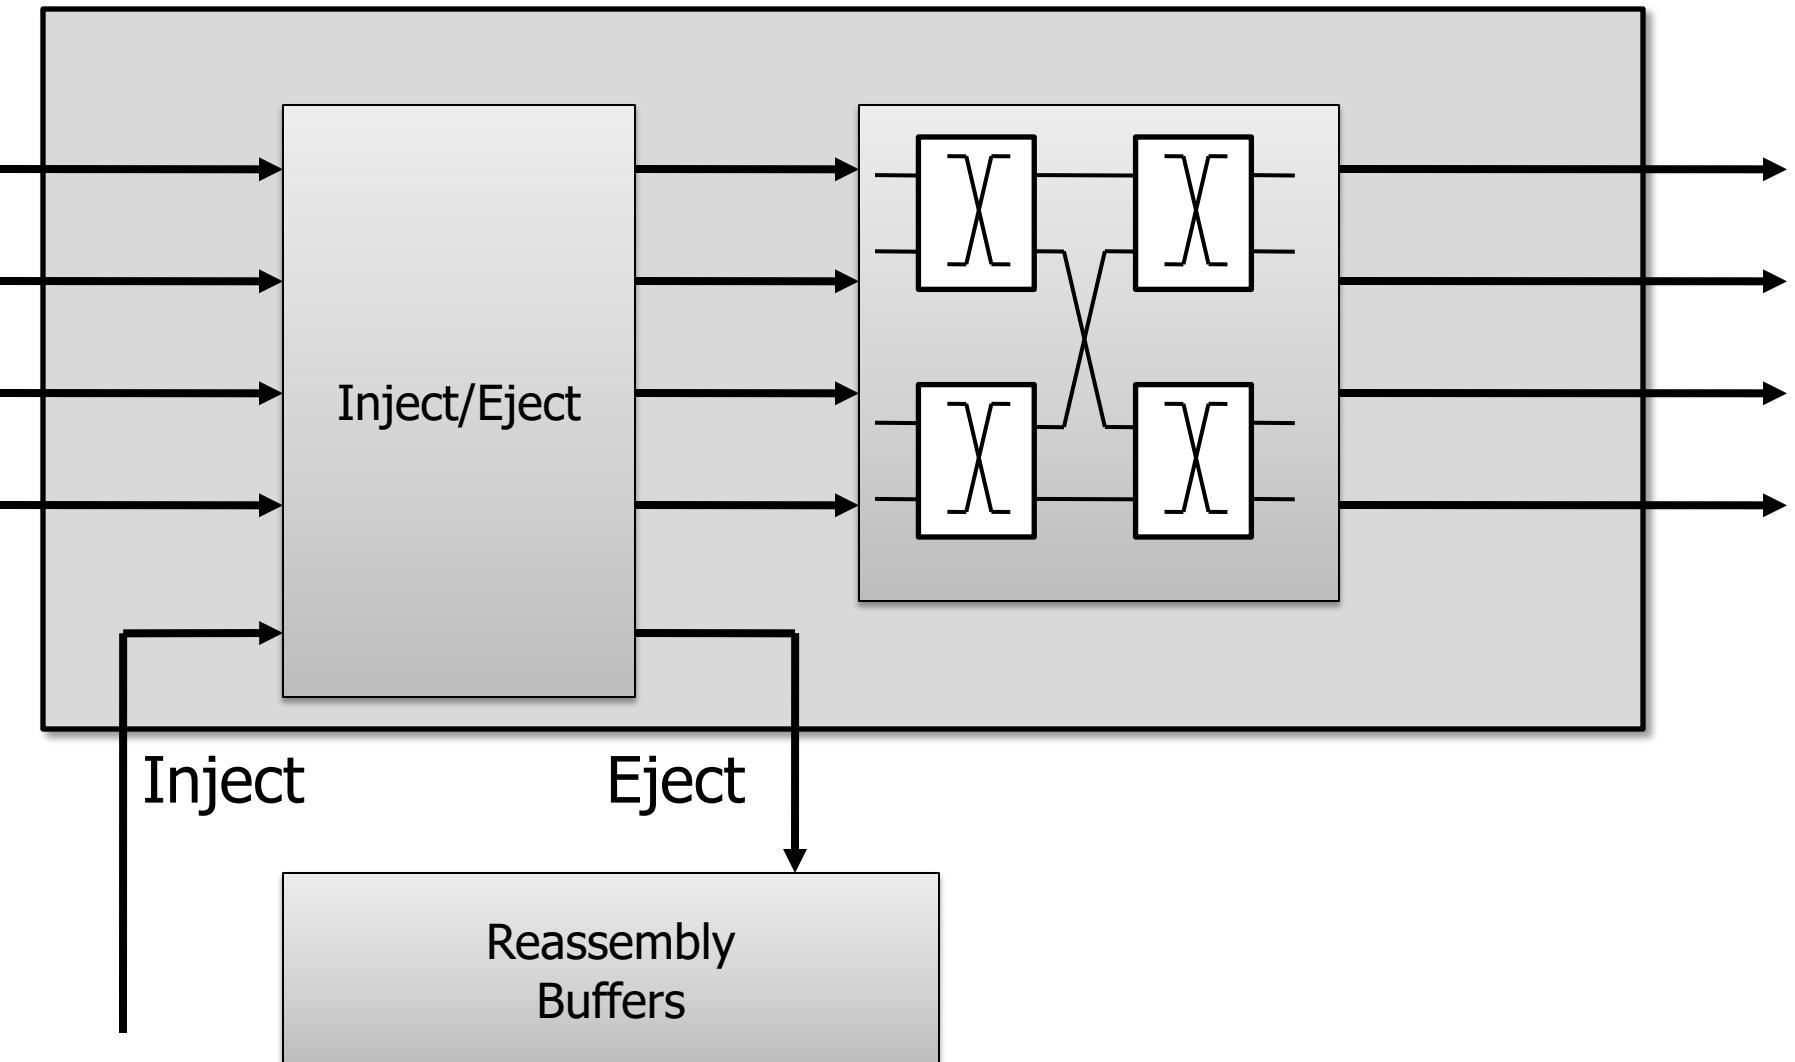
- Each block makes decisions **independently!**
  - **Deflection is a distributed decision**



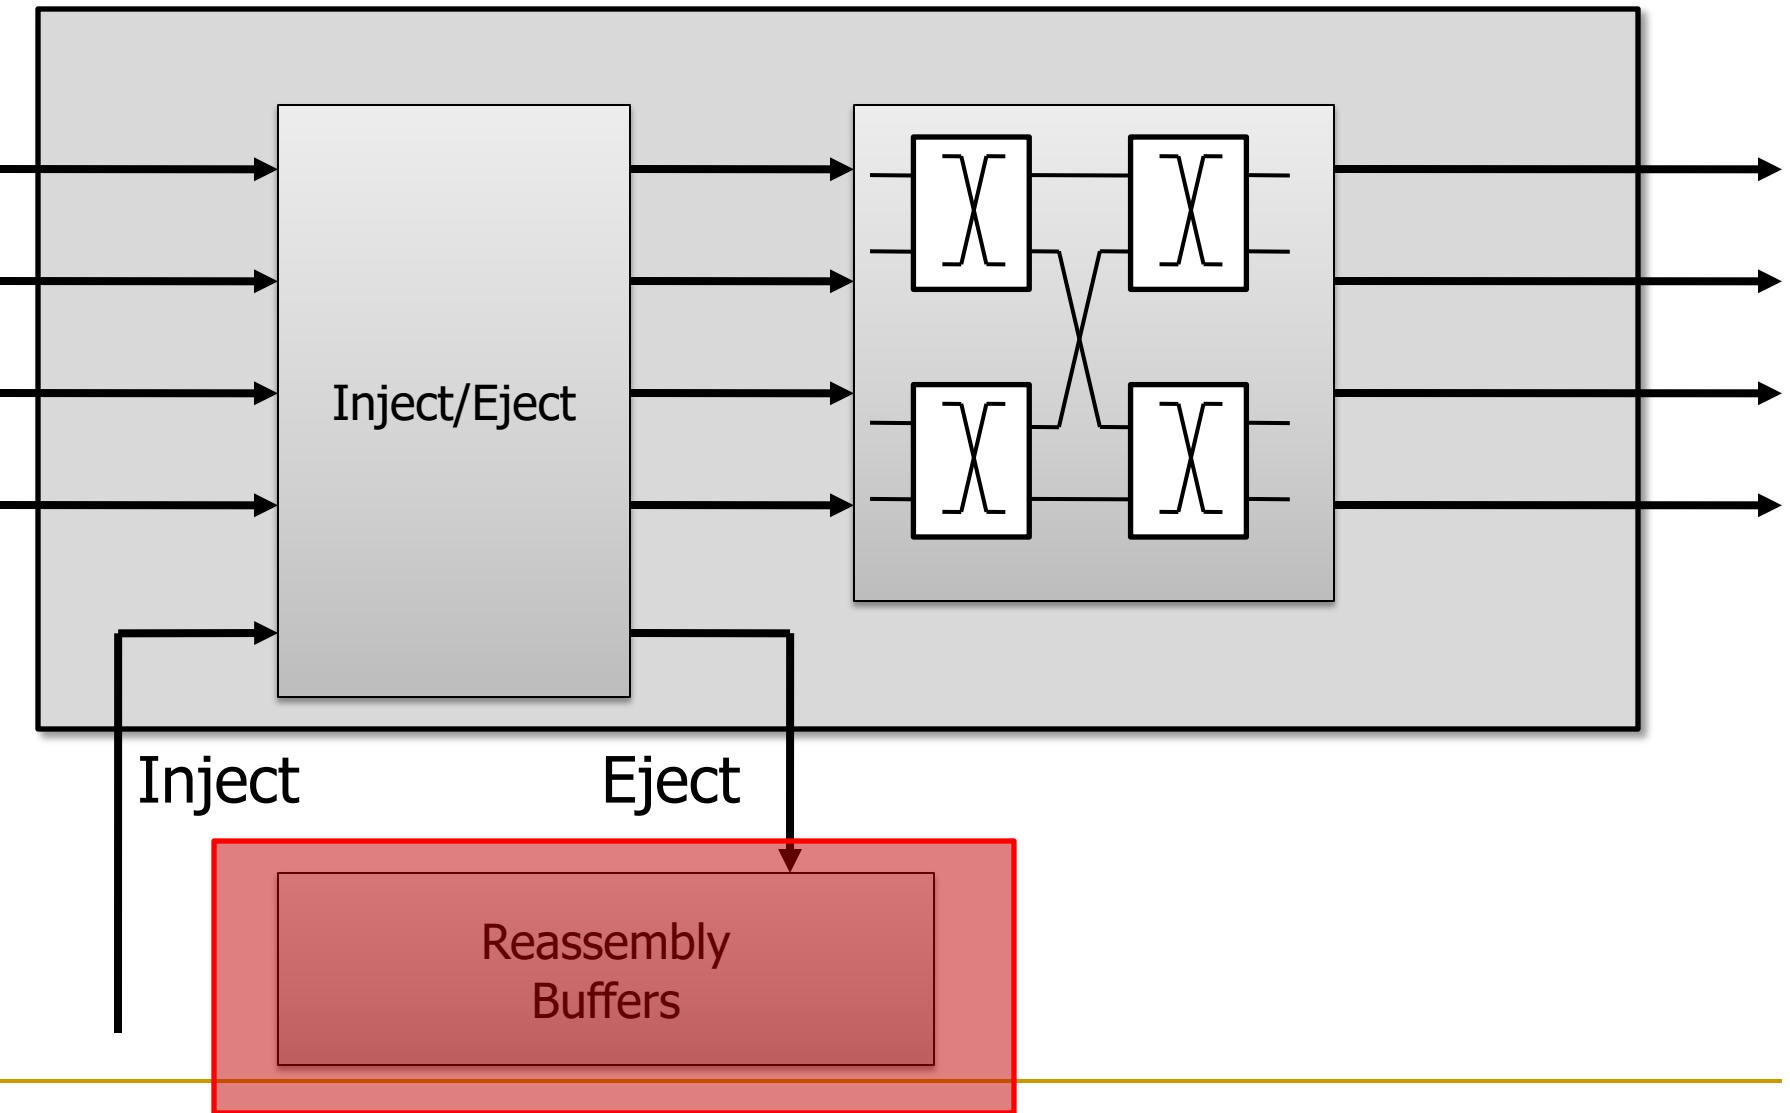
# Permutation Network Operation



# Permutation Network-based Pipeline

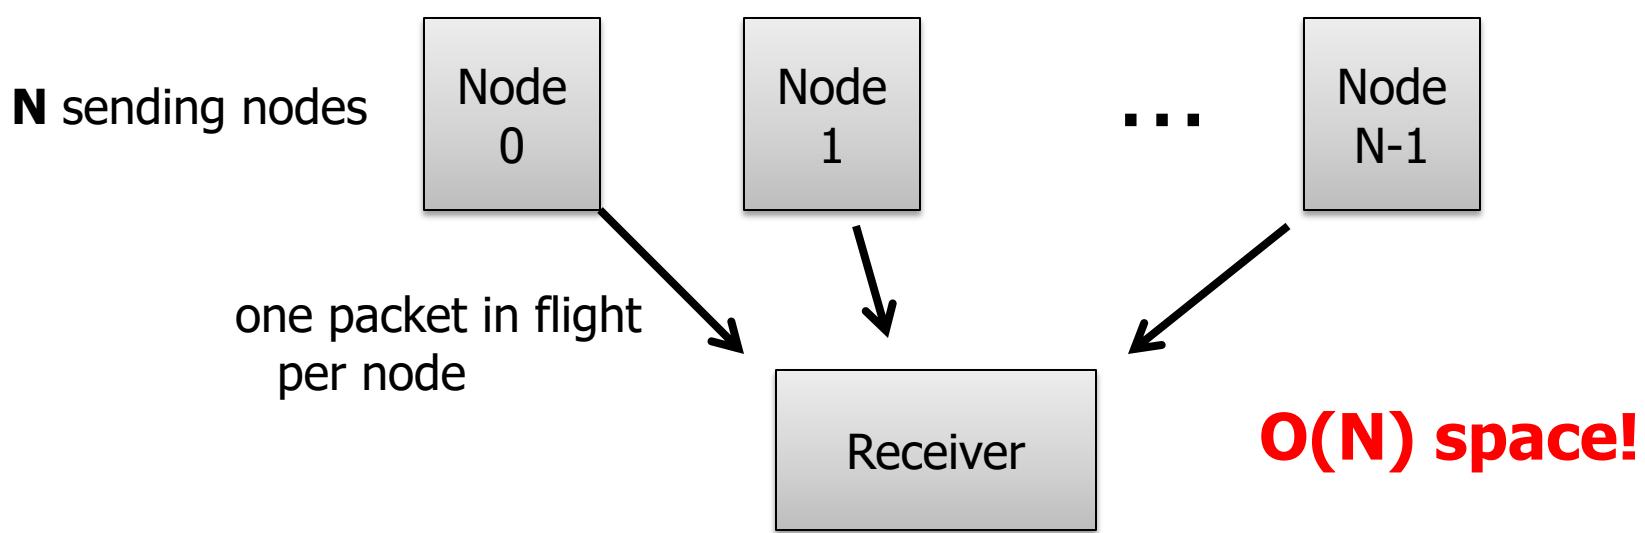


# Problem 2: Packet Reassembly



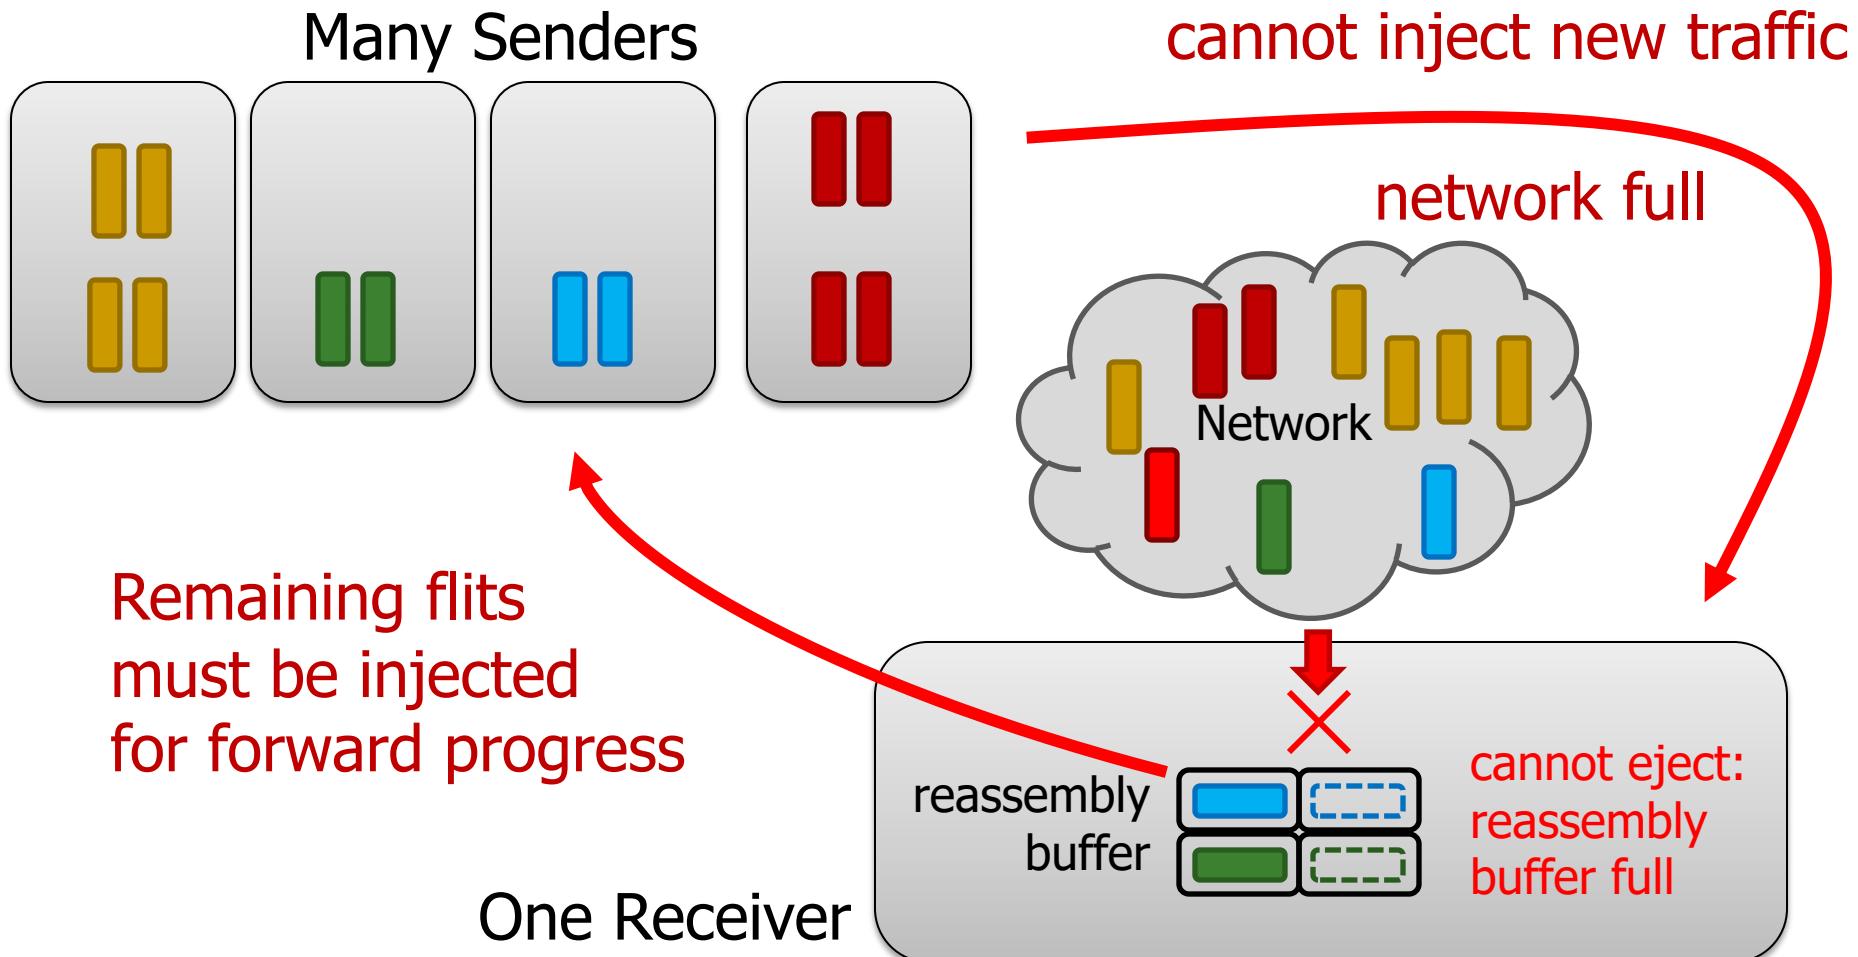
# Reassembly Buffers are Large

- **Worst case:** every node sends a packet to one receiver
- Why can't we make reassembly buffers smaller?



# Small Reassembly Buffers Cause Deadlock

- What happens when reassembly buffer is too small?

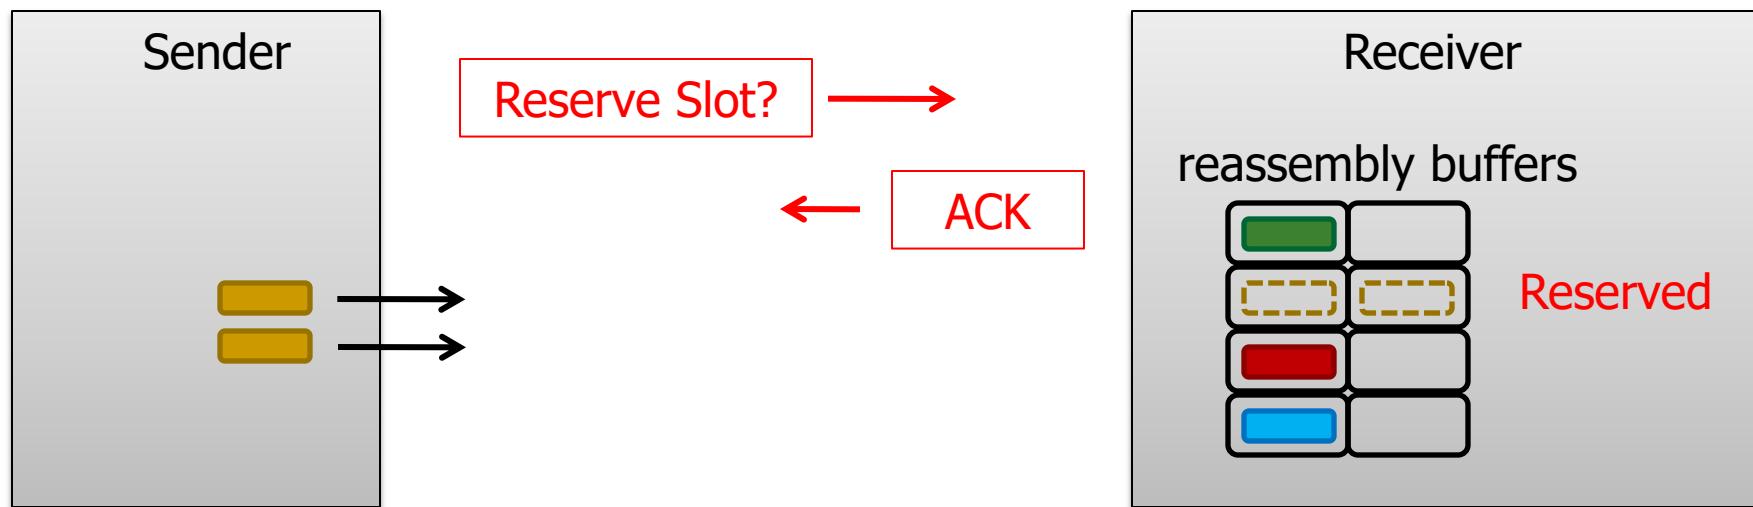


# Reserve Space to Avoid Deadlock?

- What if every sender **asks permission** from the receiver before it sends?

→ **adds additional delay to every request**

- 1. Reserve Slot
- 2. ACK
- 3. Send Packet



# Escaping Deadlock with Retransmissions

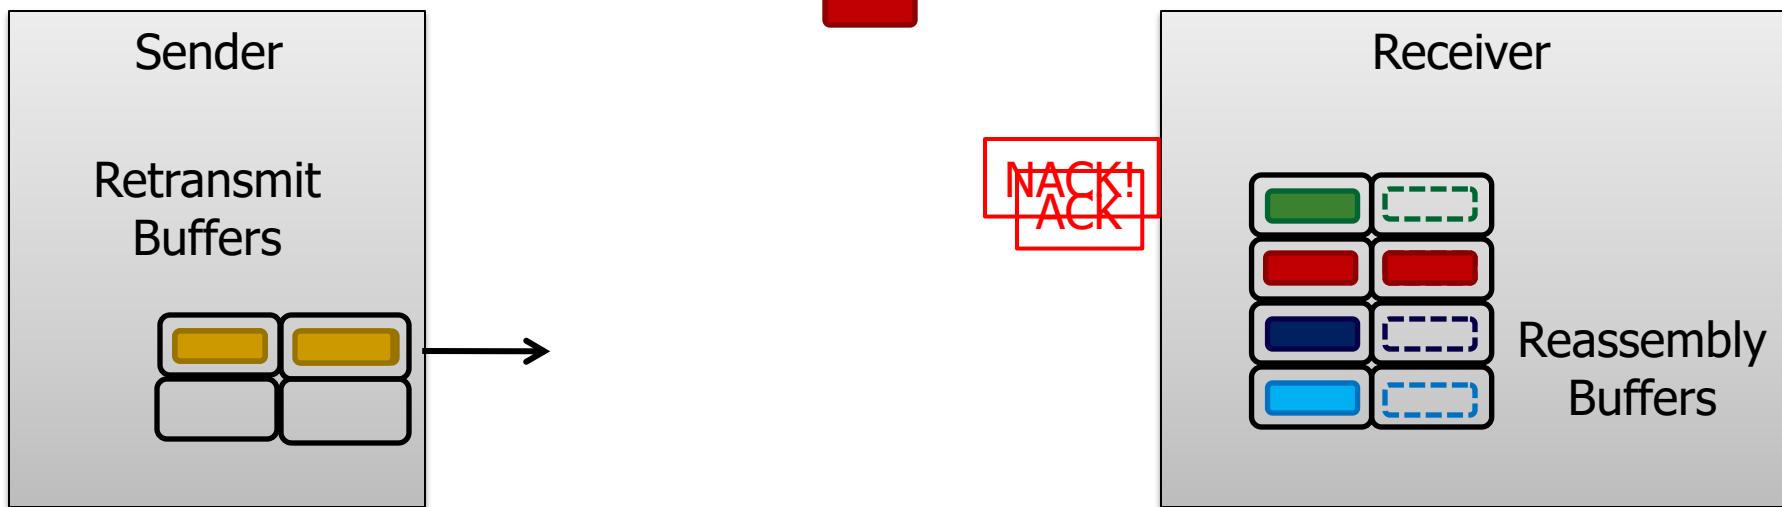
- Sender is optimistic instead: assume buffer is free
  - If not, receiver **drops** and NACKs; sender **retransmits**

→ no additional delay in best case

→ transmit buffering overhead for all packets

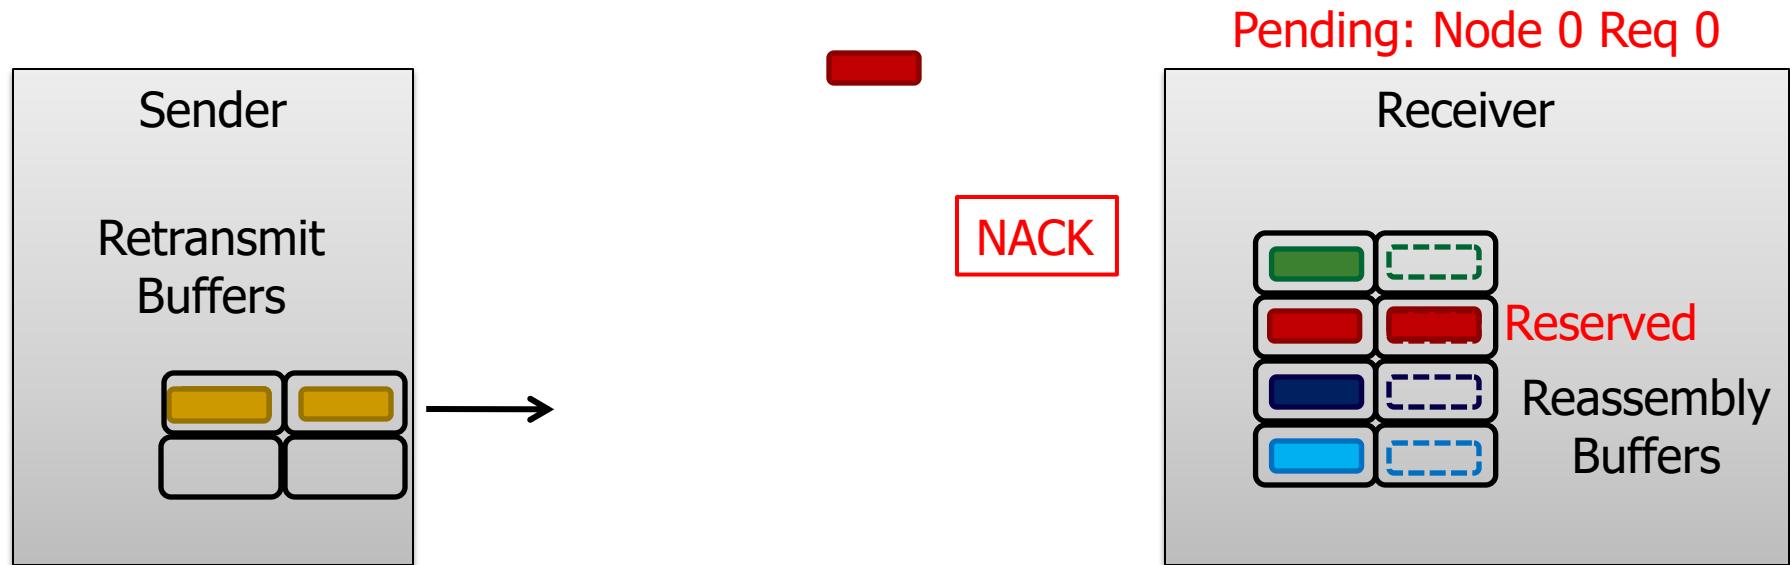
→ potentially many retransmits

- 1. Send (2 flits)
- 2. Drop, NACK
- 3. Other packet completes
- 4. Retransmit packet
- 5. ACK
- 6. Sender frees data

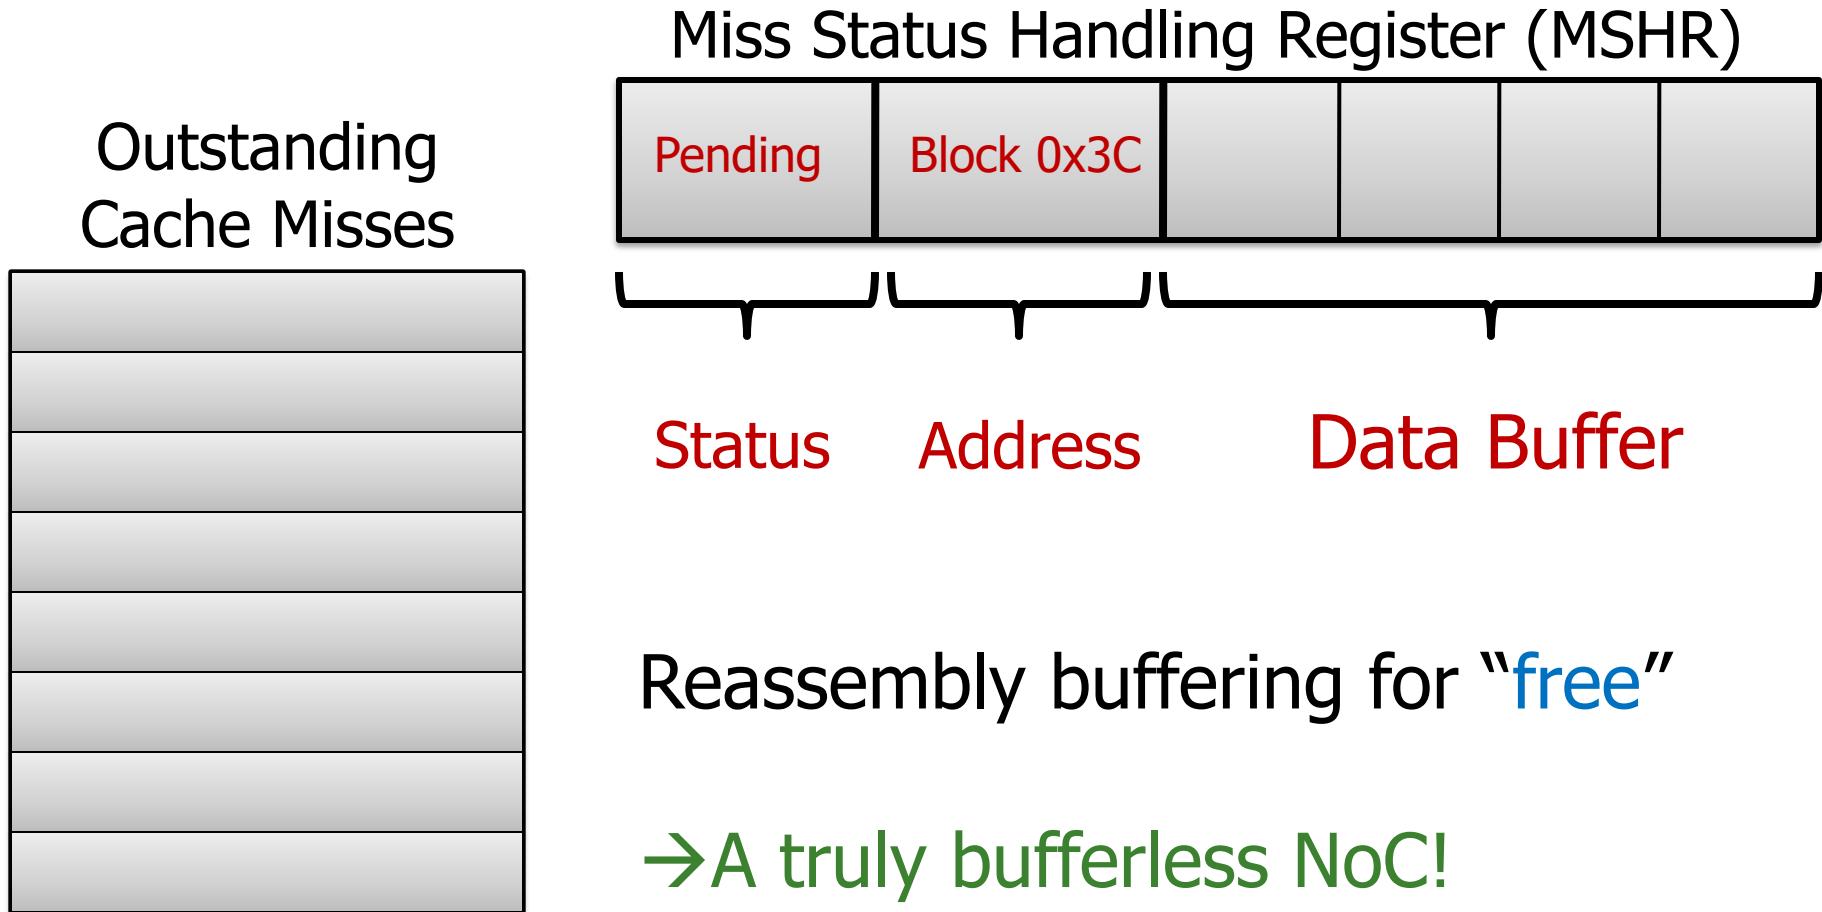


# Solution: Retransmitting Only Once

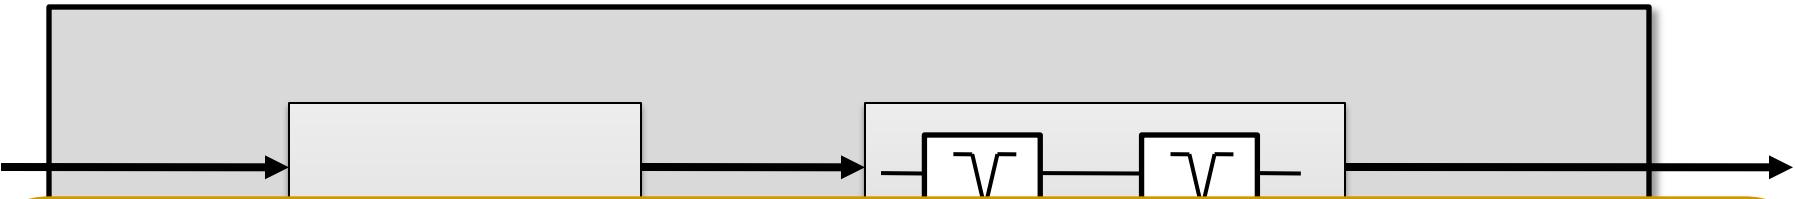
- **Key Idea:** Retransmit only when space becomes available.
  - Receiver **drops packet** if full; **notes which packet it drops**
  - When space frees up, receiver **reserves space** so retransmit is successful
  - Receiver notifies sender to retransmit



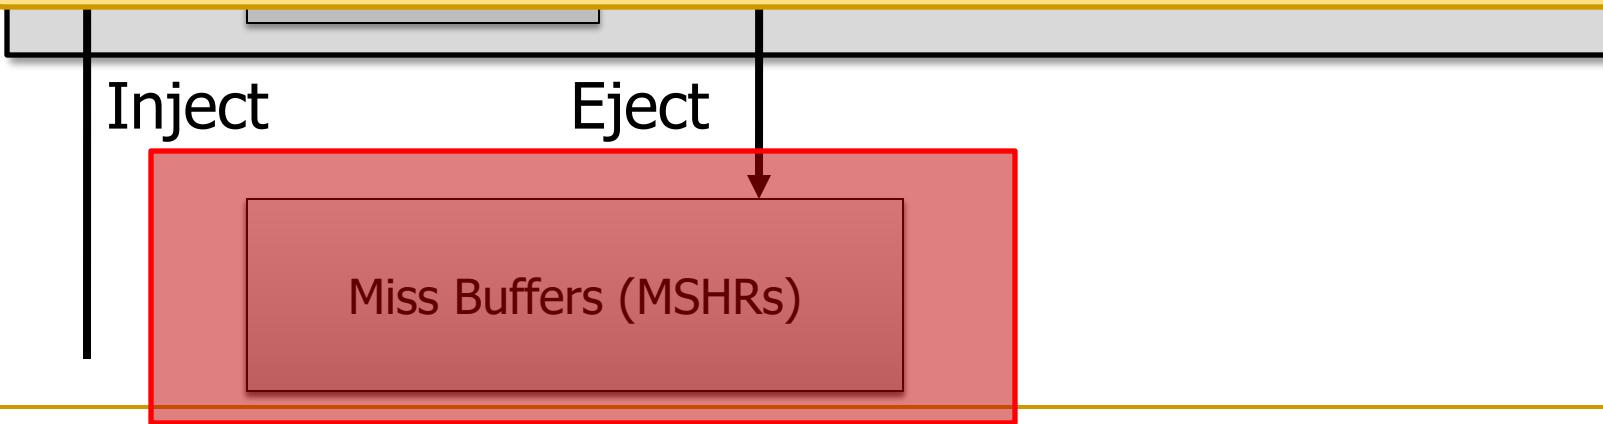
# Use MSHRs as Reassembly Buffers



# Using MSHRs as Reassembly Buffers

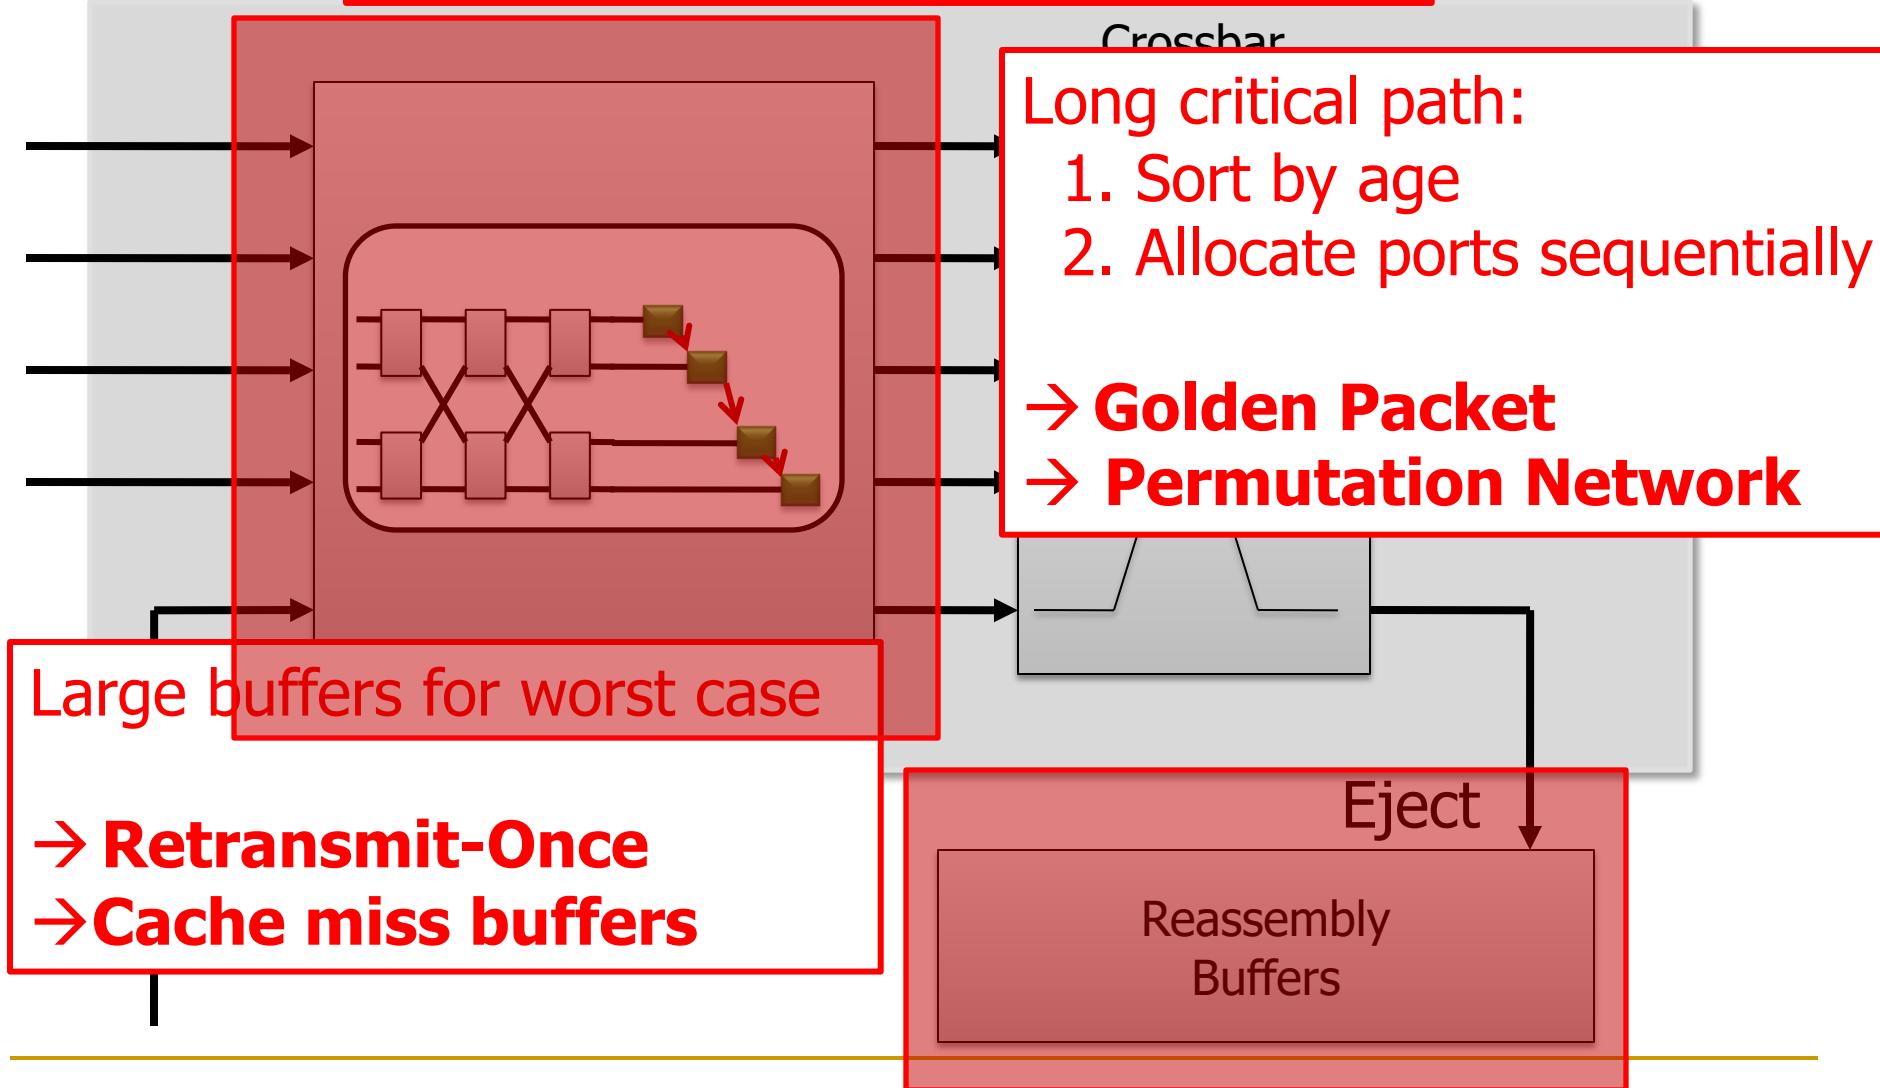


👍 Using miss buffers for reassembly makes this a **truly bufferless network.**

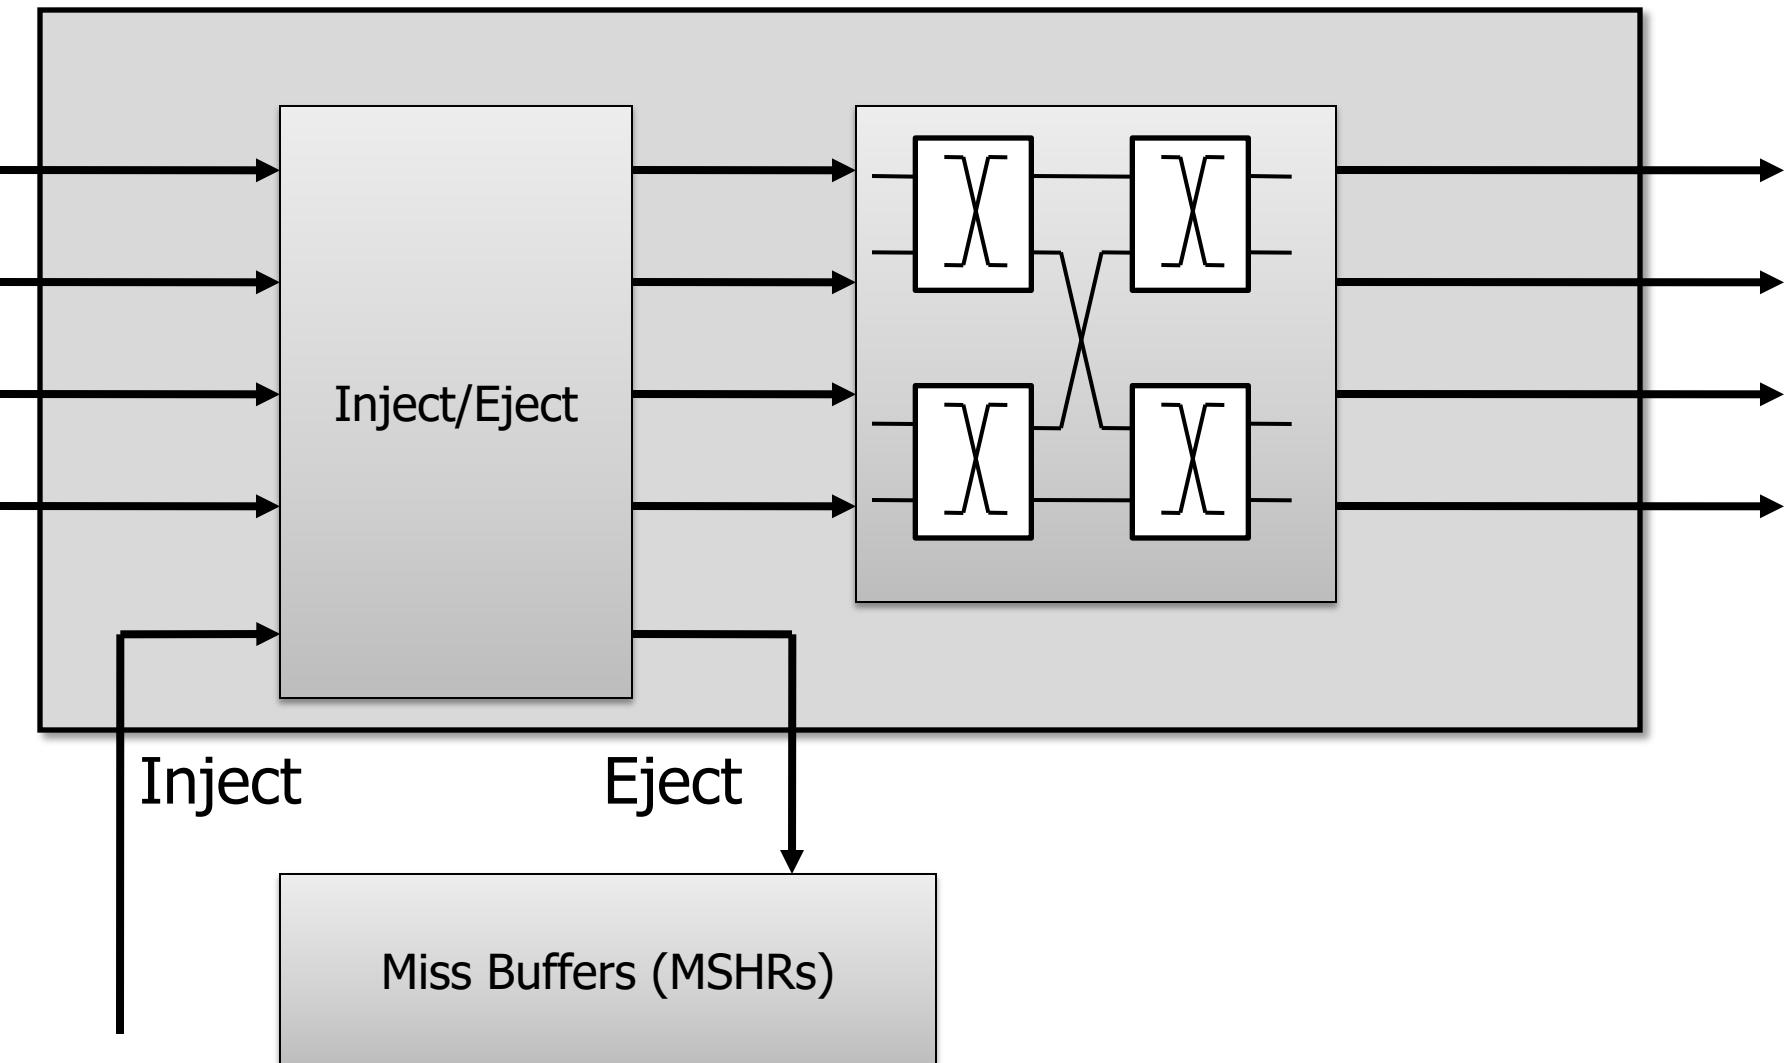


# CHIPPER: Cheap Interconnect Partially-Permuting Router

## Baseline Bufferless Deflection Router



# CHIPPER: Cheap Interconnect Partially-Permuting Router



---

# EVALUATION

# Methodology

---

- **Multiprogrammed** workloads: CPU2006, server, desktop
  - 8x8 (64 cores), 39 homogeneous and 10 mixed sets
- **Multithreaded** workloads: SPLASH-2, 16 threads
  - 4x4 (16 cores), 5 applications
- **System configuration**
  - **Buffered** baseline: 2-cycle router, 4 VCs/channel, 8 flits/VC
  - **Bufferless** baseline: 2-cycle latency, FLIT-BLESS
  - Instruction-trace driven, closed-loop, 128-entry OoO window
  - 64KB L1, perfect L2 (**stresses interconnect**), XOR mapping

# Methodology

---

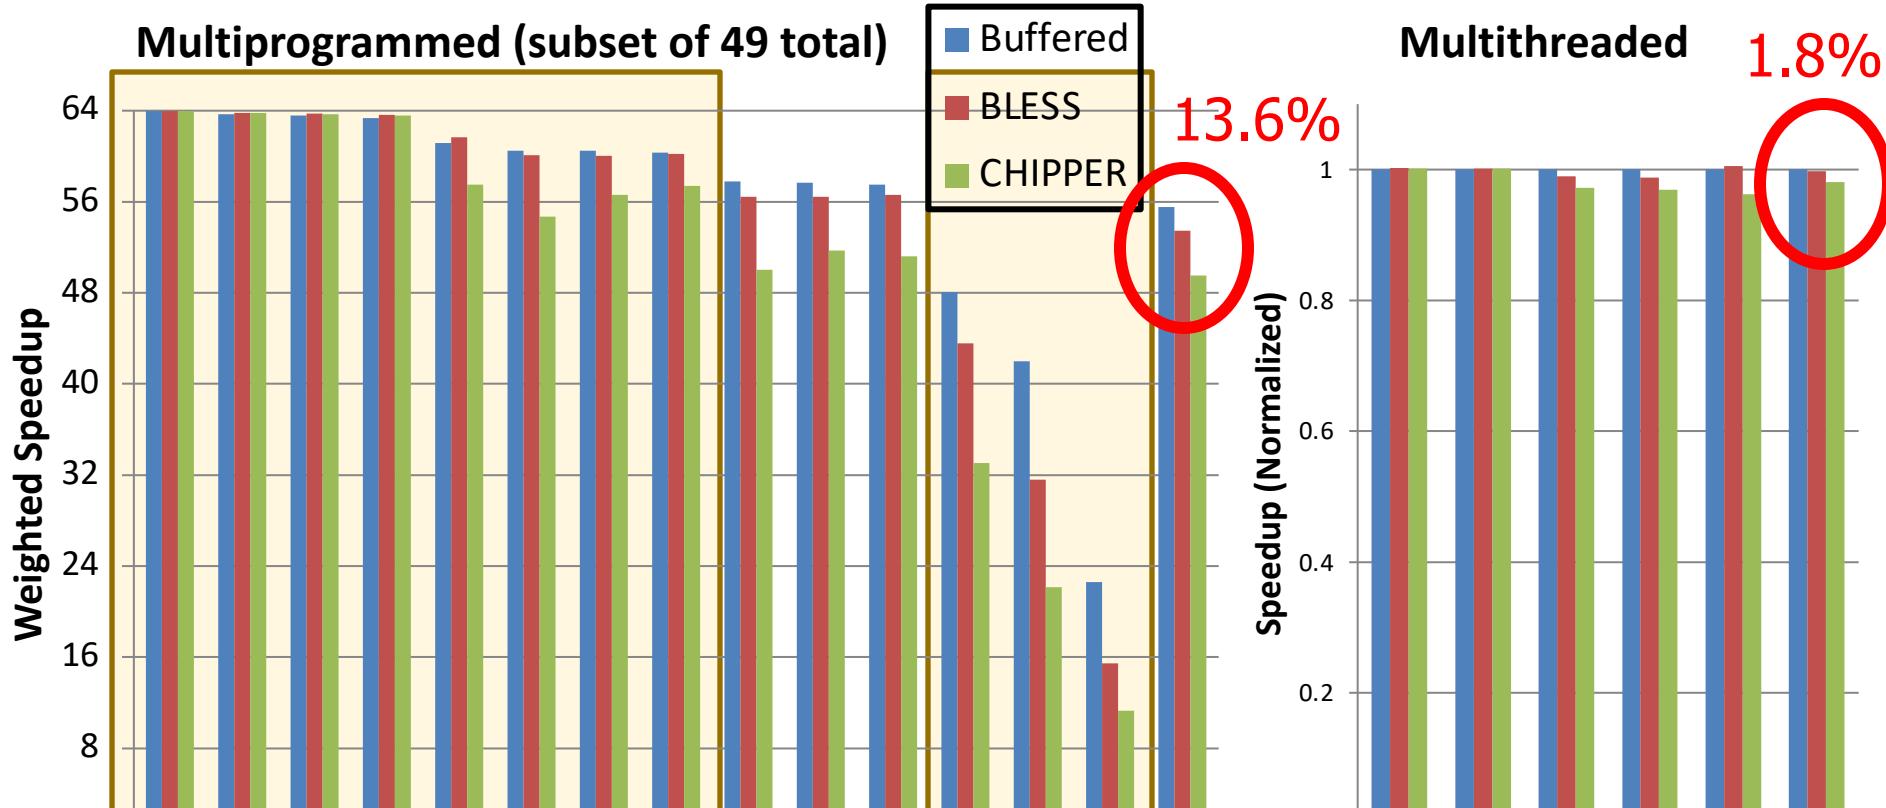
## ■ **Hardware modeling**

- **Verilog models** for CHIPPER, BLESS, buffered logic
  - Synthesized with commercial 65nm library
- **ORION** for crossbar, buffers and links

## ■ **Power**

- Static and dynamic power from hardware models
- Based on event counts in cycle-accurate simulations

# Results: Performance Degradation

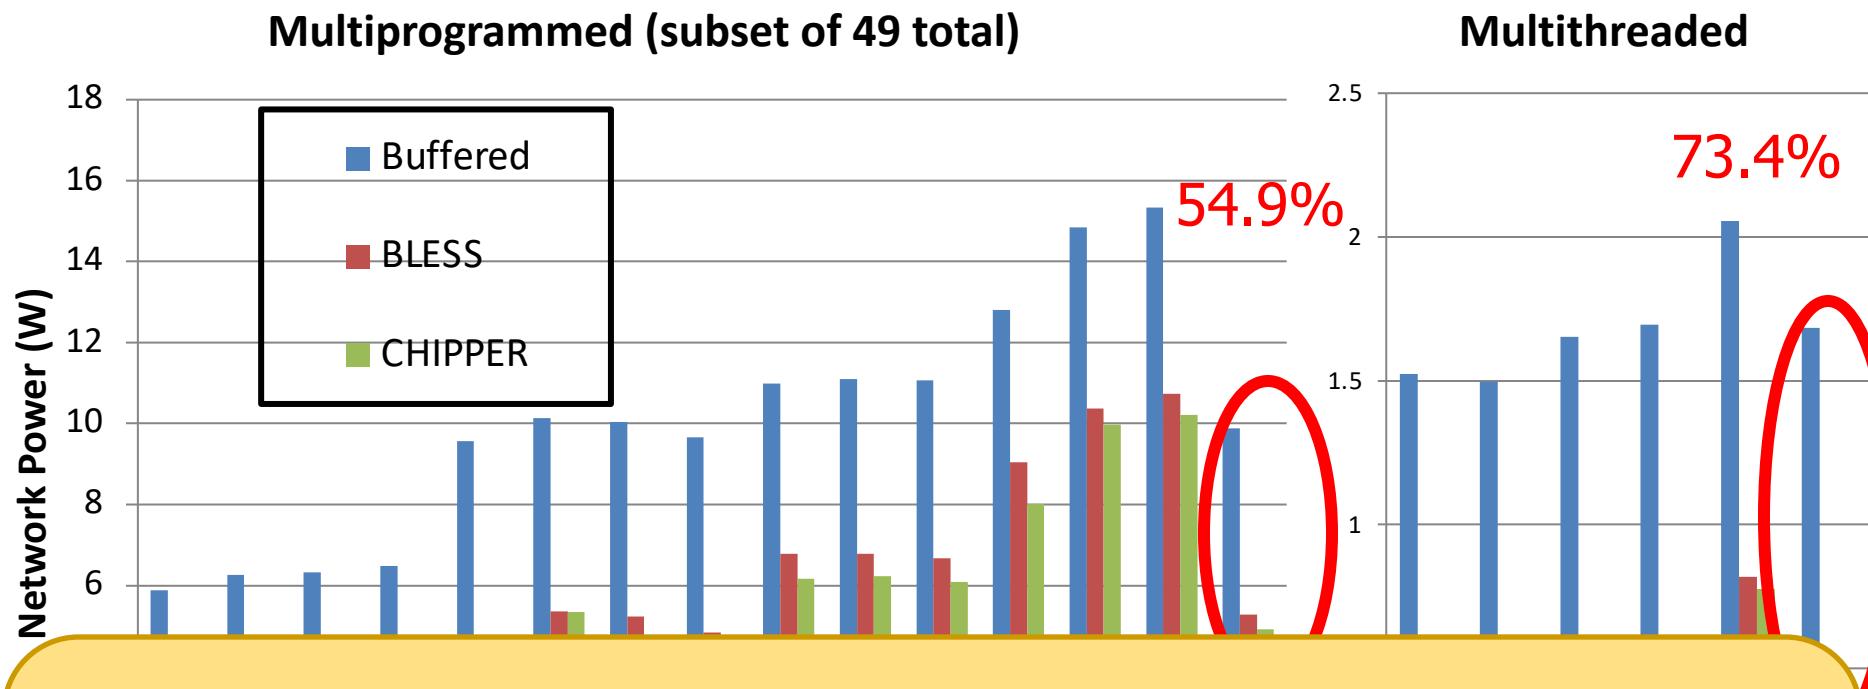


Small loss for low-to-medium-intensity workloads

5.0%

49.8%<sup>A</sup>

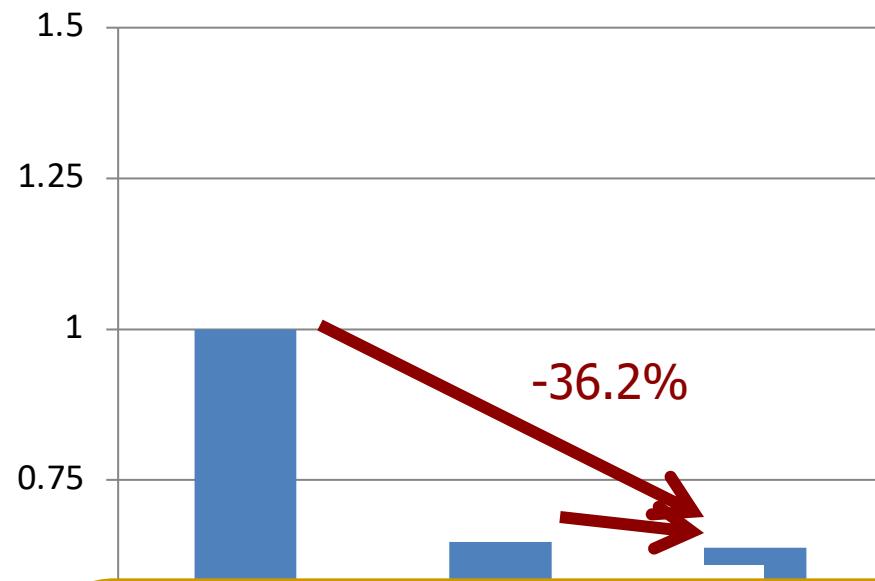
# Results: Power Reduction



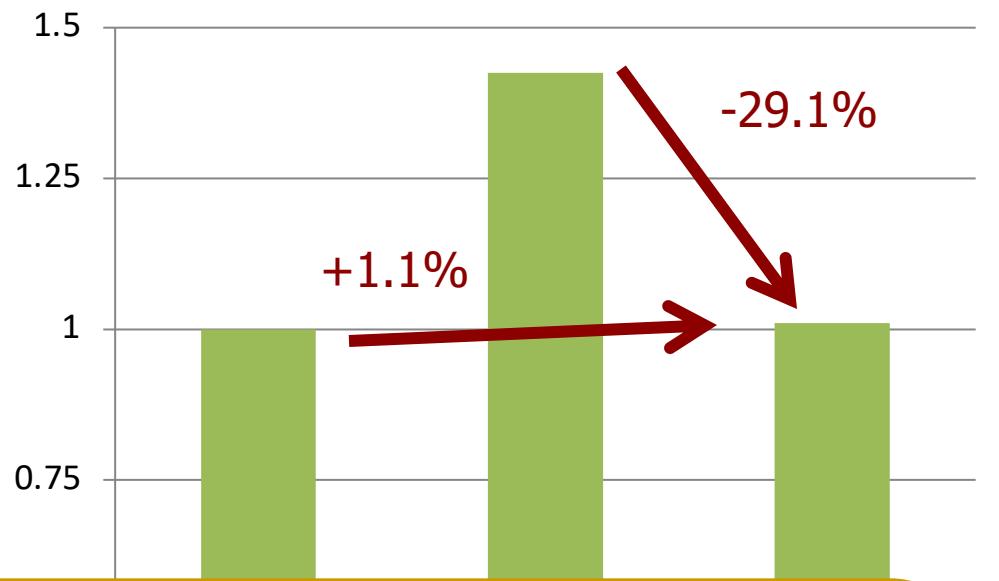
- 👍 Removing buffers → large power savings
- 👍 Slight savings from BLESS to CHIPPER

# Results: Area and Critical Path Reduction

Normalized Router Area



Normalized Critical Path



**CHIPPER maintains area savings of BLESS**



Critical path **becomes competitive** to buffered

# CHIPPER Router: Conclusions

---

- Two key issues in bufferless deflection routing
    - livelock freedom and packet reassembly
  - Bufferless deflection routers were high-complexity and impractical
    - Oldest-first prioritization → long critical path in router
    - No end-to-end flow control for reassembly → prone to deadlock with reasonably-sized reassembly buffers
  - CHIPPER is a new, practical bufferless deflection router
    - Golden packet prioritization → short critical path in router
    - Retransmit-once protocol → deadlock-free packet reassembly
    - Cache miss buffers as reassembly buffers → truly bufferless network
  - CHIPPER frequency comparable to buffered routers at much lower area and power cost, and minimal performance loss
-

# More on CHIPPER

---

- Chris Fallin, Chris Craik, and Onur Mutlu,  
**"CHIPPER: A Low-Complexity Bufferless Deflection Router"**  
*Proceedings of the 17th International Symposium on High-Performance Computer Architecture (HPCA)*, pages 144-155, San Antonio, TX, February 2011. [Slides \(pptx\)](#)
  - An extended version as *SAFARI Technical Report*, TR-SAFARI-2010-001, Carnegie Mellon University, December 2010.

## CHIPPER: A Low-complexity Bufferless Deflection Router

Chris Fallin cfallin@cmu.edu	Chris Craik craik@cmu.edu	Onur Mutlu onur@cmu.edu
---------------------------------	------------------------------	----------------------------

Computer Architecture Lab (CALCM)  
Carnegie Mellon University

# Minimally-Buffered Deflection Routing

---

- Bufferless deflection routing offers **reduced power & area**
  - But, high deflection rate hurts **performance at high load**
  - **MinBD** (Minimally-Buffered Deflection Router) introduces:
    - Side buffer to hold **only** flits that would have been deflected
    - Dual-width ejection to address ejection bottleneck
    - Two-level prioritization to avoid unnecessary deflections
  - MinBD yields **reduced power (31%) & reduced area (36%)** relative to **buffered** routers
  - MinBD yields **improved performance (8.1% at high load)** relative to **bufferless** routers → closes half of perf. gap
  - MinBD has the **best energy efficiency** of all evaluated designs with **competitive performance**
-

# Minimally-Buffered Deflection Routing

---

- Chris Fallin, Greg Nazario, Xiangyao Yu, Kevin Chang, Rachata Ausavarungnirun, and Onur Mutlu,

**"MinBD: Minimally-Buffered Deflection Routing for Energy-Efficient Interconnect"**

*Proceedings of the 6th ACM/IEEE International Symposium on Networks on Chip (NOCS), Lyngby, Denmark, May 2012.* [Slides \(pptx\)](#) [\(pdf\)](#)

***One of the five papers nominated for the Best Paper Award by the Program Committee.***

## MinBD: Minimally-Buffered Deflection Routing for Energy-Efficient Interconnect

Chris Fallin, Greg Nazario, Xiangyao Yu<sup>†</sup>, Kevin Chang, Rachata Ausavarungnirun, Onur Mutlu

Carnegie Mellon University  
`{cfallin,gnazario,kevincha,rachata,onur}@cmu.edu`

<sup>†</sup>Tsinghua University & Carnegie Mellon University  
`yxythu@gmail.com`

# MinBD: Minimally-Buffered Deflection Routing for Energy-Efficient Interconnect

Chris Fallin, Greg Nazario, Xiangyao Yu, Kevin Chang, Rachata Ausavarungnirun, and  
Onur Mutlu,

**["MinBD: Minimally-Buffered Deflection Routing for Energy-Efficient Interconnect"](#)**  
*Proceedings of the [6th ACM/IEEE International Symposium on Networks on Chip \(NOCS\)](#), Lyngby, Denmark, May 2012.* [Slides \(pptx\)](#) [\(pdf\)](#)

# Bufferless Deflection Routing

---

- **Key idea:** Packets are never buffered in the network. When two packets contend for the same link, one is **deflected**.
  - Removing **buffers** yields significant benefits
    - Reduces **power** (CHIPPER: reduces NoC power by 55%)
    - Reduces **die area** (CHIPPER: reduces NoC area by 36%)
  - But, at **high network utilization** (load), bufferless deflection routing causes **unnecessary link & router traversals**
    - Reduces network throughput and application performance
    - Increases dynamic power
  - **Goal:** Improve high-load performance of low-cost deflection networks by reducing the deflection rate.
-

# Outline: This Talk

---

- **Motivation**
- **Background:** Bufferless Deflection Routing
- **MinBD:** Reducing Deflections
  - Addressing Link Contention
  - Addressing the Ejection Bottleneck
  - Improving Deflection Arbitration
- **Results**
- **Conclusions**

# Outline: This Talk

---

- Motivation
- **Background:** Bufferless Deflection Routing
- **MinBD:** Reducing Deflections
  - Addressing Link Contention
  - Addressing the Ejection Bottleneck
  - Improving Deflection Arbitration
- Results
- Conclusions

# Issues in Bufferless Deflection Routing

---

- **Correctness:** Deliver all packets without **livelock**
  - **CHIPPER<sup>1</sup>: Golden Packet**
  - Globally prioritize one packet until delivered
- **Correctness:** Reassemble packets without **deadlock**
  - **CHIPPER<sup>1</sup>: Retransmit-Once**
- **Performance:** Avoid performance degradation at **high load**
  - **MinBD**

# Key Performance Issues

---

1. **Link contention:** no buffers to hold traffic → any link contention causes a deflection  
→ use side buffers
2. **Ejection bottleneck:** only one flit can eject per router per cycle → simultaneous arrival causes deflection  
→ eject up to 2 flits/cycle
3. **Deflection arbitration:** practical (fast) deflection arbiters deflect unnecessarily  
→ new priority scheme (silver flit)

# Outline: This Talk

---

- Motivation
- Background: Bufferless Deflection Routing
- **MinBD:** Reducing Deflections
  - Addressing Link Contention
  - Addressing the Ejection Bottleneck
  - Improving Deflection Arbitration
- Results
- Conclusions

# Outline: This Talk

---

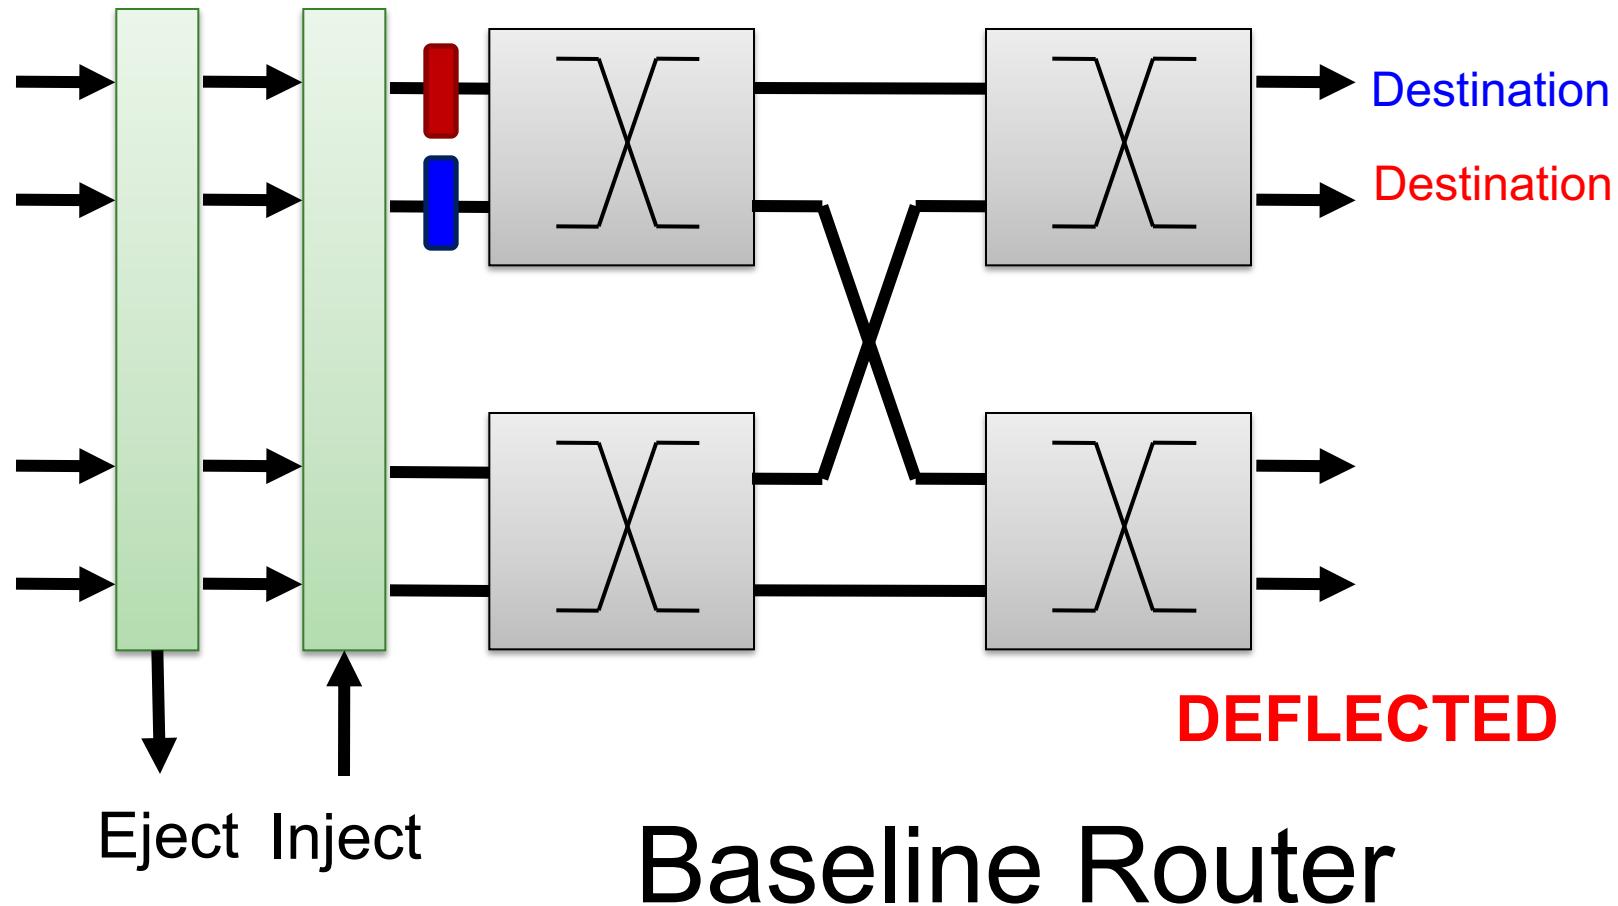
- Motivation
- Background: Bufferless Deflection Routing
- **MinBD:** Reducing Deflections
  - Addressing Link Contention
  - Addressing the Ejection Bottleneck
  - Improving Deflection Arbitration
- Results
- Conclusions

# Addressing Link Contention

---

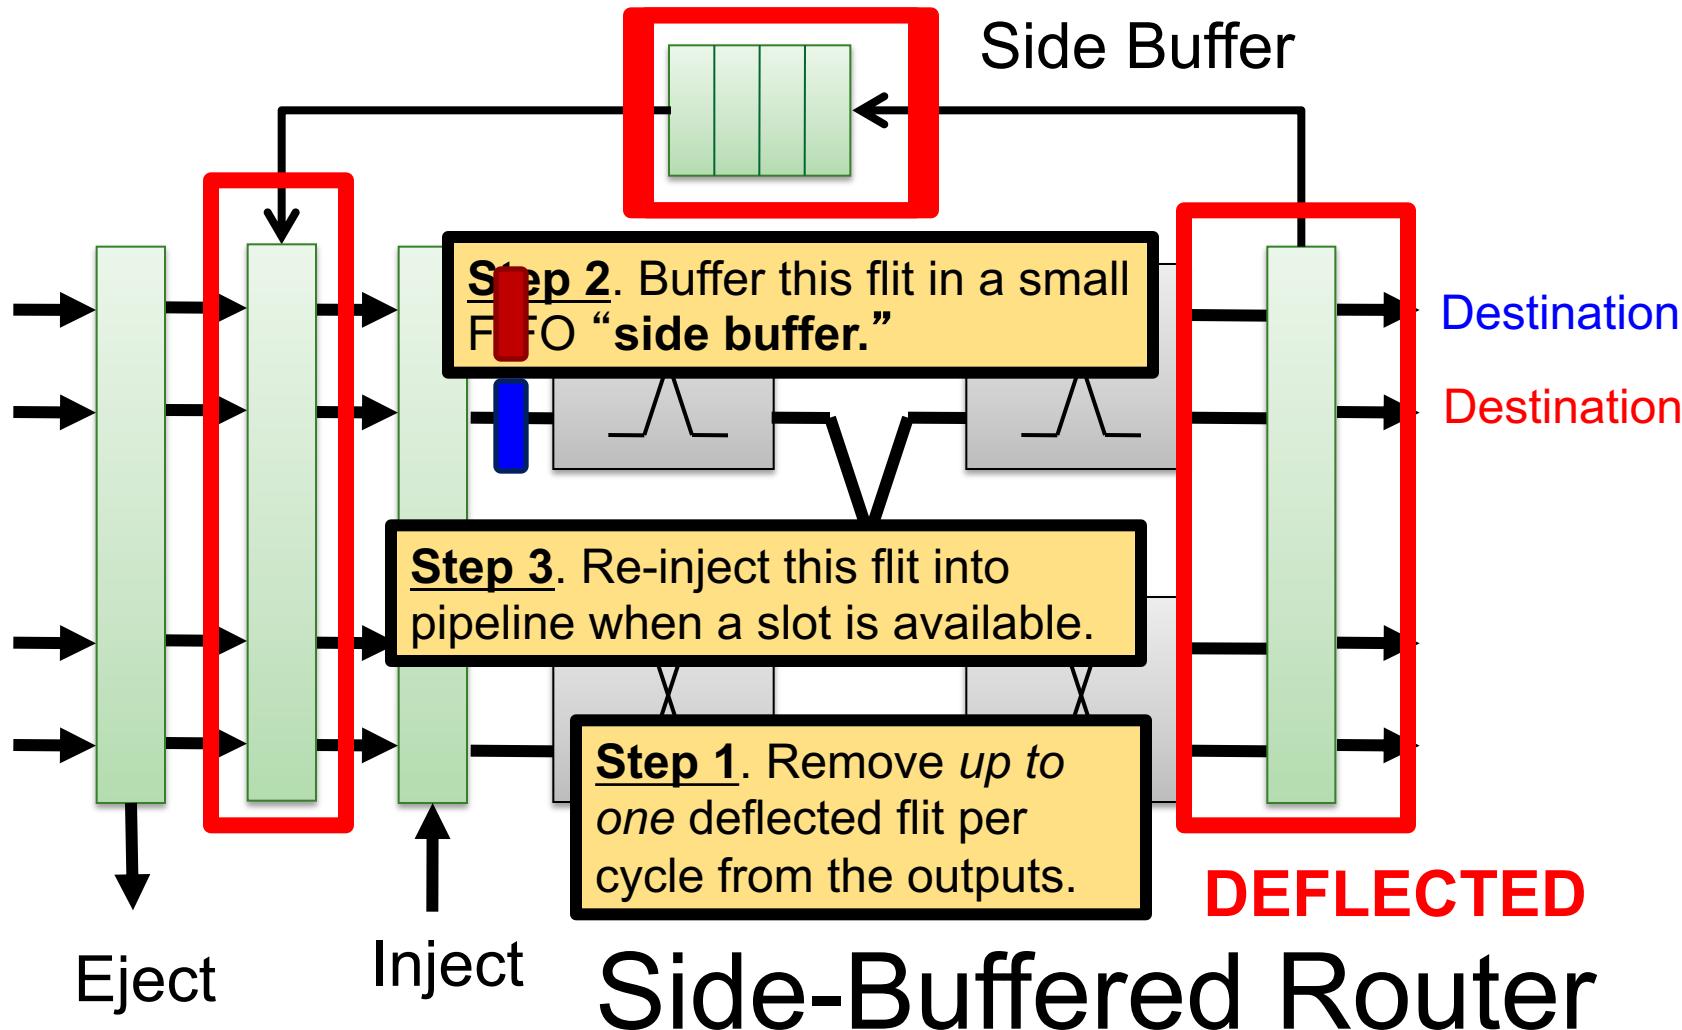
- **Problem 1:** Any link contention causes a deflection
- Buffering a flit can avoid deflection on contention
- But, **input buffers** are expensive:
  - All flits are buffered on every hop → high dynamic energy
  - Large buffers necessary → high static energy and large area
- **Key Idea 1:** add a small buffer to a bufferless deflection router to buffer **only** flits that would have been deflected

# How to Buffer Deflected Flits



<sup>1</sup> Fallin et al., “CHIPPER: A Low-complexity Bufferless Deflection Router”, HPCA 2011.

# How to Buffer Deflected Flits



# Why Could A Side Buffer Work Well?

---

- Buffer some flits and deflect other flits at per-flit level
  - Relative to **bufferless routers**, deflection rate reduces (need not deflect all contending flits)  
→ 4-flit buffer reduces deflection rate by 39%
  - Relative to **buffered routers**, buffer is more efficiently used (need not buffer all flits)  
→ similar performance with 25% of buffer space

# Outline: This Talk

---

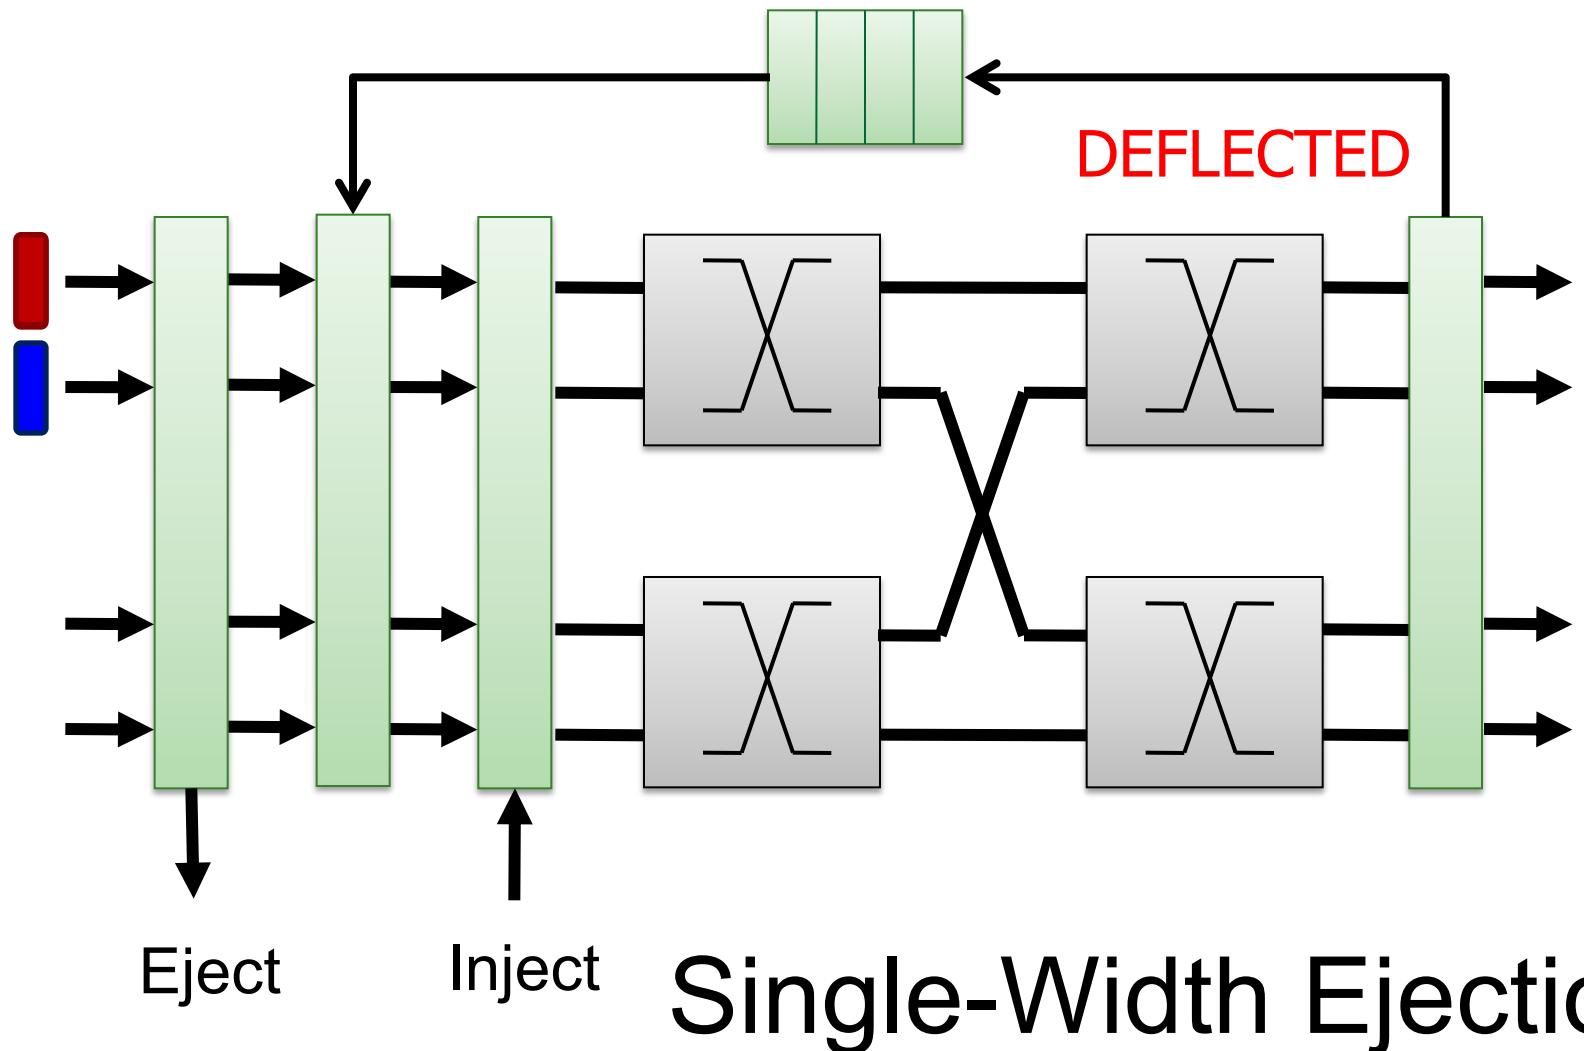
- Motivation
- Background: Bufferless Deflection Routing
- **MinBD:** Reducing Deflections
  - Addressing Link Contention
  - Addressing the Ejection Bottleneck
  - Improving Deflection Arbitration
- Results
- Conclusions

# Addressing the Ejection Bottleneck

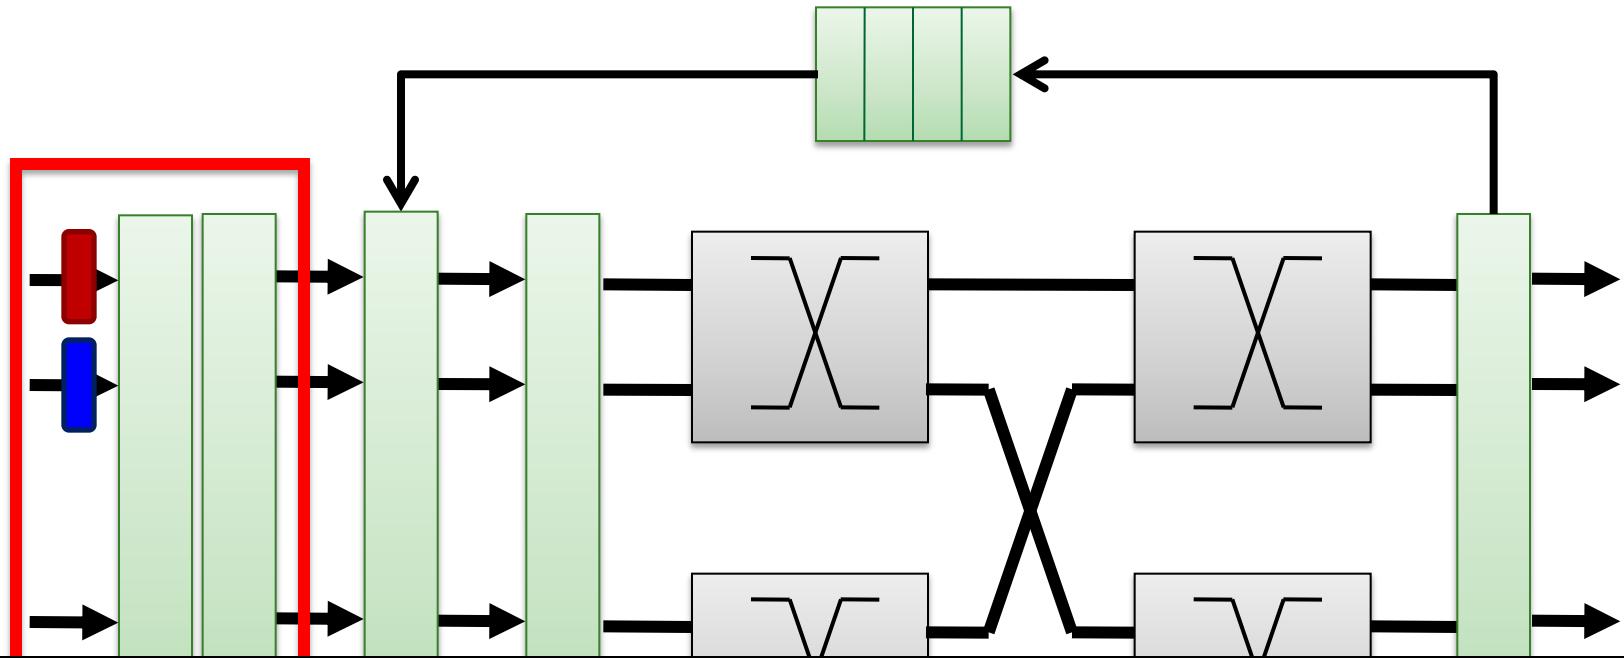
---

- **Problem 2:** Flits deflect unnecessarily because only one flit can **eject** per router per cycle
- In 20% of all ejections,  $\geq 2$  flits could have ejected
  - all but one flit must **deflect and try again**
  - these deflected flits cause additional contention
- Ejection width of 2 flits/cycle reduces **deflection rate 21%**
- **Key idea 2:** Reduce deflections due to a single-flit ejection port by allowing **two flits** to eject per cycle

# Addressing the Ejection Bottleneck



# Addressing the Ejection Bottleneck



For fair comparison, **baseline routers** have dual-width ejection for perf. (not power/area)

Eject

Inject

Dual-Width Ejection

# Outline: This Talk

---

- Motivation
- Background: Bufferless Deflection Routing
- **MinBD:** Reducing Deflections
  - Addressing Link Contention
  - Addressing the Ejection Bottleneck
  - Improving Deflection Arbitration
- Results
- Conclusions

# Improving Deflection Arbitration

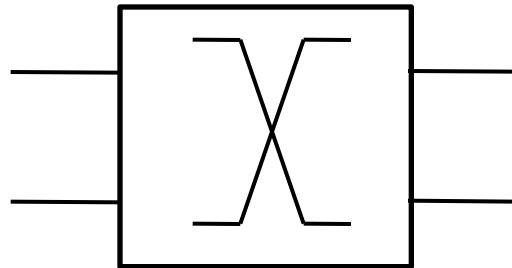
---

- **Problem 3:** Deflections occur unnecessarily because fast arbiters must use simple priority schemes
- Age-based priorities (several past works): full priority order gives fewer deflections, but requires slow arbiters
- State-of-the-art deflection arbitration (Golden Packet & two-stage permutation network)
  - Prioritize one packet globally (**ensure forward progress**)
  - Arbitrate other flits randomly (**fast critical path**)
- Random common case leads to uncoordinated arbitration

# Fast Deflection Routing Implementation

---

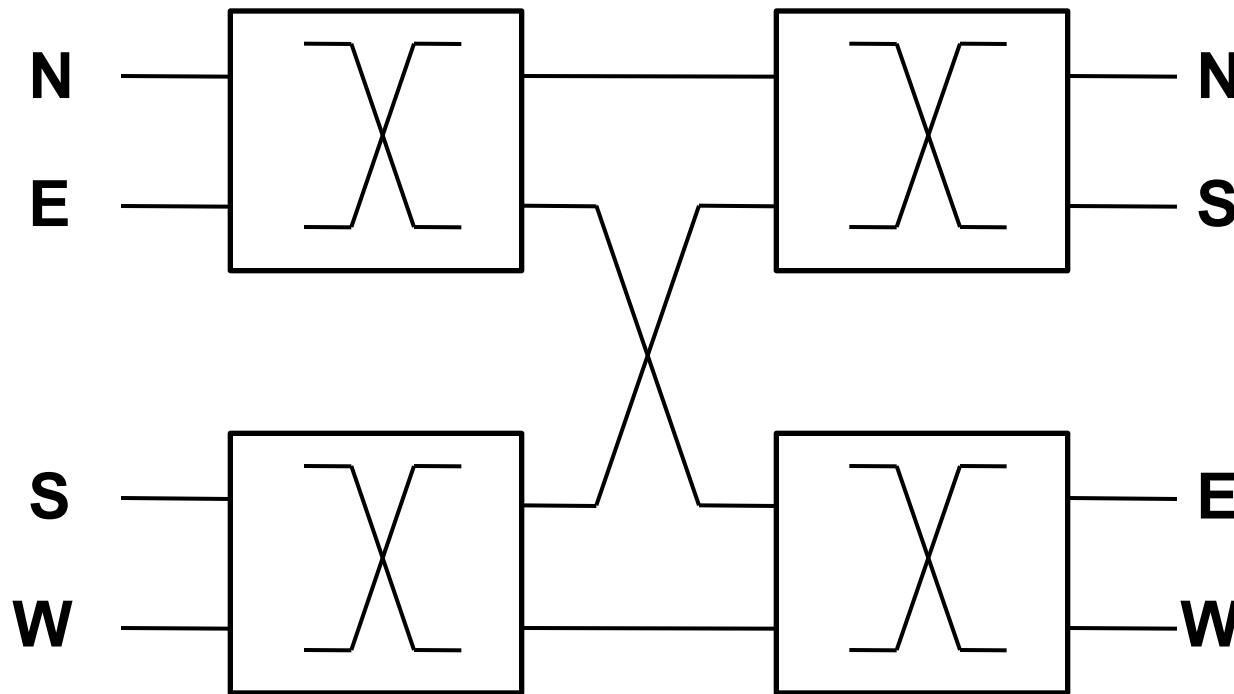
- Let's route in a two-input router first:



- Step 1:** pick a “winning” flit (Golden Packet, else random)
  - Step 2:** steer the winning flit to its desired output and deflect other flit
- **Highest-priority flit always routes to destination**

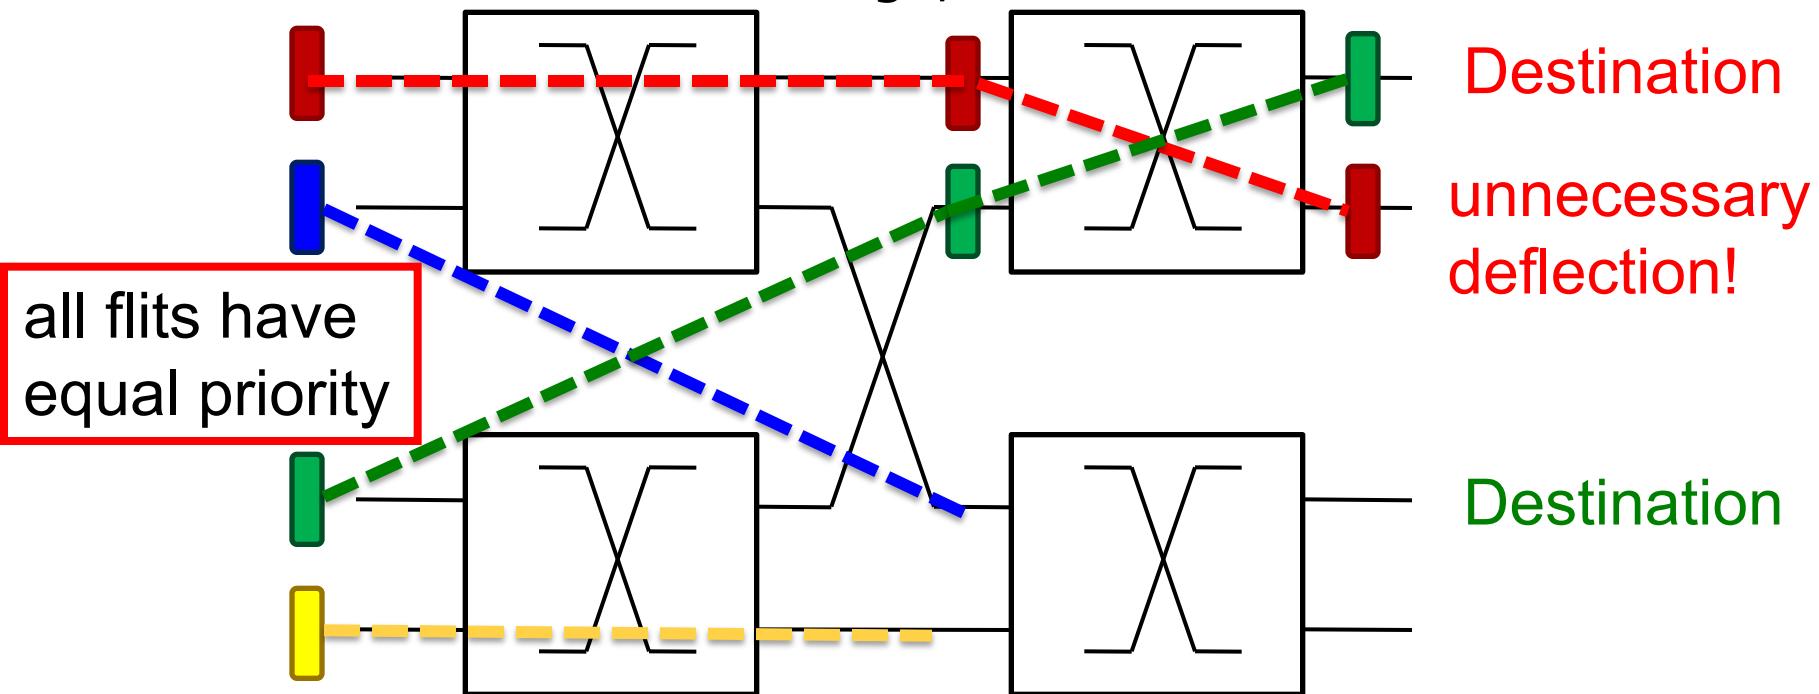
# Fast Deflection Routing with Four Inputs

- Each block makes decisions **independently**
  - **Deflection is a distributed decision**



# Unnecessary Deflections in Fast Arbiters

- How does lack of coordination cause unnecessary deflections?
  1. No flit is golden (pseudorandom arbitration)
  2. Red flit wins at first stage
  3. Green flit loses at first stage (must be deflected now)
  4. Red flit loses at second stage; Red and Green are deflected



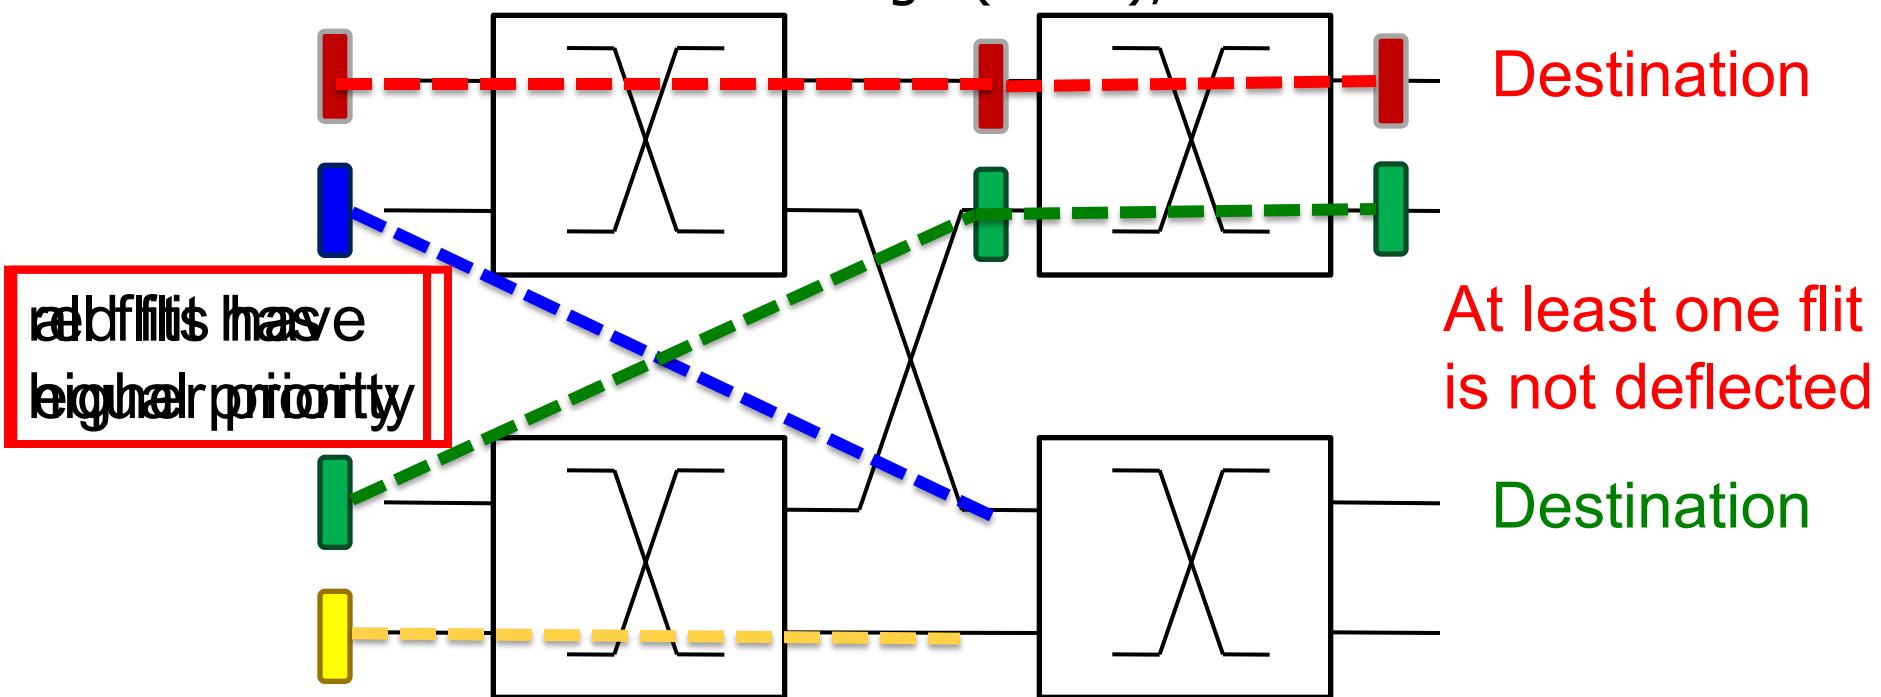
# Improving Deflection Arbitration

---

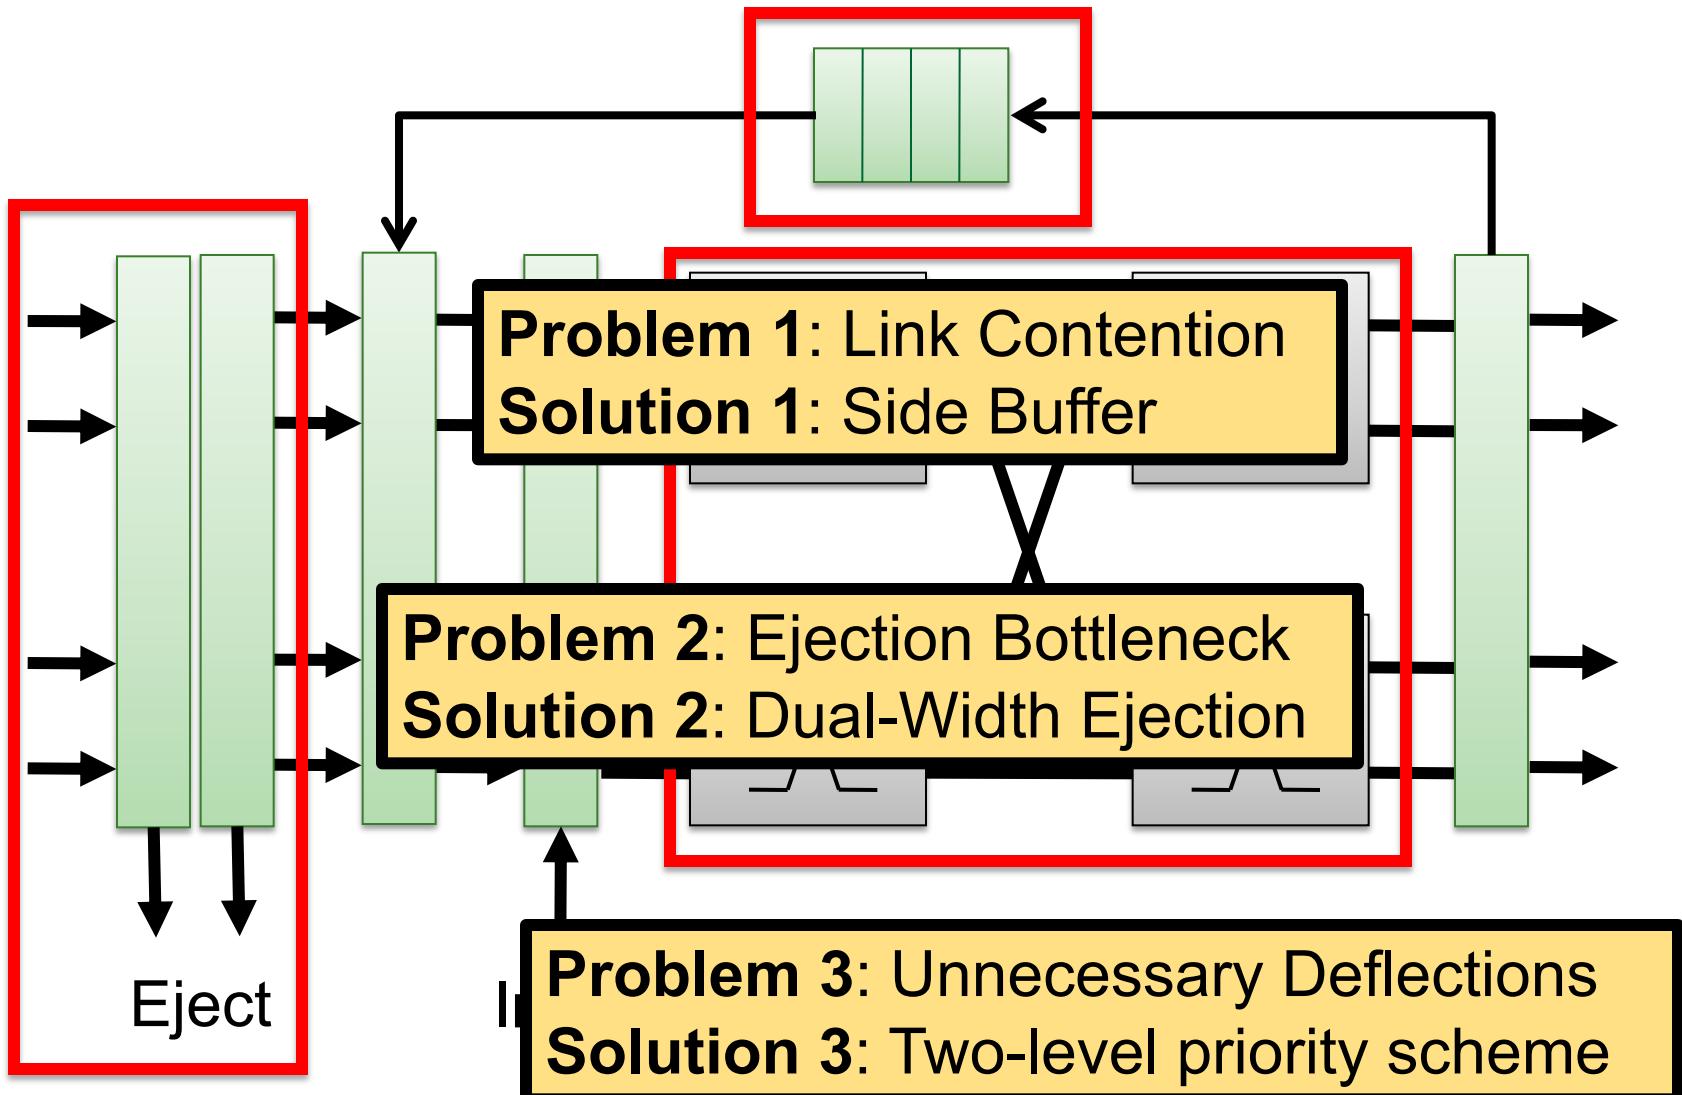
- **Key idea 3: Add a priority level** and prioritize one flit to ensure at least **one flit is not deflected in each cycle**
  
- **Highest priority:** one **Golden Packet** in network
  - Chosen in static round-robin schedule
  - Ensures correctness
  
- **Next-highest priority:** one **silver flit** per router per cycle
  - Chosen pseudo-randomly & local to one router
  - Enhances performance

# Adding A Silver Flit

- Randomly picking a silver flit ensures **one flit is not deflected**
  - No flit is golden but **Red** flit is silver
  - Red** flit wins at first stage (silver)
  - Green** flit is deflected at first stage
  - Red** flit wins at second stage (silver); not deflected



# Minimally-Buffered Deflection Router



# Outline: This Talk

---

- Motivation
- Background: Bufferless Deflection Routing
- **MinBD:** Reducing Deflections
  - Addressing Link Contention
  - Addressing the Ejection Bottleneck
  - Improving Deflection Arbitration

# Outline: This Talk

---

- **Motivation**
- **Background:** Bufferless Deflection Routing
- **MinBD:** Reducing Deflections
  - Addressing Link Contention
  - Addressing the Ejection Bottleneck
  - Improving Deflection Arbitration
- **Results**
- **Conclusions**

# Methodology: Simulated System

---

## ■ Chip Multiprocessor Simulation

- **64-core** and **16-core** models
- **Closed-loop** core/cache/NoC cycle-level model
- Directory cache coherence protocol (SGI Origin-based)
- 64KB L1, perfect L2 (stresses interconnect), XOR-mapping
- Performance metric: **Weighted Speedup**  
(similar conclusions from network-level latency)
- Workloads: multiprogrammed SPEC CPU2006
  - 75 randomly-chosen workloads
  - Binned into network-load categories by average injection rate

# Methodology: Routers and Network

---

- **Input-buffered** virtual-channel router
  - 8 VCs, 8 flits/VC [Buffered(8,8)]: large buffered router
  - 4 VCs, 4 flits/VC [Buffered(4,4)]: typical buffered router
  - 4 VCs, 1 flit/VC [Buffered(4,1)]: smallest deadlock-free router
  - All power-of-2 buffer sizes up to (8, 8) for perf/power sweep
- **Bufferless deflection** router: **CHIPPER**<sup>1</sup>
- **Bufferless-buffered hybrid** router: **AFC**<sup>2</sup>
  - Has input buffers and deflection routing logic
  - Performs coarse-grained (multi-cycle) mode switching
- **Common parameters**
  - 2-cycle router latency, 1-cycle link latency
  - 2D-mesh topology (16-node: 4x4; 64-node: 8x8)
  - Dual ejection assumed for baseline routers (for perf. only)

# Methodology: Power, Die Area, Crit. Path

---

## ■ Hardware modeling

- Verilog models for CHIPPER, MinBD, buffered control logic
  - Synthesized with commercial 65nm library
- ORION 2.0 for datapath: crossbar, muxes, buffers and links

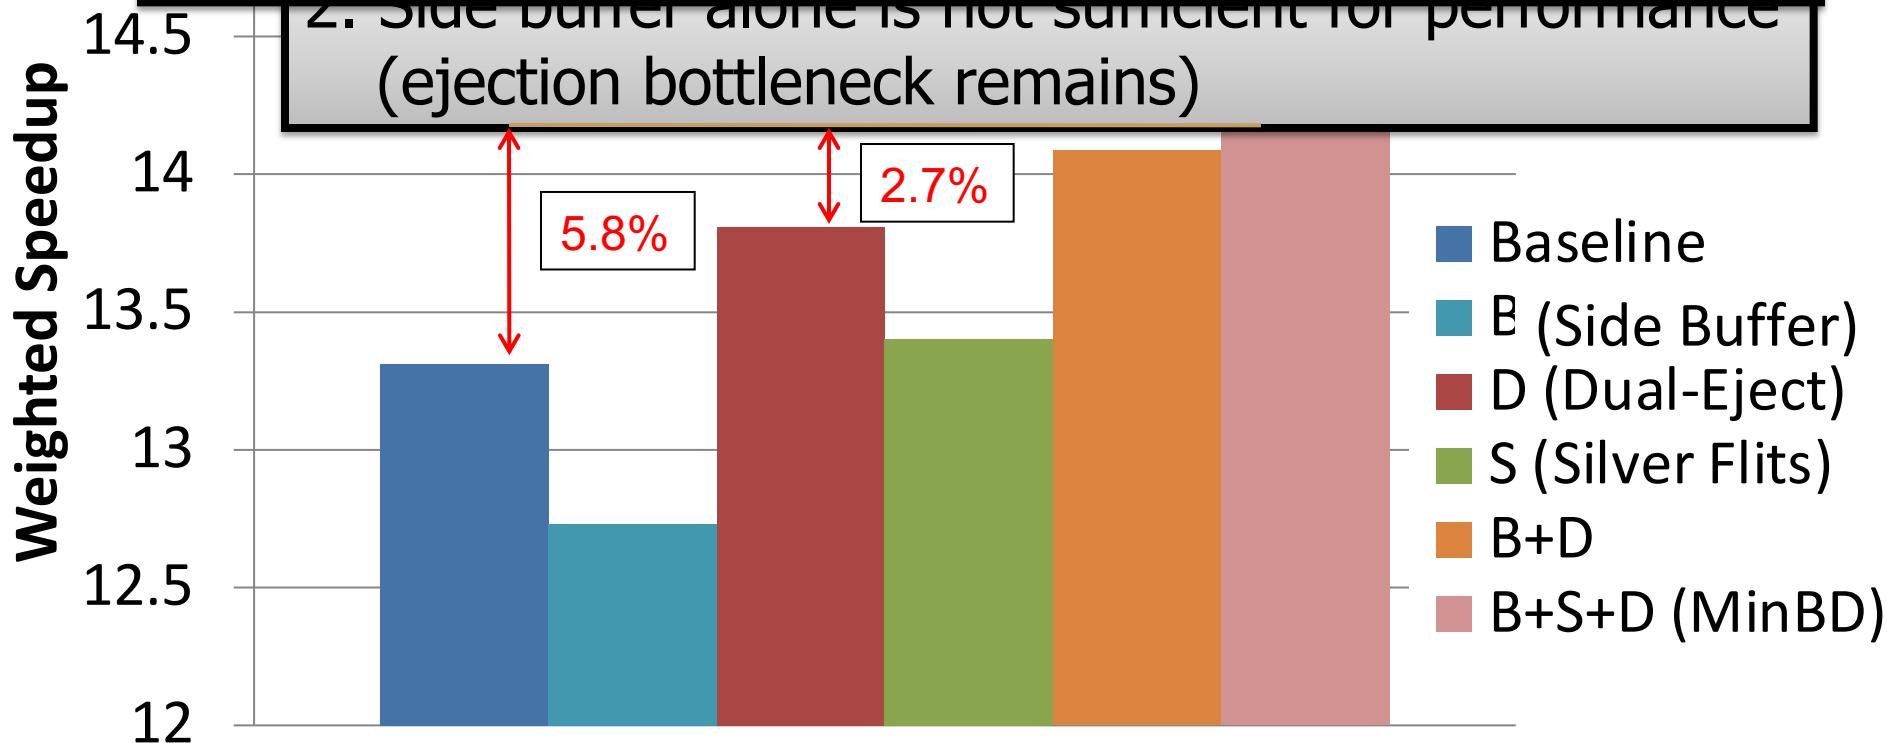
## ■ Power

- Static and dynamic power from hardware models
- Based on event counts in cycle-accurate simulations
- Broken down into buffer, link, other

# Reduced Deflections & Improved Perf.

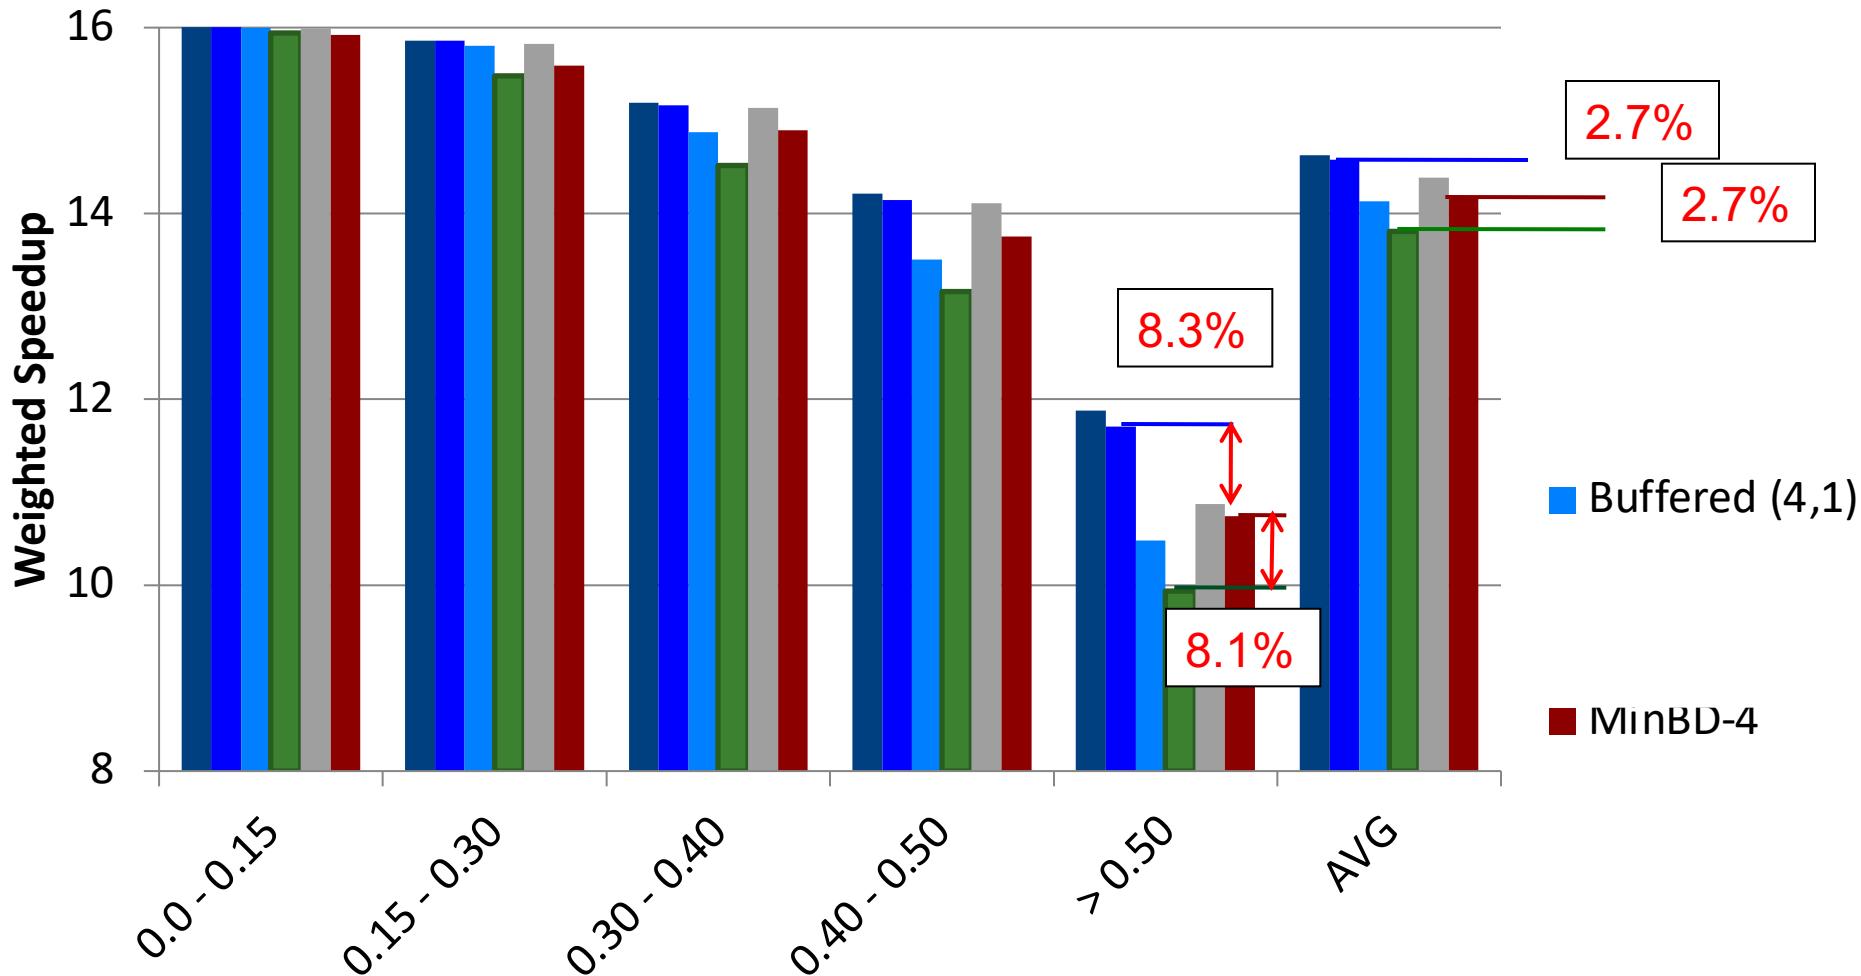
3. Overall, **5.8%** over baseline, **2.7%** over dual-eject by reducing deflections **64% / 54%**

2. Side buffer alone is not sufficient for performance (ejection bottleneck remains)



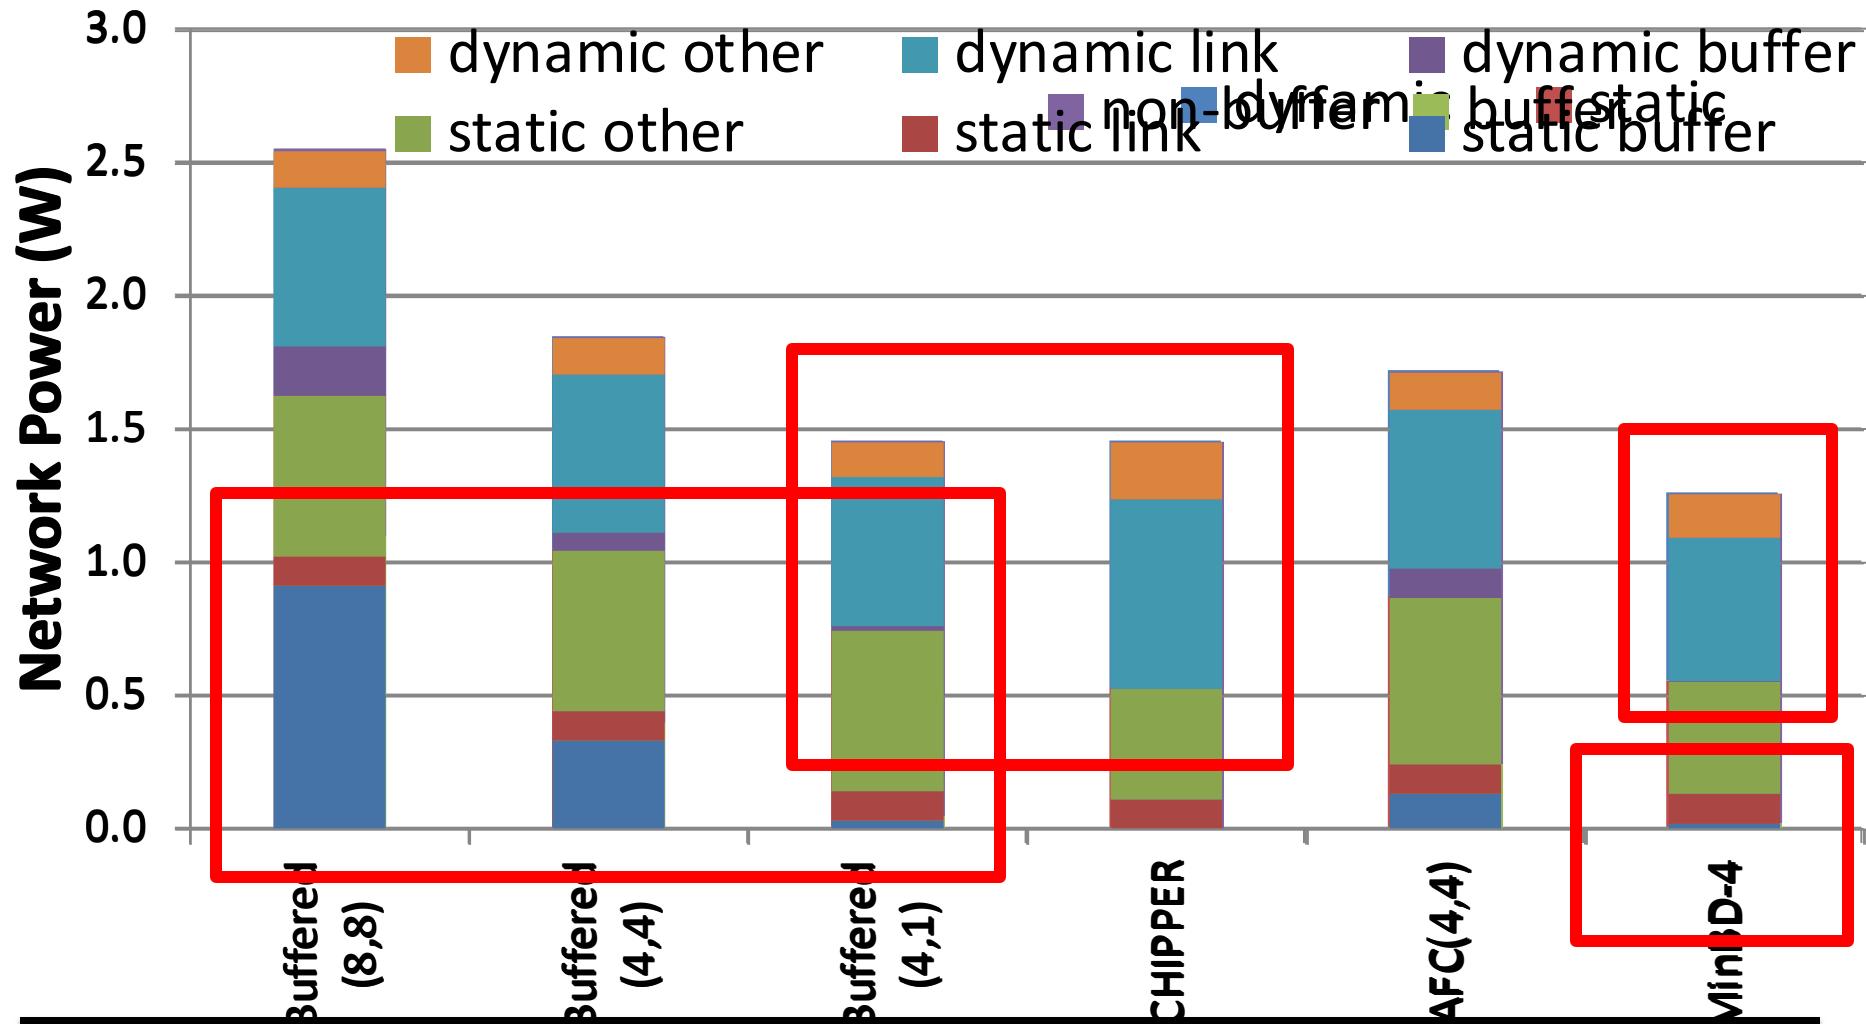
Deflection Rate	28%	17%	22%	27%	11%	10%

# Overall Performance Results



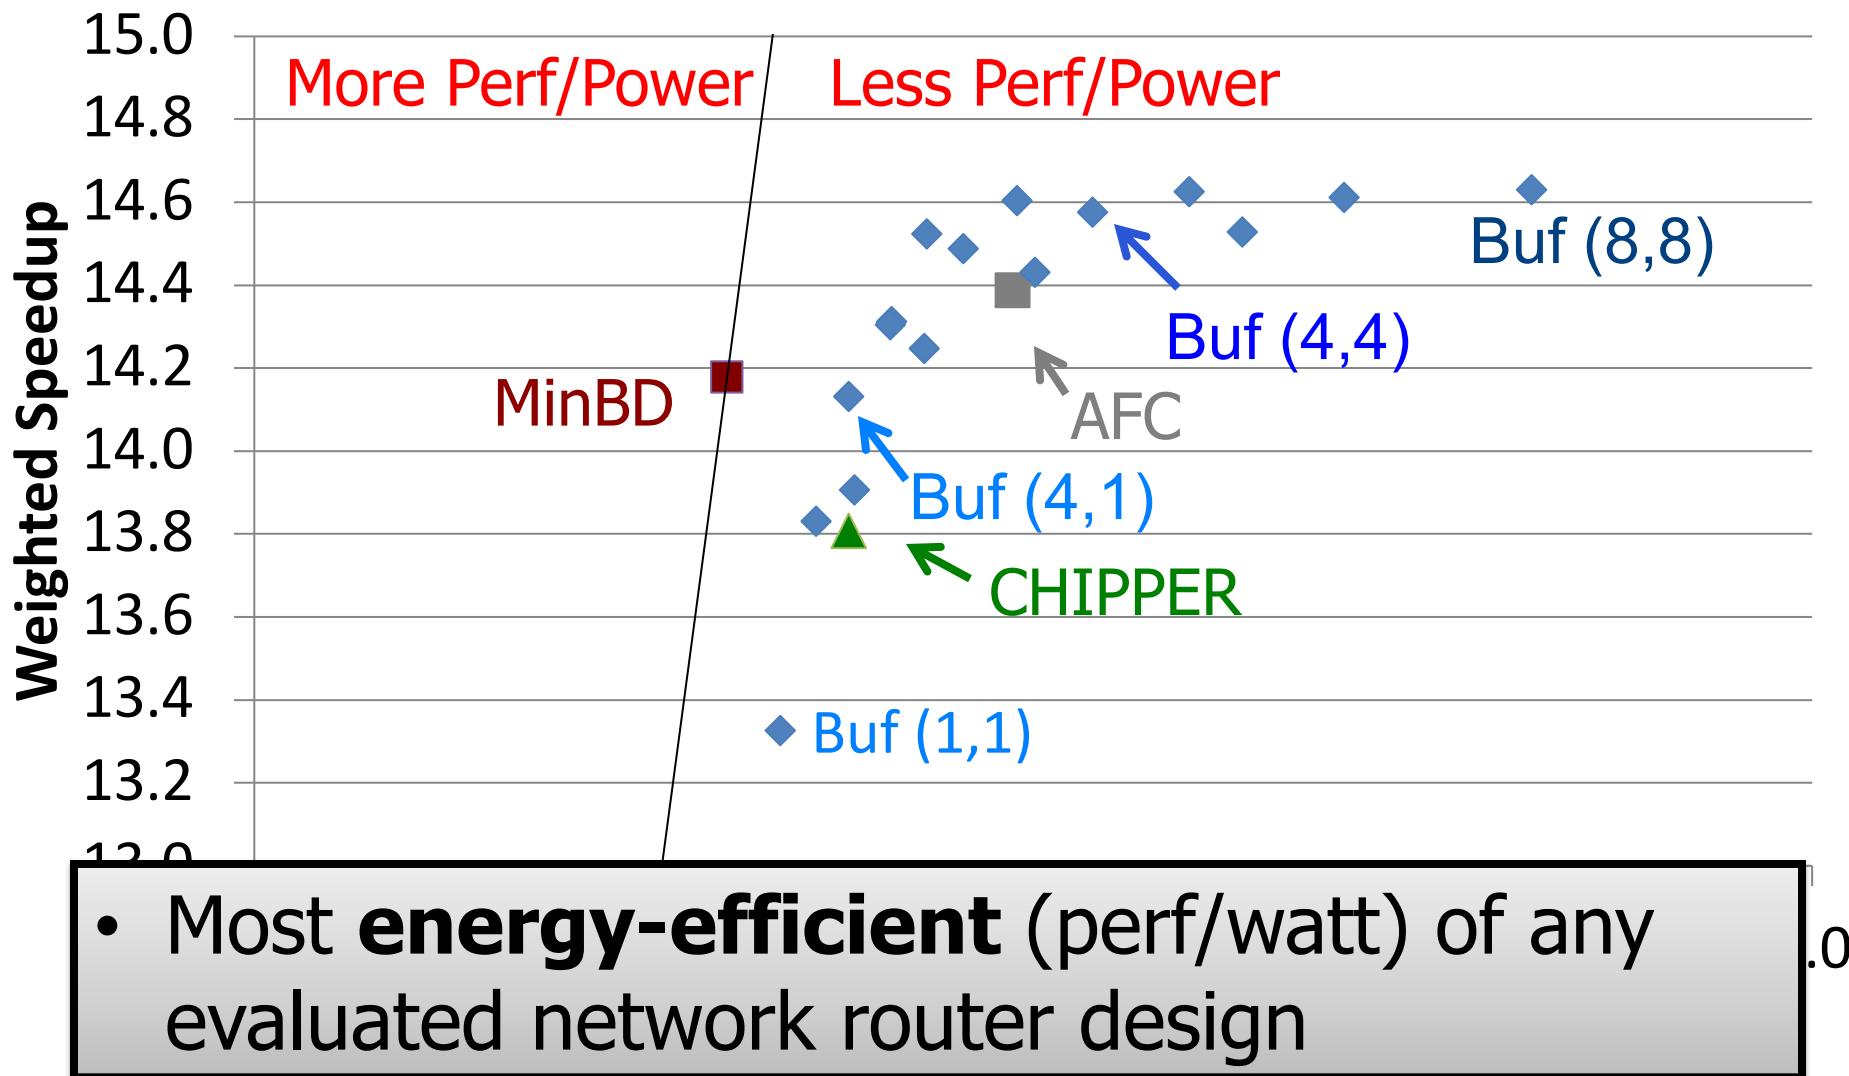
- Similar perf. to Buffered (4,1) @ 25% of buffering space
- Within **2.7%** of Buffered (4,4) (**8.3%** at high load)

# Overall Power Results

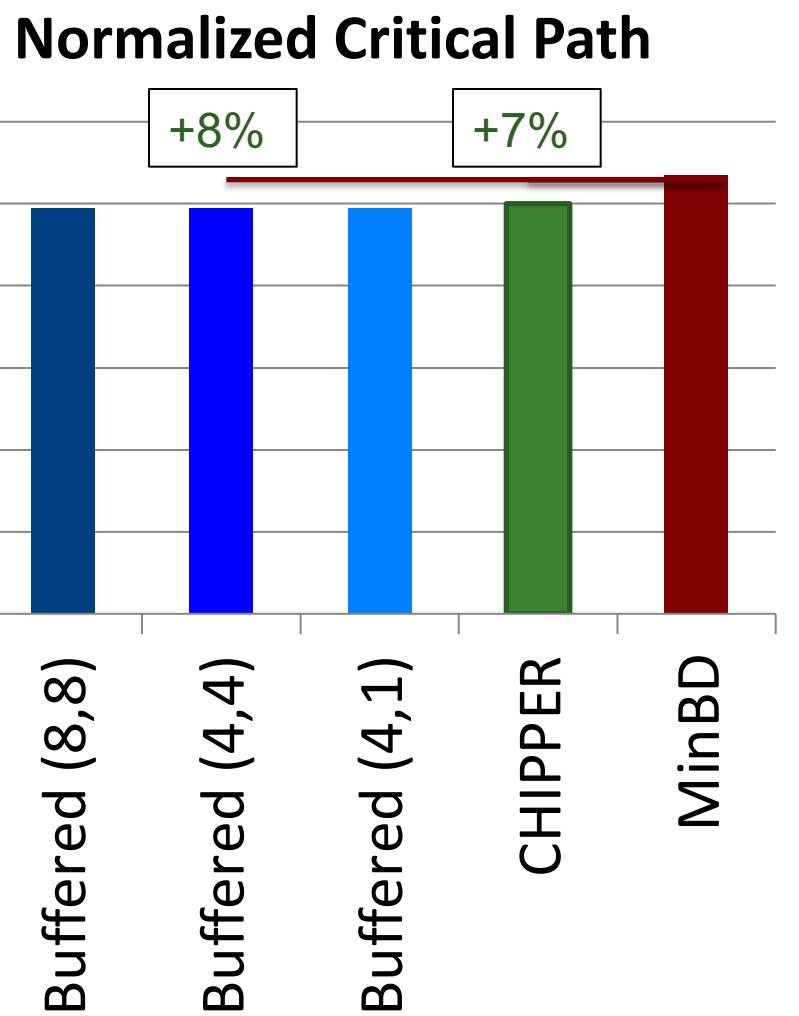
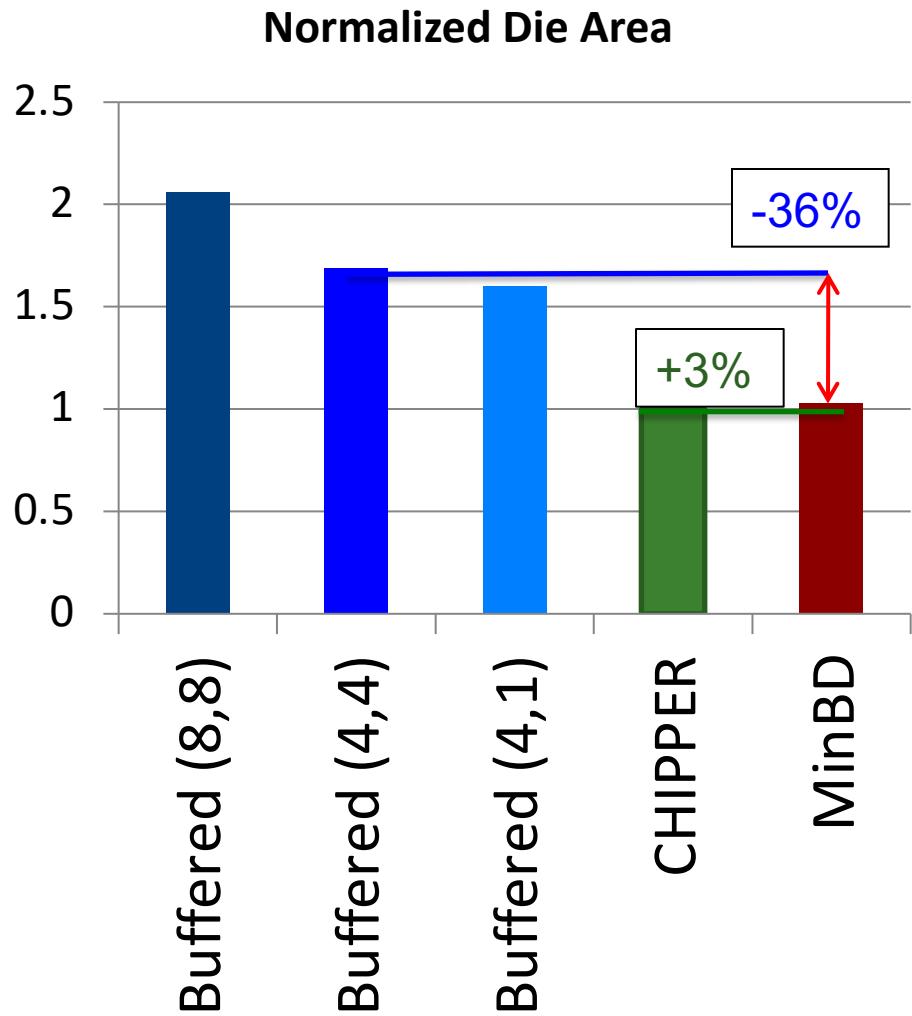


- Dynamic power increases with deflection routing
- Dynamic power reduces in MinBD relative to CHIPPER

# Performance-Power Spectrum



# Die Area and Critical Path



Only 3% area increase over CHIPPER (4-bit buffer)

- Increases by 7% over CHIPPER, 8% over Buffered (4,4)

# MinBD Router: Conclusions

---

- Bufferless deflection routing offers **reduced power & area**
  - But, high deflection rate hurts **performance at high load**
  - **MinBD** (Minimally-Buffered Deflection Router) introduces:
    - Side buffer to hold **only** flits that would have been deflected
    - Dual-width ejection to address ejection bottleneck
    - Two-level prioritization to avoid unnecessary deflections
  - MinBD yields **reduced power (31%) & reduced area (36%)** relative to **buffered** routers
  - MinBD yields **improved performance (8.1% at high load)** relative to **bufferless** routers → closes half of perf. gap
  - MinBD has the **best energy efficiency** of all evaluated designs with **competitive performance**
-

# Minimally-Buffered Deflection Routing

---

- Chris Fallin, Greg Nazario, Xiangyao Yu, Kevin Chang, Rachata Ausavarungnirun, and Onur Mutlu,

**"MinBD: Minimally-Buffered Deflection Routing for Energy-Efficient Interconnect"**

*Proceedings of the 6th ACM/IEEE International Symposium on Networks on Chip (NOCS), Lyngby, Denmark, May 2012.* [Slides \(pptx\)](#) [\(pdf\)](#)

***One of the five papers nominated for the Best Paper Award by the Program Committee.***

## MinBD: Minimally-Buffered Deflection Routing for Energy-Efficient Interconnect

Chris Fallin, Greg Nazario, Xiangyao Yu<sup>†</sup>, Kevin Chang, Rachata Ausavarungnirun, Onur Mutlu

Carnegie Mellon University  
`{cfallin,gnazario,kevincha,rachata,onur}@cmu.edu`

<sup>†</sup>Tsinghua University & Carnegie Mellon University  
`yxythu@gmail.com`

# HAT: Heterogeneous Adaptive Throttling for On-Chip Networks

Kevin Chang, Rachata Ausavarungnirun, Chris Fallin, and Onur Mutlu,  
"HAT: Heterogeneous Adaptive Throttling for On-Chip Networks"

*Proceedings of the 24th International Symposium on Computer Architecture and High Performance Computing (SBAC-PAD)*, New York, NY, October 2012. Slides (pptx) (pdf)

Carnegie Mellon University

**SAFARI**

# Executive Summary

- **Problem:** Packets contend in on-chip networks (NoCs), causing congestion, thus reducing performance
- **Observations:**
  - 1) Some applications are more sensitive to network latency than others
  - 2) Applications must be throttled differently to achieve peak performance
- **Key Idea: Heterogeneous Adaptive Throttling (HAT)**
  - 1) Application-aware source throttling
  - 2) Network-load-aware throttling rate adjustment
- **Result:** Improves performance and energy efficiency over state-of-the-art source throttling policies

# Source Throttling in Bufferless NoCs

---

- Kevin Chang, Rachata Ausavarungnirun, Chris Fallin, and Onur Mutlu,  
**"HAT: Heterogeneous Adaptive Throttling for On-Chip Networks"**  
*Proceedings of the 24th International Symposium on Computer Architecture  
and High Performance Computing (SBAC-PAD)*, New York, NY, October  
2012. [Slides \(pptx\)](#) [\(pdf\)](#)

## HAT: Heterogeneous Adaptive Throttling for On-Chip Networks

Kevin Kai-Wei Chang, Rachata Ausavarungnirun, Chris Fallin, Onur Mutlu  
Carnegie Mellon University  
`{kevincha, rachata, cfallin, onur}@cmu.edu`

# “Bufferless” Hierarchical Rings

---

- Rachata Ausavarungnirun, Chris Fallin, Xiangyao Yu, Kevin Chang, Greg Nazario, Reetuparna Das, Gabriel Loh, and Onur Mutlu,  
**["Design and Evaluation of Hierarchical Rings with Deflection Routing"](#)**  
*Proceedings of the [26th International Symposium on Computer Architecture and High Performance Computing \(SBAC-PAD\)](#)*, Paris, France, October 2014. [[Slides \(pptx\)](#)] [[\(pdf\)](#)] [[Source Code](#)]
- Describes the design and implementation of a mostly-bufferless hierarchical ring

## Design and Evaluation of Hierarchical Rings with Deflection Routing

Rachata Ausavarungnirun    Chris Fallin    Xiangyao Yu†    Kevin Kai-Wei Chang  
Greg Nazario    Reetuparna Das§    Gabriel H. Loh‡    Onur Mutlu

Carnegie Mellon University    §University of Michigan    †MIT    ‡Advanced Micro Devices, Inc.

---

# “Bufferless” Hierarchical Rings (II)

---

- Rachata Ausavarungnirun, Chris Fallin, Xiangyao Yu, Kevin Chang, Greg Nazario, Reetuparna Das, Gabriel Loh, and Onur Mutlu,  
**"A Case for Hierarchical Rings with Deflection Routing: An Energy-Efficient On-Chip Communication Substrate"**  
*Parallel Computing (PARCO)*, 2016.
  - [arXiv.org version](#), February 2016.

Achieving both High Energy Efficiency  
and High Performance in On-Chip Communication  
using Hierarchical Rings with Deflection Routing

Rachata Ausavarungnirun    Chris Fallin    Xiangyao Yu†    Kevin Kai-Wei Chang  
Greg Nazario    Reetuparna Das§    Gabriel H. Loh‡    Onur Mutlu  
Carnegie Mellon University    §University of Michigan    †MIT    ‡AMD

# A Review of Bufferless Interconnects

---

- Chris Fallin, Greg Nazario, Xiangyao Yu, Kevin Chang, Rachata Ausavarungnirun, and Onur Mutlu,  
**"Bufferless and Minimally-Buffered Deflection Routing"**  
*Invited Book Chapter in Routing Algorithms in Networks-on-Chip, pp. 241-275, Springer, 2014.*

## Chapter 1 Bufferless and Minimally-Buffered Deflection Routing

Chris Fallin, Greg Nazario, Xiangyao Yu, Kevin Chang, Rachata Ausavarungnirun, Onur Mutlu

# Summary of Eight Years of Research

---

## Energy-Efficient Deflection-based On-chip Networks: Topology, Routing, Flow Control

Rachata Ausavarungnirun<sup>b</sup>, Onur Mutlu<sup>a</sup>

*SAFARI Research Group*

<sup>a</sup>*ETH Zürich*

<sup>b</sup>*King Mongkut's University of Technology North Bangkok*

---

### Abstract

As the number of cores scales to tens and hundreds, the energy consumption of routers across various types of on-chip networks in chip multiprocessors (CMPs) increases significantly. A major source of this energy consumption comes from the input buffers inside Network-on-Chip (NoC) routers, which are traditionally designed to maximize performance. To mitigate this high energy cost, many works propose bufferless router designs that utilize deflection routing to resolve port contention. While this approach is able to maintain high performance relative to its buffered counterparts at low network traffic, the bufferless router design suffers performance degradation under high network load.

In order to maintain high performance and energy efficiency under *both* low and high network loads, this chapter discusses critical drawbacks of traditional bufferless designs and describes recent research works focusing on two major modifications to improve the overall performance of the traditional bufferless network-on-chip design. The first modification is a minimally-buffered design that introduces limited buffering inside critical parts of the on-chip network in order to reduce the number of deflections. The second modification is a hierarchical bufferless interconnect design that aims to further improve performance by limiting the number of hops each packet needs to travel while in the network. In both approaches, we discuss design tradeoffs and provide evaluation results based on common CMP configurations with various network topologies to show the effectiveness of each proposal.

*Keywords:* network-on-chip, deflection routing, topology, bufferless router, energy efficiency, high-performance computing, computer architecture, emerging technologies, latency, low-latency computing

---

# Bufferless Interconnects in Real Systems

---

## Application Defined On-chip Networks for Heterogeneous Chiplets: An Implementation Perspective

Tianqi Wang<sup>1,\*</sup>, Fan Feng<sup>1,\*</sup>, Shaolin Xiang<sup>1,\*</sup>, Qi Li<sup>1</sup>, and Jing Xia<sup>1,\*\*</sup>

<sup>1</sup>*Huawei*

THEME ARTICLE: COMMERCIAL PRODUCTS 2021

## Kunpeng 920: The First 7-nm Chiplet-Based 64-Core ARM SoC for Cloud Services

Jing Xia, Chuanning Cheng, Xiping Zhou, Yuxing Hu , and Peter Chun, *HiSilicon Technologies Company, Ltd., Shenzhen, 518129, China*

# More Readings

---

- Studies of congestion and congestion control in on-chip vs. internet-like networks
- George Nychis, Chris Fallin, Thomas Moscibroda, Onur Mutlu, and Srinivasan Seshan,  
**"On-Chip Networks from a Networking Perspective: Congestion and Scalability in Many-core Interconnects"**  
*Proceedings of the 2012 ACM SIGCOMM Conference (SIGCOMM), Helsinki, Finland, August 2012.* Slides (pptx)
- George Nychis, Chris Fallin, Thomas Moscibroda, and Onur Mutlu,  
**"Next Generation On-Chip Networks: What Kind of Congestion Control Do We Need?"**  
*Proceedings of the 9th ACM Workshop on Hot Topics in Networks (HOTNETS), Monterey, CA, October 2010.* Slides (ppt) (key)

# On-Chip vs. Off-Chip Congestion Control

---

- George Nychis, Chris Fallin, Thomas Moscibroda, Onur Mutlu, and Srinivasan Seshan,

**"On-Chip Networks from a Networking Perspective:  
Congestion and Scalability in Many-core Interconnects"**

*Proceedings of the 2012 ACM SIGCOMM*

*Conference (**SIGCOMM**)*, Helsinki, Finland, August 2012. Slides  
(pptx)

## **On-Chip Networks from a Networking Perspective: Congestion and Scalability in Many-Core Interconnects**

George Nychis<sup>†</sup>, Chris Fallin<sup>†</sup>, Thomas Moscibroda<sup>§</sup>, Onur Mutlu<sup>†</sup>, Srinivasan Seshan<sup>†</sup>

<sup>†</sup> Carnegie Mellon University  
[{gnychis,cfallin,onur,srini}@cmu.edu](mailto:{gnychis,cfallin,onur,srini}@cmu.edu)

<sup>§</sup> Microsoft Research Asia  
[moscitho@microsoft.com](mailto:moscitho@microsoft.com)

# On-Chip vs. Off-Chip Congestion Control

---

- George Nychis, Chris Fallin, Thomas Moscibroda, and Onur Mutlu,  
**"Next Generation On-Chip Networks: What Kind of Congestion Control Do We Need?"**

*Proceedings of the 9th ACM Workshop on Hot Topics in Networks (HOTNETS), Monterey, CA, October 2010.* Slides (ppt) (key)

## Next Generation On-Chip Networks: What Kind of Congestion Control Do We Need?

George Nychis†, Chris Fallin†, Thomas Moscibroda§, Onur Mutlu†

†Carnegie Mellon University

{gnychis,cfallin,onur}@cmu.edu

§Microsoft Research  
moscitho@microsoft.com

# Summary of Study [SIGCOMM 2012]

---

- Highlighted a traditional networking problem in a new context
  - Unique design requires novel solution
- Showed **congestion limits efficiency and scalability**, and that **self-throttling nature of cores prevents congestion collapse**
- Showed **on-chip congestion control requires application-awareness**
- Our **application-aware congestion controller** provided:
  - A more efficient network-layer (reduced latency)
  - Improvements in system throughput (by 27%)
  - Effectively scale the CMP (shown for up to 4096 cores)

# Heterogeneous Networks

---

- Asit K. Mishra, Onur Mutlu, and Chita R. Das,  
**"A Heterogeneous Multiple Network-on-Chip Design: An Application-Aware Approach"**  
*Proceedings of the 50th Design Automation Conference (DAC),*  
Austin, TX, June 2013. [Slides \(pptx\)](#) [Slides \(pdf\)](#)

## A Heterogeneous Multiple Network-On-Chip Design: An Application-Aware Approach

Asit K. Mishra  
Intel Corporation  
Hillsboro, OR 97124, USA  
asit.k.mishra@intel.com

Onur Mutlu  
Carnegie Mellon University  
Pittsburgh, PA 15213, USA  
onur@cmu.edu

Chita R. Das  
The Pennsylvania State University  
University Park, PA 16802, USA  
das@cse.psu.edu

# Packet Scheduling

# Packet Scheduling

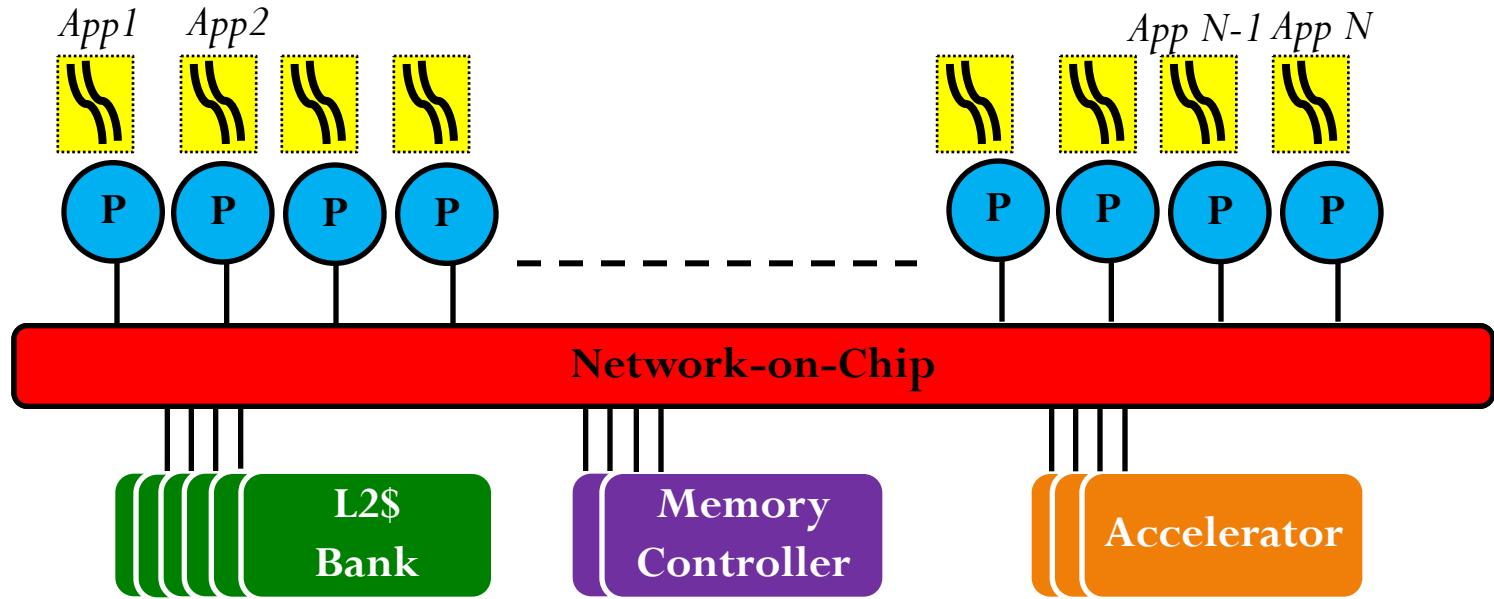
---

- Which packet to choose for a given output port?
  - Router needs to prioritize between competing flits
  - Which input port?
  - Which virtual channel?
  - Which application's packet?
- Common strategies
  - Round robin across virtual channels
  - Oldest packet first (or an approximation)
  - Prioritize some virtual channels over others
- Better policies in a multi-core environment
  - Use application characteristics

# Application-Aware Packet Scheduling

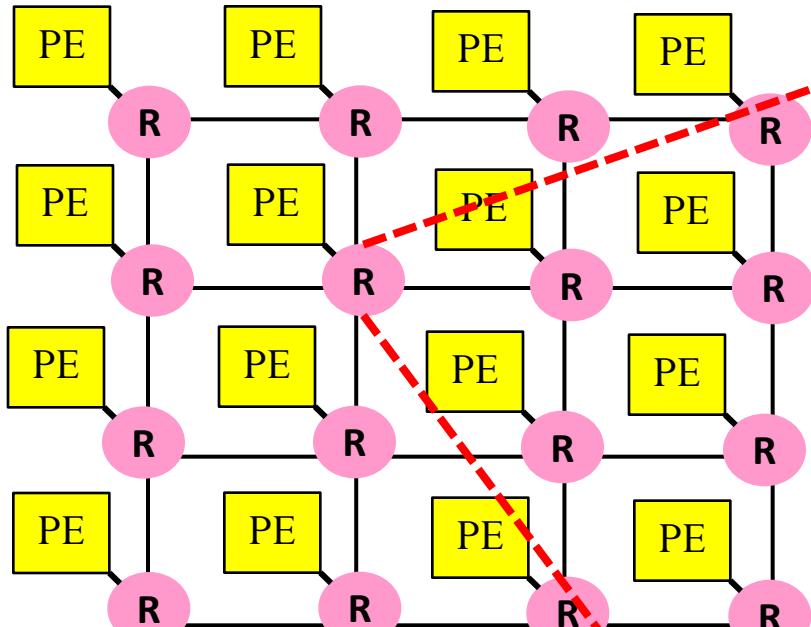
Das et al., “[Application-Aware Prioritization Mechanisms for On-Chip Networks](#),” MICRO 2009.

# The Problem: Packet Scheduling



Network-on-Chip is a **critical resource**  
shared by **multiple applications**

# The Problem: Packet Scheduling



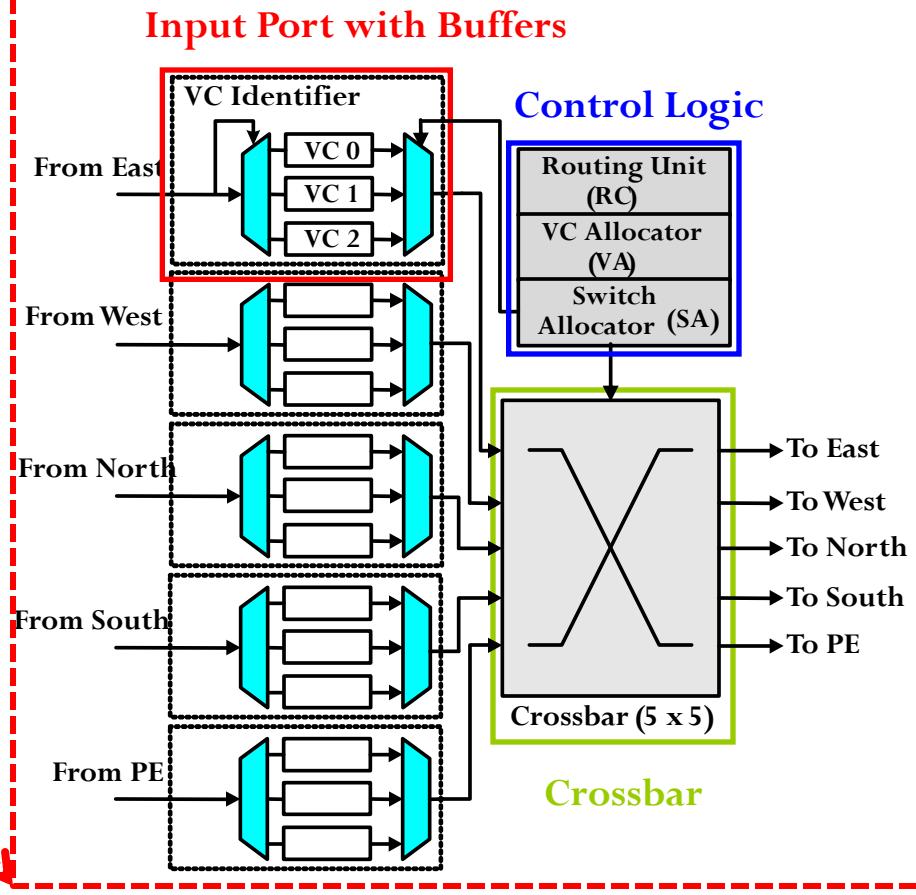
R

Routers

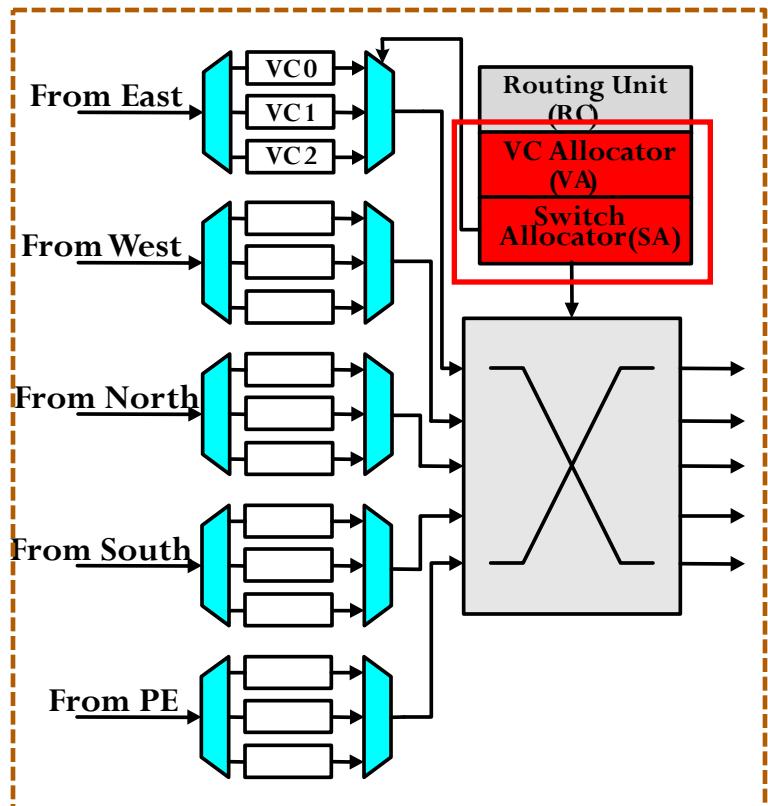
PE

Processing Element

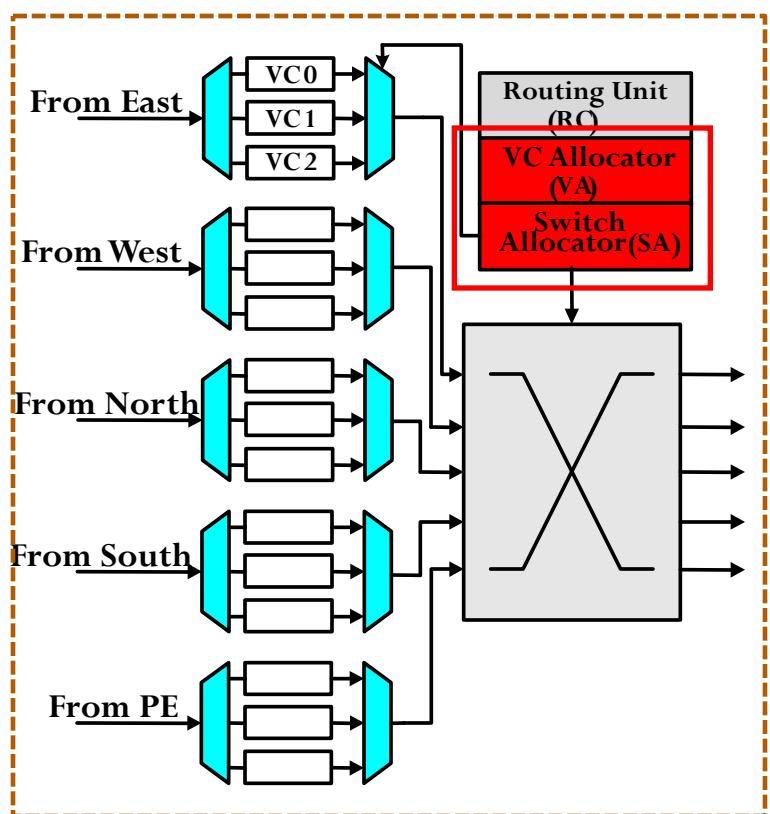
(Cores, L2 Banks, Memory Controllers etc)



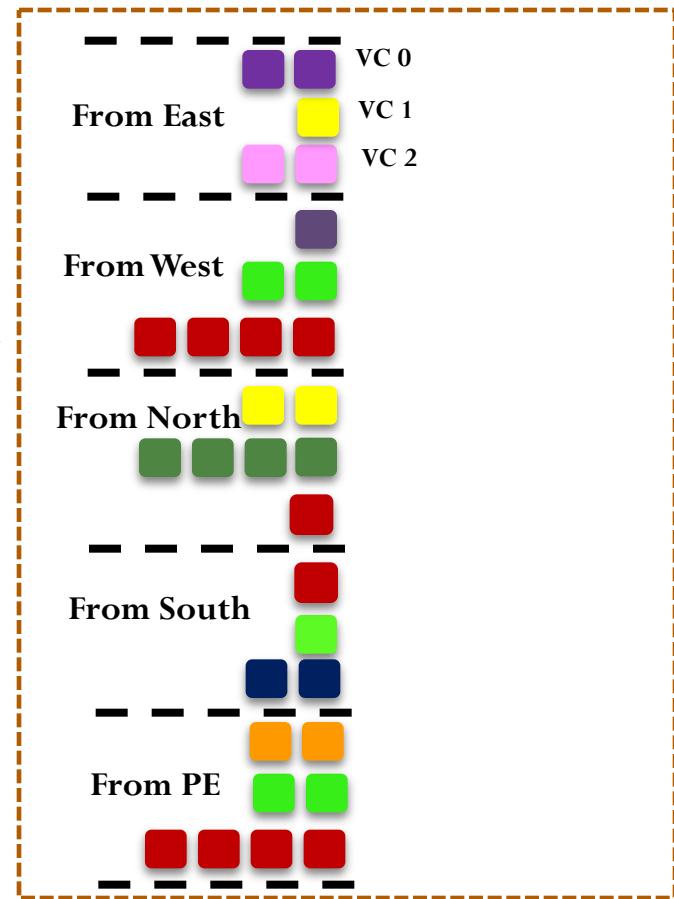
# The Problem: Packet Scheduling



# The Problem: Packet Scheduling

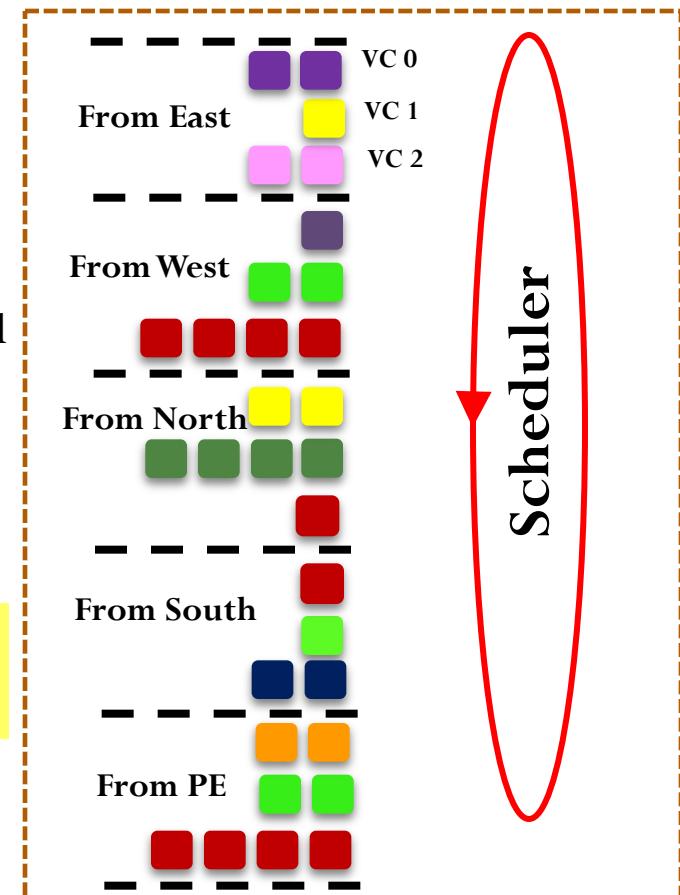
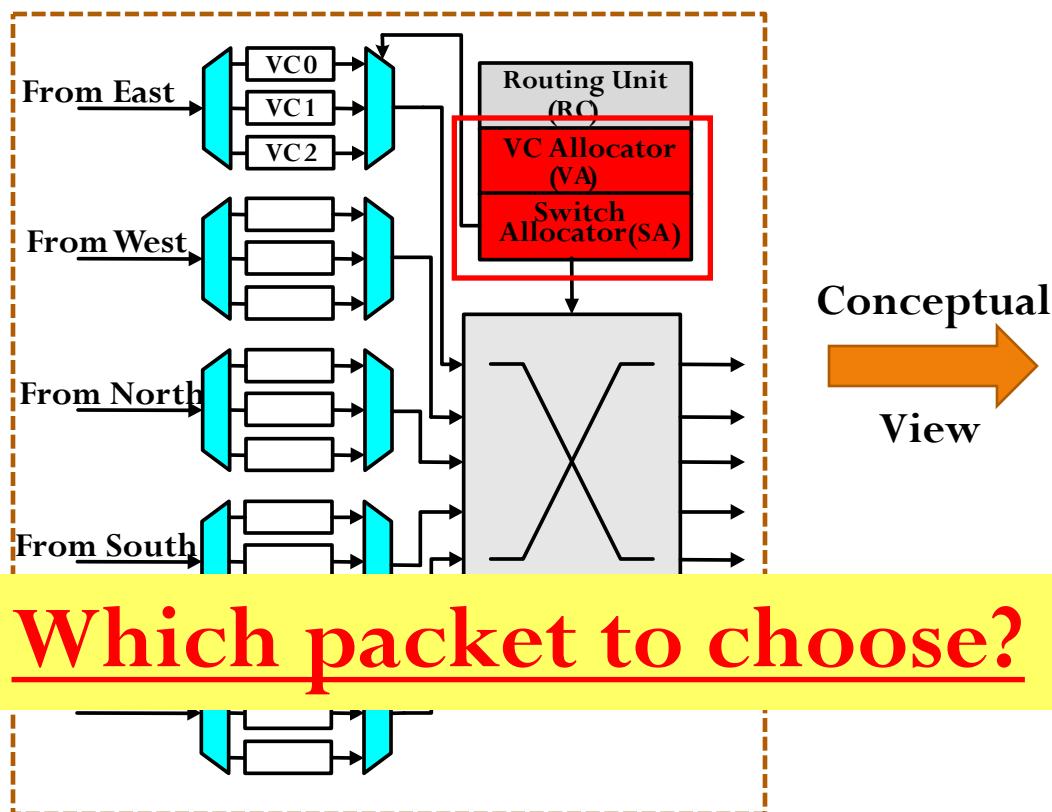


Conceptual  
View



Legend:  
VC 0: App1  
VC 1: App2  
VC 2: App3  
VC 3: App4  
VC 4: App5  
VC 5: App6  
VC 6: App7  
VC 7: App8

# The Problem: Packet Scheduling



**Legend:**

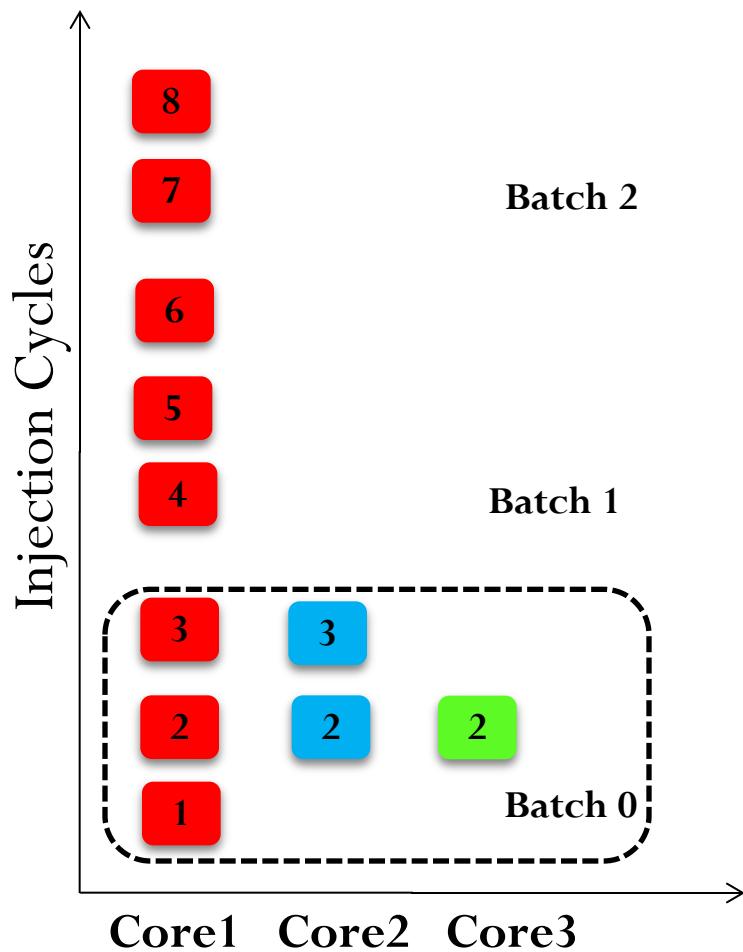
VC 0	VC 1	VC 2	App1	App2	App3	App4
App5	App6	App7	App8	App1	App2	App3

# The Problem: Packet Scheduling

---

- Existing scheduling policies
  - Round Robin
  - Age
- Problem 1: **Local** to a router
  - Lead to **contradictory decision making between routers**: packets from one application may be prioritized at one router, to be delayed at next.
- Problem 2: **Application oblivious**
  - Treat all applications **packets equally**
  - But **applications are heterogeneous**
- **Solution** : Application-aware global scheduling policies.

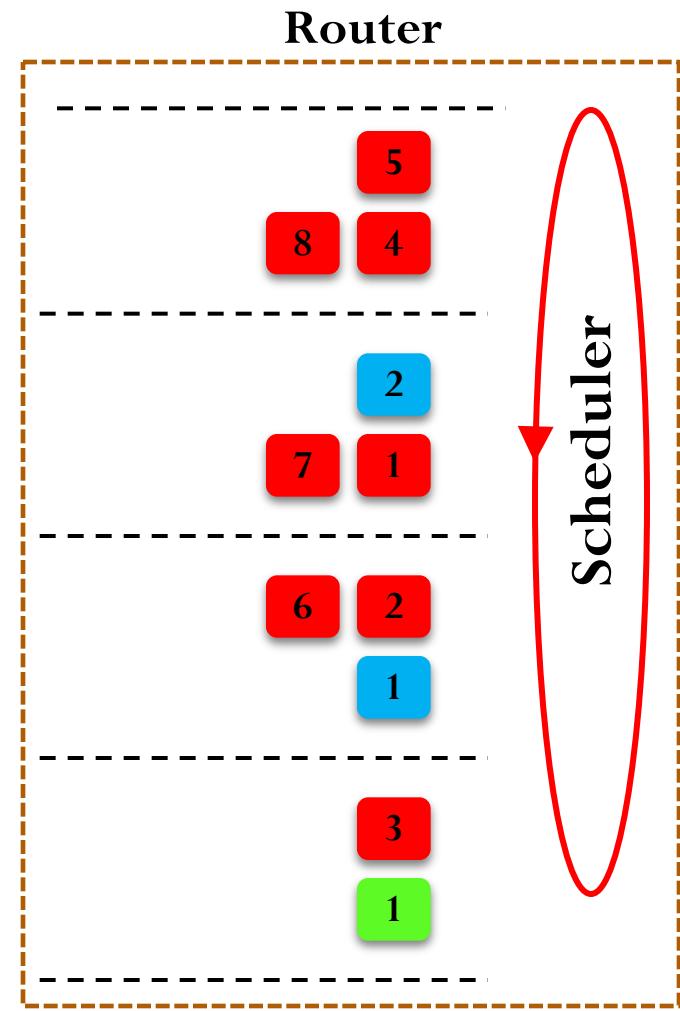
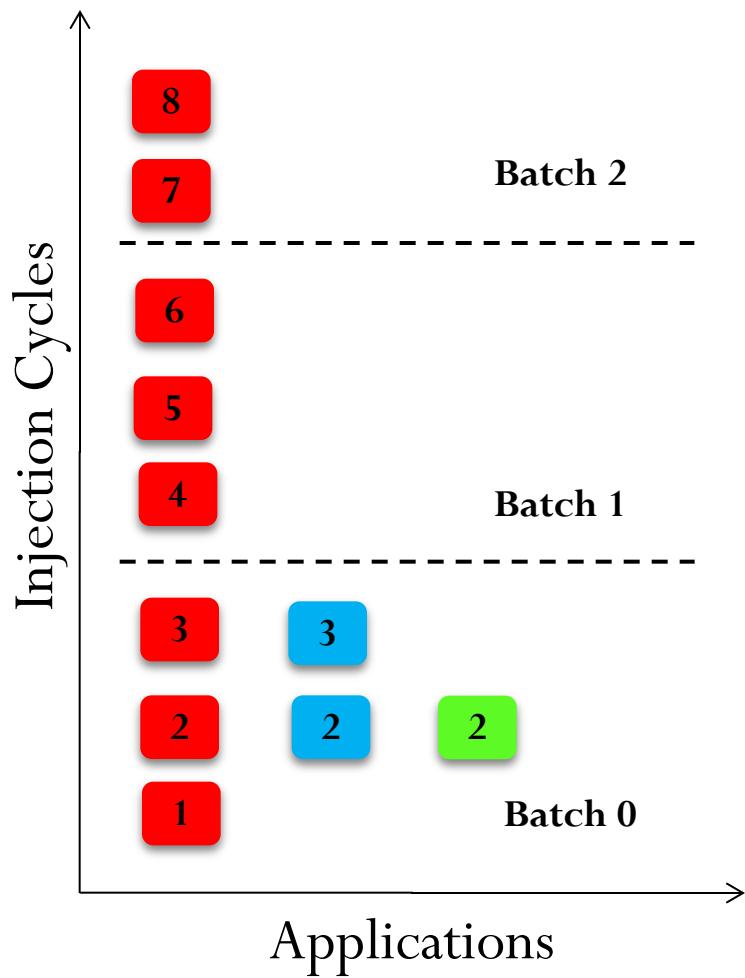
# STC Scheduling Example



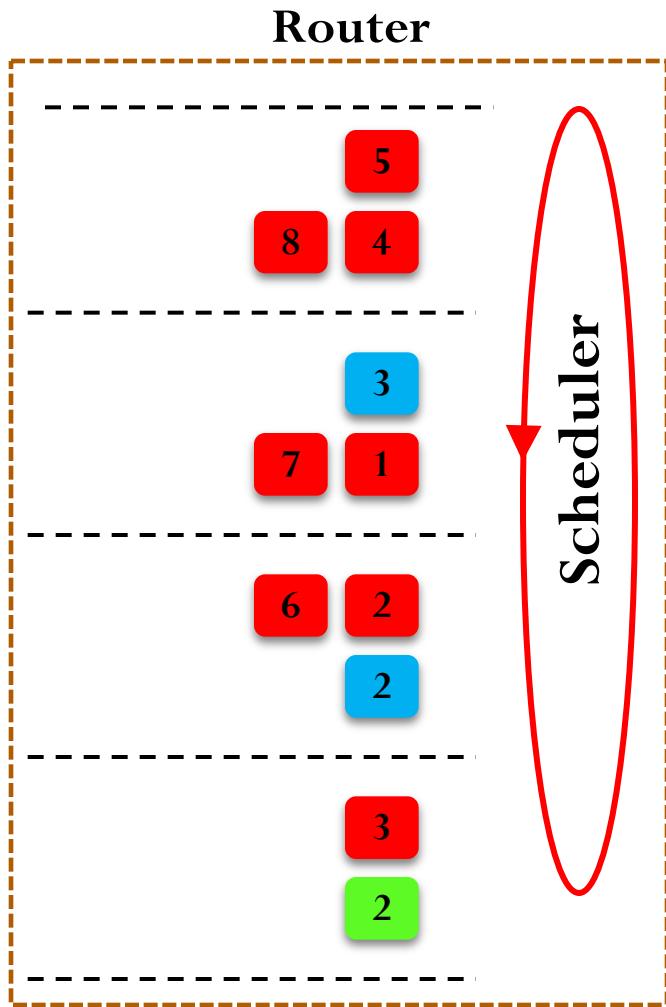
Batching interval length = 3 cycles

Ranking order = (Green > Blue > Red)

# STC Scheduling Example

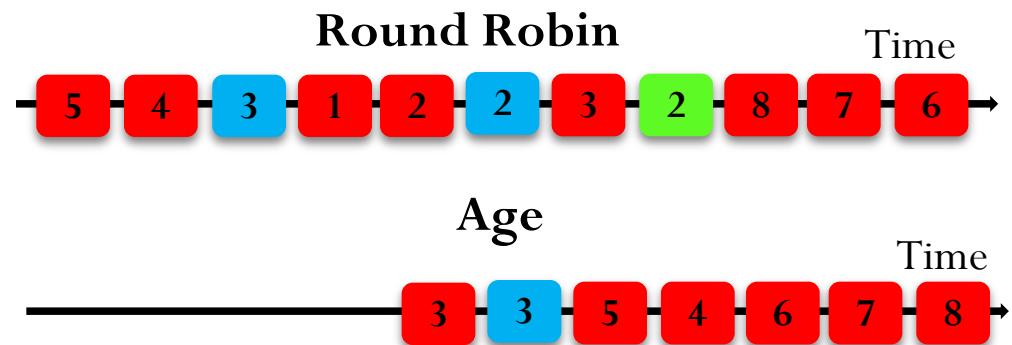
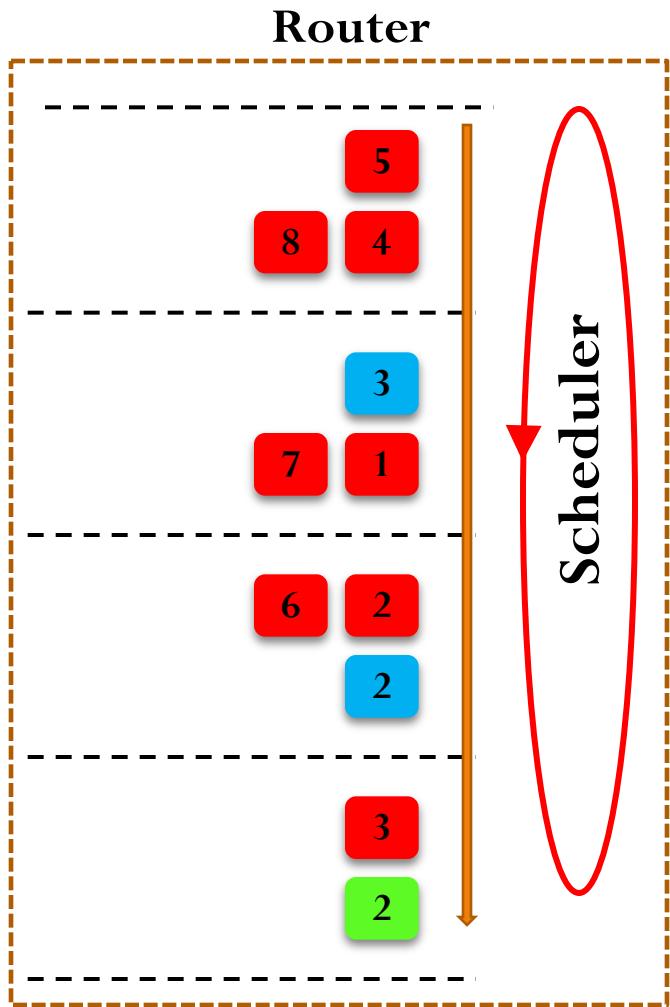


# STC Scheduling Example



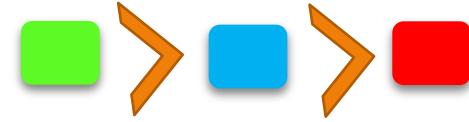
	STALL CYCLES			Avg
RR	8	6	11	8.3
Age				
STC				

# STC Scheduling Example

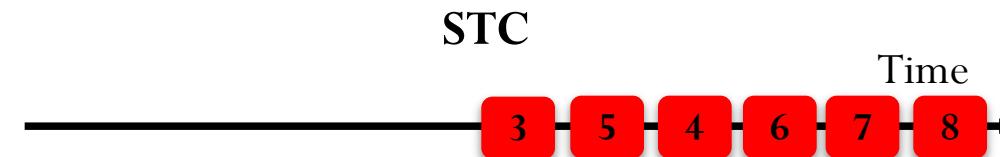
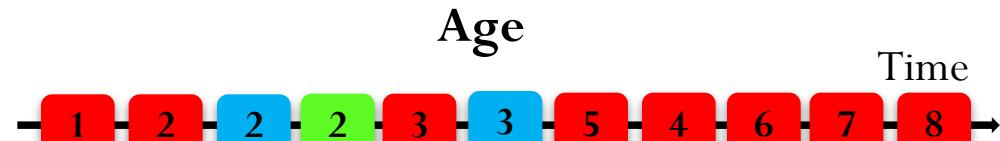
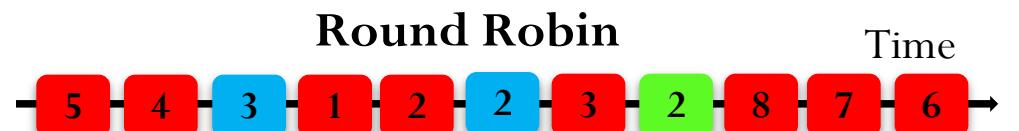
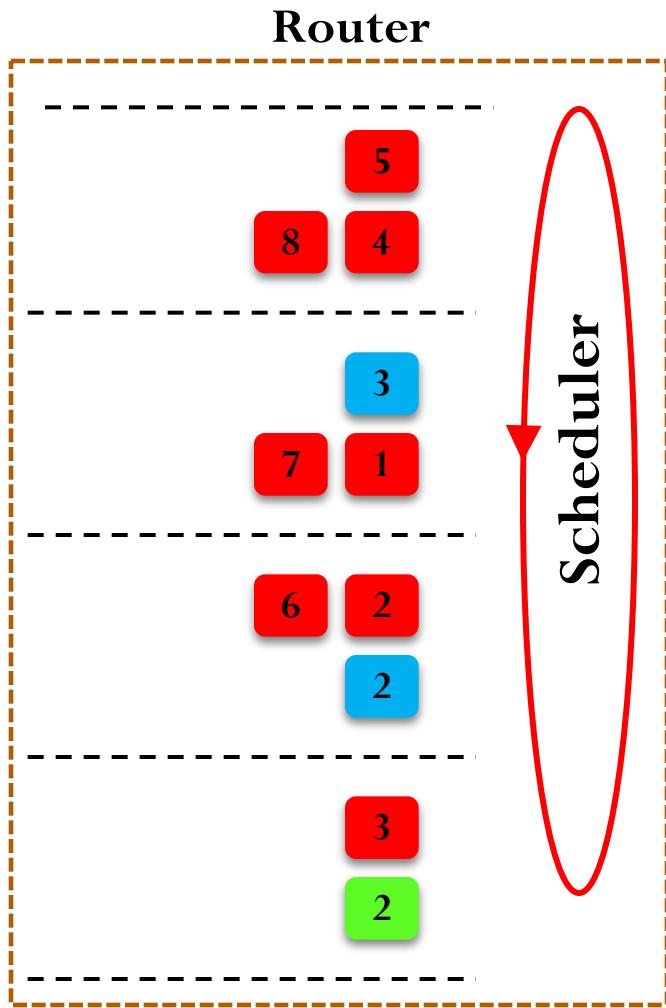


	STALL CYCLES			Avg
RR	8	6	11	8.3
Age	4	6	11	7.0
STC				

Ranking order



# STC Scheduling Example



STALL CYCLES				Avg
RR	8	6	11	8.3
Age	4	6	11	7.0
STC	1	3	11	5.0

# Application-Aware Prioritization in NoCs

---

- Reetuparna Das, Onur Mutlu, Thomas Moscibroda, and Chita R. Das,  
**"Application-Aware Prioritization Mechanisms for On-Chip Networks"**

*Proceedings of the 42nd International Symposium on Microarchitecture (MICRO), pages 280-291, New York, NY, December 2009.* Slides (pptx)

## Application-Aware Prioritization Mechanisms for On-Chip Networks

Reetuparna Das<sup>§</sup>   Onur Mutlu<sup>†</sup>   Thomas Moscibroda<sup>‡</sup>   Chita R. Das<sup>§</sup>

§Pennsylvania State University  
{rdas,das}@cse.psu.edu

†Carnegie Mellon University  
onur@cmu.edu

‡Microsoft Research  
moscitho@microsoft.com

# Slack-Based Packet Scheduling

---

- Reetuparna Das, Onur Mutlu, Thomas Moscibroda, and Chita R. Das,  
**"Aergia: Exploiting Packet Latency Slack in On-Chip Networks"**  
*Proceedings of the 37th International Symposium on Computer Architecture (ISCA)*, pages 106-116, Saint-Malo, France, June 2010. [Slides \(pptx\)](#)  
**One of the 11 computer architecture papers of 2010 selected as Top Picks by IEEE Micro.**

## Aérgia: Exploiting Packet Latency Slack in On-Chip Networks

Reetuparna Das<sup>§</sup>   Onur Mutlu<sup>†</sup>   Thomas Moscibroda<sup>‡</sup>   Chita R. Das<sup>§</sup>

§Pennsylvania State University  
[{rdas,das}@cse.psu.edu](mailto:{rdas,das}@cse.psu.edu)

†Carnegie Mellon University  
[onur@cmu.edu](mailto:onur@cmu.edu)

‡Microsoft Research  
[moscitho@microsoft.com](mailto:moscitho@microsoft.com)

# Slack-Based Packet Scheduling

---

- The notion of “packet slack”
  - Slack of a packet is the number of cycles it can be delayed in a router without (significantly) reducing application’s performance
  - Local network slack
- Source of slack: Memory-Level Parallelism (MLP) or other latency tolerance mechanisms
  - Latency of an application’s packet hidden from application due to **overlap** with latency of pending cache miss requests or other long-latency operations
- Key idea of slack-based packet scheduling:
  - Estimate the slack of each packet
  - Prioritize packets with lower slack

# Slowdown Estimation in NoCs

---

- Xiyue Xiang, Saugata Ghose, Onur Mutlu, and Nian-Feng Tzeng,  
**"A Model for Application Slowdown Estimation in On-Chip Networks and Its Use for Improving System Fairness and Performance"**

*Proceedings of the 34th IEEE International Conference on Computer Design (ICCD), Phoenix, AZ, USA, October 2016.*  
[Slides (pptx) (pdf)]

## A Model for Application Slowdown Estimation in On-Chip Networks and Its Use for Improving System Fairness and Performance

Xiyue Xiang<sup>†</sup>

<sup>†</sup>*University of Louisiana at Lafayette*

Saugata Ghose<sup>‡</sup>

<sup>‡</sup>*Carnegie Mellon University*

Onur Mutlu<sup>§‡</sup>

Nian-Feng Tzeng<sup>†</sup>

<sup>§</sup>*ETH Zürich*

# Handling Multicast and Hotspot Issues

---

- Xiyue Xiang, Wentao Shi, Saugata Ghose, Lu Peng, Onur Mutlu, and Nian-Feng Tzeng,

## **"Carpool: A Bufferless On-Chip Network Supporting Adaptive Multicast and Hotspot Alleviation"**

*Proceedings of the International Conference on Supercomputing (ICS), Chicago, IL, USA, June 2017.*

[[Slides \(pptx\)](#) ([pdf](#))]

## **Carpool: A Bufferless On-Chip Network Supporting Adaptive Multicast and Hotspot Alleviation**

Xiyue Xiang<sup>†</sup> Wentao Shi<sup>★</sup> Saugata Ghose<sup>‡</sup> Lu Peng<sup>★</sup> Onur Mutlu<sup>§‡</sup> Nian-Feng Tzeng<sup>†</sup>  
<sup>†</sup>University of Louisiana at Lafayette    <sup>★</sup>Louisiana State University    <sup>‡</sup>Carnegie Mellon University    <sup>§</sup>ETH Zürich

# Heterogeneous Networks

---

- Asit K. Mishra, Onur Mutlu, and Chita R. Das,  
**"A Heterogeneous Multiple Network-on-Chip Design: An Application-Aware Approach"**  
*Proceedings of the 50th Design Automation Conference (DAC),*  
Austin, TX, June 2013. [Slides \(pptx\)](#) [Slides \(pdf\)](#)

## A Heterogeneous Multiple Network-On-Chip Design: An Application-Aware Approach

Asit K. Mishra  
Intel Corporation  
Hillsboro, OR 97124, USA  
asit.k.mishra@intel.com

Onur Mutlu  
Carnegie Mellon University  
Pittsburgh, PA 15213, USA  
onur@cmu.edu

Chita R. Das  
The Pennsylvania State University  
University Park, PA 16802, USA  
das@cse.psu.edu

# Low-Cost QoS in On-Chip Networks (I)

---

- Boris Grot, Stephen W. Keckler, and Onur Mutlu,  
**"Preemptive Virtual Clock: A Flexible, Efficient, and Cost-effective QOS Scheme for Networks-on-Chip"**  
*Proceedings of the 42nd International Symposium on Microarchitecture (MICRO)*, pages 268-279, New York, NY, December 2009. [Slides \(pdf\)](#)

## Preemptive Virtual Clock: A Flexible, Efficient, and Cost-effective QOS Scheme for Networks-on-Chip

Boris Grot

Stephen W. Keckler

Onur Mutlu<sup>†</sup>

Department of Computer Sciences  
The University of Texas at Austin  
{bgrot, skeckler@cs.utexas.edu}

<sup>†</sup>Computer Architecture Laboratory (CALCM)  
Carnegie Mellon University  
onur@cmu.edu

# Low-Cost QoS in On-Chip Networks (II)

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"Kilo-NOC: A Heterogeneous Network-on-Chip Architecture for Scalability and Service Guarantees"**

*Proceedings of the 38th International Symposium on Computer Architecture (ISCA), San Jose, CA, June 2011.* Slides (pptx)

***One of the 12 computer architecture papers of 2011 selected as Top Picks by IEEE Micro.***

## Kilo-NOC: A Heterogeneous Network-on-Chip Architecture for Scalability and Service Guarantees

Boris Grot<sup>1</sup>

bgrot@cs.utexas.edu

Joel Hestness<sup>1</sup>

hestness@cs.utexas.edu

Stephen W. Keckler<sup>1,2</sup>

skeckler@nvidia.com

Onur Mutlu<sup>3</sup>

onur@cmu.edu

<sup>1</sup>The University of Texas at Austin  
Austin, TX

<sup>2</sup>NVIDIA  
Santa Clara, CA

<sup>3</sup>Carnegie Mellon University  
Pittsburgh, PA

# Low-Cost QoS in On-Chip Networks (III)

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"A QoS-Enabled On-Die Interconnect Fabric for Kilo-Node Chips"**  
*IEEE Micro*, Special Issue: *Micro's Top Picks from 2011 Computer Architecture Conferences (MICRO TOP PICKS)*, Vol. 32, No. 3, May/June 2012.
- 

## A QoS-ENABLED ON-DIE INTERCONNECT FABRIC FOR KILO-NODE CHIPS

---

TO MEET RAPIDLY GROWING PERFORMANCE DEMANDS AND ENERGY CONSTRAINTS, FUTURE CHIPS WILL LIKELY FEATURE THOUSANDS OF ON-DIE RESOURCES. EXISTING NETWORK-ON-CHIP SOLUTIONS WEREN'T DESIGNED FOR SCALABILITY AND WILL BE UNABLE TO MEET FUTURE INTERCONNECT DEMANDS. A HYBRID NETWORK-ON-CHIP ARCHITECTURE CALLED KILO-NoC CO-OPTIMIZES TOPOLOGY, FLOW CONTROL, AND QUALITY OF SERVICE TO ACHIEVE SIGNIFICANT GAINS IN EFFICIENCY.

---

# Kilo-NoC: Topology-Aware QoS

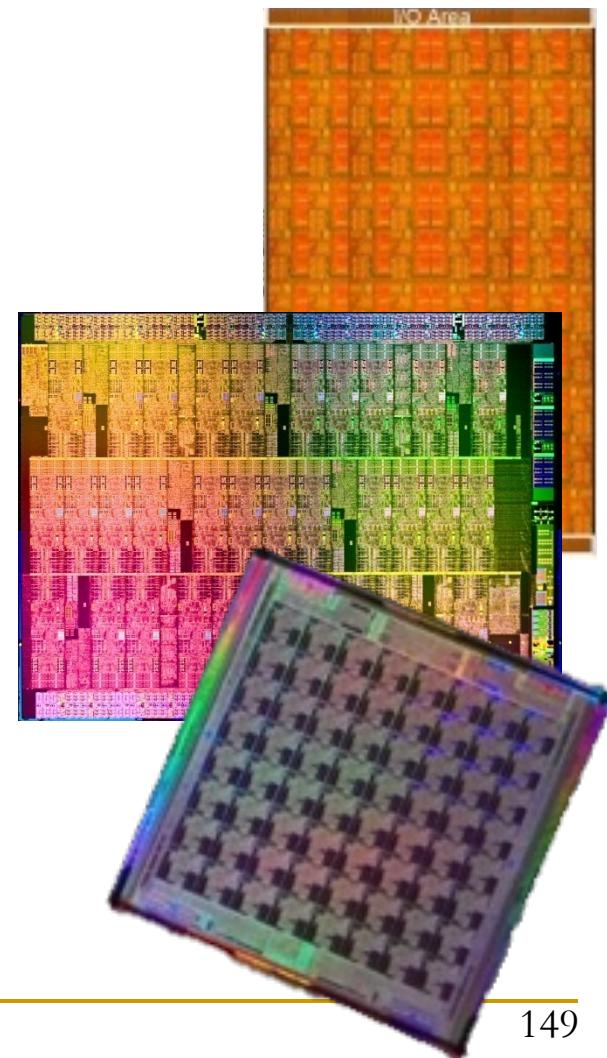
Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"Kilo-NOC: A Heterogeneous Network-on-Chip Architecture for Scalability and Service Guarantees"**

*Proceedings of the 38th International Symposium on Computer Architecture (ISCA), San Jose, CA, June 2011.* Slides (pptx)

# Motivation

---

- Extreme-scale chip-level integration
  - Cores
  - Cache banks
  - Accelerators
  - I/O logic
  - Network-on-chip (NOC)
- 10-100 cores today
- 1000+ agents in the near future



# Kilo-NOC requirements

---

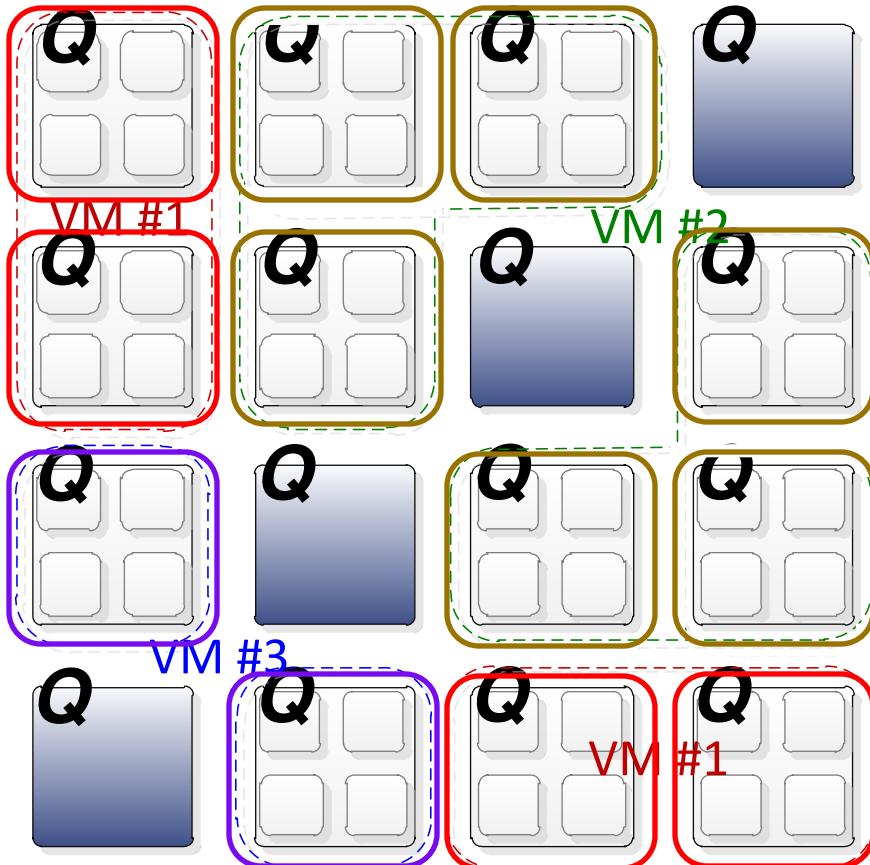
- High efficiency
  - Area
  - Energy
- Good performance
- Strong service guarantees (QoS)

# Topology-Aware QoS

---

- Problem: QoS support in each router is expensive (in terms of buffering, arbitration, bookkeeping)
  - E.g., Grot et al., “[Preemptive Virtual Clock: A Flexible, Efficient, and Cost-effective QOS Scheme for Networks-on-Chip](#),” MICRO 2009.
- Goal: Provide QoS guarantees at low area and power cost
- Idea:
  - Isolate shared resources in a region of the network, support QoS within that area
  - Design the topology so that applications can access the region without interference

# Baseline QOS-enabled CMP



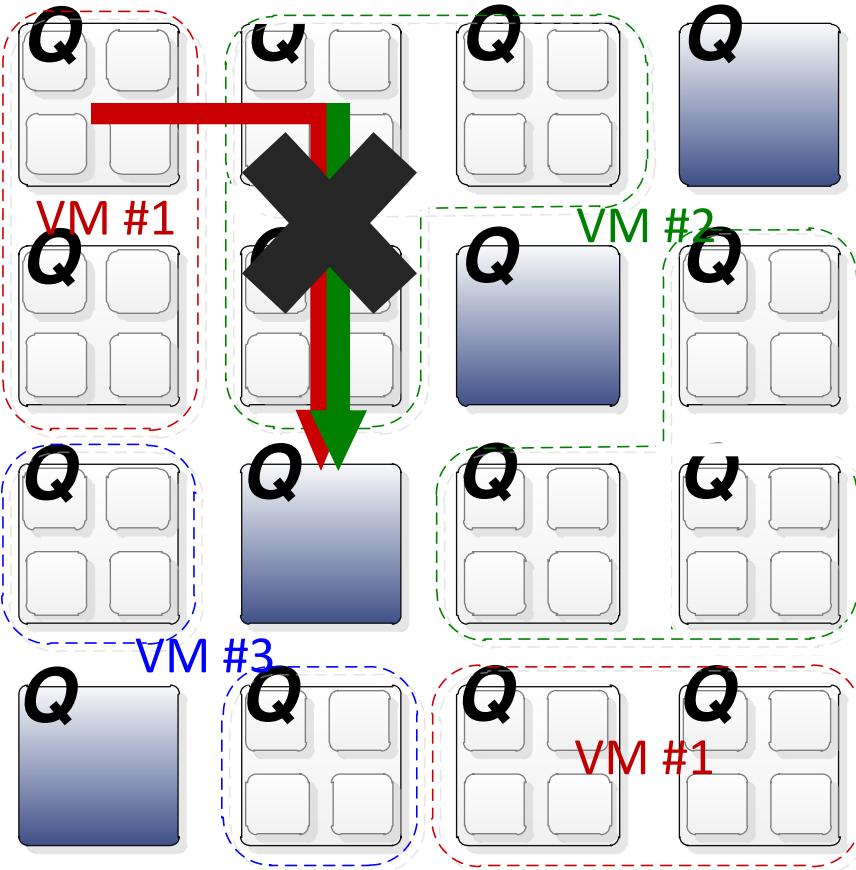
Multiple VMs  
sharing a die

■ Shared resources  
(e.g., memory controllers)

□ VM-private resources  
(cores, caches)

■ Q QOS-enabled router

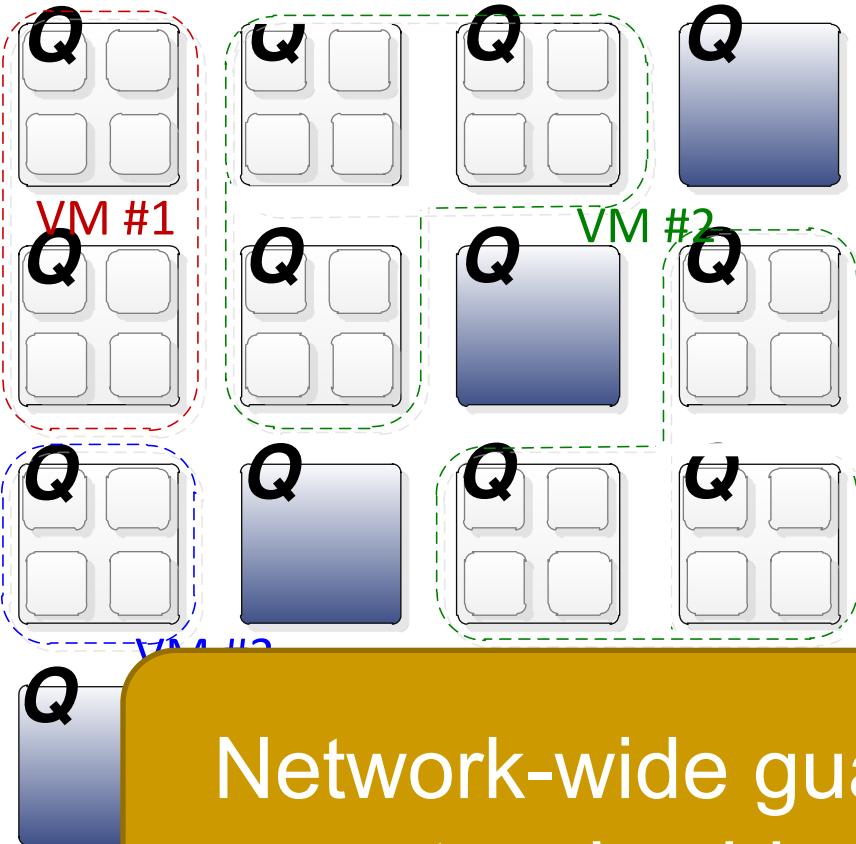
# Conventional NOC QOS



Contention scenarios:

- Shared resources
  - memory access
- Intra-VM traffic
  - shared cache access
- Inter-VM traffic
  - VM page sharing

# Conventional NOC QOS



Contention scenarios:

- Shared resources
  - memory access
- Intra-VM traffic
  - shared cache access
- Inter-VM traffic
  - VM page sharing

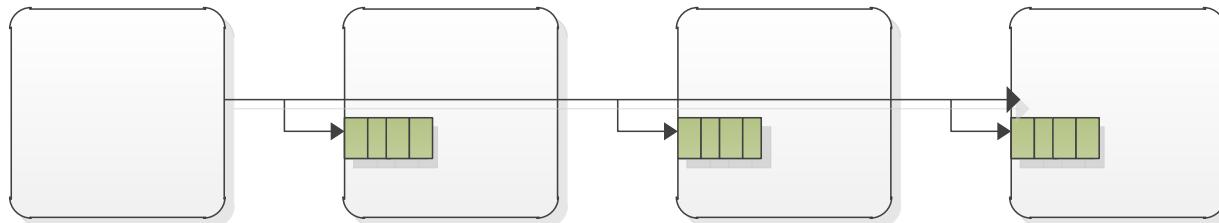
Network-wide guarantees *without* network-wide QOS support

# Kilo-NOC QOS

---

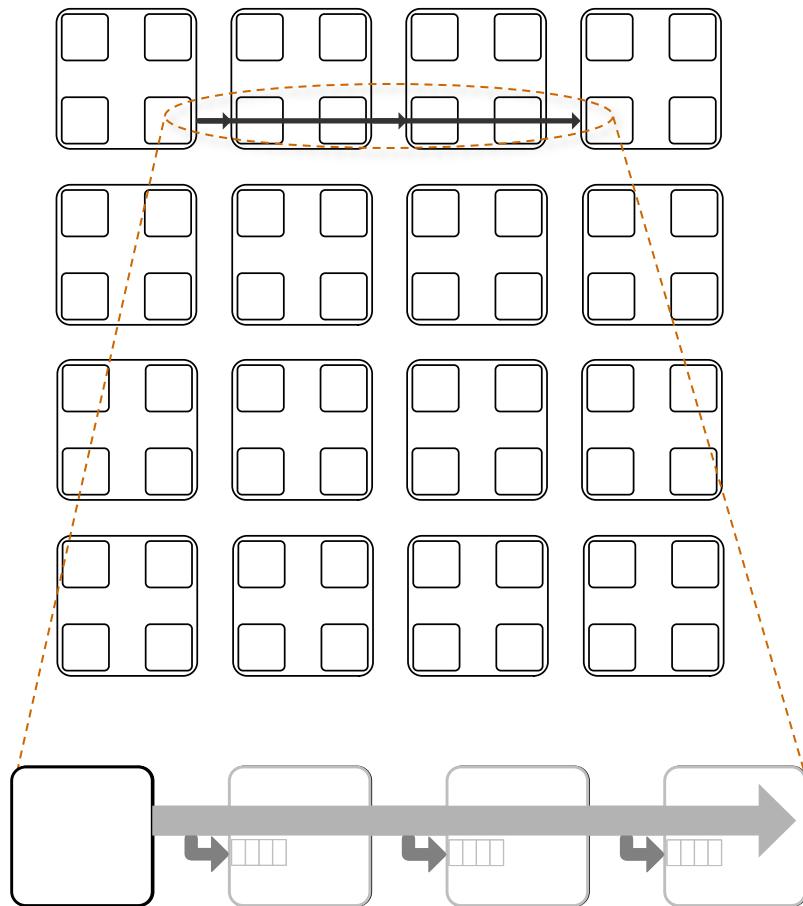
- Insight: leverage rich network connectivity
  - Naturally reduce interference among flows
    - Limit the extent of hardware QOS support
- Requires a low-diameter topology
  - This work: Multidrop Express Channels (MECS)

*Grot et al., HPCA  
2009*



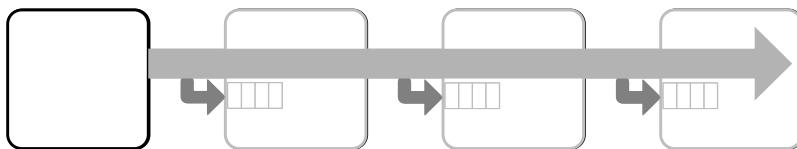
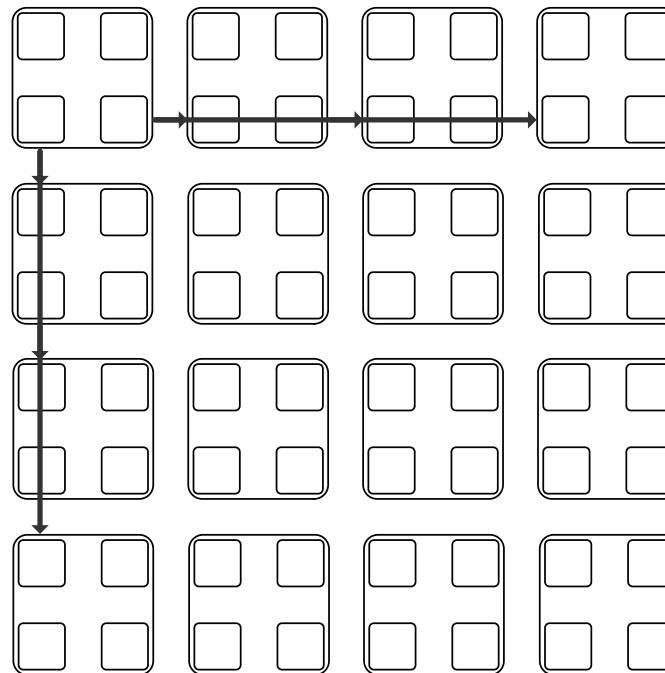
# Multidrop Express Channels (MECS)

---



# Multidrop Express Channels (MECS)

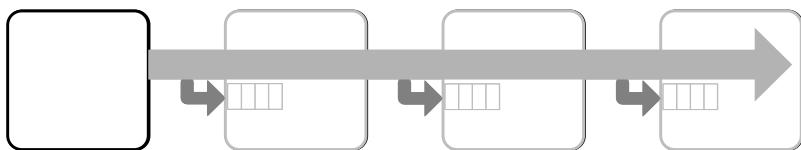
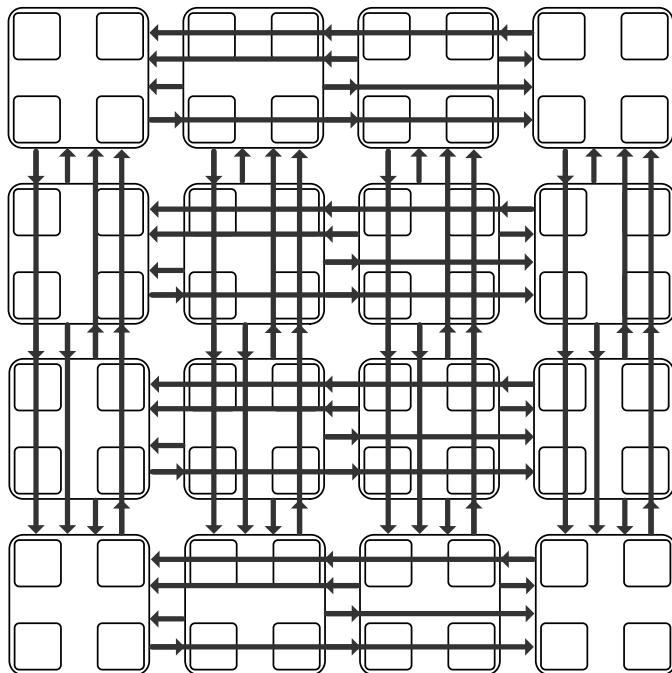
---



---

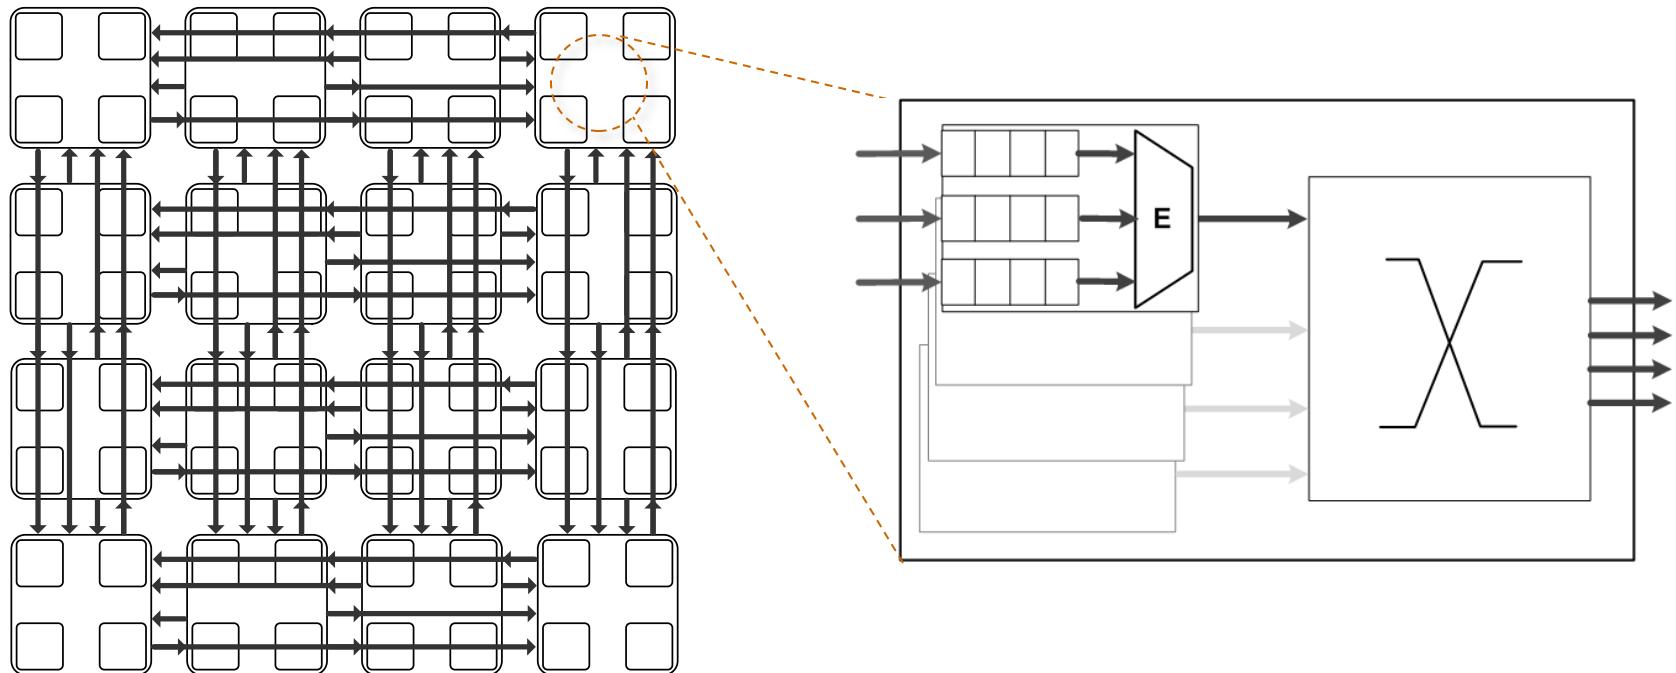
# Multidrop Express Channels (MECS)

---



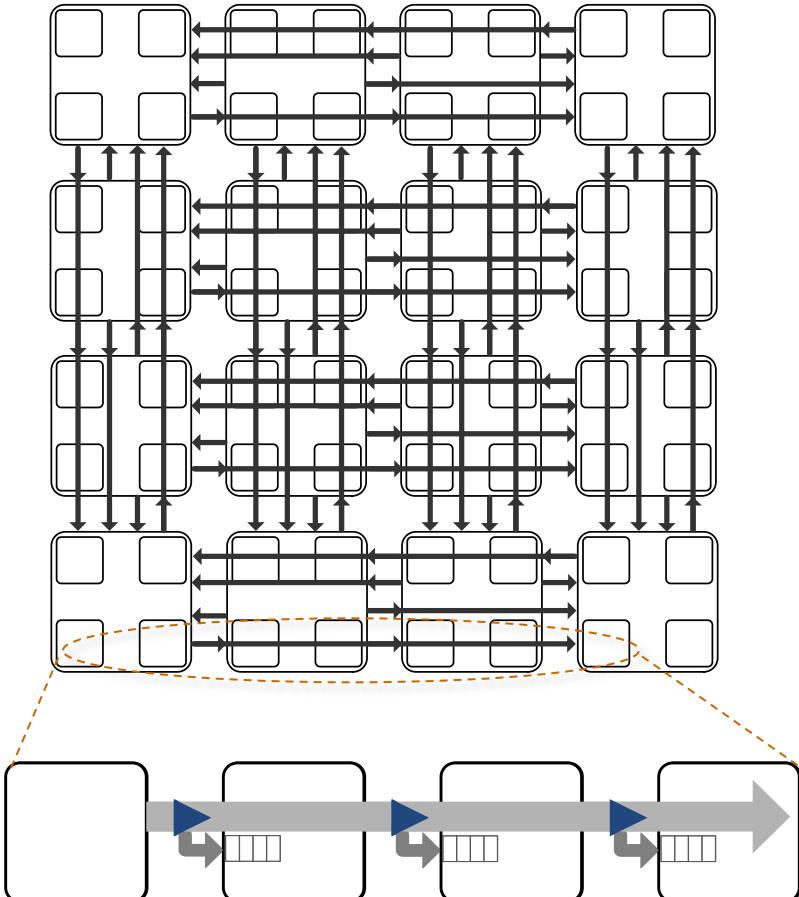
# Multidrop Express Channels (MECS)

---



# Multidrop Express Channels (MECS)

---



## ■ Pros

- One-to-many topology
- Low diameter: 2 hops
- $k$  channels row/column
- Asymmetric

## ■ Cons

- Asymmetric
- Increased control (arbitration) complexity

# Multi-Drop Express Channels (MECS)

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"Express Cube Topologies for On-Chip Interconnects"**  
*Proceedings of the 15th International Symposium on High-Performance Computer Architecture (HPCA)*, pages 163-174, Raleigh, NC, February 2009. [Slides \(ppt\)](#)

## Express Cube Topologies for On-Chip Interconnects

Boris Grot

Joel Hestness

Stephen W. Keckler

Onur Mutlu<sup>†</sup>

Department of Computer Sciences

The University of Texas at Austin

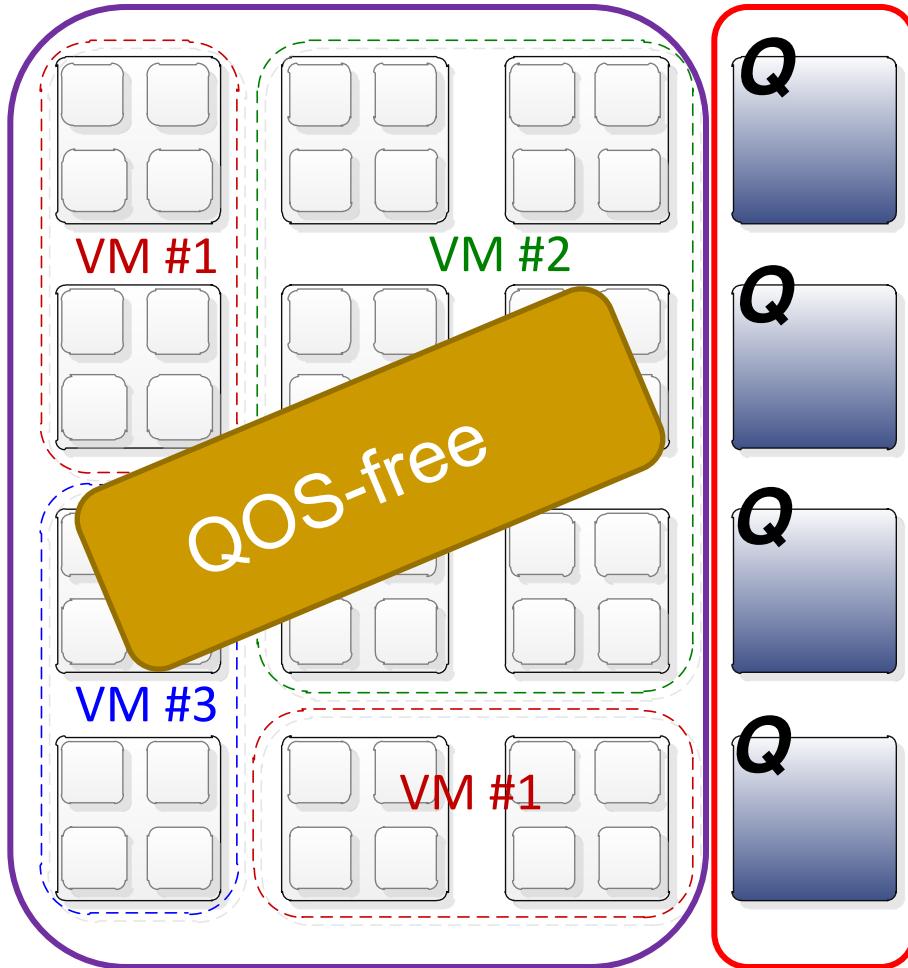
{bgrot, hestness, skeckler}@cs.utexas.edu

<sup>†</sup>Computer Architecture Laboratory (CALCM)

Carnegie Mellon University

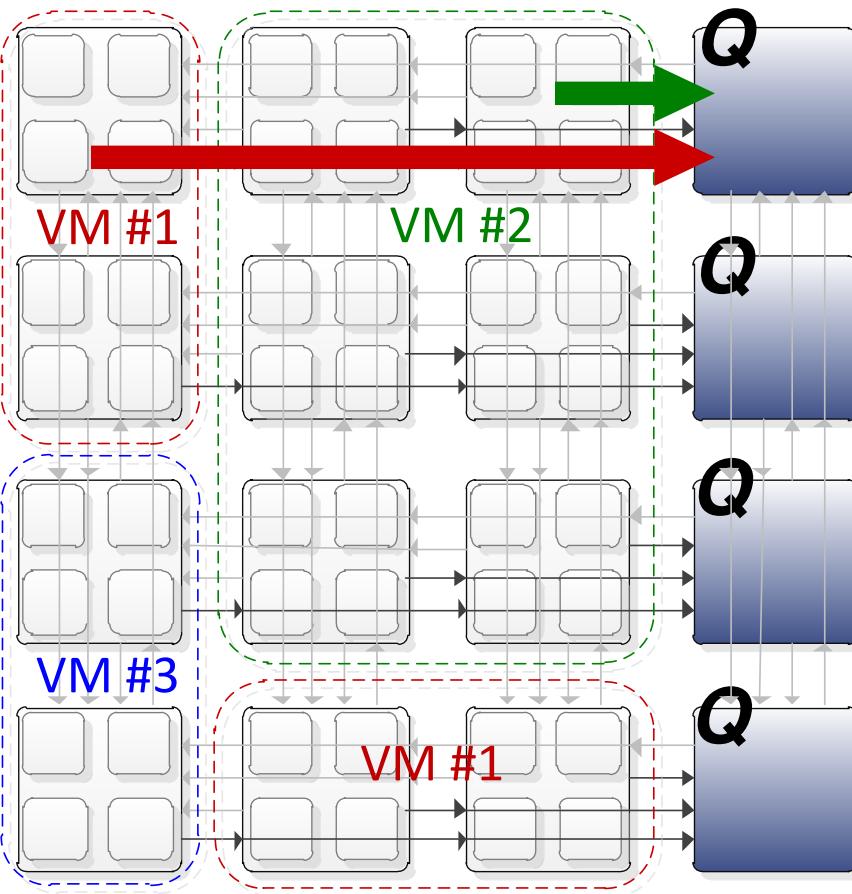
onur@cmu.edu

# Topology-Aware QOS



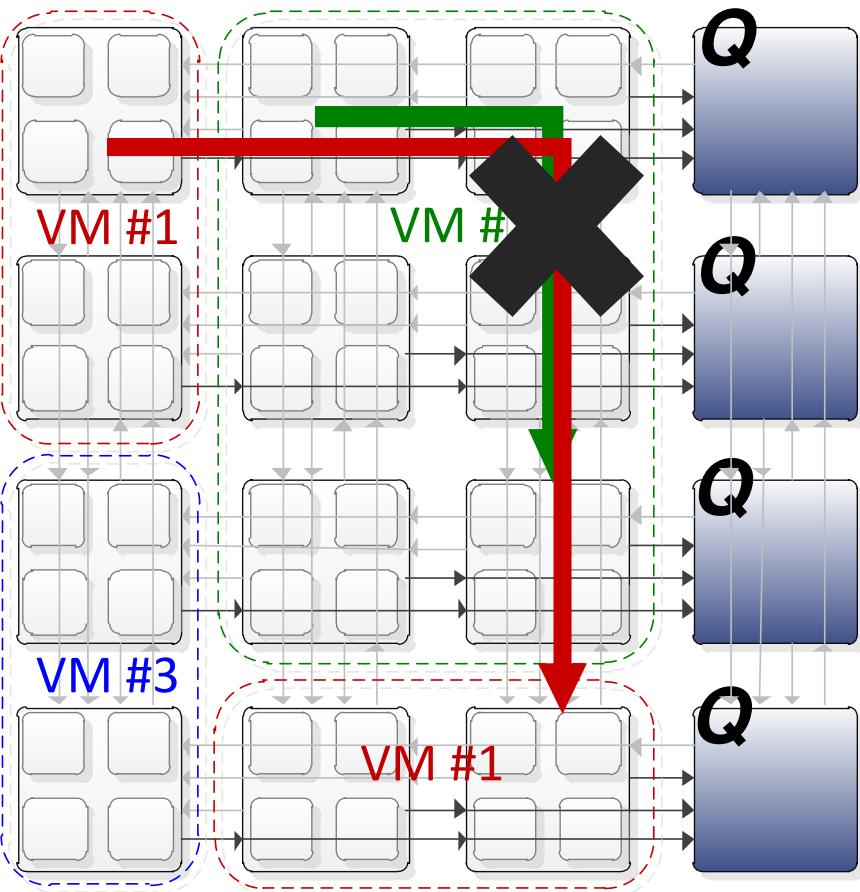
- Dedicated, QOS-enabled regions
  - Rest of die: QOS-free
- Richly-connected topology
  - Traffic isolation
- Special routing rules
  - Manage interference

# Topology-Aware QOS



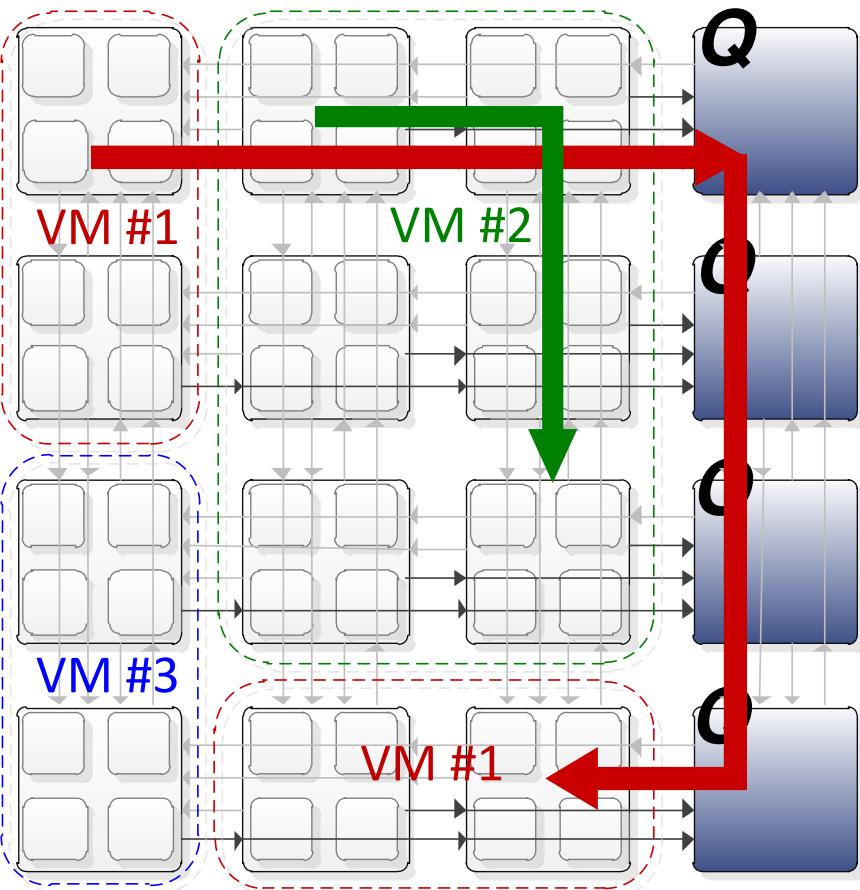
- Dedicated, QOS-enabled regions
  - Rest of die: QOS-free
- Richly-connected topology
  - Traffic isolation
- Special routing rules
  - Manage interference

# Topology-Aware QOS



- Dedicated, QOS-enabled regions
  - Rest of die: QOS-free
- Richly-connected topology
  - Traffic isolation
- Special routing rules
  - Manage interference

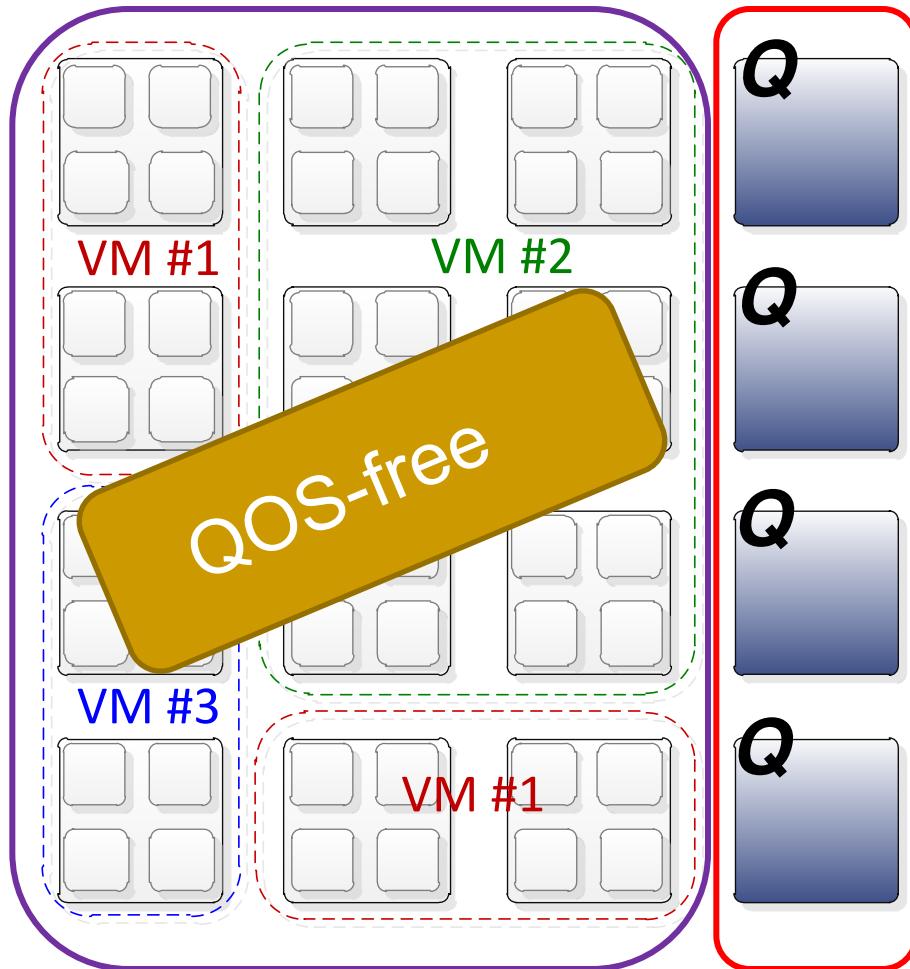
# Topology-Aware QOS



- Dedicated, QOS-enabled regions
  - Rest of die: QOS-free
- Richly-connected topology
  - Traffic isolation
- Special routing rules
  - Manage interference

# Kilo-NOC view

---

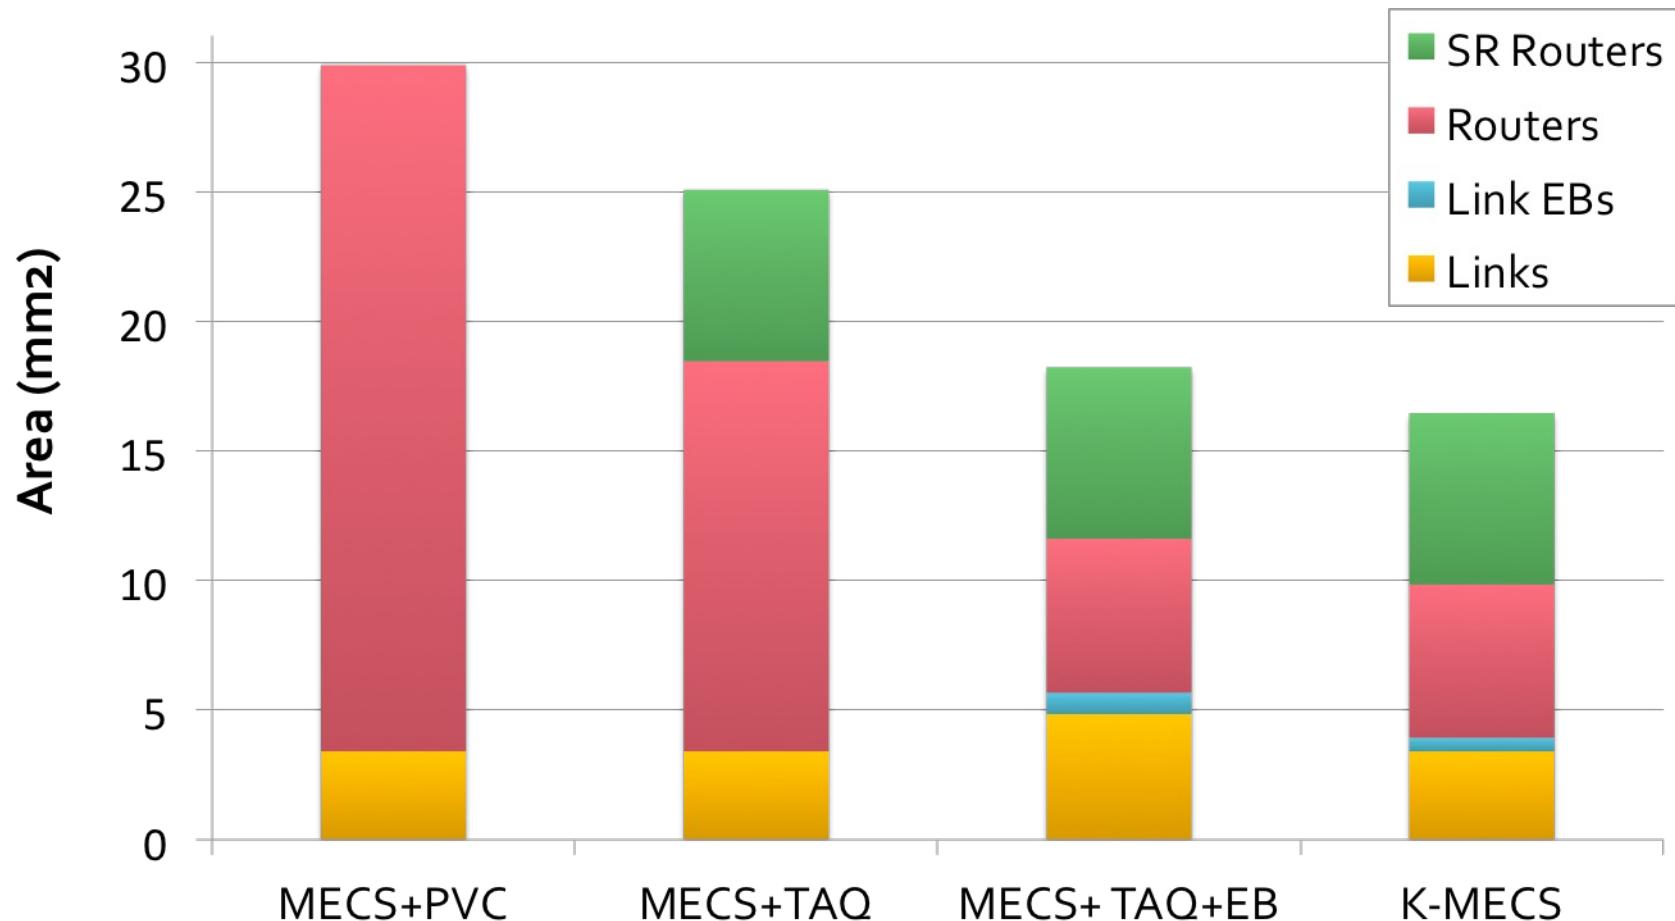


- Topology-aware QOS support
  - Limit QOS complexity to a fraction of the die
- Optimized flow control
  - Reduce buffer requirements in QOS-free regions

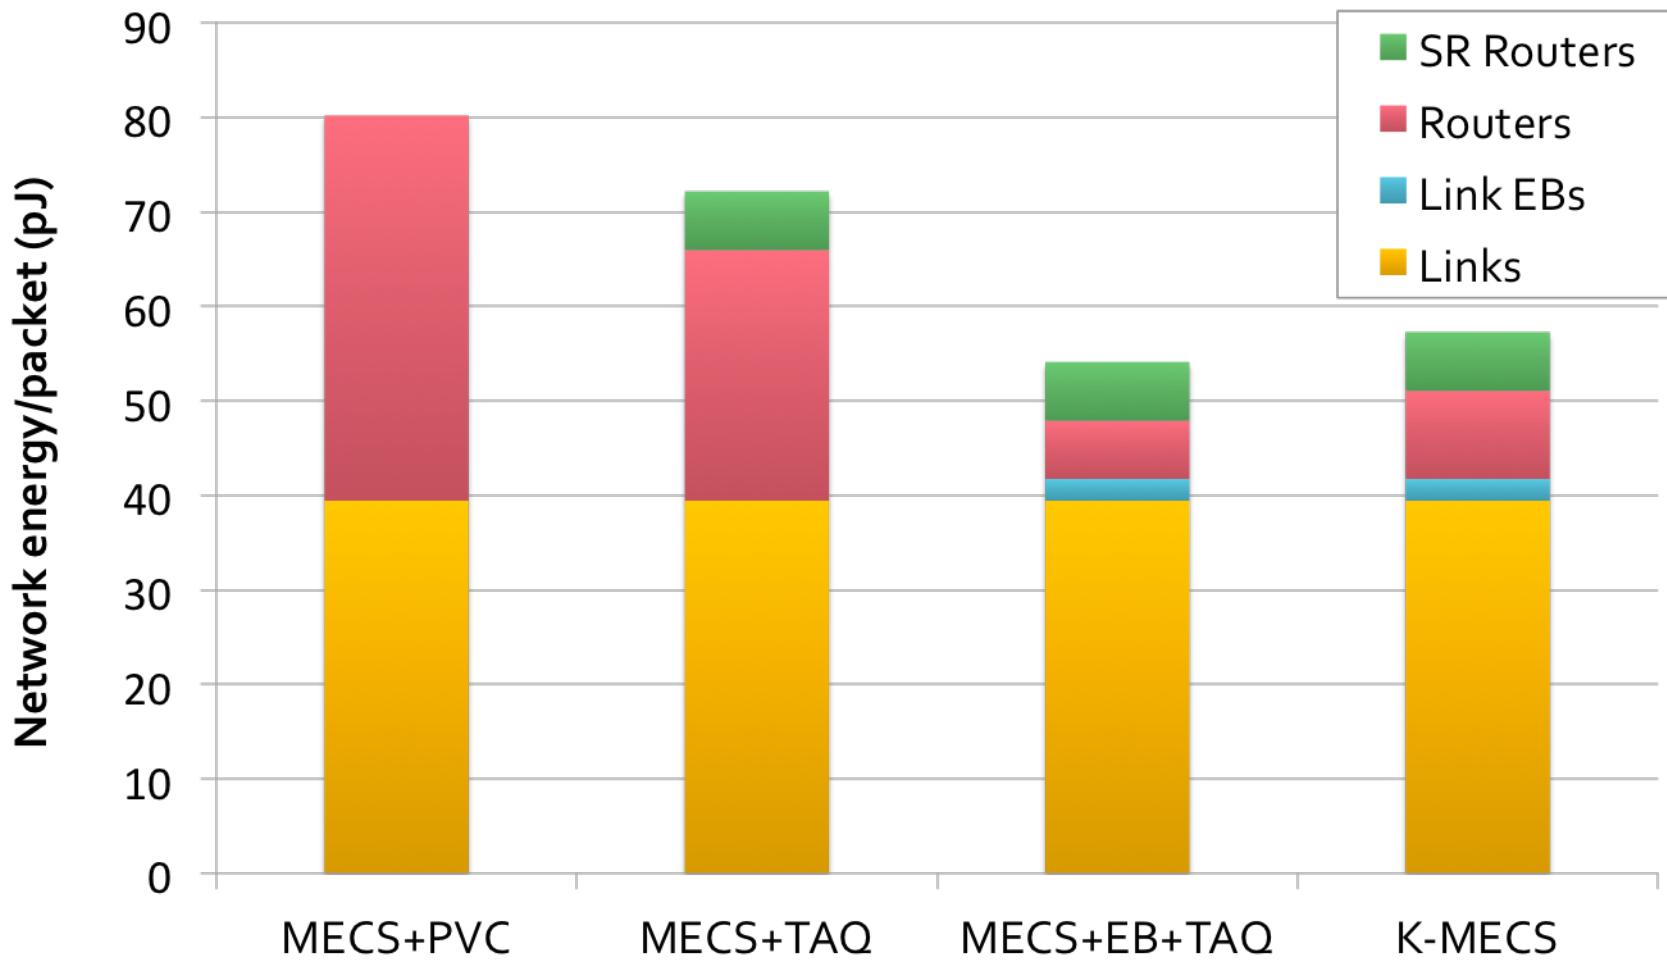
# Evaluation Methodology

Parameter	Value
Technology	15 nm
Vdd	0.7 V
System	1024 tiles: 256 concentrated nodes (64 shared resources)
<b>Networks:</b>	
MECS+PVC	VC flow control, QOS support (PVC) at each node
MECS+TAQ	VC flow control, QOS support only in shared regions
MECS+TAQ+EB	EB flow control outside of SRs, Separate <i>Request</i> and <i>Reply</i> networks
K-MECS	Proposed organization: TAQ + hybrid flow control

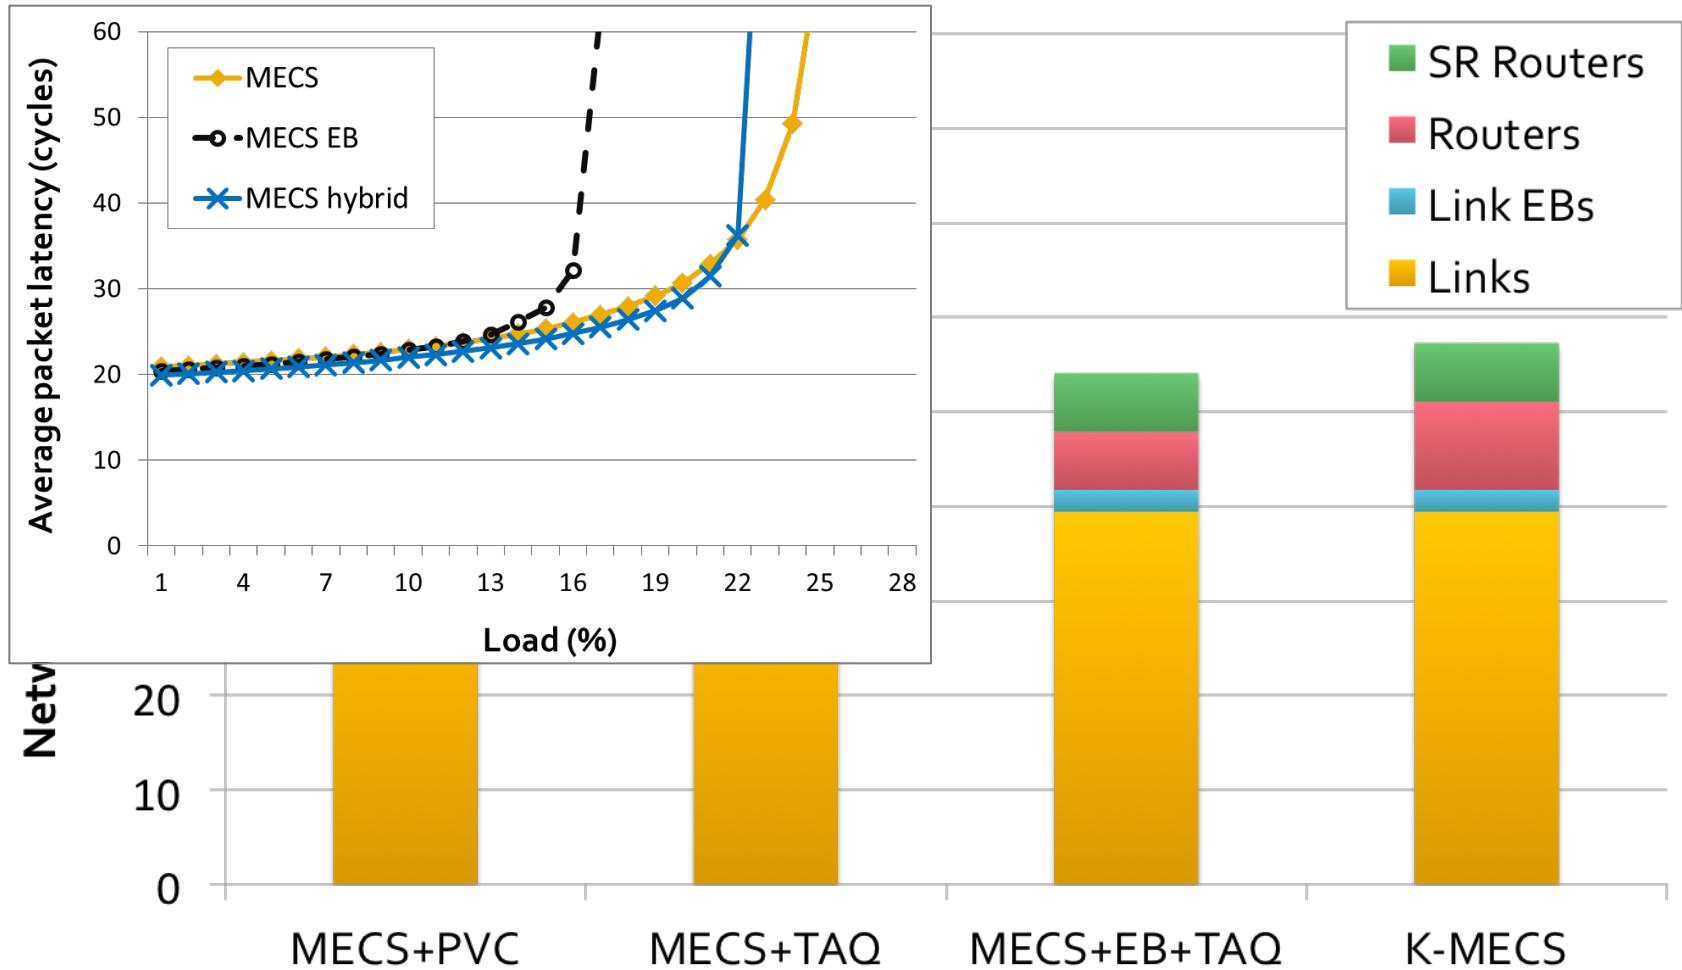
# Area comparison



# Energy comparison



# Energy comparison



# Summary

Kilo-NOC: a heterogeneous NOC architecture for kilo-node substrates

- Topology-aware QOS
  - Limits QOS support to a fraction of the die
  - Leverages low-diameter topologies
  - Improves NOC area- and energy-efficiency
  - Provides strong guarantees

# More on Kilo-NoC (I)

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"Kilo-NOC: A Heterogeneous Network-on-Chip Architecture for Scalability and Service Guarantees"**

*Proceedings of the 38th International Symposium on Computer Architecture (ISCA), San Jose, CA, June 2011.* Slides (pptx)

***One of the 12 computer architecture papers of 2011 selected as Top Picks by IEEE Micro.***

## Kilo-NOC: A Heterogeneous Network-on-Chip Architecture for Scalability and Service Guarantees

Boris Grot<sup>1</sup>

bgrot@cs.utexas.edu

Joel Hestness<sup>1</sup>

hestness@cs.utexas.edu

Stephen W. Keckler<sup>1,2</sup>

skeckler@nvidia.com

Onur Mutlu<sup>3</sup>

onur@cmu.edu

<sup>1</sup>The University of Texas at Austin  
Austin, TX

<sup>2</sup>NVIDIA  
Santa Clara, CA

<sup>3</sup>Carnegie Mellon University  
Pittsburgh, PA

# More on Kilo-NoC (II)

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"A QoS-Enabled On-Die Interconnect Fabric for Kilo-Node Chips"**  
*IEEE Micro*, Special Issue: *Micro's Top Picks from 2011 Computer Architecture Conferences (MICRO TOP PICKS)*, Vol. 32, No. 3, May/June 2012.
- 

## A QoS-ENABLED ON-DIE INTERCONNECT FABRIC FOR KILO-NODE CHIPS

---

TO MEET RAPIDLY GROWING PERFORMANCE DEMANDS AND ENERGY CONSTRAINTS, FUTURE CHIPS WILL LIKELY FEATURE THOUSANDS OF ON-DIE RESOURCES. EXISTING NETWORK-ON-CHIP SOLUTIONS WEREN'T DESIGNED FOR SCALABILITY AND WILL BE UNABLE TO MEET FUTURE INTERCONNECT DEMANDS. A HYBRID NETWORK-ON-CHIP ARCHITECTURE CALLED KILO-NOC CO-OPTIMIZES TOPOLOGY, FLOW CONTROL, AND QUALITY OF SERVICE TO ACHIEVE SIGNIFICANT GAINS IN EFFICIENCY.

---

# Multi-Drop Express Channels

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"Express Cube Topologies for On-Chip Interconnects"**  
*Proceedings of the 15th International Symposium on High-Performance Computer Architecture (HPCA)*, pages 163-174, Raleigh, NC, February 2009. [Slides \(ppt\)](#)

## Express Cube Topologies for On-Chip Interconnects

Boris Grot

Joel Hestness

Stephen W. Keckler

Onur Mutlu<sup>†</sup>

Department of Computer Sciences

The University of Texas at Austin

{bgrot, hestness, skeckler}@cs.utexas.edu

<sup>†</sup>Computer Architecture Laboratory (CALCM)

Carnegie Mellon University

onur@cmu.edu

# Computer Architecture

## Lecture 26: On-Chip Networks

Prof. Onur Mutlu

ETH Zürich

Fall 2023

26 January 2024

# Backup Slides for Additional Information

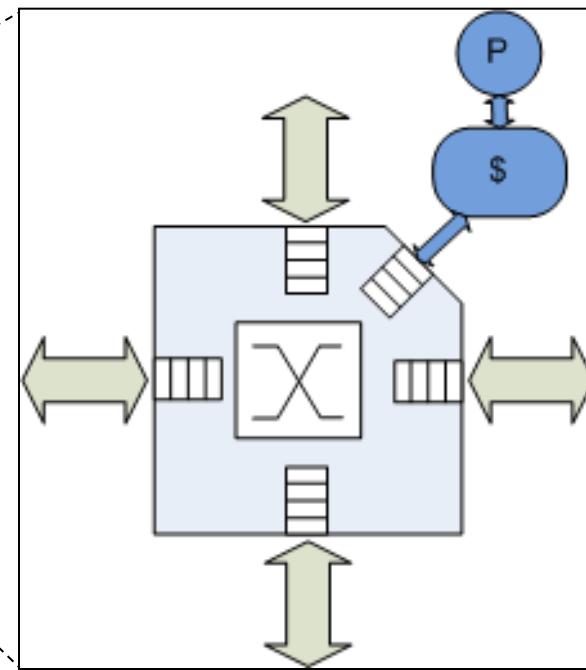
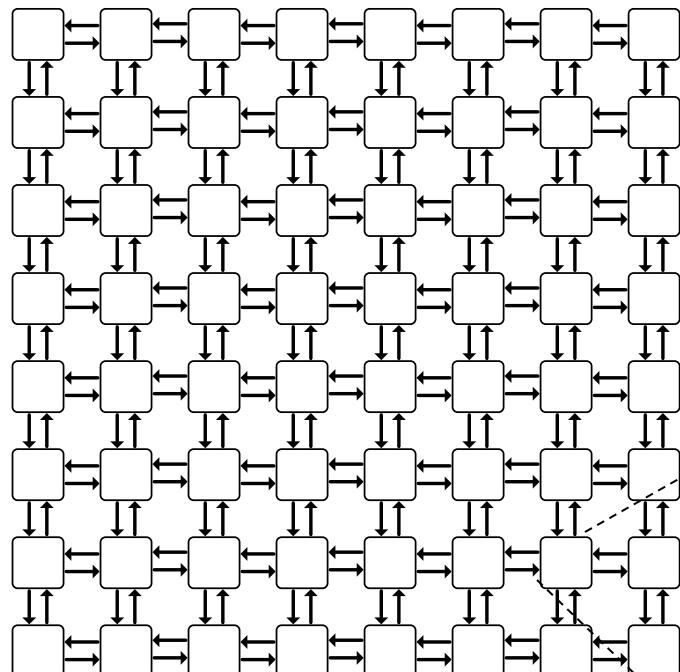
# Express-Cube Topologies

Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"Express Cube Topologies for On-Chip Interconnects"**

*Proceedings of the 15th International Symposium on High-Performance Computer Architecture (HPCA), pages 163-174, Raleigh, NC, February 2009.*  
Slides (ppt)

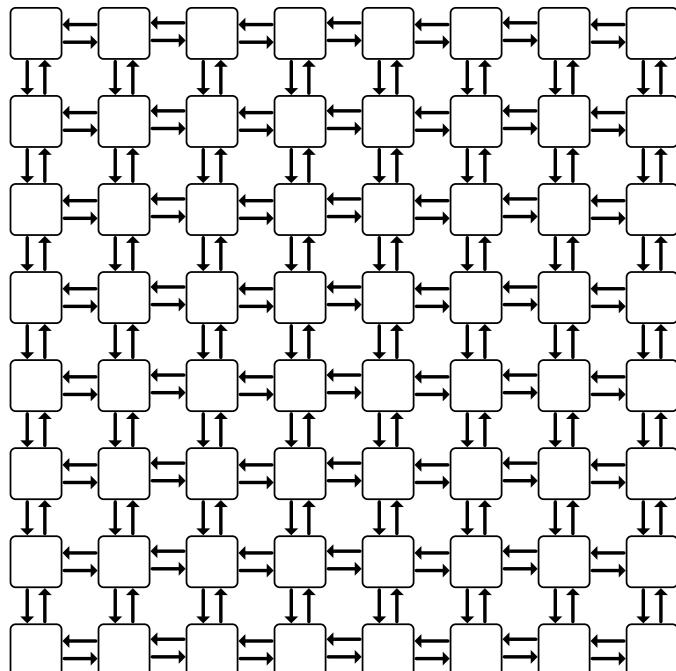
# 2-D Mesh

---



# 2-D Mesh

---



## □ Pros

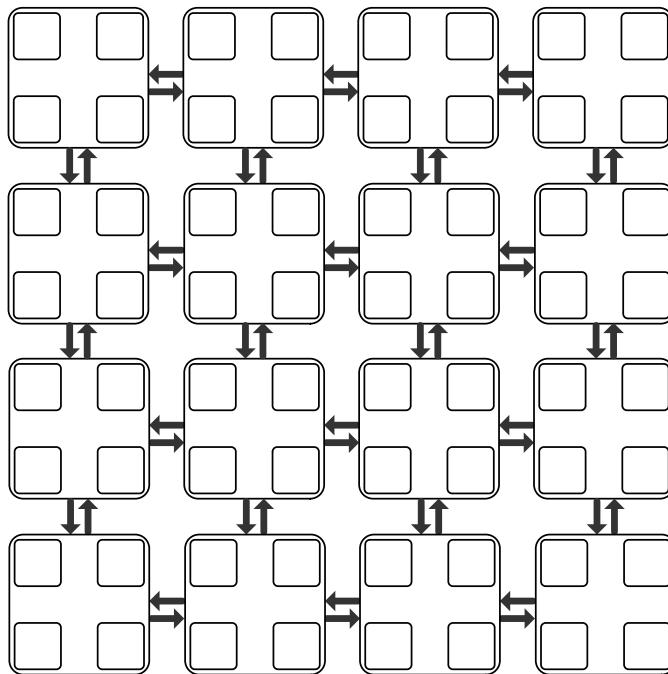
- Low design & layout complexity
- Simple, fast routers

## □ Cons

- Large diameter
- Energy & latency impact

# Concentration (*Balfour & Dally, ICS '06*)

---



## □ Pros

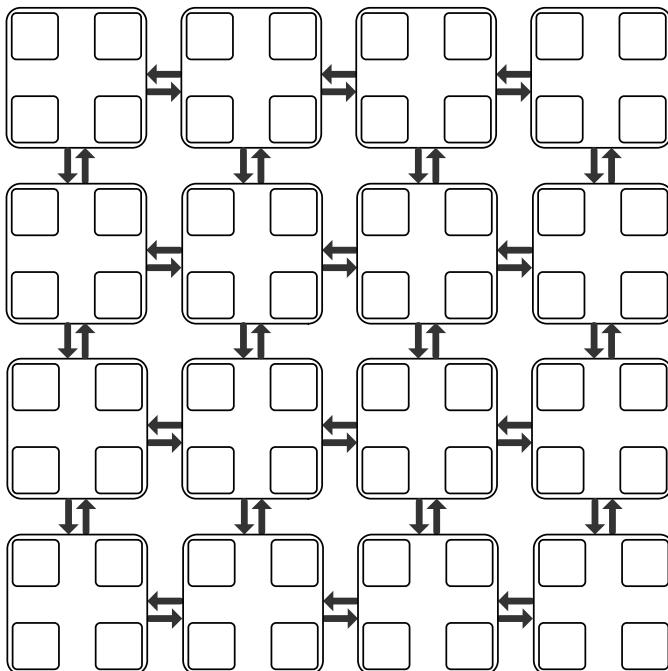
- Multiple *terminals* attached to a router node
- Fast nearest-neighbor communication via the crossbar
- Hop count reduction proportional to *concentration* degree

## □ Cons

- Benefits limited by crossbar complexity

# Concentration

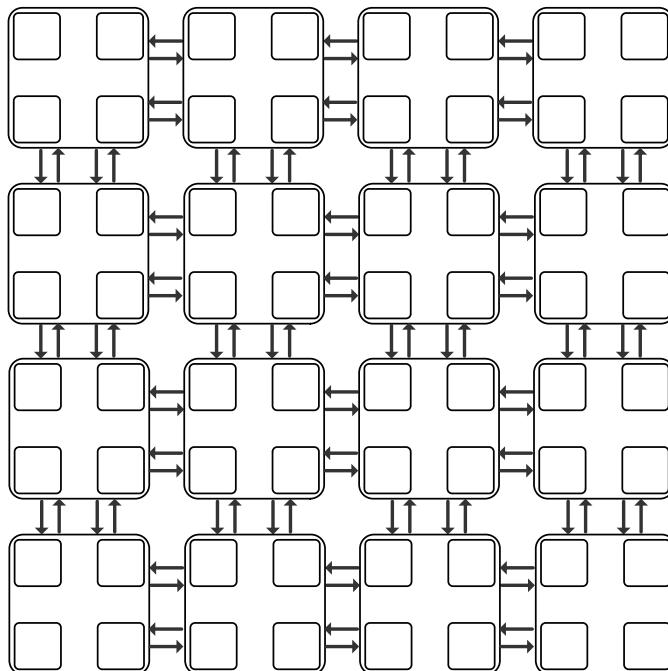
---



- Side-effects
  - Fewer channels
  - Greater channel width

# Replication

---



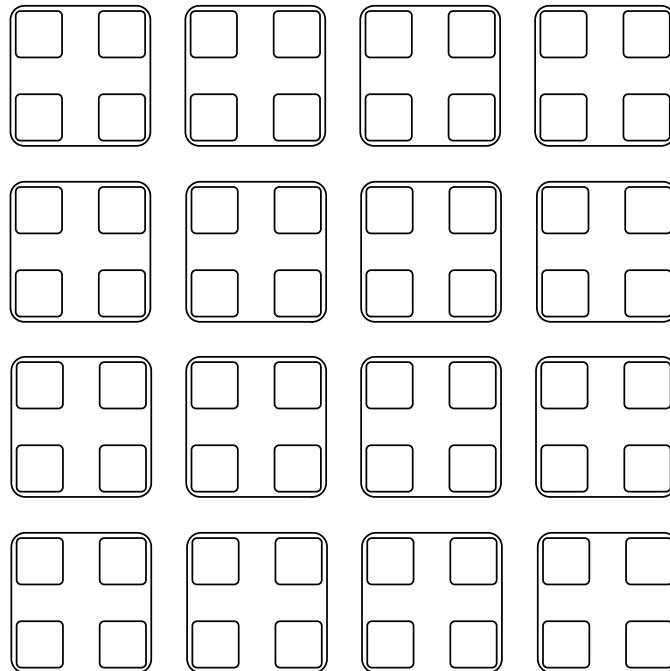
CMesh-X2

## □ Benefits

- Restores bisection channel count
- Restores channel width
- Reduced crossbar complexity

# Flattened Butterfly (*Kim et al., Micro '07*)

---

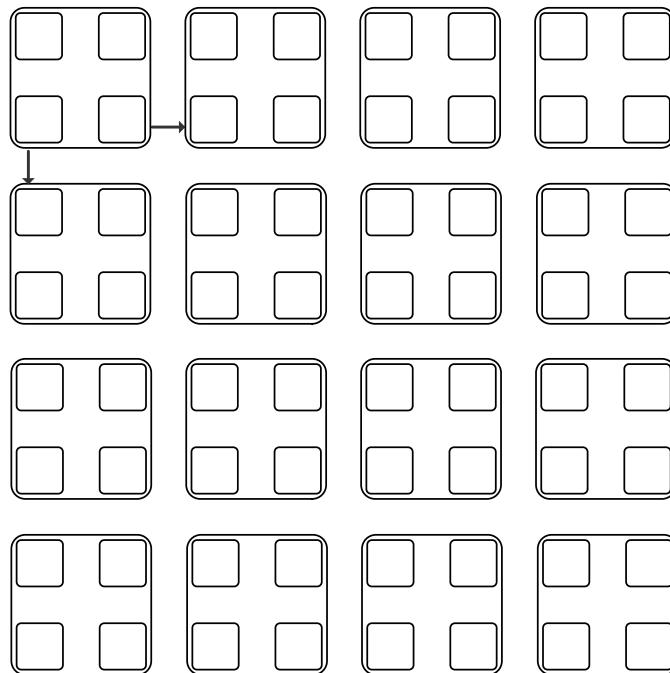


## □ Objectives:

- Improve connectivity
- Exploit the wire budget

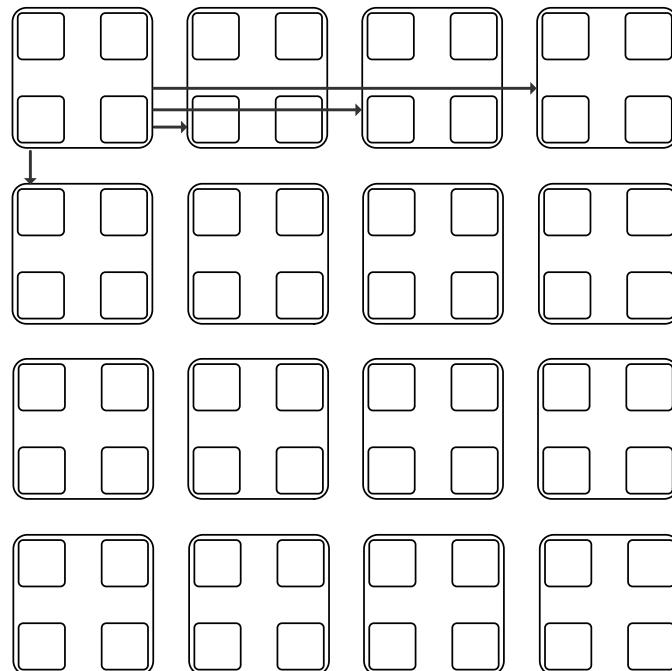
# Flattened Butterfly (*Kim et al., Micro '07*)

---



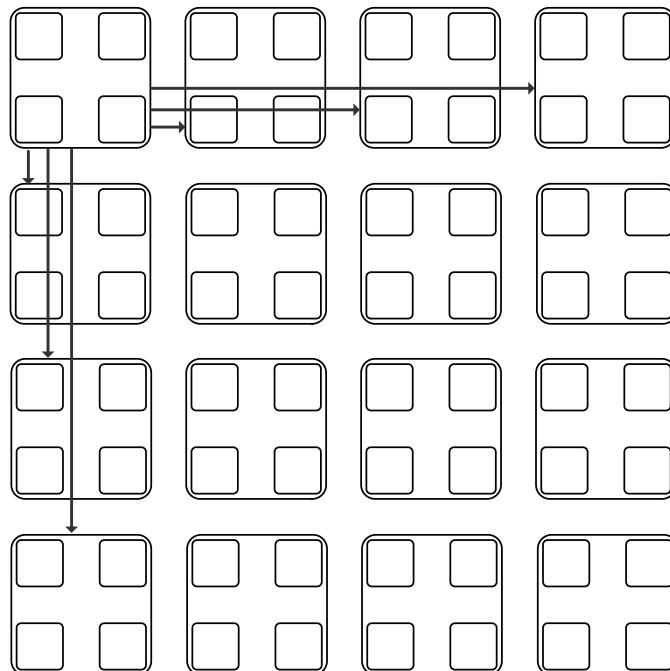
# Flattened Butterfly (*Kim et al., Micro '07*)

---



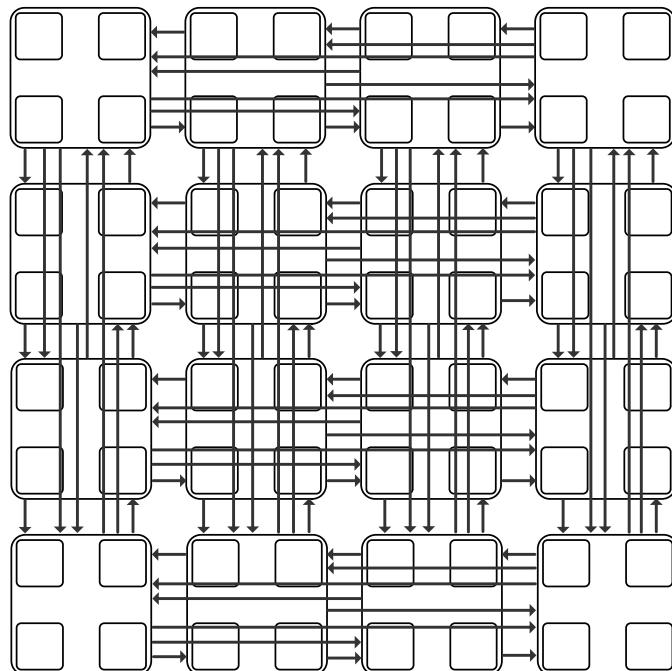
# Flattened Butterfly (*Kim et al., Micro '07*)

---



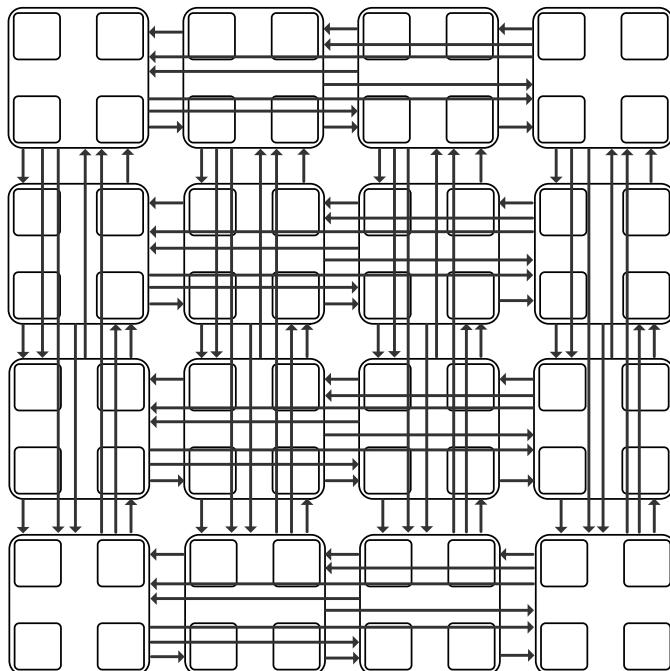
# Flattened Butterfly (*Kim et al., Micro '07*)

---



# Flattened Butterfly (*Kim et al., Micro '07*)

---



## □ Pros

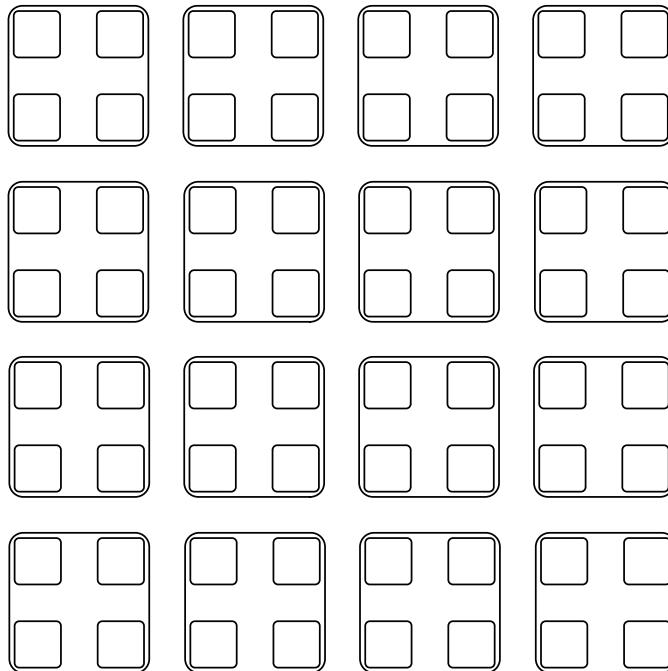
- Excellent connectivity
- Low diameter: 2 hops

## □ Cons

- High channel count:  
 $k^2/2$  per row/column
- Low channel utilization
- Increased control  
(arbitration) complexity

# Multidrop Express Channels (MECS)

---

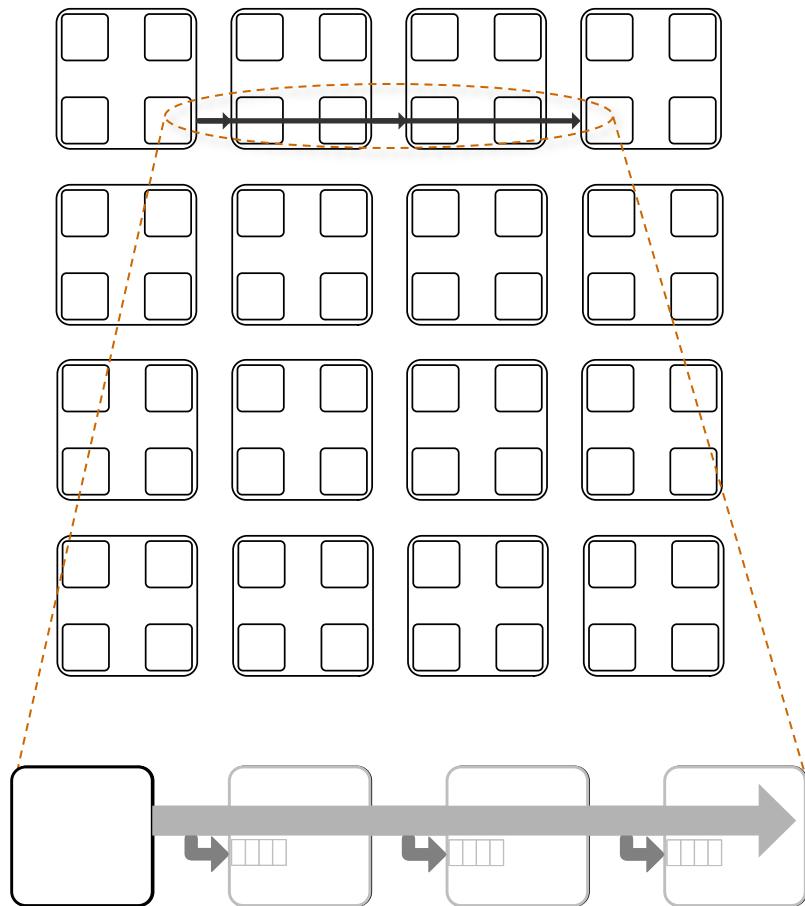


## □ Objectives:

- Connectivity
- More scalable channel count
- Better channel utilization

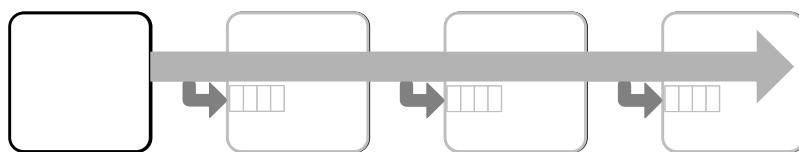
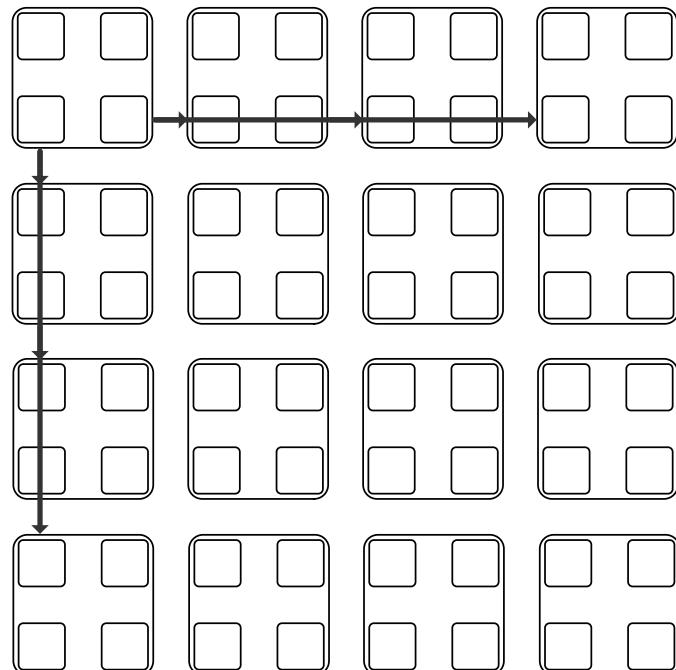
# Multidrop Express Channels (MECS)

---



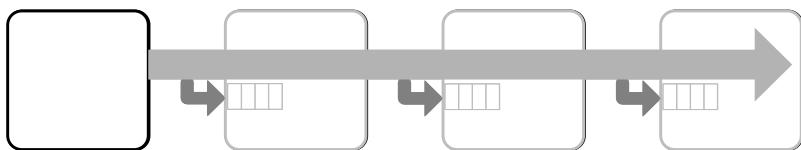
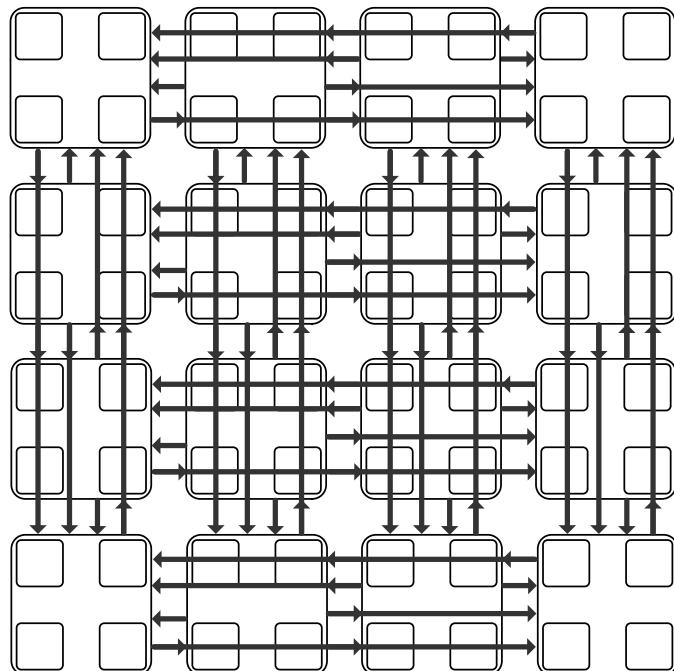
# Multidrop Express Channels (MECS)

---

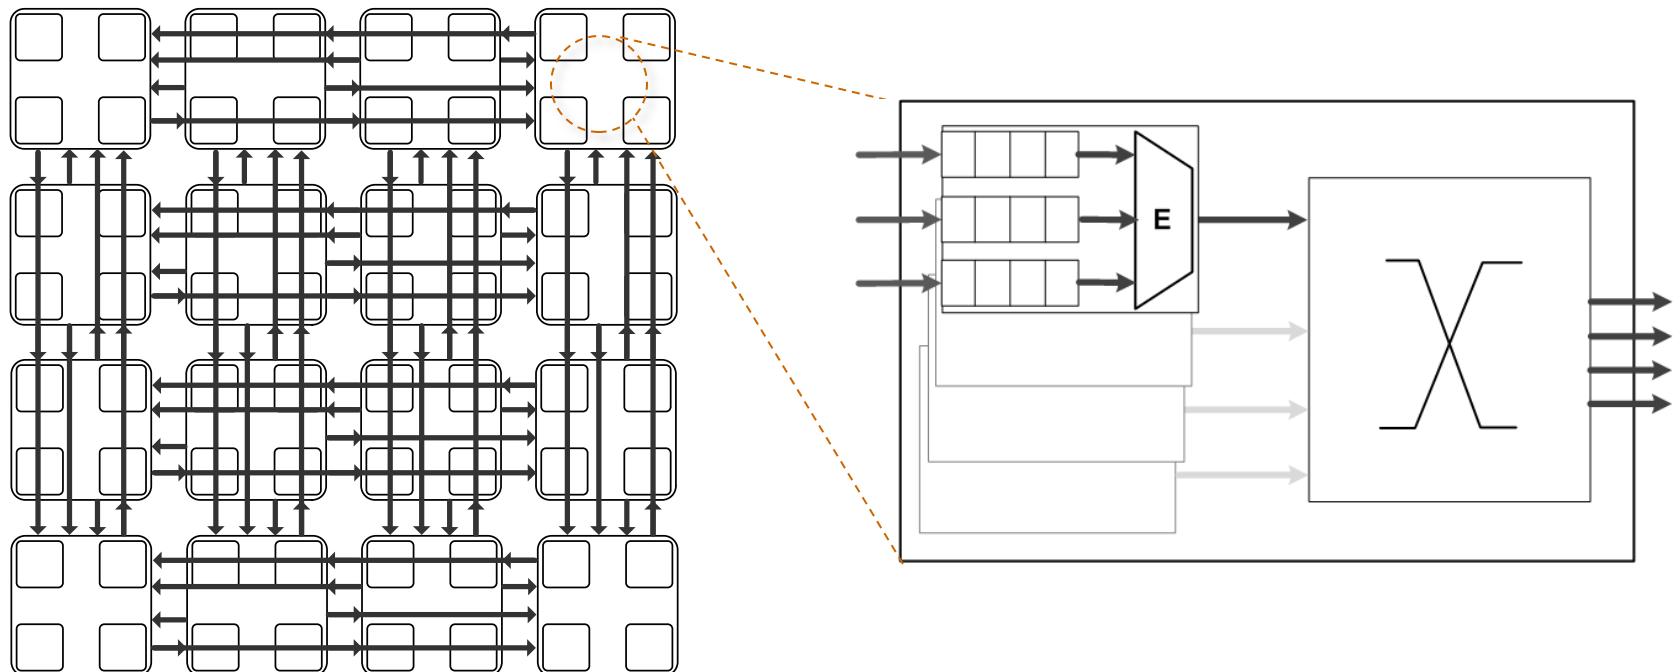


# Multidrop Express Channels (MECS)

---

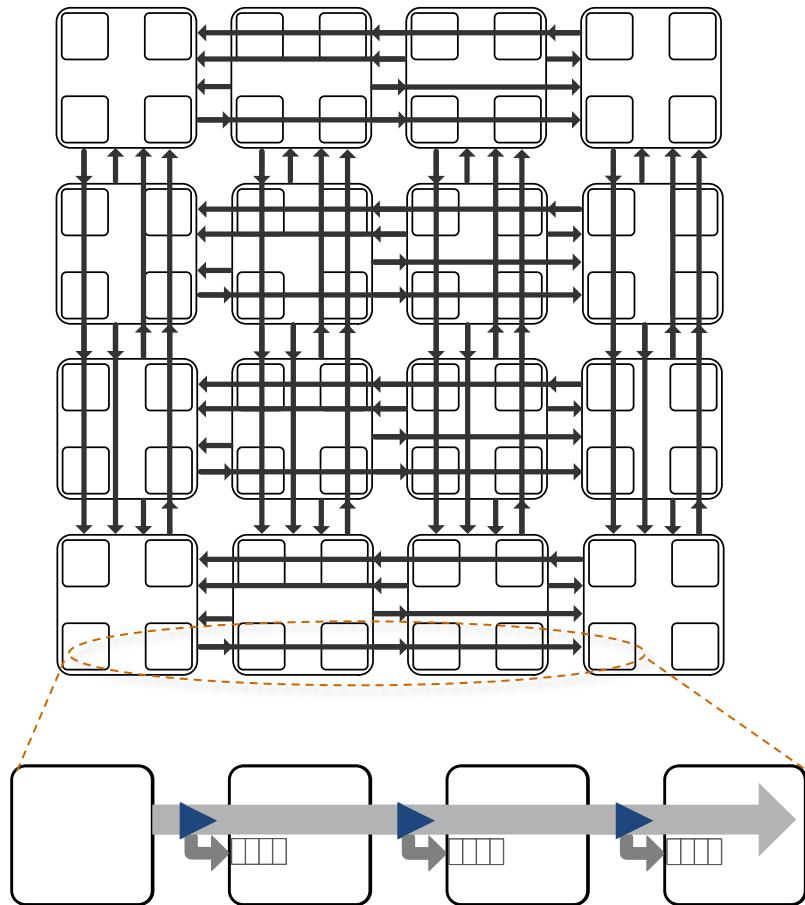


# Multidrop Express Channels (MECS)

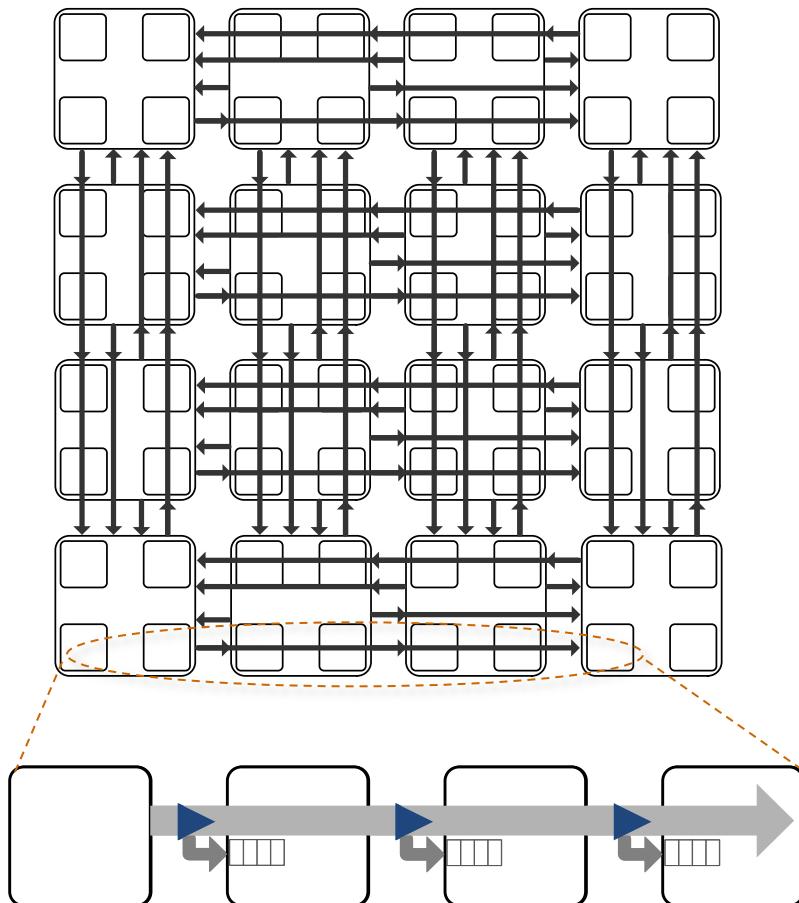


# Multidrop Express Channels (MECS)

---



# Multidrop Express Channels (MECS)



## □ Pros

- One-to-many topology
- Low diameter: 2 hops
- $k$  channels row/column
- Asymmetric

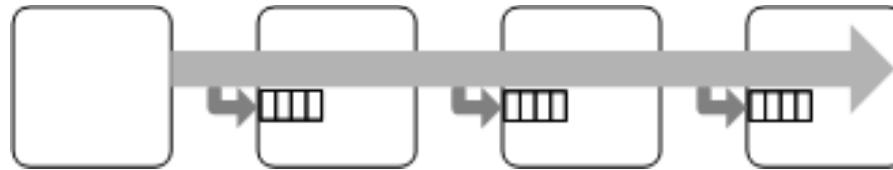
## □ Cons

- Asymmetric
- Increased control (arbitration) complexity

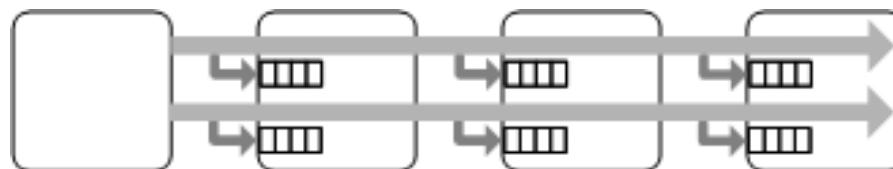
# Partitioning: a GEC Example

---

**MECS**



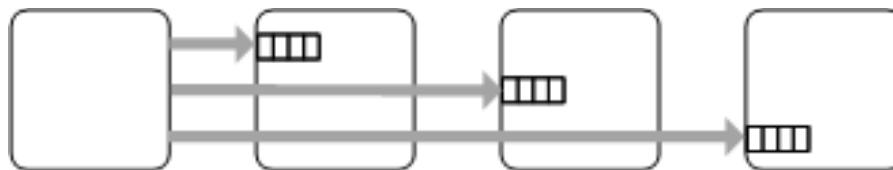
**MECS-X2**



**Partitioned  
MECS**



**Flattened  
Butterfly**



# Analytical Comparison

---

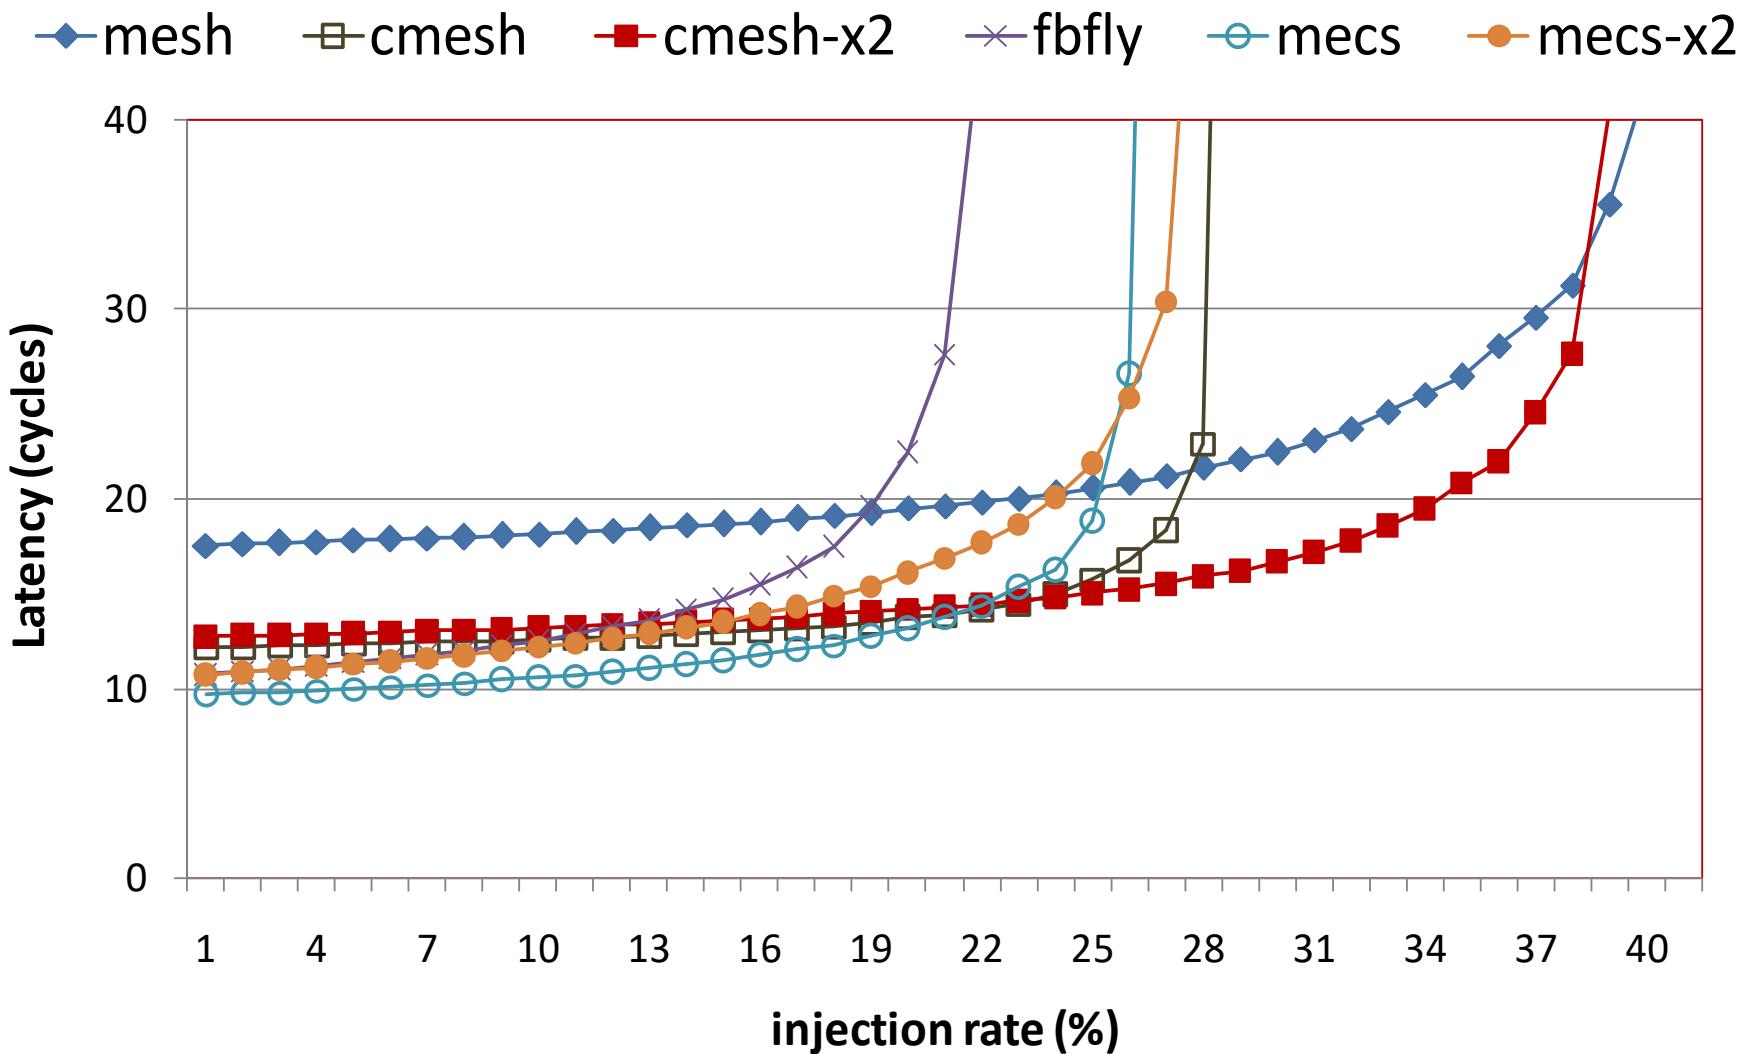
	CMesh		FBfly		MECS	
<b>Network Size</b>	64	256	64	256	64	256
<b>Radix (conctr' d)</b>	4	8	4	8	4	8
<b>Diameter</b>	6	14	2	2	2	2
<b>Channel count</b>	2	2	8	32	4	8
<b>Channel width</b>	576	1152	144	72	288	288
<b>Router inputs</b>	4	4	6	14	6	14
<b>Router outputs</b>	4	4	6	14	4	4

# Experimental Methodology

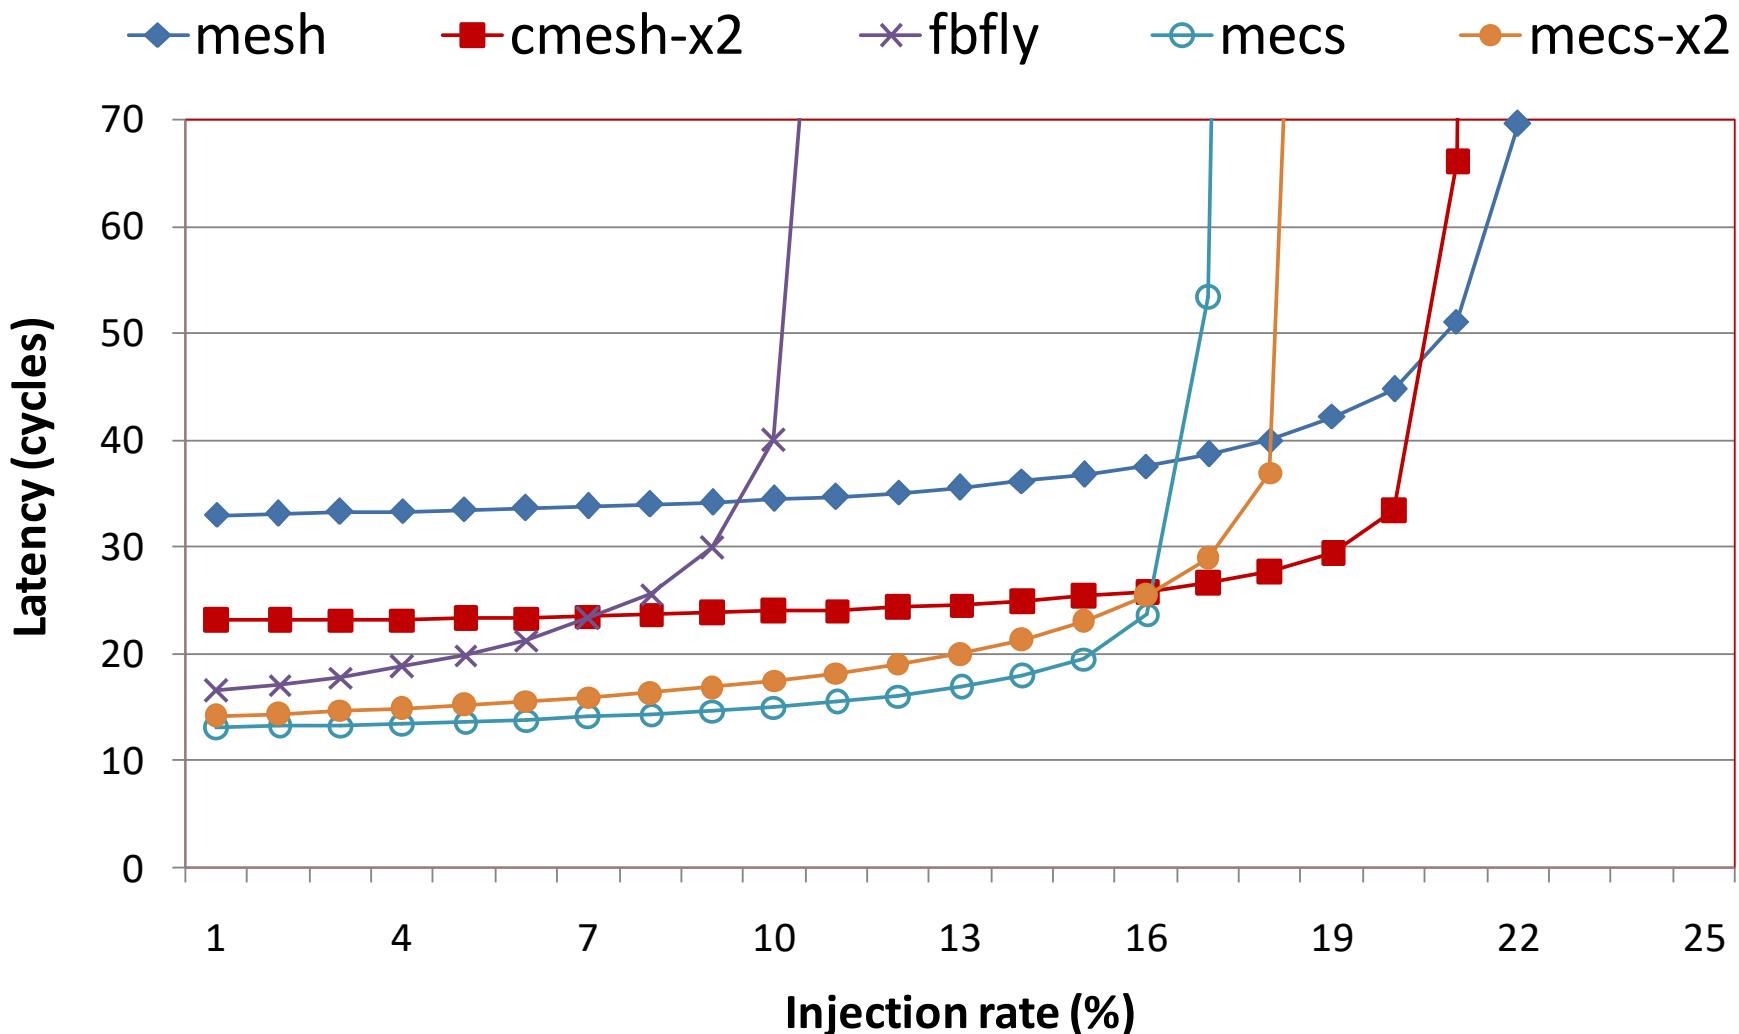
---

<b>Topologies</b>	Mesh, CMesh, CMesh-X2, FBFly, MECS, MECS-X2
<b>Network sizes</b>	64 & 256 terminals
<b>Routing</b>	DOR, adaptive
<b>Messages</b>	64 & 576 bits
<b>Synthetic traffic</b>	Uniform random, bit complement, transpose, self-similar
<b>PARSEC benchmarks</b>	Blackscholes, Bodytrack, Canneal, Ferret, Fluidanimate, Freqmine, Vip, x264
<b>Full-system config</b>	M5 simulator, Alpha ISA, 64 OOO cores
<b>Energy evaluation</b>	Orion + CACTI 6

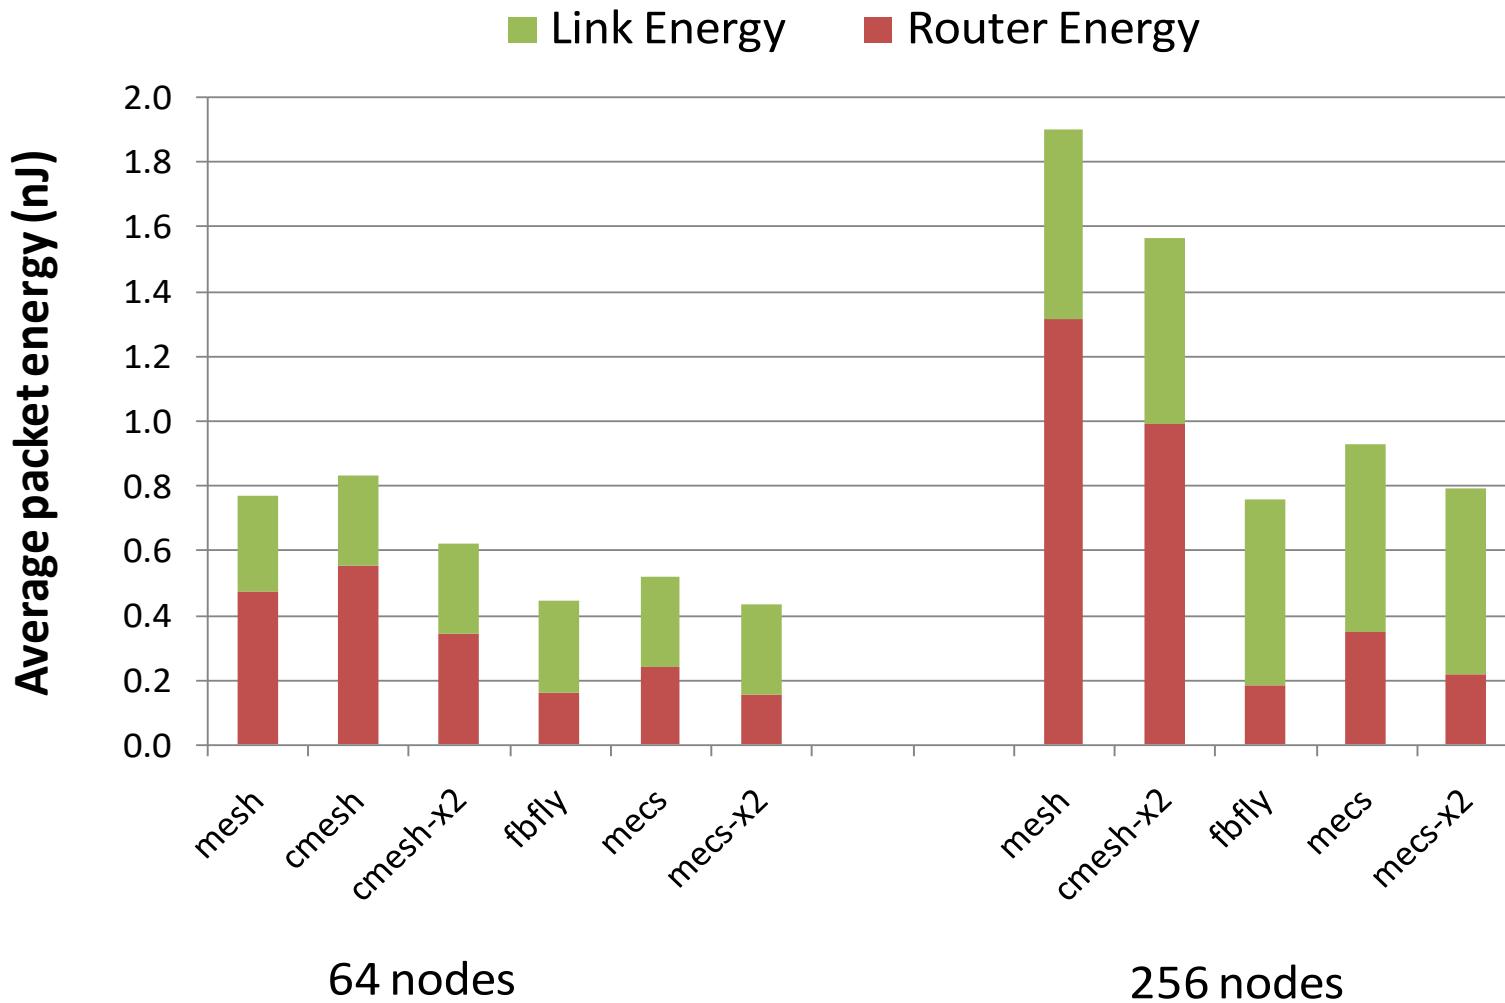
# 64 nodes: Uniform Random



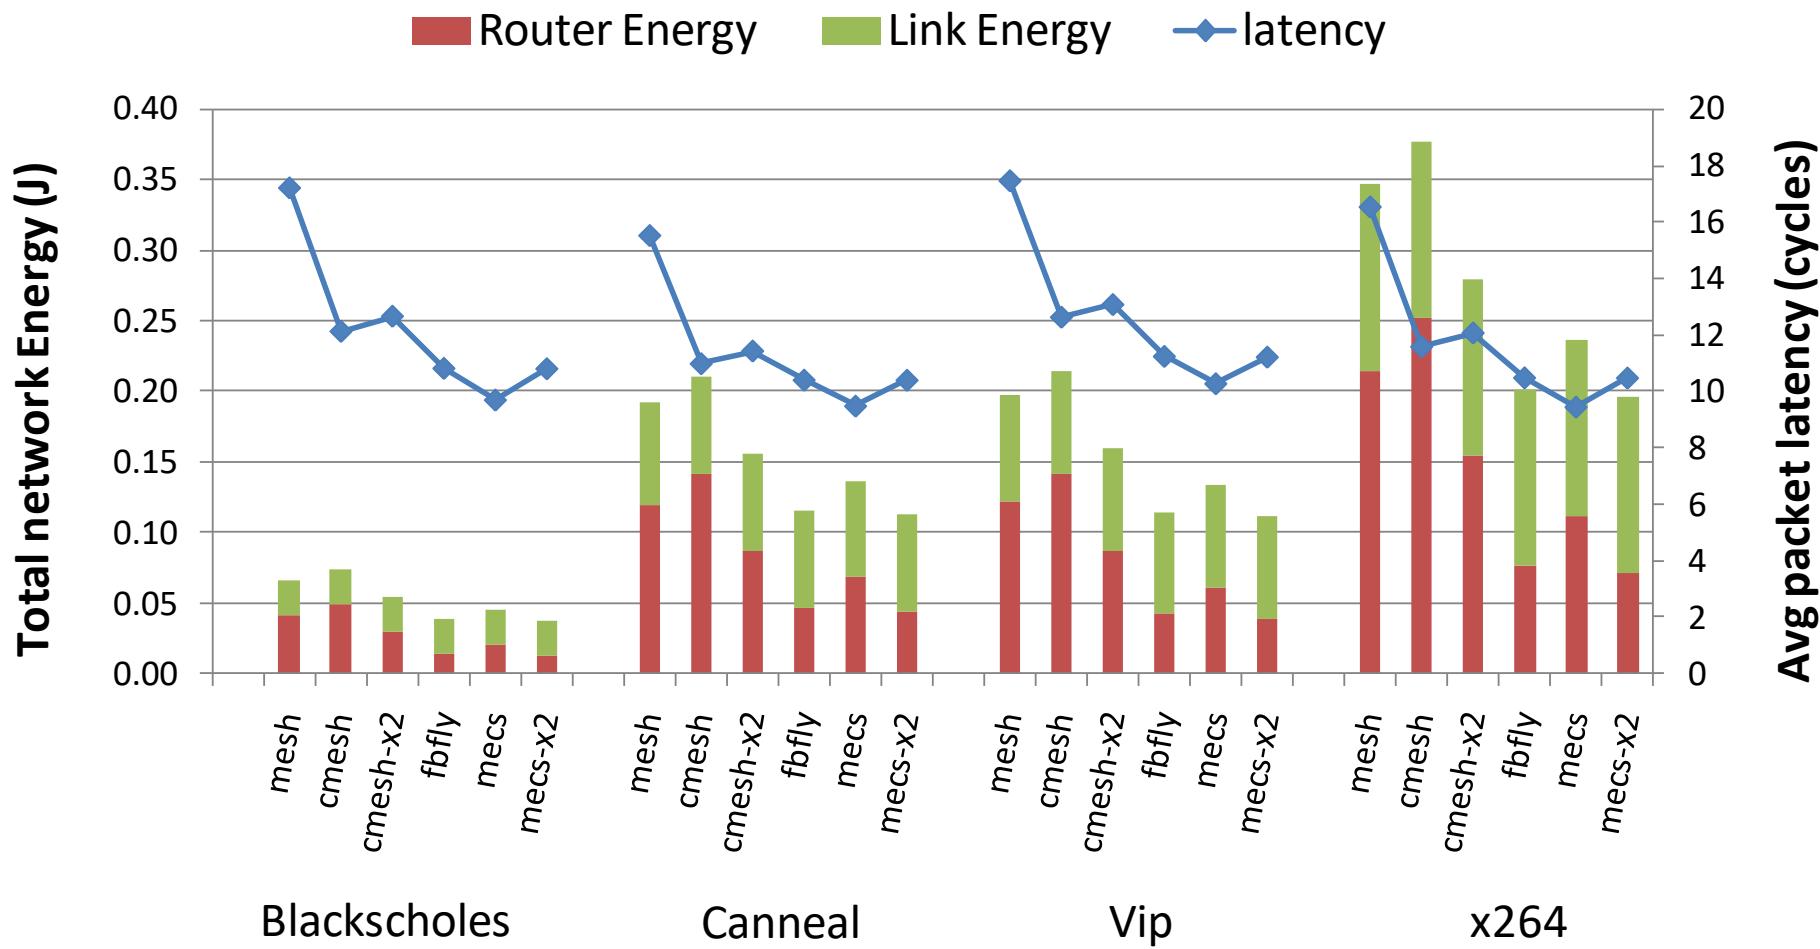
# 256 nodes: Uniform Random



# Energy (100K pkts, Uniform Random)



# 64 Nodes: PARSEC



# Summary

---

## ❑ MECS

- A new one-to-many topology
- Good fit for planar substrates
- Excellent connectivity
- Effective wire utilization

## ❑ Generalized Express Cubes

- Framework & taxonomy for NOC topologies
- Extension of the k-ary n-cube model
- Useful for understanding and exploring on-chip interconnect options
- Future: expand & formalize

# Scalability: Express Cube Topologies

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"Express Cube Topologies for On-Chip Interconnects"**  
*Proceedings of the 15th International Symposium on High-Performance Computer Architecture (HPCA)*, pages 163-174,  
Raleigh, NC, February 2009. [Slides \(ppt\)](#)

## Express Cube Topologies for On-Chip Interconnects

Boris Grot

Joel Hestness

Stephen W. Keckler

Onur Mutlu<sup>†</sup>

Department of Computer Sciences  
The University of Texas at Austin  
`{bgrot, hestness, skeckler}@cs.utexas.edu`

<sup>†</sup>Computer Architecture Laboratory (CALCM)  
Carnegie Mellon University  
`onur@cmu.edu`

# Interconnect Readings

# Application-Aware Prioritization in NoCs

---

- Reetuparna Das, Onur Mutlu, Thomas Moscibroda, and Chita R. Das,  
**"Application-Aware Prioritization Mechanisms for On-Chip Networks"**

*Proceedings of the 42nd International Symposium on Microarchitecture (MICRO), pages 280-291, New York, NY, December 2009.* Slides (pptx)

## Application-Aware Prioritization Mechanisms for On-Chip Networks

Reetuparna Das<sup>§</sup>   Onur Mutlu<sup>†</sup>   Thomas Moscibroda<sup>‡</sup>   Chita R. Das<sup>§</sup>

§Pennsylvania State University  
{rdas,das}@cse.psu.edu

†Carnegie Mellon University  
onur@cmu.edu

‡Microsoft Research  
moscitho@microsoft.com

# Slack-Based Packet Scheduling

---

- Reetuparna Das, Onur Mutlu, Thomas Moscibroda, and Chita R. Das,  
**"Aergia: Exploiting Packet Latency Slack in On-Chip Networks"**  
*Proceedings of the 37th International Symposium on Computer Architecture (ISCA)*, pages 106-116, Saint-Malo, France, June 2010. [Slides \(pptx\)](#)  
**One of the 11 computer architecture papers of 2010 selected as Top Picks by IEEE Micro.**

## Aérgia: Exploiting Packet Latency Slack in On-Chip Networks

Reetuparna Das<sup>§</sup>   Onur Mutlu<sup>†</sup>   Thomas Moscibroda<sup>‡</sup>   Chita R. Das<sup>§</sup>

§Pennsylvania State University  
[{rdas,das}@cse.psu.edu](mailto:{rdas,das}@cse.psu.edu)

†Carnegie Mellon University  
[onur@cmu.edu](mailto:onur@cmu.edu)

‡Microsoft Research  
[moscitho@microsoft.com](mailto:moscitho@microsoft.com)

# Low-Cost QoS in On-Chip Networks (I)

---

- Boris Grot, Stephen W. Keckler, and Onur Mutlu,  
**"Preemptive Virtual Clock: A Flexible, Efficient, and Cost-effective QOS Scheme for Networks-on-Chip"**  
*Proceedings of the 42nd International Symposium on Microarchitecture (MICRO)*, pages 268-279, New York, NY, December 2009. [Slides \(pdf\)](#)

## Preemptive Virtual Clock: A Flexible, Efficient, and Cost-effective QOS Scheme for Networks-on-Chip

Boris Grot

Stephen W. Keckler

Onur Mutlu<sup>†</sup>

Department of Computer Sciences  
The University of Texas at Austin  
{bgrot, skeckler@cs.utexas.edu}

<sup>†</sup>Computer Architecture Laboratory (CALCM)  
Carnegie Mellon University  
onur@cmu.edu

# Low-Cost QoS in On-Chip Networks (II)

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"Kilo-NOC: A Heterogeneous Network-on-Chip Architecture for Scalability and Service Guarantees"**

*Proceedings of the 38th International Symposium on Computer Architecture (ISCA), San Jose, CA, June 2011.* Slides (pptx)

***One of the 12 computer architecture papers of 2011 selected as Top Picks by IEEE Micro.***

## Kilo-NOC: A Heterogeneous Network-on-Chip Architecture for Scalability and Service Guarantees

Boris Grot<sup>1</sup>

bgrot@cs.utexas.edu

Joel Hestness<sup>1</sup>

hestness@cs.utexas.edu

Stephen W. Keckler<sup>1,2</sup>

skeckler@nvidia.com

Onur Mutlu<sup>3</sup>

onur@cmu.edu

<sup>1</sup>The University of Texas at Austin  
Austin, TX

<sup>2</sup>NVIDIA  
Santa Clara, CA

<sup>3</sup>Carnegie Mellon University  
Pittsburgh, PA

# Low-Cost QoS in On-Chip Networks (III)

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"A QoS-Enabled On-Die Interconnect Fabric for Kilo-Node Chips"**  
*IEEE Micro*, Special Issue: *Micro's Top Picks from 2011 Computer Architecture Conferences (MICRO TOP PICKS)*, Vol. 32, No. 3, May/June 2012.
- 

## A QoS-ENABLED ON-DIE INTERCONNECT FABRIC FOR KILO-NODE CHIPS

---

TO MEET RAPIDLY GROWING PERFORMANCE DEMANDS AND ENERGY CONSTRAINTS, FUTURE CHIPS WILL LIKELY FEATURE THOUSANDS OF ON-DIE RESOURCES. EXISTING NETWORK-ON-CHIP SOLUTIONS WEREN'T DESIGNED FOR SCALABILITY AND WILL BE UNABLE TO MEET FUTURE INTERCONNECT DEMANDS. A HYBRID NETWORK-ON-CHIP ARCHITECTURE CALLED KILO-NoC CO-OPTIMIZES TOPOLOGY, FLOW CONTROL, AND QUALITY OF SERVICE TO ACHIEVE SIGNIFICANT GAINS IN EFFICIENCY.

---

# Throttling Based Fairness in NoCs

---

- Kevin Chang, Rachata Ausavarungnirun, Chris Fallin, and Onur Mutlu,  
**"HAT: Heterogeneous Adaptive Throttling for On-Chip Networks"**  
*Proceedings of the 24th International Symposium on Computer Architecture and High Performance Computing (SBAC-PAD)*, New York, NY, October 2012. Slides (pptx) (pdf)

## HAT: Heterogeneous Adaptive Throttling for On-Chip Networks

Kevin Kai-Wei Chang, Rachata Ausavarungnirun, Chris Fallin, Onur Mutlu  
Carnegie Mellon University  
`{kevincha, rachata, cfallin, onur}@cmu.edu`

# Scalability: Express Cube Topologies

---

- Boris Grot, Joel Hestness, Stephen W. Keckler, and Onur Mutlu,  
**"Express Cube Topologies for On-Chip Interconnects"**  
*Proceedings of the 15th International Symposium on High-Performance Computer Architecture (HPCA)*, pages 163-174,  
Raleigh, NC, February 2009. [Slides \(ppt\)](#)

## Express Cube Topologies for On-Chip Interconnects

Boris Grot

Joel Hestness

Stephen W. Keckler

Onur Mutlu<sup>†</sup>

Department of Computer Sciences  
The University of Texas at Austin  
`{bgrot, hestness, skeckler}@cs.utexas.edu`

<sup>†</sup>Computer Architecture Laboratory (CALCM)  
Carnegie Mellon University  
`onur@cmu.edu`

# Scalability: Slim NoC

---

- Maciej Besta, Syed Minhaj Hassan, Sudhakar Yalamanchili, Rachata Ausavarungnirun, Onur Mutlu, Torsten Hoefler,

**"Slim NoC: A Low-Diameter On-Chip Network Topology for High Energy Efficiency and Scalability"**

*Proceedings of the 23rd International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), Williamsburg, VA, USA, March 2018.*

[Slides (pptx) (pdf)] [Lightning Session Slides (pptx) (pdf)]

[Poster (pdf)]

## Slim NoC: A Low-Diameter On-Chip Network Topology for High Energy Efficiency and Scalability

Maciej Besta<sup>1</sup>      Syed Minhaj Hassan<sup>2</sup>      Sudhakar Yalamanchili<sup>2</sup>  
Rachata Ausavarungnirun<sup>3</sup>      Onur Mutlu<sup>1,3</sup>      Torsten Hoefler<sup>1</sup>

---

<sup>1</sup>ETH Zürich

<sup>2</sup>Georgia Institute of Technology

<sup>3</sup>Carnegie Mellon University

# Bufferless Deflection Routing in NoCs

---

- Thomas Moscibroda and Onur Mutlu,  
**"A Case for Bufferless Routing in On-Chip Networks"**  
*Proceedings of the 36th International Symposium on Computer Architecture (ISCA)*, pages 196-207, Austin, TX, June 2009. Slides (pptx)

## A Case for Bufferless Routing in On-Chip Networks

Thomas Moscibroda  
Microsoft Research  
[moscitho@microsoft.com](mailto:moscitho@microsoft.com)

Onur Mutlu  
Carnegie Mellon University  
[onur@cmu.edu](mailto:onur@cmu.edu)

# Minimally-Buffered Deflection Routing

---

- Chris Fallin, Greg Nazario, Xiangyao Yu, Kevin Chang, Rachata Ausavarungnirun, and Onur Mutlu,

**"MinBD: Minimally-Buffered Deflection Routing for Energy-Efficient Interconnect"**

*Proceedings of the 6th ACM/IEEE International Symposium on Networks on Chip (NOCS), Lyngby, Denmark, May 2012.* [Slides \(pptx\)](#) [\(pdf\)](#)

***One of the five papers nominated for the Best Paper Award by the Program Committee.***

## MinBD: Minimally-Buffered Deflection Routing for Energy-Efficient Interconnect

Chris Fallin, Greg Nazario, Xiangyao Yu<sup>†</sup>, Kevin Chang, Rachata Ausavarungnirun, Onur Mutlu

Carnegie Mellon University  
`{cfallin,gnazario,kevincha,rachata,onur}@cmu.edu`

<sup>†</sup>Tsinghua University & Carnegie Mellon University  
`yxythu@gmail.com`

# “Bufferless” Hierarchical Rings

---

- Rachata Ausavarungnirun, Chris Fallin, Xiangyao Yu, Kevin Chang, Greg Nazario, Reetuparna Das, Gabriel Loh, and Onur Mutlu,  
**["Design and Evaluation of Hierarchical Rings with Deflection Routing"](#)**  
*Proceedings of the [26th International Symposium on Computer Architecture and High Performance Computing \(SBAC-PAD\)](#)*, Paris, France, October 2014. [[Slides \(pptx\)](#)] [[\(pdf\)](#)] [[Source Code](#)]
- Describes the design and implementation of a mostly-bufferless hierarchical ring

## Design and Evaluation of Hierarchical Rings with Deflection Routing

Rachata Ausavarungnirun    Chris Fallin    Xiangyao Yu†    Kevin Kai-Wei Chang  
Greg Nazario    Reetuparna Das§    Gabriel H. Loh‡    Onur Mutlu

Carnegie Mellon University    §University of Michigan    †MIT    ‡Advanced Micro Devices, Inc.

---

# “Bufferless” Hierarchical Rings (II)

---

- Rachata Ausavarungnirun, Chris Fallin, Xiangyao Yu, Kevin Chang, Greg Nazario, Reetuparna Das, Gabriel Loh, and Onur Mutlu,  
**"A Case for Hierarchical Rings with Deflection Routing: An Energy-Efficient On-Chip Communication Substrate"**  
*Parallel Computing (PARCO)*, 2016.
  - [arXiv.org version](#), February 2016.

Achieving both High Energy Efficiency  
and High Performance in On-Chip Communication  
using Hierarchical Rings with Deflection Routing

Rachata Ausavarungnirun    Chris Fallin    Xiangyao Yu†    Kevin Kai-Wei Chang  
Greg Nazario    Reetuparna Das§    Gabriel H. Loh‡    Onur Mutlu  
Carnegie Mellon University    §University of Michigan    †MIT    ‡AMD

# A Review of Bufferless Interconnects

---

- Chris Fallin, Greg Nazario, Xiangyao Yu, Kevin Chang, Rachata Ausavarungnirun, and Onur Mutlu,  
**"Bufferless and Minimally-Buffered Deflection Routing"**  
*Invited Book Chapter in Routing Algorithms in Networks-on-Chip, pp. 241-275, Springer, 2014.*

## Chapter 1 Bufferless and Minimally-Buffered Deflection Routing

Chris Fallin, Greg Nazario, Xiangyao Yu, Kevin Chang, Rachata Ausavarungnirun, Onur Mutlu

# Summary of Eight Years of Research

---

## Energy-Efficient Deflection-based On-chip Networks: Topology, Routing, Flow Control

Rachata Ausavarungnirun<sup>b</sup>, Onur Mutlu<sup>a</sup>

*SAFARI Research Group*

<sup>a</sup>*ETH Zürich*

<sup>b</sup>*King Mongkut's University of Technology North Bangkok*

---

### Abstract

As the number of cores scales to tens and hundreds, the energy consumption of routers across various types of on-chip networks in chip multiprocessors (CMPs) increases significantly. A major source of this energy consumption comes from the input buffers inside Network-on-Chip (NoC) routers, which are traditionally designed to maximize performance. To mitigate this high energy cost, many works propose bufferless router designs that utilize deflection routing to resolve port contention. While this approach is able to maintain high performance relative to its buffered counterparts at low network traffic, the bufferless router design suffers performance degradation under high network load.

In order to maintain high performance and energy efficiency under *both* low and high network loads, this chapter discusses critical drawbacks of traditional bufferless designs and describes recent research works focusing on two major modifications to improve the overall performance of the traditional bufferless network-on-chip design. The first modification is a minimally-buffered design that introduces limited buffering inside critical parts of the on-chip network in order to reduce the number of deflections. The second modification is a hierarchical bufferless interconnect design that aims to further improve performance by limiting the number of hops each packet needs to travel while in the network. In both approaches, we discuss design tradeoffs and provide evaluation results based on common CMP configurations with various network topologies to show the effectiveness of each proposal.

*Keywords:* network-on-chip, deflection routing, topology, bufferless router, energy efficiency, high-performance computing, computer architecture, emerging technologies, latency, low-latency computing

---

# On-Chip vs. Off-Chip Congestion Control

---

- George Nychis, Chris Fallin, Thomas Moscibroda, Onur Mutlu, and Srinivasan Seshan,  
**"On-Chip Networks from a Networking Perspective:  
Congestion and Scalability in Many-core Interconnects"**  
*Proceedings of the 2012 ACM SIGCOMM  
Conference (SIGCOMM), Helsinki, Finland, August 2012.* Slides  
(pptx)

## On-Chip Networks from a Networking Perspective: Congestion and Scalability in Many-Core Interconnects

George Nychis<sup>†</sup>, Chris Fallin<sup>†</sup>, Thomas Moscibroda<sup>§</sup>, Onur Mutlu<sup>†</sup>, Srinivasan Seshan<sup>†</sup>

<sup>†</sup> Carnegie Mellon University  
[{gnychis,cfallin,onur,srini}@cmu.edu](mailto:{gnychis,cfallin,onur,srini}@cmu.edu)

<sup>§</sup> Microsoft Research Asia  
[moscitho@microsoft.com](mailto:moscitho@microsoft.com)

# On-Chip vs. Off-Chip Congestion Control

---

- George Nychis, Chris Fallin, Thomas Moscibroda, and Onur Mutlu,  
**"Next Generation On-Chip Networks: What Kind of Congestion Control Do We Need?"**

*Proceedings of the 9th ACM Workshop on Hot Topics in Networks (HOTNETS), Monterey, CA, October 2010.* Slides (ppt) (key)

## Next Generation On-Chip Networks: What Kind of Congestion Control Do We Need?

George Nychis†, Chris Fallin†, Thomas Moscibroda§, Onur Mutlu†

†Carnegie Mellon University

{gnychis,cfallin,onur}@cmu.edu

§Microsoft Research  
moscitho@microsoft.com

# Summary of Study [SIGCOMM 2012]

---

- Highlighted a traditional networking problem in a new context
  - Unique design requires novel solution
- Showed **congestion limits efficiency and scalability**, and that **self-throttling nature of cores prevents congestion collapse**
- Showed **on-chip congestion control requires application-awareness**
- Our **application-aware congestion controller** provided:
  - A more efficient network-layer (reduced latency)
  - Improvements in system throughput (by 27%)
  - Effectively scale the CMP (shown for up to 4096 cores)

# Slowdown Estimation in NoCs

---

- Xiyue Xiang, Saugata Ghose, Onur Mutlu, and Nian-Feng Tzeng,  
**"A Model for Application Slowdown Estimation in On-Chip Networks and Its Use for Improving System Fairness and Performance"**

*Proceedings of the 34th IEEE International Conference on Computer Design (ICCD), Phoenix, AZ, USA, October 2016.*  
[Slides (pptx) (pdf)]

## A Model for Application Slowdown Estimation in On-Chip Networks and Its Use for Improving System Fairness and Performance

Xiyue Xiang<sup>†</sup>

<sup>†</sup>*University of Louisiana at Lafayette*

Saugata Ghose<sup>‡</sup>

<sup>‡</sup>*Carnegie Mellon University*

Onur Mutlu<sup>§‡</sup>

Nian-Feng Tzeng<sup>†</sup>

<sup>§</sup>*ETH Zürich*

# Handling Multicast and Hotspot Issues

---

- Xiyue Xiang, Wentao Shi, Saugata Ghose, Lu Peng, Onur Mutlu, and Nian-Feng Tzeng,

## **"Carpool: A Bufferless On-Chip Network Supporting Adaptive Multicast and Hotspot Alleviation"**

*Proceedings of the International Conference on Supercomputing (ICS), Chicago, IL, USA, June 2017.*

[[Slides \(pptx\)](#) ([pdf](#))]

## **Carpool: A Bufferless On-Chip Network Supporting Adaptive Multicast and Hotspot Alleviation**

Xiyue Xiang<sup>†</sup> Wentao Shi<sup>★</sup> Saugata Ghose<sup>‡</sup> Lu Peng<sup>★</sup> Onur Mutlu<sup>§‡</sup> Nian-Feng Tzeng<sup>†</sup>  
<sup>†</sup>University of Louisiana at Lafayette    <sup>★</sup>Louisiana State University    <sup>‡</sup>Carnegie Mellon University    <sup>§</sup>ETH Zürich

# Heterogeneous Networks

---

- Asit K. Mishra, Onur Mutlu, and Chita R. Das,  
**"A Heterogeneous Multiple Network-on-Chip Design: An Application-Aware Approach"**  
*Proceedings of the 50th Design Automation Conference (DAC),*  
Austin, TX, June 2013. [Slides \(pptx\)](#) [Slides \(pdf\)](#)

## A Heterogeneous Multiple Network-On-Chip Design: An Application-Aware Approach

Asit K. Mishra  
Intel Corporation  
Hillsboro, OR 97124, USA  
asit.k.mishra@intel.com

Onur Mutlu  
Carnegie Mellon University  
Pittsburgh, PA 15213, USA  
onur@cmu.edu

Chita R. Das  
The Pennsylvania State University  
University Park, PA 16802, USA  
das@cse.psu.edu

# More Readings

---

- Studies of congestion and congestion control in on-chip vs. internet-like networks
- George Nychis, Chris Fallin, Thomas Moscibroda, Onur Mutlu, and Srinivasan Seshan,  
**"On-Chip Networks from a Networking Perspective: Congestion and Scalability in Many-core Interconnects"**  
*Proceedings of the 2012 ACM SIGCOMM Conference (SIGCOMM), Helsinki, Finland, August 2012.* Slides (pptx)
- George Nychis, Chris Fallin, Thomas Moscibroda, and Onur Mutlu,  
**"Next Generation On-Chip Networks: What Kind of Congestion Control Do We Need?"**  
*Proceedings of the 9th ACM Workshop on Hot Topics in Networks (HOTNETS), Monterey, CA, October 2010.* Slides (ppt) (key)

# Source Throttling in Bufferless NoCs

---

- Kevin Chang, Rachata Ausavarungnirun, Chris Fallin, and Onur Mutlu,  
**"HAT: Heterogeneous Adaptive Throttling for On-Chip Networks"**  
*Proceedings of the 24th International Symposium on Computer Architecture  
and High Performance Computing (SBAC-PAD)*, New York, NY, October  
2012. [Slides \(pptx\)](#) [\(pdf\)](#)

## HAT: Heterogeneous Adaptive Throttling for On-Chip Networks

Kevin Kai-Wei Chang, Rachata Ausavarungnirun, Chris Fallin, Onur Mutlu  
Carnegie Mellon University  
`{kevincha, rachata, cfallin, onur}@cmu.edu`

# HAT: Heterogeneous Adaptive Throttling for On-Chip Networks

Kevin Chang, Rachata Ausavarungnirun, Chris Fallin, and Onur Mutlu,

**"HAT: Heterogeneous Adaptive Throttling for On-Chip Networks"**

*Proceedings of the 24th International Symposium on Computer Architecture and High Performance Computing (SBAC-PAD), New York, NY, October 2012.* Slides (pptx) (pdf)

Carnegie Mellon University

**SAFARI**

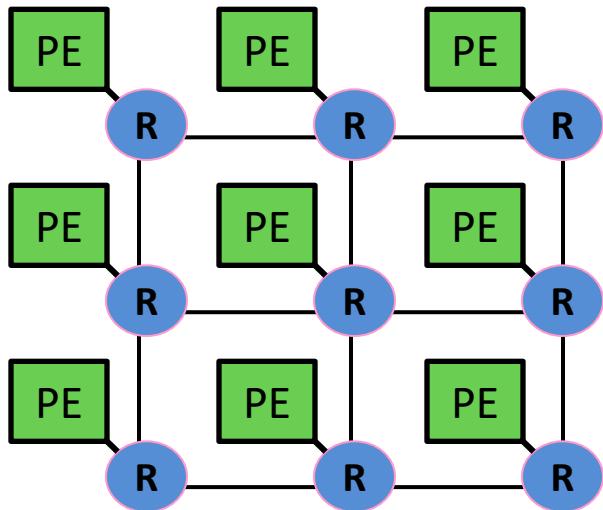
# Executive Summary

- **Problem:** Packets contend in on-chip networks (NoCs), causing congestion, thus reducing performance
- **Observations:**
  - 1) Some applications are more sensitive to network latency than others
  - 2) Applications must be throttled differently to achieve peak performance
- **Key Idea: Heterogeneous Adaptive Throttling (HAT)**
  - 1) Application-aware source throttling
  - 2) Network-load-aware throttling rate adjustment
- **Result:** Improves performance and energy efficiency over state-of-the-art source throttling policies

# Outline

- **Background and Motivation**
- Mechanism
- Prior Works
- Results

# On-Chip Networks



- Connect **cores, caches, memory controllers, etc**
- **Packet switched**
- **2D mesh:** Most commonly used topology
- Primarily serve **cache misses and memory requests**
- **Router designs**
  - Buffered: **Input buffers** to hold contending packets
  - Bufferless: **Misroute (deflect)** contending packets



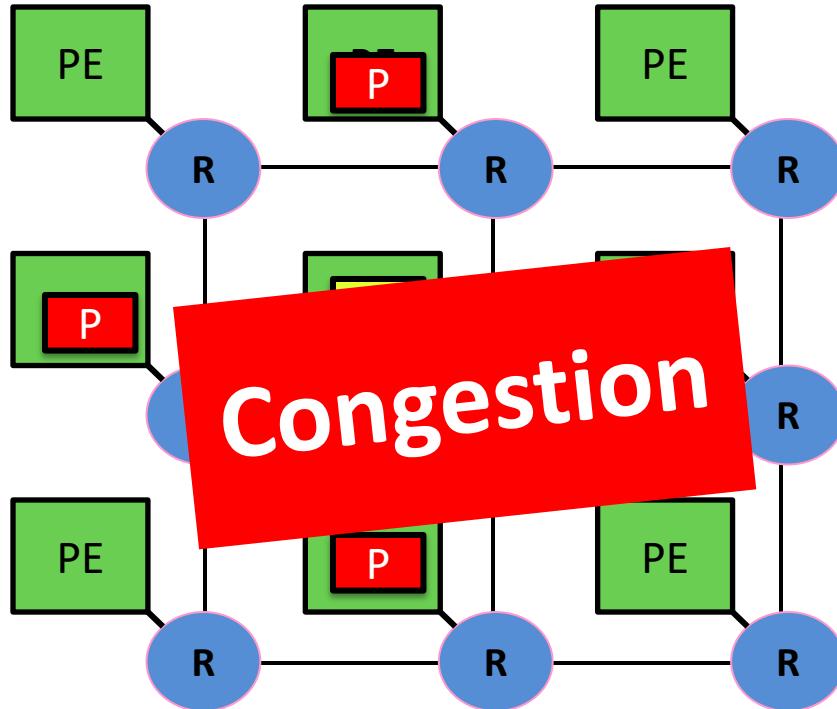
Router



Processing Element

(Cores, L2 Banks, Memory Controllers, etc)

# Network Congestion Reduces Performance



Limited shared resources  
(buffers and links)

- Design constraints: power, chip area, and timing

**Network congestion:**

- ↓ Network throughput
- ↓ Application performance

Router      Packet

Processing Element  
(Cores, L2 Banks, Memory Controllers, etc)

# Goal

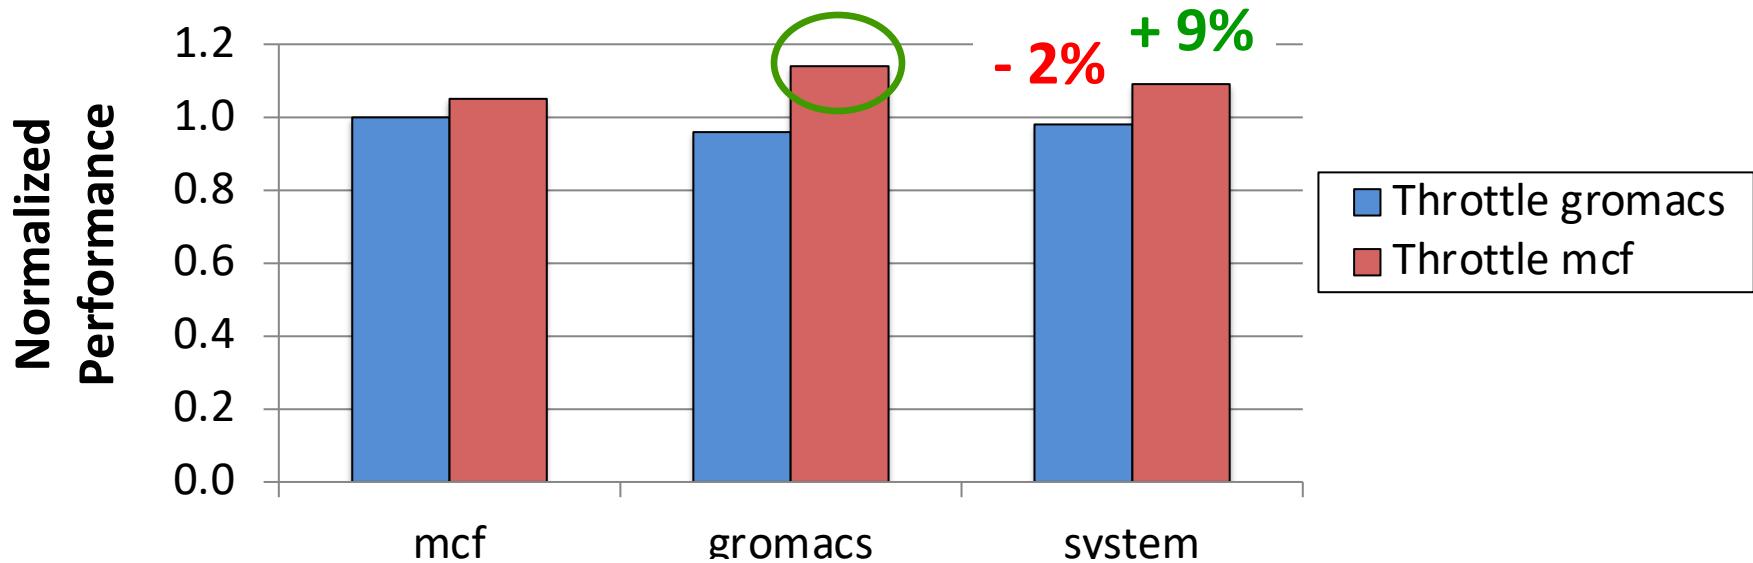
- **Improve performance in a highly congested NoC**
- Reducing network load decreases network congestion, hence improves performance
- **Approach: source throttling to reduce network load**
  - Temporarily delay new traffic injection
- **Naïve mechanism: throttle every single node**

# Key Observation #1

Different applications respond differently to changes in **network latency**

gromacs: network-non-intensive

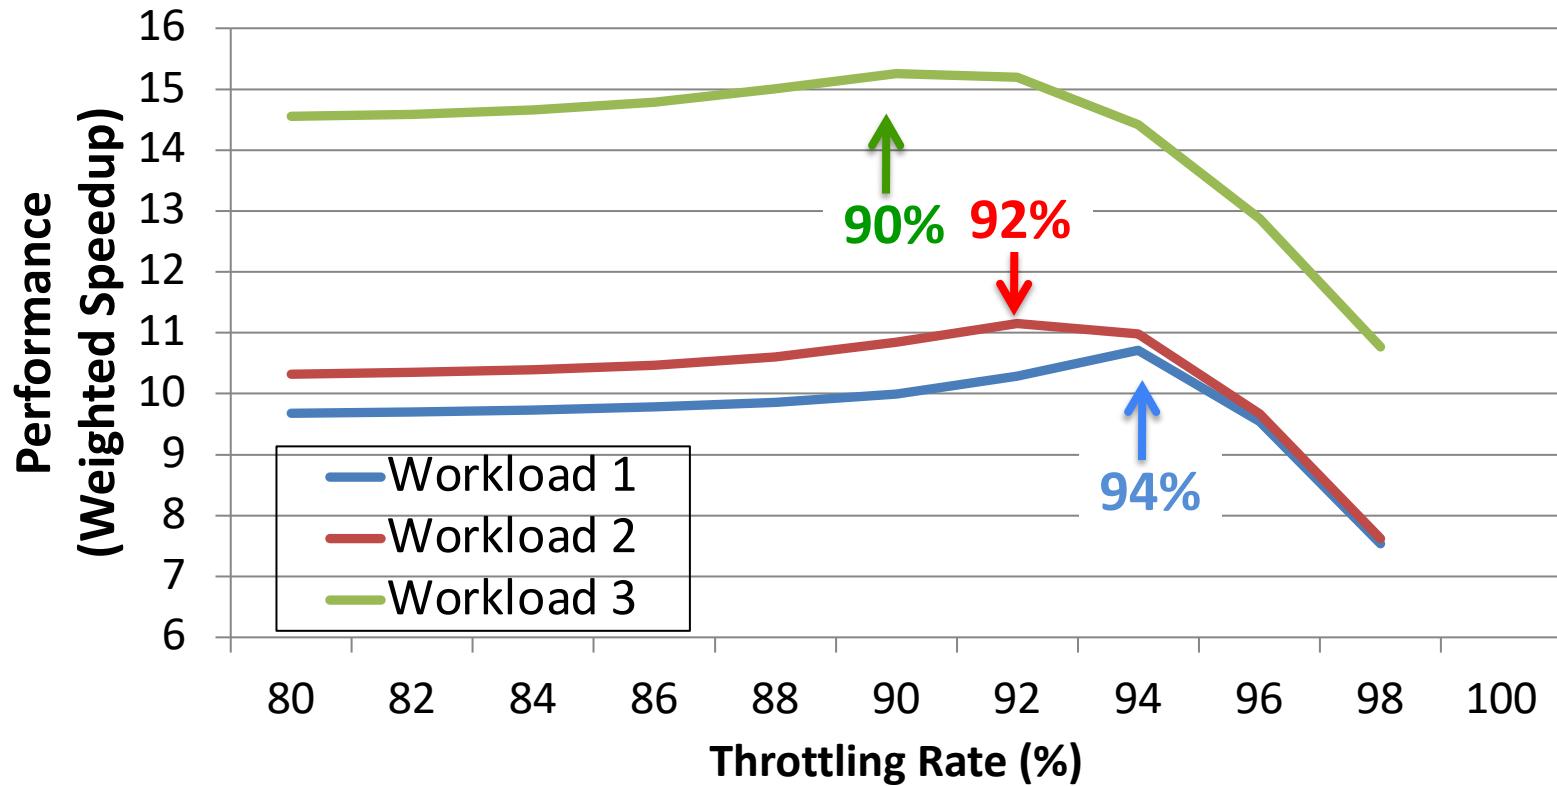
mcf: network-intensive



Throttling **network-intensive** applications benefits system performance more

# Key Observation #2

Different workloads achieve peak performance at different throttling rates



Dynamically adjusting throttling rate yields better performance than a single static rate

# Outline

- Background and Motivation
- **Mechanism**
- Prior Works
- Results

# Heterogeneous Adaptive Throttling (HAT)

## 1. Application-aware throttling:

Throttle **network-intensive** applications that interfere with **network-non-intensive** applications

## 2. Network-load-aware throttling rate adjustment:

Dynamically adjusts throttling rate to adapt to different workloads

# Heterogeneous Adaptive Throttling (HAT)

## 1. Application-aware throttling:

Throttle **network-intensive** applications that interfere with **network-non-intensive** applications

## 2. Network-load-aware throttling rate adjustment:

Dynamically adjusts throttling rate to adapt to different workloads

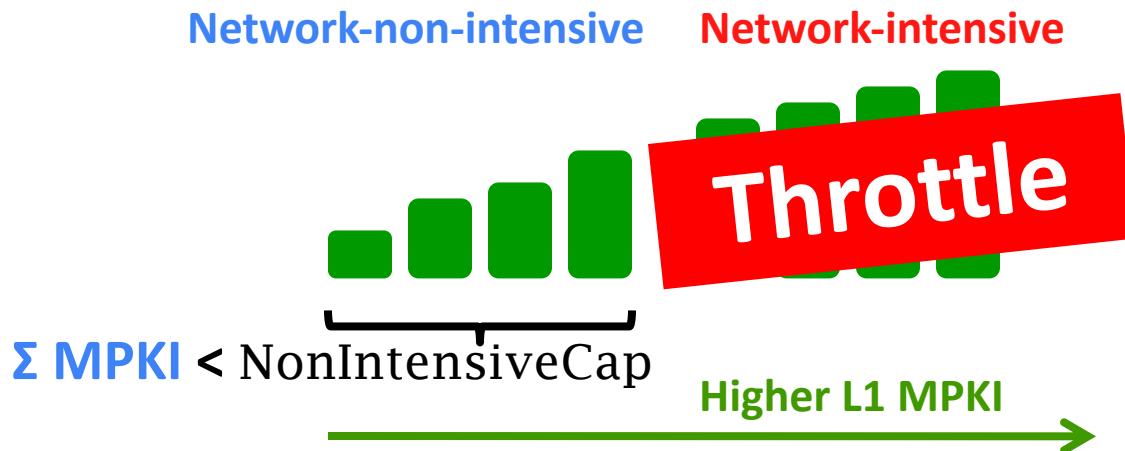
# Application-Aware Throttling

## 1. Measure Network Intensity

Use **L1 MPKI** (misses per thousand instructions) to estimate network intensity

## 2. Classify Application

Sort applications by L1 MPKI



## 3. Throttle network-intensive applications

# Heterogeneous Adaptive Throttling (HAT)

## 1. Application-aware throttling:

Throttle **network-intensive** applications that interfere with **network-non-intensive** applications

## 2. Network-load-aware throttling rate adjustment:

**Dynamically** adjusts throttling rate to adapt to different workloads

# Dynamic Throttling Rate Adjustment

- For a given **network design**, peak performance tends to occur at a **fixed network load point**
- **Dynamically** adjust throttling rate to achieve that network load point

# Dynamic Throttling Rate Adjustment

- **Goal:** maintain network load at a peak performance point

1. **Measure network load**
2. **Compare and adjust throttling rate**

If **network load > peak point**:

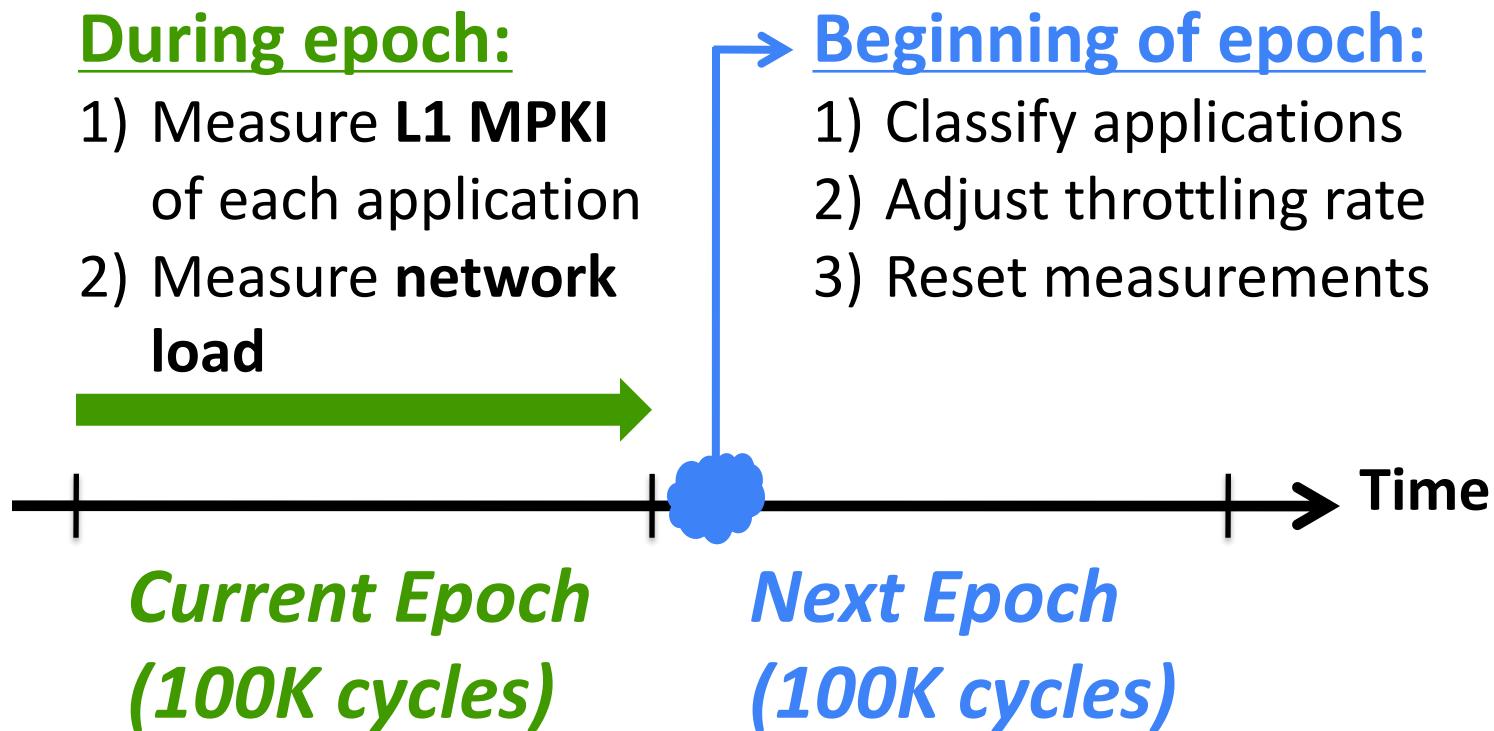
Increase throttling rate

**elif network load ≤ peak point**:

Decrease throttling rate

# Epoch-Based Operation

- Continuous HAT operation is expensive
- **Solution:** performs HAT at epoch granularity



# Outline

- Background and Motivation
- Mechanism
- **Prior Works**
- Results

# Prior Source Throttling Works

- **Source throttling for bufferless NoCs**  
[Nychis+ Hotnets'10, SIGCOMM'12]
  - Application-aware throttling based on starvation rate
  - Does not adaptively adjust throttling rate
  - “Heterogeneous Throttling”
- **Source throttling off-chip buffered networks**  
[Thottethodi+ HPCA'01]
  - Dynamically trigger throttling based on fraction of buffer occupancy
  - Not application-aware: fully block packet injections of every node
  - “Self-tuned Throttling”

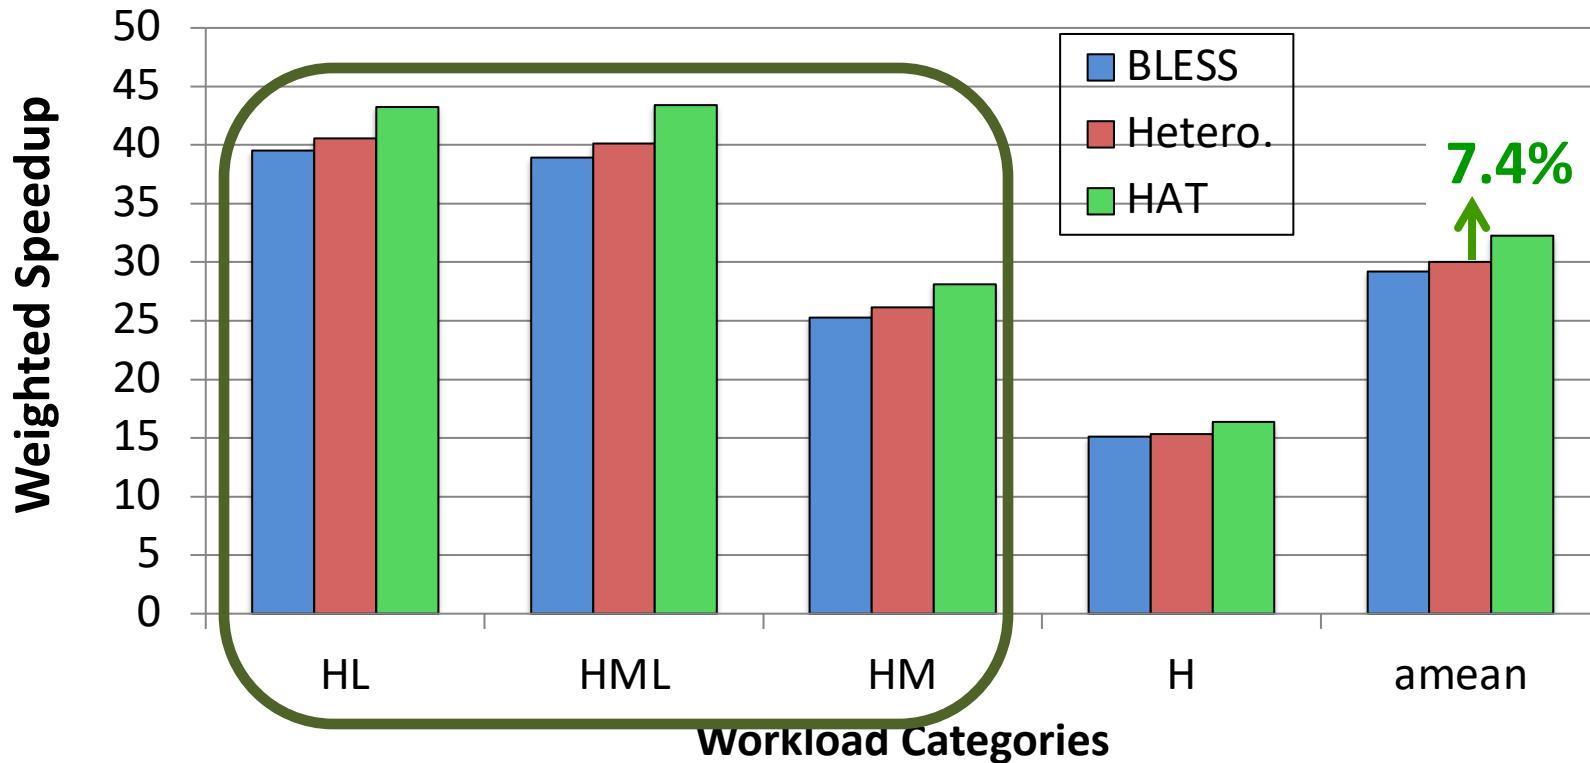
# Outline

- Background and Motivation
- Mechanism
- Prior Works
- **Results**

# Methodology

- **Chip Multiprocessor Simulator**
  - **64-node** multi-core systems with a **2D-mesh topology**
  - Closed-loop core/cache/NoC cycle-level model
  - 64KB L1, perfect L2 (always hits to stress NoC)
- **Router Designs**
  - **Virtual-channel buffered** router: 4 VCs, 4 flits/VC [Dally+ IEEE TPDS'92]
  - **Bufferless deflection** routers: **BLESS** [Moscibroda+ ISCA'09]
- **Workloads**
  - 60 multi-core workloads: SPEC CPU2006 benchmarks
  - Categorized based on their network intensity
    - **Low/Medium/High** intensity categories
- **Metrics:** Weighted Speedup (perf.), perf./Watt (energy eff.), and maximum slowdown (fairness)

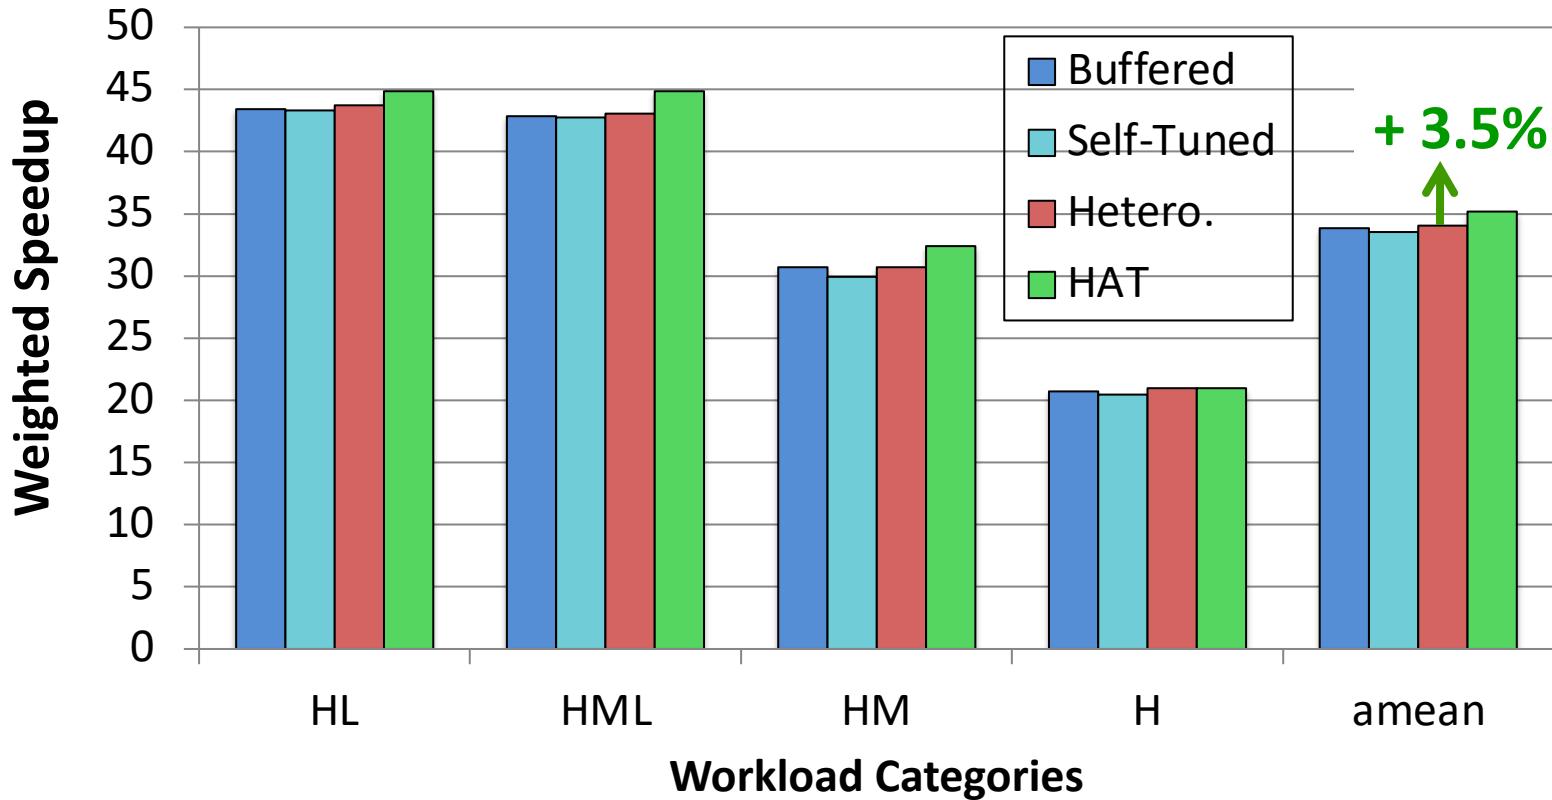
# Performance: Bufferless NoC (BLESS)



**HAT** provides better performance improvement than past work

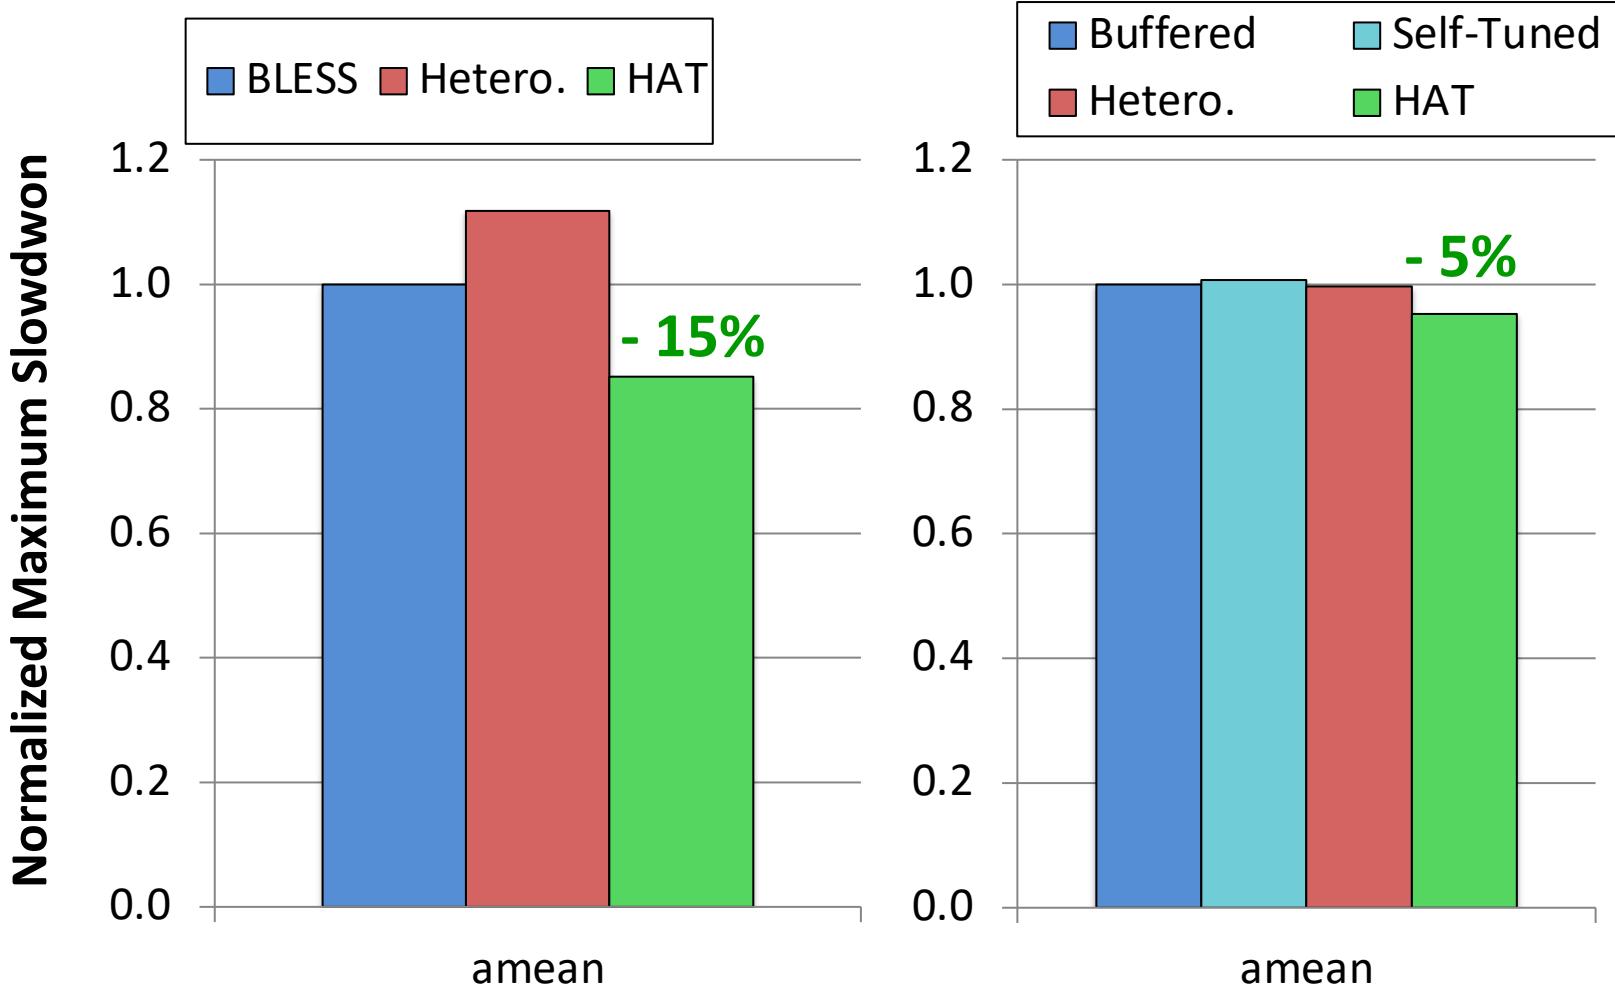
Highest improvement on **heterogeneous** workload mixes  
- L and M are more **sensitive** to network latency

# Performance: Buffered NoC



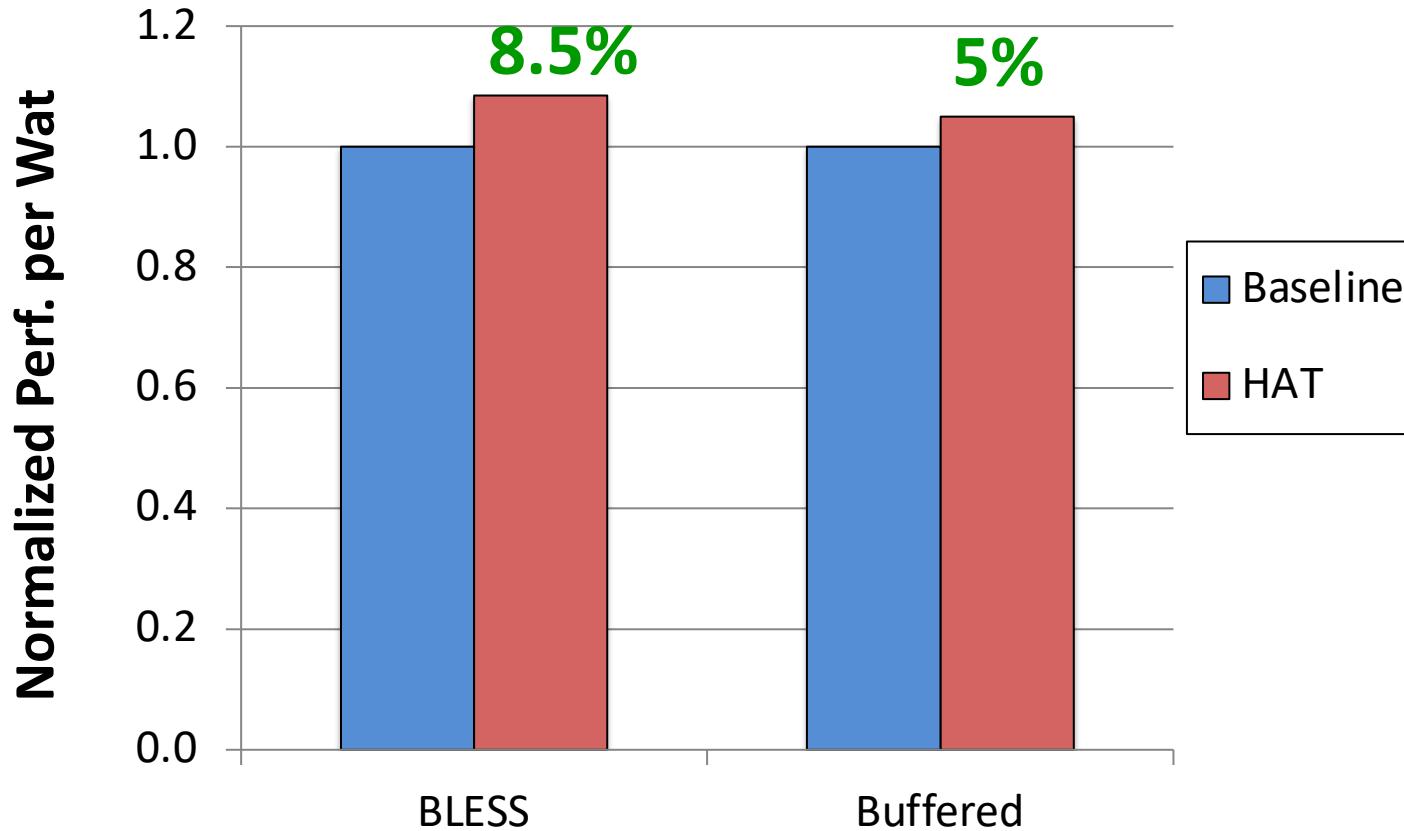
Congestion is much lower in Buffered NoC, but HAT still provides performance benefit

# Application Fairness



**HAT provides better fairness than prior works**

# Network Energy Efficiency



HAT increases energy efficiency by  
reducing congestion

# Other Results in Paper

- Performance on **CHIPPER**
- Performance on **multithreaded** workloads
- Parameters sensitivity sweep of **HAT**

# Conclusion

- **Problem:** Packets contend in on-chip networks (NoCs), causing congestion, thus reducing performance
- **Observations:**
  - 1) Some applications are more sensitive to network latency than others
  - 2) Applications must be throttled differently to achieve peak performance
- **Key Idea: Heterogeneous Adaptive Throttling (HAT)**
  - 1) Application-aware source throttling
  - 2) Network-load-aware throttling rate adjustment
- **Result:** Improves performance and energy efficiency over state-of-the-art source throttling policies

# Source Throttling in Bufferless NoCs

---

- Kevin Chang, Rachata Ausavarungnirun, Chris Fallin, and Onur Mutlu,  
**"HAT: Heterogeneous Adaptive Throttling for On-Chip Networks"**  
*Proceedings of the 24th International Symposium on Computer Architecture  
and High Performance Computing (SBAC-PAD)*, New York, NY, October  
2012. [Slides \(pptx\)](#) [\(pdf\)](#)

## HAT: Heterogeneous Adaptive Throttling for On-Chip Networks

Kevin Kai-Wei Chang, Rachata Ausavarungnirun, Chris Fallin, Onur Mutlu  
Carnegie Mellon University  
`{kevincha, rachata, cfallin, onur}@cmu.edu`

# Slack-Driven Packet Scheduling

Reetuparna Das, Onur Mutlu, Thomas Moscibroda, and Chita R. Das,  
**"Aergia: Exploiting Packet Latency Slack in On-Chip Networks"**

*Proceedings of the 37th International Symposium on Computer Architecture (ISCA)*, pages 106-116, Saint-Malo, France, June 2010. Slides (pptx)

# Packet Scheduling in NoC

---

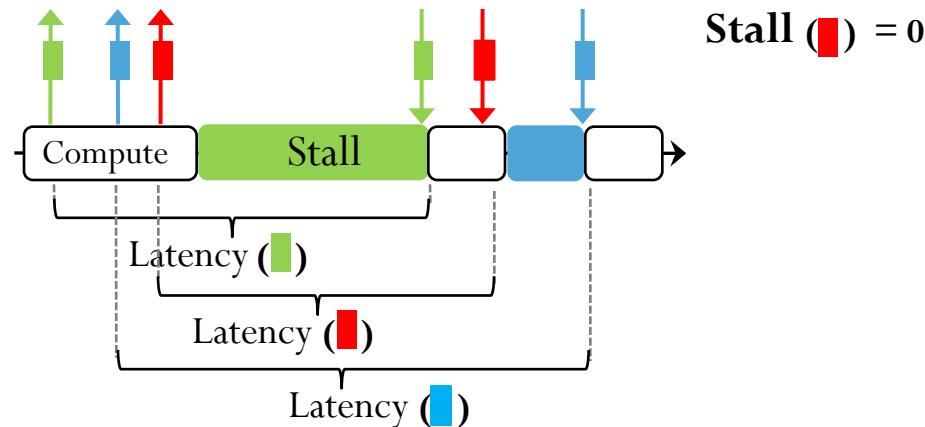
- Existing scheduling policies
  - Round robin
  - Age
- Problem
  - Treat all packets equally
  - Application-oblivious

All packets are not the same...!!!


- Packets have **different criticality**
  - Packet is critical if latency of a packet affects application's performance
  - Different criticality due to memory level parallelism (MLP)

# MLP Principle

---



Packet Latency  $\neq$  Network Stall Time

Different Packets have different criticality due to MLP

Criticality(green) > Criticality(blue) > Criticality(red)

# Outline

---

- Introduction
  - Packet Scheduling
  - Memory Level Parallelism
- Aérgia
  - Concept of Slack
  - Estimating Slack
- Evaluation
- Conclusion

# What is Aérgia?

---



- Aérgia is the spirit of laziness in Greek mythology
- Some packets can afford to **slack!**

# Outline

---

- Introduction
  - Packet Scheduling
  - Memory Level Parallelism
- Aérgia
  - Concept of Slack
  - Estimating Slack
- Evaluation
- Conclusion

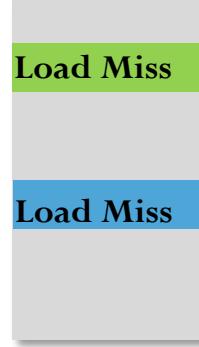
# Slack of Packets

---

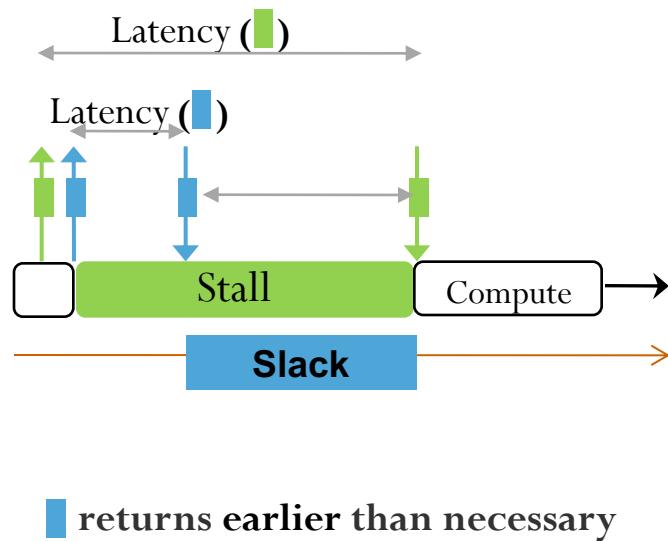
- What is slack of a packet?
  - Slack of a packet is number of cycles it can be delayed in a router without (significantly) reducing application's performance
  - Local network slack
- Source of slack: Memory-Level Parallelism (MLP)
  - Latency of an application's packet hidden from application due to overlap with latency of pending cache miss requests
- Prioritize packets with lower slack

# Concept of Slack

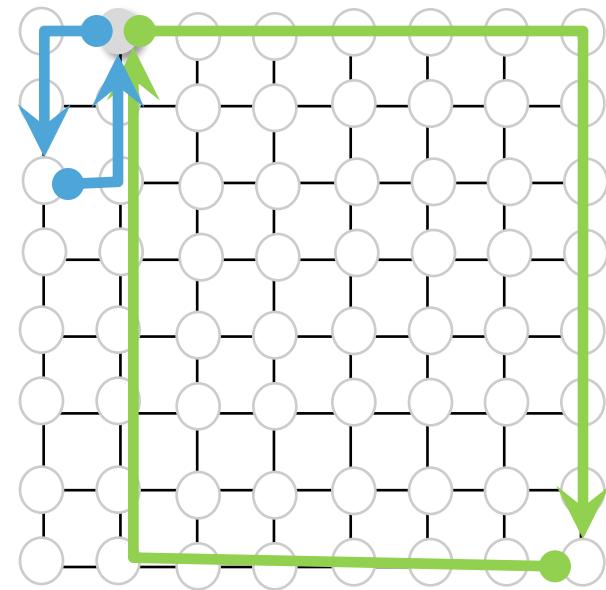
## Instruction Window



## Execution Time



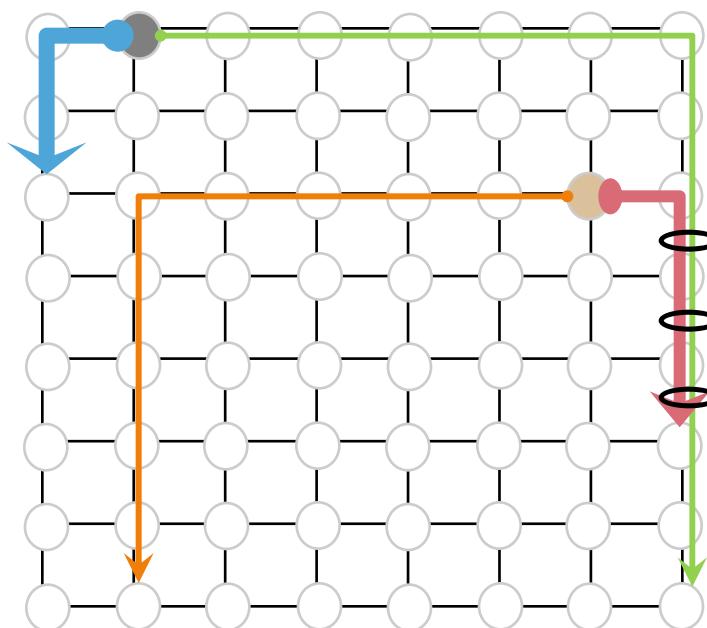
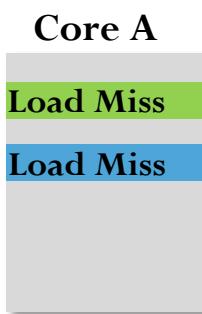
## Network-on-Chip



$$\text{Slack } (\textcolor{blue}{\square}) = \text{Latency } (\textcolor{green}{\square}) - \text{Latency } (\textcolor{blue}{\square}) = 26 - 6 = 20 \text{ hops}$$

Packet(\textcolor{blue}{\square}) can be delayed for available slack cycles without reducing performance!

# Prioritizing using Slack



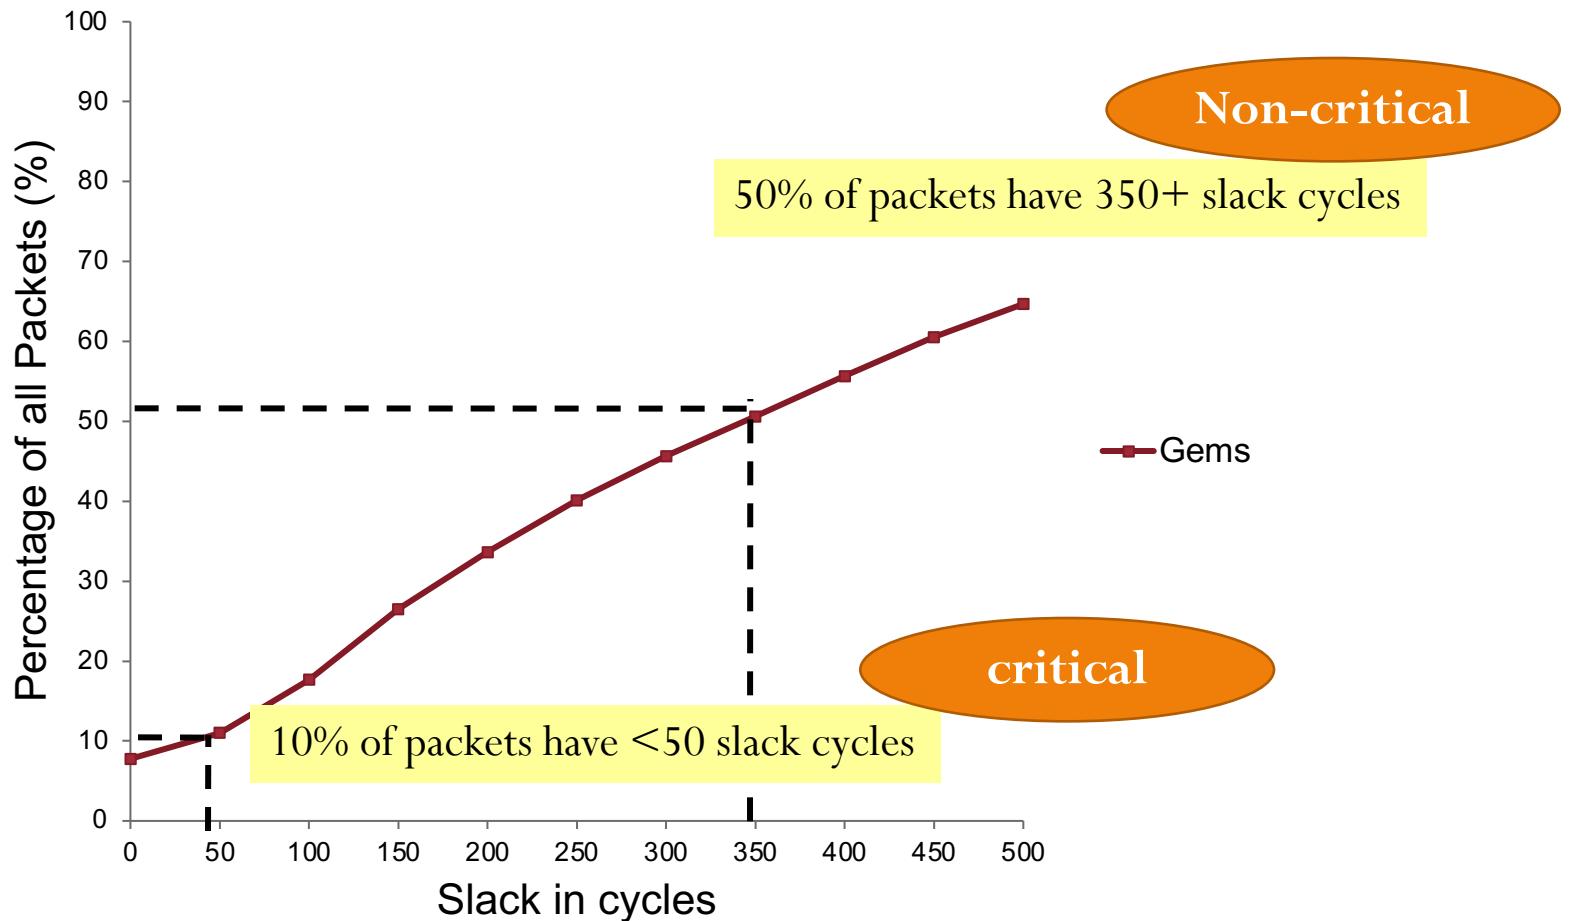
Packet	Latency	Slack
█	13 hops	0 hops
█	3 hops	10 hops

○ Interference at 3 hops

$\text{Slack}(\text{█}) > \text{Slack}(\text{█})$

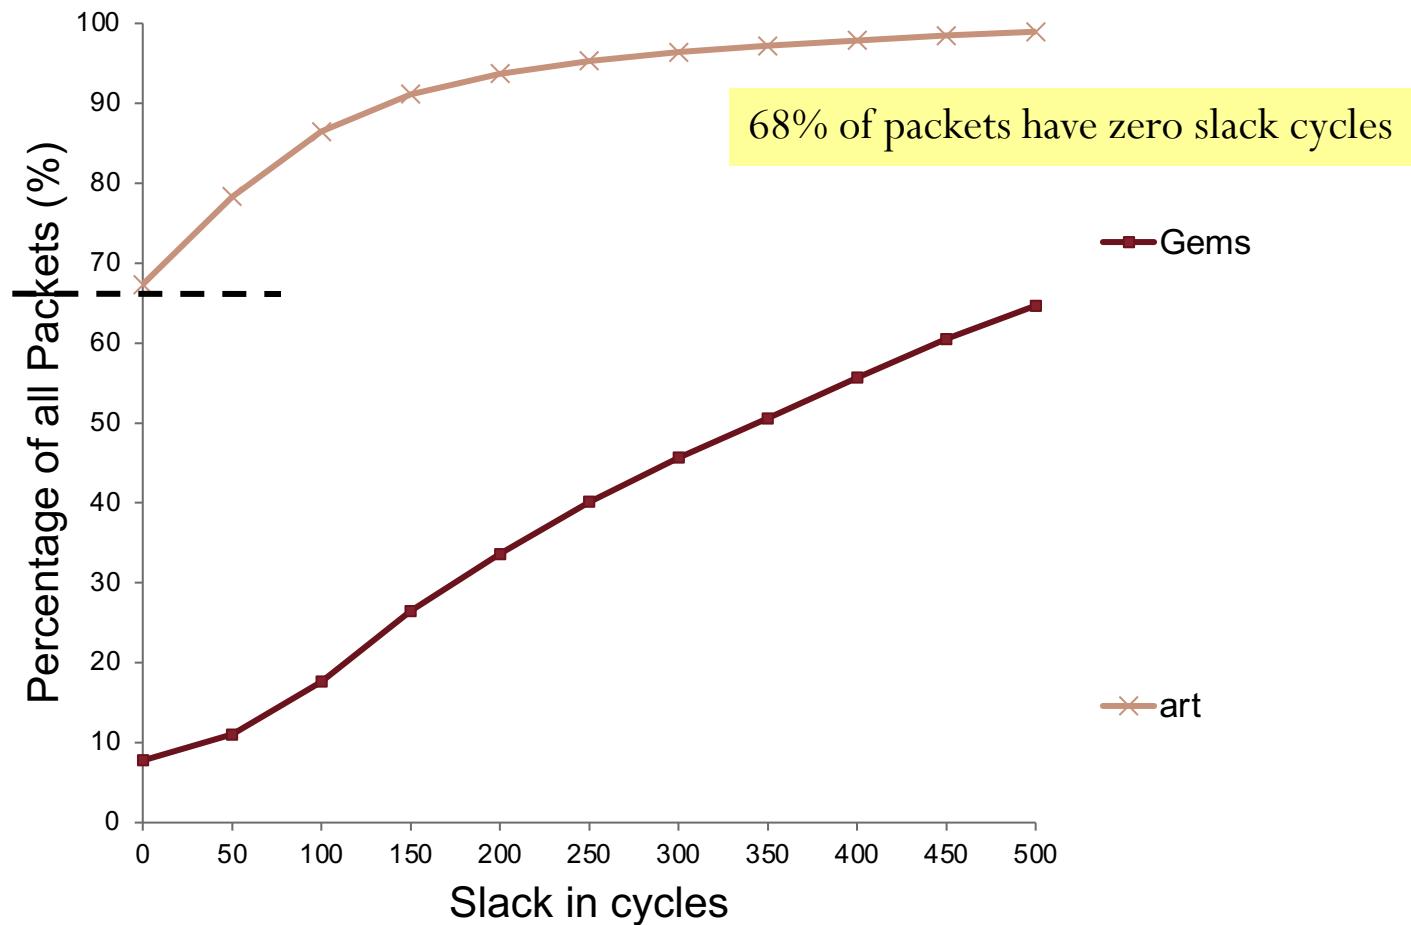
Prioritize █

# Slack in Applications



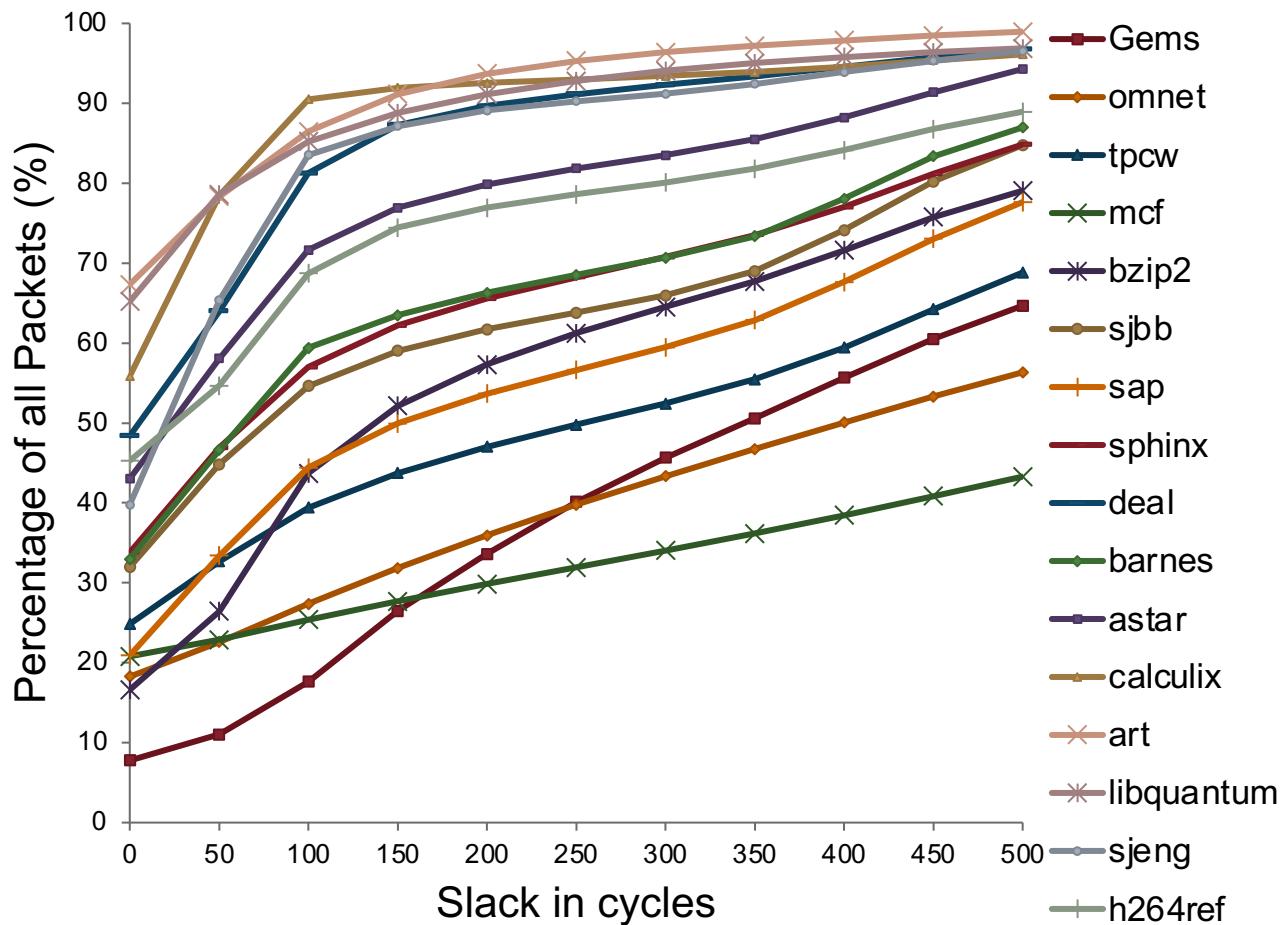
# Slack in Applications

---



# Diversity in Slack

---



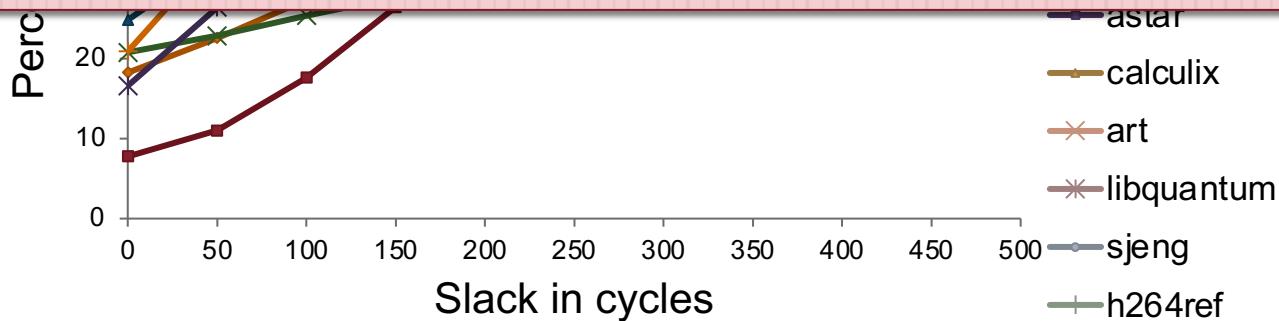
# Diversity in Slack



Slack varies **between** packets of **different** applications



Slack varies **between** packets of **a single** application



# Outline

---

- Introduction
  - Packet Scheduling
  - Memory Level Parallelism
- Aérgia
  - Concept of Slack
  - Estimating Slack
- Evaluation
- Conclusion

# Estimating Slack Priority

---

$$\text{Slack}(P) = \text{Max}(\text{Latencies of } P\text{'s Predecessors}) - \text{Latency of } P$$

Predecessors( $P$ ) are the packets of outstanding cache miss requests when  $P$  is issued

- Packet latencies not known when issued
  
- Predicting latency of any packet  $Q$ 
  - Higher latency if  $Q$  corresponds to an L2 miss
  - Higher latency if  $Q$  has to travel farther number of hops

# Estimating Slack Priority

---

- Slack of P = Maximum Predecessor Latency – Latency of P
- $\text{Slack}(P) = \begin{array}{|c|c|c|} \hline \text{PredL2} & \text{MyL2} & \text{HopEstimate} \\ (2 \text{ bits}) & (1 \text{ bit}) & (2 \text{ bits}) \\ \hline \end{array}$

**PredL2:** Set if any predecessor packet is servicing L2 miss

**MyL2:** Set if P is NOT servicing an L2 miss

**HopEstimate:** Max (# of hops of Predecessors) – hops of P

# Estimating Slack Priority

---

- How to predict L2 hit or miss at core?
  - *Global Branch Predictor* based L2 Miss Predictor
    - Use Pattern History Table and 2-bit saturating counters
  - *Threshold* based L2 Miss Predictor
    - If #L2 misses in “M” misses  $\geq$  “T” threshold then next load is a L2 miss.
- Number of miss predecessors?
  - List of outstanding L2 Misses
- Hops estimate?
  - Hops  $\Rightarrow \Delta X + \Delta Y$  distance
  - Use predecessor list to calculate slack hop estimate

# Starvation Avoidance

---

- Problem: Starvation
  - Prioritizing packets can lead to starvation of lower priority packets
- Solution: Time-Based Packet Batching
  - New batches are formed at every T cycles
  - Packets of older batches are prioritized over younger batches

# Putting it all together

---

- Tag header of the packet with priority bits before injection

**Priority (P) =**



- Priority(P)?
  - P's batch *(highest priority)*
  - P's Slack
  - Local Round-Robin *(final tie breaker)*

# Outline

---

- Introduction
  - Packet Scheduling
  - Memory Level Parallelism
- Aérgia
  - Concept of Slack
  - Estimating Slack
- Evaluation
- Conclusion

# Evaluation Methodology

---

- 64-core system
  - x86 processor model based on Intel Pentium M
  - 2 GHz processor, 128-entry instruction window
  - 32KB private L1 and 1MB per core shared L2 caches, 32 miss buffers
  - 4GB DRAM, 320 cycle access latency, 4 on-chip DRAM controllers
- Detailed Network-on-Chip model
  - 2-stage routers (with speculation and look ahead routing)
  - Wormhole switching (8 flit data packets)
  - Virtual channel flow control (6 VCs, 5 flit buffer depth)
  - 8x8 Mesh (128 bit bi-directional channels)
- Benchmarks
  - Multiprogrammed scientific, server, desktop workloads (35 applications)
  - 96 workload combinations

# Qualitative Comparison

---

- **Round Robin & Age**

- Local and application oblivious
  - Age is biased towards heavy applications

- **Globally Synchronized Frames (GSF)**

[Lee et al., ISCA 2008]

- Provides **bandwidth fairness** at the expense of **system performance**
  - Penalizes heavy and bursty applications

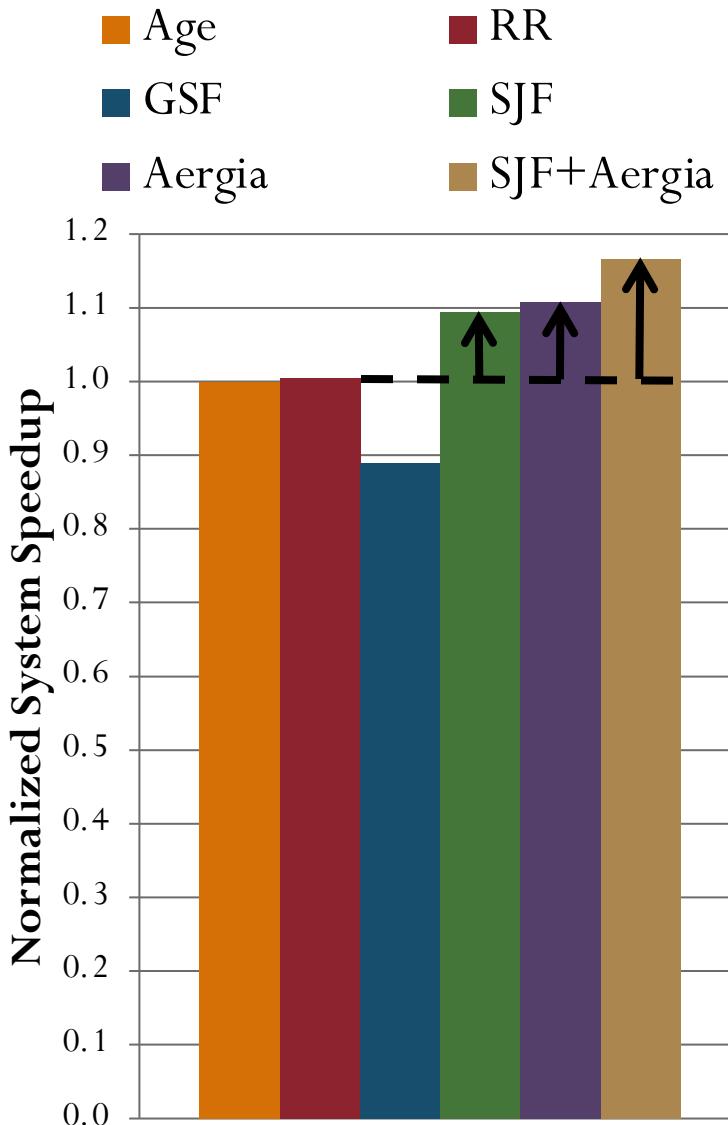
- **Application-Aware Prioritization Policies (SJF)**

[Das et al., MICRO 2009]

- **Shortest-Job-First Principle**
  - Packet scheduling policies which prioritize network sensitive applications which inject lower load

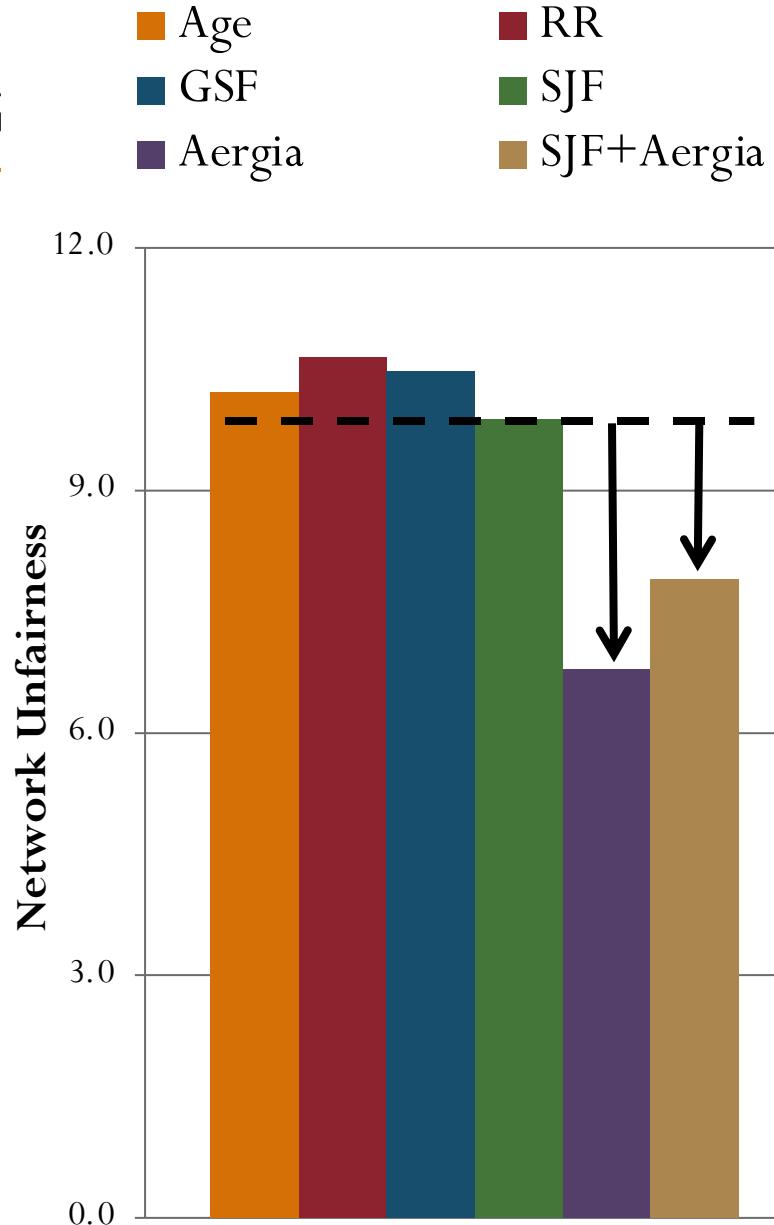
# System Performance

- SJF provides 8.9% improvement in weighted speedup
- Aérgia improves system throughput by 10.3%
- Aérgia+SJF improves system throughput by 16.1%



# Network Unfairness

- SJF does not imbalance network fairness
- Aergia improves network unfairness by 1.5X
- SJF+Aergia improves network unfairness by 1.3X



# Conclusions & Future Directions

---

- Packets have different criticality, yet existing packet scheduling policies **treat all packets equally**
- We propose a new approach to packet scheduling in NoCs
  - We define **Slack** as a key measure that characterizes the relative importance of a packet.
  - We propose **Aérgia** a novel architecture to accelerate low slack critical packets
- Result
  - Improves system performance: 16.1%
  - Improves network fairness: 30.8%

# Slack-Based Packet Scheduling

---

- Reetuparna Das, Onur Mutlu, Thomas Moscibroda, and Chita R. Das,  
**"Aergia: Exploiting Packet Latency Slack in On-Chip Networks"**  
*Proceedings of the 37th International Symposium on Computer Architecture (ISCA)*, pages 106-116, Saint-Malo, France, June 2010. [Slides \(pptx\)](#)  
**One of the 11 computer architecture papers of 2010 selected as Top Picks by IEEE Micro.**

## Aérgia: Exploiting Packet Latency Slack in On-Chip Networks

Reetuparna Das<sup>§</sup>   Onur Mutlu<sup>†</sup>   Thomas Moscibroda<sup>‡</sup>   Chita R. Das<sup>§</sup>

§Pennsylvania State University  
[{rdas,das}@cse.psu.edu](mailto:{rdas,das}@cse.psu.edu)

†Carnegie Mellon University  
[onur@cmu.edu](mailto:onur@cmu.edu)

‡Microsoft Research  
[moscitho@microsoft.com](mailto:moscitho@microsoft.com)